



UBISOFT

OPTIMIZING BUILDS ON WINDOWS

SOME PRACTICAL CONSIDERATIONS

Alexandre Ganea, Ubisoft
alexandre.ganea@ubisoft.com

2019 Bay Area LLVM Developers' Meeting, Oct.22-23

SUMMARY

PART 1

PREAMBLE

PART 2

EXPERIMENTS

PART 3

PROPOSAL

PART 4

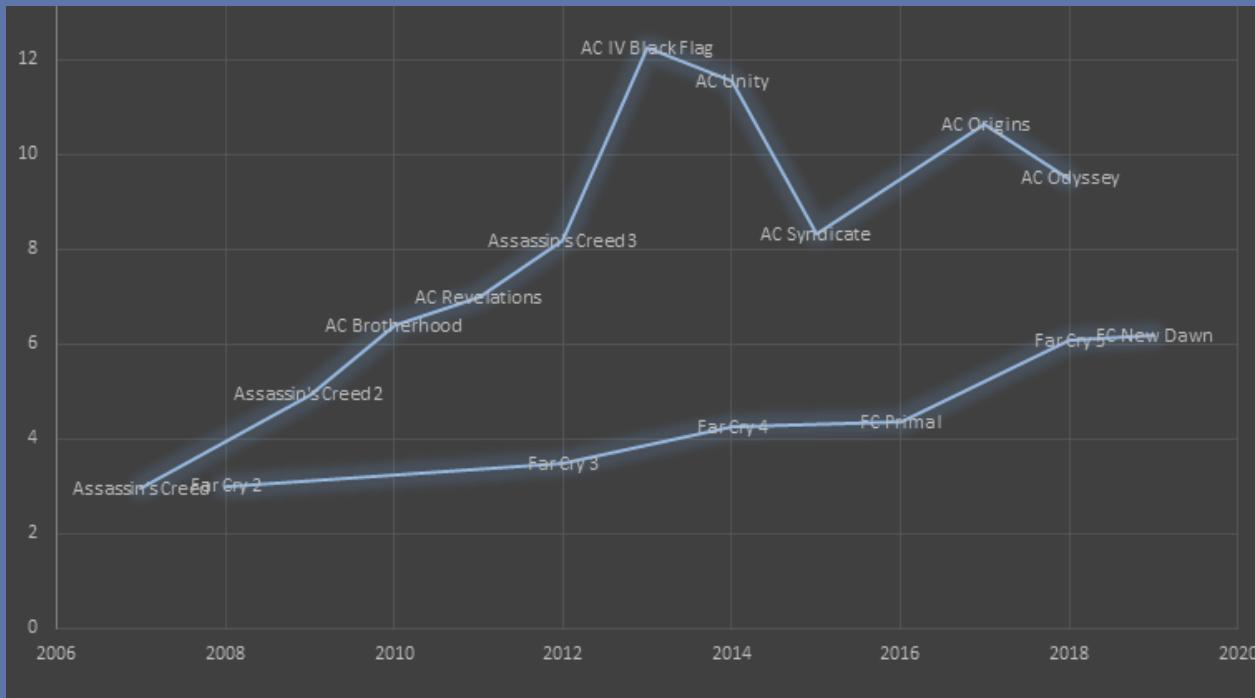
NEXT STEPS



A collage of six movie posters from the 2010s, each featuring a different protagonist or character. From left to right: 1. A woman in detailed Roman-style armor. 2. A pig wearing a suit and tie, holding a sword. 3. A person in a futuristic, glowing suit. 4. A woman in a blue hoodie. 5. A person in a dark, metallic gas mask. 6. A person with a heavily damaged, bandaged face.

PART 1
PREAMBLE:
CHALLENGES

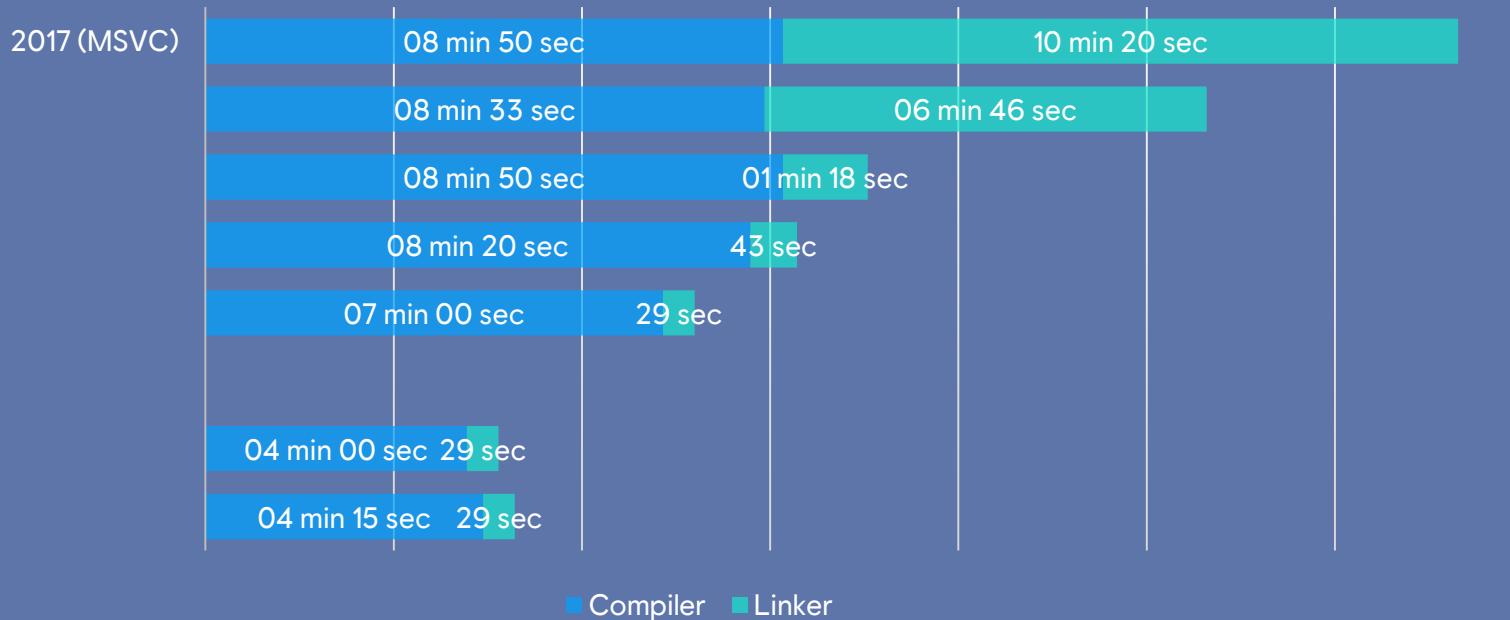
Lines of Code (Assassins' Creed, Far Cry)



Game production constraints @ Ubisoft

Concurrent AAA games LoC/game	20 – 25 30 - 50 M	Editor Build	20,000 .CPP 25,000 .H
Programmers/title Code Changes/day	100 – 250 100 – 150 (peak:400)		23 GB .OBJ 9 GB .DEBUG\$T
Build targets/platform Platforms/Game	5 – 6 4+		10 M TYPE RECORDS 42 M SYMBOLS
Code workspace Data workspace	70 - 100 GB 100 - 200 GB		300 M .EXE 2 GB .PDB
Game builds/day Stripped Build Final Build	100 – 150 1 - 6 GB 50 - 90 GB	Windows 10 Fastbuild, distributed Always Unity builds	

AAA GAME, CLEAN REBUILD X64 EDITOR RELEASE (FASTBUILD)



A photograph of a group of people in a mountainous landscape. In the foreground, a person wearing a helmet and a large backpack is seen from behind, looking out over a valley. Another person stands further back on the left. In the distance, a town is nestled in a valley, and a winding road leads through the mountains. The sky is blue with some clouds.

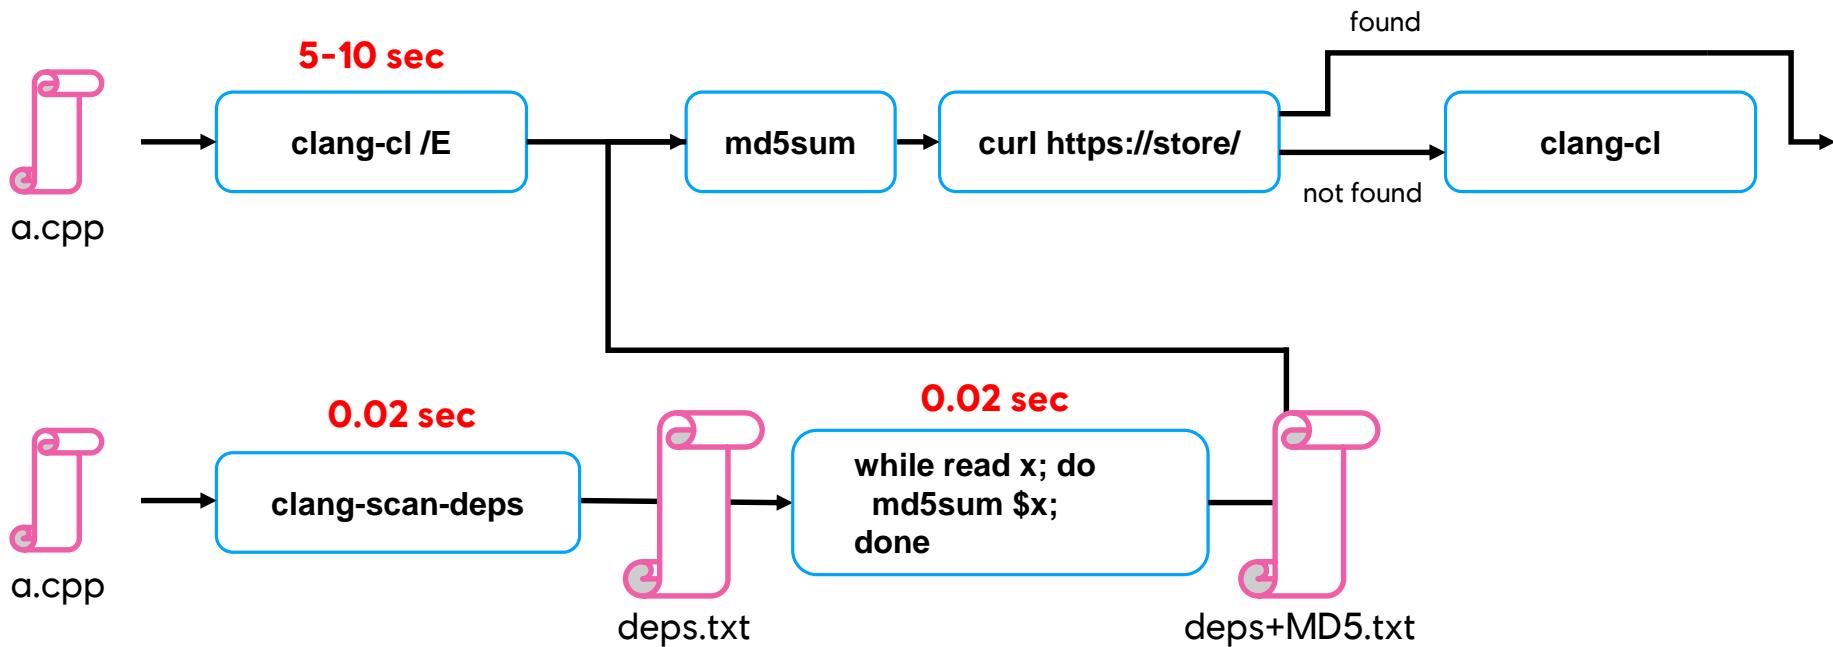
PART 2

EXPERIMENTS

2.1

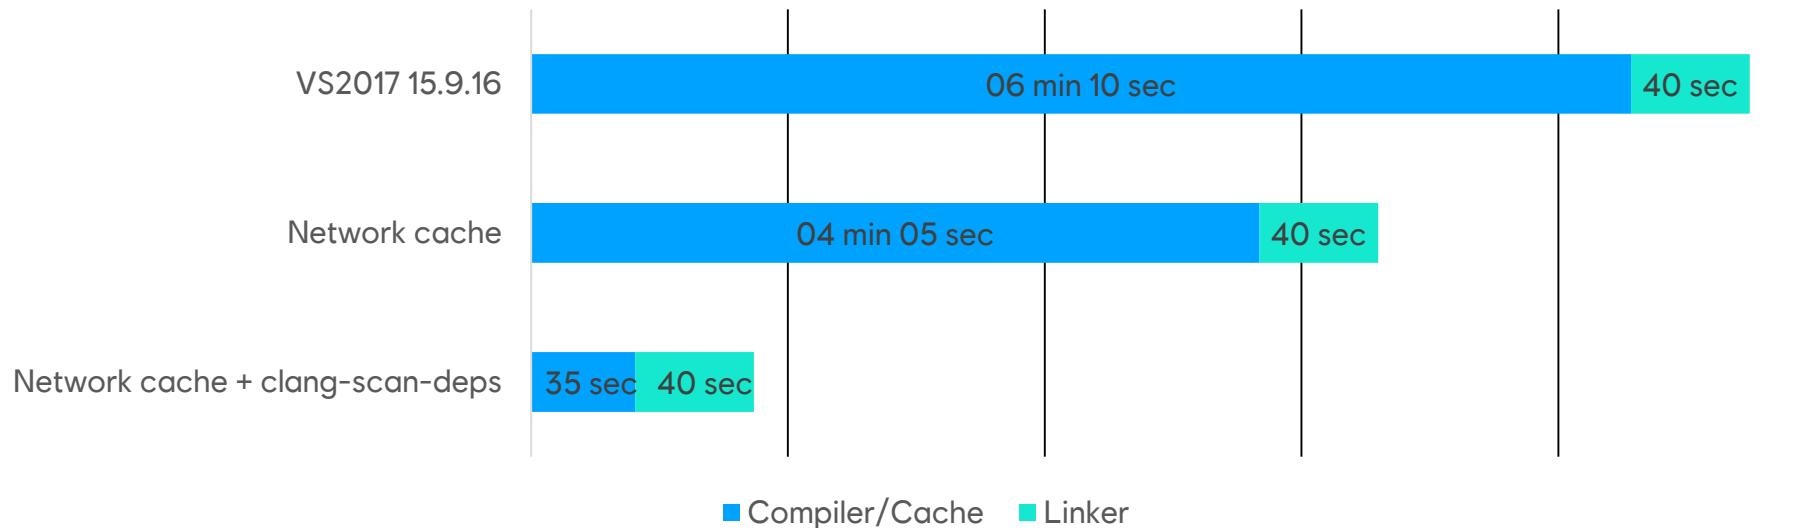
Clang-scan-deps & Fastbuild cache

FASTBUILD CACHE READ ALGORITHM



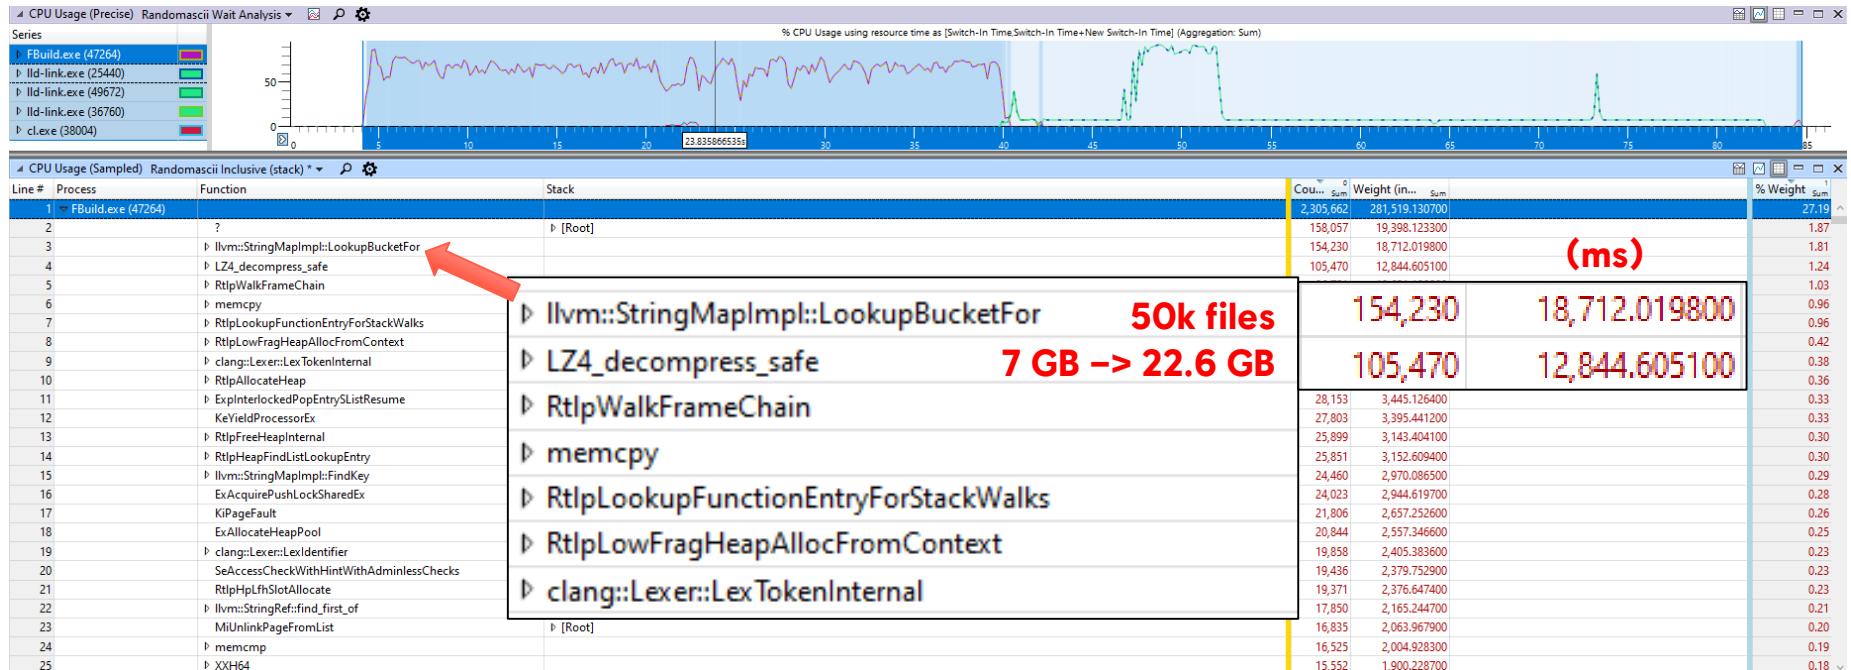
100% NETWORK CACHE HITS

AAA GAME, X64 EDITOR RELEASE (FASTBUILD)



clang-scan-deps + network cache

LLD (MSVC OBJS + ghash)

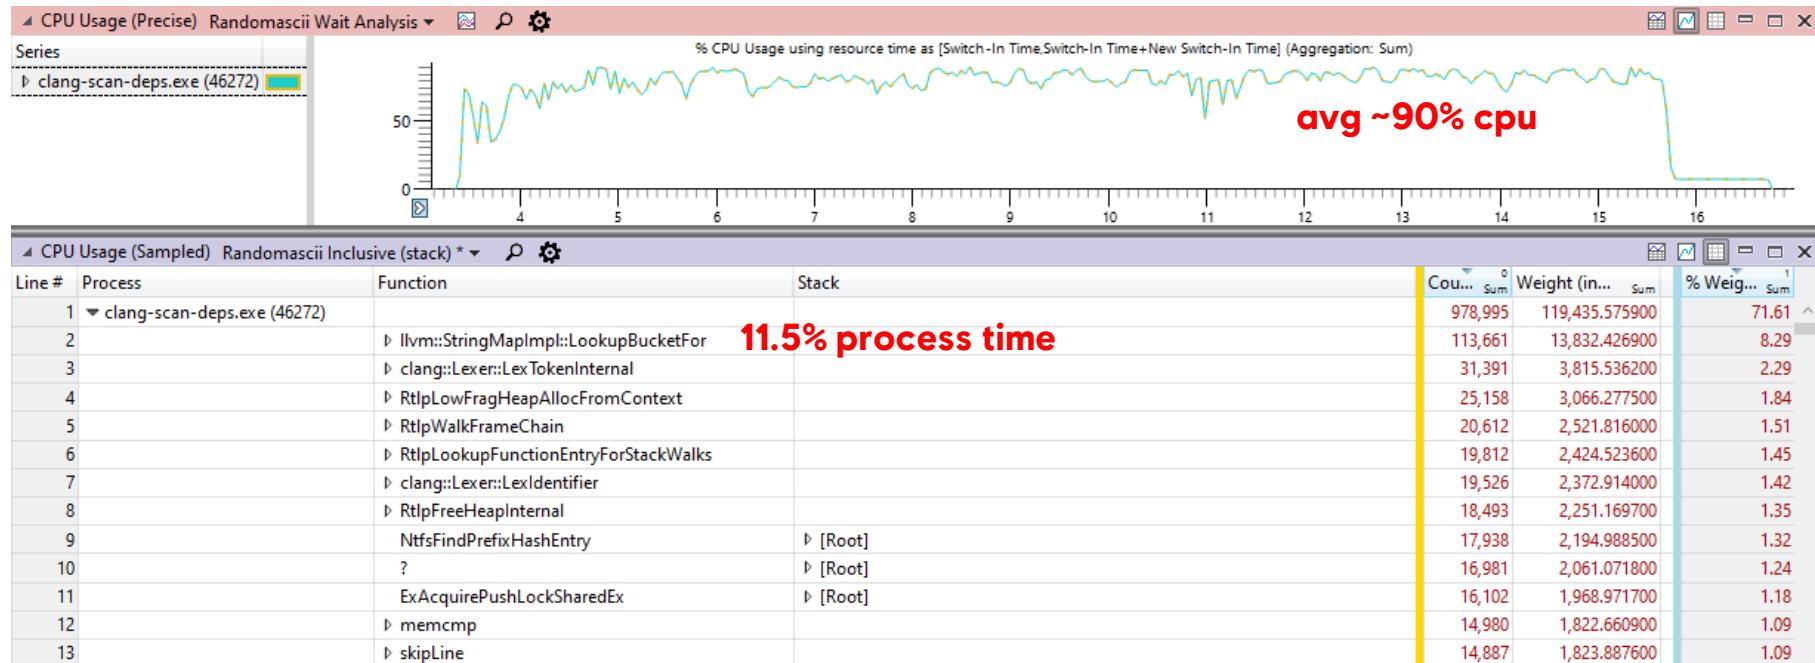


Intel Xeon W-2135 @ 3.7 GHz, 128 GB, NVMe SSD, 1Gbps Network

2.2

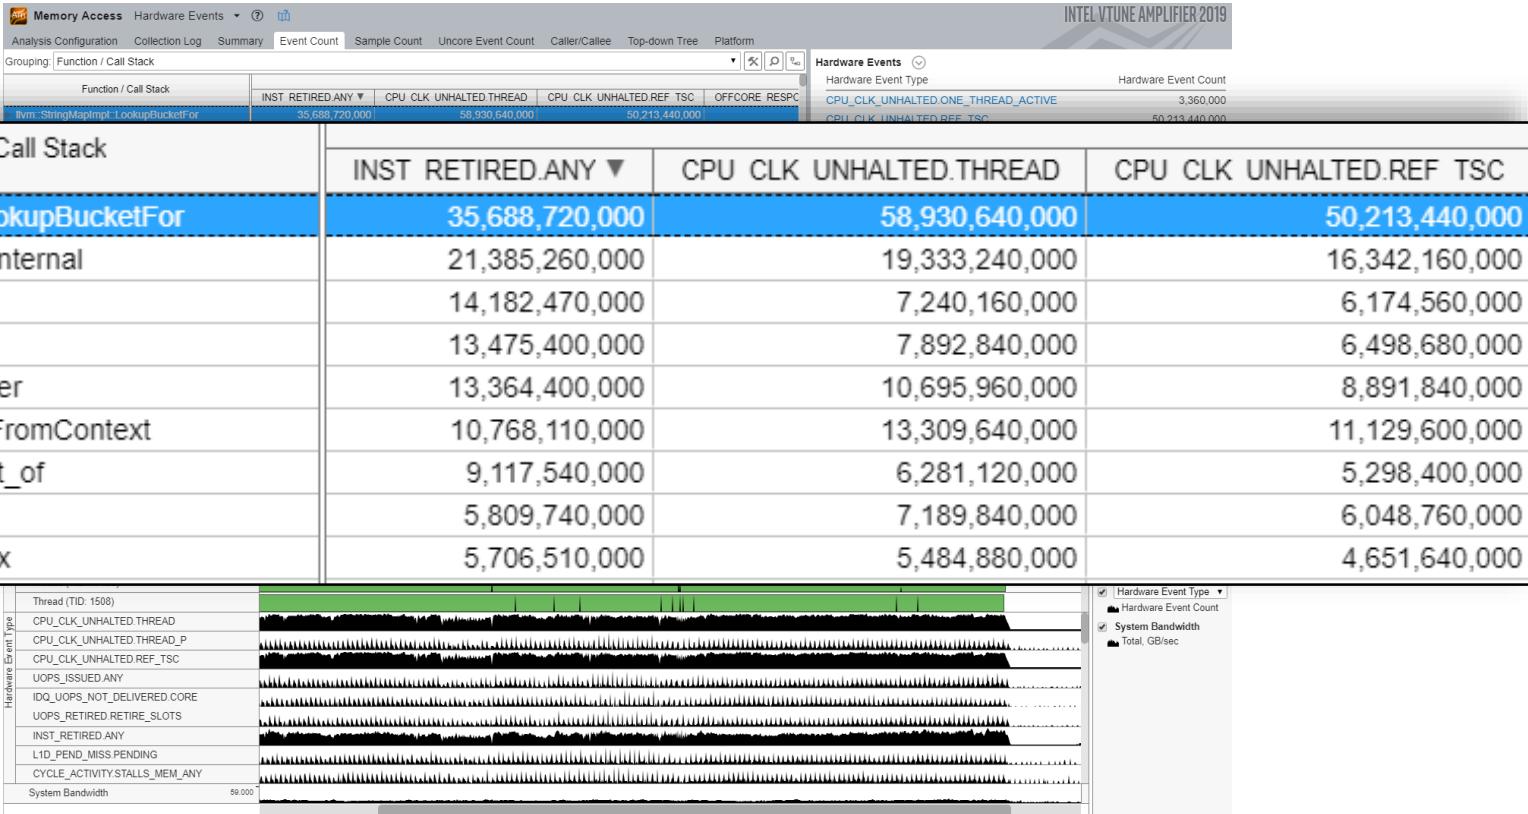
StringMap

CLANG-SCAN-DEPS STANDALONE (50K FILES)

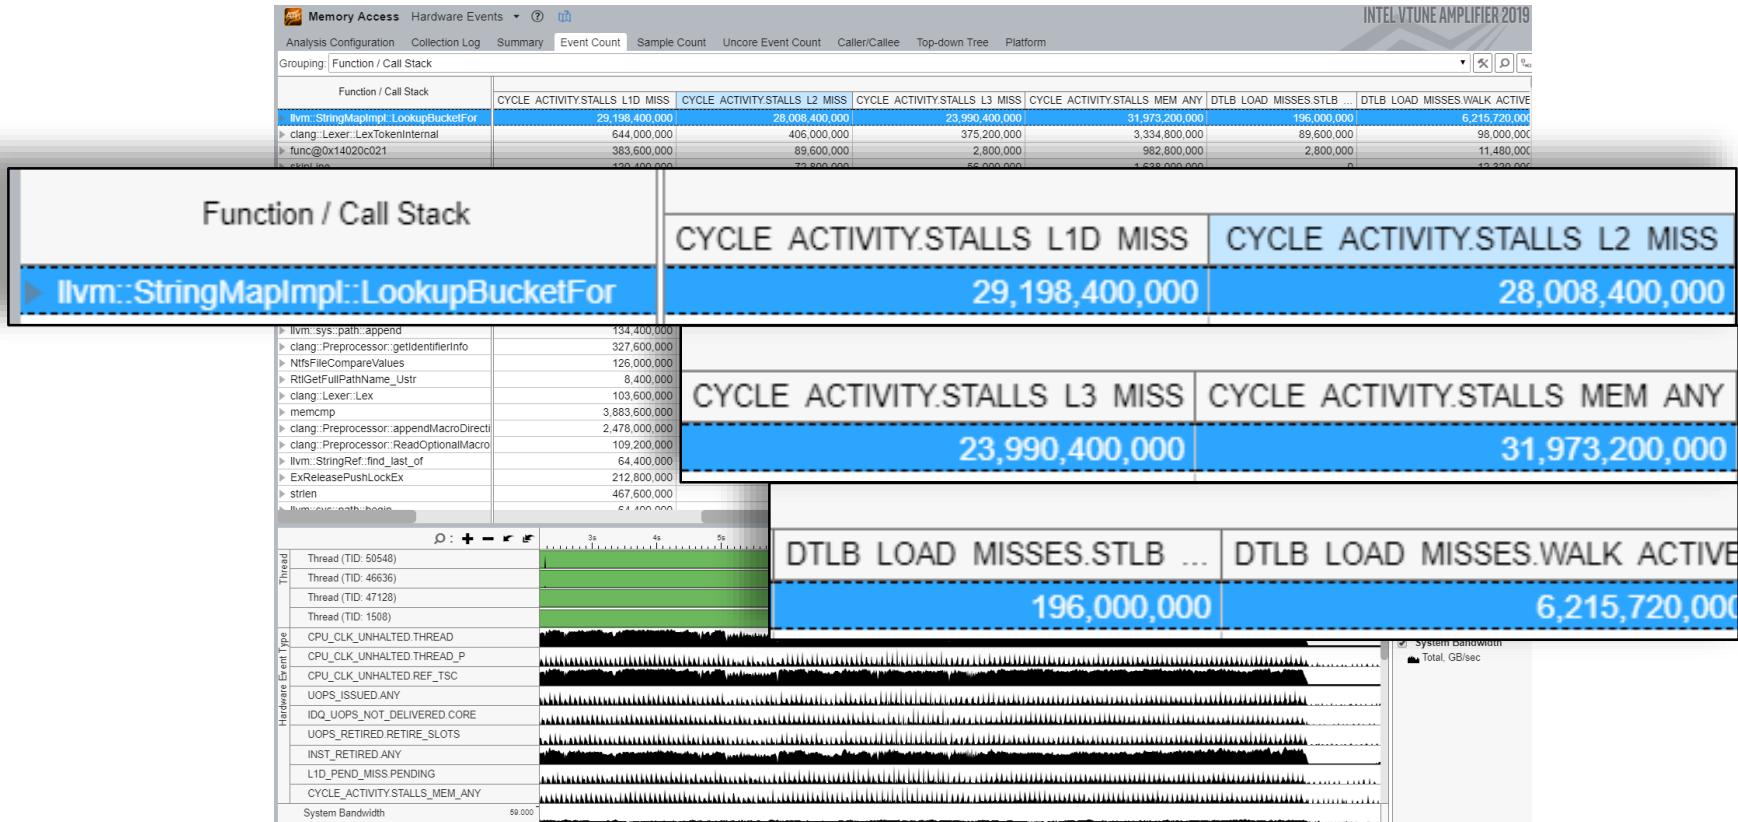


Intel Xeon W-2135 @ 3.7 GHz (6-core), 128 GB, NVMe SSD

STRINGMAP



STRINGMAP



DOWN THE RABBIT HOLE

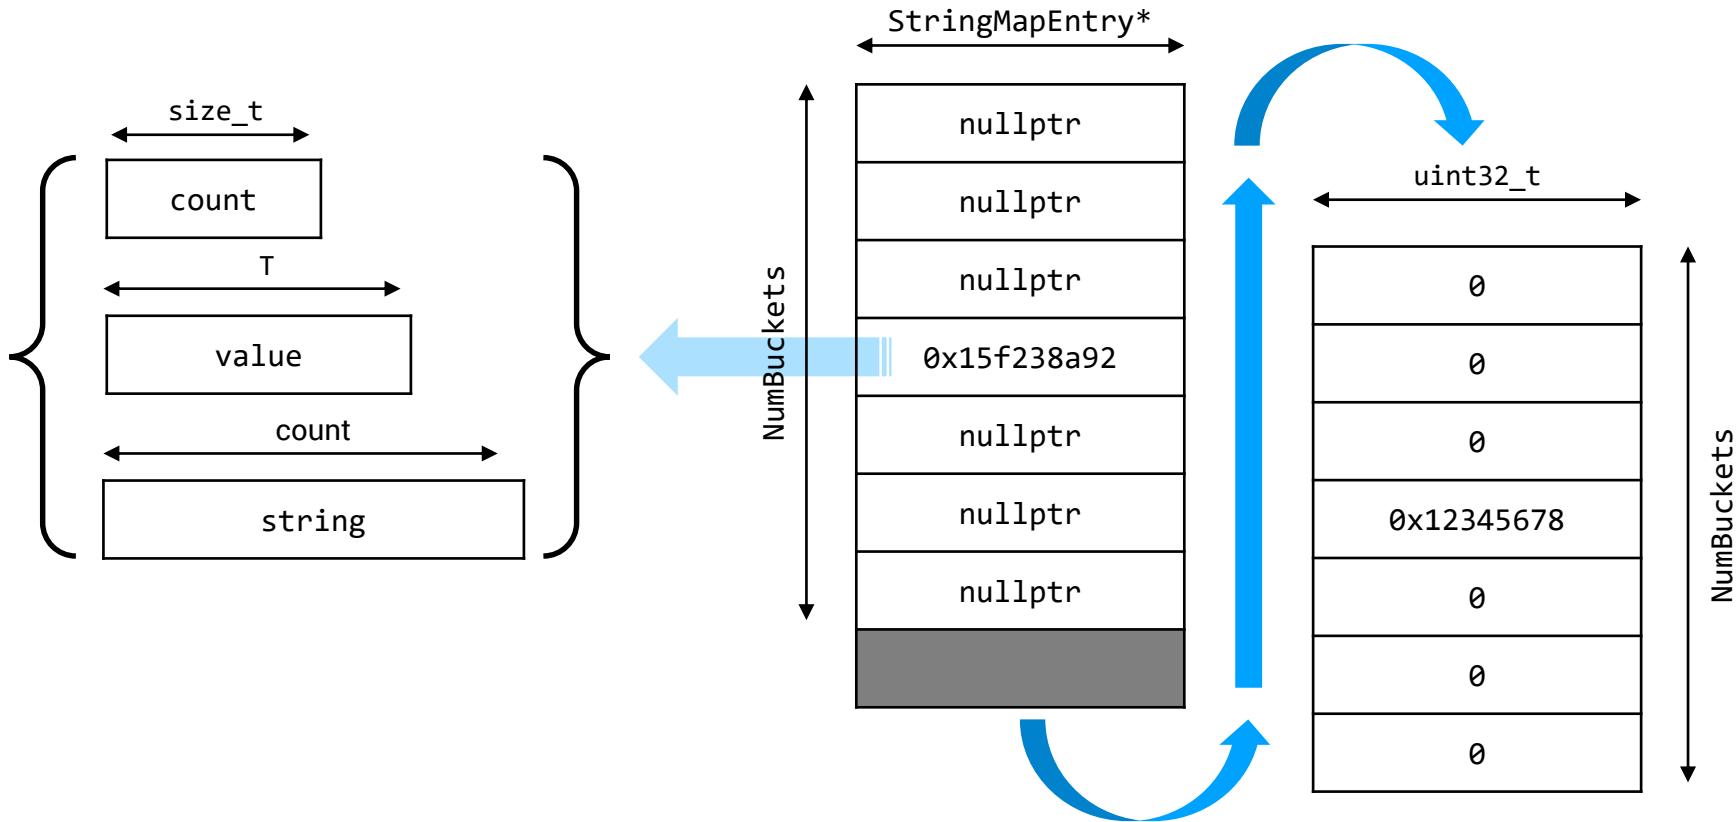
`sizeof(std::error_code) -> 16 bytes`

`sizeof(llvm::ErrorOr<DirectoryEntry&>) -> 24 bytes`

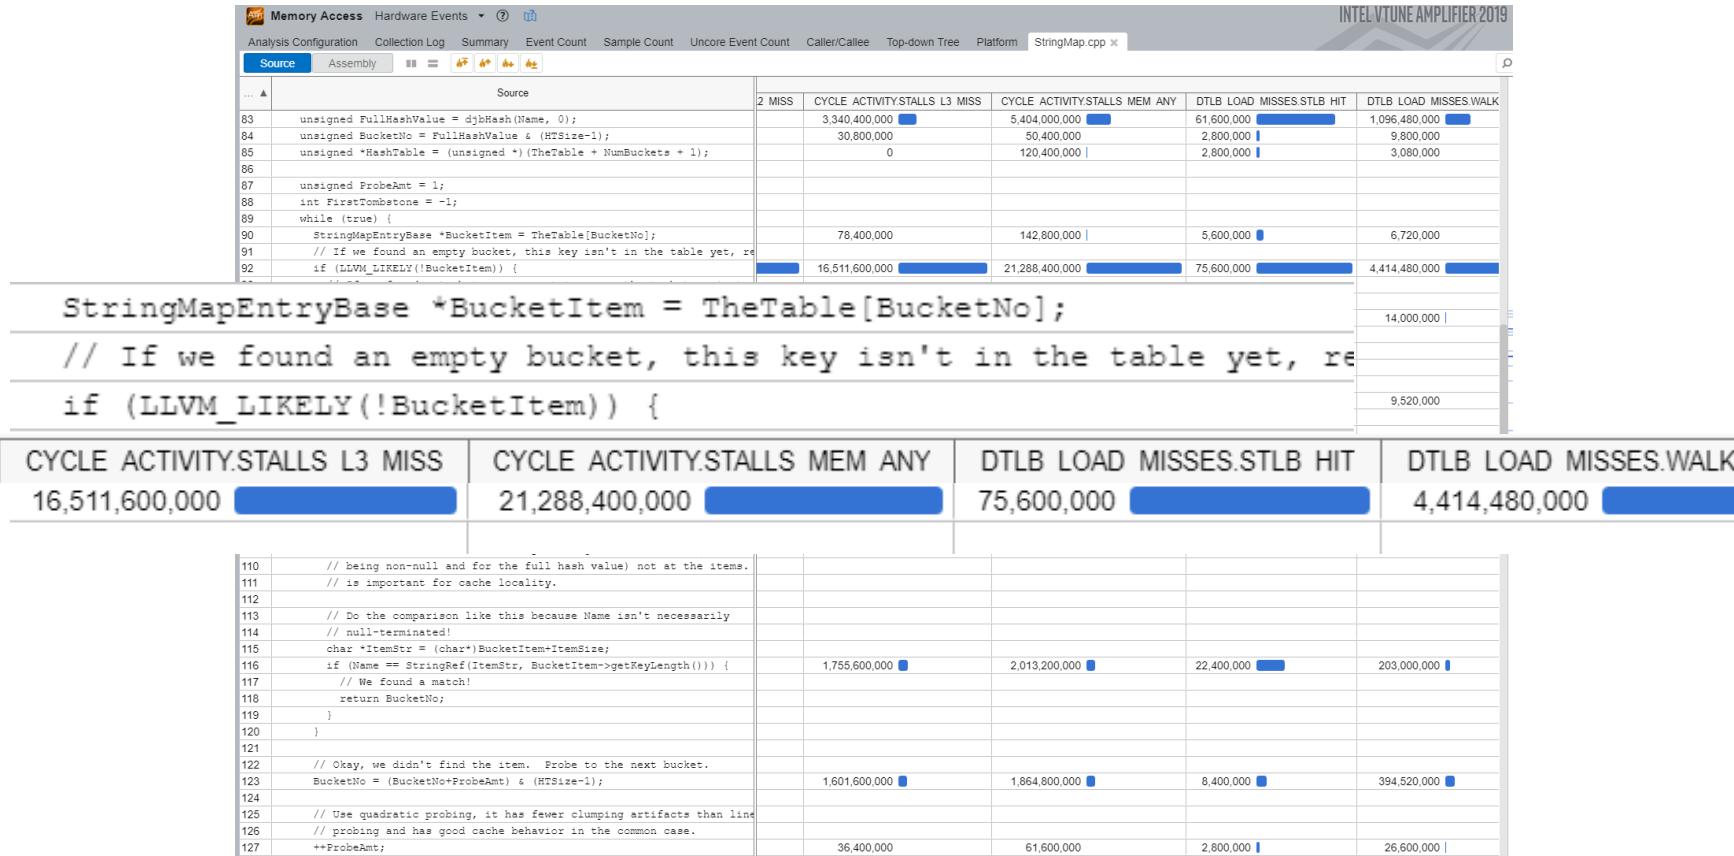
`sizeof(llvm::StringMapEntry<llvm::ErrorOr<DirectoryEntry&>>) -> 32 bytes (+string contents)`

llvm::StringMapEntry<llvm::ErrorOr<clang::DirectoryEntry > >						
	Members	Functions				
	Field	Type	Offset	Bit offset	Size	Padding
-	Base: llvm::StringMapEntryBase	llvm::StringMapEntryBase	0		8 bytes	0
	└__ StrLen	uint64	0		8 bytes	
-	second	llvm::ErrorOr<clang::DirectoryEntry &>	8		24 bytes	7
+	└__ ErrorStorage	llvm::AlignedCharArrayUnion<std::e...	8		16 bytes	0
+	└__ TStorage	llvm::AlignedCharArrayUnion<std::r...	8		8 bytes	0
	└__ HasError	bool	24	0	1 bit	
		Padding	17	7	7	

STRINGMAP: MEMORY LAYOUT

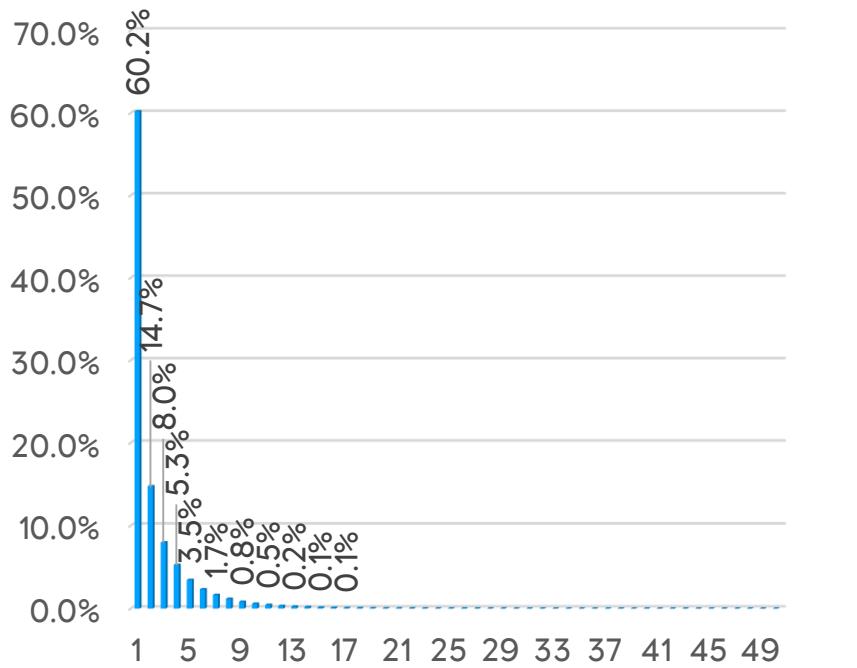


STRINGMAP (VTUNE)



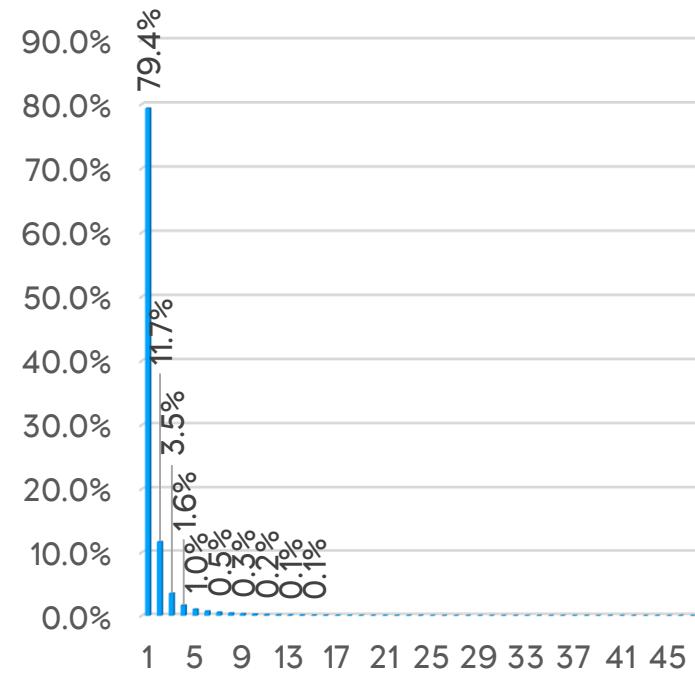
STRINGMAP STATS

Hash collisions / call



187 M samples

Cachelines hit / call



DenseMap<uint64_t, T> + xxHash64()
+ StringSaver

```
DenseMap<__int128,T> + XXH128()
+ StringSaver
```

2.4

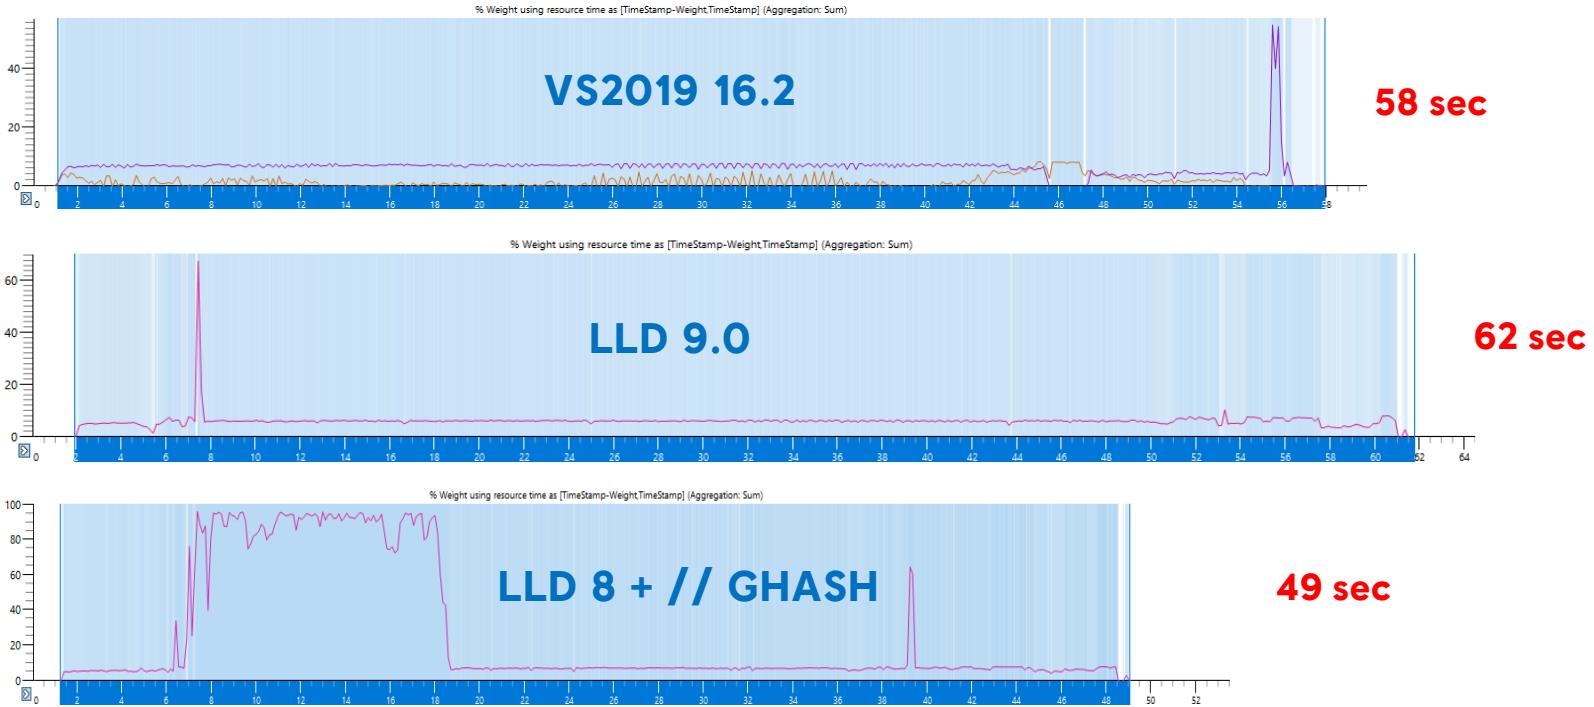
Multithreading LLD

(COFF driver)

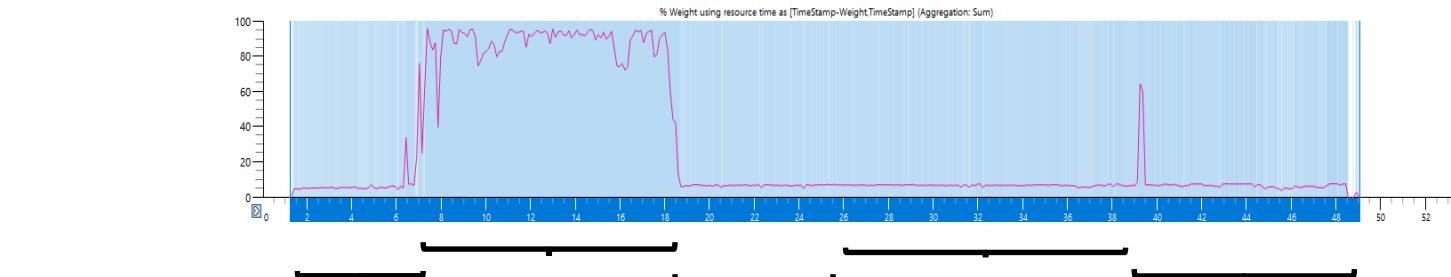


LINK AAA GAME, X64 EDITOR RELEASE

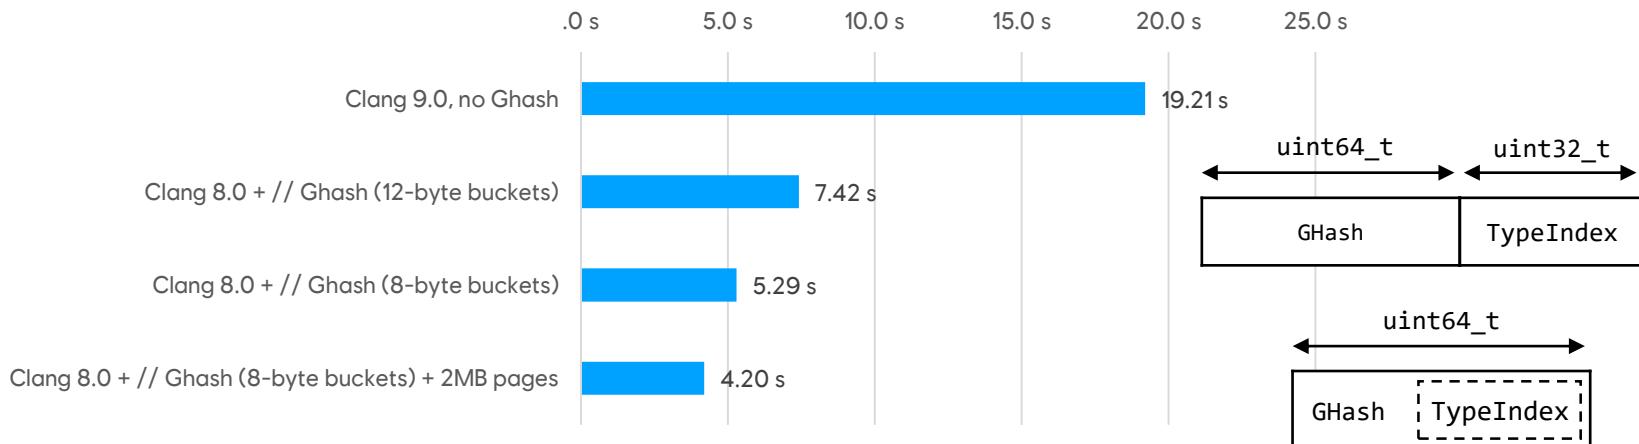
(22.8GB MSVC OBJS)



Intel Xeon W-2135 @ 3.7 GHz (6-core), 128 GB, NVMe SSD



Read sections & layout EXE
 Debug Info Type Global Hash
 Type Merge
 Symbols Merge
 Layout PDB



2.3

Process Creation



COMPILING WITH CLANG 9.0

	ninja.exe	0.29	30,332 K	31,500 K	4,276,340 K 23620	
	clang-cl.exe		2,180 K	10,248 K	4,355,108 K 28944	
	clang-cl.exe	8.14	275,144 K	300,784 K	4,640,804 K 31480	
	clang-cl.exe		2,192 K	10,252 K	4,355,108 K 7564	
	clang-cl.exe	8.25	256,096 K	281,312 K	4,622,340 K 26956	
	clang-cl.exe		2,208 K	10,272 K	4,355,108 K 8064	
	clang-cl.exe	7.46	167,672 K	195,076 K	4,562,456 K 35484	
	clang-cl.exe		2,200 K	10,264 K	4,355,108 K 14292	
	clang-cl.exe	8.10	143,192 K	170,672 K	4,513,596 K 23720	
	clang-cl.exe		2,160 K	10,220 K	4,355,108 K 14232	
	clang-cl.exe	8.70	167,492 K	186,804 K	4,532,180 K 7480	
	clang-cl.exe		2,164 K	10,224 K	4,355,108 K 30776	
	clang-cl.exe	6.45	144,060 K	161,972 K	4,516,752 K 20112	
	clang-cl.exe		2,160 K	10,224 K	4,355,108 K 29408	
	clang-cl.exe	7.21	102,432 K	124,076 K	4,464,760 K 15044	
	clang-cl.exe		2,260 K	10,296 K	4,355,108 K 16404	
	clang-cl.exe	8.40	101,684 K	122,984 K	4,463,836 K 19840	
	clang-cl.exe		2,248 K	10,288 K	4,364,900 K 31124	
	clang-cl.exe	7.14	91,104 K	108,872 K	4,465,128 K 17324	
	clang-cl.exe		2,156 K	10,220 K	4,355,108 K 25252	
	clang-cl.exe	6.09	86,904 K	103,776 K	4,451,788 K 15116	
	clang-cl.exe		2,156 K	10,220 K	4,355,108 K 26800	
	clang-cl.exe	6.30	69,224 K	86,408 K	4,430,528 K 34252	
	clang-cl.exe		2,148 K	10,212 K	4,355,108 K 34484	
	clang-cl.exe	5.29	66,224 K	83,432 K	4,430,780 K 24736	
	clang-cl.exe		2,156 K	10,220 K	4,355,108 K 30684	
	clang-cl.exe	5.91	58,804 K	75,688 K	4,430,440 K 26188	
	clang-cl.exe		2,160 K	10,228 K	4,355,108 K 27188	
	clang-cl.exe	6.64	45,792 K	63,144 K	4,413,180 K 29616	
	clang-cl.exe		2,216 K	10,272 K	4,355,108 K 32748	
	clang-cl.exe	0.27	17,004 K	32,716 K	4,377,492 K 32024	
	clang-cl.exe		2,160 K	10,224 K	4,355,108 K 25944	
	clang-cl.exe	1.58	12,936 K	27,728 K	4,386,724 K 34904	
	clang-cl.exe		0.29	2,172 K	10,232 K	4,355,108 K 15304
	clang-cl.exe	< 0.01	936 K	108 K	88,984 K 19824	

CLANG CC1 IN PROCMON

Time of Day	Process Name	PID	Operation	F	Result	Detail
3:57:07.2331413 PM	clang-cl.exe	50920	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.6718750 seconds, Kernel Time: 0.1250000 seconds, Private Bytes: 35,127,296, Pe
3:57:07.3101077 PM	clang-cl.exe	48240	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.6406250 seconds, Kernel Time: 0.0781250 seconds, Private Bytes: 32,358,400, Pe
3:57:07.3173499 PM	clang-cl.exe	38632	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0156250 seconds, Kernel Time: 0.0156250 seconds, Private Bytes: 1,900,544, Pe
3:57:07.3385507 PM	clang-cl.exe	1052	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0312500 seconds, Private Bytes: 1,888,256, Pe
3:57:07.3467944 PM	clang-cl.exe	34696	Process Exit		SUCCESS	Exit Status: 0, User Time: 1.2343750 seconds, Kernel Time: 0.1406250 seconds, Private Bytes: 54,820,864, Pe
3:57:07.3617581 PM	clang-cl.exe	33684	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0468750 seconds, Private Bytes: 1,937,408, Pe
3:57:07.4323399 PM	clang-cl.exe	38904	Process Exit		SUCCESS	Exit Status: 0, User Time: 1.2343750 seconds, Kernel Time: 0.0781250 seconds, Private Bytes: 55,709,696, Pe
3:57:07.4873310 PM	clang-cl.exe	36756	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0468750 seconds, Private Bytes: 1,933,312, Pe
3:57:						
3:57:						Exit Status: 0, User Time: 0.0312500 seconds, Kernel Time: 0.0312500 seconds, Private Bytes: 1,884,160, Pe
3:57:						Exit Status: 0, User Time: 0.0156250 seconds, Kernel Time: 0.0156250 seconds, Private Bytes: 1,937,408, Pe
3:57:						Exit Status: 0, User Time: 2.2500000 seconds, Kernel Time: 0.1875000 seconds, Private Bytes: 63,782,912,
3:57:						Exit Status: 0, User Time: 0.0312500 seconds, Kernel Time: 0.0156250 seconds, Private Bytes: 1,937,408, Pe
3:57:						Exit Status: 0, User Time: 0.9375000 seconds, Kernel Time: 0.1875000 seconds, Private Bytes: 45,416,448,
3:57:						Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0937500 seconds, Private Bytes: 1,892,352, Pe
3:57:						Exit Status: 0, User Time: 2.281250 seconds, Kernel Time: 0.2916250 seconds, Private Bytes: 76,128,256,
3:57:						Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0625000 seconds, Private Bytes: 1,908,736, Pe
3:57:						
3:57:10.5163850 PM	clang-cl.exe	24864	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0312500 seconds, Kernel Time: 0.0156250 seconds, Private Bytes: 1,937,408, Pe
3:57:10.5314471 PM	clang-cl.exe	22880	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.9375000 seconds, Kernel Time: 0.1875000 seconds, Private Bytes: 45,416,448, Pe
3:57:10.5574145 PM	clang-cl.exe	36944	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0937500 seconds, Private Bytes: 1,892,352, Pe
3:57:10.6401201 PM	clang-cl.exe	13372	Process Exit		SUCCESS	Exit Status: 0, User Time: 2.2812500 seconds, Kernel Time: 0.2343750 seconds, Private Bytes: 76,128,256, Pe
3:57:10.7026652 PM	clang-cl.exe	51780	Process Exit		SUCCESS	Exit Status: 0, User Time: 0.0000000 seconds, Kernel Time: 0.0625000 seconds, Private Bytes: 1,908,736, Pe

93 ms

MAKING CC1 REENTRANT

clang/lib/driver/Job.cpp

clang/tools/driver/driver.cpp

```
int main(int argc_, const char **argv_) {
    noteBottomOfStack();
    llvm::InitLLVM X(argc_, argv_);
    SmallVector<const char *, 256> argv(argv_, argv_ + argc_);

    if (llvm::sys::Process::FixupStandardFileDescriptors())
        return 1;

    llvm::InitializeAllTLS();
    return ClangDriverMain(argv);
}

int ClangDriverMain(SmallVectorImpl<const char *>& argv) {
    static LLVM_THREAD_LOCAL bool EnterPE = true;
    if (EnterPE) {
        llvm::sys::DynamicLibrary::AddSymbol("ClangDriverMain", (void*)(i...
        EnterPE = false;
    } else {
        llvm::cl::ResetAllOptionOccurrences();
    }

    auto TargetAndMode = ToolChain::getTargetAndModeFromProgramName(argv);
}
```

```
int Command::Execute(ArrayRef<llvm::Optional<StringRef>> Redirects,
                     std::string *ErrMsg, bool *ExecutionFailed) const {
    [...]
    typedef int (*ClangDriverMainFunc)(SmallVectorImpl<const char *> &);

    ClangDriverMainFunc ClangDriverMain = nullptr;
    [...]
    if (ClangDriverMain) {
        [...]
        llvm::CrashRecoveryContext CRC;
        CRC.EnableExceptionHandler = true;

        const void *PrettyState = llvm::SavePrettyStackState();

        int Ret = 0;
        auto ExecuteClangMain = [&]() { Ret = ClangDriverMain(Argv); };

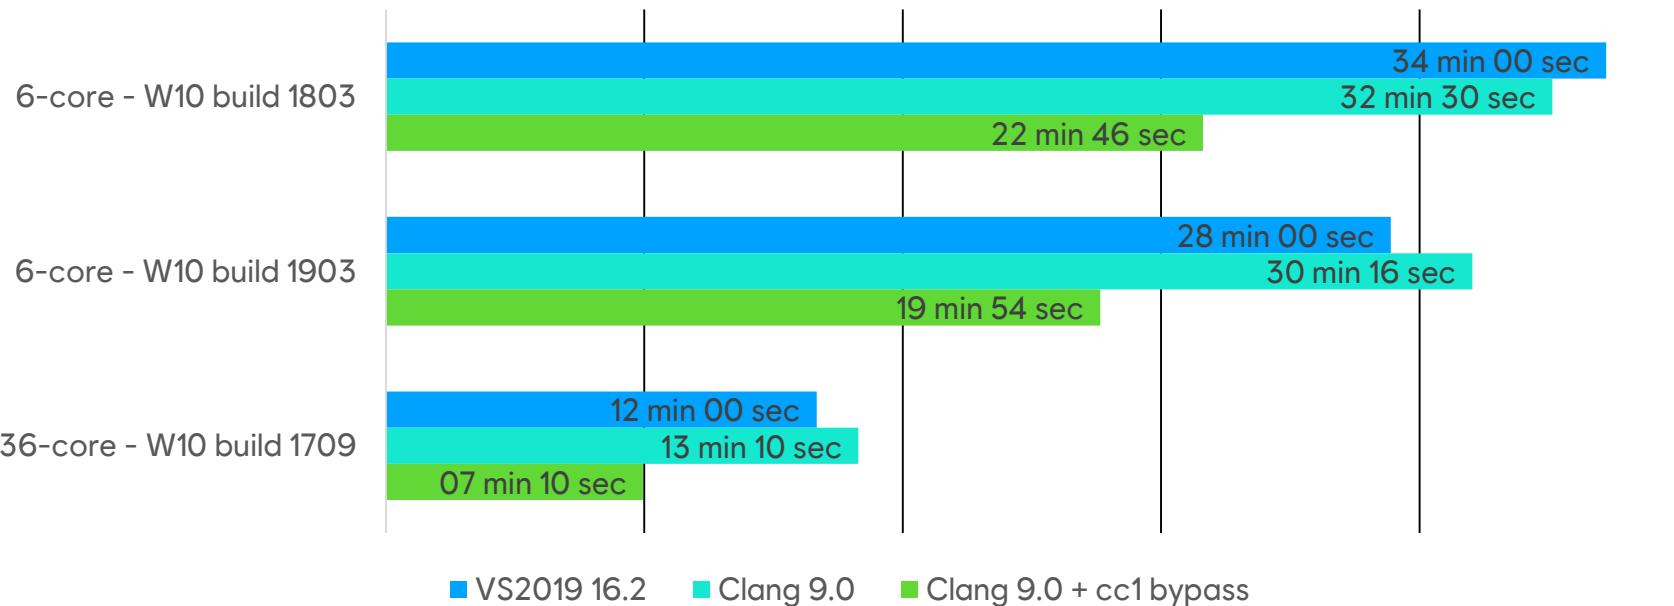
        if (!CRC.RunSafely(ExecuteClangMain)) {
            llvm::RestorePrettyStackState(PrettyState);
            return CRC.RetCode;
        }
        return Ret;
    } else {
        auto Args = llvm::toStringRefArray(Argv.data());
        return llvm::sys::ExecuteAndWait(Executable, Args, Env, Redirects,
                                         /*secondsToWait*/ 0,
                                         /*memoryLimit*/ 0, ErrMsg,
                                         ExecutionFailed);
    }
}
```



CLANG DRIVER & CC1 MERGED

ninja.exe	0.25	101,580 K	101,348 K	4,356,428 K	8220
clang-cl.exe	8.79	261,936 K	283,840 K	4,638,060 K	51044
clang-cl.exe	8.70	341,664 K	364,680 K	4,702,980 K	12980
clang-cl.exe	8.21	225,288 K	251,888 K	4,599,624 K	28236
clang-cl.exe	6.31	201,144 K	227,548 K	4,564,128 K	7480
clang-cl.exe	5.89	273,580 K	289,808 K	4,671,332 K	50408
clang-cl.exe	6.83	176,852 K	203,424 K	4,548,028 K	17296
clang-cl.exe	8.74	165,520 K	186,028 K	4,532,868 K	48632
clang-cl.exe	8.87	173,024 K	192,700 K	4,547,932 K	7276
clang-cl.exe	7.61	116,856 K	139,108 K	4,479,316 K	49772
clang-cl.exe	6.42	108,768 K	130,916 K	4,479,288 K	30200
clang-cl.exe	5.40	75,856 K	96,612 K	4,448,360 K	33116
clang-cl.exe	6.02	64,732 K	85,696 K	4,429,460 K	46696
clang-cl.exe	5.70	52,180 K	70,856 K	4,423,088 K	24380
clang-cl.exe	3.99	33,452 K	51,036 K	4,395,836 K	21816
clang-cl.exe	1.47	12,748 K	27,664 K	4,378,196 K	48468
clang-cl.exe	1.77	14,000 K	29,160 K	4,378,236 K	26784
clang-cl.exe	2.56	19,448 K	35,604 K	4,378,832 K	52520

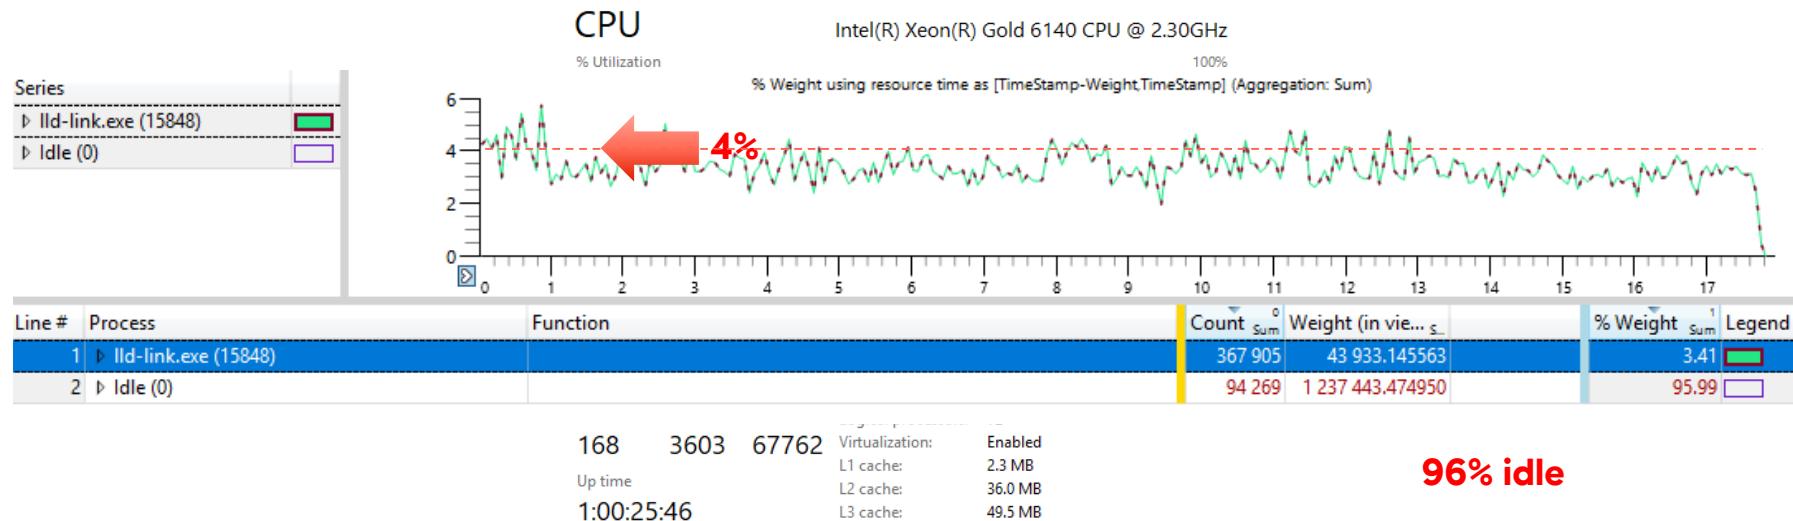
BYPASSING THE CC1 PROCESS CLEAN REBUILD LLVM, CLANG & LLD



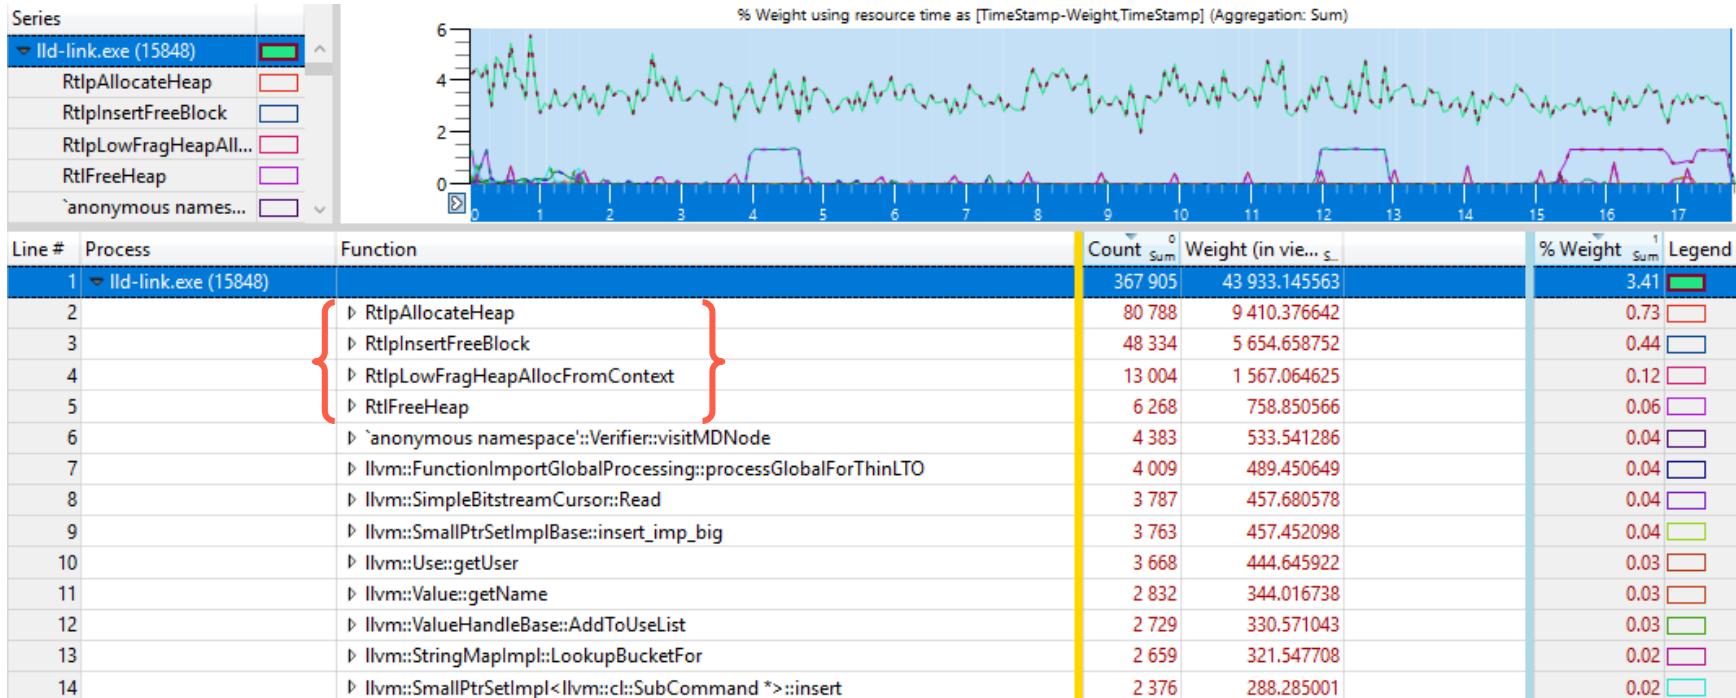
2.5

CRT Allocator

LINKING RAINBOW6: SIEGE WITH THINLTO :-(



THINLTO: ALLOCATOR CONTENTION



REPLACING THE CRT ALLOCATOR

llvm/lib/Support/Windows/Memory.inc

```
#include "rpmalloc/rpmalloc.c"

extern "C" {
    _ACRTIMP __CRTRESTRICT void *malloc(size_t size) {
        return rpmalloc(size);
    }

    _ACRTIMP void free(void *p) { rpfree(p); }

    _ACRTIMP __CRTRESTRICT void *calloc(size_t n, size_t elem_size) {
        return $PFD_FRELOAD=malloc($PFD_FRELOAD=n * elem_size);
    }

    _ACRTIMP __CRTRESTRICT void *realloc(void *ptr, size_t size) {
        return rprealloc(ptr, size);
    }
}

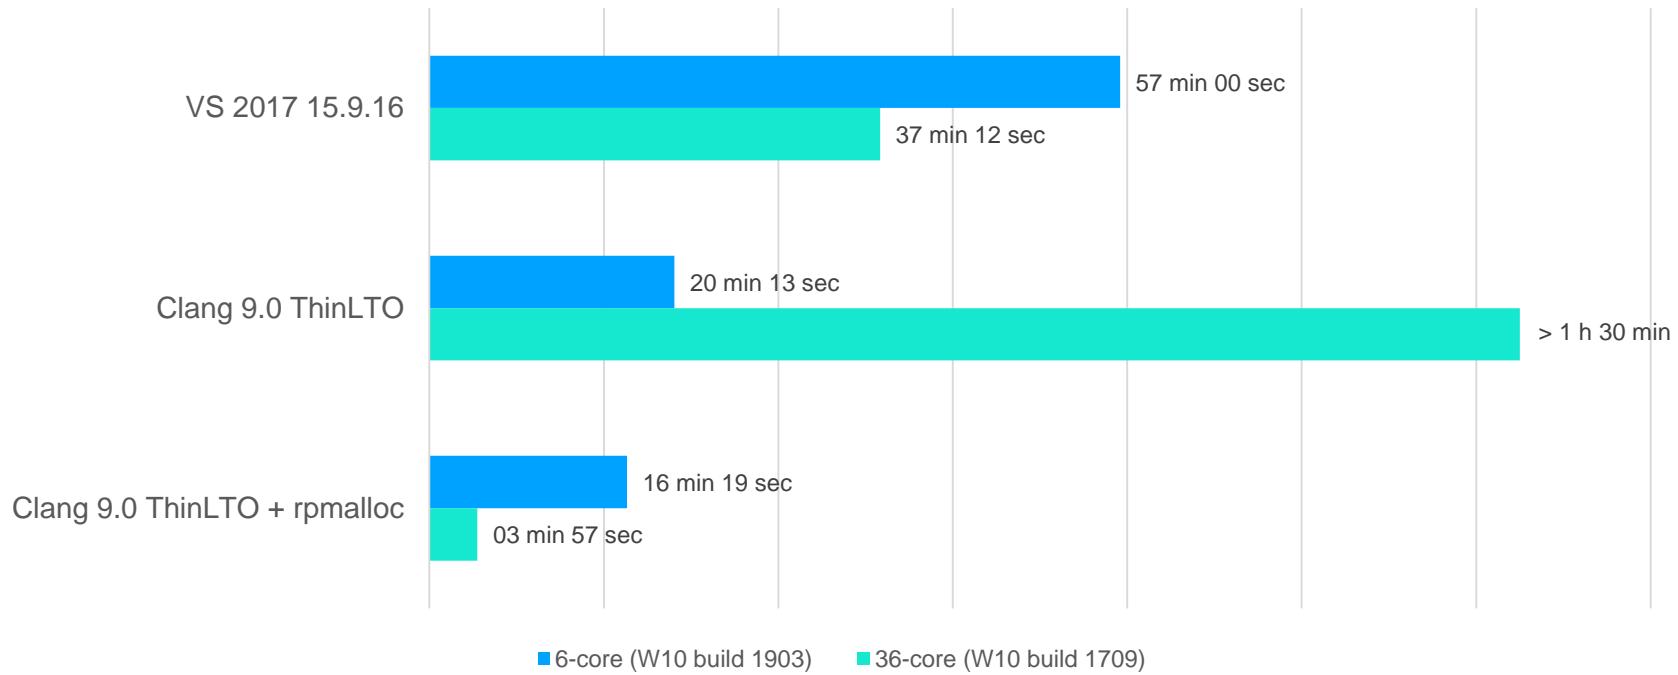
// Bypass CRT debug allocator
#ifndef _DEBUG
void *operator new(sizeof(0) n) noexcept(false) { return malloc(n); }
void __CRTDECL operator delete(void *const block) noexcept { free(block); }

void *operator new[](std::size_t s) throw(std::bad_alloc) { return malloc(s); }
void operator delete[](void *p) throw() { free(p); }
#endif
```

<https://github.com/mjansson/rpmalloc>



THINLTO (CLEAN REBUILD) RAINBOW 6: SIEGE, PC GAME PROFILE



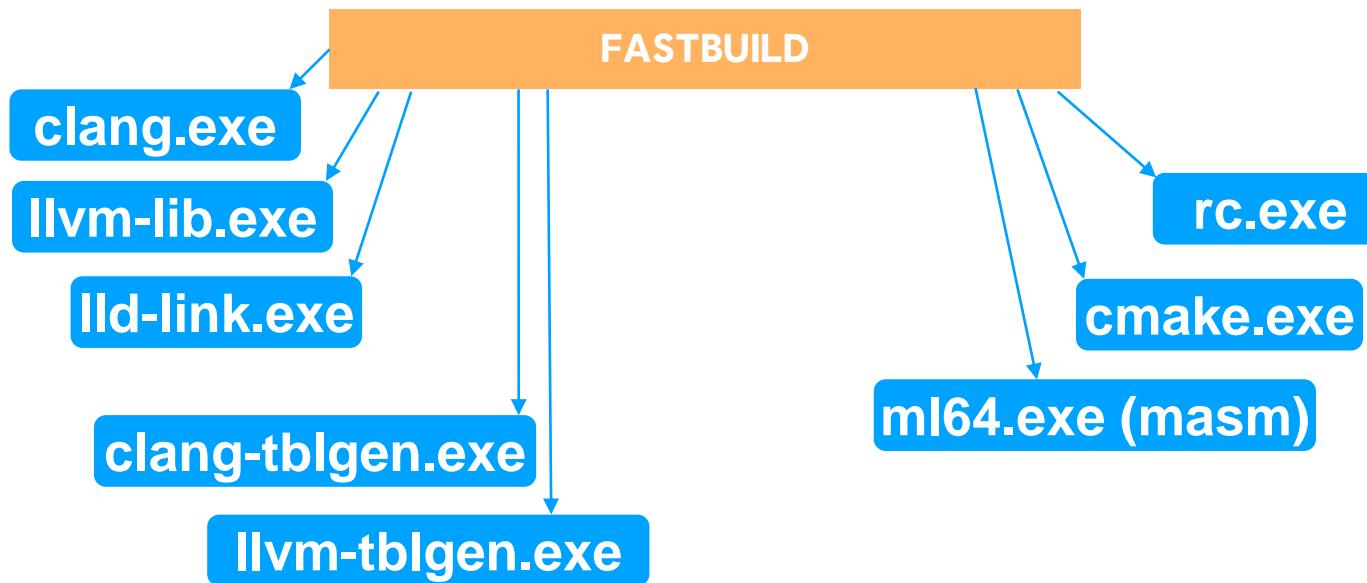


PART 3

PROPOSAL

PROOF-OF-CONCEPT

PREVIOUS BUILD PROCESS



Maybe there's a better way

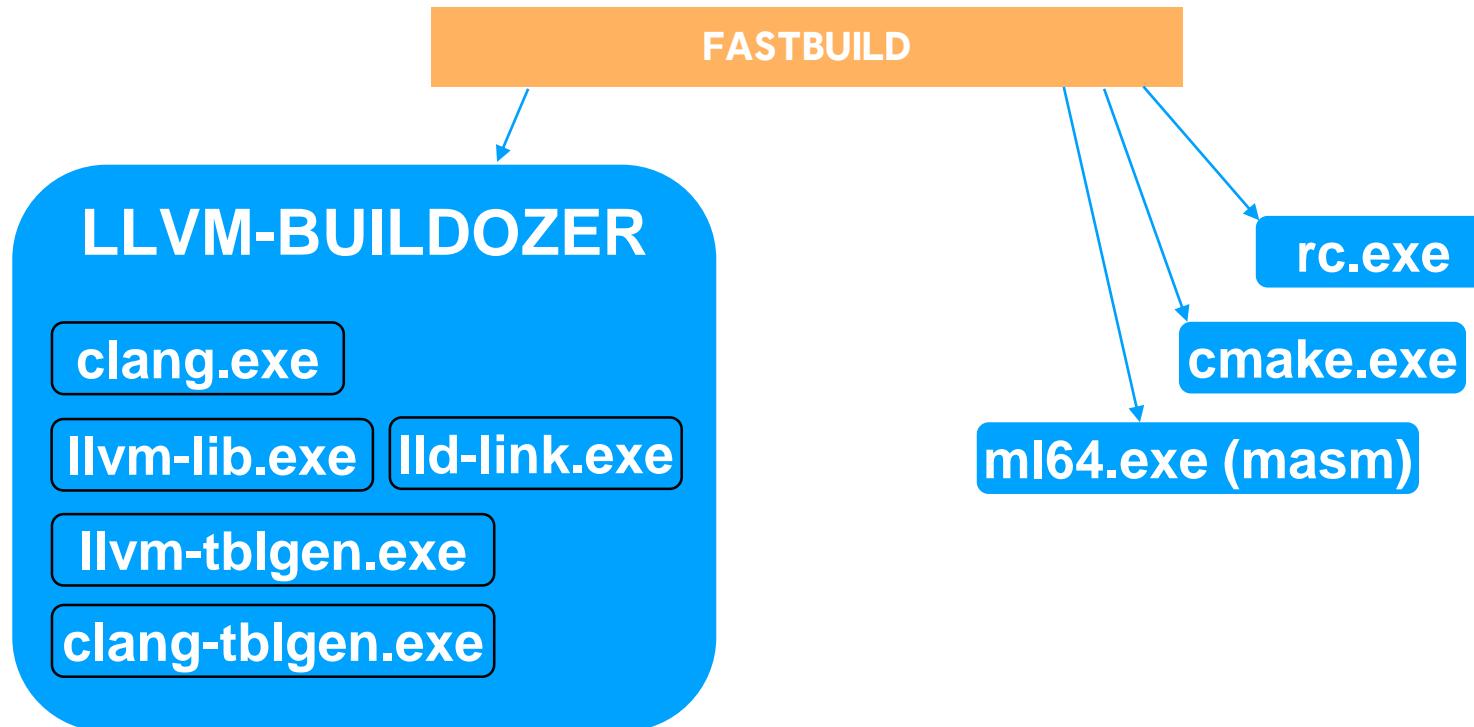


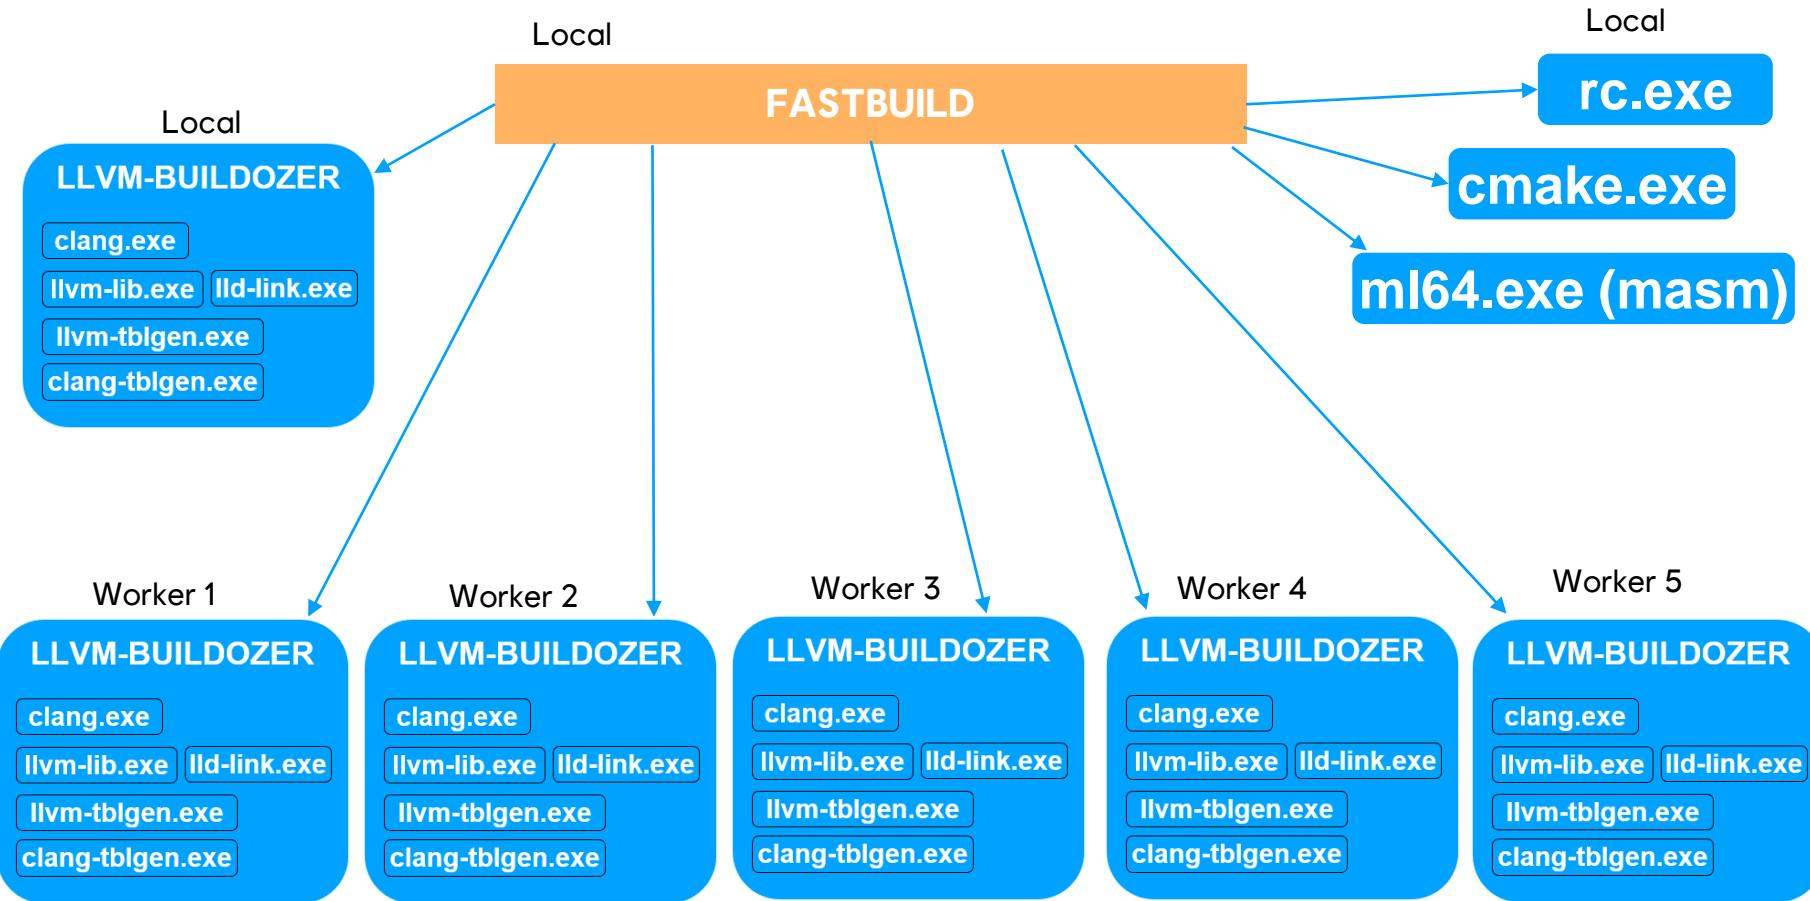
LLVM-BUILDOZER



Image Credit: Caterpillar

BUILDING WITH BUILDOZER





RUNNING THE DOZER

```
int buildozer::ImportEXE(llvm::StringRef EXE) {  
    [...]  
    HINSTANCE H = LoadLibraryA(EXE.data());  
    if (!H)  
        return 0;
```

“**LoadLibrary** can also be used to load other executable modules.[...]
However, do not use **LoadLibrary** to run an .exe file.
Instead, use the [CreateProcess](#) function.” (MSDN)

RUNNING THE DOZER

```
int buildozer::ImportEXE(llvm::StringRef EXE) {
[..]
HINSTANCE H = LoadLibraryA(EXE.data());
if (!H)
    return 0;

RemapImportAddressTable(H);
InitDebInfo();
PatchRPMalloc(M);
InitializeStaticTLS(H);
InitializeCRT(M);
FindEntryPoints(M);
[..]
}
```



RUNNING THE DOZER

```
Pool.emplace(NumWorkers, [&]() {
    while (true) {
        buildozer::WorkUnit *WU = AcquireWork(..);
        if (!WU)
            break;

        int Mod = IdentifyMOD(WU);

        llvm::CrashRecoveryContext CRC;
        CRC.RunSafely([&] {
            buildozer::Launch(Mod, WU->Directory, WU->Arguments);
        });
        [..]
    }
});
```

Pool.join();



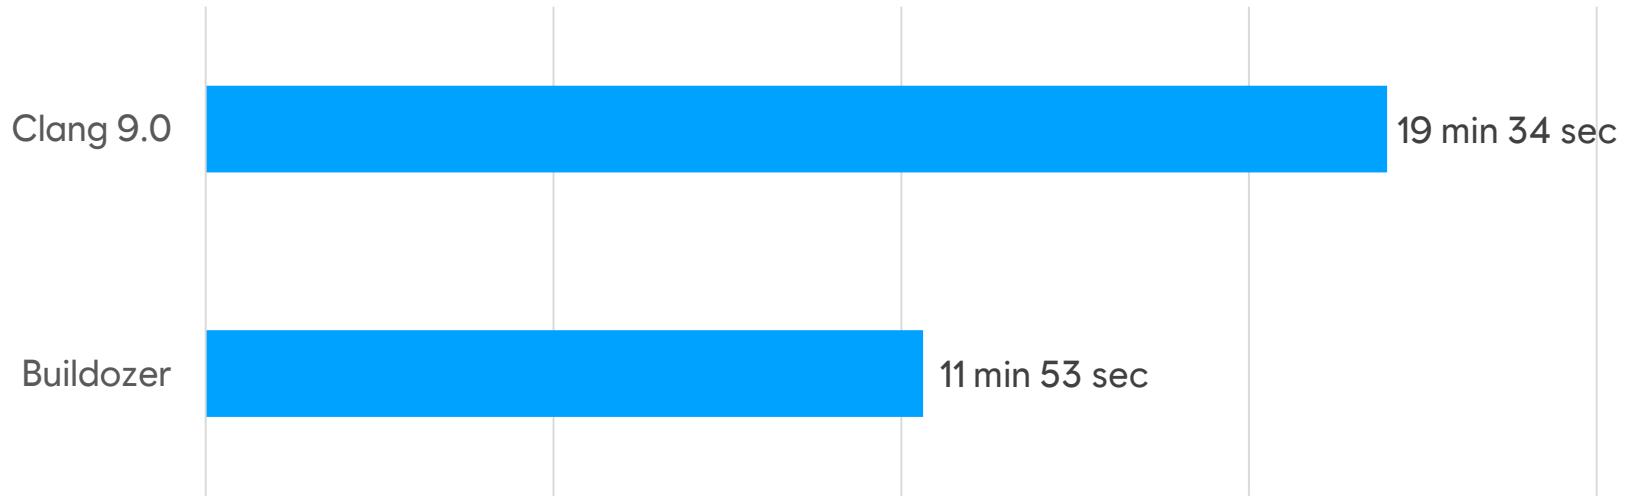
RUNNING THE DOZER

The screenshot shows a debugger interface with two main panes: 'Threads' and 'Call Stack'. The 'Threads' pane on the left lists several threads, each with a location in the code. The 'Call Stack' pane on the right shows the call stack for a specific thread, with blue arrows pointing from the stack frames to the corresponding locations in the 'Threads' pane.

Threads	Location
mainThread	clang-cl.exe!clang::ExecuteCompilerInvocation(clang::CompilerInstance ...
er.exe!thread_start<unsigned int>	clang-cl.exe!llvm::sys::fs::openNativeFileInternal
er.exe!thread_start<unsigned int>	clang-cl.exe!llvm::MachineFunction::CreateMachineInstr
er.exe!thread_start<unsigned int>	[Inline Frame] clang-cl.exe!llvm::PointerIntPair<void *,1,int,II>
er.exe!thread_start<unsigned int>	clang-cl.exe!'anonymous namespace'::MicrosoftCXXNameM
er.exe!thread_start<unsigned int>	clang-cl.exe!llvm::sys::fs::openNativeFileInternal
er.exe!thread_start<unsigned int>	clang-cl.exe!llvm::PseudoSourceValue::isAliased
er.exe!thread_start<unsigned int>	[Inline Frame] clang-cl.exe!llvm::DenseMapInfo<const llvm::...
er.exe!thread_start<unsigned int>	[Inline Frame] clang-cl.exe!llvm::SmallVectorBase::size
er.exe!thread_start<unsigned int>	clang-cl.exe!llvm::CallBase::getIntrinsicID
er.exe!thread_start<unsigned int>	clang-cl.exe!llvm::sys::fs::openNativeFileInternal
TnnWorkerThread	clang-cl.exe!llvm::sys::fs::openNativeFileInternal
TnnWorkerThread	ntr!NtWaitForWorkViaWorkerFactory

Call Stack
Name
clang-cl.exe!clang::ExecuteCompilerInvocation(clang::CompilerInstance ...)
clang-cl.exe!cc1_main(llvm::ArrayRef<const char *> Argv, const char * A...)
[Inline Frame] clang-cl.exe!ExecuteCC1Tool(llvm::ArrayRef<const char *> ...)
clang-cl.exe!ClangDriverMain(llvm::SmallVectorImpl<const char *> & ar...)
[Inline Frame] clang-cl.exe!clang::driver::Command::Execute;<unnamed-...>
clang-cl.exe!llvm::function_ref<void ()>::callback_fn<'lambda at F:\svn\c...>
clang-cl.exe!llvm::CrashRecoveryContext::RunSafely(llvm::function_ref<v...)
clang-cl.exe!clang::driver::Command::Execute(llvm::ArrayRef<llvm::Opti...)
clang-cl.exe!clang::driver::Compilation::ExecuteCommand(const clang::...)
clang-cl.exe!clang::driver::Compilation::ExecuteJobs(const clang::driver::J...)
clang-cl.exe!clang::driver::Driver::ExecuteCompilation(clang::driver::Com...)
clang-cl.exe!ClangDriverMain(llvm::SmallVectorImpl<const char *> & ar...)
clang-cl.exe!main(int argc, const char ** argv_=0x000001a4e61806f0) Li...
buildozer.exe!buildozer::Launch(int ModuleIndex, llvm::StringRef CurDir...)
[Inline Frame] buildozer.exe!ExecuteDatabase;<unnamed-tag>::operator...>
buildozer.exe!fun<void ()>::callback_fn<'lambda at F:\svn\c...>

Local build, AAA game, x64 Editor Release



Intel Xeon W-2135 @ 3.7 GHz (6-core), 128 GB, NVMe SSD

The background of the slide is a dark, atmospheric forest scene at night. On the left, there's a large tree trunk with glowing orange mushrooms growing from its base. In the center-right, a small campfire is burning, with a few glowing embers scattered on the ground around it. The overall mood is mysterious and magical.

PART 4

NEXT STEPS

SHORT TERM

- Remove OS jitter (in-RAM file content & stat cache)
- OBJ cache (in-RAM)
- Clang-LLD in-memory bridge
- Incrementally link along the way
- Incrementally compile along the way (SN Systems' Program Repository)
- Remote API for distribution & caching

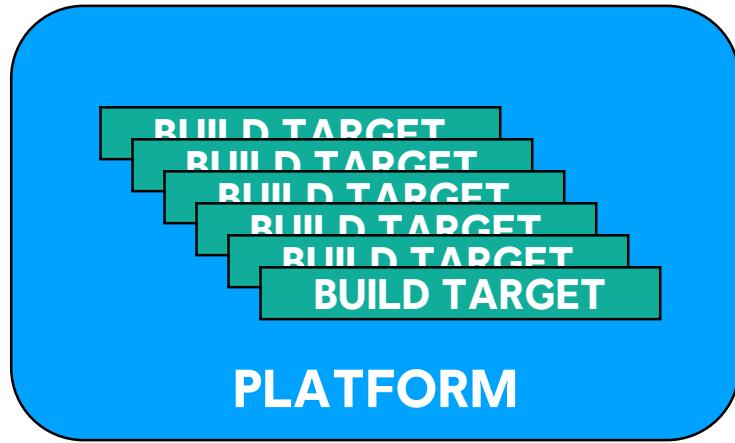


LONG TERM

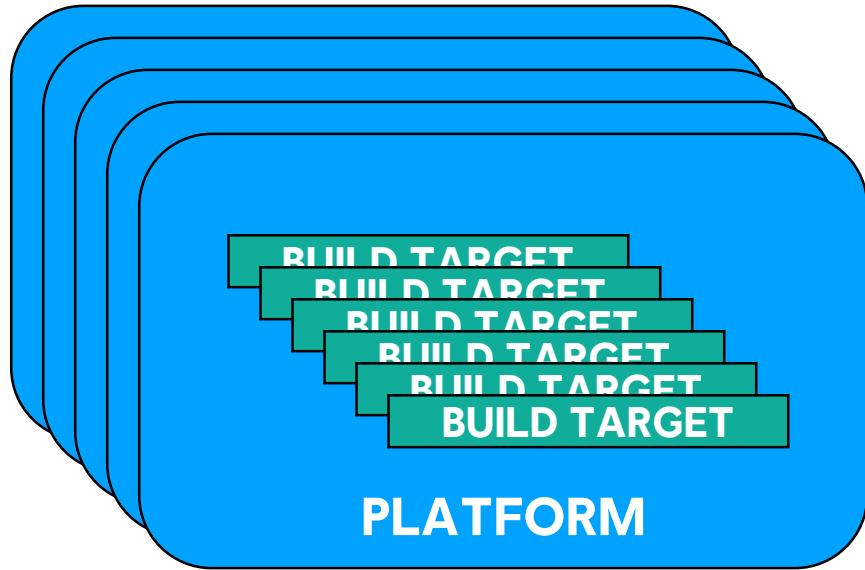
BUILD TARGET



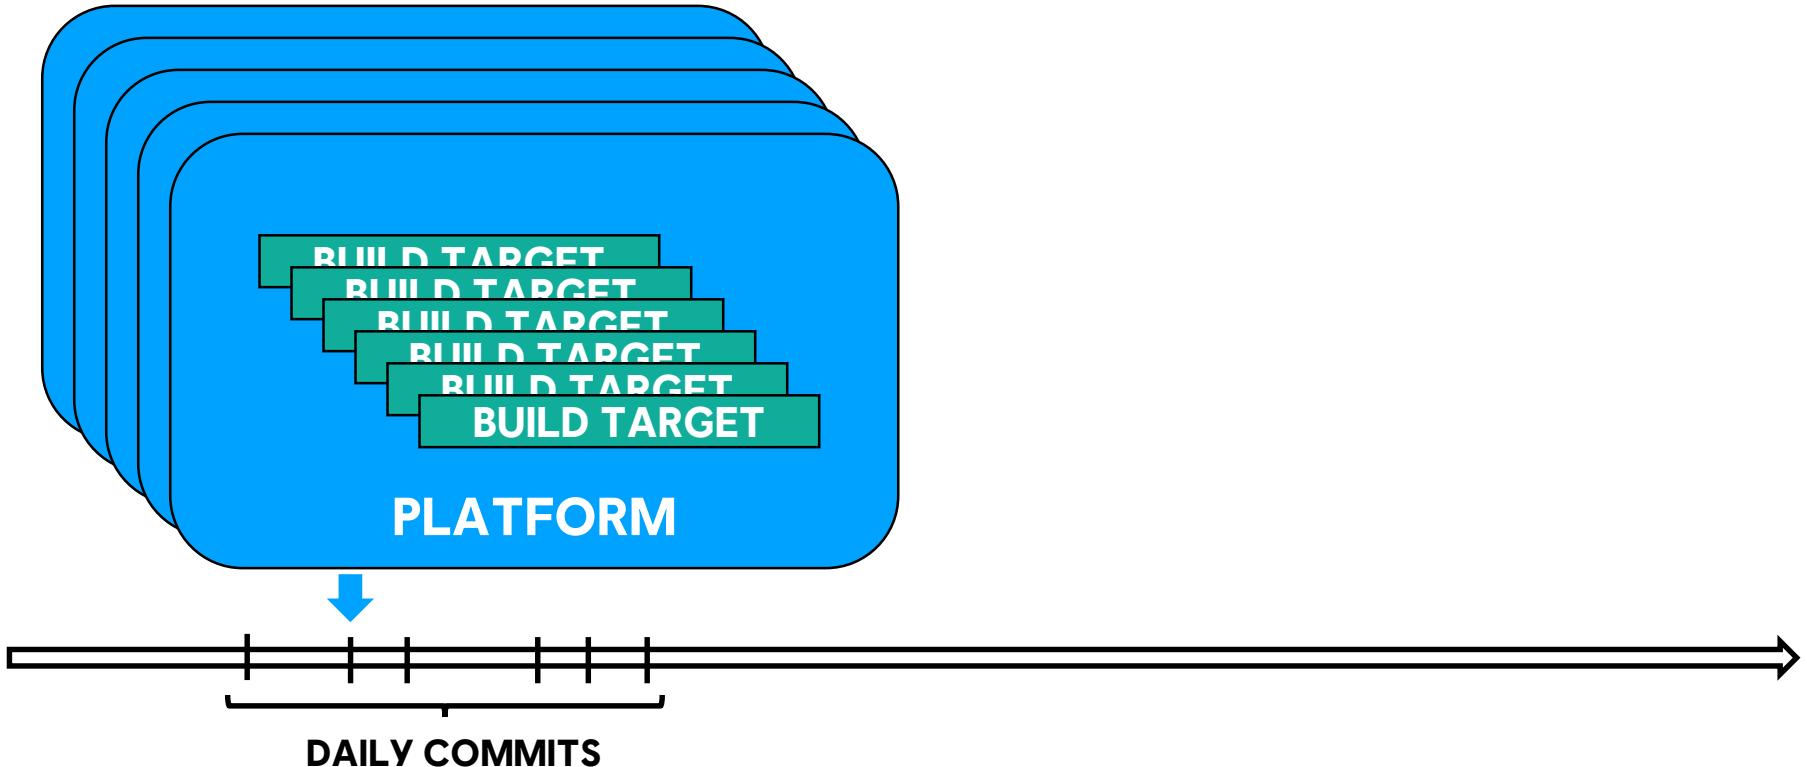
LONG TERM



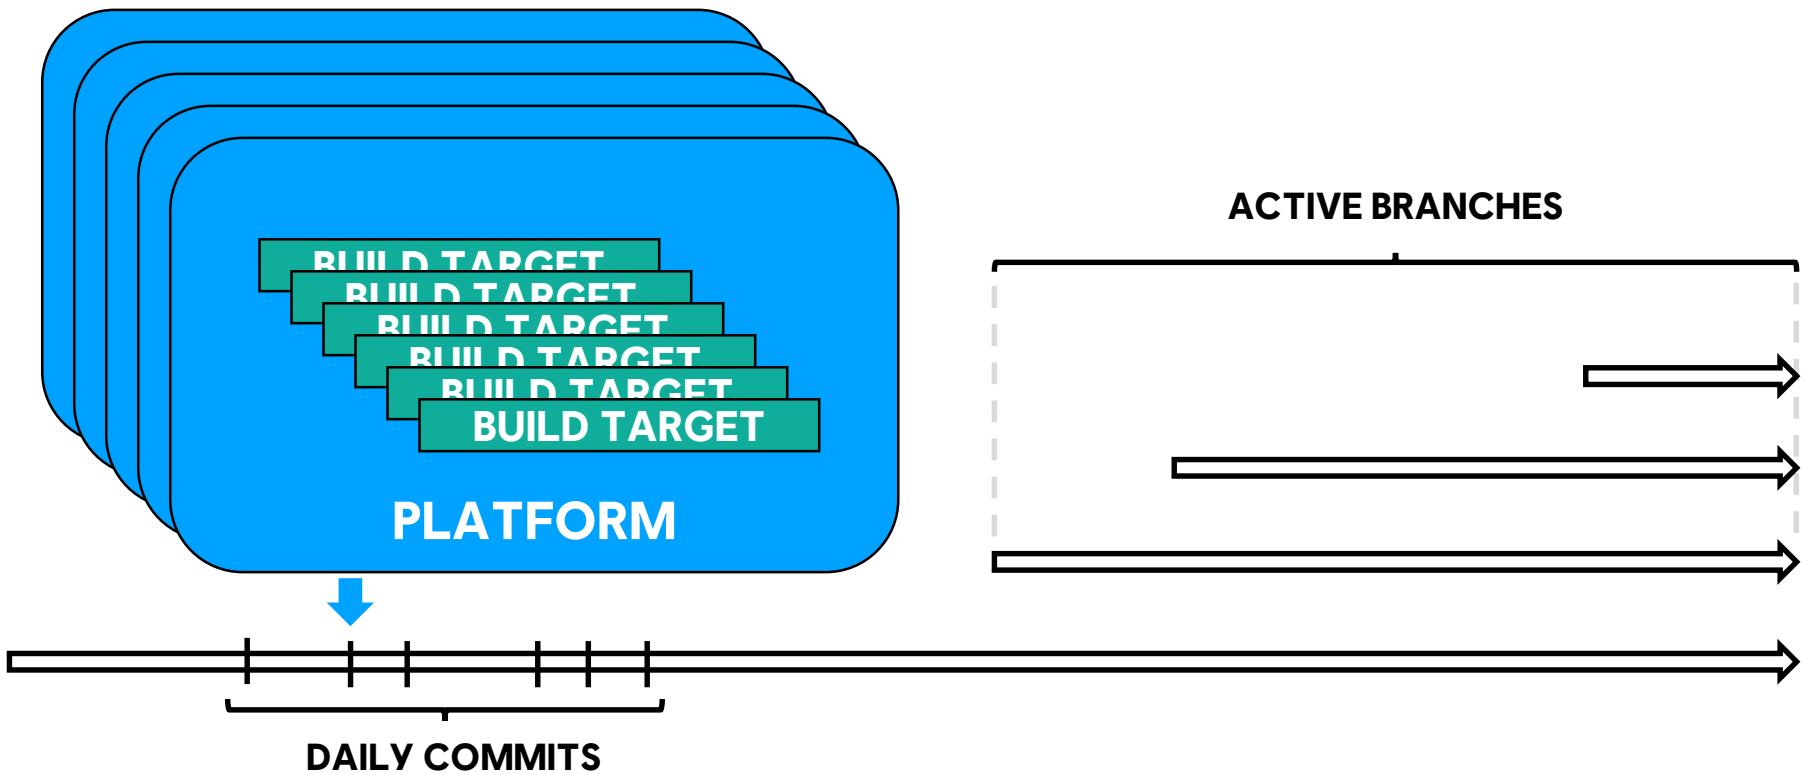
LONG TERM



LONG TERM



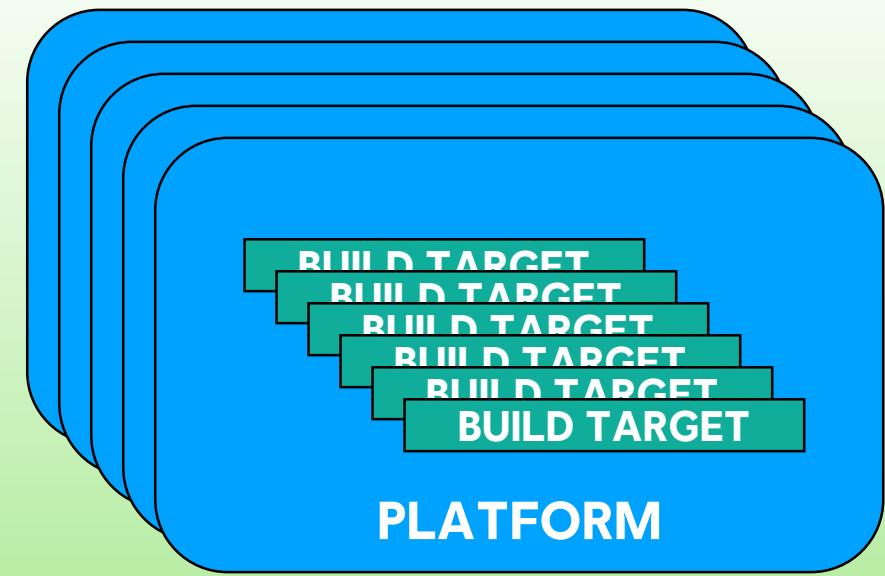
LONG TERM



LONG TERM

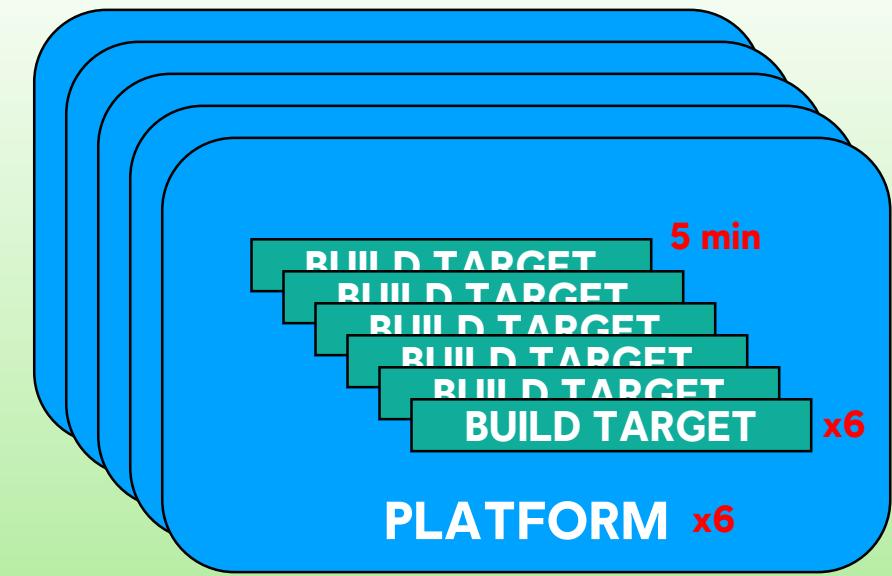
GAME PRODUCTION

ACTIVE BRANCHES



LONG TERM

x20 GAME PRODUCTION



ACTIVE BRANCHES x4



Is there a better way?





THANK YOU



The image is a collage of six vertical video game character portraits from Ubisoft games. From left to right: 1. A woman in red and gold Roman-style armor. 2. A pig-headed man in a blue suit and monocle. 3. A futuristic soldier in a black suit and helmet. 4. A woman in a blue hoodie. 5. A person in a dark, heavily armored suit. 6. A woman in a traditional East Asian outfit with a mask.

Q&A

Alexandre Ganea, Ubisoft
alexandre.ganea@ubisoft.com