

2019

LLVM DEV MEETING



Orchestrating a brighter world



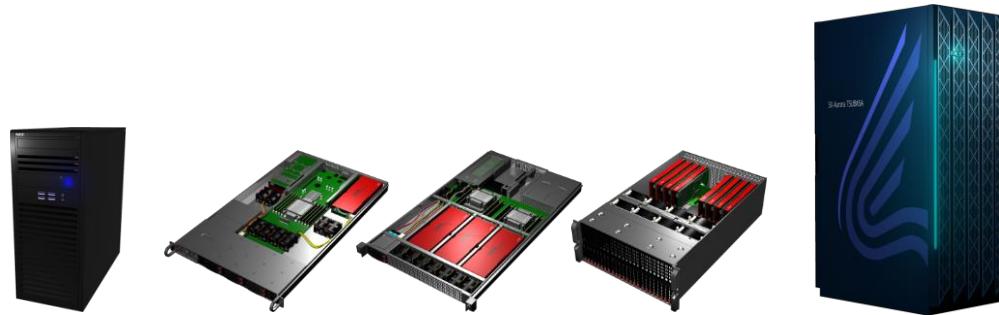
NEC SX-Aurora as a Scalable Vector Playground

Kazuhisa Ishizaka, Yoshiyuki Ohno, Yuta Ideguchi, Erich Focht, Simon Moll

simon.moll@emea.nec.com

NEC Corporation

High performance computer ranged from deskside to cloud and supercomputer



Vector Processor on a Vector Engine PCIe accelerator card



Vector Engine (VE)
(PCIe accelerator card)



Vector Processor with 6 HBM2s

Theoretical Performance

Computation:

FP32: 4.30 TFlops

FP64: 2.15 TFlops

Memory:

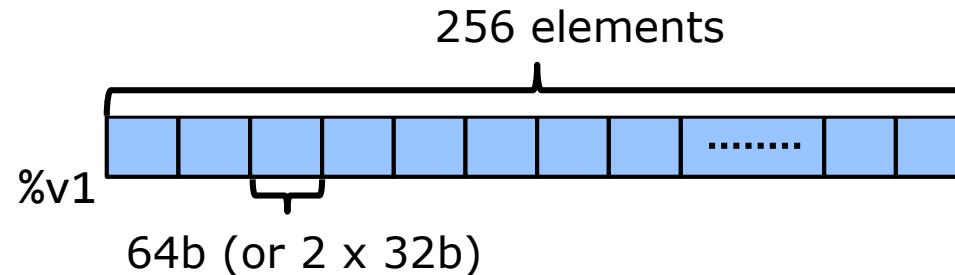
Capacity: 48GB

Bandwidth: 1.2 TB/s

Basics

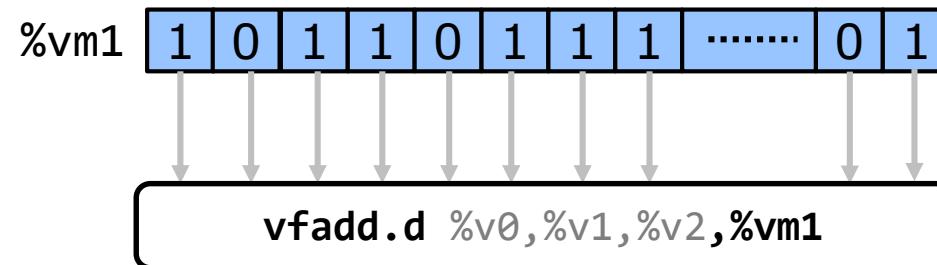
- **Wide vector registers (256 x 64bit)**

- 64 Vector registers ($\%vi$)



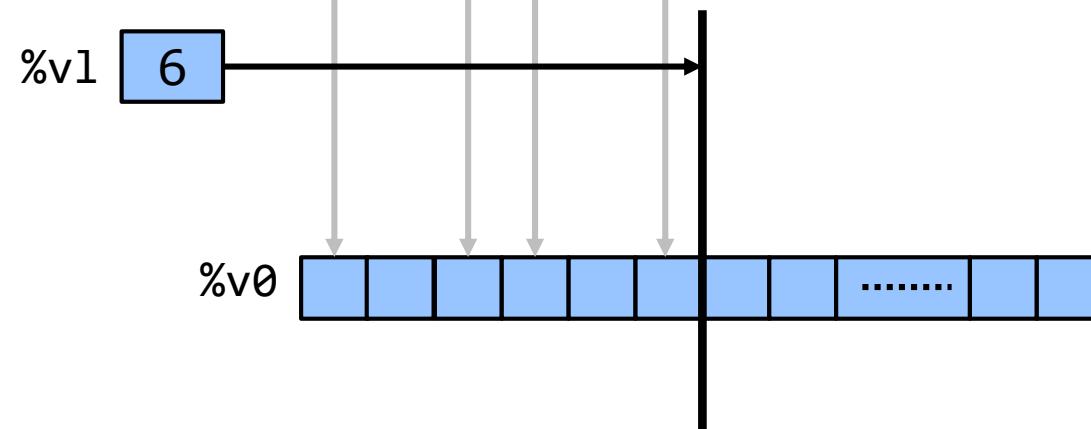
- **Full vector predication**

- 16 Vector mask registers ($\%vmi$)
- *Explicit* operand



- **(Active) Vector Length Register**

- One, global, VL Register
- *Implicit* dependence



1) Implicit dependency through VL register

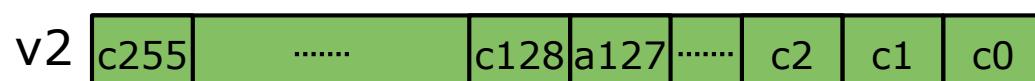
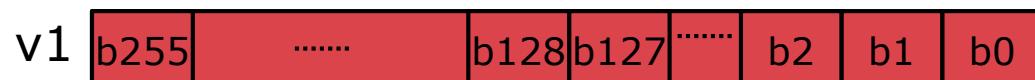
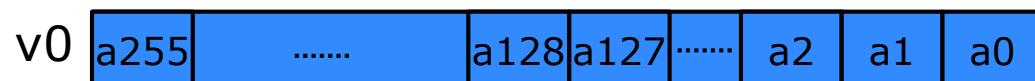
```
1v1 %s37  
vld %v3,8,%s0  
vld %v4,8,%s1  
vfmad.d %v3,%v3,%s2,%v4
```

define VL

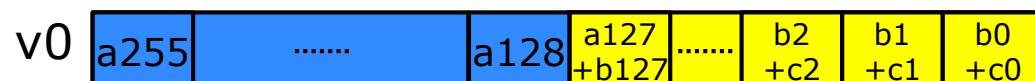
use VL

How do we implement implicit def-use?

2) Partial update of a destination vector register



vfadd.d %v0,%v1,%v2 (VL=128)



not updated

updated($v1 + v2$)

How do we introduce partial update?

1. Vector IR based on LLVM-VP

- VL as parameter

```
$v3 = vfadd.d $v4,$v5,$pt,$vl  
# for i=0,256  
#   v3[i] = i < vl ? v4[i] + v5[i] : pt[i]
```

+

2. Automatic LVL generation in backend

- Inserting LVL instruction from \$vl argument in IR
- Minimizing LVL instruction by current VL inference



Status and Roadmap

Status

- LLVM-VE is available at <https://github.com/sx-aurora-dev/llvm>
- Scalar code backend + vector intrinsics
- Application: TensorFlow for SX-Aurora

Roadmap

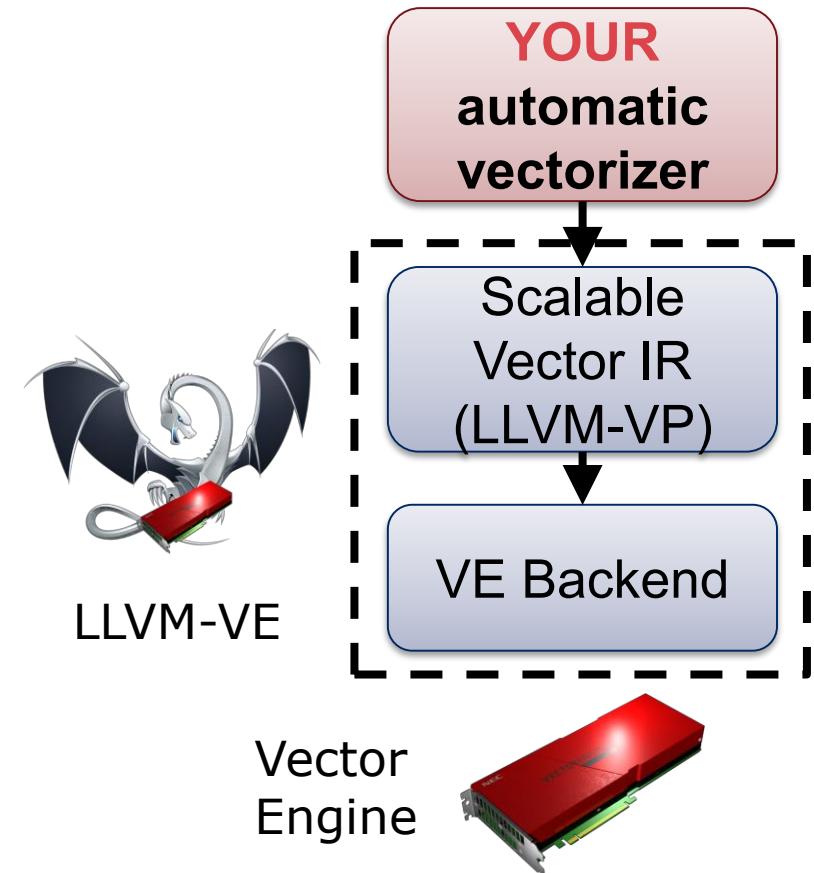
- Upstreaming! <https://reviews.llvm.org/D69103>
- Vector predication with LLVM-VP (D57504)

Welcome collaborators

- for automatic vectorization, etc.

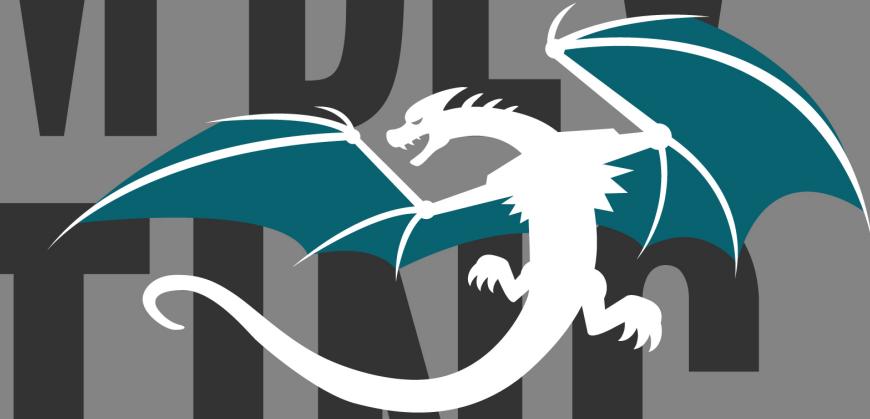
Come see us at the poster session!

"NEC SX-Aurora as a Scalable Vector Playground"



2019

LLVM DEV MEETING



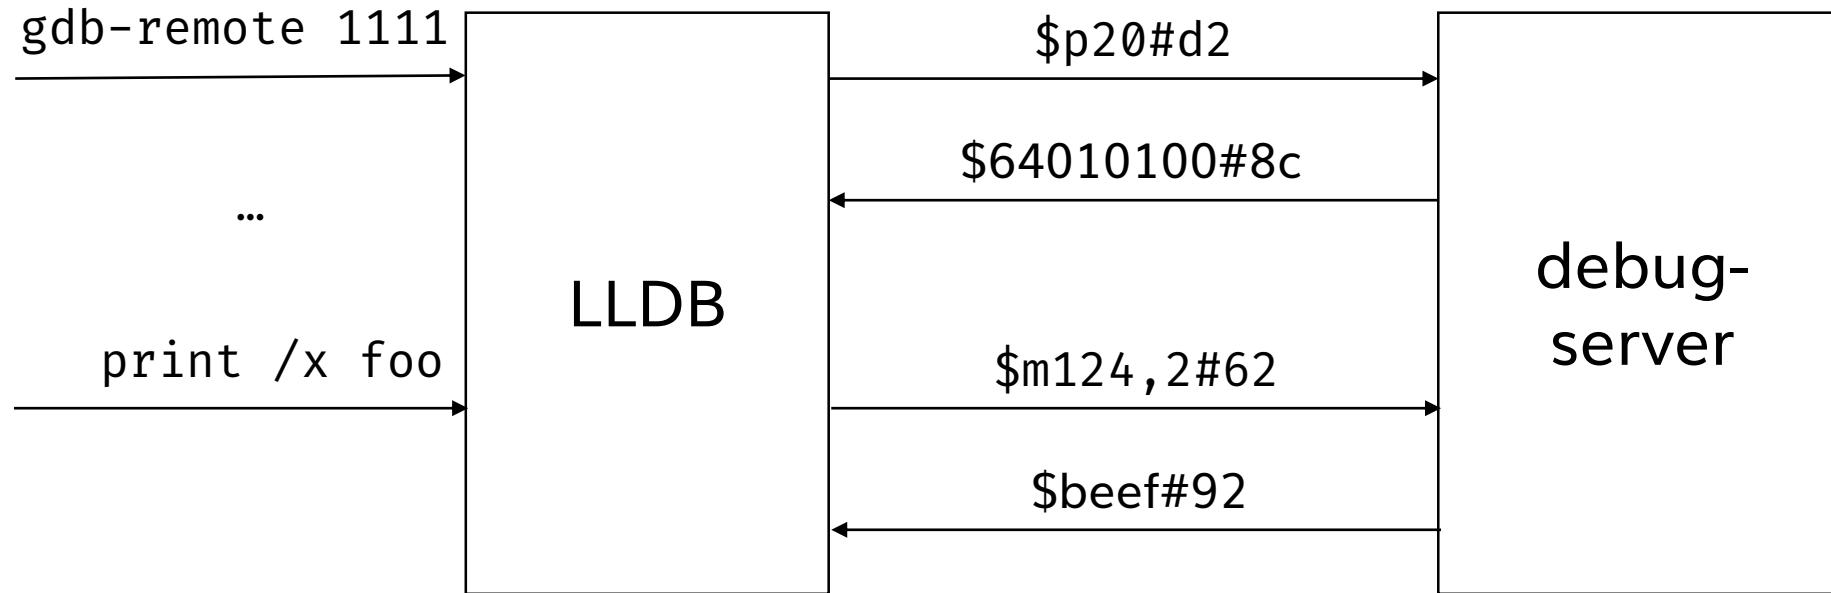


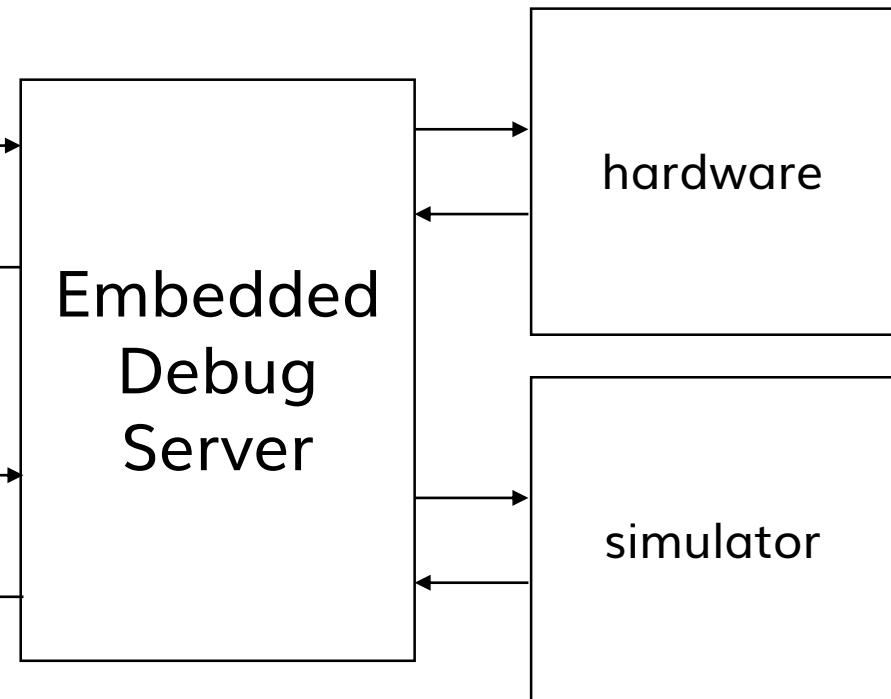
A Unified Debug Server for Deeply Embedded Systems and LLDB

Simon Cook



Copyright © 2019 Embecosm.
Freely available under a Creative Commons license.



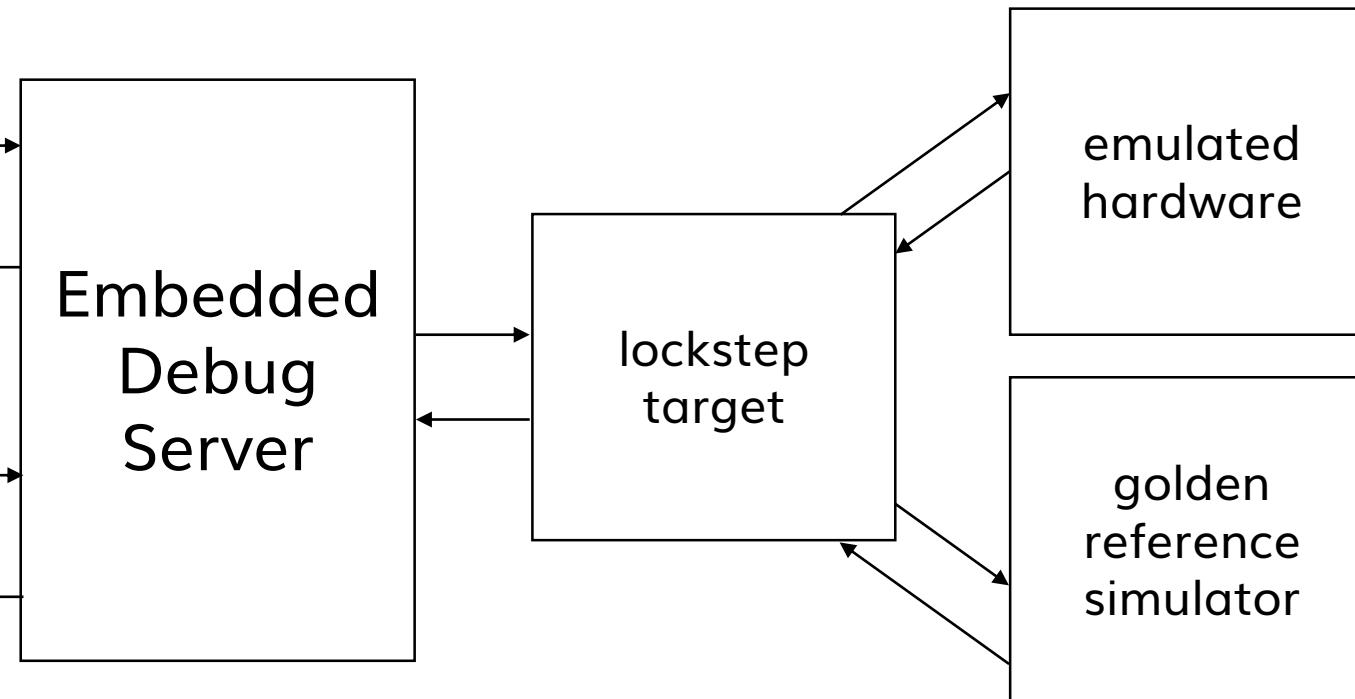


```

class ExampleTarget : public Target {
public:
    // Timers
    uint64_t getCycleCount() const;
    uint64_t getInstrCount() const;

    // Read-write memory and registers
    std::size_t readRegister(const int reg,
                           uint_reg_t &value);
    std::size_t writeRegister(const int reg,
                           const uint_reg_t value);
    std::size_t read(const uint_addr_t addr,
                     uint8_t *buffer,
                     const std::size_t size);
    std::size_t write(const uint_addr_t addr,
                     const uint8_t *buffer,
                     const std::size_t size);

    // Execution control
    bool prepare(const std::vector<ResumeType> &actions);
    bool resume(void) override;
    WaitRes wait(std::vector<ResumeRes> &results);
    ...
};
```



```

std::size_t
LockstepTarget::readRegister(const int reg,
                             uint_reg_t &value) {
    uint_reg_t vL;
    uint_reg_t vR;
    std::size_t rL = _l->readRegister(reg, vL);
    std::size_t rR = _r->readRegister(reg, vR);

    // Report inconsistency server side.
    if ((rL != rR) || (vL != vR))
        std::cerr << "Lockstep error: register"
                  << " inconsistency."
                  << std::endl;

    // TODO: Allow this to be configured.
    // For now we just choose the left value
    // since readRegister cannot fail.
    value = vL;
    return rL;
}

```

Lockstep

```
$ lldb  
(lldb) gdb-remote 51000  
Process 1 stopped  
* thread #1, stop reason = signal SIGTRAP  
frame #0: 0x0001015c  
-> 0x1015c: addi a5, zero, 5  
0x10160: mv t6, a5  
0x10164: mv a5, zero  
0x10168: mv a0, a5  
(lldb) si  
Process 1 stopped  
* thread #1, stop reason = instruction step into  
frame #0: 0x00010160  
-> 0x10160: mv t6, a5  
0x10164: mv a5, zero  
0x10168: mv a0, a5  
0x1016c: lw s0, 12(sp)  
(lldb) si  
Process 1 stopped  
* thread #1, stop reason = signal SIGSYS  
frame #0: 0x00010164  
-> 0x10164: mv a5, zero  
0x10168: mv a0, a5  
0x1016c: lw s0, 12(sp)  
0x10170: addi sp, sp, 16
```

SIGSYS returned, indicating divergence

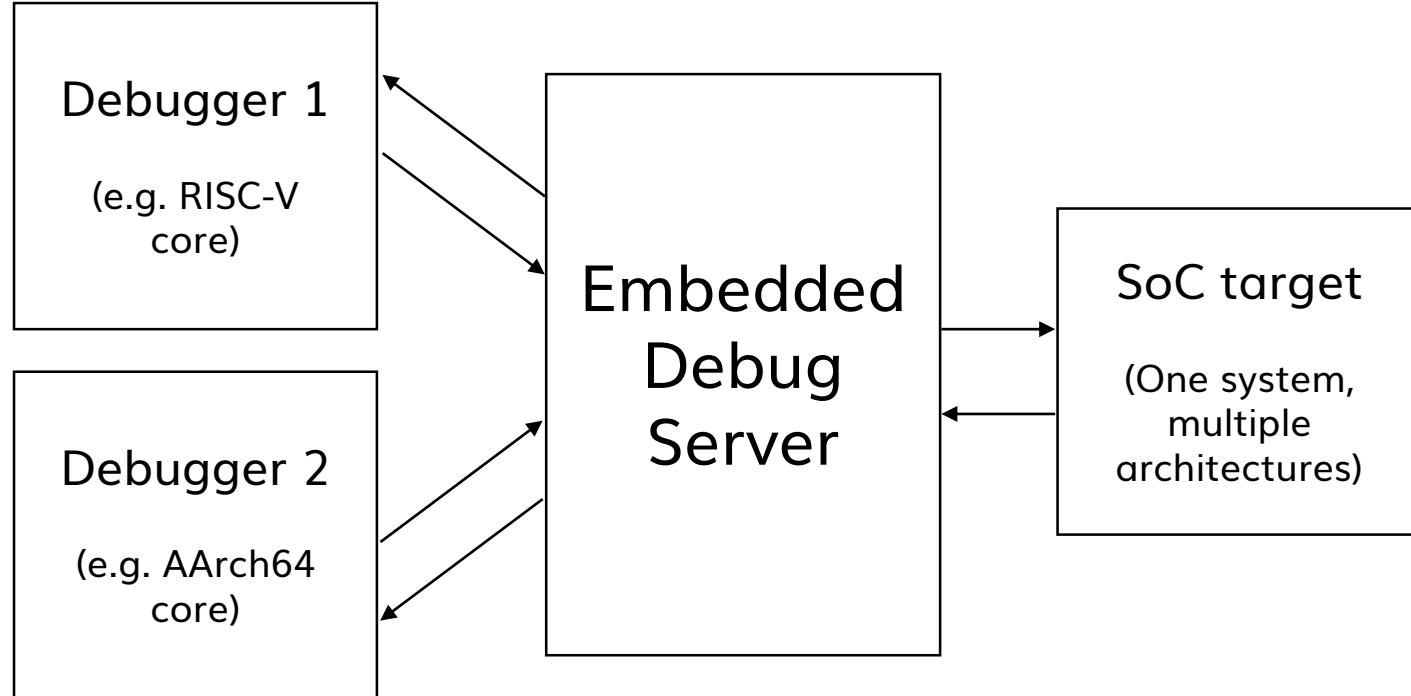
```
(lldb) register read --all  
general:  
x0 = 0x00000000  
x1 = 0x000100d8  
x2 = 0xfffffffff0  
x3 = 0x00011da8  
< ... >  
x28 = 0x00000000  
x29 = 0x00000000  
x30 = 0x00000000  
pc = 0x00010164  
1 registers were unavailable.
```

x31 not shown, since both targets disagree on value

```
lockstep-left:  
x31-left = 0x00000005 ←  
32 registers were unavailable.
```

```
lockstep-right:  
x31-right = 0x00000004 ←  
32 registers were unavailable.
```

both targets x31 shown, allowing investigation into divergence





Questions?

www.embecosm.com



Copyright © 2019 Embecosm.
Freely available under a Creative Commons license.

2019

LLVM DEV MEETING



Speculative Compilation in ORC



ORC

- LLVM Modular Just in Time Compilation Library
- Custom compilers, program representations...
- Supports concurrent compilation

JIT Variants

- Eager JIT - high startup time, zero compiler interactions at runtime
- Lazy JIT - ⚡ startup time, compilation overhead on first call

Can we do better? Can we have benefits of two worlds?



Let's guess it!

```
void Driving(Signal S)

    switch(S) {
        case red:stop();tweet();

        break;

        case yellow:

        like_reply_to_a_tweet();break;

        case green :think_next_tweet();

        break;

    }
}
```

What if we guess the signal's outcome and do action!

Likewise, we guess control flow path and compile the *likely* functions before calling them.

Let's guess it!

```
void Driving(Signal S)

    switch(S) {
        case red:stop();tweet();
        break;

        case yellow:
            like_reply_to_a_tweet();break;

        case green :think_next_tweet();
        break;
    }
}
```

What if we guess the signal's outcome and do action!

Likewise, we guess control flow path and compile the *likely* functions before calling them.



Speculative Decisions

- Compile only the most likely next executable functions
- Speculate based on CFG edge probabilities and hot blocks heuristics
- Implemented as `SpeculateQuery` function objects, you can try your own ideas 😊
 - `Map<Function, LikelyFunctionSymbols> SpeculateQuery(Function& F);`
 - ➔ “Jump into JIT through IR Instrumentation”

Mix'n'Match

```
ObjectLinkingLayer LinkLayer;  
  
IRCompileLayer<...> CompileLayer(LinkLayer, ConcurrentCompiler);  
  
IRSspeculationLayer SpeculateLayer(..., CompileLayer, Speculator, SpeculateQuery)  
CompileOnDemandLayer<...> CODLayer(SpeculateLayer, ...);  
CODLayer.addModule(Mod, MemMgr, SymResolver);  
  
auto FooSym = CODLayer.findSymbol("foo", true);  
  
auto Foo = reinterpret_cast<int(*)()>(FooSym.getAddress());  
  
int Result = Foo(); // <-- Call foo's stub.
```



Before

```
define dso_local void @Driving(i32 %Signal)

switch i32 %Signal, label %exit[i32 0, label %red
    i32 1, label %yellow i32 2, label %green]

red:

call void @stop()  call void @tweet()

yellow:

call void @like_reply_to_a_tweet()

green:

call void @think_next_tweet()
```

After...

```
@__orc_speculator = external global %Class.Speculator

declare void @_orc_speculate_for(%Class.Speculator* %0, i64 %1)

define dso_local void @Driving(i32 %0) #0 {

    call void @_orc_speculate_for(%Class.Speculator* @_orc_speculator,
        i64 ptrtoint (i32 ()* @Driving to i64)) // Jump into JIT ↗

    ...

    switch i32 %3, label %7 [ i32 0, label %Red
        i32 1, label %Yellow
        i32 2, label %Green]
    ...

}
```

Performance

7x Speed-Down 💔💔

Multiple Jumps into JIT ✗

ExecutionSession::Lookup's are not free ✓

We want Performance



- Relatively easy fix, guard the orc_speculate_for call
- Jump into JIT only on the first call
- This will give us - what we want

```
@__orc_speculate.guard.for.main = internal local_unnamed_addr global i8 0, align 1

define dso_local void @Driving(i32 %0) {

__orc_speculate.decision.block:

%guard.value = load i8, i8* @_orc_speculate.guard.for.main

%compare.to.speculate = icmp eq i8 %guard.value, 0

br i1 %compare.to.speculate, label %__orc_speculate.block, label %program.entry

__orc_speculate.block:

call void @_orc_speculate_for(%Class.Speculator* @_orc_speculator,
i64 ptrtoint (i32 ()* @Driving to i64))

store i8 1, i8* @_orc_speculate.guard.for.main

br label %program.entry
```

Performance

We see significant speedup with our proof-of-concept speculative jit

For SPEC 403.gcc benchmark, reduce exec time from 17.4 seconds to
10.5 seconds (4 threads) 😊

What's Next?

- Finish dynamic profiling support to collect branch probability information
- Reduce the scope of speculation region in a function
- Implementing more SpeculateQueries
- Performance tuning

Thank you so much - Lang Hames and David Blaikie

Thank you LLVM Foundation and Conf Sponsors.

2019

LLVM DEV
MEETING



Loom

Weaving Instrumentation for Program Analysis

Brian Kidney (Presenter)
Jonathan Anderson
Memorial University

But Instrumentation is done, right?

Why another instrumentation tool

- There are lots of instrumentation tools
 - Intel Pin, XRay, CSI, DTrace...
- Most tools focus on performance
- We needed something different
 - No custom compiler frontend or backend
 - “Non-traditional” instrumentation
 - The ability to transform code when needed
- Our focus was security
 - We wanted something more general-purpose

What if you wanted to instrument every
PAM Authentication?

Instrumenting PAM

And you could do it with this:

```
strategy: callout
dtrace: userspace
functions:
- callee: [entry]
metadata:
  name: auth
  id: 1
name: pam_authenticate
```

Instrumenting PAM

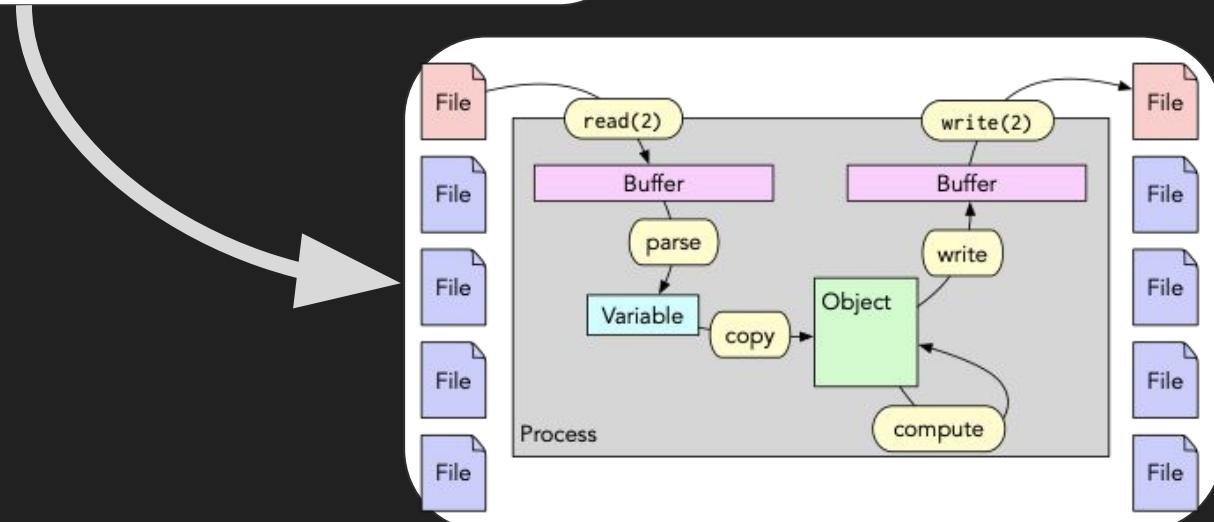
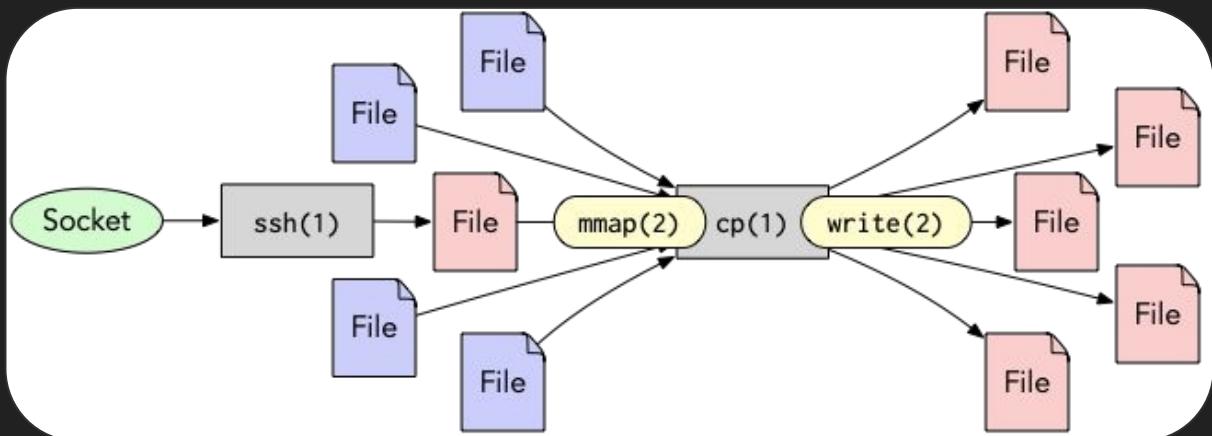
And you could do it with this:

```
strategy: callout
dtrace: userspace
functions:
- callee: [entry]
metadata:
  name: auth
  id: 1
name: pam_authenticate
```

And you get this:

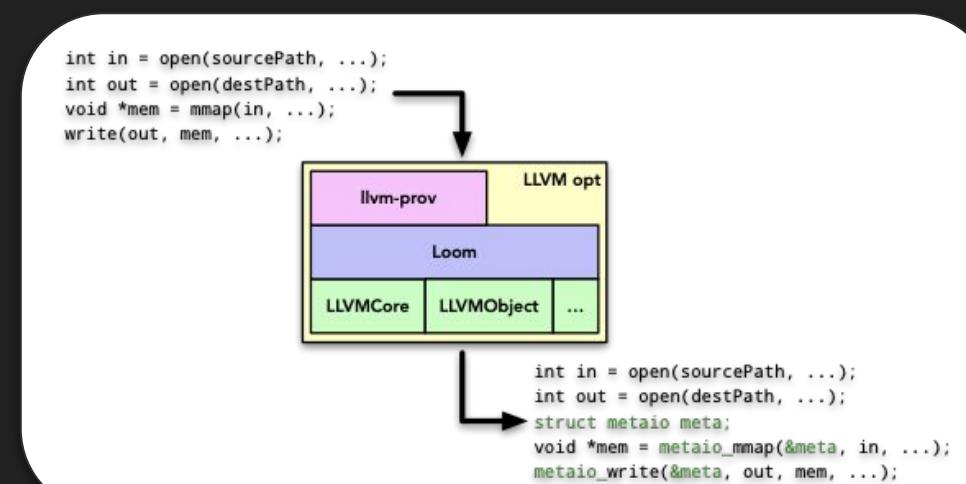
```
dtrace: script './pam.d' matched 6 probes
CPU ID          FUNCTION:NAME
  2 61725        none:dt-probe
Pam Authentication - execname: sshd, user: bkidney
  0 61725        none:dt-probe
Pam Authentication - execname: sshd, user: steve
  2 61725        none:dt-probe
Pam Authentication - execname: sshd, user: badguy
```

What if you wanted to transform system calls for provenance instrumentation?



llvm-prov

- Too complex an application for the policy file
- Transforms system call API to new API
- Domain specific logic to determine code of interest
- Loom emits code transformation



Loom

- Simple Policy files to apply instrumentation without code modification
- Framework to build custom tools when you need more

Loom

- Simple Policy files to apply instrumentation without code modification
- Framework to build custom tools when you need more

Currently supports instrumenting:

- Struct fields
- Functions
 - Callee / Caller
- Global Variables
- Pointer Instructions

Supported Outputs:

- Logging
 - KTrace, DTrace, Text, Json, XML
- Code transformation

Loom

- Simple Policy files to apply instrumentation without code modification
- Framework to build custom tools when you need more
- Work has started on instrumentation language
 - DAG matching for code transformation

Currently supports instrumenting:

- Struct fields
- Functions
 - Callee / Caller
- Global Variables
- Pointer Instructions

Supported Outputs:

- Logging
 - KTrace, DTrace, Text, Json, XML
- Code transformation

Come see more
during poster session!

2019

LLVM DEV
MEETING



2019

LLVM DEV MEETING

Supporting 64 bit pointers in RISCV 32 bit LLVM backend

Reshabh Sharma

Background:

Prof. Taylor's [Bespoke Silicon Group](#) is developing a GP-GPU based on RISC-V 32 bit ISA (RV32)

Why 32 bit ISA?

It is good for very high energy efficiency and density.

What is a good first thing you would expect from a GP-GPU?

Access to 64 bit addressable DRAM

We provide access to 64 bit addresses in DRAM using custom load and store instructions.

LDW rd, rs1, rs2

SDW rd, rs1, rs2

Where rs1 and rs2 always store 32 bit halves of a 64 bit address

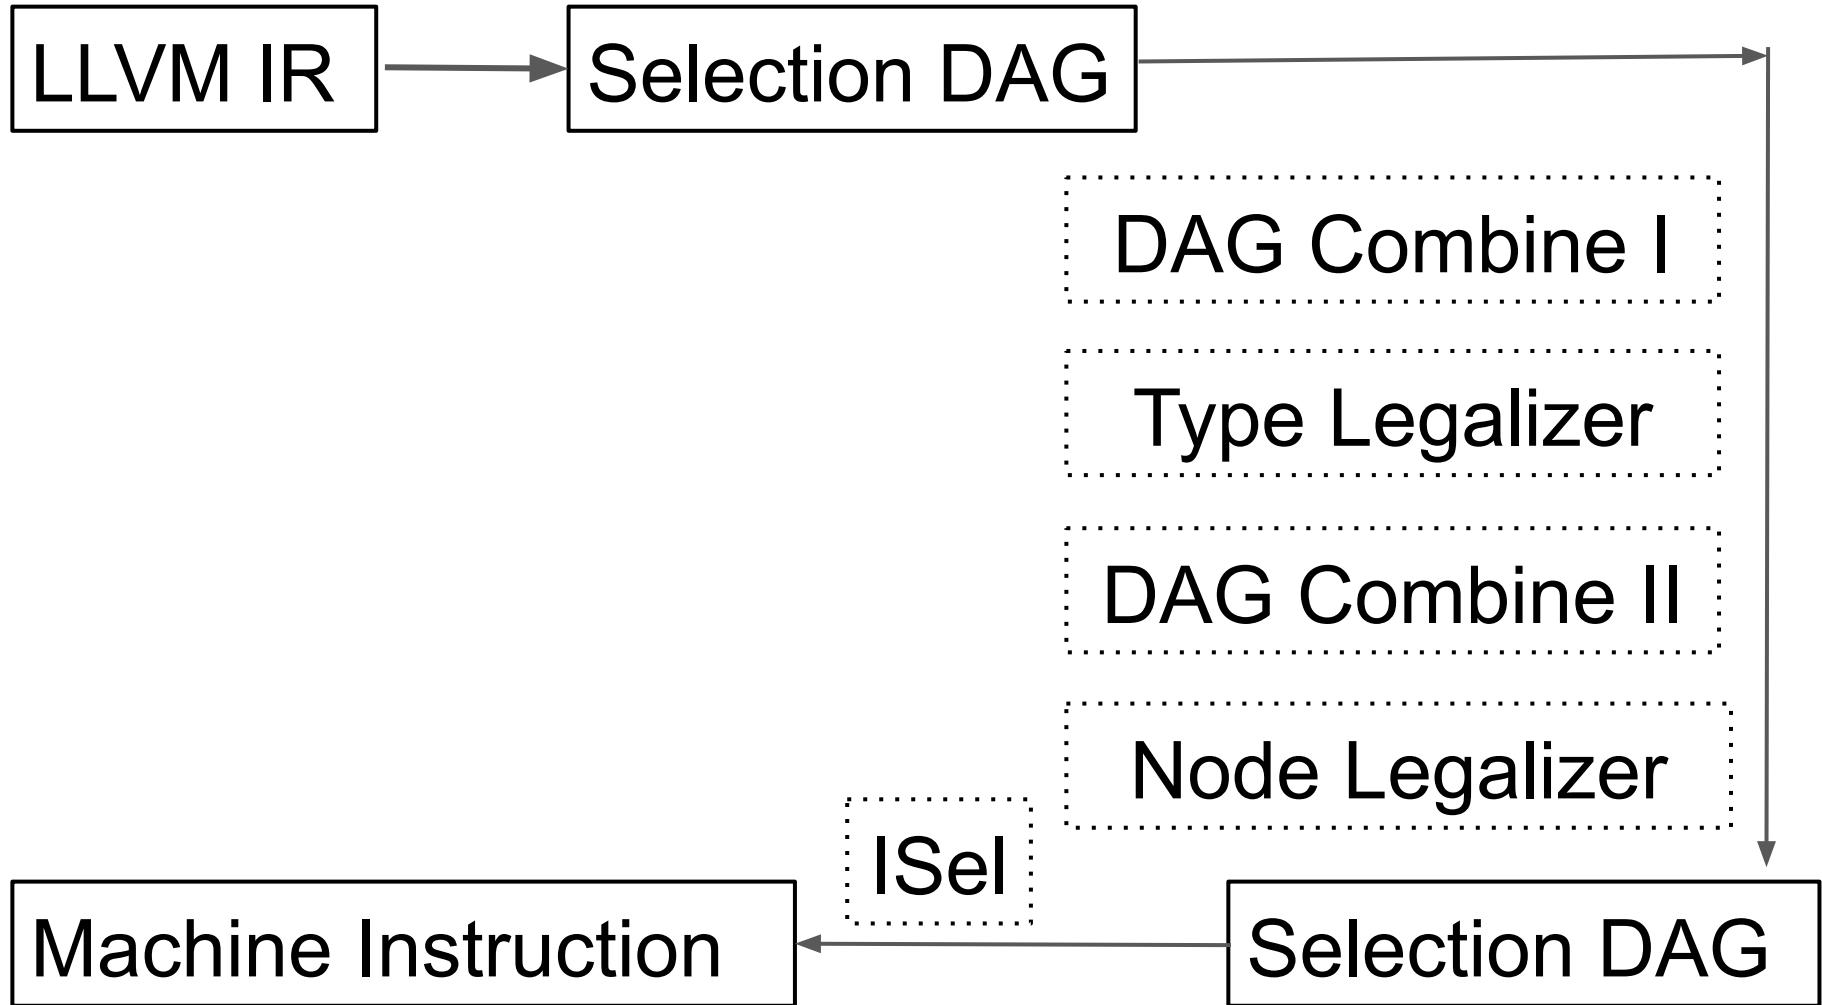




Image credits:

Dragon illustration: Vintage vector created by stockgiu - www.freepik.com

Smiley image: The logo belong to the awesome band nirvana

The global address node with 64 bit address fails the legalize.

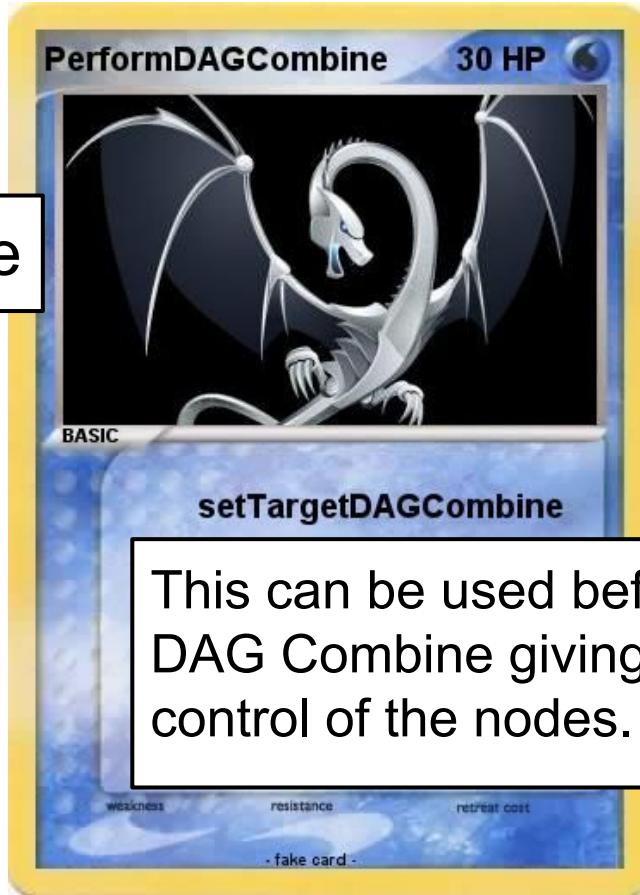


Some hacks

GlobalAddress node now passes the legalize but fails when it interacts with the store and load nodes.

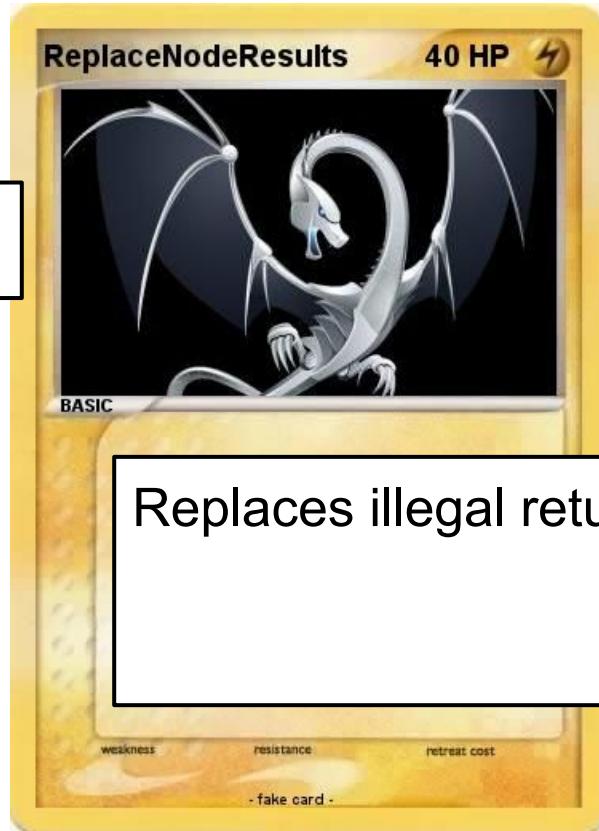
LLVM gives us many good ways to manipulate the nodes before ISel

PerformDAGCombine



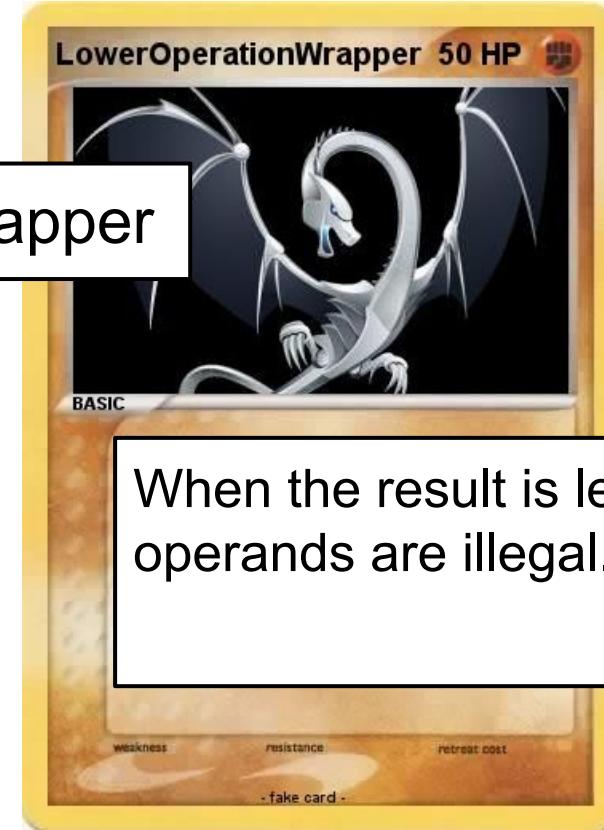
LLVM gives us many good ways to manipulate the nodes before ISel

ReplaceNodeResults



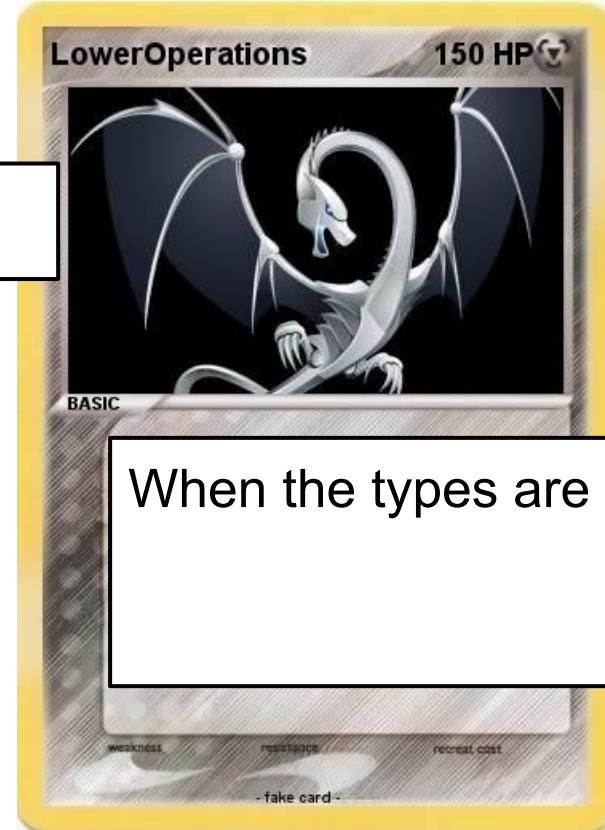
LLVM gives us many good ways to manipulate the nodes before ISel

LowerOperationWrapper



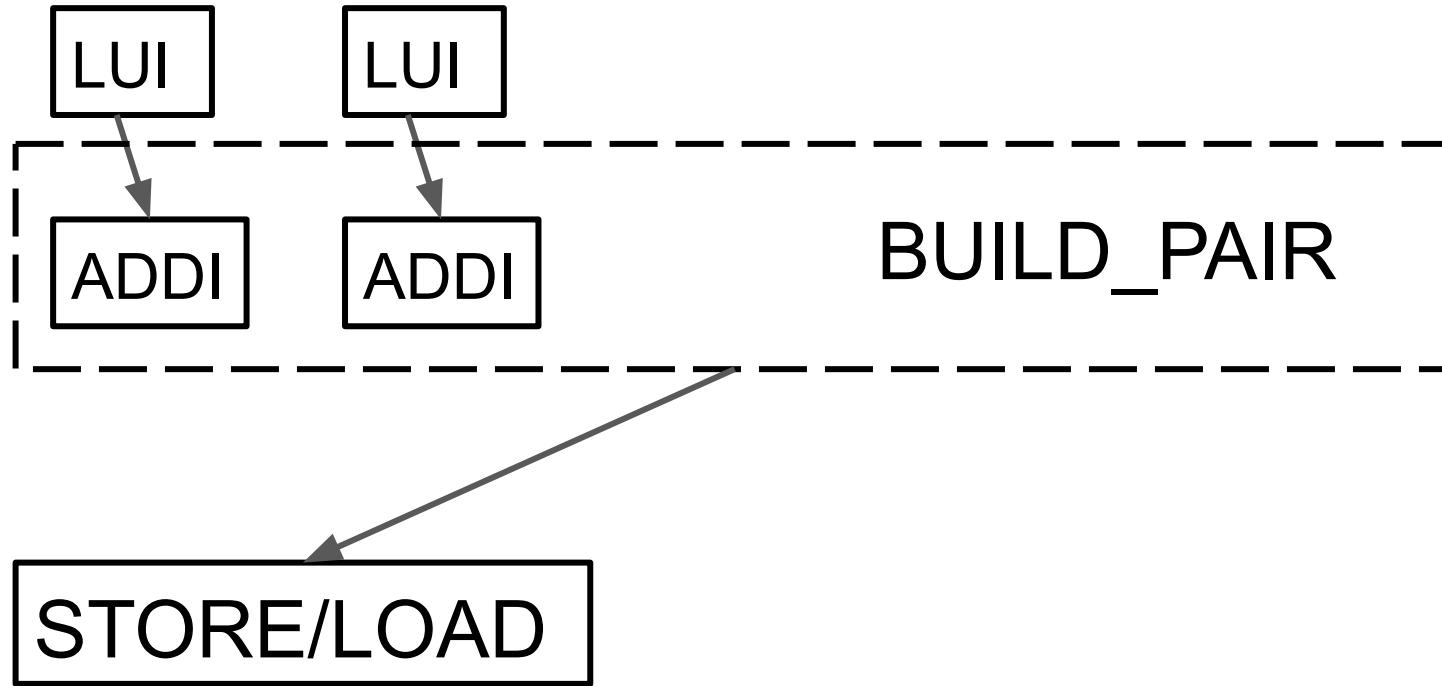
LLVM gives us many good ways to manipulate the nodes before ISel

LowerOperations



How we lowered global address?

GlobalAddress 0x64-bit-address



How we managed to handle store and load with GlobalAddress node?

The only option we had was to lower at the farthest point possible.

Load node is lowered before the Type Legalizer at DAG Combine 1

Store node is lowered at the Node Legalizer

Thank you!

2019

LLVM DEV
MEETING



Clang Interface Stubs

Syntax Directed Stub Library Generation

Puyan Lotfi
Facebook

Interface Stubs

```
echo "" | clang -shared -fPIC -x c - -o - | llvm-objdump -section-headers  
a.out: file format ELF64-x86-64
```

Sections:

Idx	Name	Size	Address	Type
0		00000000	0000000000000000	
1	.dynsym	00000108	0000000000001d0	
2	.dynstr	0000008f	0000000000002d8	
3	.symtab	00000498	0000000000000000	
4	.strtab	00000178	0000000000000000	
5	.shstrtab	000000cc	0000000000000000	
6	.gnu.hash	0000003c	000000000000190	
7	.gnu.version	00000016	000000000000368	
8	.gnu.version_r	00000020	000000000000380	
9	.rela.dyn	000000a8	0000000000003a0	
10	.init	00000017	000000000000448	TEXT
11	.plt	00000010	000000000000460	TEXT
12	.plt.got	00000008	000000000000470	TEXT
13	.text	000000c6	000000000000480	TEXT
14	.fini	00000009	000000000000548	TEXT
15	.eh_frame_hdr	00000024	000000000000554	DATA
16	.eh_frame	0000007c	000000000000578	DATA
17	.init_array	00000008	000000000200e40	
18	.fini_array	00000008	000000000200e48	
19	.dynamic	00000190	000000000200e50	
20	.got	00000020	000000000200fe0	DATA
21	.got.plt	00000018	000000000201000	DATA
22	.data	00000008	000000000201018	DATA
23	.bss	00000008	000000000201020	BSS
24	.comment	0000008f	0000000000000000	

Interface Stubs

```
echo "" | clang -shared -fPIC -x c - -o - | llvm-objdump -section-headers  
a.out: file format ELF64-x86-64
```

Sections:

Idx	Name	Size	Address	Type
0		00000000	0000000000000000	
1	.dynsym	00000108	0000000000001d0	
2	.dynstr	0000008f	0000000000002d8	
3	.symtab	00000498	0000000000000000	
4	.strtab	00000178	0000000000000000	
5	.shstrtab	000000cc	0000000000000000	
6	.gnu.hash	0000003c	000000000000190	
7	.gnu.version	00000016	000000000000368	
8	.gnu.version_r	00000020	000000000000380	
9	.rela.dyn	000000a8	0000000000003a0	
10	.init	00000017	000000000000448	TEXT
11	.plt	00000010	000000000000460	TEXT
12	.plt.got	00000008	000000000000470	TEXT
13	.text	000000c6	000000000000480	TEXT
14	.fini	00000009	000000000000548	TEXT
15	.eh_frame_hdr	00000024	000000000000554	DATA
16	.eh_frame	0000007c	000000000000578	DATA
17	.init_array	00000008	000000000200e40	
18	.fini_array	00000008	000000000200e48	
19	.dynamic	00000190	000000000200e50	
20	.got	00000020	000000000200fe0	DATA
21	.got.plt	00000018	000000000201000	DATA
22	.data	00000008	000000000201018	DATA
23	.bss	00000008	000000000201020	BSS
24	.comment	0000008f	0000000000000000	

Motivation

- Generation of lean SDKs
 - No Code
 - Explicit Symbol Exposure
- Code as Source of Truth: Syntax Directed
 - Make use of visibility attributes (ie `__attribute__(__visibility__("hidden"))`)

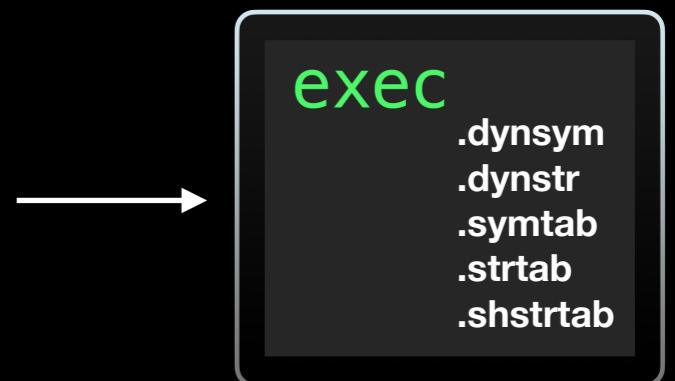
Motivation

- Generation of lean SDKs
 - No Code
 - Explicit Symbol Exposure
- Code as Source of Truth: Syntax Directed
 - Make use of visibility attributes (ie `__attribute__((__visibility__("hidden")))`)

clang -emit-interface-stubs -o libfoo.so a.cpp c.cpp sq.cpp

```
#define hidden __attribute__(( \
    __visibility__("hidden")))

hidden int b;
int red() { return b; }
hidden void green() { }
hidden void blue() { }
int a;
hidden int c;
```



Prior Art & Approaches

- Microsoft's Import Libraries
 - Generate stub code from compiler & linker
 - Syntax directed through `__declspec(dllexport/dllexport)`
- Apple's TAPI (Text API)
 - Header Scanning & stub generation
 - `.so/.dylib/.dll` scanning & stripping

Clang Interface Stubs

- Repurposes visibility attribute to direct symbol exposure using code syntax
 - Fine grain control (internal SPI vs external API)
- Aggregates exposure across compilation units via text
- Supports ELF
- Yields smaller SDK and faster link times

Clang Driver Pipeline

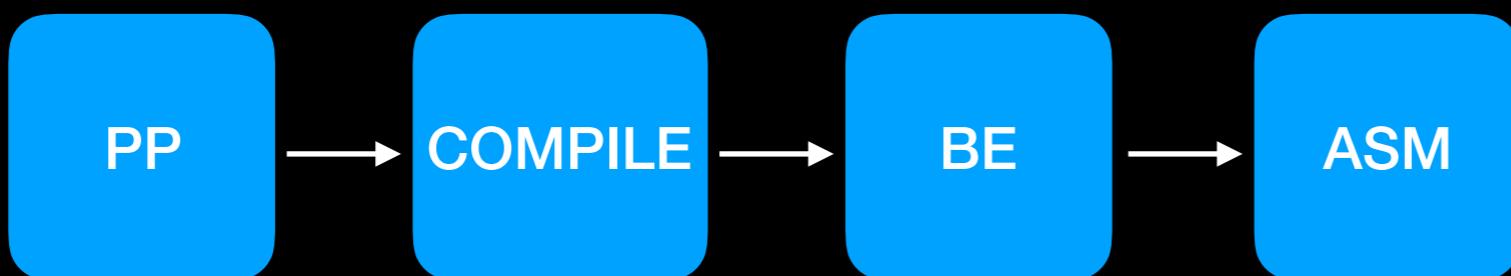
- Traditional Pipeline: Preprocess, Compile, Backend, Assemble, and Link Phases



Clang Driver Pipeline

- Traditional Pipeline: Preprocess, Compile, Backend, Assemble, and Link Phases

clang -c



Clang Driver Pipeline

- Traditional Pipeline: Preprocess, Compile, Backend, Assemble, and Link Phases

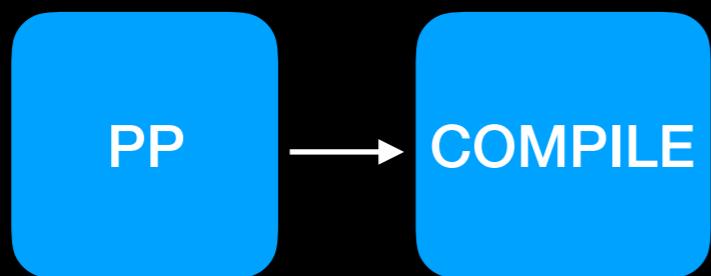
clang -S



Clang Driver Pipeline

- Traditional Pipeline: Preprocess, Compile, Backend, Assemble, and Link Phases

```
clang -fsyntax_only
```



Clang Driver Pipeline

- Traditional Pipeline: Preprocess, Compile, Backend, Assemble, and Link Phases

clang -E



Driver Pipeline

- Clang Interface Stubs Pipeline: Preprocess, Compile, and Merge phases
- Compile Phase: Generates intermediate text (.ifs files)
- Merge Phase: Invokes llvm-ifs to consume & merge .ifs files to produce ELF .so



Driver Pipeline

- Clang Interface Stubs Pipeline: Preprocess, Compile, and Merge phases
- Compile Phase: Generates intermediate text (.ifs files)
- Merge Phase: Invokes llvm-ifs to consume & merge .ifs files to produce ELF .so
- Compile Phase invokes InterfaceStubFunctionsConsumer (clang -cc1)
 - Walks the AST scanning for visible decls

```
#define weak \
    __attribute__((__weak__))
#define hidden __attribute__(( \
    __visibility__("hidden")))

hidden int b;
int red() { return b; }
weak void green() { }
hidden void blue() { }
int a;
hidden int c;
```

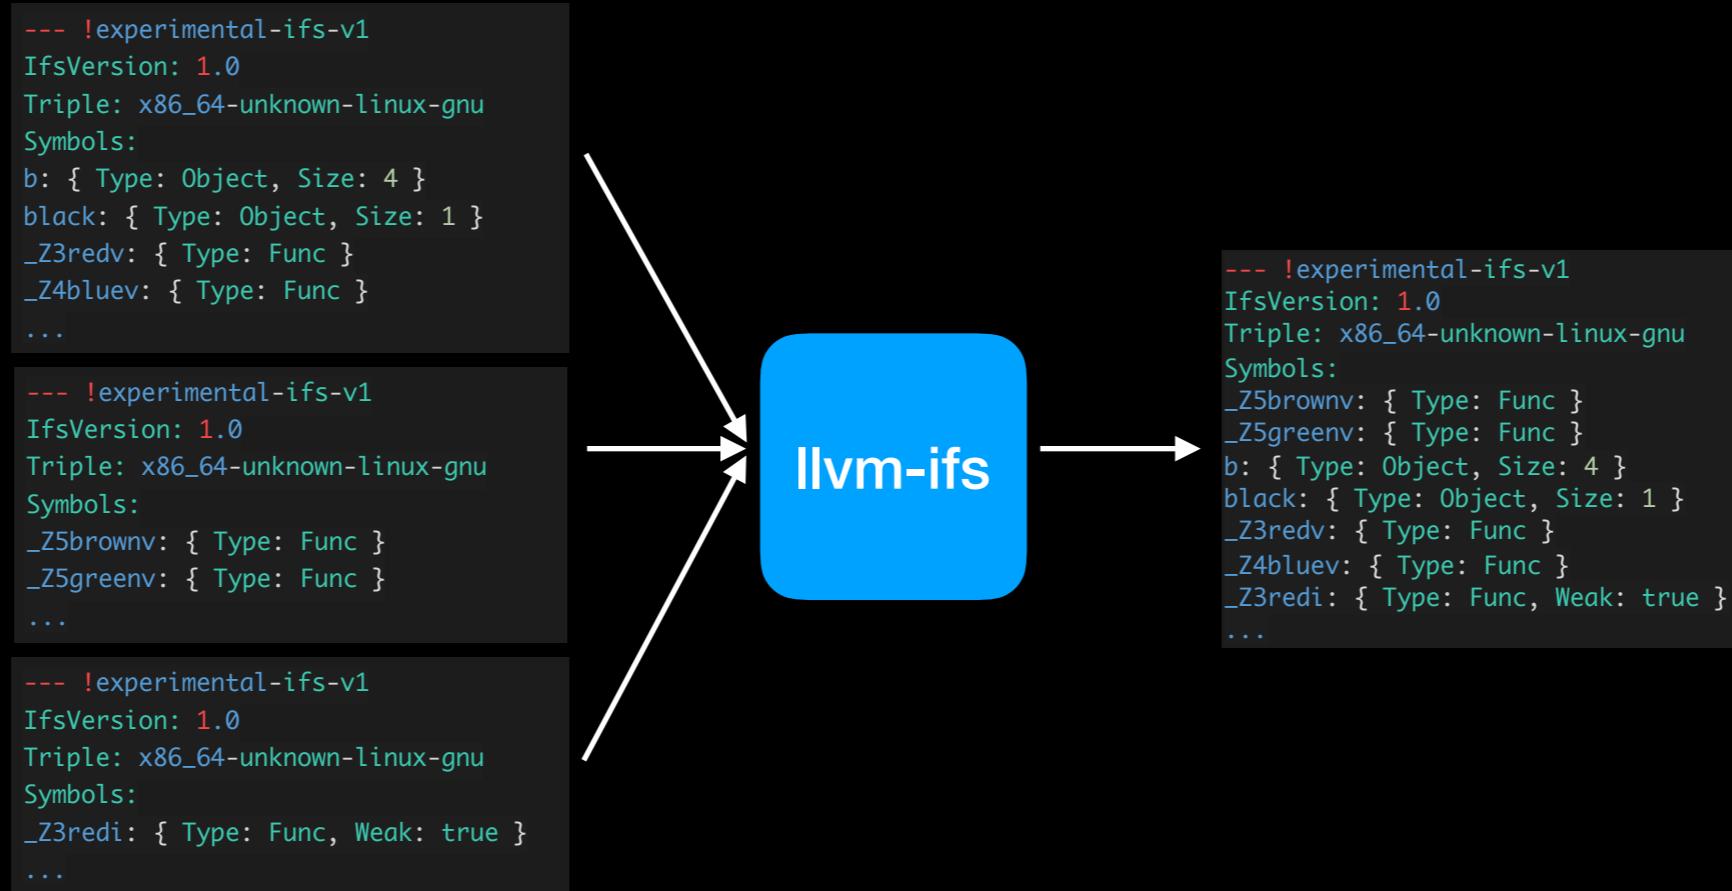


```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z3redv: { Type: Func }
_Z5greenv: { Type: Func,
             Weak: true }
a: { Type: Object, Size: 4 }
...
```

IFs Text Format

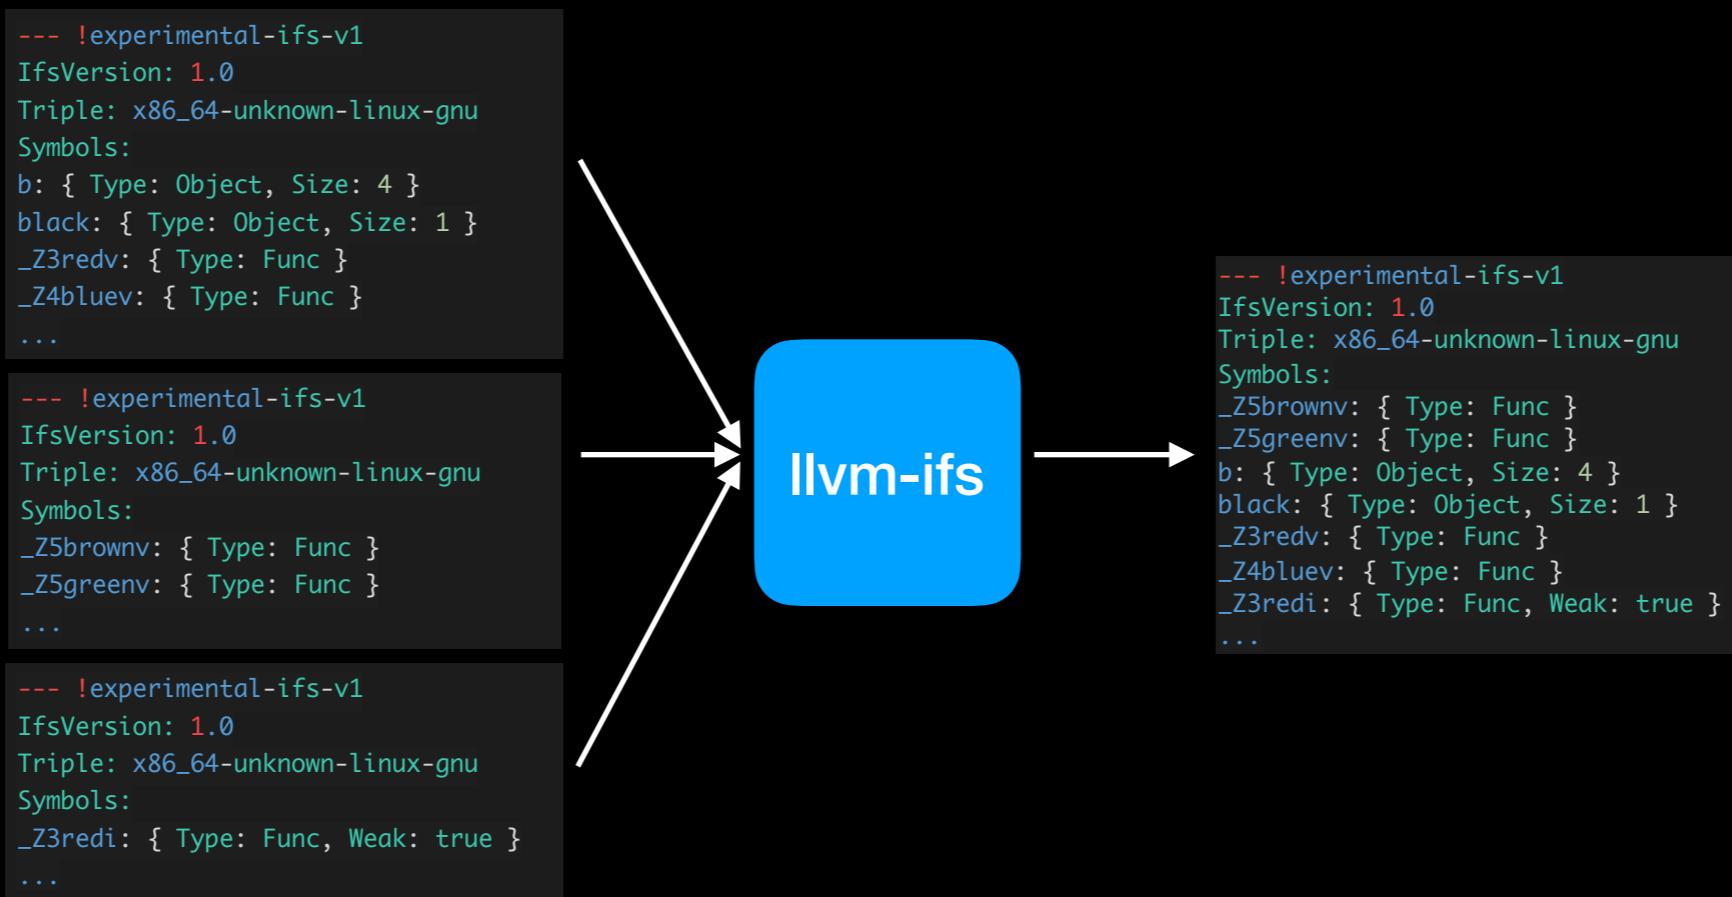
```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z5brownv: { Type: Func }
_Z5greenv: { Type: Func }
b: { Type: Object, Size: 4 }
black: { Type: Object, Size: 1 }
_Z3redv: { Type: Func }
_Z4bluev: { Type: Func }
_Z3redi: { Type: Func, Weak: true }
...
```

IFS Text Format



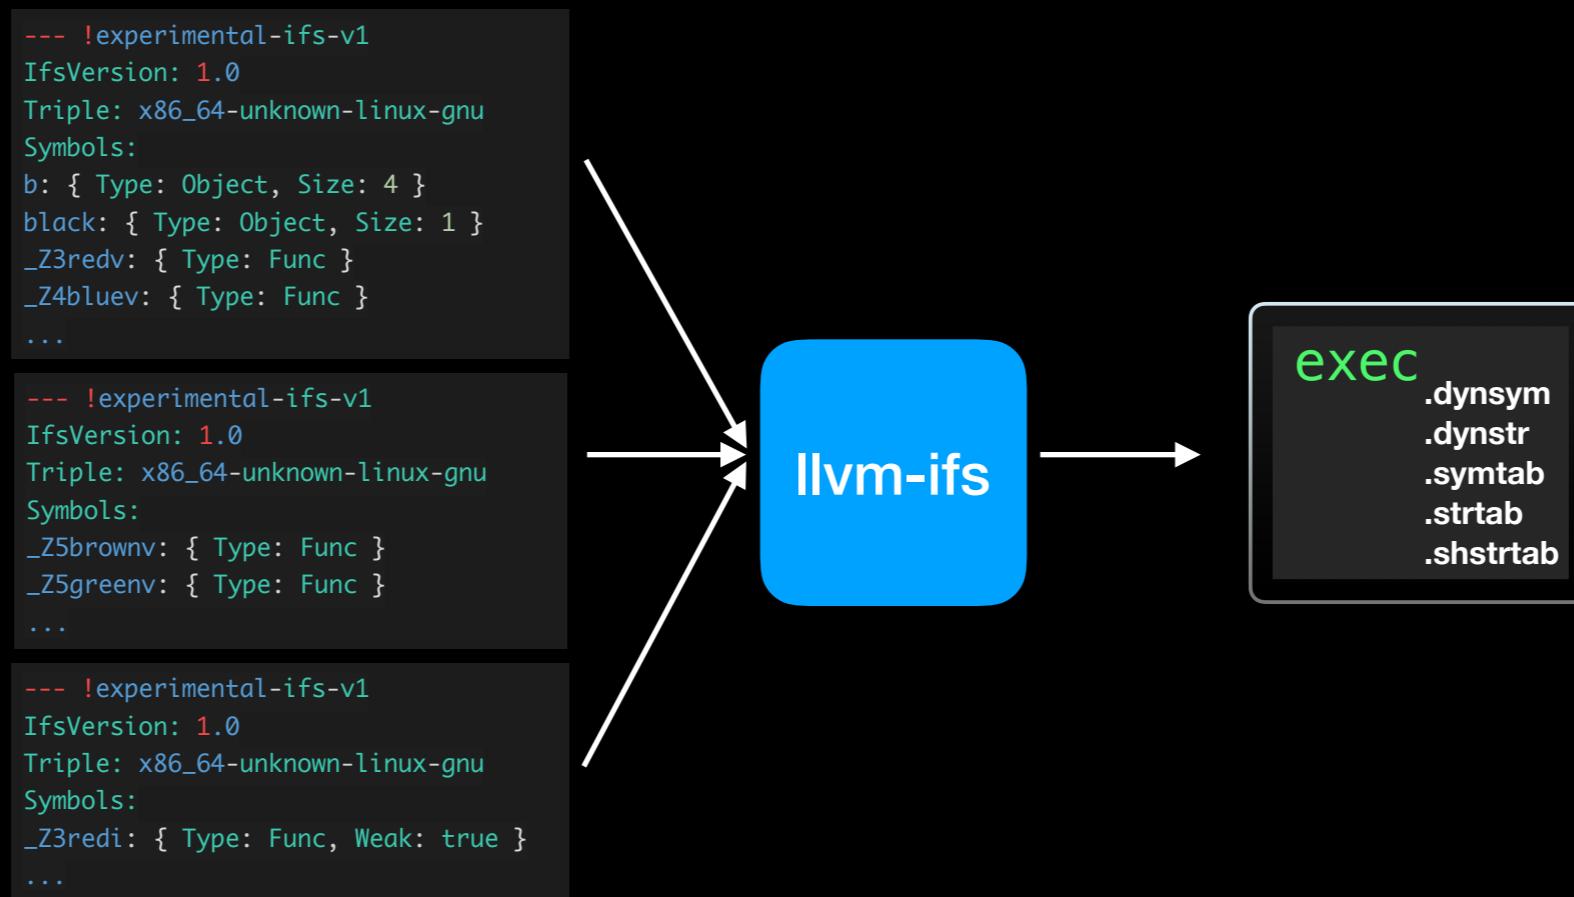
llvm-ifs

- A new llvm tool for consuming IFS files: produces a merged IFS file, or ELF shared object (.dynsym, .dynstr, .syntab only)



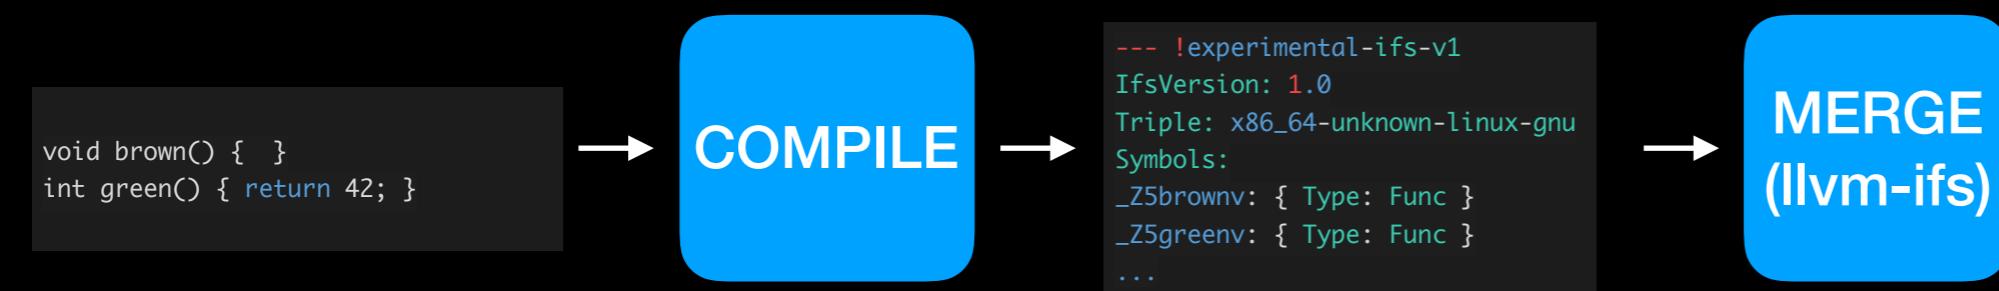
llvm-ifs

- A new llvm tool for consuming IFS files: produces a merged IFS file, or ELF shared object (.dynsym, .dynstr, .symtab only)



Example

```
clang -emit-interface-stubs -o libfoo.so a.cpp c.cpp sq.cpp
```



Example

`clang -emit-interface-stubs -o libfoo.so a.cpp c.cpp sq.cpp`

```
#define hidden __attribute__(( \
__visibility__("hidden")))

hidden int b;
hidden char black;
int orange() { return b; }
void blue() { }
hidden int c;
```

→ COMPILE →

```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z3orangev: { Type: Func }
_Z4bluev: { Type: Func }
...
```

```
void brown() { }
int green() { return 42; }
```

→ COMPILE →

```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z5brownv: { Type: Func }
_Z5greenv: { Type: Func }
...
```

```
__attribute__((__weak__))
int red(int num) {
    return num * num;
}
```

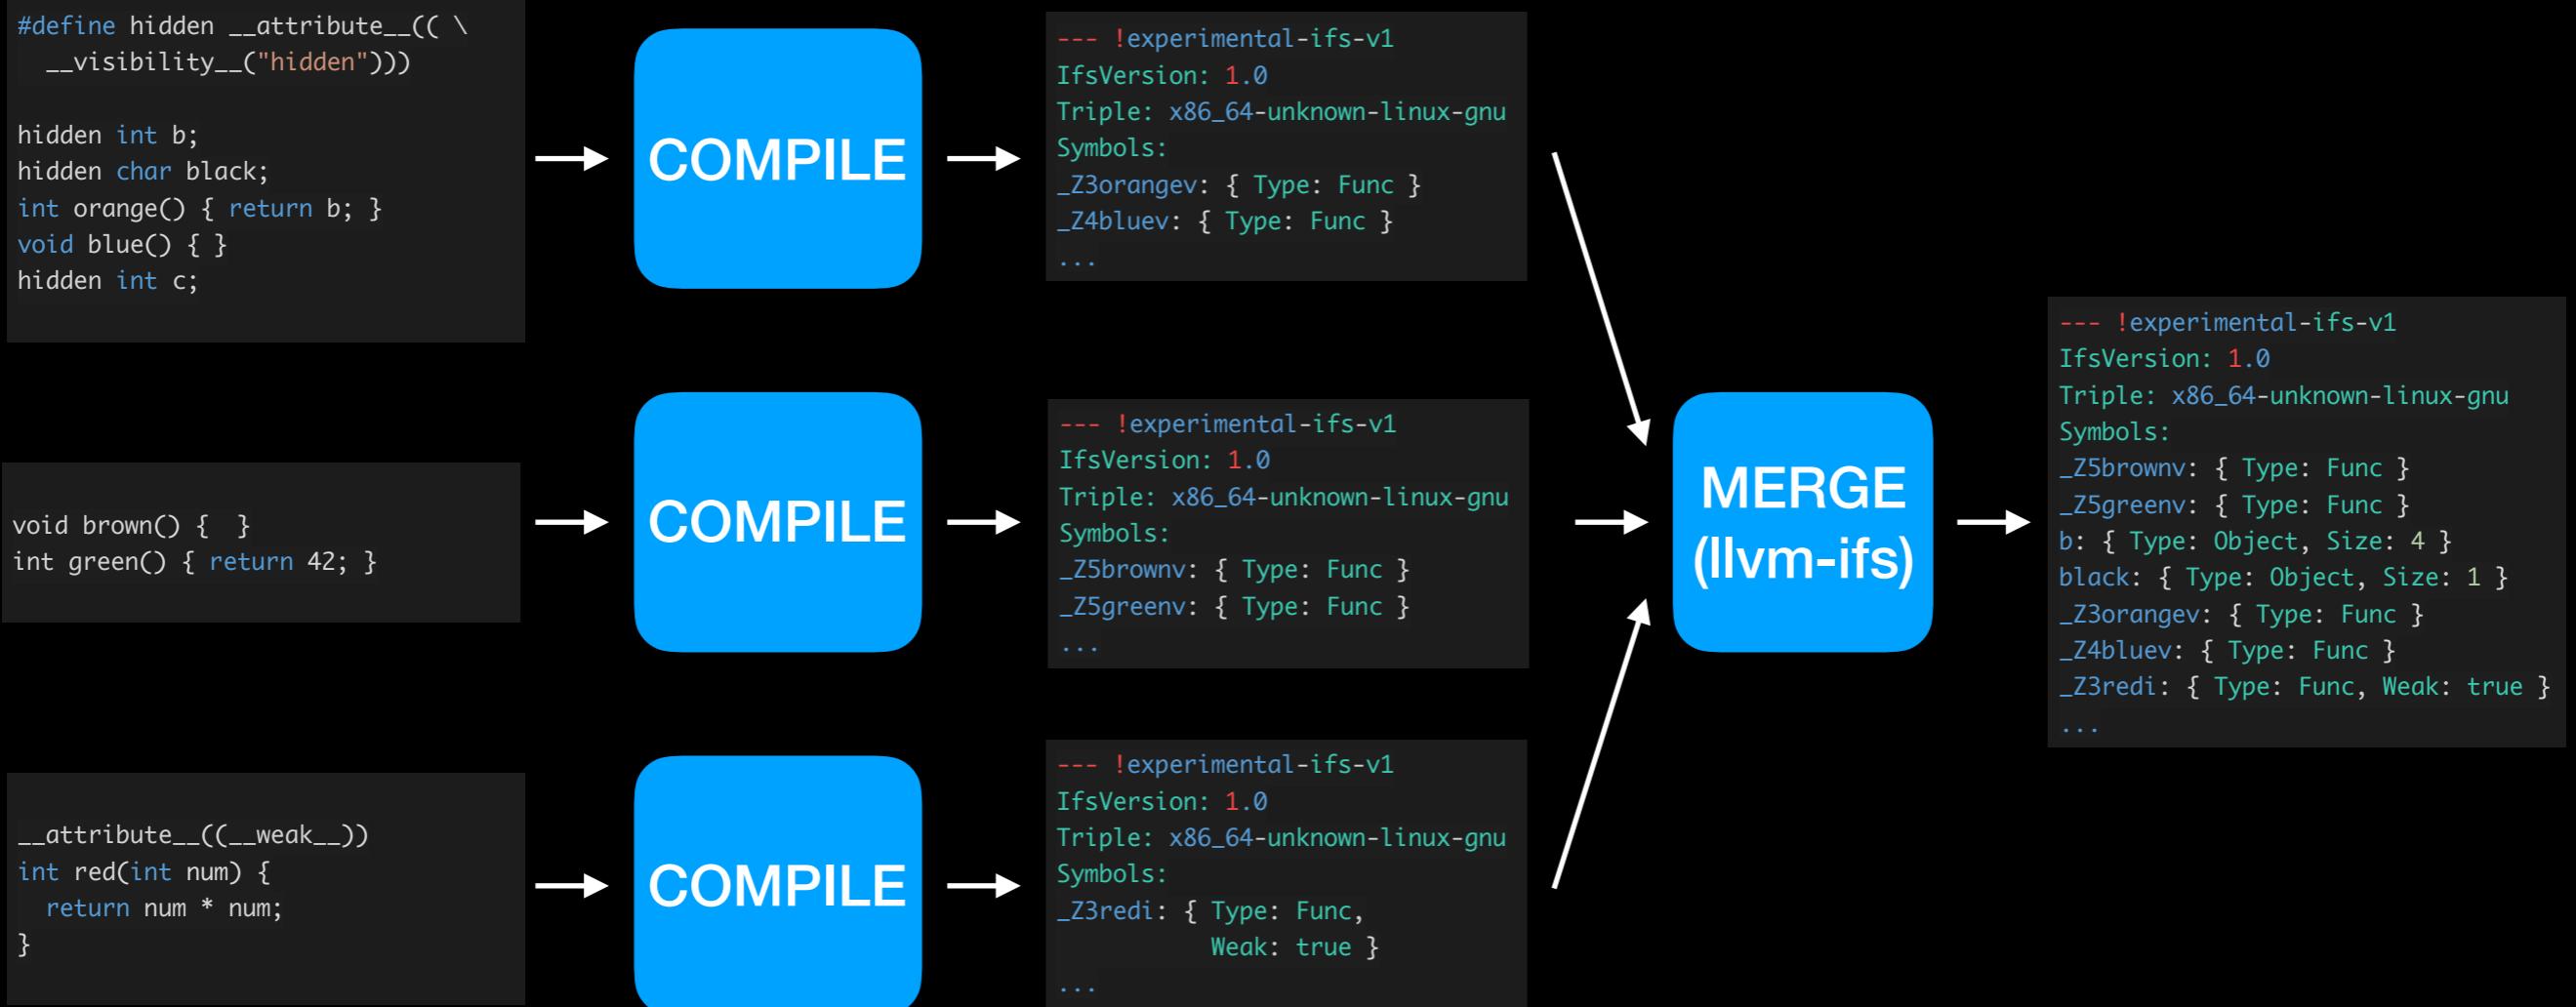
→ COMPILE →

```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z3redi: { Type: Func,
            Weak: true }
...
```

MERGE
(llvm-ifs)

Example

`clang -emit-interface-stubs -o libfoo.so a.cpp c.cpp sq.cpp`



Example

`clang -emit-interface-stubs -o libfoo.so a.cpp c.cpp sq.cpp`

```
#define hidden __attribute__(( \
__visibility__("hidden")))

hidden int b;
hidden char black;
int orange() { return b; }
void blue() { }
hidden int c;
```

→ COMPILE →

```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z3orangev: { Type: Func }
_Z4bluev: { Type: Func }
...
```

```
void brown() { }
int green() { return 42; }
```

→ COMPILE →

```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z5brownv: { Type: Func }
_Z5greenv: { Type: Func }
...
```

```
__attribute__((__weak__))
int red(int num) {
    return num * num;
}
```

→ COMPILE →

```
--- !experimental-ifs-v1
IfsVersion: 1.0
Triple: x86_64-unknown-linux-gnu
Symbols:
_Z3redi: { Type: Func,
            Weak: true }
...
```

MERGE
(llvm-ifs)

exec
.dynsym
.dynstr
.syntab
.strtab
.shstrtab

Challenges

- Refactoring for alternate pipeline setup required changes in `Driver::BuildActions` and `getCompilationPhases`
- Handling corner cases in the driver setup
- Handling Templates, Specializations, and non-trivial decls
- Converging on the text format required many iterations

Future Work

- Hardening IFS ASTConsumer and llvm-ifs
- Generating interface stubs for libc++ builds
- Inline Assembly
- Support for new formats: MS Import Libs, Darwin TAPI

2019

LLVM DEV
MEETING





FLANG UPDATE

Steve Scalpone, LLVM Developers' Meeting, October 23, 2019

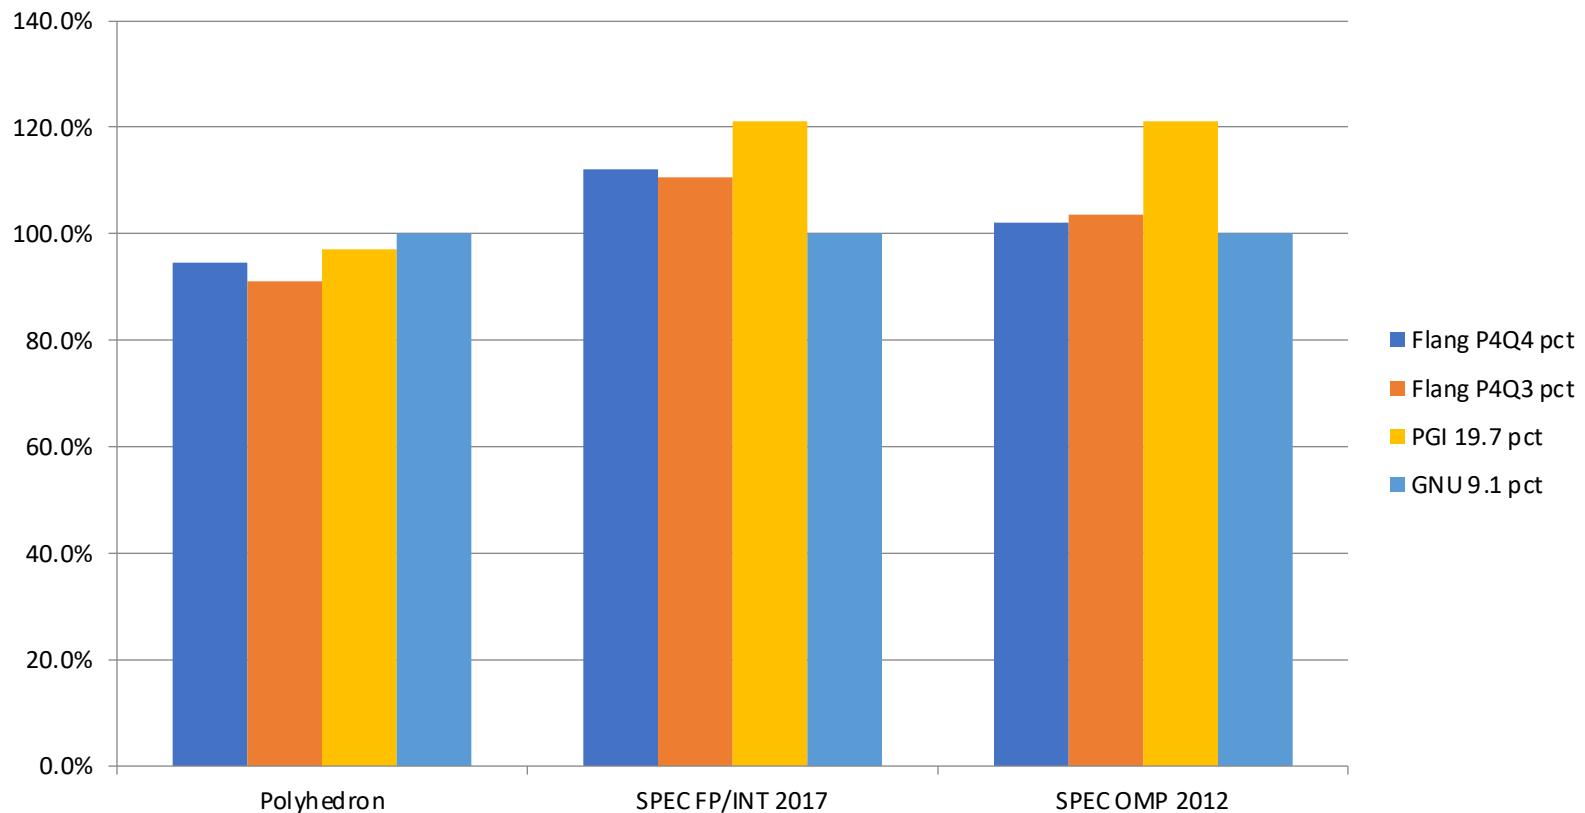
THE FLANG PROJECT

Now affectionately known as "Old Flang"

- ▶ A multi-year collaboration between the NNSA labs and NVIDIA/PGI
 - ▶ Open-source Fortran compiler since May 2017
 - ▶ Arm and AMD are shipping compilers built using Flang
 - ▶ Wide variety of contributors; most recently Microsoft signed the CCLA
- ▶ The end of the CLA
 - ▶ No longer require a CCLA or CLA to contribute
 - ▶ Adding Arm and other community members as committers
 - ▶ Relicensing from Apache 2.0 to Apache 2.0 with LLVM Exception

FLANG PERFORMANCE

All runs on dual-socket Intel Xeon Skylake



Performance measured September 2019 and are considered ESTIMATES per SPEC run and reporting rules.
Two 20 core Skylake Intel® Gold 6148 CPU @ 2.40GHz CPUs @ 2.4GHz w/ 256GB memory. SPEC® is a registered trademark
of the Standard Performance Evaluation Corporation (www.spec.org).

THE FLANG PROJECT

Some call it “f18” or “new flang”

New Flang compiler was started from scratch

Modern & well-documented design written in C++17

NNSA, ECP, NVIDIA, Arm and others have contributed

It's an LLVM subproject despite not yet having files in the repo

Developed 100% in open source guided by LLVM coding standards

THE FLANG PROJECT

Using some interesting technologies

- ▶ FIR implementation is as an MLIR dialect
 - ▶ MLIR replaced the original f18/FIR infrastructure
 - ▶ Working with the MLIR team to extend MLIR as needed for FIR
- ▶ OpenMP IRBuilder drives language-independent lowering
 - ▶ Defining an OpenMP dialect in MLIR
 - ▶ Optz and transformations may be in OpenMP IRBuilder, MLIR framework, or LLVM

THE FLANG PROJECT

SOURCE CODE

<https://github.com/flang-compiler>

MAILING LIST

flang-dev@lists.llvm.org

PROJECT CALL

Every other Wednesday 8:30am Pacific
See the mailing list for details

TECHNICAL CALL

Every other Monday 8:30 Pacific
See the mailing list for details

2019

LLVM DEV MEETING



Dead Virtual Function Elimination (VFE) in LLVM



An example we'd like to optimise

```
struct A {  
    virtual int foo() { return 1; }  
    virtual int bar() { return 2; }  
};  
  
struct B : A {  
    virtual int foo() { return 3; }  
    virtual int bar() { return 4; }  
};  
  
A* make_A() { return new A(); }  
B* make_B() { return new B(); }  
  
int call_1(A* p) { return p->foo(); }  
int call_2(B* p) { return p->bar(); }
```

An example we'd like to optimise

```
struct A {  
    virtual int foo() { return 1; }  
    virtual int bar() { return 2; }  
};
```

```
struct B : A {  
    virtual int foo() { return 3; }  
    virtual int bar() { return 4; }  
};
```

```
A* make_A() { return new A(); }  
B* make_B() { return new B(); }
```

```
int call_1(A* p) { return p->foo(); }  
int call_2(B* p) { return p->bar(); }
```

Could call A::foo or B::foo

Can only call B::bar

An example we'd like to optimise

```
struct A {  
    virtual int foo() { return 1; }  
    virtual int bar() { return 2; }  
};
```

No calls to this function

```
struct B : A {  
    virtual int foo() { return 3; }  
    virtual int bar() { return 4; }  
};
```

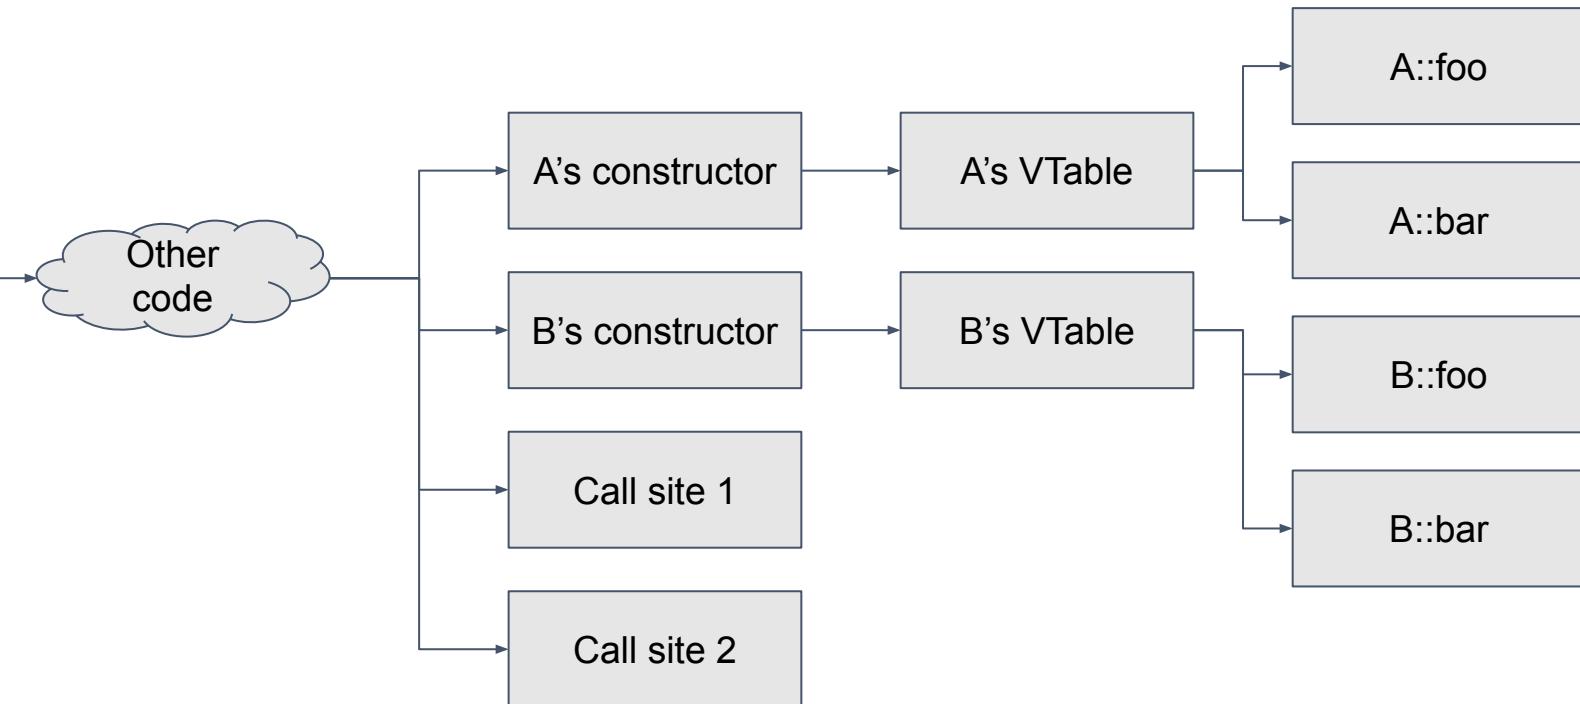
Could call A::foo or B::foo

```
A* make_A() { return new A(); }  
B* make_B() { return new B(); }
```

```
int call_1(A* p) { return p->foo(); }  
int call_2(B* p) { return p->bar(); }
```

Can only call B::bar

What does this looks like to GlobalDCE?



Type metadata

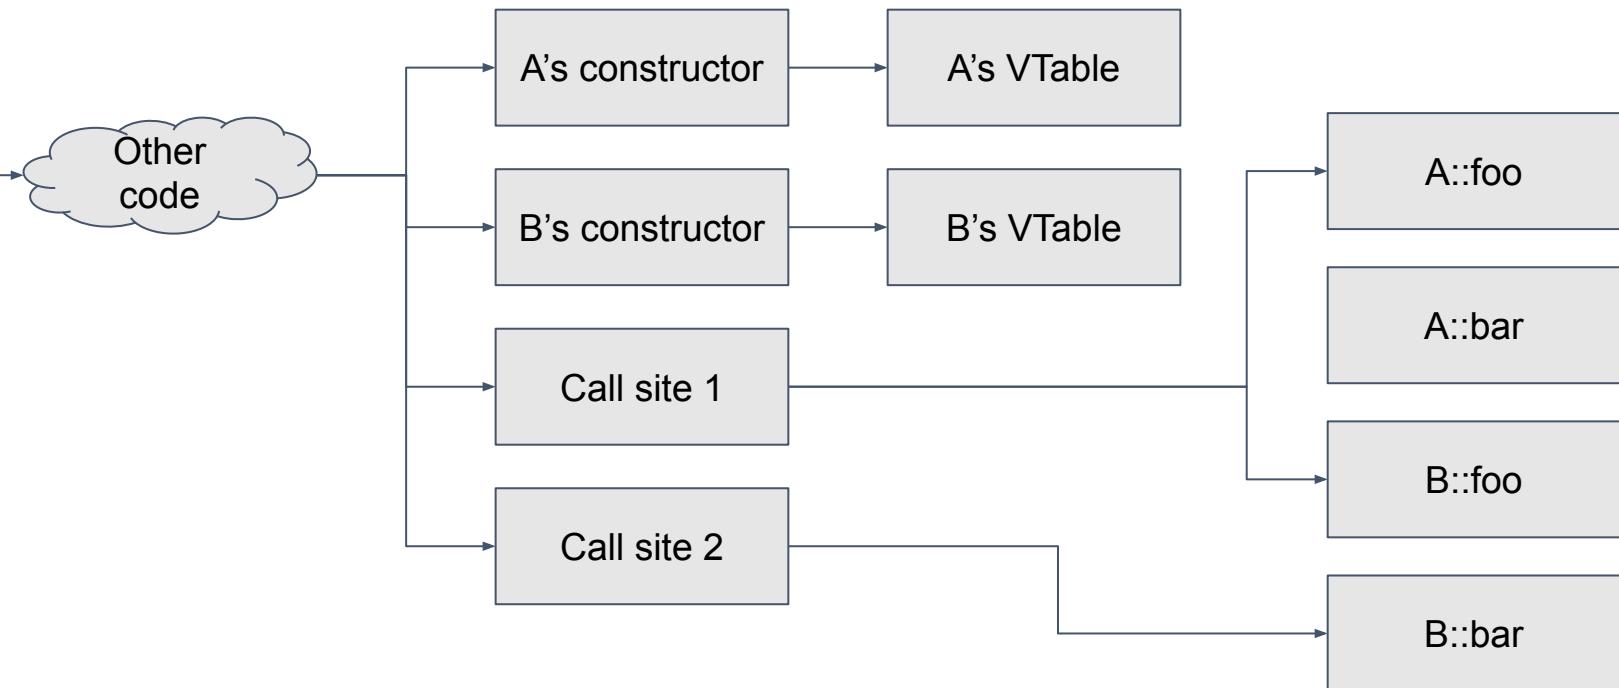
- Optionally added to vtables, virtual call sites
- Provides link between call sites and vtable slots
- Currently used for devirtualisation, control flow integrity
- Needs some changes to call sites to allow VFE

```
@_ZTV1A = constant { [4 x i8*] } ..., !type !0
@_ZTV1B = constant { [4 x i8*] } ..., !type !0, !type !4

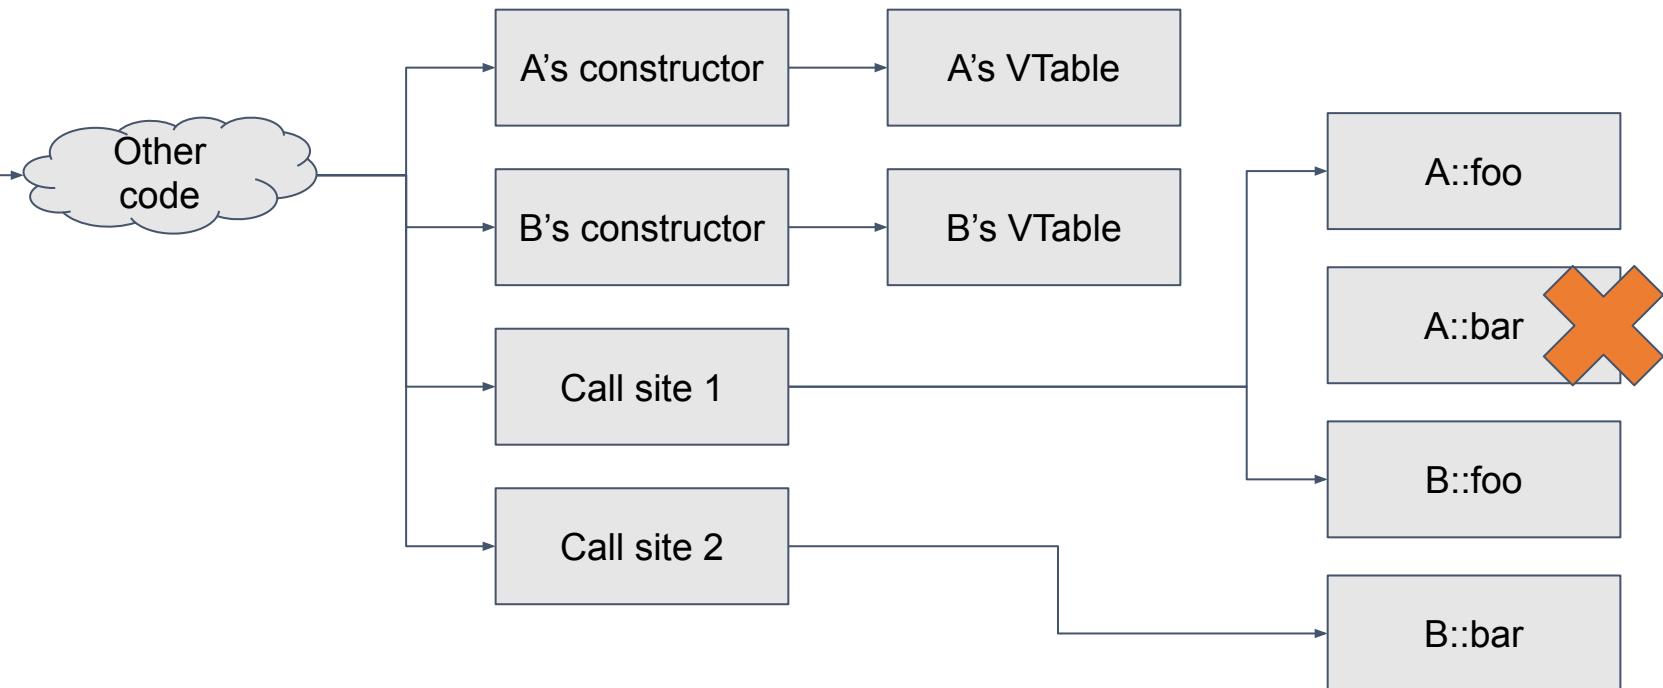
!0 = !{i64 16, !"__ZTS1A"}
!4 = !{i64 16, !"__ZTS1B"}

%vtable1 = load i8*, i8** %0, align 8, !tbaa !9
%1 = tail call { i8*, i1 } @llvm.type.checked.load(
                        i8* %vtable1, i32 8, metadata !"__ZTS1B")
```

How does this change GlobalDCE's view?



How does this change GlobalDCE's view?

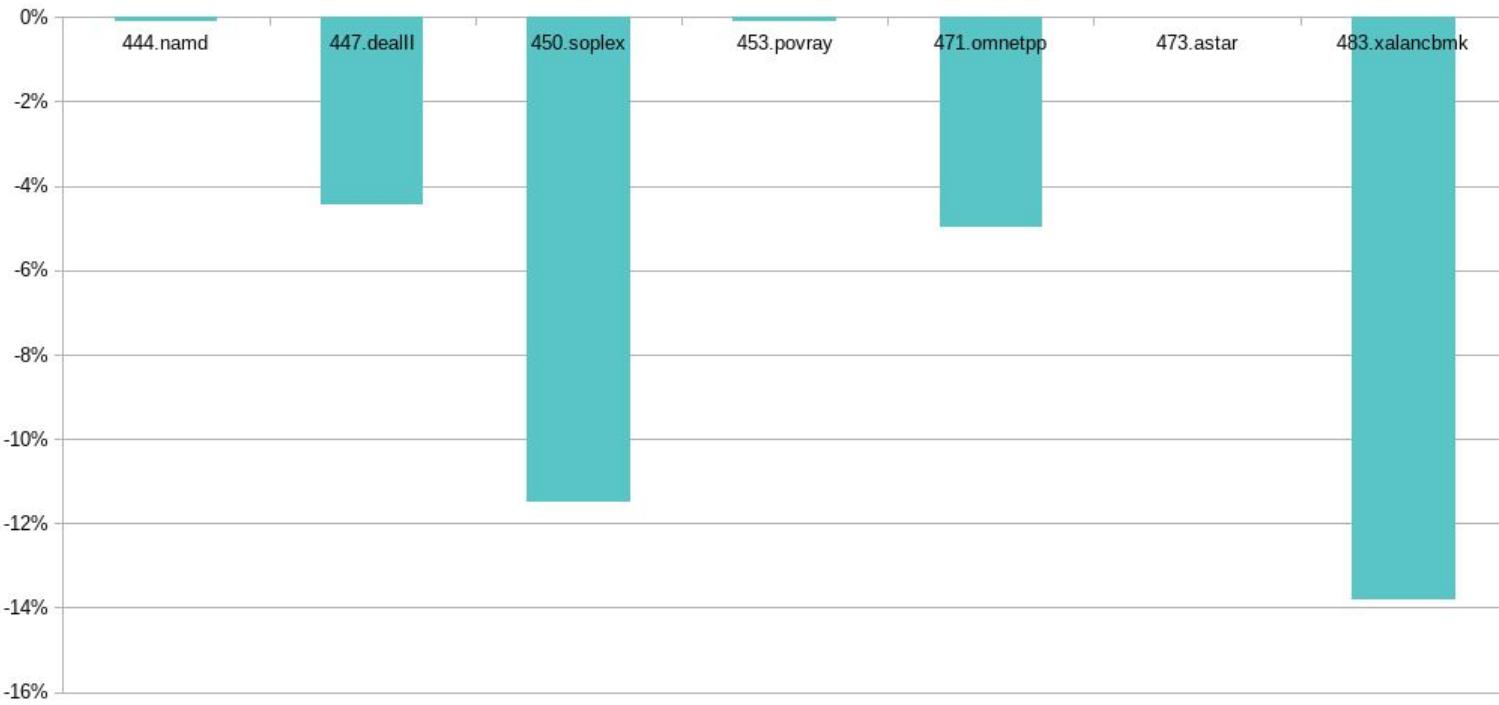


When is this optimisation valid?

- Need to see every possible call site involving a class
- Generally requires `-fvisibility=hidden` and LTO
- All base classes (with some exceptions) must also be hidden
 - Otherwise, more virtual calls could be outside the LTO unit
- Added new `!vcall_visibility` metadata to represent this

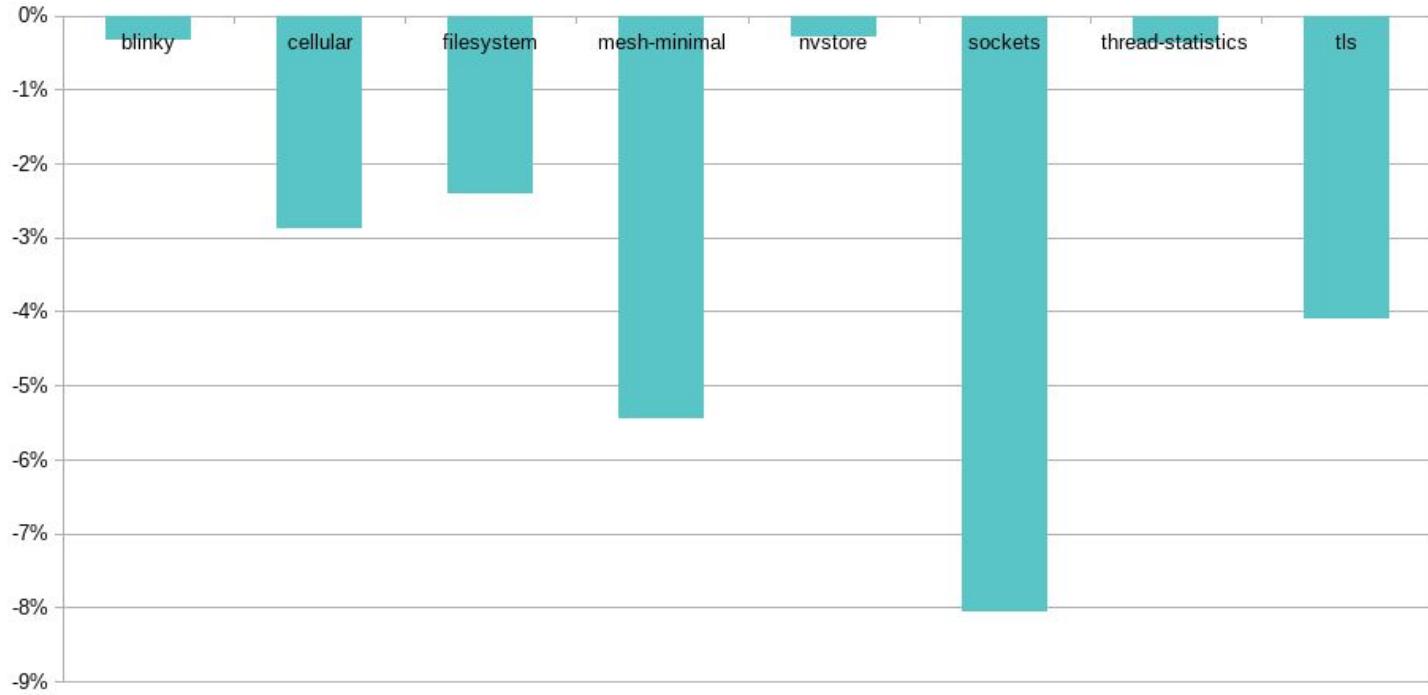
Benchmark - SPEC2006 (C++ subset)

Code size change, compared to -Oz -flto -fvisibility=hidden



Benchmark - mbed-os examples

Code size change, compared to -Oz -flto -fvisibility=hidden



Future work

- ThinLTO
- Dead RTTI elimination

2019

LLVM DEV
MEETING



Making a Language Cross Platform

Libraries and Tooling

Gwen Mittertreiner

Facebook

Goals

- Make you think about cross platform support
- High level, but still have practical advice.

Paths and File Systems

Paths and File Systems

- On Unix, paths are easy to parse right? /usr/bin/easy/peasy

Paths and File Systems

- On Unix, paths are easy to parse right? /usr/bin/easy/peasy
- On Windows, paths are easy

Paths and File Systems

- On Unix, paths are easy to parse right? /usr/bin/easy/peasy
- On Windows, paths are easy to get wrong.

Paths and File Systems

- On Unix, paths are easy to parse right? /usr/bin/easy/peasy
- On Windows, paths are easy to get wrong.
- Which of these are valid paths on Windows?
 - A:\foo\bar
 - B:foo\bar
 - C:/foo/bar
 - D:/foo\bar
 - \\?\E:\foo\bar
 - \\.\F:\foo\bar
 - \??\G:\foo\bar
 - \usr\bin\clang
 - \\?\UNC\abc\xyz
 - \\?\GLOBALROOT\??\UNC\abc\xyz

Paths and File Systems

- On Unix, paths are easy to parse right? /usr/bin/easy/peasy
- On Windows, paths are easy to get wrong.
- Which of these are valid paths on Windows?
 - A:\foo\bar
 - B:foo\bar
 - C:/foo/bar
 - D:/foo\bar
 - \\?\E:\foo\bar
 - \\.\F:\foo\bar
 - \??\G:\foo\bar
 - \usr\bin\clang
 - \\?\UNC\abc\xyz
 - \\?\GLOBALROOT\??\UNC\abc\xyz
- All of them!

Paths and File Systems

- On Unix, paths are easy to parse right? /usr/bin/easy/peasy
- On Windows, paths are easy to get wrong.
- Which of these are valid paths on Windows?
 - A:\foo\bar
 - B:foo\bar
 - C:/foo/bar
 - D:/foo\bar
 - \\?\E:\foo\bar
 - \\.\F:\foo\bar
 - ??\G:\foo\bar
 - \usr\bin\clang
 - \\?\UNC\abc\xyz
 - \\?\GLOBALROOT\??\UNC\abc\xyz
- All of them! * (mostly)

*<https://googleprojectzero.blogspot.com/2016/02/the-definitive-guide-on-win32-to-nt.html>

Paths and File Systems

Paths and File Systems

- Let system APIs do the lifting

Paths and File Systems

- Let system APIs do the lifting
- Avoid matching paths by `streq`

Building and Testing

Building and Testing

- Keep cross compiling in mind through out

Building and Testing

- Keep cross compiling in mind through out
- Be very careful with Unix commands

Building and Testing

- Keep cross compiling in mind through out
- Be very careful with Unix commands
- Linker formats

Building and Testing

- Keep cross compiling in mind through out
- Be very careful with Unix commands
- Linker formats
 - Different formats support different things
 - Weak Symbols
 - Symbol Replacement
 - Thread Local Storage
 - DLL Storage
 - Two level namespaces

Porting Existing Libraries

Porting Existing Libraries

- UTF-8 is pretty standard now...

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(
- POSIX apis are a handy tool

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(
- POSIX apis are a handy tool
 - Windows supports many POSIX APIs

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(
- POSIX apis are a handy tool
 - Windows supports many POSIX APIs
 - With some limitations

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(
- POSIX apis are a handy tool
 - Windows supports many POSIX APIs
 - With some limitations
 - Android doesn't support some POSIX APIs

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(
- POSIX apis are a handy tool
 - Windows supports many POSIX APIs
 - With some limitations
 - Android doesn't support some POSIX APIs
 - Depends on API level

Porting Existing Libraries

- UTF-8 is pretty standard now...
 - Except Windows and Apple's Cocoa any sufficiently old platform! Those use UTF-16.
- Two Approaches:
 - Define a character type that's different sizes on different platforms
 - Convert everything to a single encoding
- Be okay with things not making sense :(
- POSIX apis are a handy tool
 - Windows supports many POSIX APIs
 - With some limitations
 - Android doesn't support some POSIX APIs
 - Depends on API level
 - Reimplementing the original target's API set

Thanks!

- I hope you'll all go and make someone porting your project happy.

2019

LLVM DEV
MEETING





Grafter: A Clang tool for fusing tree traversals

Laith Sakka

Kirshanthan Sundararajah

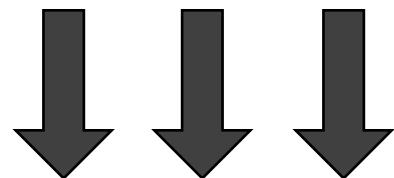
Milind Kulkarni

What is Grafter?



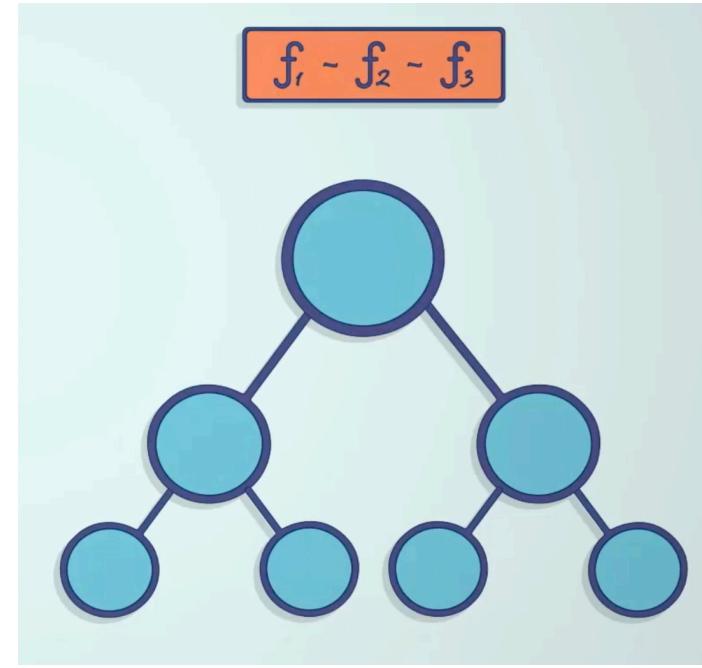
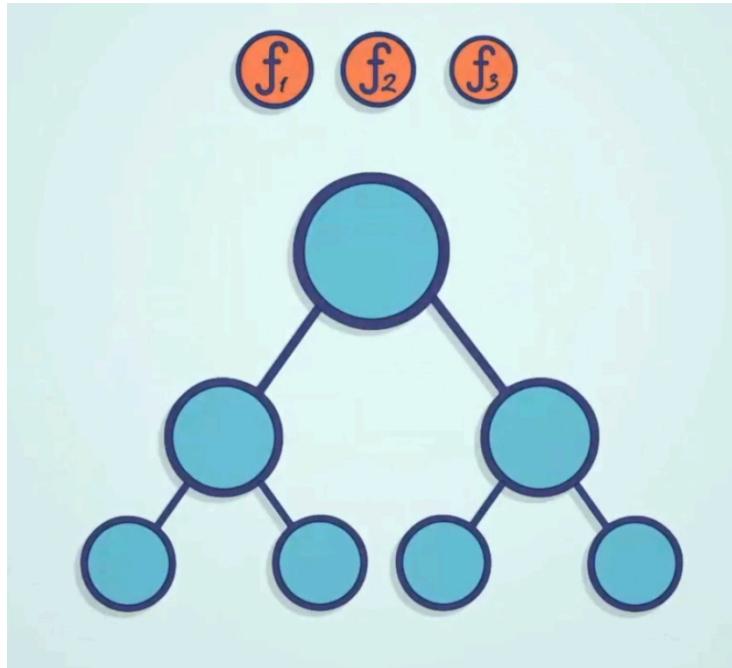
Grafter performs source to source transformations to fuse recursive functions that traverse trees

```
computeHeight(render_tree);  
computeWidth(render_tree);  
computePos(render_tree);
```



```
computeHeight_computeWidth_computePos(render_tree);
```

What is Grafter?



Sound, Fine-Grained Traversal Fusion for Heterogeneous Trees (PLDI 2019)

<https://dl.acm.org/citation.cfm?id=3314221.3314626>

<https://www.youtube.com/watch?v=j2henSFtZds>

<https://github.com/laithsakka/Grafter>

This talk is not about Grafter!
Its about utilizing Clang to implement Grafter

1. Embedded DSL in C++
2. Static analysis
3. Code generation

Source to source

1. Easier to understand tree/graph programs on the source level.
2. Analyze dependences in tree/graph programs
3. Define high level transformations as source level rewrites

Embedded DSL in C++

- Annotate components
- Verify annotated components against a set of rules.

```
Annotations class __tree_structure__ ValueNode : public Node {  
public:  
    int Value;  
    __tree_child__ Node *Left, *Right;  
  
    __tree_traversal__ void search(int Key, bool ValidCall) override {  
        Found = false;  
        if (Key == Value) {  
            Found = true;  
            return;  
        }  
        Left->search(Key, Key < Value);  
        Right->search(Key, Key >= Value);  
        Found = Left->Found || Right->Found;  
    }  
}
```

No aliasing

this->Left = this->Right;

this->Left = new ValueNode();

No condition calls

if(...)
| Left->insert(...);

Left->insert(...);

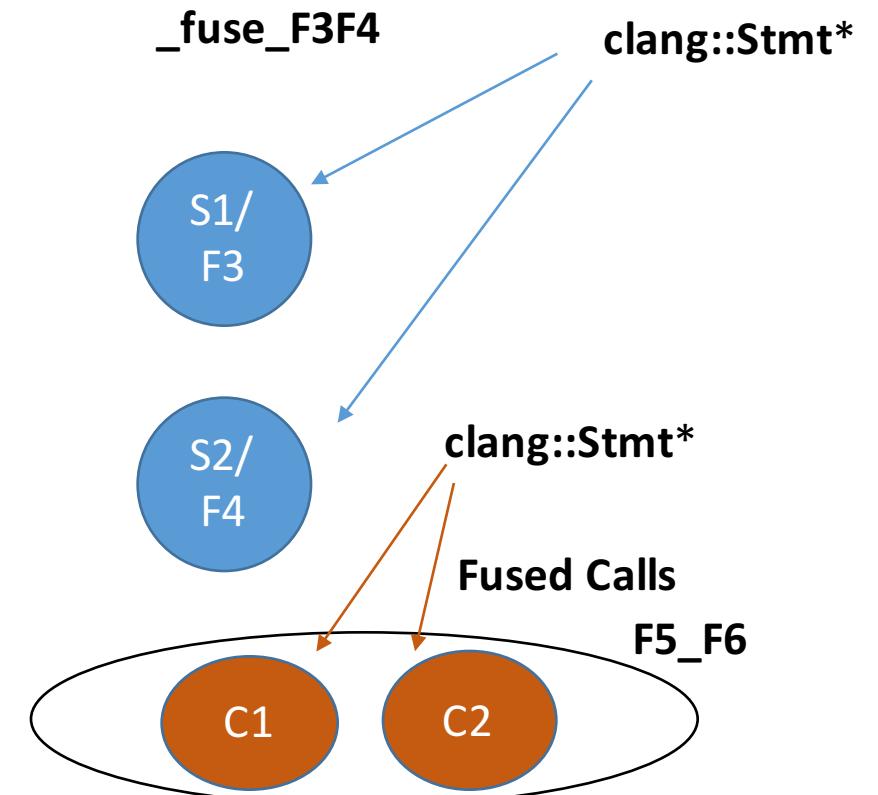
Static analysis

We used clang AST to collect information needed for Grafter static analysis

1. A function is represented as a sequence of Clang::Stmt* and Clang::CallExpr* nodes.

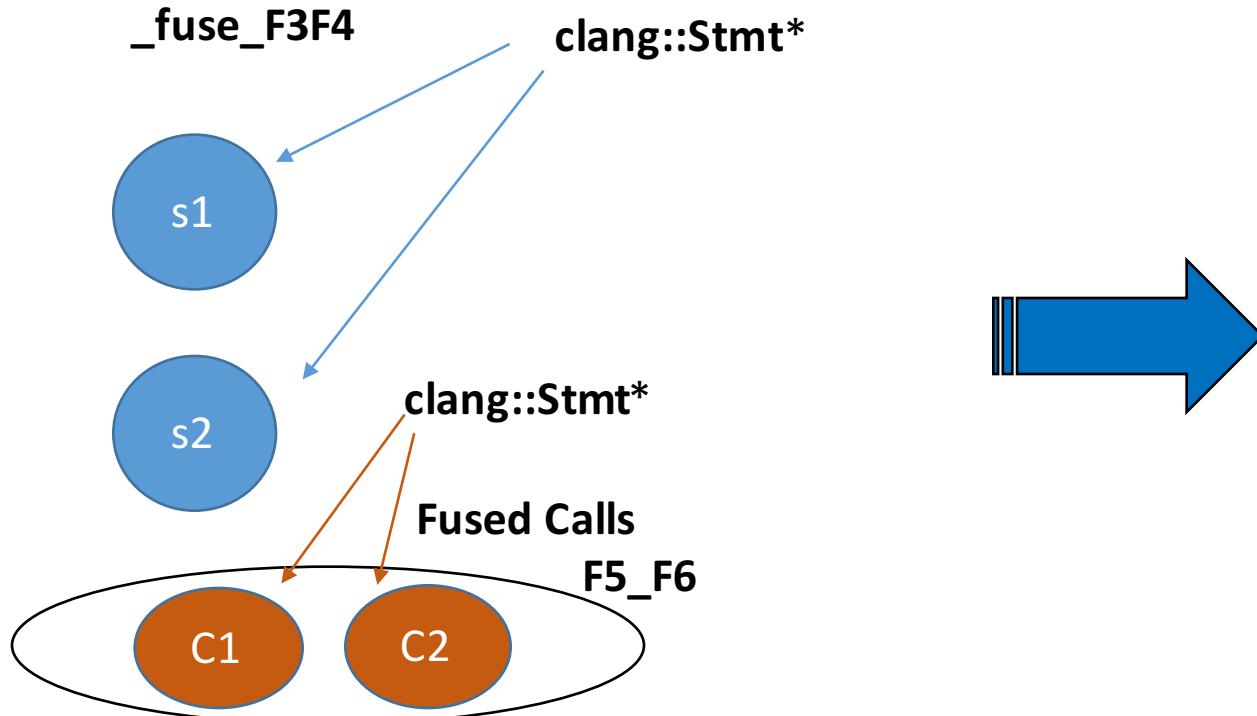
Different schedules are achieved by reordering statements and collapsing calls.

2. Understand accesses of statements and build call graphs to analyze dependences.



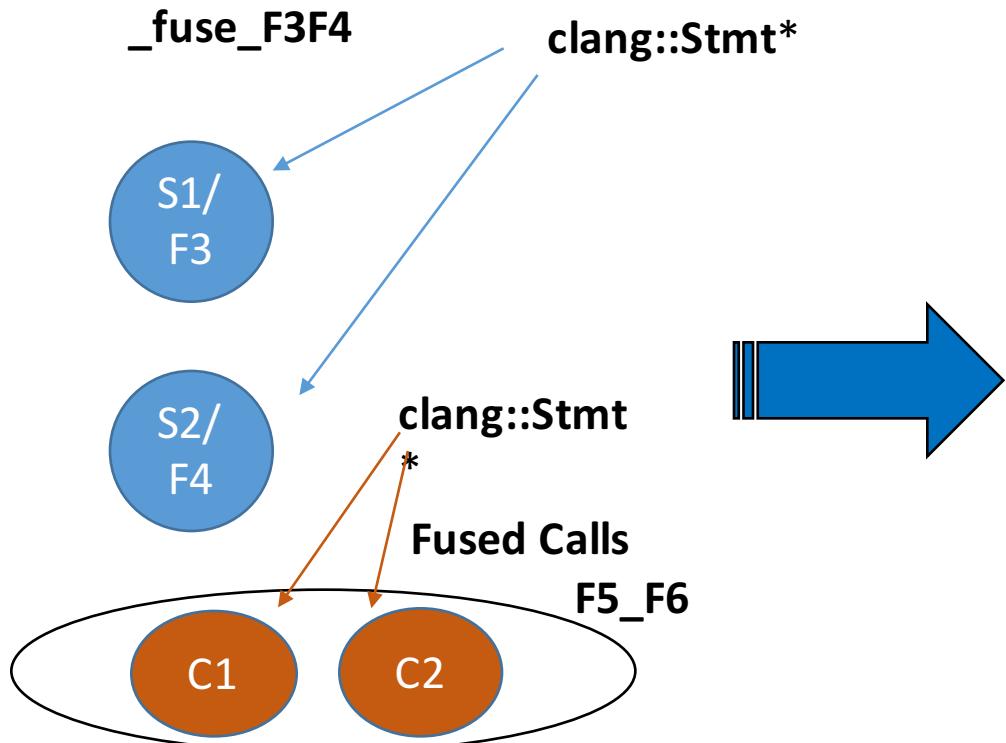
Code Generation

- Grafter build the fused functions incrementally following a set of rewrite rules while tracking the original source code



```
3 void _fuse__F3F4(TextBox *_r, int active_flags) {  
4     TextBox *_r_f0 = (TextBox *)(_r);  
5     TextBox *_r_f1 = (TextBox *)(_r);  
6     if (active_flags & 0b11) /*call*/ {  
7         unsigned int call_flags = 0;  
8         call_flags <= 1;  
9         call_flags |= (0b01 & (active_flags >> 1));  
10        call_flags <= 1;  
11        call_flags |= (0b01 & (active_flags >> 0));  
12        _r_f0->Next->_stub1(call_flags);  
13    }  
14    if (active_flags & 0b1) {  
15        _r_f0->Width = _r_f0->Text.Length;  
16        _r_f0->TotalWidth = _r_f0->Next->Width + _r_f0  
17            ->Width;  
18    }  
19    if (active_flags & 0b10) {  
20        _r_f1->Height = _r_f1->Text.Length * (_r_f1->  
21            Width / CHAR_WIDTH) + 1;  
22        _r_f1->MaxHeight = _r_f1->Height;  
23        if (_r_f1->Next->Height > _r_f1->Height) {  
24            _r_f1->MaxHeight = _r_f1->Next->Height;  
25        }  
26    }  
27};
```

Code generation example



Parameters

Track active traversals

Call fused functions

Computations

```
3 void _fuse__F3F4(TextBox *_r, int active_flags) {  
4     TextBox *_r_f0 = (TextBox *)(_r);  
5     TextBox *_r_f1 = (TextBox *)(_r);  
6     if (active_flags & 0b11) /*call*/ {  
7         unsigned int call_flags = 0;  
8         call_flags <<= 1;  
9         call_flags |= (0b01 & (active_flags >> 1));  
10        call_flags <<= 1;  
11        call_flags |= (0b01 & (active_flags >> 0));  
12        _r_f0->Next->__stub1(call_flags);  
13    }  
14    if (active_flags & 0b1) {  
15        _r_f0->Width = _r_f0->Text.Length;  
16        _r_f0->TotalWidth = _r_f0->Next->Width + _r_f0->Width;  
17    }  
18    if (active_flags & 0b10) {  
19        _r_f1->Height = _r_f1->Text.Length * (_r_f1->Width / CHAR_WIDTH) + 1;  
20        _r_f1->MaxHeight = _r_f1->Height;  
21        if (_r_f1->Next->Height > _r_f1->Height) {  
22            _r_f1->MaxHeight = _r_f1->Next->Height;  
23        }  
24    }  
25};
```

Code generation example

- Replace original calls with calls to fused functions

```
int main(){
    Element *ElementsList = ...;
    ElementsList->computeWidth();
    ElementsList->computeHeight();
}
```



```
void TextBox::__stub1(int active_flags) {
    _fuse__F3F4(this, active_flags);
}
void Group::__stub1(int active_flags) {
    _fuse__F5F6(this, active_flags);
}
void End::__stub1(int active_flags) {
    _fuse__F1F2(this, active_flags);
}
int main() {
    Group *ElementsList;
    //ElementsList->computeWidth();
    //ElementsList->computeHeight();
    ElementsList->__stub1(0b11);
}
```

It does scale..

- We run Grafter on programs with more than 50 functions
- Generating thousands lines of code
- Achieve significant speedups

Conclusions

- Clang can be useful in doing **performance-oriented** domain-specific source-level transformations
 - Easy to implement an embedded DSL in C++
 - Collect source-level information needed for static analysis
 - Generate new source code while tracking the input program

2019

LLVM DEV MEETING



Improving Your TABLEGEN Description

Javed Absar



GRAPHCORE

WHAT IS ‘TABLEGEN’ ?

- DSL invented for LLVM
- Used extensively
- Describe records of –
 - Instructions, registers, intrinsics, attributes, scheduler model, warnings, ...

THIS TALK IS **NOT...**

- Introduction to TableGen
- Complete Guide to TableGen



BUT...

Look at interesting features that you may

- Not know exist
- Know, but haven't used
- Used, but not this way

'FOREACH' AND 'CONCAT'

Listing One-by-One

```
def F9Dwarf : DwarfMapping<28>;  
def F11Dwarf : DwarfMapping<29>;  
def F13Dwarf : DwarfMapping<30>;  
def F15Dwarf : DwarfMapping<31>;
```

```
def F16Dwarf : DwarfMapping<68>;  
def F18Dwarf : DwarfMapping<69>;  
def F20Dwarf : DwarfMapping<70>;  
def F22Dwarf : DwarfMapping<71>;
```



'FOREACH' AND 'CONCAT'

Listing One-by-One

```
def F9Dwarf : DwarfMapping<28>;  
def F11Dwarf : DwarfMapping<29>;  
def F13Dwarf : DwarfMapping<30>;  
def F15Dwarf : DwarfMapping<31>;
```

```
def F16Dwarf : DwarfMapping<68>;  
def F18Dwarf : DwarfMapping<69>;  
def F20Dwarf : DwarfMapping<70>;  
def F22Dwarf : DwarfMapping<71>;
```

Generating using 'foreach'

```
1. def RegMap {  
2.   list<list<int>> val = [[9,28], [11,29], [13,30],  
   [15, 31], [16, 68], [18, 69],  
   [20, 70], [22, 71]];  
3. }  
  
4. foreach N = RegMap.val in {  
5.   def F#!head(N)#Dwarf :  
     DwarfMapping<!head(!tail(N))>;  
6. }
```



'FOLD'

Listing One-by-One

```
def {
    list<int> Ascending = [1, 2, 3, 4, 5, 6, 7, 8, 9];
    list<int> Descending = [9, 8, 7, 6, 5, 4, 3, 2, 1];
}
```



'FOLD'

Listing One-by-One

```
1. def {  
2.   list<int> Ascending = [1, 2, 3, 4, 5, 6, 7, 8, 9];  
  
3.   // list<int> Descending = [9, 8, 7, 6, 5, 4, 3, 2, 1];  
4.   list<int> Descending = reverse<Ascending>.val;  
5. }
```

Generating using 'foldl'

```
1. class reverse<list<int> L> {  
2.   list<int> val = !foldl([ ]<int>, L, a, b, [b] # a);  
3. }
```



'SORT' USING TABLEGEN

Merge Sort Algorithm

```
def mergesort(L):
    if len(L) < 2:
        return L
    else:
        h = len(L) div 2
        return merge(mergesort(L[:h]), mergesort(L[h:]))
```



'SORT' USING TABLEGEN

Merge Sort Algorithm

```
def mergesort(L):
    if len(L) < 2:
        return L
    else:
        h = len(L) div 2
        return merge(mergesort(L[:h]), mergesort(L[h:]))
```

Using TableGen

```
1. class mergeSort<list<int> L> {
2.     list<int> val = !if(!lt(!size(L),2), L,
3.                           merge<
4.                               mergeSort< split<L>.lower >.val,
5.                               mergeSort< split<L>.upper >.val
6.                           >.val );
7. }
```



'SORT' USING TABLEGEN

Using TableGen

```
1. class split<list<int> L> {  
2.   int sizeW = !size(L);  
  
3.   int sizeL = !srI(!add(sizeW,1),1);  
4.   int sizeU = !add(sizeW,!mul(sizeL,-1));  
  
5.   list<int> lower = !foldl([ ]<int>, L, a, b, !if(!lt(!size(a),sizeL),  
6.                             a # [b],  
7.                             a));  
8.   list<int> upper = !foldl(L, L, a, b, !if(!gt(!size(a), sizeU),  
9.                            !tail(a),  
10.                           a));  
11. }
```



'SORT' USING TABLEGEN

Merge Function (Functional Programming)

```
def merge(a, b):
    if len(a) == 0: return b
    elif len(b) == 0: return a
    elif a[0] < b[0]:
        return [a[0]] + merge(a[1:], b)
    else:
        return [b[0]] + merge(a, b[1:])
```



'SORT' USING TABLEGEN

Merge Function (Functional Programming)

```
def merge(a, b):
    if len(a) == 0: return b
    elif len(b) == 0: return a
    elif a[0] < b[0]:
        return [a[0]] + merge(a[1:], b)
    else:
        return [b[0]] + merge(a, b[1:])
```

Using TableGen

```
1. class merge<list<int> a, list<int> b> {
2.     list<int> val = !if(!eq(!size(a), 0), b,
3.                         !if(!eq(!size(b),0), a,
4.                             !if(!lt(!head(a), !head(b)), [!head(a)] # merge<!tail(a), b>.val,
5.                                 [!head(b)] # merge<a, !tail(b)>.val));
6. }
```



Source:

STATIC POLYMORPHISM WITH TABLEGEN

Test Code

```
1. class GeoObj { string draw = "Invalid"; }

2. class Line<int N> : GeoObj { string draw = "Line:" # N; }
3. class Rectangle<int N> : GeoObj { string draw = "Rectangle:" # N; }

4. class myDraw<list<GeoObj> objs> {
5.     list<string> drawings = !foreach(obj, objs, obj.draw);
6. }

7. def main {
8.     GeoObj o1 = Line<1>; GeoObj o2 = Line<2>; GeoObj o3 = Rectangle<3>;
9.     list<string> D = myDraw<[o1, o2, o3]>.drawings;
10. }
```



STATIC POLYMORPHISM WITH TABLEGEN

Test Code

```
1. class GeoObj { string draw = "Invalid"; }
2. class Line<int N> : GeoObj { string draw = "Line:" # N; }
3. class Rectangle<int N> : GeoObj { string draw = "Rectangle:" # N; }
4. class myDraw<list<GeoObj> objs> {
5.     list<string> drawings = !foreach(obj, objs, obj.draw);
6. }

7. def main {
8.     GeoObj o1 = Line<1>; GeoObj o2 = Line<2>; GeoObj o3 = Rectangle<3>;
9.     list<string> D = myDraw<[o1, o2, o3]>.drawings;
10. }
```

Generated Definition

```
$ ./bin/llvm-tblgen static_polymorphism.td
def main {
...
    list<string> D = ["Line:1", "Line:2", "Rectangle:3"];
}
```



'COND'

Using Nested-'IF's

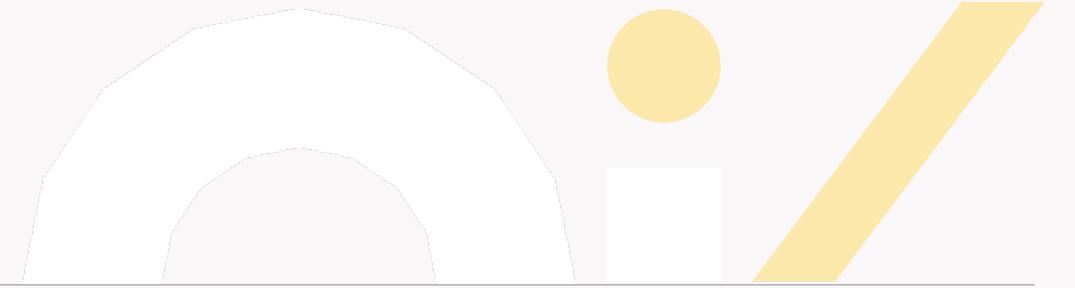
```
class getSubRegs<int size> {
    list<SubRegIndex> ret2 = [sub0, sub1];
    list<SubRegIndex> ret3 = [sub0, sub1, sub2];
    list<SubRegIndex> ret4 = [sub0, sub1, sub2, sub3];
    list<SubRegIndex> ret5 = [sub0, sub1, sub2, sub3, sub4];
    list<SubRegIndex> ret8 = [sub0, sub1, sub2, sub3, sub4, sub5,
    sub6, sub7];
    list<SubRegIndex> ret16 = [sub0, sub1, sub2, sub3,
        sub4, sub5, sub6, sub7,
        sub8, sub9, sub10, sub11,
        sub12, sub13, sub14, sub15];
    list<SubRegIndex> ret = !if(!eq(size, 2), ret2,
        !if(!eq(size, 3), ret3,
            !if(!eq(size, 4), ret4,
                !if(!eq(size, 5), ret5,
                    !if(!eq(size, 8), ret8,
                        ret16))))); }
```

Using 'COND'

```
class getSubRegs<int size> {
    list<SubRegIndex> ret2 = [sub0, sub1];
    list<SubRegIndex> ret3 = [sub0, sub1, sub2];
    list<SubRegIndex> ret4 = [sub0, sub1, sub2, sub3];
    list<SubRegIndex> ret5 = [sub0, sub1, sub2, sub3, sub4];
    list<SubRegIndex> ret8 = [sub0, sub1, sub2, sub3, sub4, sub5,
    sub6, sub7];
    list<SubRegIndex> ret_opt = !cond(!eq(size, 2): ret2,
        !eq(size, 3): ret3,
        !eq(size, 4): ret4,
        !eq(size, 5): ret5,
        !eq(size, 8): ret8,
        !eq(1, 1): ret16);
}
```



CONCLUSION



Take a look at the video and apply it to your context!

2019

LLVM DEV MEETING

