

HARDWARE LOOPS IN THE IPU BACKEND

Janek van Oirschot



GRAPHCORE

INTELLIGENCE PROCESSING UNIT (IPU)

- Chip with many cores (tiles) with inter-core communication
- Each tile has its own memory
- Each tile has multiple threads
- 2 execution pipelines (can run in parallel)
 - Main (for integer instructions + registers)
 - Aux (for float instructions + registers)
- Granularity of this talk: a single tile



HARDWARE LOOPS?

Currently, the LLVM IPU backend targets 2 instructions as hardware loops

- rpt (repeat) instruction
- brnzdec (branch if non-zero, decrement) instruction

RPT INSTRUCTION

```
rpt $reg, zimm8  
rpt zimmX, zimm8
```

Op #1: Trip count of the loop

Op #2: Number of loop body instruction bundle

```
void foo (int *out, int *in, unsigned size) {  
    for (int i = 0; i < size; ++i) {  
        *out += in[i];  
    }  
}
```

```
foo:  
    ld32 $m3, $m0, $m15, 0  
    {  
        rpt $m2, 2  
        fnop  
    }  
    {  
        ld32step $m4, $m15, $m1+=, 1  
        fnop  
    }  
    {  
        add $m3, $m3, $m4  
        fnop  
    }  
    {  
        st32 $m3, $m0, $m15, 0  
        fnop  
    }  
    br $m10
```

BRNZDEC INSTRUCTION

```
brnzdec $reg, zimm19
```

Op #1: Value to be tested and decremented
Op #2: Target address of branch

```
void foo (int *out, int *in, unsigned size) {
    for (int i = 0; i < size; ++i) {
        *out += in[i];
    }
}
```

```
foo:
    brz $m2, .LBB0_3
    add $m2, $m2, -1
    ld32 $m3, $m0, $m15, 0
.LBB0_2:
    ld32step $m4, $m15, $m1+=, 1
    add $m3, $m3, $m4
    st32 $m3, $m0, $m15, 0
    brnzdec $m2, .LBB0_2
.LBB0_3:
    br $m10
```

RPT VS BRNZDEC

rpt

- Trip count depends on IPU architecture
- Requires instructions in body to be bundled
- Maximum number of bundles in loop body is 256 bundles (constrained by operand bits)
- Implicitly skips over loop body if trip count == 0
- Control (e.g., branch) and System (e.g., trap) instructions **cannot** be executed in the body
 - -> no loop nests within

brnzdec

- Trip count limited to 32 bits available for registers
- Maximum number of instructions constrained by available bits for instruction encoding: 19 bits == 0x7FFF left shifted twice for 4 byte alignment: 0x1FFFC

RPT

rpt

- Trip count depends on IPU architecture
- Requires instructions in body to be bundled
- Maximum number of bundles in loop body is 256 bundles (constrained by operand bits)
- Implicitly skips over loop body if trip count == 0
- Control (e.g., branch) and System (e.g., trap) instructions **cannot** be executed in the body
 - -> no loop nests within

```
foo:  
    ld32 $m3, $m0, $m15, 0  
    {  
        rpt $m2, 2  
        fnop  
    }  
    {  
        ld32step $m4, $m15, $m1+=, 1  
        fnop  
    }  
    {  
        add $m3, $m3, $m4  
        fnop  
    }  
    {  
        st32 $m3, $m0, $m15, 0  
        fnop  
    }  
    br $m10
```

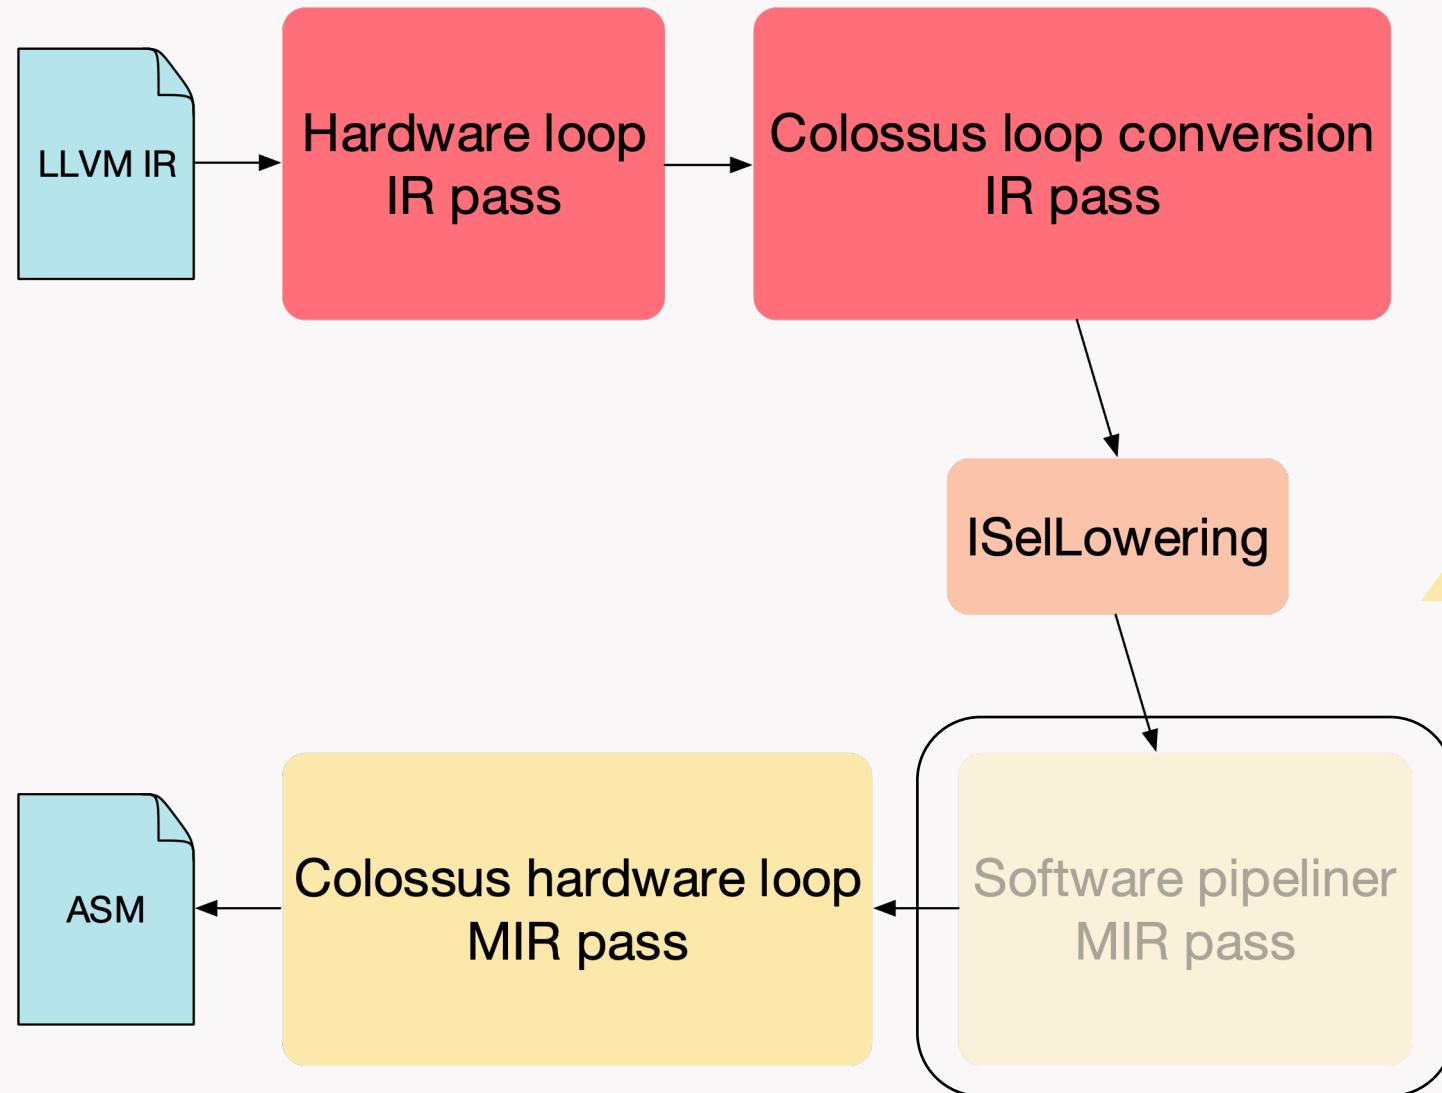
BRNZDEC

```
foo:  
    brz $m2, .LBB0_3  
    add $m2, $m2, -1  
    ld32 $m3, $m0, $m15, 0  
.LBB0_2:  
    ld32step $m4, $m15, $m1+=, 1  
    add $m3, $m3, $m4  
    st32 $m3, $m0, $m15, 0  
    brnzdec $m2, .LBB0_2  
.LBB0_3:  
    br $m10
```

brnzdec

- Trip count limited to 32 bits available for registers
- Maximum number of instructions constrained by available bits for instruction encoding: 19 bits == 0xFFFF left shifted twice for 4 byte alignment: 0x1FFFFC

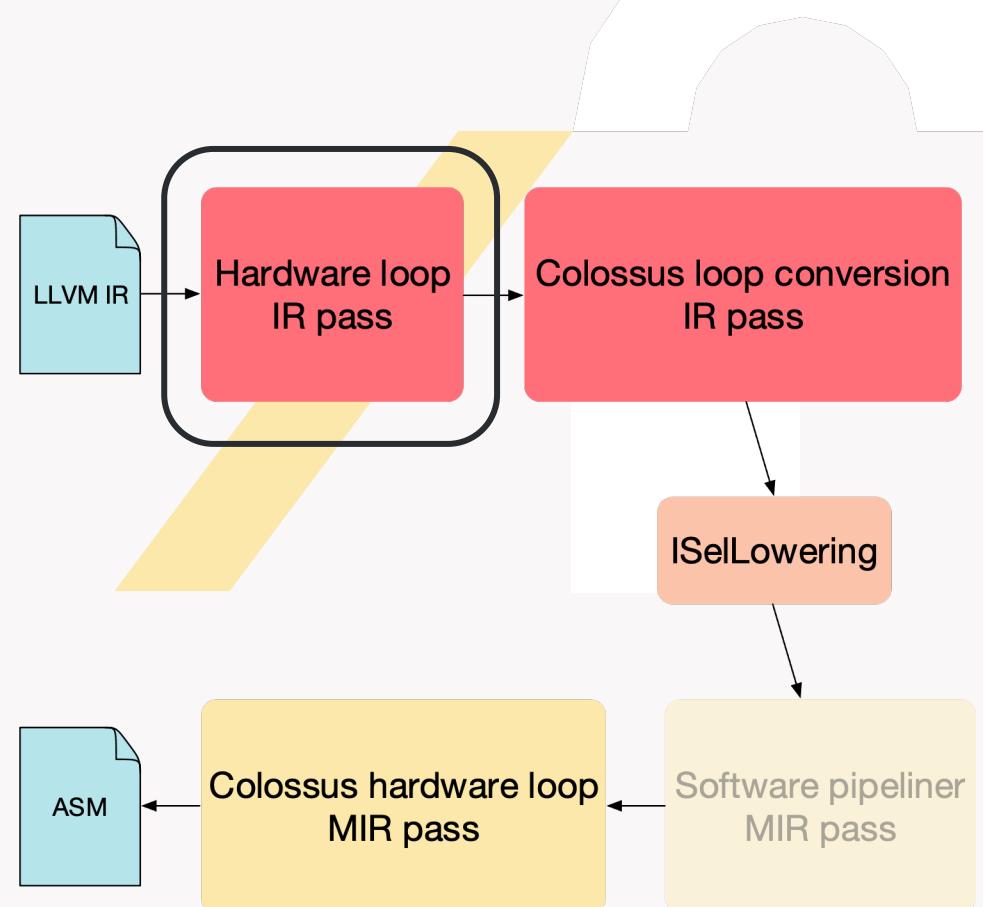
HIGH LEVEL VIEW



HARDWARE LOOP IR PASS

```
bool isHardwareLoopProfitable(Loop *L,  
ScalarEvolution &SE,  
AssumptionCache &AC,  
TargetLibraryInfo *LibInfo,  
HardwareLoopInfo &HWLoopInfo);
```

- Used to be an IPU custom pass but now we use the LLVM target independent Hardware Loop IR pass
- Run as part of CodeGenPrepare passes, as opposed to addPreISel passes of other targets, any later and the assumptions are removed from IR
- IsHardwareLoopProfitable TargetTransformInfo hook
 - Processing for assumptions
 - Emit different hardware loop settings depending on whether we want a rpt or brnzdec



HARDWARE LOOP IR PASS

```
auto emitRptHwLoop = [&HWLoopInfo, &i32, &i64](unsigned bitwidth) {  
    HWLoopInfo.IsNestingLegal = true;  
    HWLoopInfo.CounterInReg = false;  
    HWLoopInfo.CountType = bitwidth == 64 ? i64 : i32;  
    HWLoopInfo.LoopDecrement = ConstantInt::get(HWLoopInfo.CountType, 1);  
    HWLoopInfo.PerformEntryTest = true;  
};  
  
auto emitBrnzdecHwLoop = [&HWLoopInfo,&i32,&i64](unsigned bitwidth) {  
    HWLoopInfo.IsNestingLegal = true;  
    HWLoopInfo.CounterInReg = true;  
    HWLoopInfo.CountType = bitwidth == 64 ? i64 : i32;  
    HWLoopInfo.LoopDecrement = ConstantInt::get(HWLoopInfo.CountType, 1);  
};
```

HARDWARE LOOP IR PASS



HARDWARE LOOP IR PASS

- rpt is prioritized over brnzdec
- The constraints checked at this stage are trip count fit, loop nest level, and some trivial architecture constrains

```
if (!hasSubLoop && (isWithinRptRange || isKBValid) && ...) {  
    emitRptHwLoop(bitwidth);  
    return true;  
}  
  
if (isWithinBrnzdecRange || isKBValid_Brnzdec) {  
    emitBrnzdecHwLoop(bitwidth);  
    return true;  
}
```

HARDWARE LOOP IR PASS

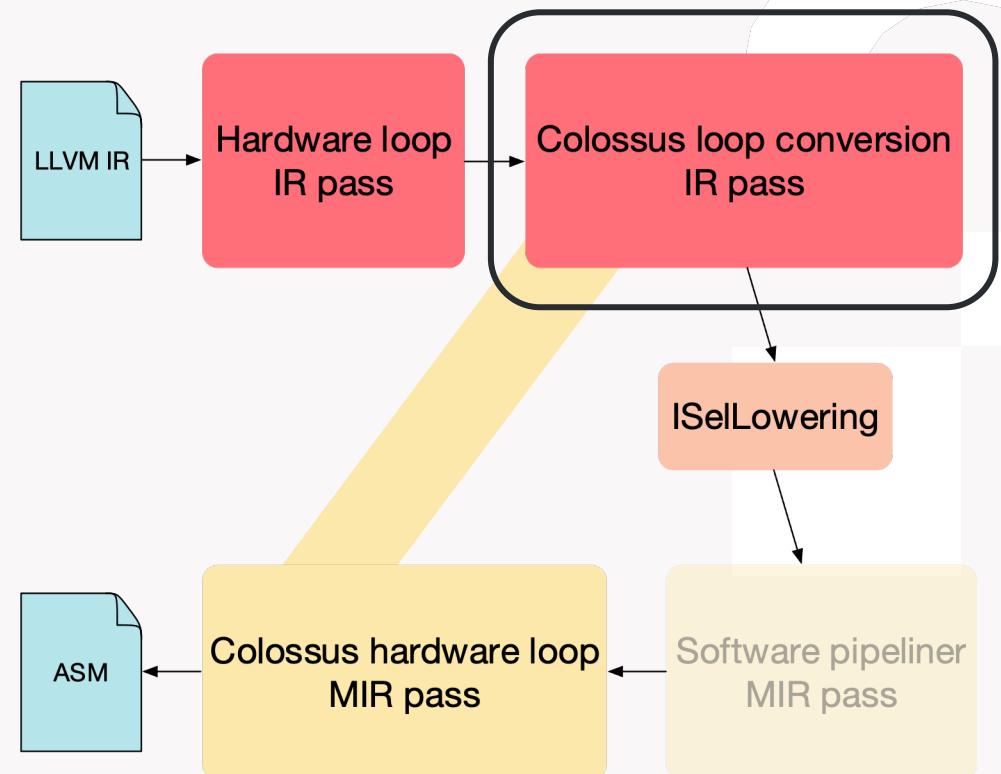
LLVM IR

Hardware loop
IR pass

LLVM IR with
hardware
loop
intrinsics

LOOP CONVERSION IR PASS

- Run as part of addPreISel passes
- Mechanical pass that converts LLVM target independent hardware loop intrinsics in IPU specific ones
- The IPU specific intrinsics are from the time when the hardware loop IR pass was target specific



LOOP CONVERSION IR PASS

```
define void @foo(i32* %out, i32* %in, i32 %size) {
entry:
    %0 = call i1 @llvm.test.set.loop.iterations.i32(i32 %size)
    br i1 %0, label %preheader, label %exit
preheader:
    %pre = load i32, i32* %out, align 4
    br label %body
body:
    %lsr.iv = phi i32* [ %in, %preheader ], [ %gep, %body ]
    %1 = phi i32 [ %add, %body ], [ %pre, %preheader ]
    %2 = load i32, i32* %lsr.iv, align 4
    %add = add nsw i32 %1, %2
    store i32 %add, i32* %out, align 4
    %gep = getelementptr i32, i32* %lsr.iv, i32 1
    %3 = call i1 @llvm.loop.decrement.i32(i32 1)
    br i1 %3, label %body, label %exit
exit:
    ret void
}
```

```
define void @foo(i32* %out, i32* %in, i32 %size) {
entry:
    %guard = call i32 @llvm.colossus.cloop.guard(i32 %size, i32 1)
    %guard.tr = trunc i32 %guard to i1
    br i1 %guard.tr, label %preheader, label %exit
preheader:
    %pre = load i32, i32* %out, align 4
    %begin = call i32 @llvm.colossus.cloop.begin(i32 %size, i32 1)
    br label %body
body:
    %lsr.iv = phi i32* [ %in, %preheader ], [ %gep, %body ]
    %0 = phi i32 [ %add, %body ], [ %pre, %preheader ]
    %1 = phi i32 [ %begin, %preheader ], [ %end.iv, %body ]
    %2 = load i32, i32* %lsr.iv, align 4
    %add = add nsw i32 %0, %2
    store i32 %add, i32* %out, align 4
    %gep = getelementptr i32, i32* %lsr.iv, i32 1
    %end = call { i32, i32 } @llvm.colossus.cloop.end(i32 %1, i32 1)
    %end.iv = extractvalue { i32, i32 } %end, 0
    %end.cc = extractvalue { i32, i32 } %end, 1
    %end.cc.tr = trunc i32 %end.cc to i1
    br i1 %end.cc.tr, label %body, label %exit
exit:
    ret void
}
```

LOOP CONVERSION IR PASS

```
define void @foo(i32* %out, i32* %in, i32 %size) {
entry:
  %0 = call i1 @llvm.test.set.loop.iterations.i32(i32 %size)
  br i1 %0, label %preheader, label %exit
preheader:
  %pre = load i32, i32* %out, align 4
  br label %body
body:
  %lsr.iv = phi i32* [ %in, %preheader ], [ %gep, %body ]
  %1 = phi i32 [ %add, %body ], [ %pre, %preheader ]
  %2 = load i32, i32* %lsr.iv, align 4
  %add = add nsw i32 %1, %2
  store i32 %add, i32* %out, align 4
  %gep = getelementptr i32, i32* %lsr.iv, i32 1
  %3 = call i1 @llvm.loop.decrement.i32(i32 1)
  br i1 %3, label %body, label %exit
exit:
  ret void
}
```

```
define void @foo(i32* %out, i32* %in, i32 %size) {
entry:
  %guard = call i32 @llvm.colossus.cloop.guard(i32 %size, i32 1)
  %guard.tr = trunc i32 %guard to i1
  br i1 %guard.tr, label %preheader, label %exit
preheader:
  %pre = load i32, i32* %out, align 4
  %begin = call i32 @llvm.colossus.cloop.begin(i32 %size, i32 1)
  br label %body
body:
  %lsr.iv = phi i32* [ %in, %preheader ], [ %gep, %body ]
  %0 = phi i32 [ %add, %body ], [ %pre, %preheader ]
  %1 = phi i32 [ %begin, %preheader ], [ %end.iv, %body ]
  %2 = load i32, i32* %lsr.iv, align 4
  %add = add nsw i32 %0, %2
  store i32 %add, i32* %out, align 4
  %gep = getelementptr i32, i32* %lsr.iv, i32 1
  %end = call { i32, i32 } @llvm.colossus.cloop.end(i32 %1, i32 1)
  %end.iv = extractvalue { i32, i32 } %end, 0
  %end.cc = extractvalue { i32, i32 } %end, 1
  %end.cc.tr = trunc i32 %end.cc to i1
  br i1 %end.cc.tr, label %body, label %exit
exit:
  ret void
}
```

LOOP CONVERSION IR PASS

```
define void @foo(i32* %out, i32* %in, i32 %size) {
entry:
    %0 = call i1 @llvm.test.set.loop.iterations.i32(i32 %size)
    br i1 %0, label %preheader, label %exit
preheader:
    %pre = load i32, i32* %out, align 4
    br label %body
body:
    %lsr.iv = phi i32* [ %in, %preheader ], [ %gep, %body ]
    %1 = phi i32 [ %add, %body ], [ %pre, %preheader ]
    %2 = load i32, i32* %lsr.iv, align 4
    %add = add nsw i32 %1, %2
    store i32 %add, i32* %out, align 4
    %gep = getelementptr i32, i32* %lsr.iv, i32 1
    %3 = call i1 @llvm.loop.decrement.i32(i32 1)
    br i1 %3, label %body, label %exit
exit:
    ret void
}
```

```
enum class metadata : uint16_t {
    none = 0,
    tripCount0KForRpt = 1,
};
```

```
define void @foo(i32* %out, i32* %in, i32 %size) {
entry:
    %guard = call i32 @llvm.colossus.cloop.guard(i32 %size, [i32 1])
    %guard.tr = trunc i32 %guard to i1
    br i1 %guard.tr, label %preheader, label %exit
preheader:
    %pre = load i32, i32* %out, align 4
    %begin = call i32 @llvm.colossus.cloop.begin(i32 %size, [i32 1])
    br label %body
body:
    %lsr.iv = phi i32* [ %in, %preheader ], [ %gep, %body ]
    %0 = phi i32 [ %add, %body ], [ %pre, %preheader ]
    %1 = phi i32 [ %begin, %preheader ], [ %end.iv, %body ]
    %2 = load i32, i32* %lsr.iv, align 4
    %add = add nsw i32 %0, %2
    store i32 %add, i32* %out, align 4
    %gep = getelementptr i32, i32* %lsr.iv, i32 1
    %end = call { i32, i32 } @llvm.colossus.cloop.end(i32 %1, [i32 1])
    %end.iv = extractvalue { i32, i32 } %end, 0
    %end.cc = extractvalue { i32, i32 } %end, 1
    %end.cc.tr = trunc i32 %end.cc to i1
    br i1 %end.cc.tr, label %body, label %exit
exit:
    ret void
}
```

LOOP CONVERSION IR PASS

LLVM IR with
hardware
loop
intrinsics

Colossus loop conversion
IR pass

LLVM IR with
IPU hardware
loop intrinsics

LOWERING TO ISD NODES

i32 colossus_cloop_guard(i32, i32)

CLOOP_GUARD_BRANCH(bb, i32, i32)

Basic block from branch
being replaced

Trip count

Metadata

LLVM IR

ISD nodes

LOWERING TO ISD NODES

i32 colossus_cloop_begin(i32, i32)

i32 CLOOP_BEGIN_VALUE(i32, i32)
Induction var Metadata

CLOOP_BEGIN_TERMINATOR(i32, i32)

Metadata

ISD nodes

LOWERING TO ISD NODES

(i32, i32) `colossus_cloop_end(i32, i32)`

i32

CLOOP_END_VALUE(i32, i32)

Induction var

Metadata

CLOOP_END_BRANCH(bb, i32, i32)

Basic block from branch
being replaced

Metadata

LLVM IR

ISD nodes

INSTRUCTION SELECTION

Almost 1:1 mapping from ISD node with same-name pseudo instructions

```
def ColossusCloopBeginValue :  
    SDNode<"ColossusISD::CL0OP_BEGIN_VALUE",  
    ...>;
```

```
def CL0OP_BEGIN_VALUE : Pseudo<(outs MR:$op0),  
    (ins MR:$op1, imm16zi:$meta),  
    ...>;
```

```
def : Pat<(i32 (ColossusCloopBeginValue i32MR:$op0, imm16zi:$meta)),  
    (i32 (CL0OP_BEGIN_VALUE i32MR:$op0, imm16zi:$meta))>;
```

INSTRUCTION SELECTION

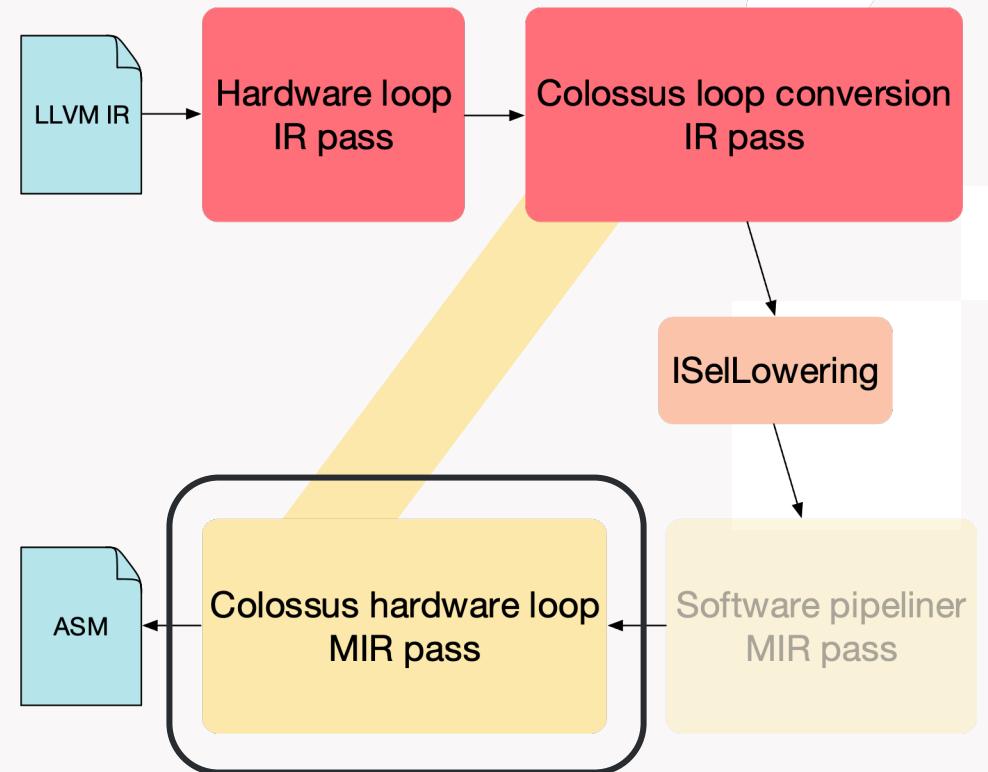
LLVM IR with
IPU hardware
loop intrinsics

ISelLowering

LLVM MIR with
IPU hardware
loop pseudo
instructions

HARDWARE LOOP MIR PASS

- Run as part of the addPreEmitPass passes
 - After register allocation but just before emitting assembly
 - Removes all pseudo instructions and emits final hardware loop instruction
 - Late in pass pipeline because the rpt instruction bundle count can only be accurately computed near the end of all code transformations



HARDWARE LOOP MIR PASS

- Run as part of the addPreEmitPass passes
 - After register allocation but just before emitting assembly
- Removes all pseudo instructions and emits final hardware loop instruction
 - Late in pass pipeline because the rpt instruction bundle count can only be accurately computed near the end of all code transformations

```
void foo (int *out, int *in, unsigned size) {
    for (int i = 0; i < size; ++i) {
        *out += in[i];
    }
}
```

```
bb.0.entry:
    successors: %bb.1(0x50000000), %bb.3(0x30000000)
    liveins: $m0, $m1, $m2
    CLOOP_GUARD_BRANCH $m2, %bb.3, 1

bb.1.for.body.preheader:
; predecessors: %bb.0
successors: %bb.2(0x80000000); %bb.2(100.00%)
liveins: $m0, $m1, $m2
$m2 = CLOOP_BEGIN_VALUE killed $m2(tied-def 0), 1
$m3 = LD32_ZI $m0, $mzero, 0, 0
CLOOP_BEGIN_TERMINATOR $m2, 1

bb.2.for.body:
; predecessors: %bb.1, %bb.2
successors: %bb.2(0x7c000000), %bb.3(0x04000000)
liveins: $m0, $m1, $m2, $m3
$m1, $m4 = LD32STEP_SI $mzero, killed $m1(tied-def 0),
1, 0
$m2 = CLOOP_END_VALUE killed $m2(tied-def 0), 1
$m3 = nsw ADD killed $m3, killed $m4, 0
ST32_ZI $m3, $m0, $mzero, 0, 0
CLOOP_END_BRANCH $m2, %bb.2, 1

bb.3.for.cond.cleanup:
; predecessors: %bb.0, %bb.2
RTN $lr, 0
```

HARDWARE LOOP MIR PASS

After hardware loop mir pass

- Run as part of the addPreEmitPass passes
 - After register allocation but just before emitting assembly
 - Removes all pseudo instructions and emits final hardware loop instruction
 - Late in pass pipeline because the rpt instruction bundle count can only be accurately computed near the end of all code transformations

```
void foo (int *out, int *in, unsigned size) {
    for (int i = 0; i < size; ++i) {
        *out += in[i];
    }
}
```

```
bb.0.entry:
    successors: %bb.1(0x80000000); %bb.1(100.00%)
    liveins: $m0, $m1, $m2

bb.1.for.body.preheader:
; predecessors: %bb.0
    successors: %bb.2(0x80000000); %bb.2(100.00%)
    liveins: $m0, $m1, $m2
    $m3 = LD32_ZI $m0, $mzero, 0, 0

bb.2.for.body:
; predecessors: %bb.1
    successors: %bb.3(0x80000000); %bb.3(100.00%)
    liveins: $m0, $m1, $m2, $m3
    RPT $m2, 2, 1 {
        $a13 = SETZI_A 0, 1
    }
    $m1, $m4 = LD32STEP_SI $mzero, killed $m1(tied-def 0),
1, 1 {
        $a13 = SETZI_A 0, 1
    }
    $m3 = nsw ADD killed $m3, killed $m4, 1 {
        $a13 = SETZI_A 0, 1
    }
    ST32_ZI $m3, $m0, $mzero, 0, 1 {
        $a13 = SETZI_A 0, 1
    }

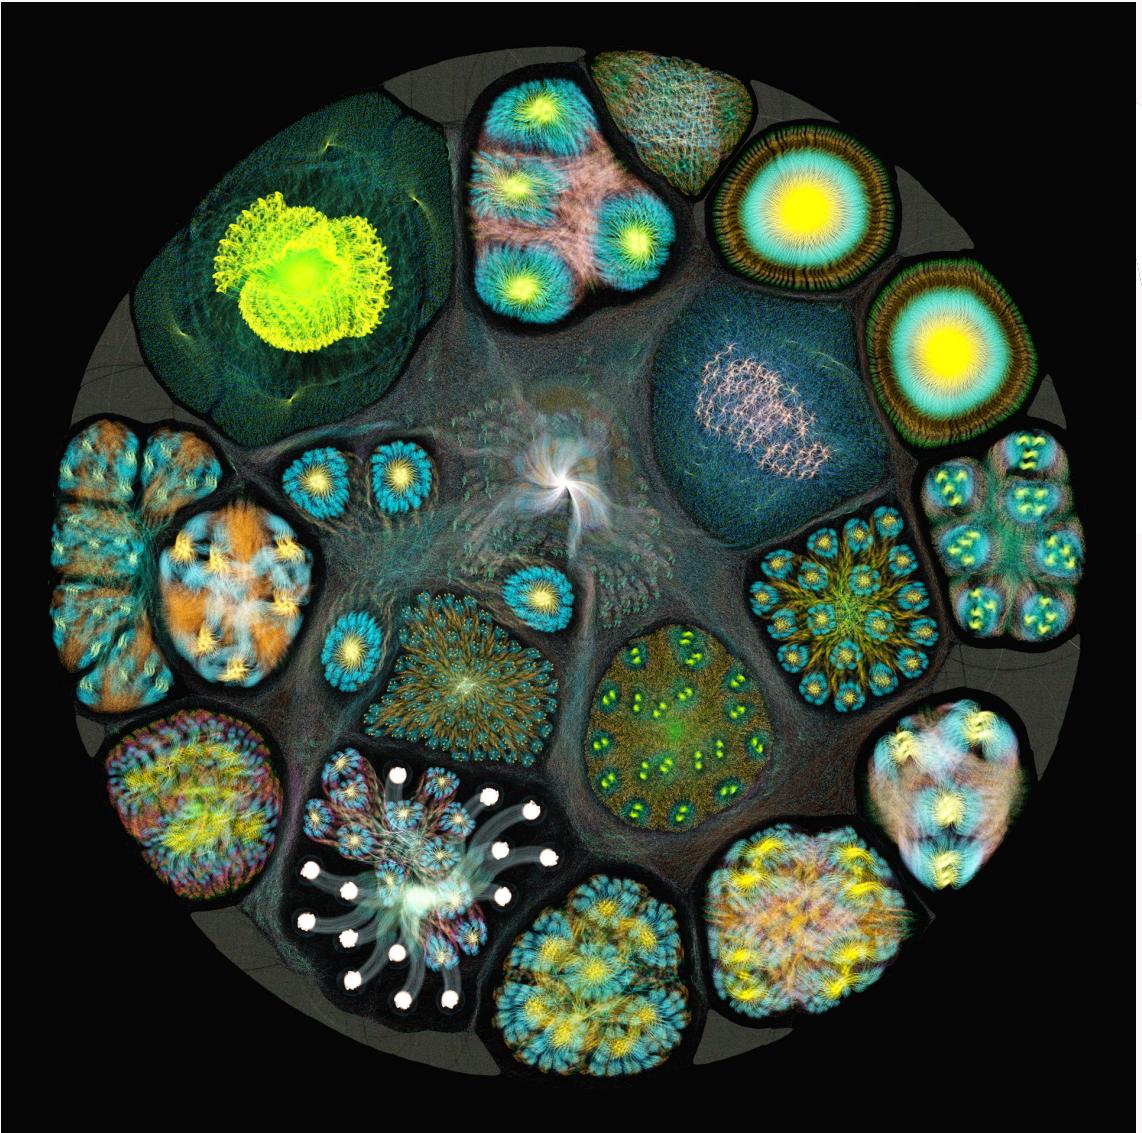
bb.3.for.cond.cleanup:
; predecessors: %bb.2
    RTN $lr, 0
```

```
bb.0.entry:  
successors: %bb.1(0x50000000), %bb.3(0x30000000)  
liveins: $m0, $m1, $m2  
CL0OP_GUARD_BRANCH $m2, %bb.3, 1  
  
bb.1.for.body.preheader:  
; predecessors: %bb.0  
successors: %bb.2(0x80000000); %bb.2(100.00%)  
liveins: $m0, $m1, $m2  
$m2 = CL0OP_BEGIN_VALUE killed $m2(tied-def 0), 1  
$m3 = LD32_ZI $m0, $mzero, 0, 0  
CL0OP_BEGIN_TERMINATOR $m2, 1  
  
bb.2.for.body:   
; predecessors: %bb.1, %bb.2  
successors: %bb.2(0x7c000000), %bb.3(0x04000000)  
liveins: $m0, $m1, $m2, $m3  
$m1, $m4 = LD32STEP_SI $mzero, killed $m1(tied-def 0),  
$m2 = CL0OP_END_VALUE killed $m2(tied-def 0), 1  
$m3 = nsw ADD killed $m3, killed $m4, 0  
ST32_ZI $m3, $m0, $mzero, 0, 0  
CL0OP_END_BRANCH $m2, %bb.2, 1  
  
bb.3.for.cond.cleanup:  
; predecessors: %bb.0, %bb.2  
RTN $lr, 0
```

```
bb.0.entry:  
successors: %bb.1(0x80000000); %bb.1(100.00%)  
liveins: $m0, $m1, $m2  
  
bb.1.for.body.preheader:  
; predecessors: %bb.0  
successors: %bb.2(0x80000000); %bb.2(100.00%)  
liveins: $m0, $m1, $m2  
$m3 = LD32_ZI $m0, $mzero, 0, 0  
  
bb.2.for.body:  
; predecessors: %bb.1  
successors: %bb.3(0x80000000); %bb.3(100.00%)  
liveins: $m0, $m1, $m2, $m3  
RPT $m2, 2, 1 {  
    $a13 = SETZI_A 0, 1  
}  
$m1, $m4 = LD32STEP_SI $mzero, killed $m1(tied-def 0),  
1, 1 {  
    $a13 = SETZI_A 0, 1  
}  
$m3 = nsw ADD killed $m3, killed $m4, 1 {  
    $a13 = SETZI_A 0, 1  
}  
ST32_ZI $m3, $m0, $mzero, 0, 1 {  
    $a13 = SETZI_A 0, 1  
}  
  
bb.3.for.cond.cleanup:  
; predecessors: %bb.2  
RTN $lr, 0
```

HARDWARE LOOP USERS

- Most of our users are internal, working on Poplibs, and considered “power users”
 - Comfortable and aware of IPU architecture
 - Will often use assembly for optimal solution
- Try to convince them to use C++ without performance loss
 - `builtin_assume`
 - `rptsize_t`
 - [WIP] pragma



HARDWARE LOOP USERS: ASSUME

- Allow user to use `__builtin_assume` calls to insert assumptions on trip count

```
void foo (int *out, int *in, unsigned size) {  
    for (int i = 0; i < size; ++i) {  
        *out += in[i];  
    }  
}
```

Emits brnzdec

```
void assume_foo (int *out, int *in, unsigned size) {  
    __builtin_assume(size < 0xFFFF);  
    for (int i = 0; i < size; ++i) {  
        *out += in[i];  
    }  
}
```

Emits rpt

HARDWARE LOOP USERS: RPTSIZE_T

- Unsigned int type to be used as trip count to target rpt instructions
 - Bit width will be rpt instruction trip count width (and change according to architecture, if changed)
- C++ only
 - Uses _BitInt(n) (formerly _ExtInt(n))
 - Overloads operations with other (unsigned) int types

```
void bar (int *out, int *a, unsigned n) {
    for (int i = 0; i < n; ++i) {
        *out += a[i];
    }
}

void rptsizebar1 (int *out, int *a, rptsize_t n) {
    for (int i = 0; i < n; ++i) {
        *out += a[i];
    }
}

void rptsizebar2 (int *out, int *a, rptsize_t n) {
    for (rptsize_t i = 0; i < n; ++i) {
        *out += a[i];
    }
}
```

HARDWARE LOOP USERS: RPTSIZE_T

- Unsigned int type to be used as trip count to target rpt instructions
 - Bit width will be rpt instruction trip count width (and change according to architecture, if changed)
- C++ only
 - Uses _BitInt(n) (formerly _ExtInt(n))
 - Overloads operations with other (unsigned) int types

```
_Z3barPiS_j:
    brz $m2, .LBB0_3
    add $m2, $m2, -1
    ld32 $m3, $m0, $m15, 0
.LBB0_2:
    ld32step $m4, $m15, $m1+=, 1
    add $m3, $m3, $m4
    st32 $m3, $m0, $m15, 0
    brnzdec $m2, .LBB0_2
.LBB0_3:
    br $m10
_Z11rptsizebar1PiS_6ap_intILj12ELb0EE:
    ldz16 $m2, $m2, $m15, 0
    ld32 $m3, $m0, $m15, 0
{
    rpt $m2, 2
    fnop
}
{
    ld32step $m4, $m15, $m1+=, 1
    fnop
}
{
    add $m3, $m3, $m4
    fnop
}
{
    st32 $m3, $m0, $m15, 0
    fnop
}
br $m10
```

HARDWARE LOOP USERS: PRAGMA

- Not all loops have a trip count explicitly defined in a variable or have exit conditions dependent on an iterating integer (e.g., `!container.empty()`)
- Hence, we're thinking of adding pragma support for said loops such that power users can still force rpt instructions

pragma hwloop arg	semantics
<code>force_any</code>	ignore loop count & profitability (try brnzdec if not possible)
<code>force_rpt</code>	ignore loop count & profitability (error if not possible)
<code>force_brnzdec</code>	ignore rpt (implies brnzdec)
<code>none</code>	disable rpt & brnzdec

THANK YOU

Janek van Oirschot
janekvo@graphcore.ai