

# MCA Daemon

## Hybrid Throughput Analysis Beyond Basic Blocks

Min-Yih “Min” Hsu, David Gens, Michael Franz. University of California, Irvine

# Outline

# Outline

## Motivation

# Outline

Motivation

MCA Daemon (MCAD)

# Outline

Motivation

MCA Daemon (MCAD)

Future Plans

# Outline

Motivation

MCA Daemon (MCAD)

Future Plans

Epilogue

# Outline

Motivation

MCA Daemon (MCAD)

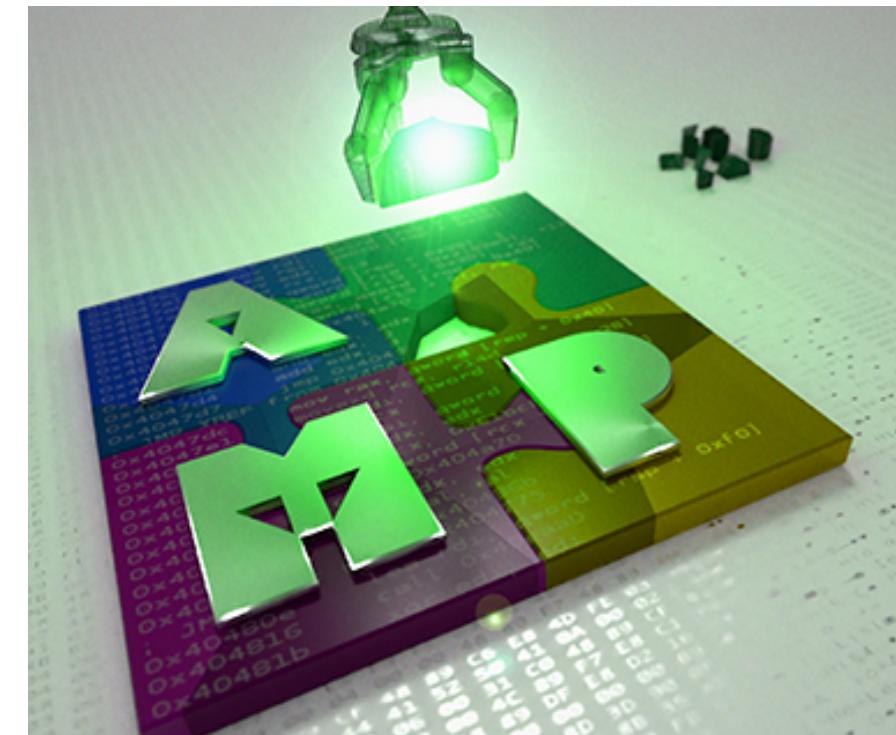
Future Plans

Epilogue

# Genesis: Assured Micro Patching (AMP)

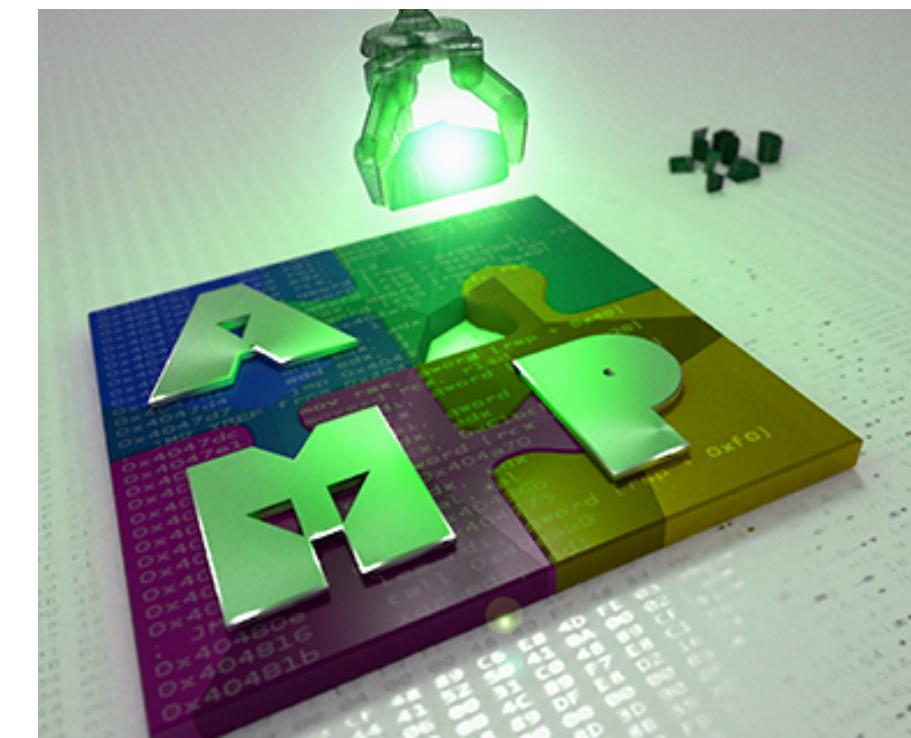
# Genesis: Assured Micro Patching (AMP)

- A research project initiated by United States DARPA to assure the correctness of **binary patching** with little or *no* source code



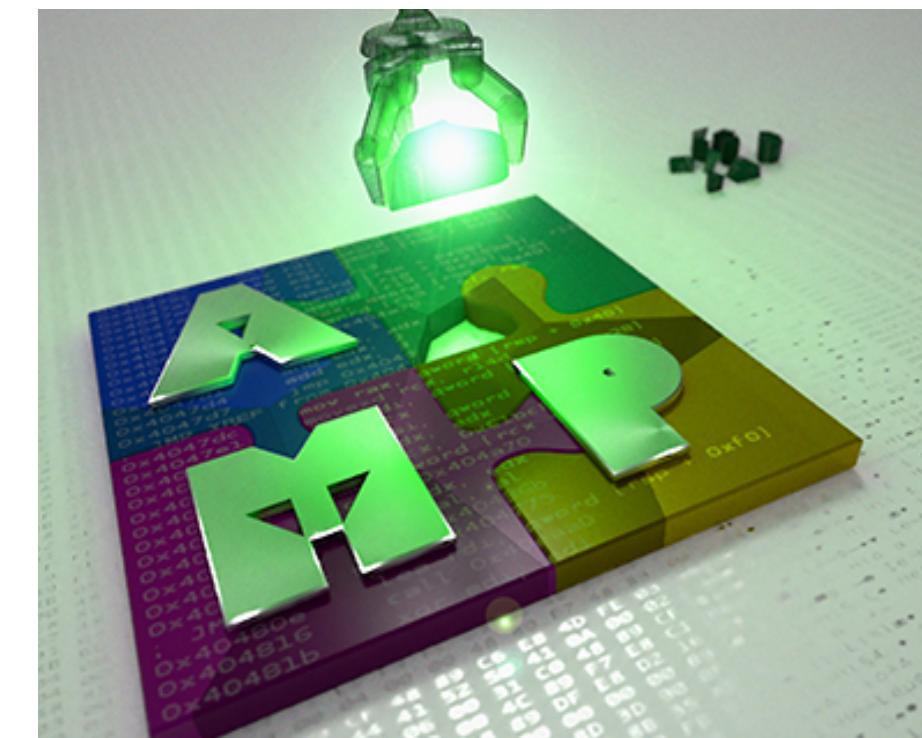
# Genesis: Assured Micro Patching (AMP)

- A research project initiated by United States DARPA to assure the correctness of **binary patching** with little or *no* source code
  - Including *functional* and *timing* aspects



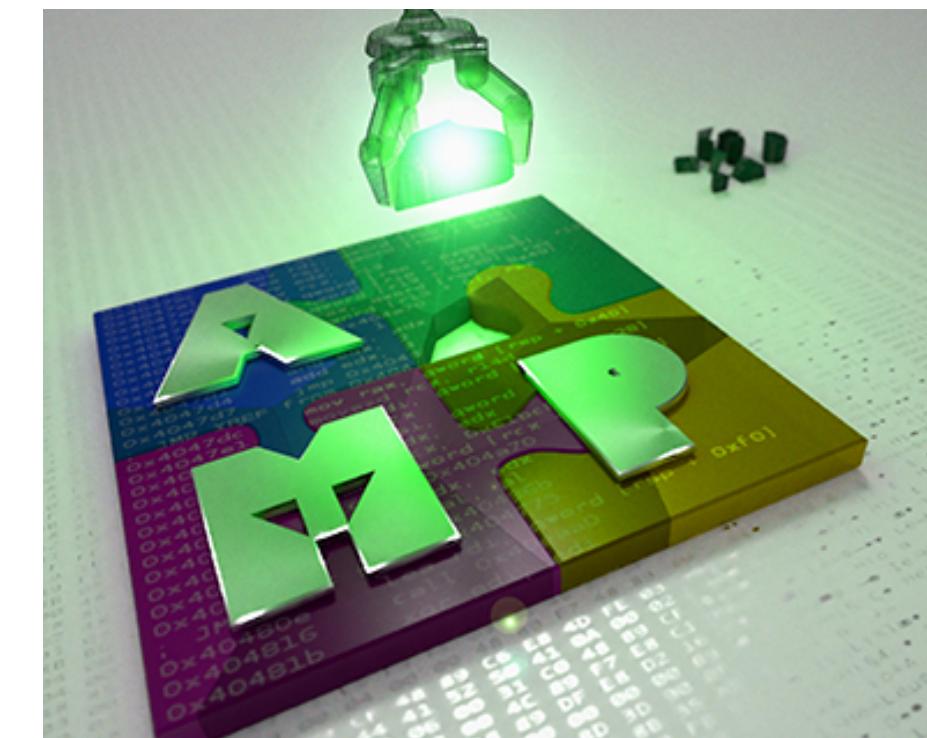
# Genesis: Assured Micro Patching (AMP)

- A research project initiated by United States DARPA to assure the correctness of **binary patching** with little or *no* source code
  - Including *functional* and *timing* aspects
  - Focuses on **small (micro)** binary patches



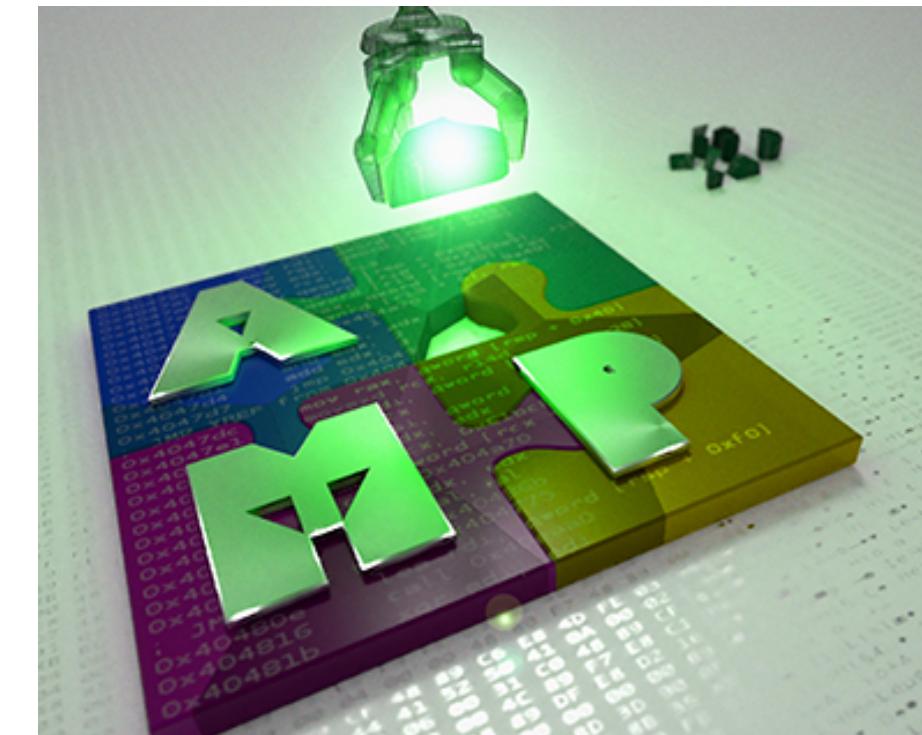
# Genesis: Assured Micro Patching (AMP)

- A research project initiated by United States DARPA to assure the correctness of **binary patching** with little or *no* source code
  - Including *functional* and *timing* aspects
  - Focuses on **small (micro)** binary patches
  - Focuses on **embedded systems**



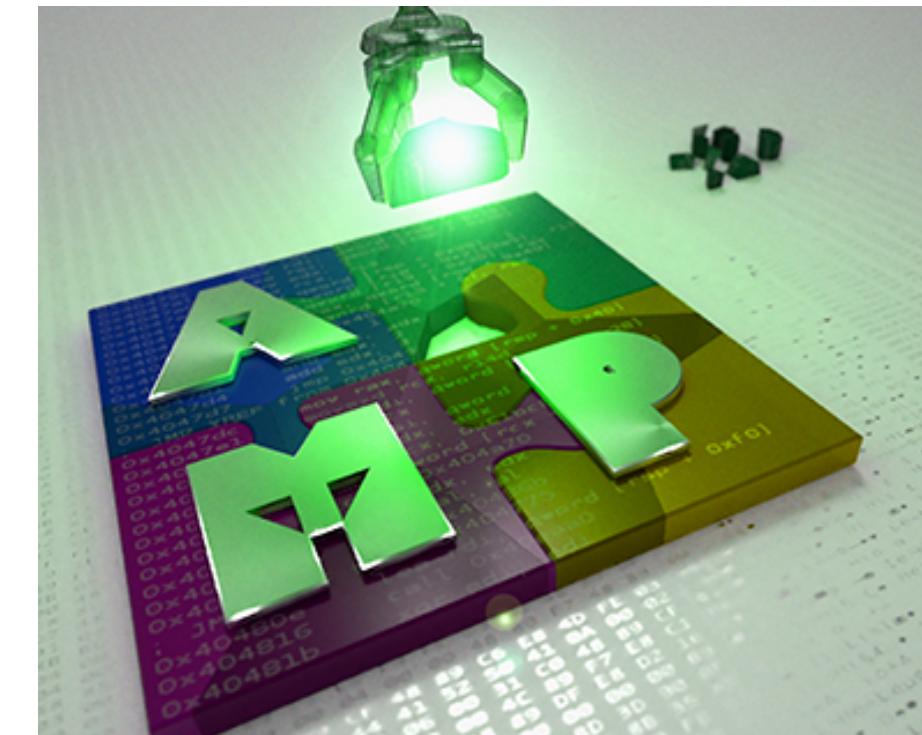
# Genesis: Assured Micro Patching (AMP)

- A research project initiated by United States DARPA to assure the correctness of **binary patching** with little or *no* source code
  - Including *functional* and *timing* aspects
  - Focuses on **small (micro)** binary patches
  - Focuses on **embedded systems**
- UCI was studying the timing impacts of binary patches



# Genesis: Assured Micro Patching (AMP)

- A research project initiated by United States DARPA to assure the correctness of **binary patching** with little or *no* source code
  - Including *functional* and *timing* aspects
  - Focuses on **small (micro)** binary patches
  - Focuses on **embedded systems**
- UCI was studying the timing impacts of binary patches
  - Example: After the firmware on a truck is binary-patched to prevent brakes from locking up, we need to make sure **latencies** do not degrade terribly



# Timing impacts of binary patches

## Problem definition

Original  
Program

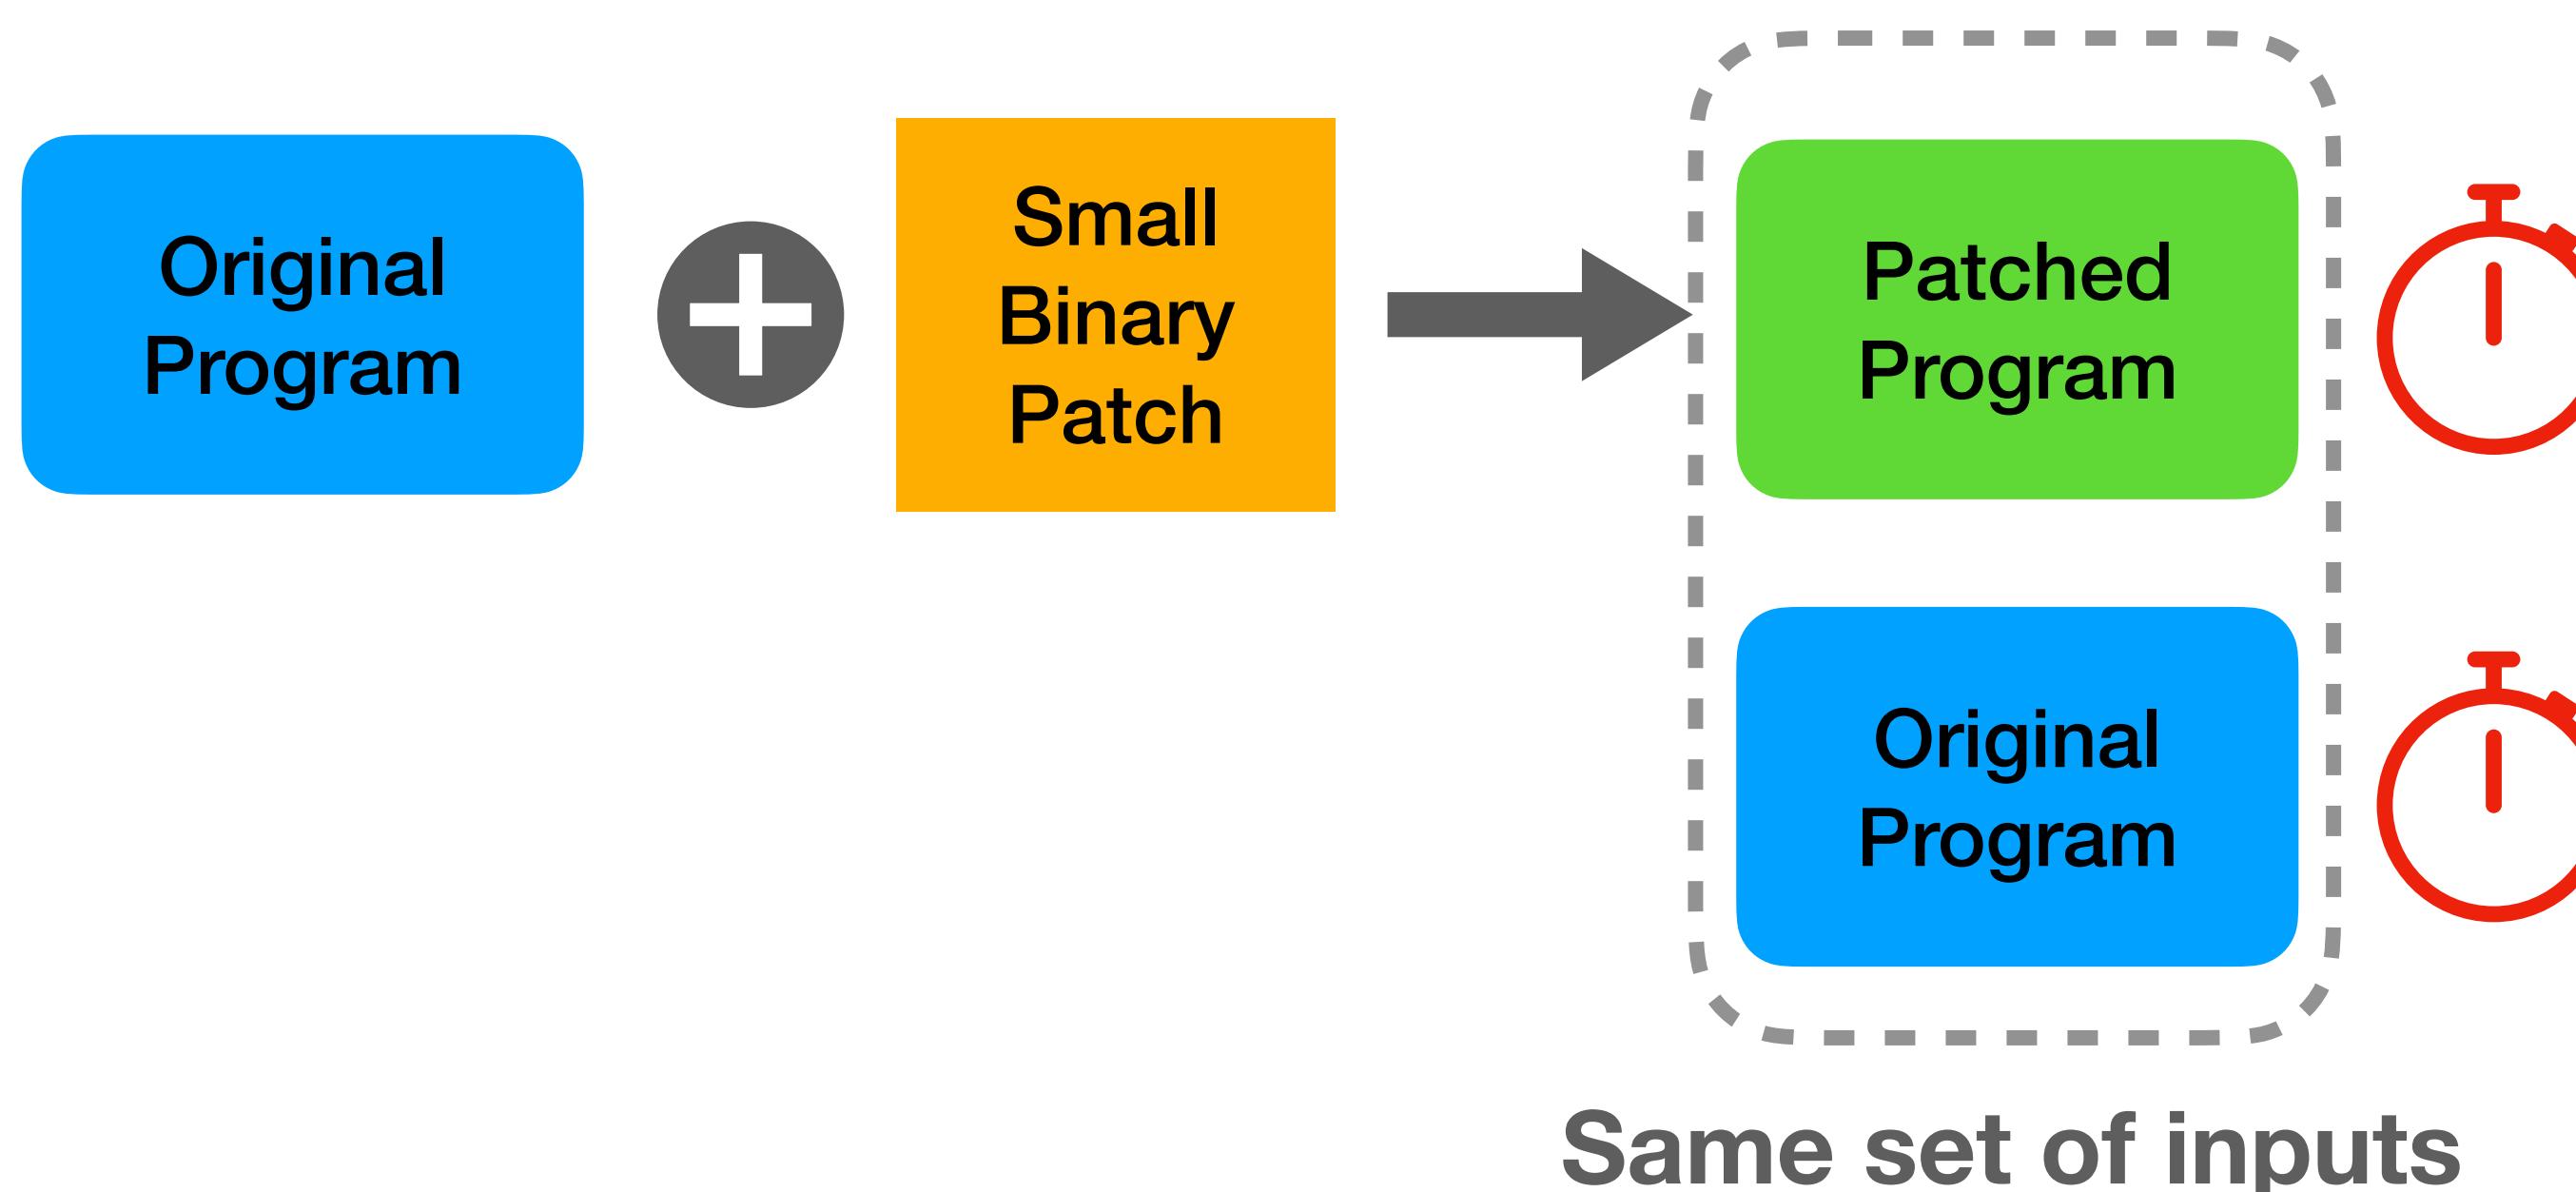
# Timing impacts of binary patches

## Problem definition



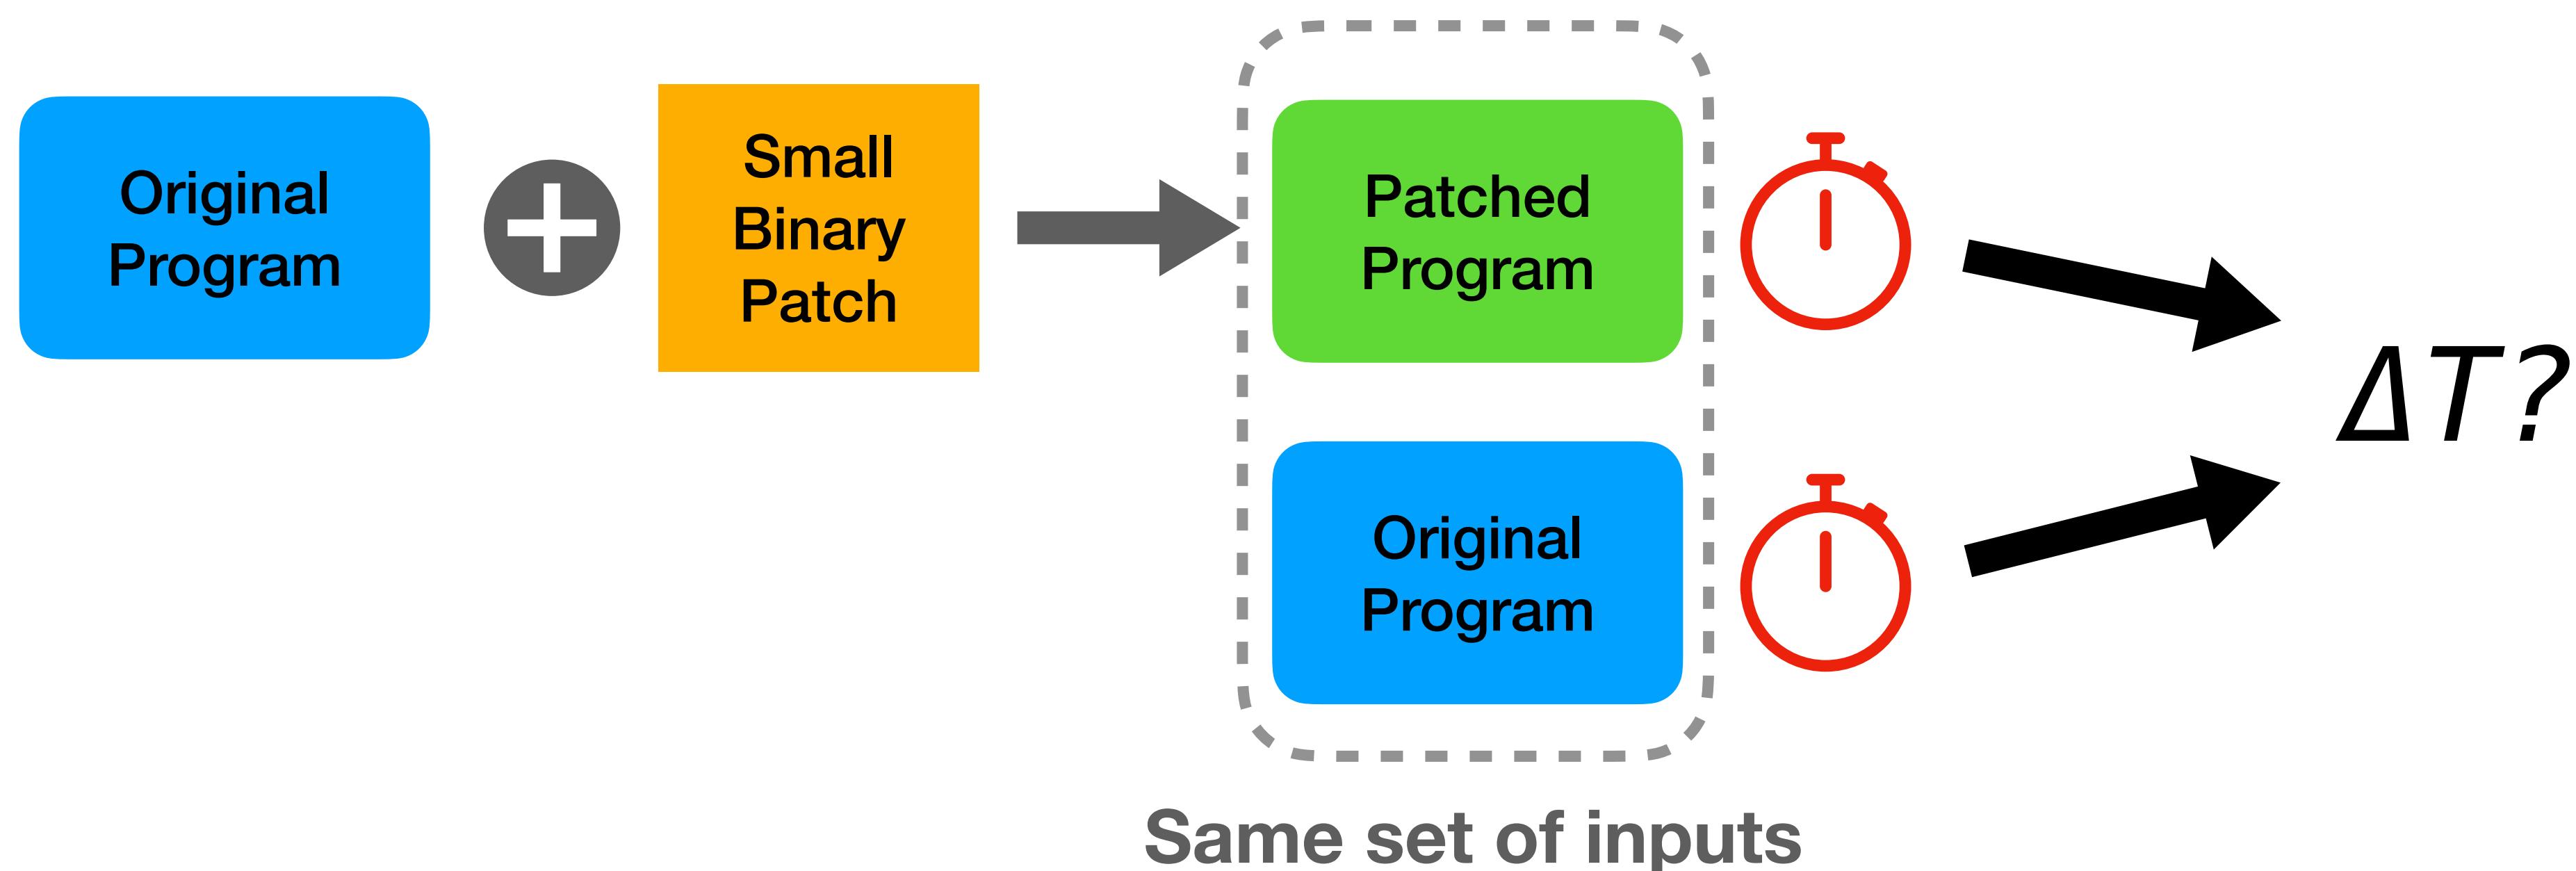
# Timing impacts of binary patches

## Problem definition

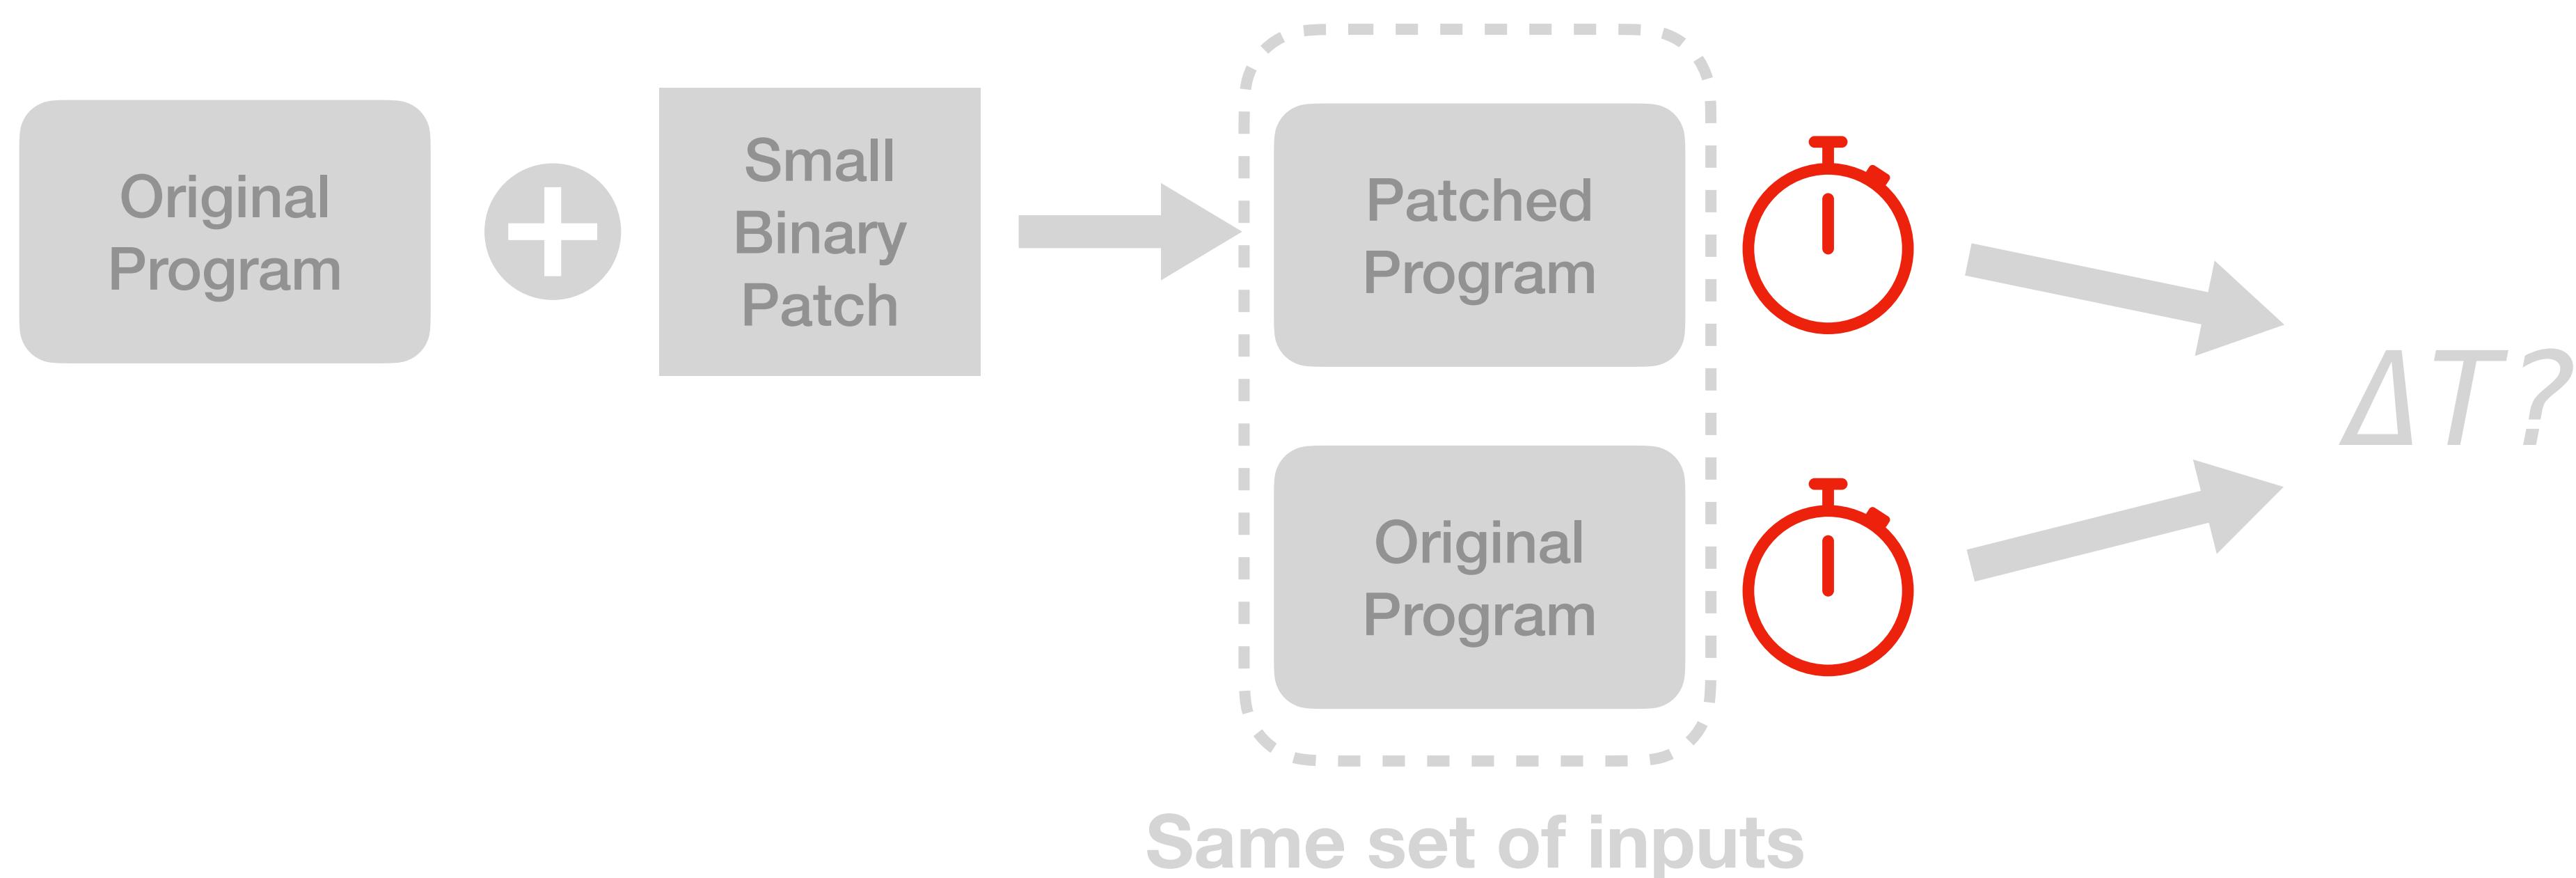


# Timing impacts of binary patches

## Problem definition



# Execution time assessment



# Execution time assessment

## Interesting use cases

# Execution time assessment

## Interesting use cases

- Predicting program run time in **remote** environments or time-sensitive applications
  - Examples: firmware in cars or satellite (e.g. Kepler space telescope by NASA)

# Execution time assessment

## Interesting use cases

- Predicting program run time in **remote** environments or time-sensitive applications
  - Examples: firmware in cars or satellite (e.g. Kepler space telescope by NASA)
- Performance analysis
  - **Insights** into performance bottlenecks

# Execution time assessment

## Interesting use cases

- Predicting program run time in **remote** environments or time-sensitive applications
  - Examples: firmware in cars or satellite (e.g. Kepler space telescope by NASA)
- Performance analysis
  - **Insights** into performance bottlenecks
    - Examples: Potential CPU pipeline stalling, GPU memory bank conflicts

# **Execution time assessment**

## **Previous efforts**

# Execution time assessment

## Previous efforts

- **Static** approaches

# Execution time assessment

## Previous efforts

- **Static** approaches
  - Throughput analysis: predicting the cycle counts for linear code (e.g. basic block, loop) statically

# Execution time assessment

## Previous efforts

- **Static** approaches
  - Throughput analysis: predicting the cycle counts for linear code (e.g. basic block, loop) statically
  - Examples: IACA, OSACA, uiCA, LLVM MCA, Ithemal

# Execution time assessment

## Previous efforts

- **Static** approaches
  - Throughput analysis: predicting the cycle counts for linear code (e.g. basic block, loop) statically
    - Examples: IACA, OSACA, uiCA, LLVM MCA, Ithemal
- **Dynamic** approaches
  - Cycle-accurate simulators / emulators

# Execution time assessment

## Previous efforts

- **Static** approaches
  - Throughput analysis: predicting the cycle counts for linear code (e.g. basic block, loop) statically
    - Examples: IACA, OSACA, uiCA, LLVM MCA, Ithemal
- **Dynamic** approaches
  - Cycle-accurate simulators / emulators
    - Examples: gem5, gpgpu-sim

# **Execution time assessment**

## **Challenges**

**Static**

**Dynamic**

# Execution time assessment

## Challenges



# Execution time assessment

## Challenges



**Static**

**Dynamic**

- Complete execution traces
- Higher fidelity on hardware details

# Execution time assessment

## Challenges



### Static

- Poor handling on branches & function calls
  - Small scope (only few blocks)
- Lack of run-time information

### Dynamic

- Complete execution traces
- Higher fidelity on hardware details

# Execution time assessment

## Challenges

Precision

Low

High

### Static

- Poor handling on branches & function calls
  - Small scope (only few blocks)
- Lack of run-time information

### Dynamic

- Complete execution traces
- Higher fidelity on hardware details

Fast

Slow

Turnaround

# Execution time assessment

## Challenges

Precision

Low

High

### Static

- Poor handling on branches & function calls
  - Small scope (only few blocks)
- Lack of run-time information
- Faster analysis speed (due to coarser granularity)
- Easier integration with other tools

### Dynamic

- Complete execution traces
- Higher fidelity on hardware details

Fast

Slow

Turnaround

# Execution time assessment

## Challenges

Precision

Low

High

### Static

- Poor handling on branches & function calls
  - Small scope (only few blocks)
- Lack of run-time information
- Faster analysis speed (due to coarser granularity)
- Easier integration with other tools

### Dynamic

- Complete execution traces
- Higher fidelity on hardware details
- Usually require non-trivial setup
- Slow simulation speed

Fast

Slow

Turnaround

# Execution time assessment

## Challenges

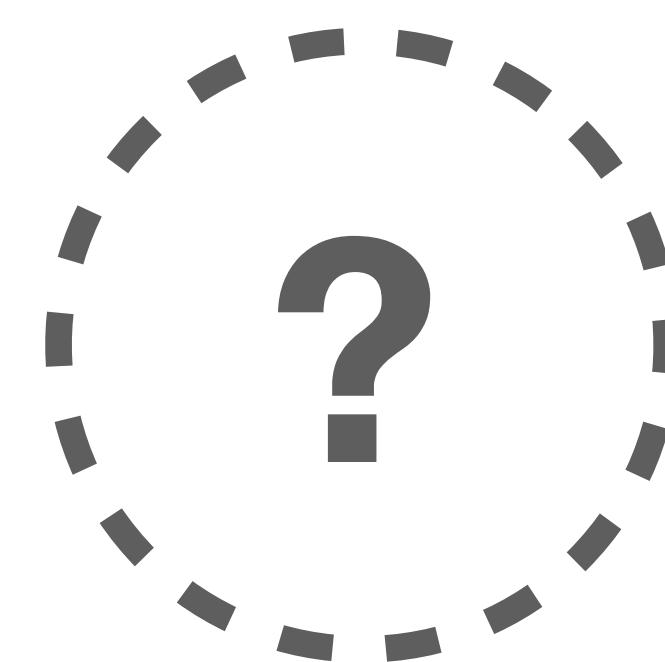
Precision

Low

High

### Static

- Poor handling on branches & function calls
  - Small scope (only few blocks)
- Lack of run-time information
- Faster analysis speed (due to coarser granularity)
- Easier integration with other tools



### Dynamic

- Complete execution traces
- Higher fidelity on hardware details
- Usually require non-trivial setup
- Slow simulation speed

Fast

Slow

Turnaround

# Outline

Motivation

MCA Daemon (MCAD)

Future Plans

Epilogue

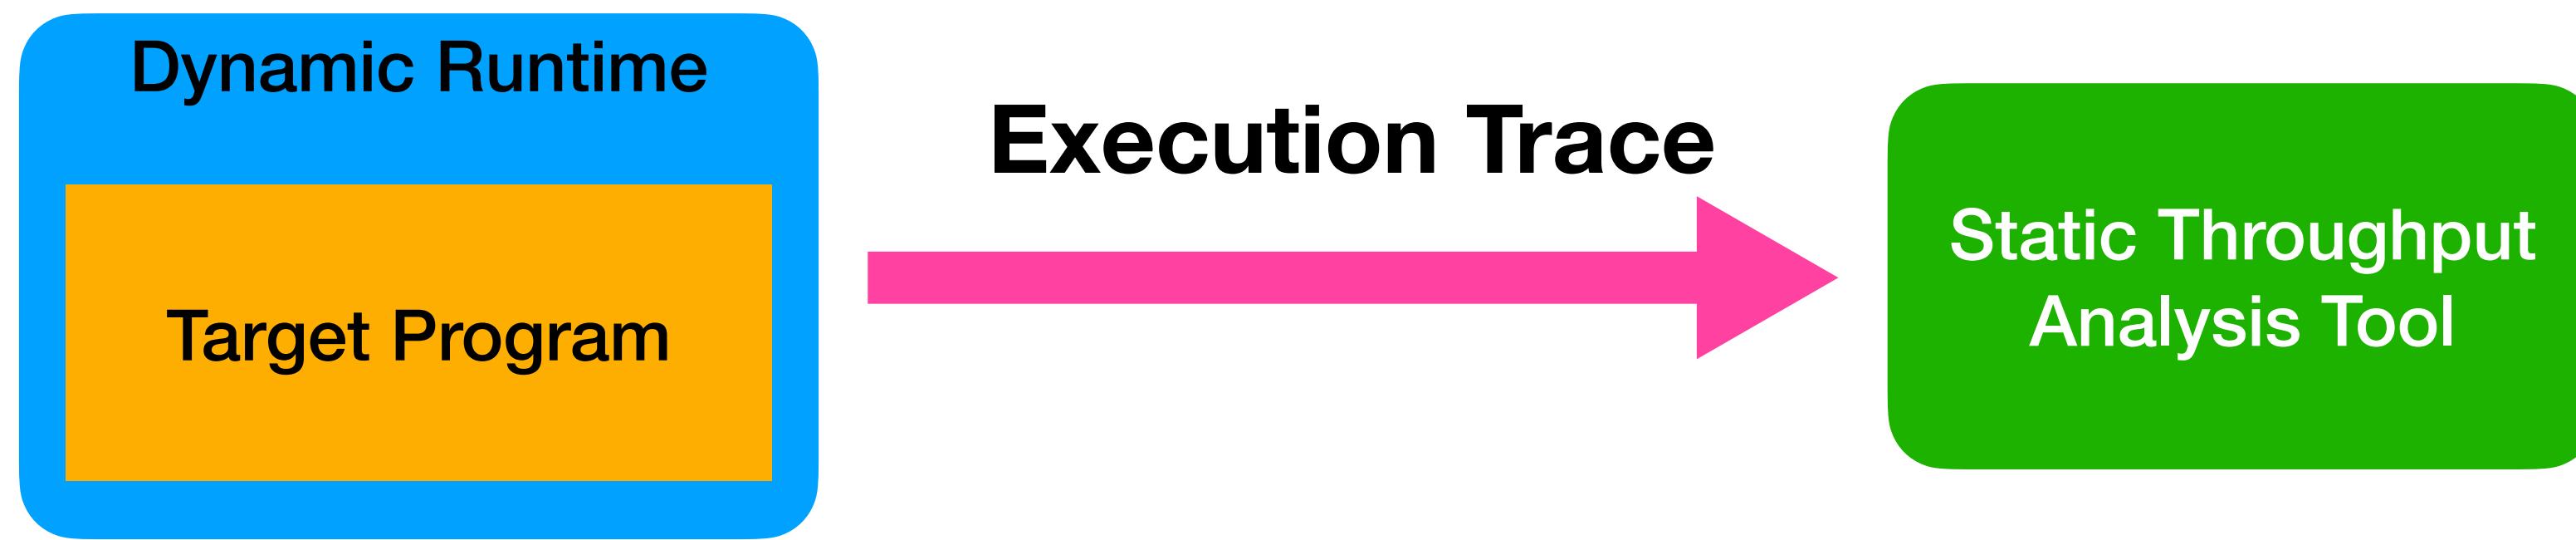
# MCA Daemon (MCAD)

## High-level concept



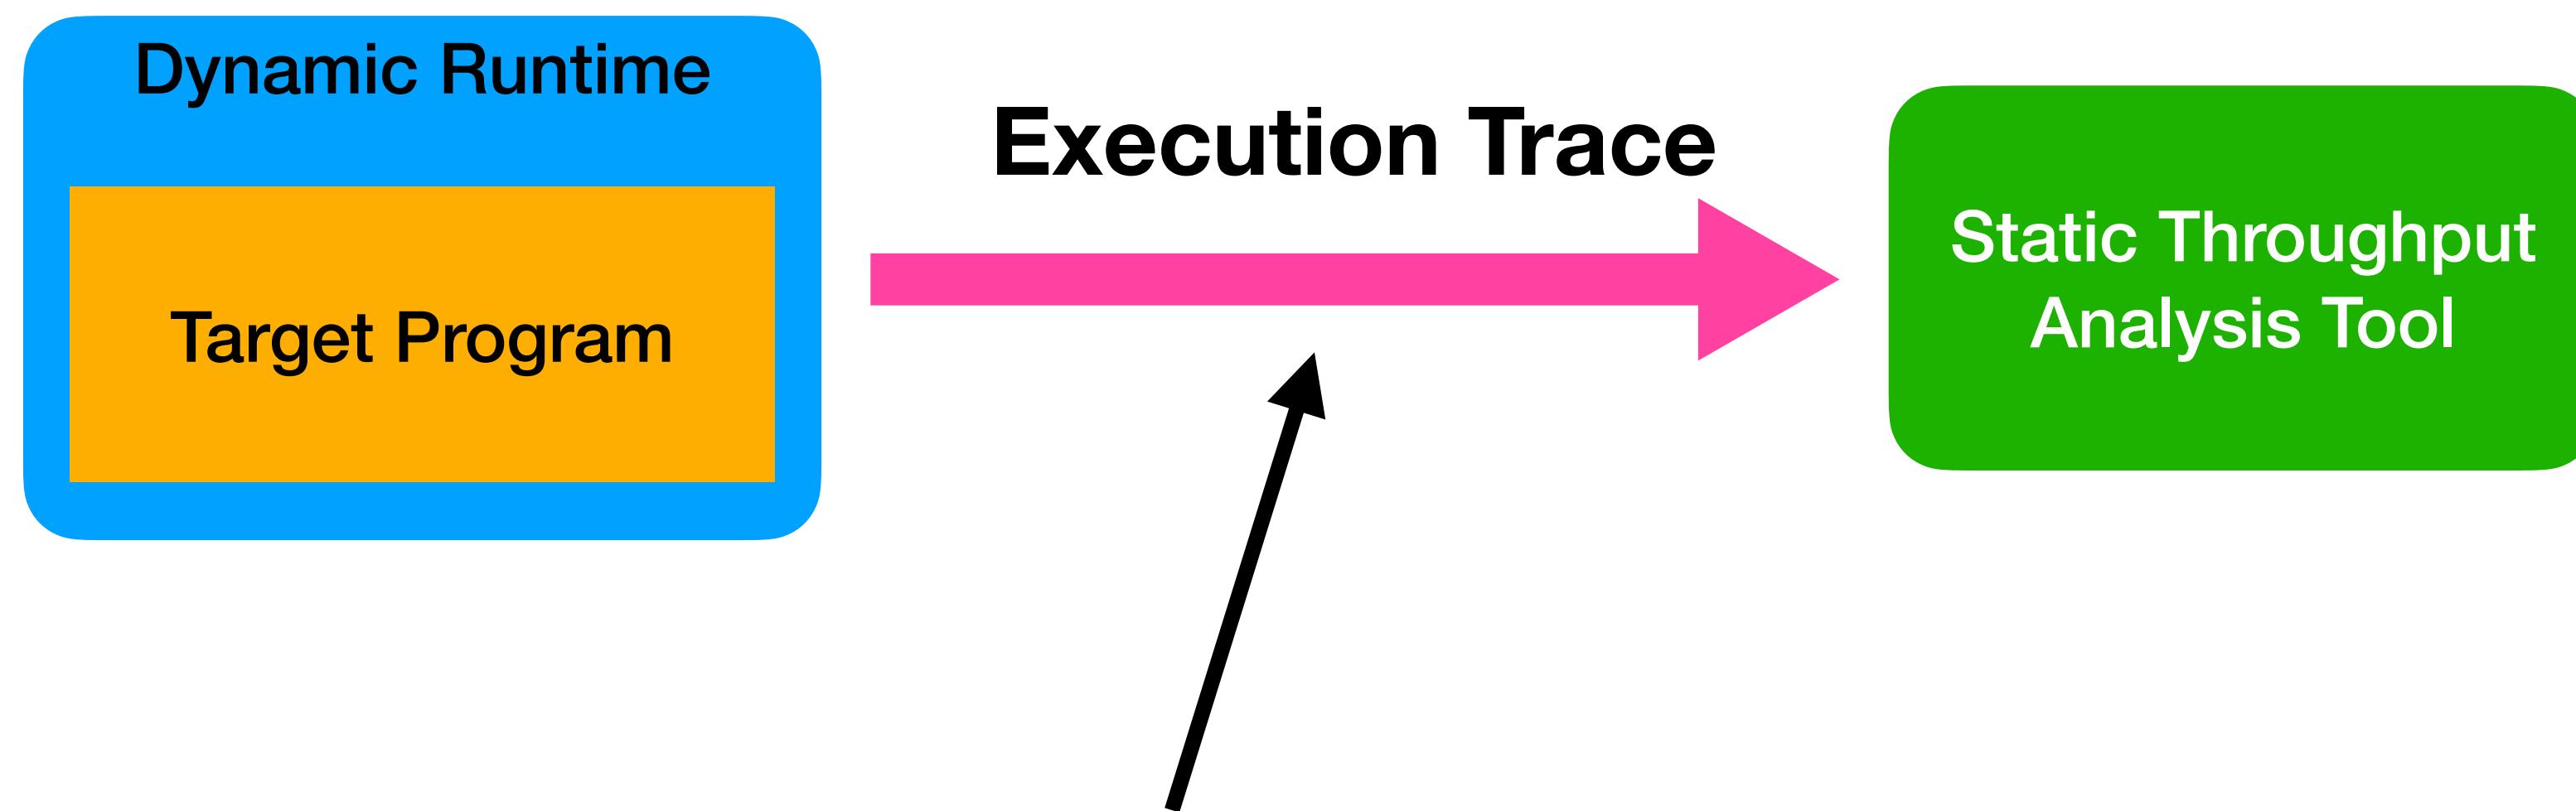
# MCA Daemon (MCAD)

## High-level concept



# MCA Daemon (MCAD)

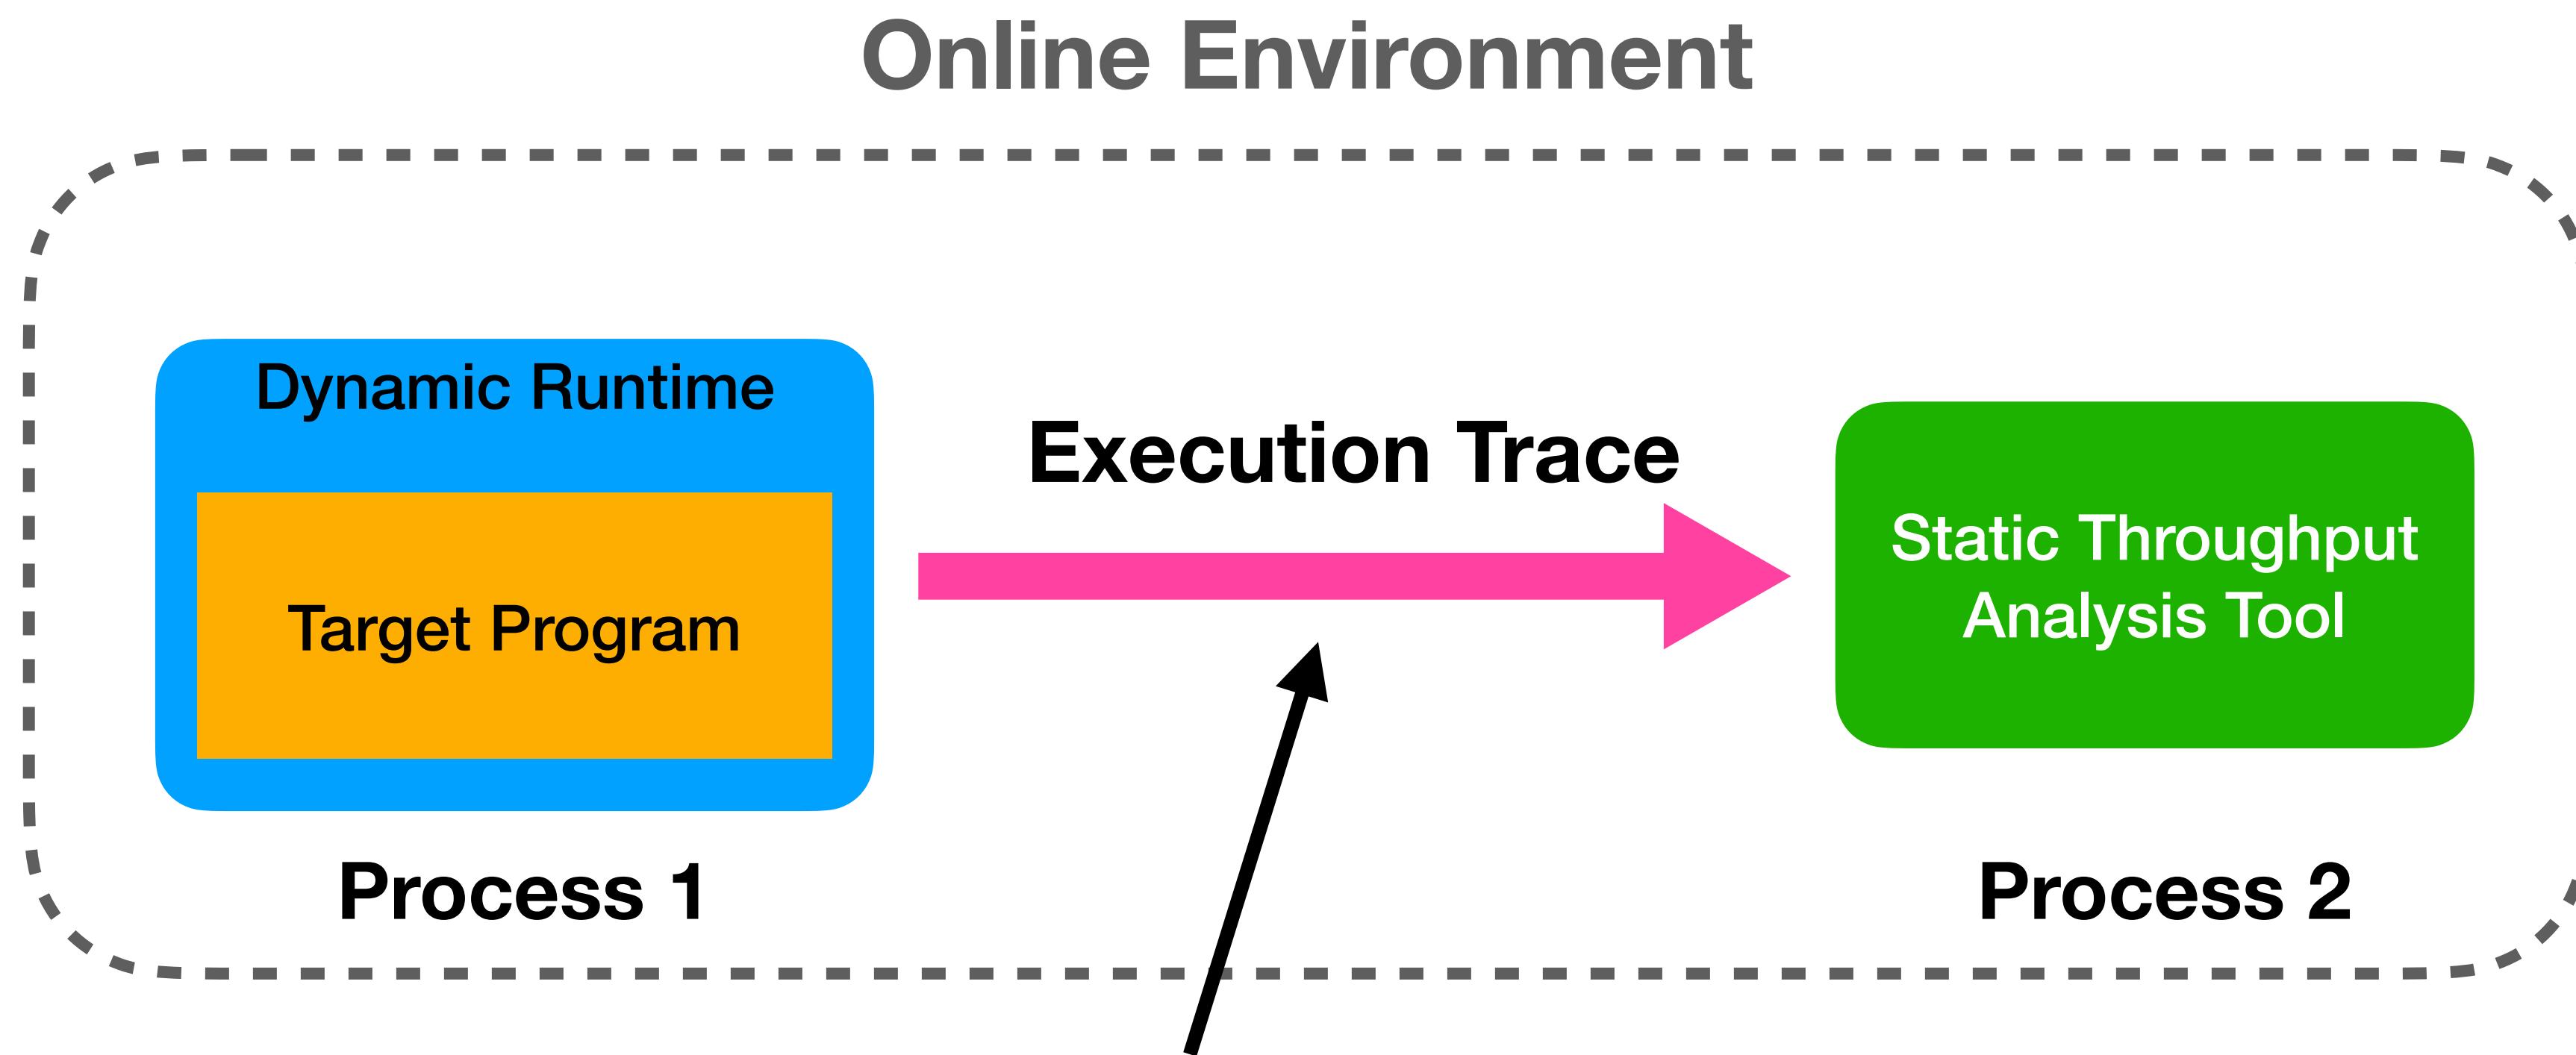
## High-level concept



- The instructions that just got executed
- Run-time values (e.g. register values)

# MCA Daemon (MCAD)

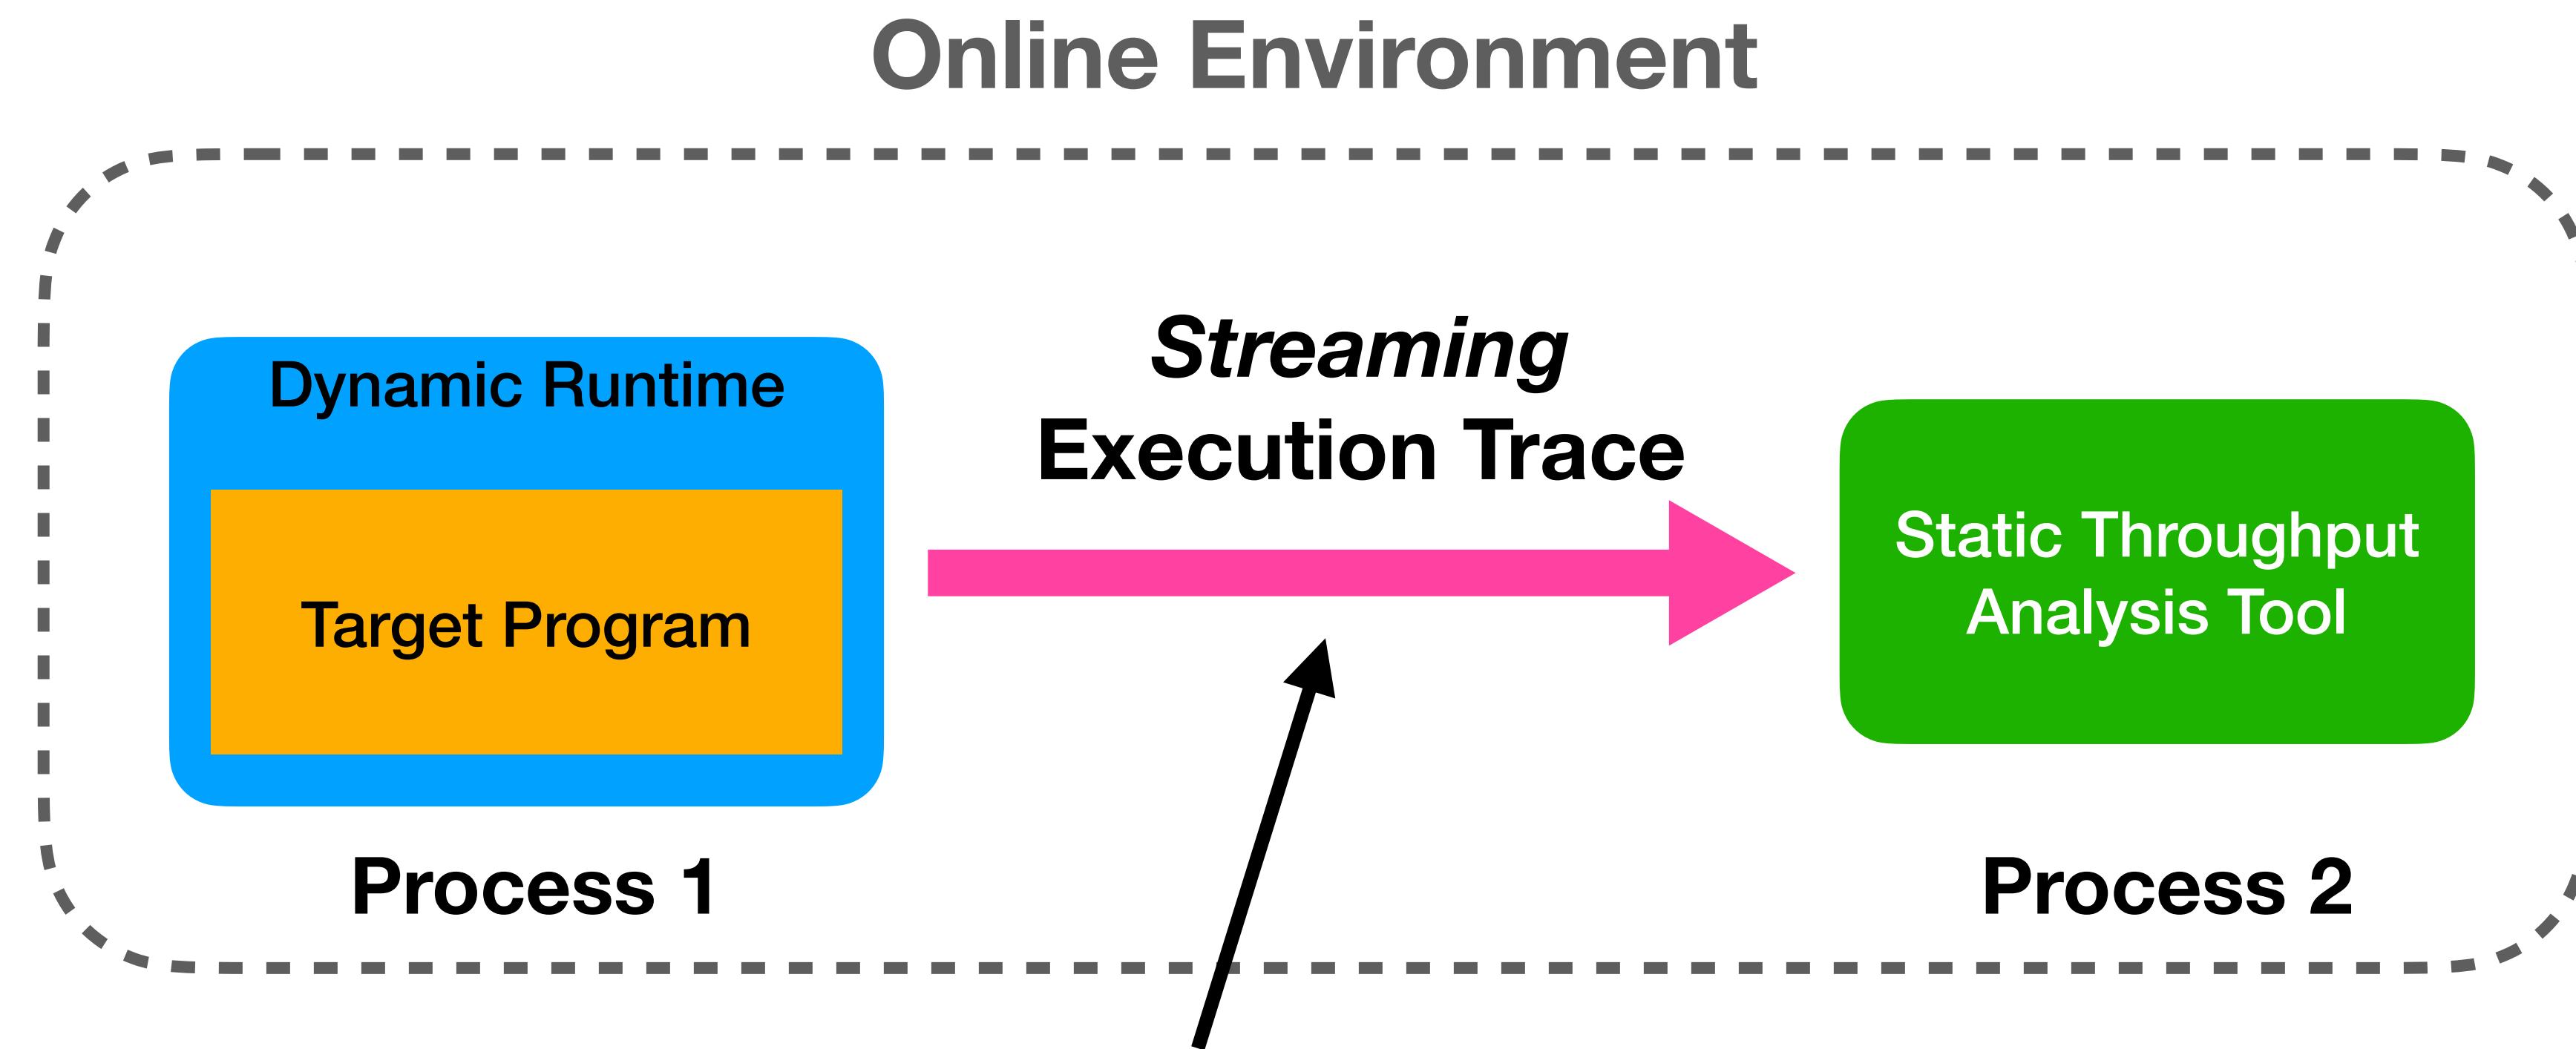
## High-level concept



- The instructions that just got executed
- Run-time values (e.g. register values)

# MCA Daemon (MCAD)

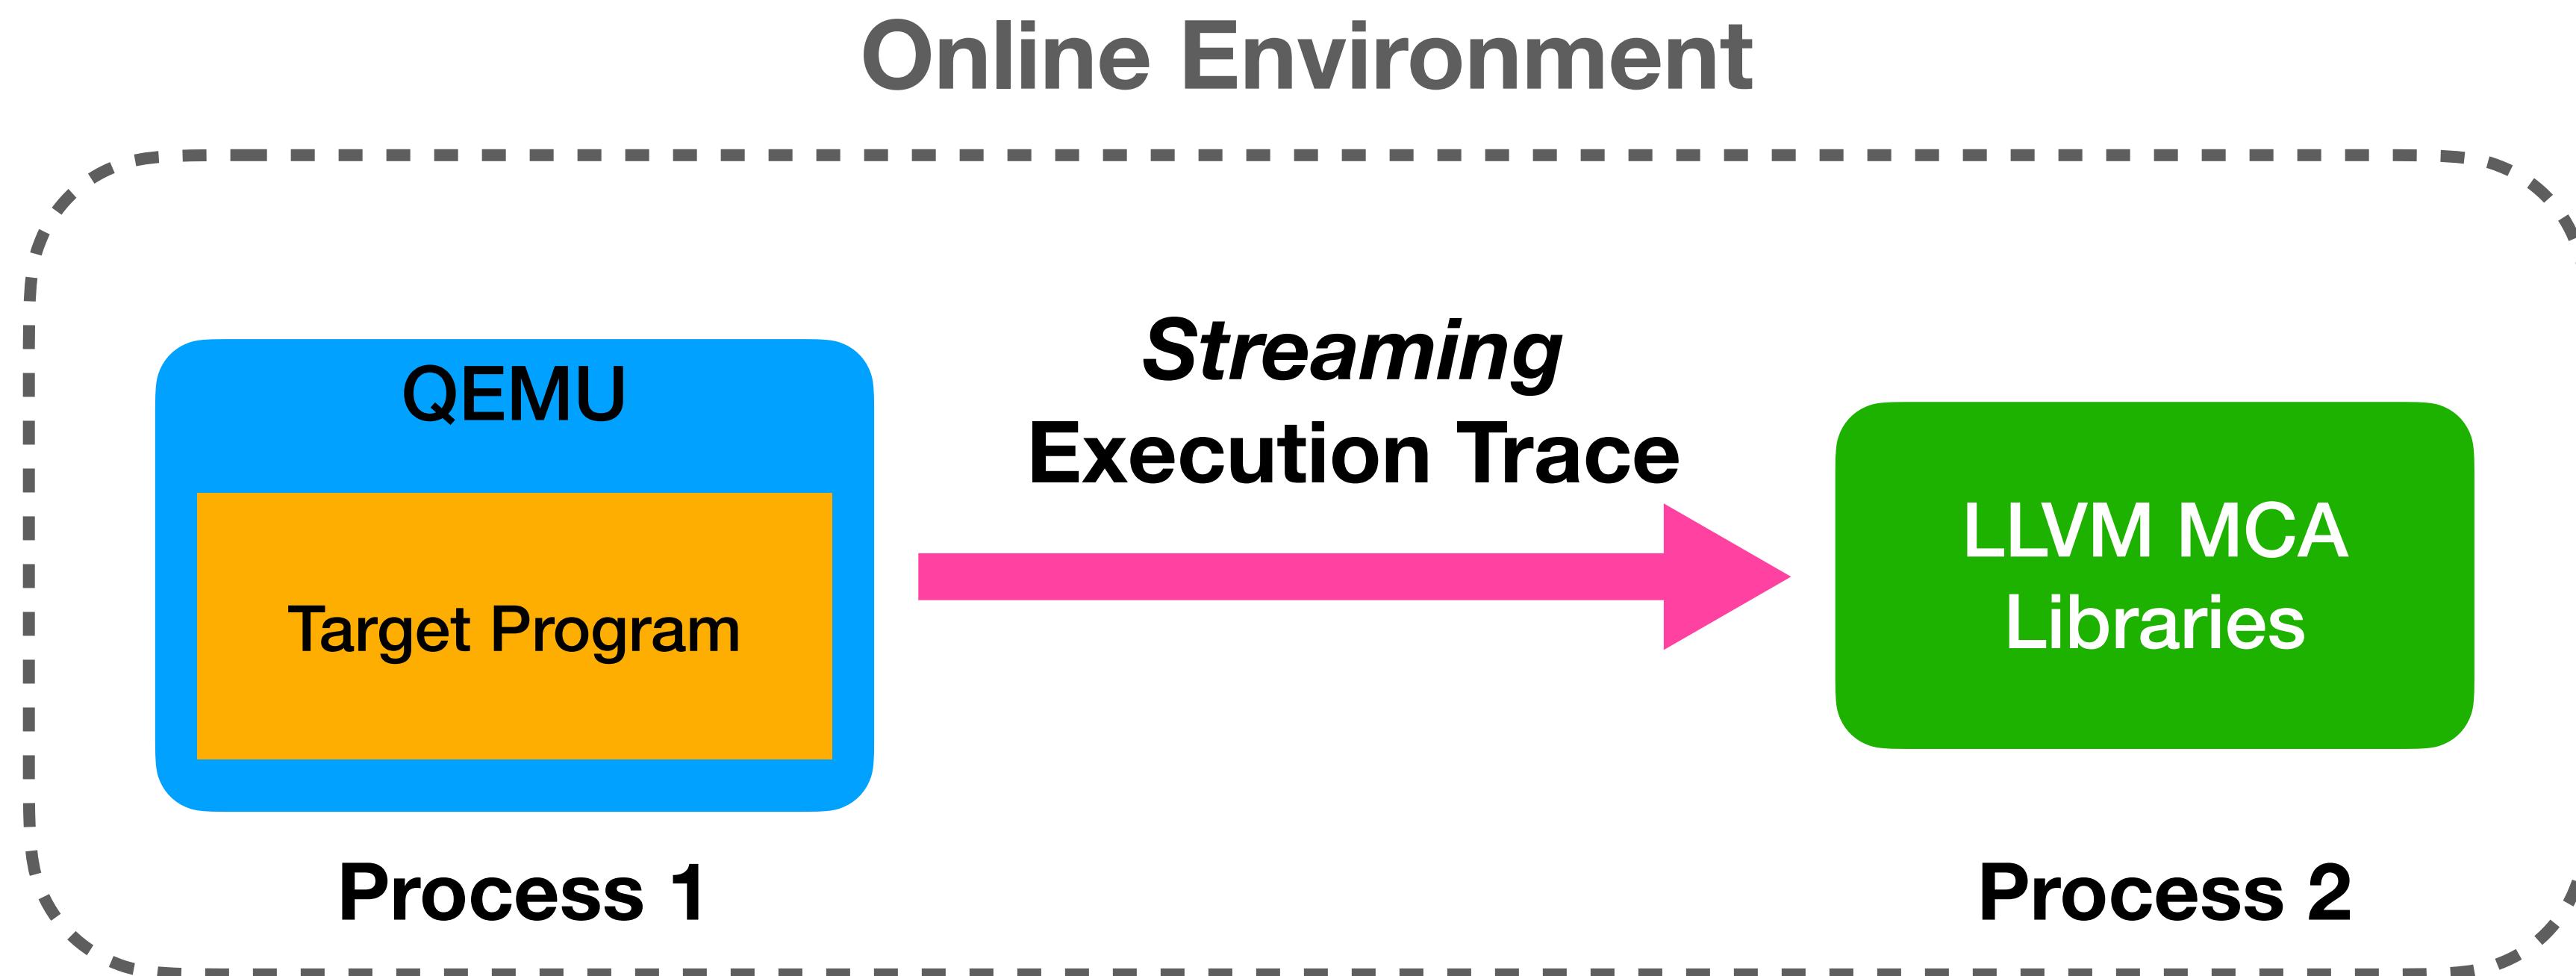
## High-level concept



- The instructions that just got executed
- Run-time values (e.g. register values)

# MCA Daemon (MCAD)

## High-level concept



# Introduction to LLVM MCA

# Introduction to LLVM MCA

- A tool (`llvm-mca`) and **library** (`libLLVMMCA`) for predicting cycle counts and potential performance hazards in a sequence of assembly code

# Introduction to LLVM MCA

- A tool (`llvm-mca`) and **library** (`libLLVMMCA`) for predicting cycle counts and potential performance hazards in a sequence of assembly code
- Using instruction scheduling data (e.g. instruction latency) provided by each LLVM target
  - New ISA (with proper scheduling info) can be supported *out of the box*

# Introduction to LLVM MCA

- A tool (`llvm-mca`) and **library** (`libLLVMMCA`) for predicting cycle counts and potential performance hazards in a sequence of assembly code
- Using instruction scheduling data (e.g. instruction latency) provided by each LLVM target
  - New ISA (with proper scheduling info) can be supported *out of the box*
  - Accounting for modern processor features: super scalar, out-of-order etc.

# Introduction to LLVM MCA

- A tool (`llvm-mca`) and **library** (`libLLVMMCA`) for predicting cycle counts and potential performance hazards in a sequence of assembly code
- Using instruction scheduling data (e.g. instruction latency) provided by each LLVM target
  - New ISA (with proper scheduling info) can be supported *out of the box*
  - Accounting for modern processor features: super scalar, out-of-order etc.
  - Implemented via lightweight simulation
    - Abstract real CPU pipeline stages into a small handful of stages

# Introduction to MCA

## An example

`test/tools/llvm-mca/X86/BtVer2/dot-product.s`

```
vmulps    %xmm0,  %xmm1,  %xmm2
vhaddps   %xmm2,  %xmm2,  %xmm3
vhaddps   %xmm3,  %xmm3,  %xmm4
```

# Introduction to MCA

## An example

`test/tools/llvm-mca/X86/BtVer2/dot-products.s`

```
vmulps    %xmm0,  %xmm1,  %xmm2
vhaddps   %xmm2,  %xmm2,  %xmm3
vhaddps   %xmm3,  %xmm3,  %xmm4
```

```
llvm-mca -mtriple=x86_64 -mcpu=btver2 \
-iterations=300 dot-products.s
```

# Introduction to MCA

## An example

test/tools/llvm-mca/X86/BtVer2/dot-products.s

```
vmulps    %xmm0,  %xmm1,  %xmm2  
vhaddps   %xmm2,  %xmm2,  %xmm3  
vhaddps   %xmm3,  %xmm3,  %xmm4
```

```
llvm-mca -mtriple=x86_64 -mcpu=btver2 \  
-iterations=300 dot-products.s
```

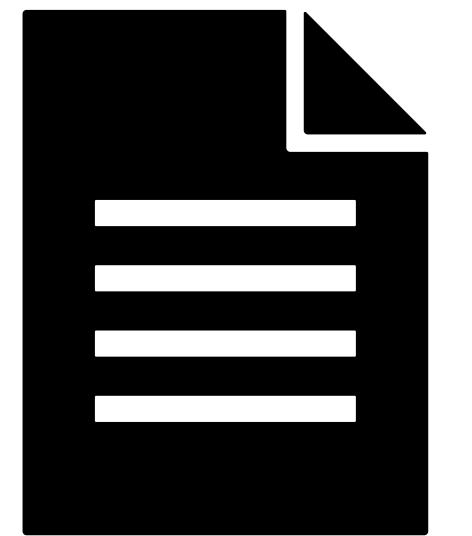
### Summary

Iterations:	300
Instructions:	900
Total Cycles:	610
Total uOps:	900
Dispatch Width:	2
uOps Per Cycle:	1.48
IPC:	1.48
Block RThroughput:	2.0

**llvm-mca**

**MCAD**

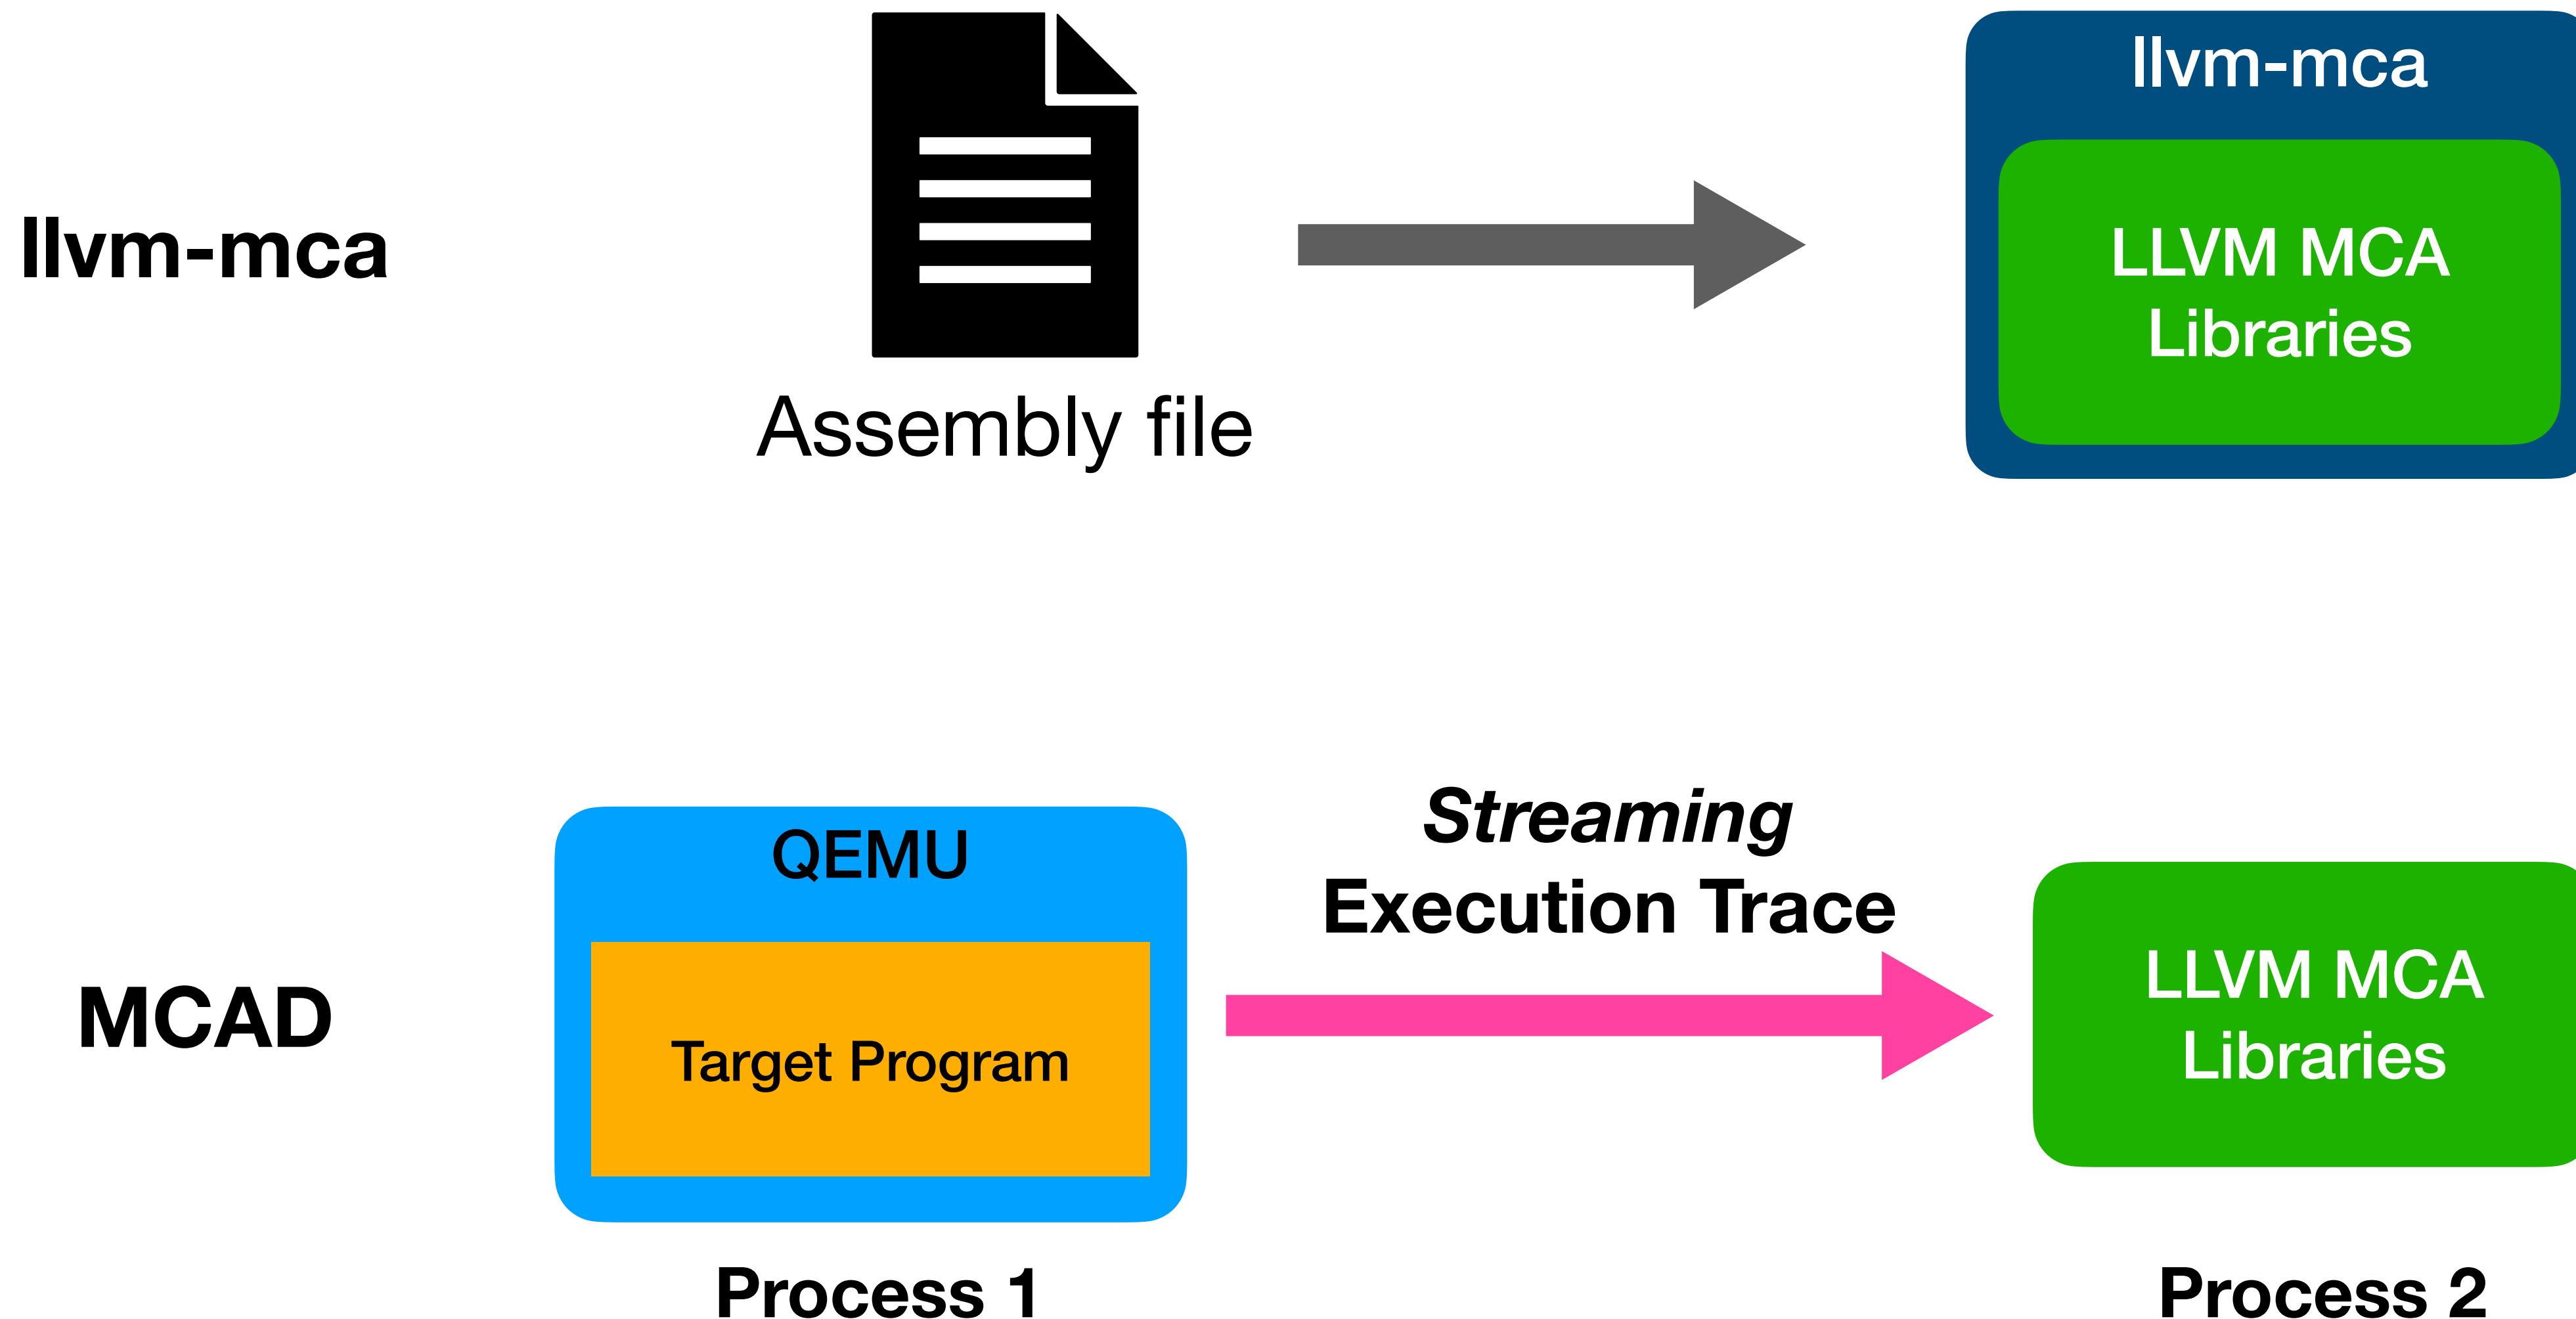
**llvm-mca**



Assembly file



**MCAD**



# MCA Daemon (MCAD)

## Highlights

# MCA Daemon (MCAD)

## Highlights

- Combine the advantages of dynamic & static throughput analysis

# MCA Daemon (MCAD)

## Highlights

- Combine the advantages of dynamic & static throughput analysis
- Augment the analysis region *beyond* basic blocks
  - MCAD is able to analyze the **entire** program execution trace

# MCA Daemon (MCAD)

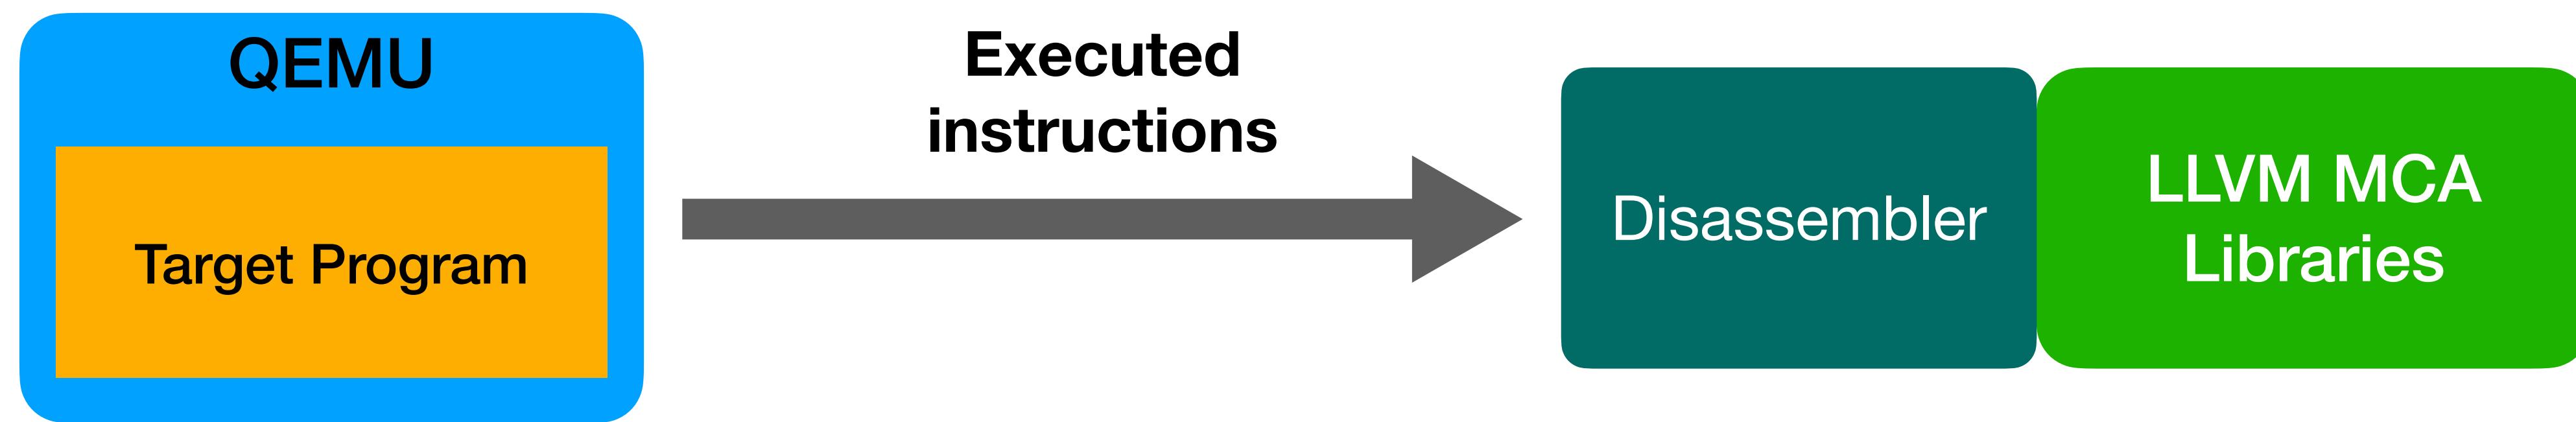
## Highlights

- Combine the advantages of dynamic & static throughput analysis
- Augment the analysis region *beyond* basic blocks
  - MCAD is able to analyze the **entire** program execution trace
- Throughput analysis is happening **in parallel / on-the-fly** with the target program execution

# Implementation

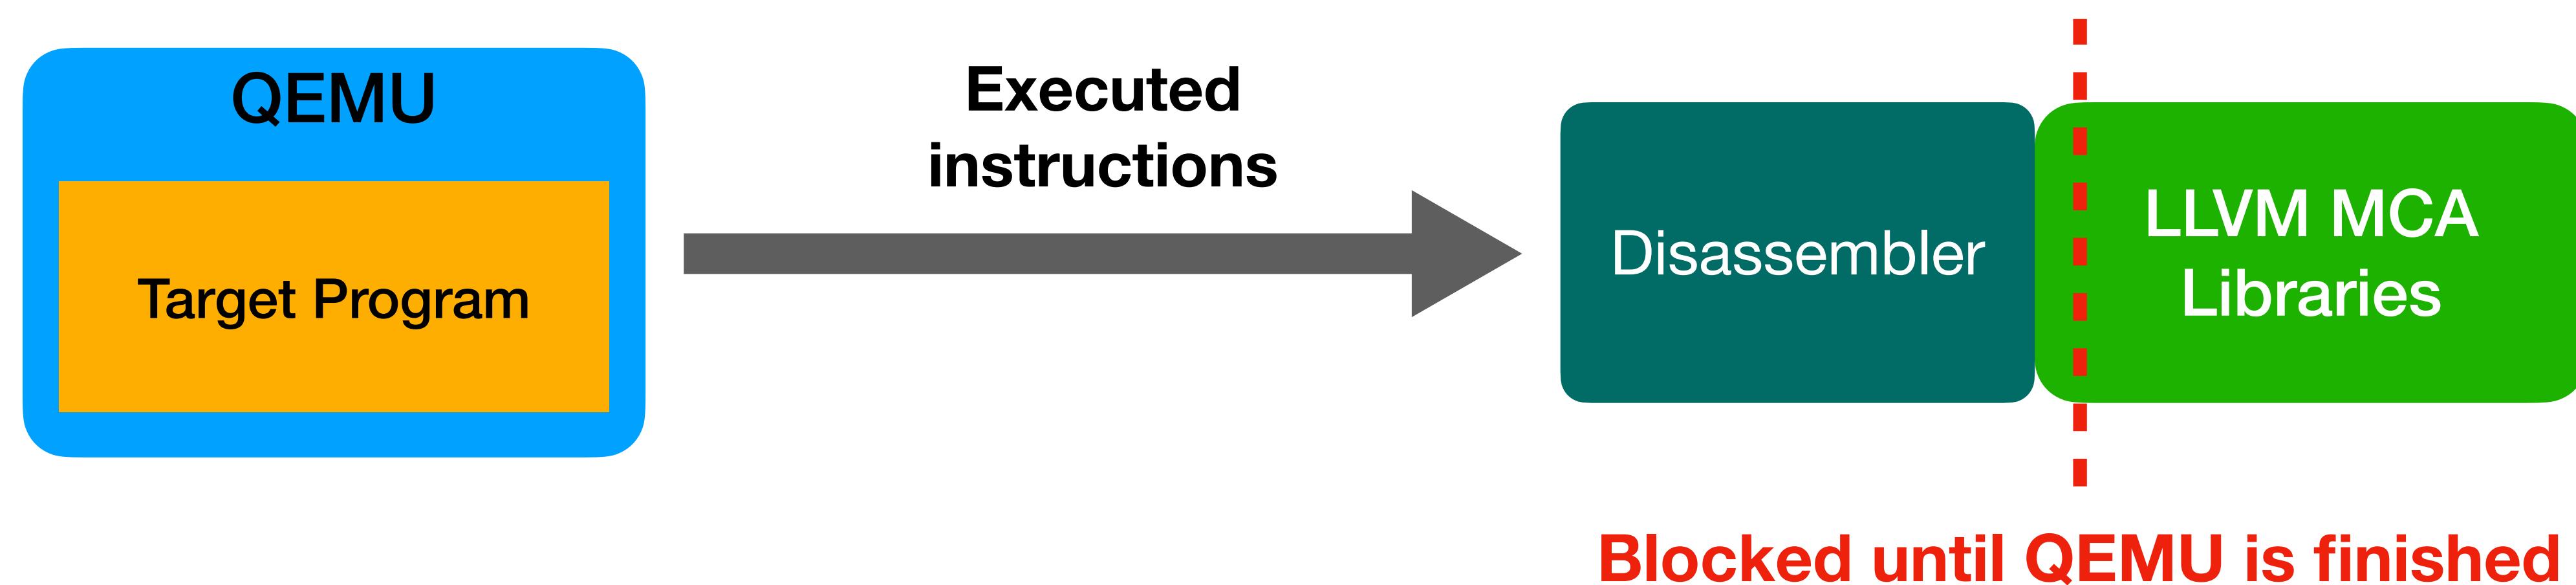
# Analyze execution traces using MCA

## Using unmodified MCA libraries

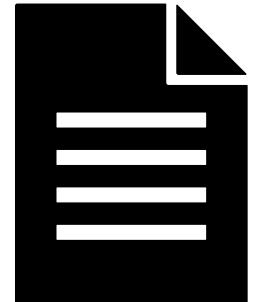


# Analyze execution traces using MCA

## Challenge: Sequential workflow

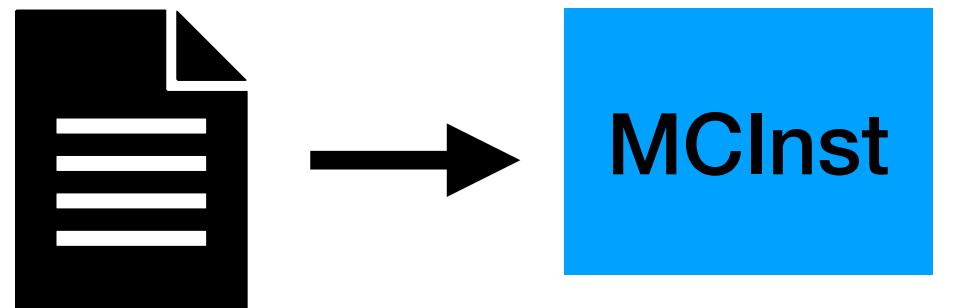


# MCA internal



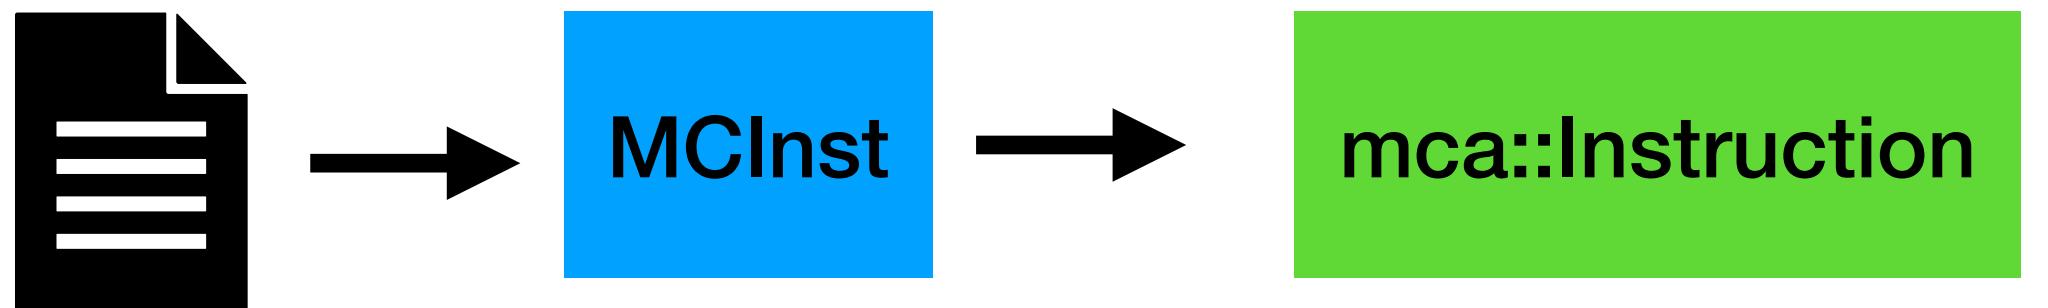
**Assembly file**

# MCA internal



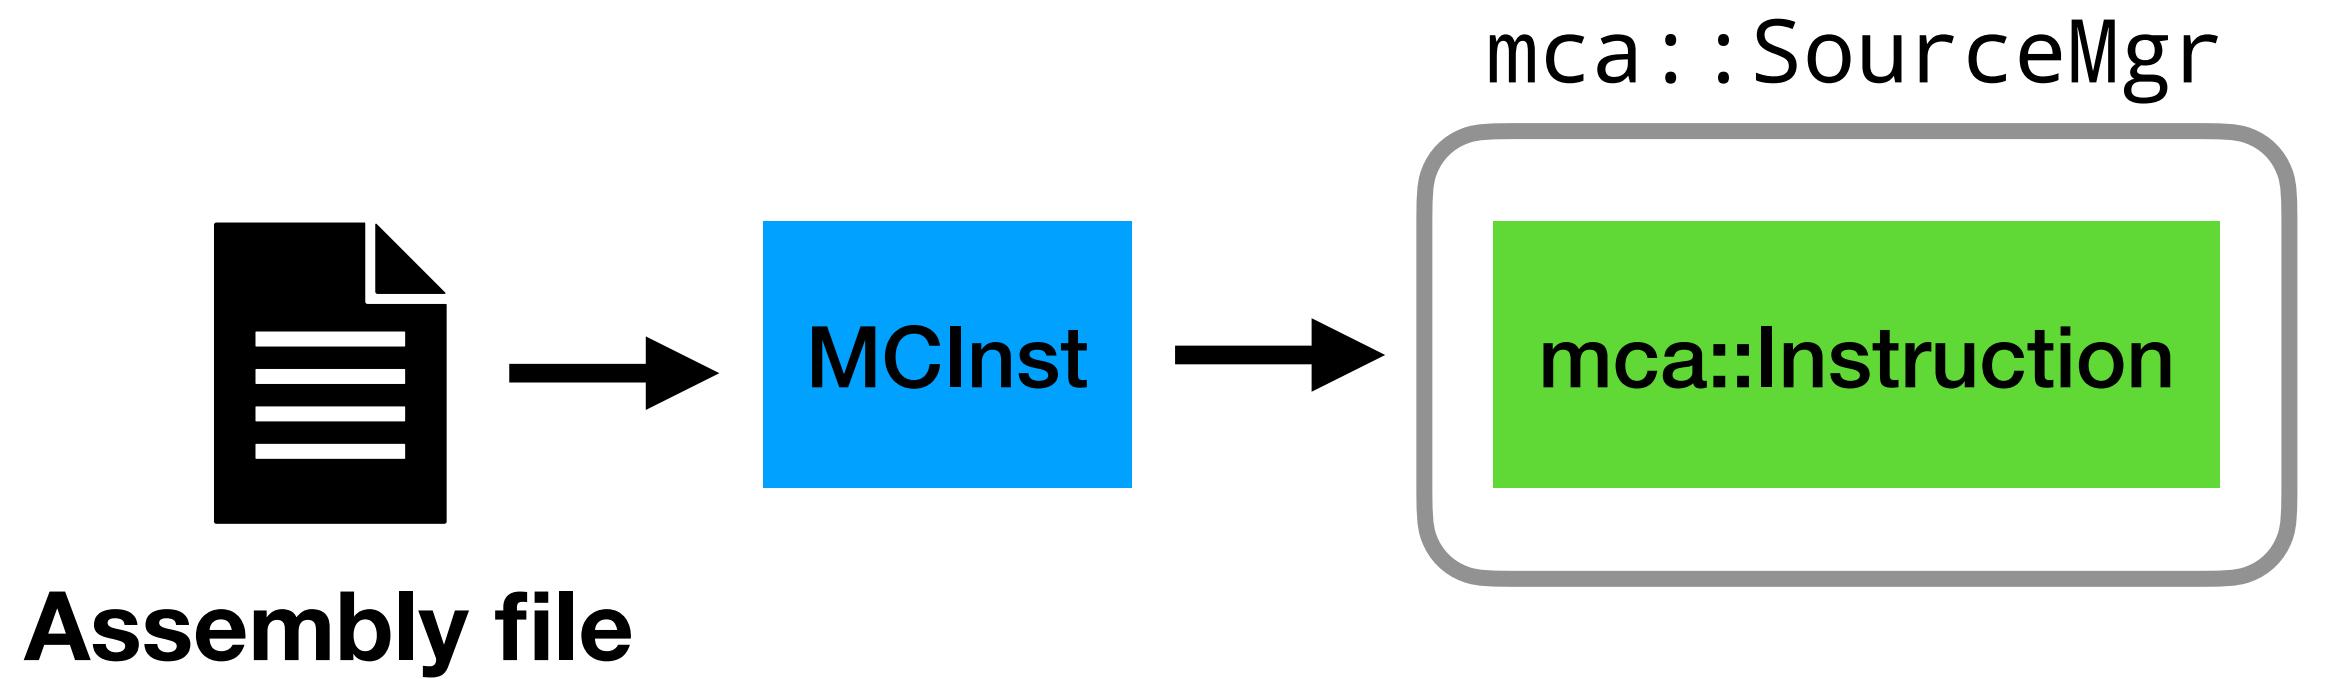
**Assembly file**

# MCA internal

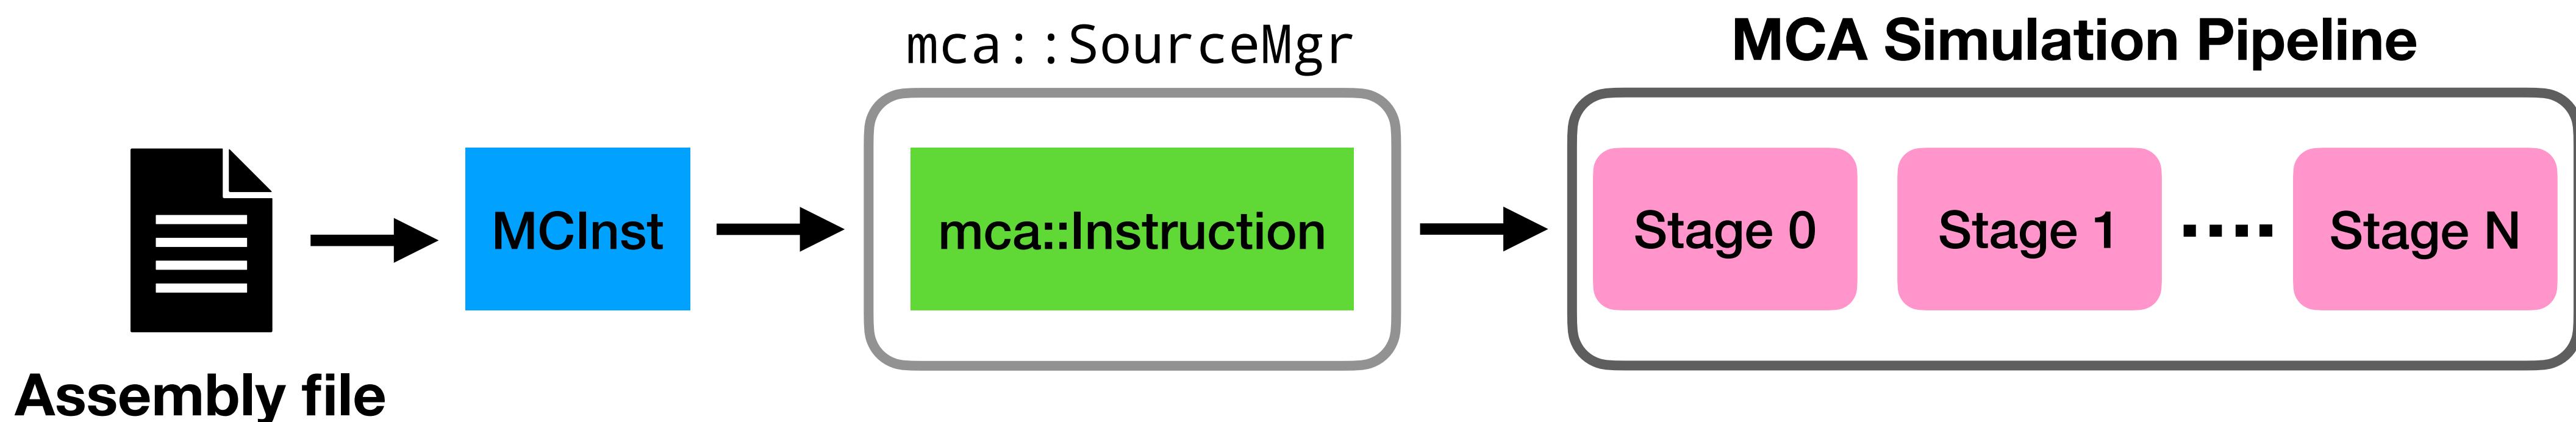


**Assembly file**

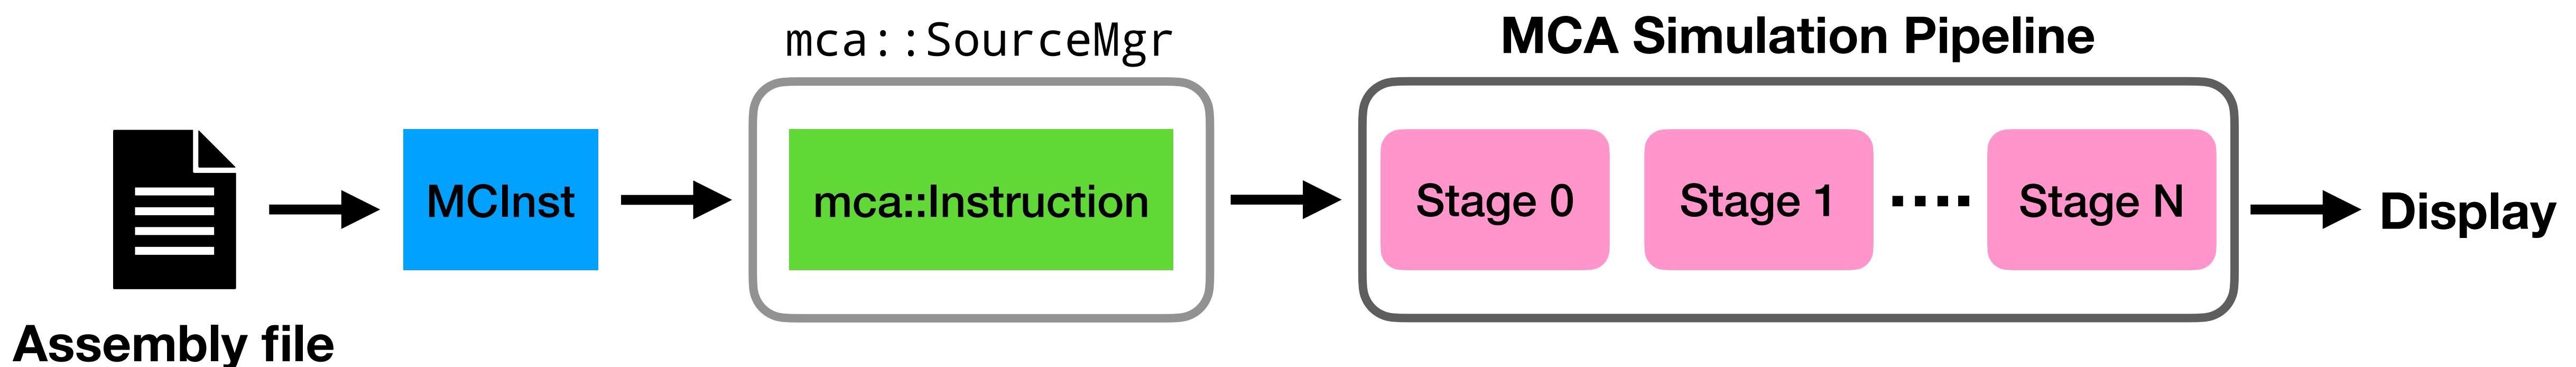
# MCA internal



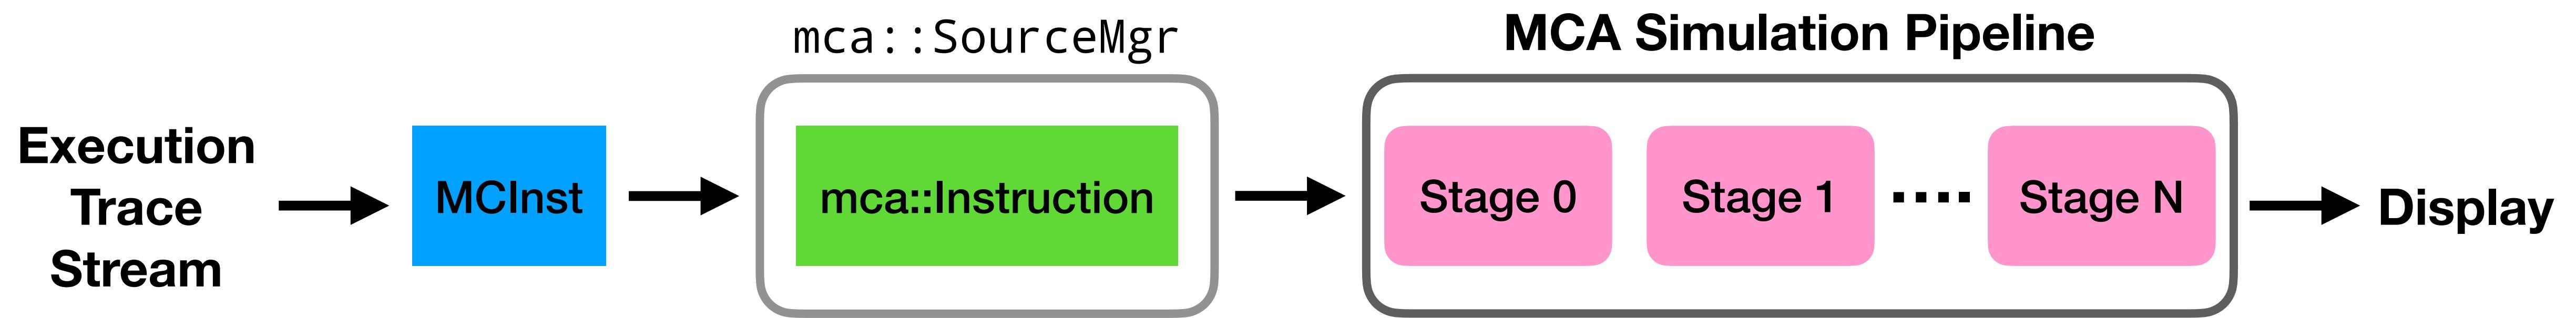
# MCA internal



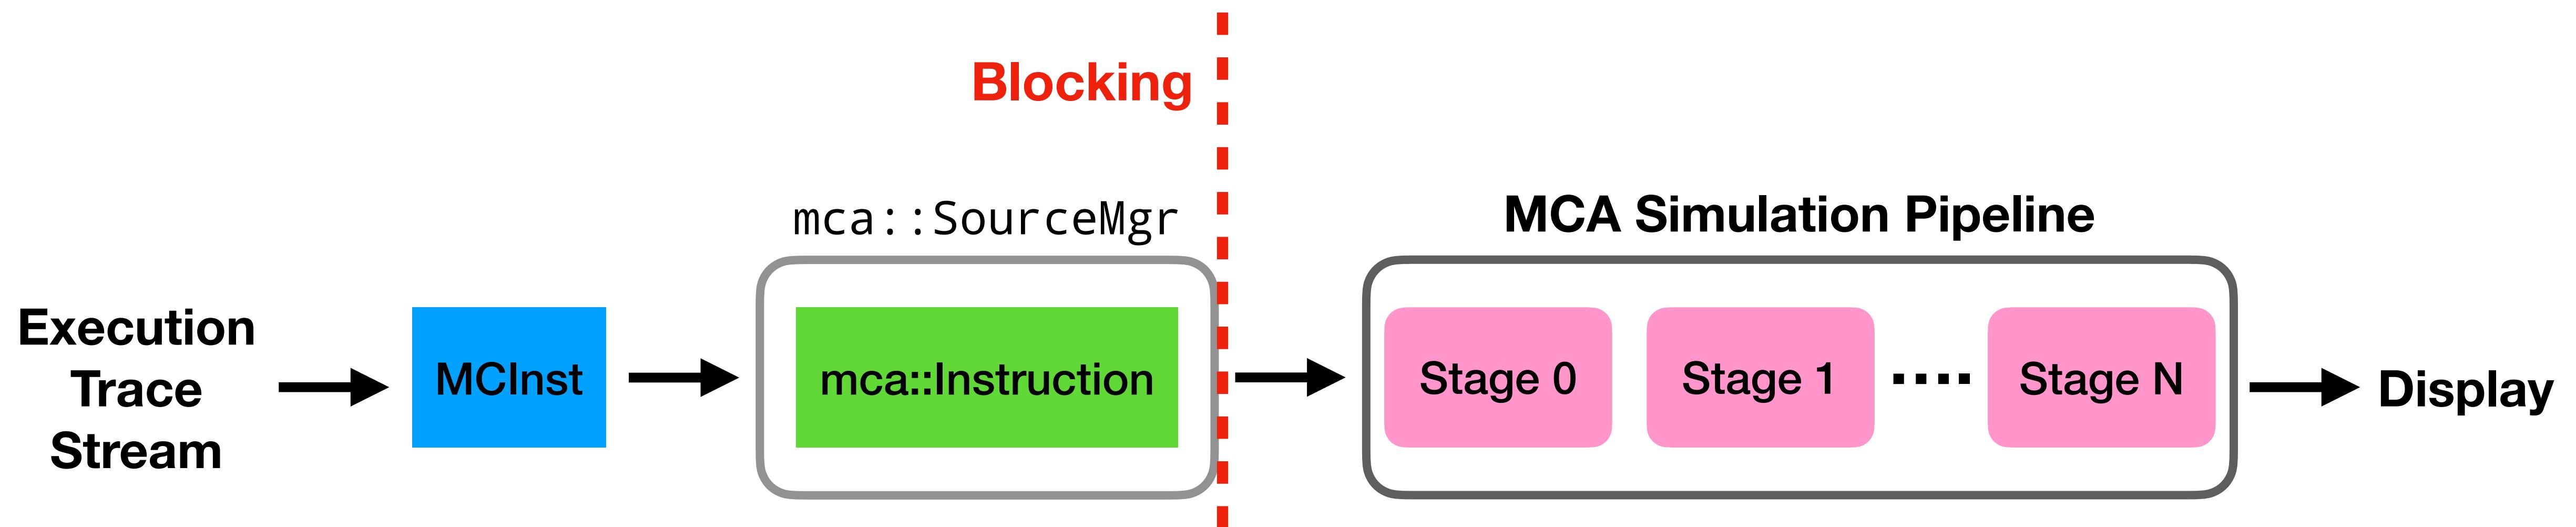
# MCA internal



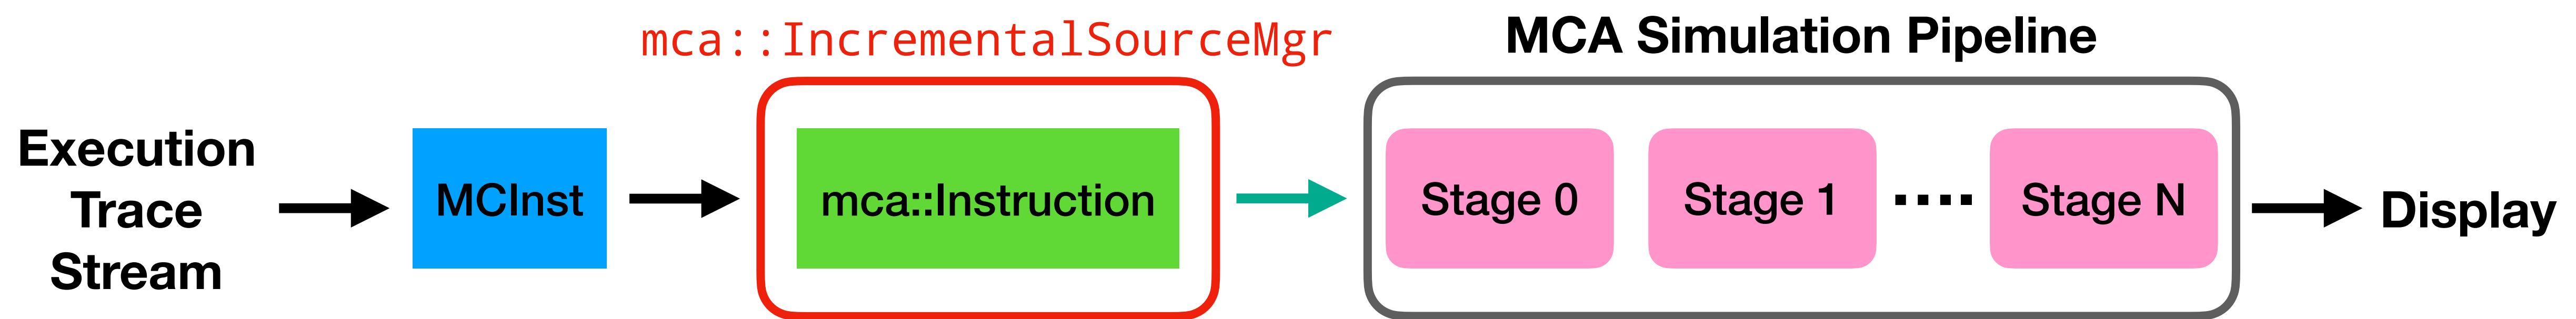
# MCA with execution trace stream as input



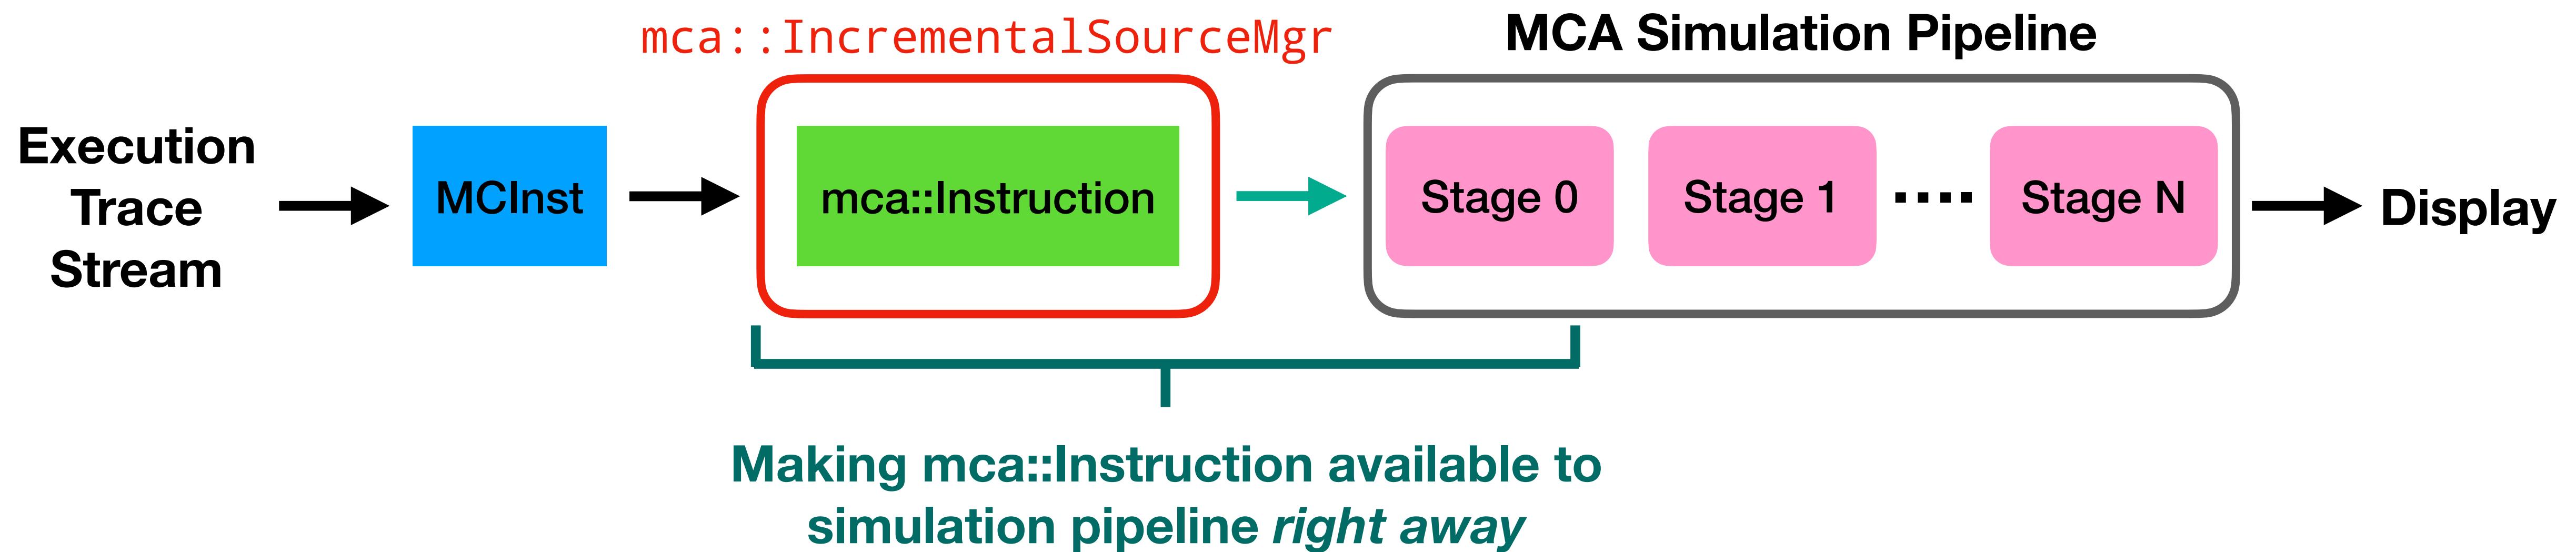
# MCA with execution trace stream as input



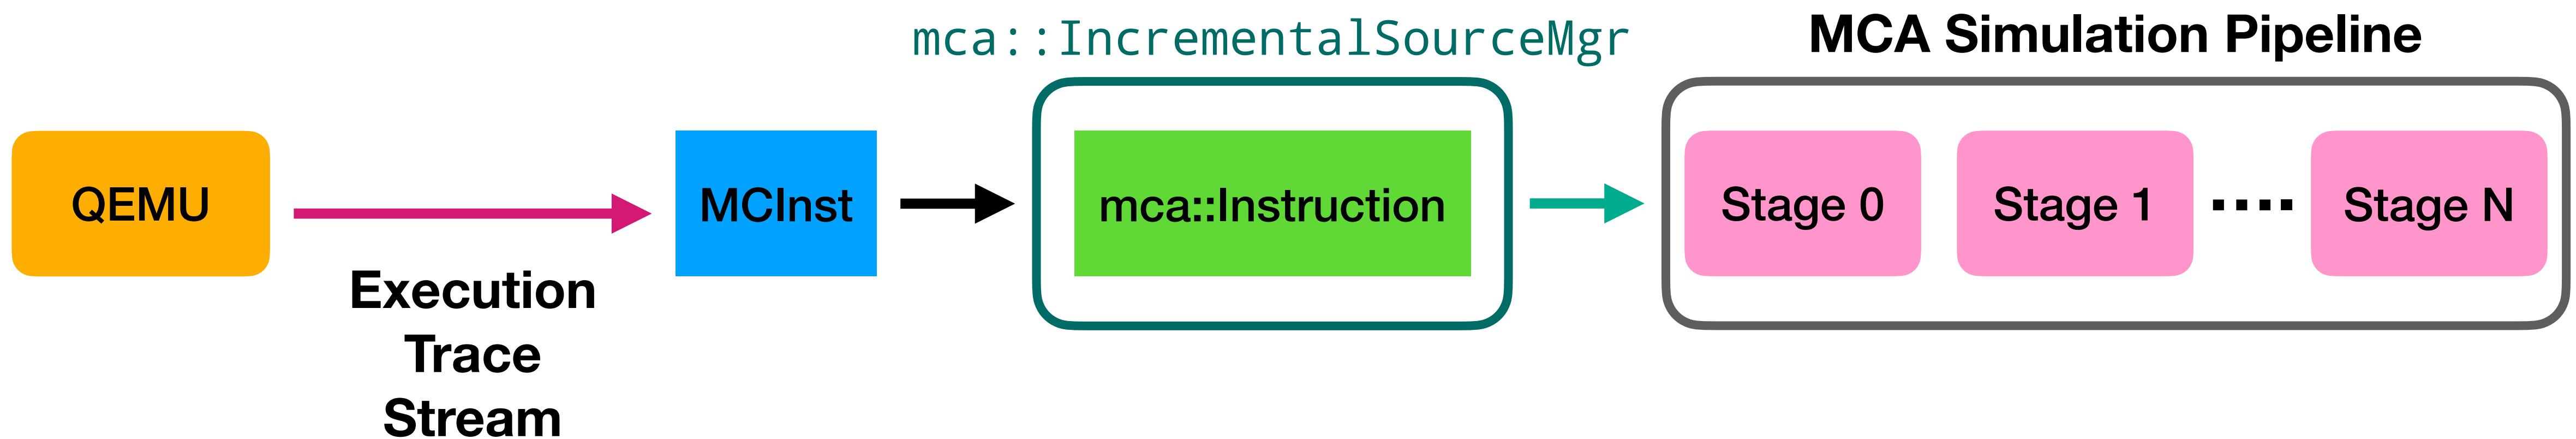
# Incremental SourceMgr



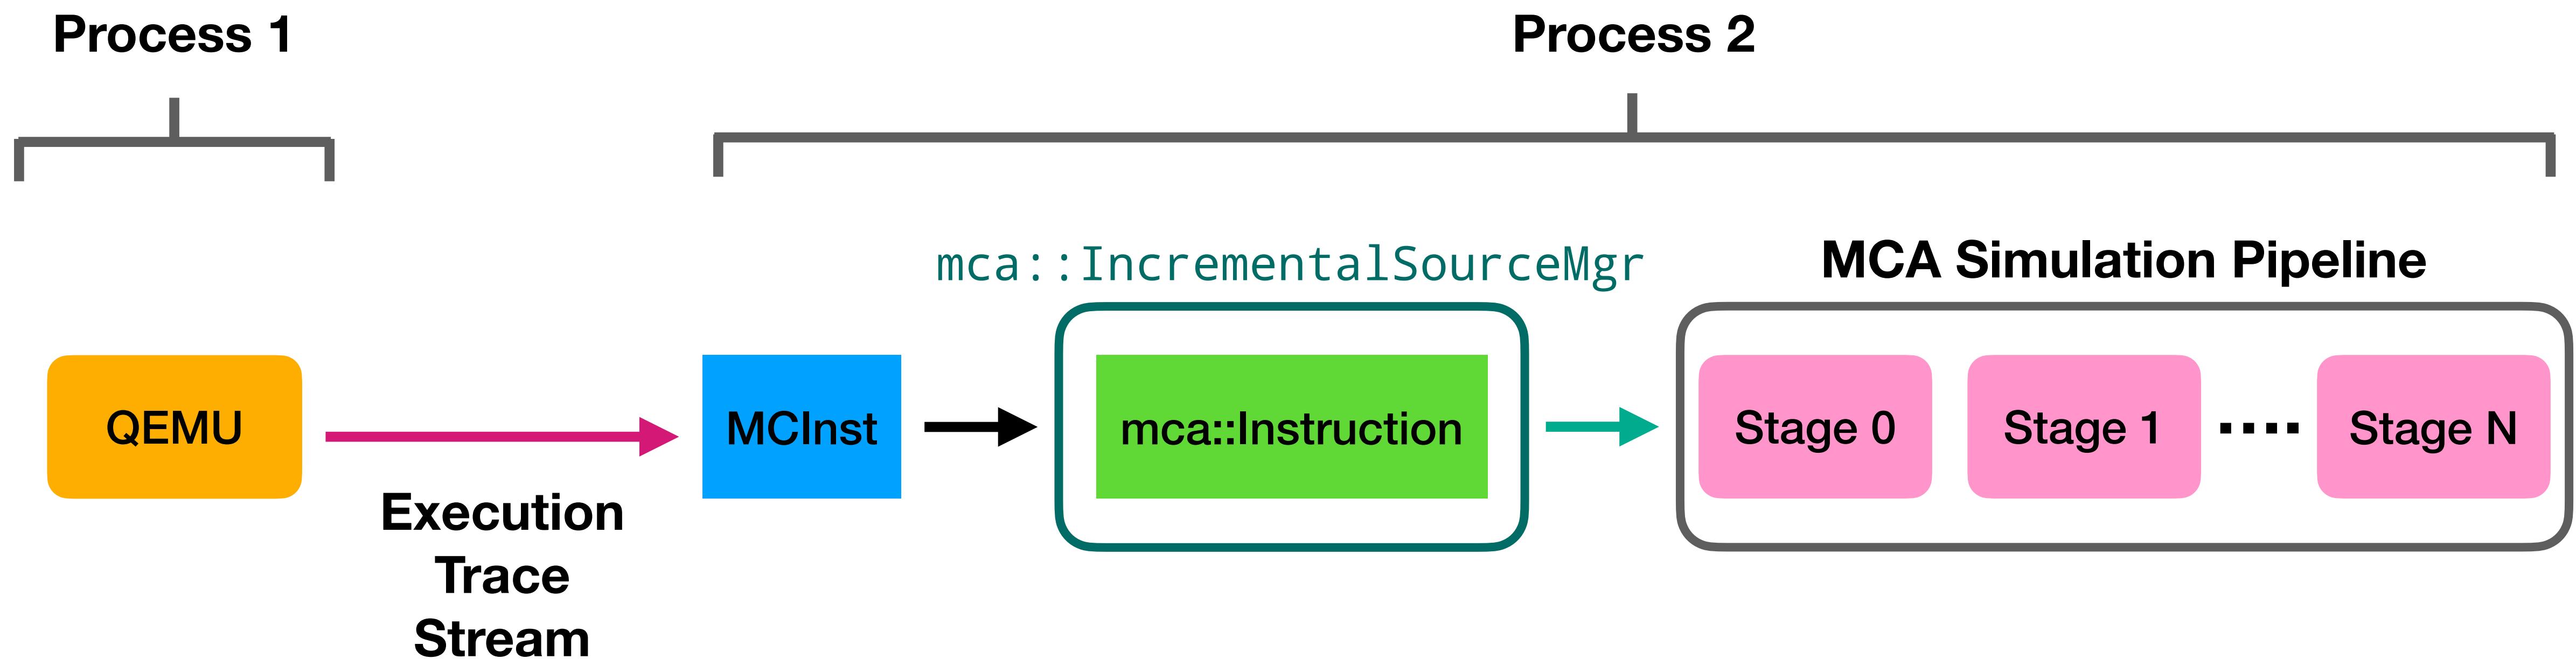
# Incremental SourceMgr



# Incremental SourceMgr

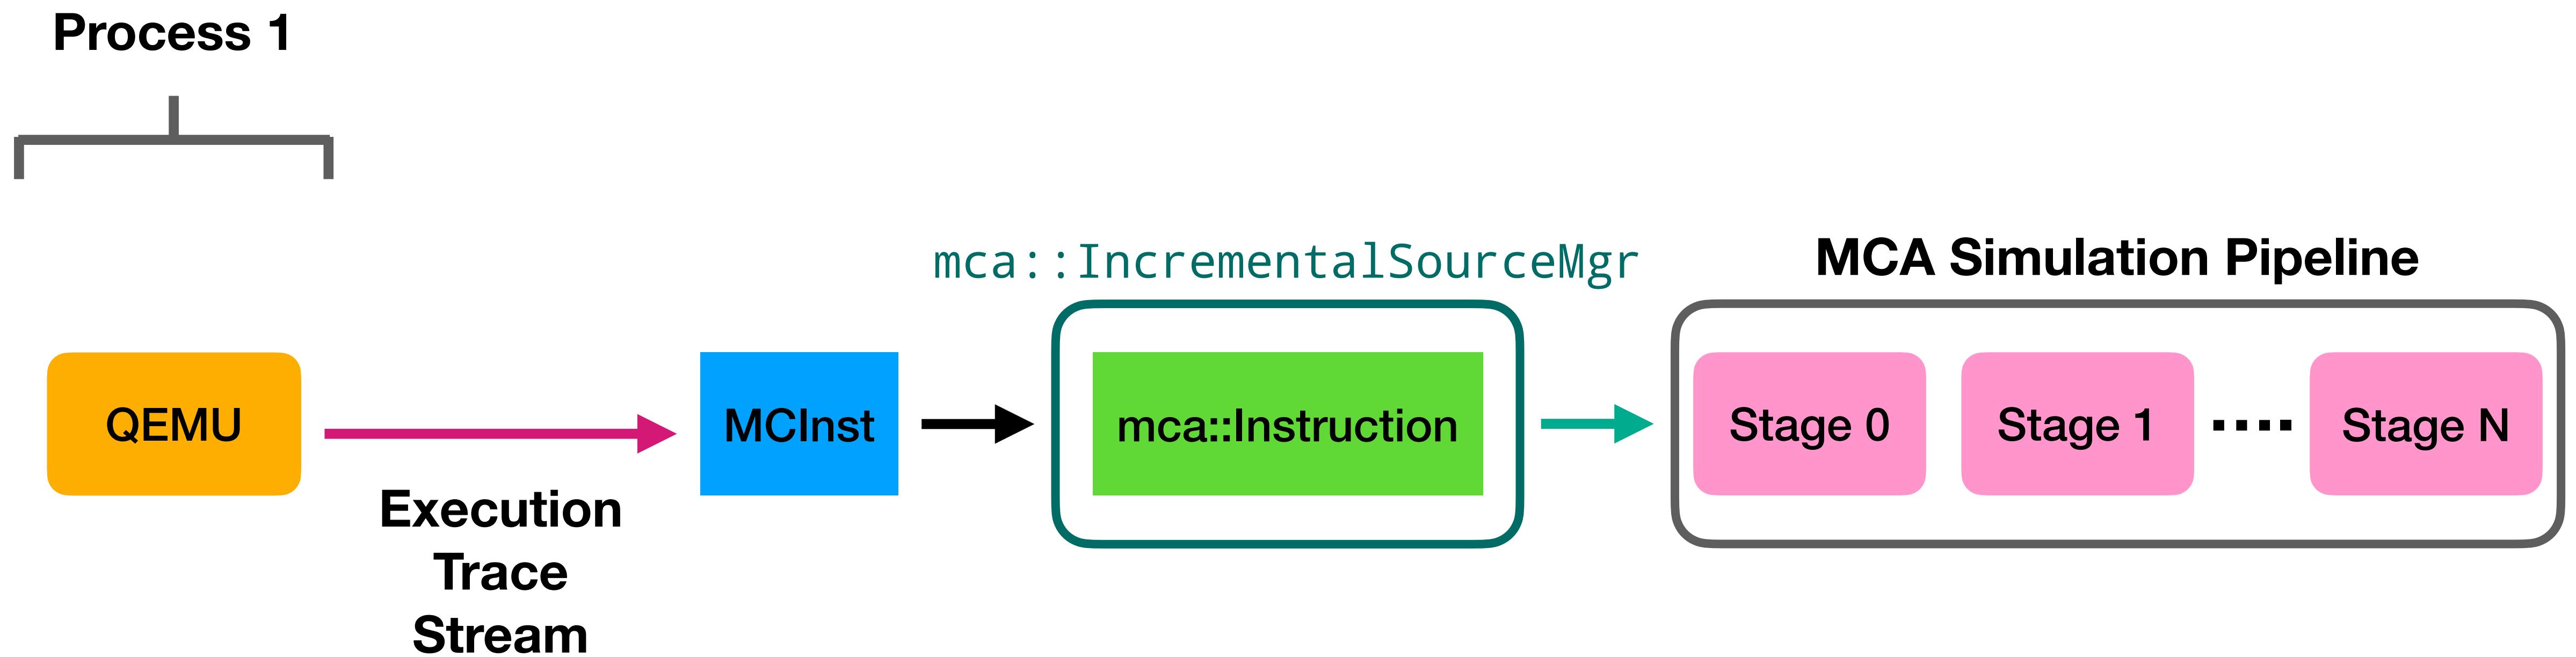


# Incremental SourceMgr



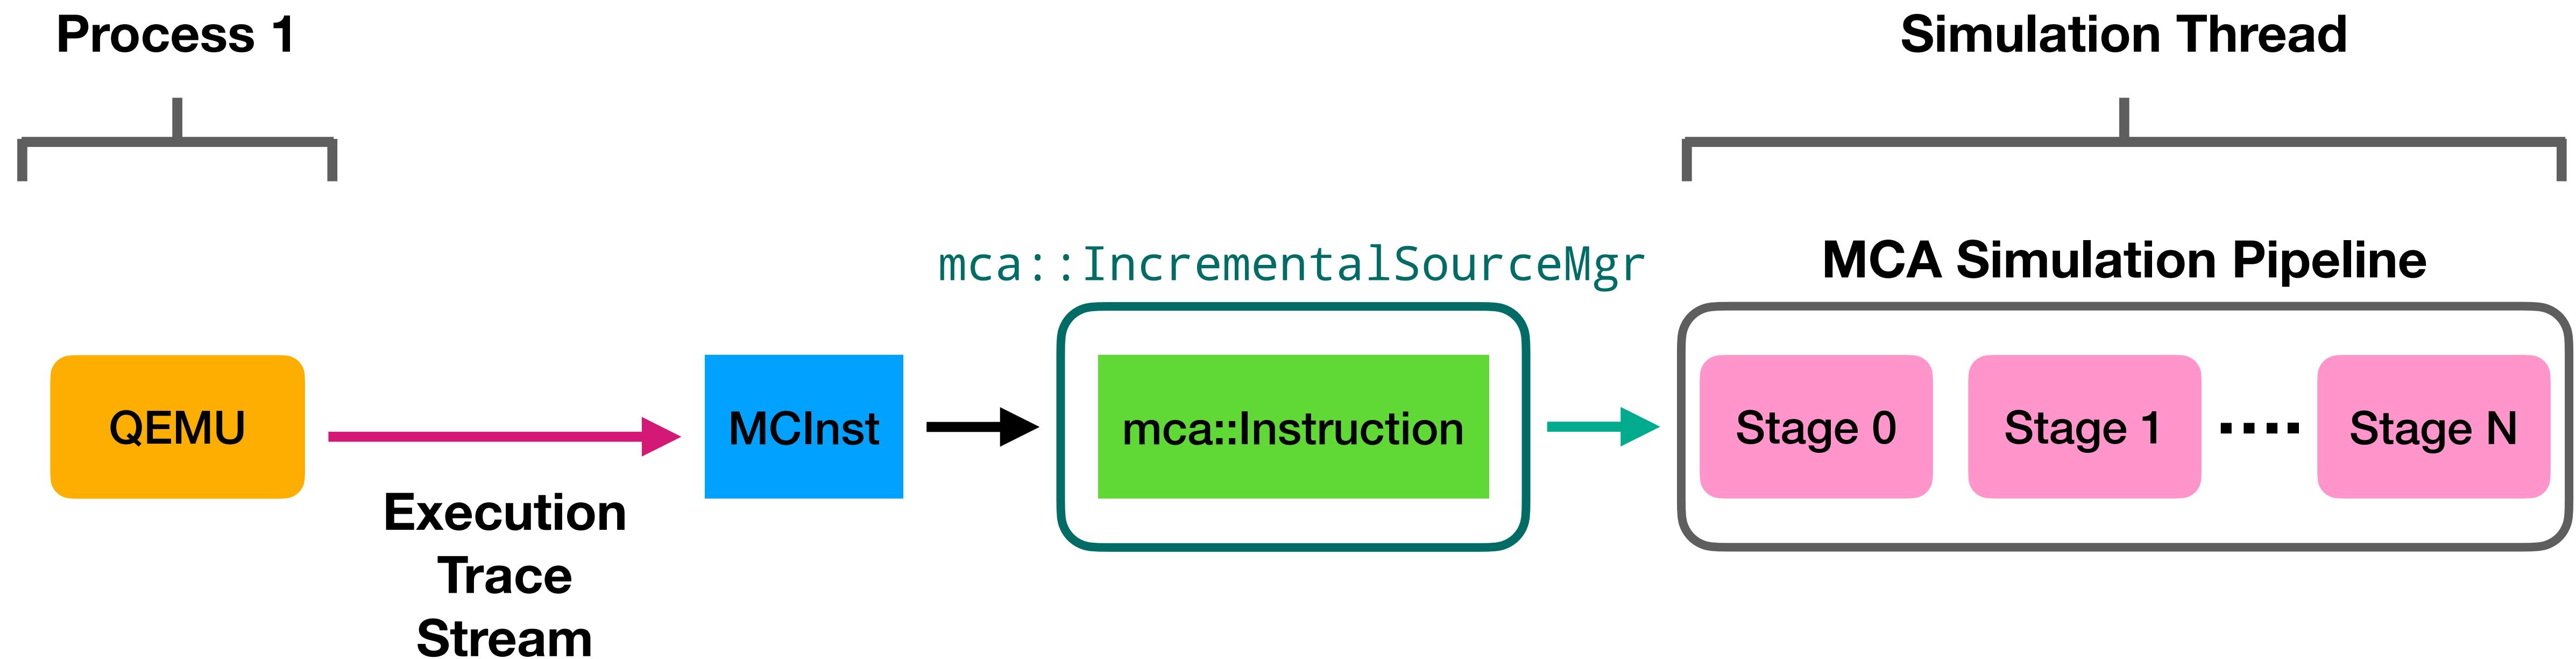
# Incremental SourceMgr

## Implement with threads



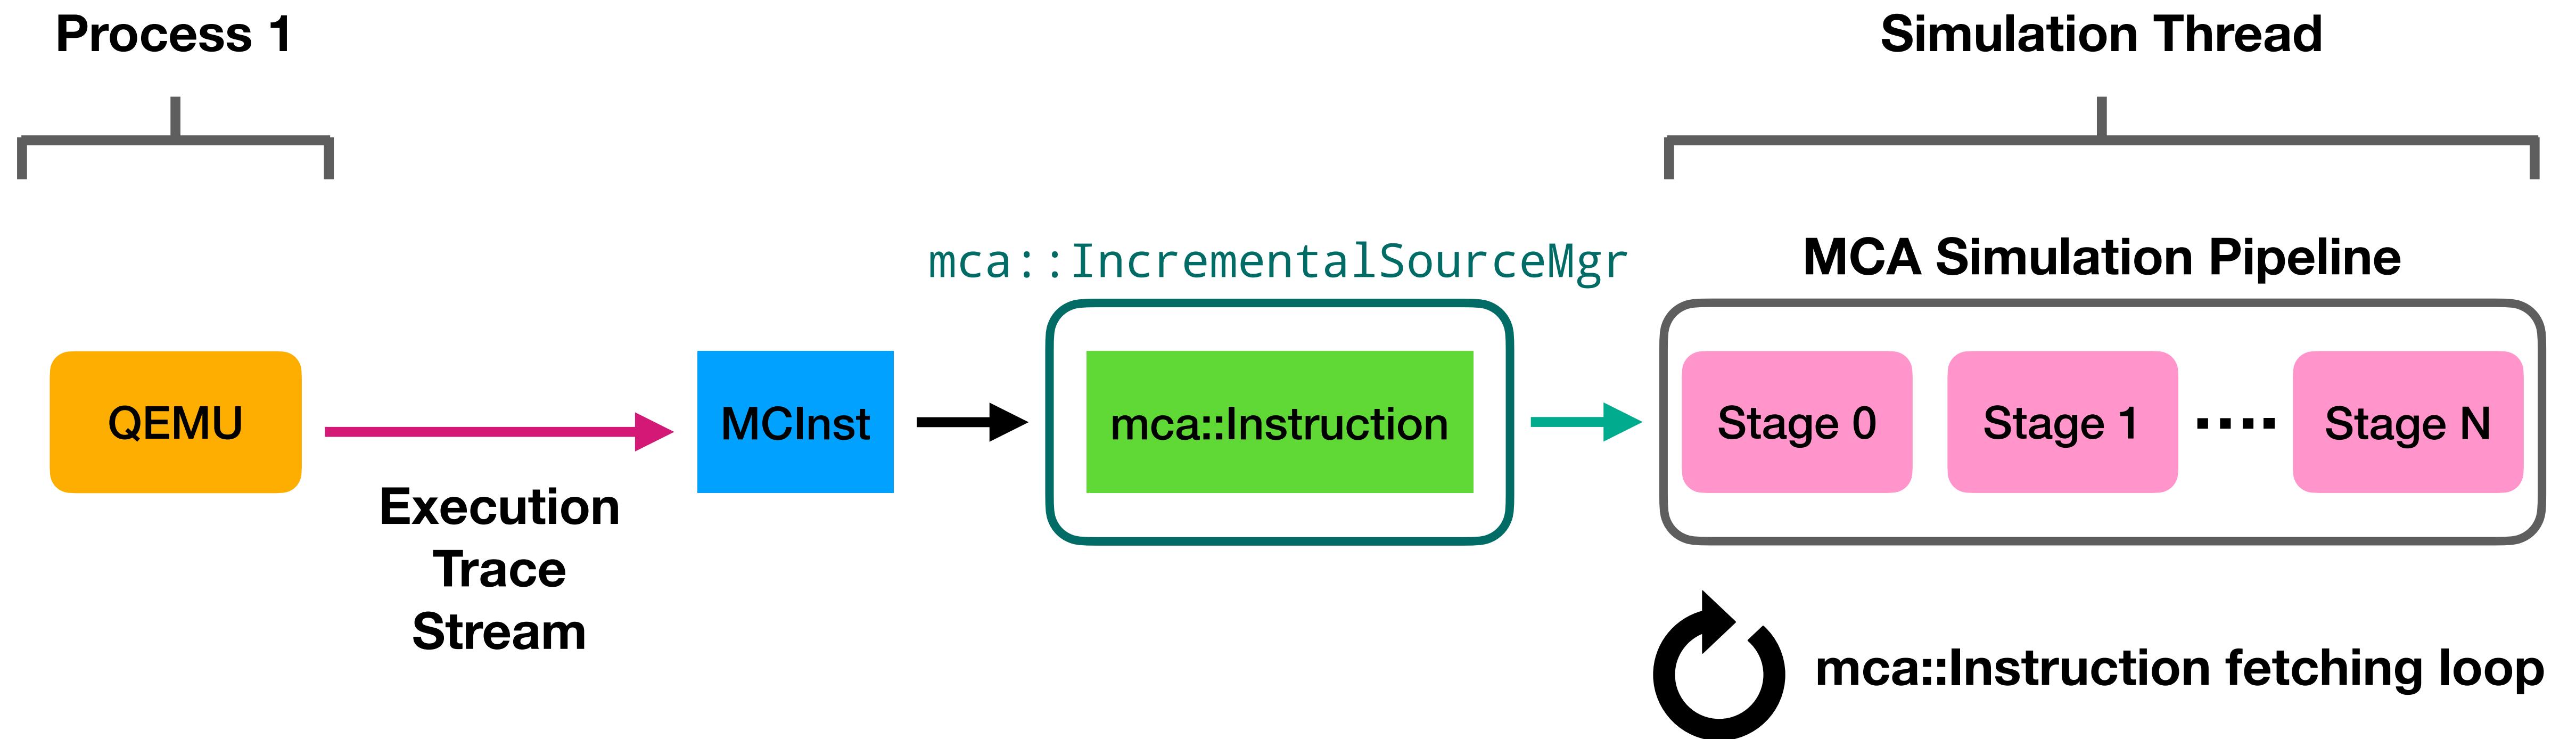
# Incremental SourceMgr

## Implement with threads



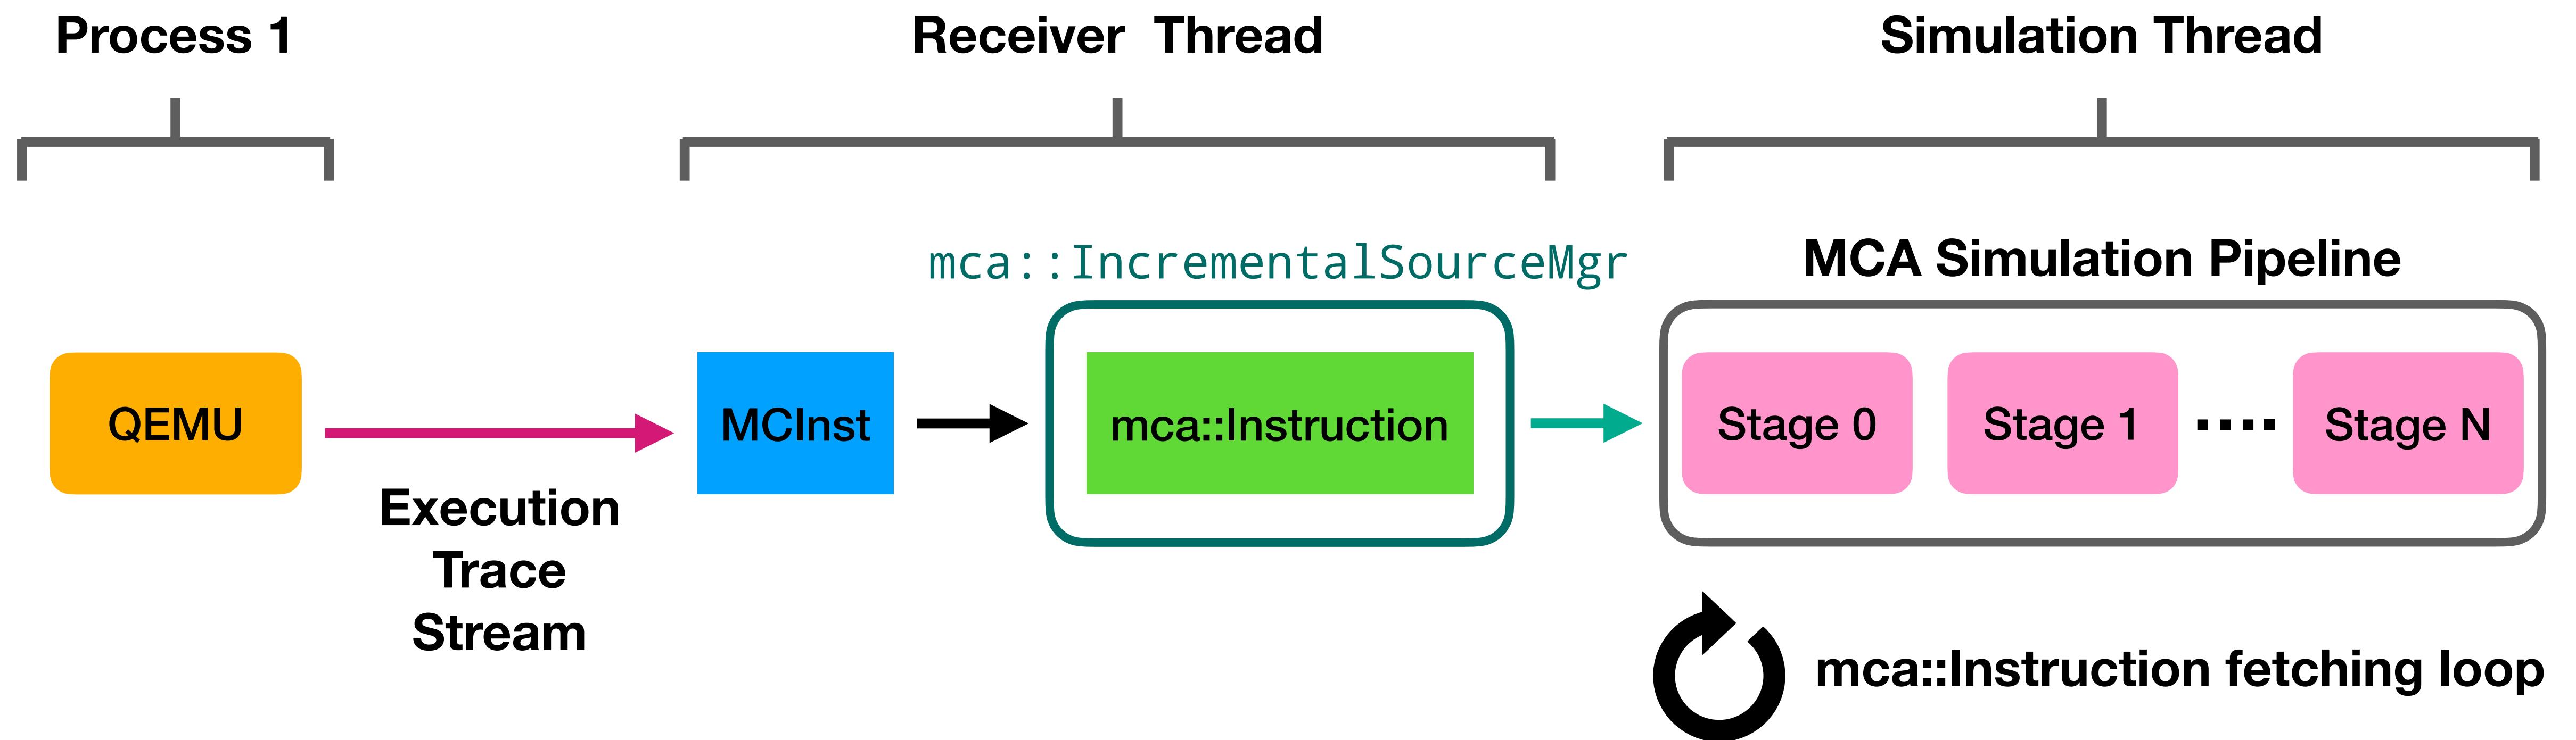
# Incremental SourceMgr

## Implement with threads



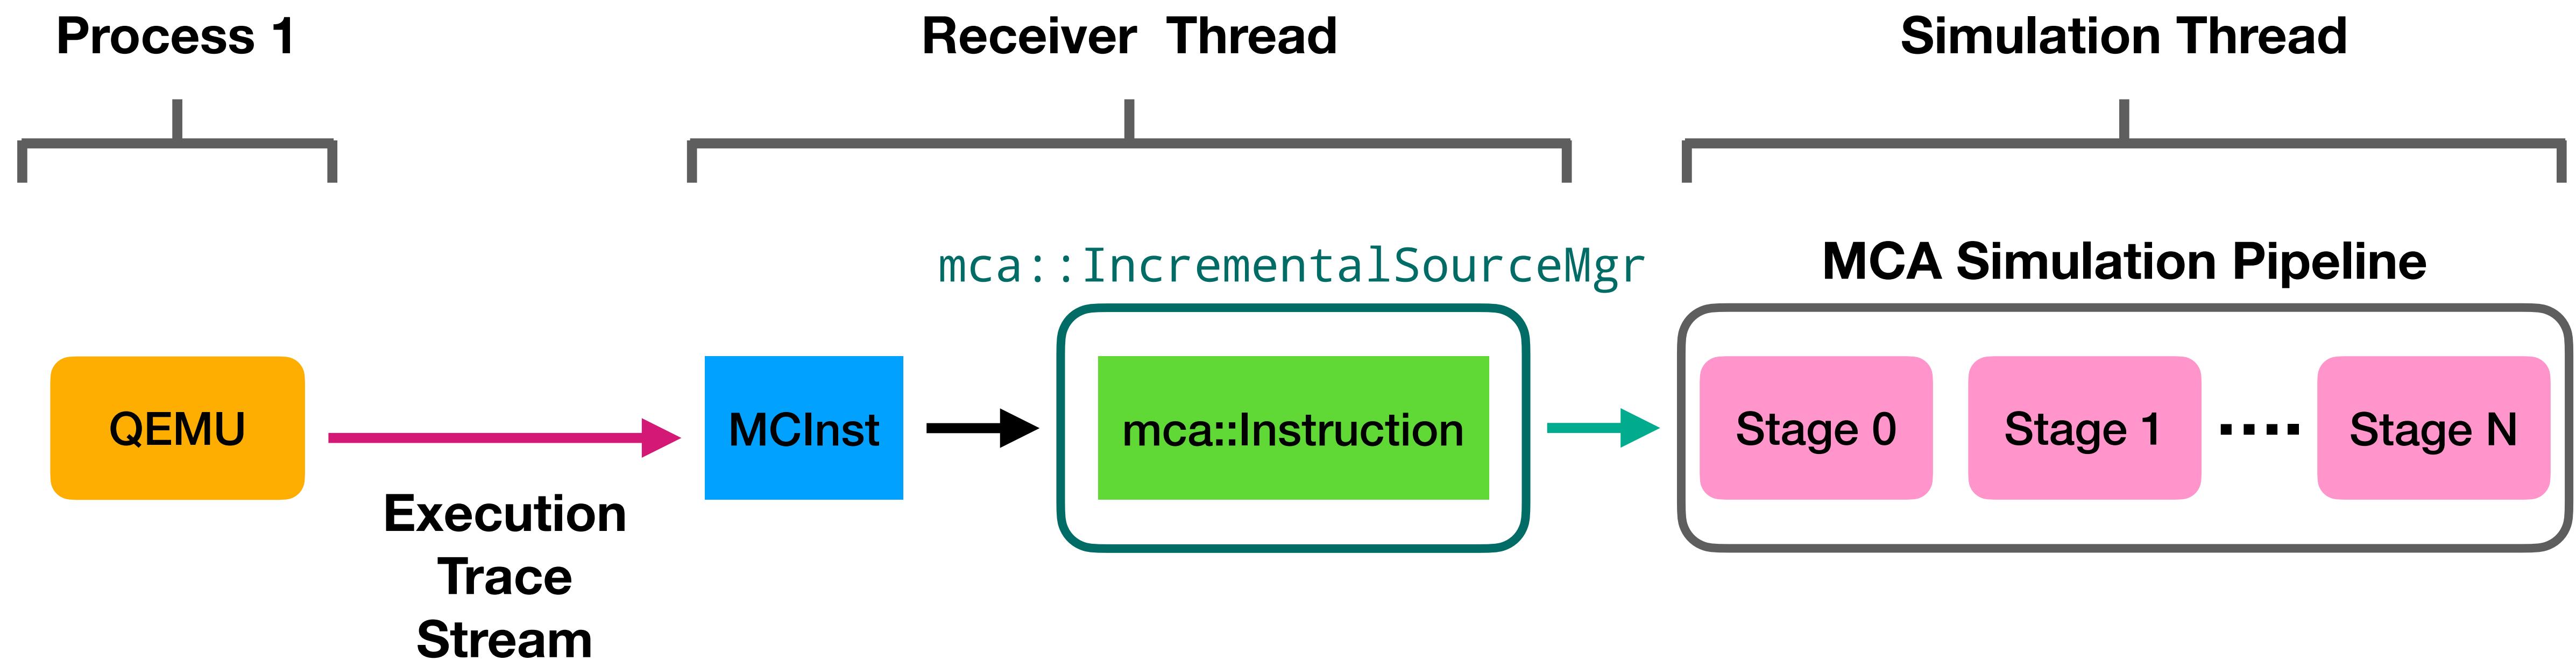
# Incremental SourceMgr

## Implement with threads



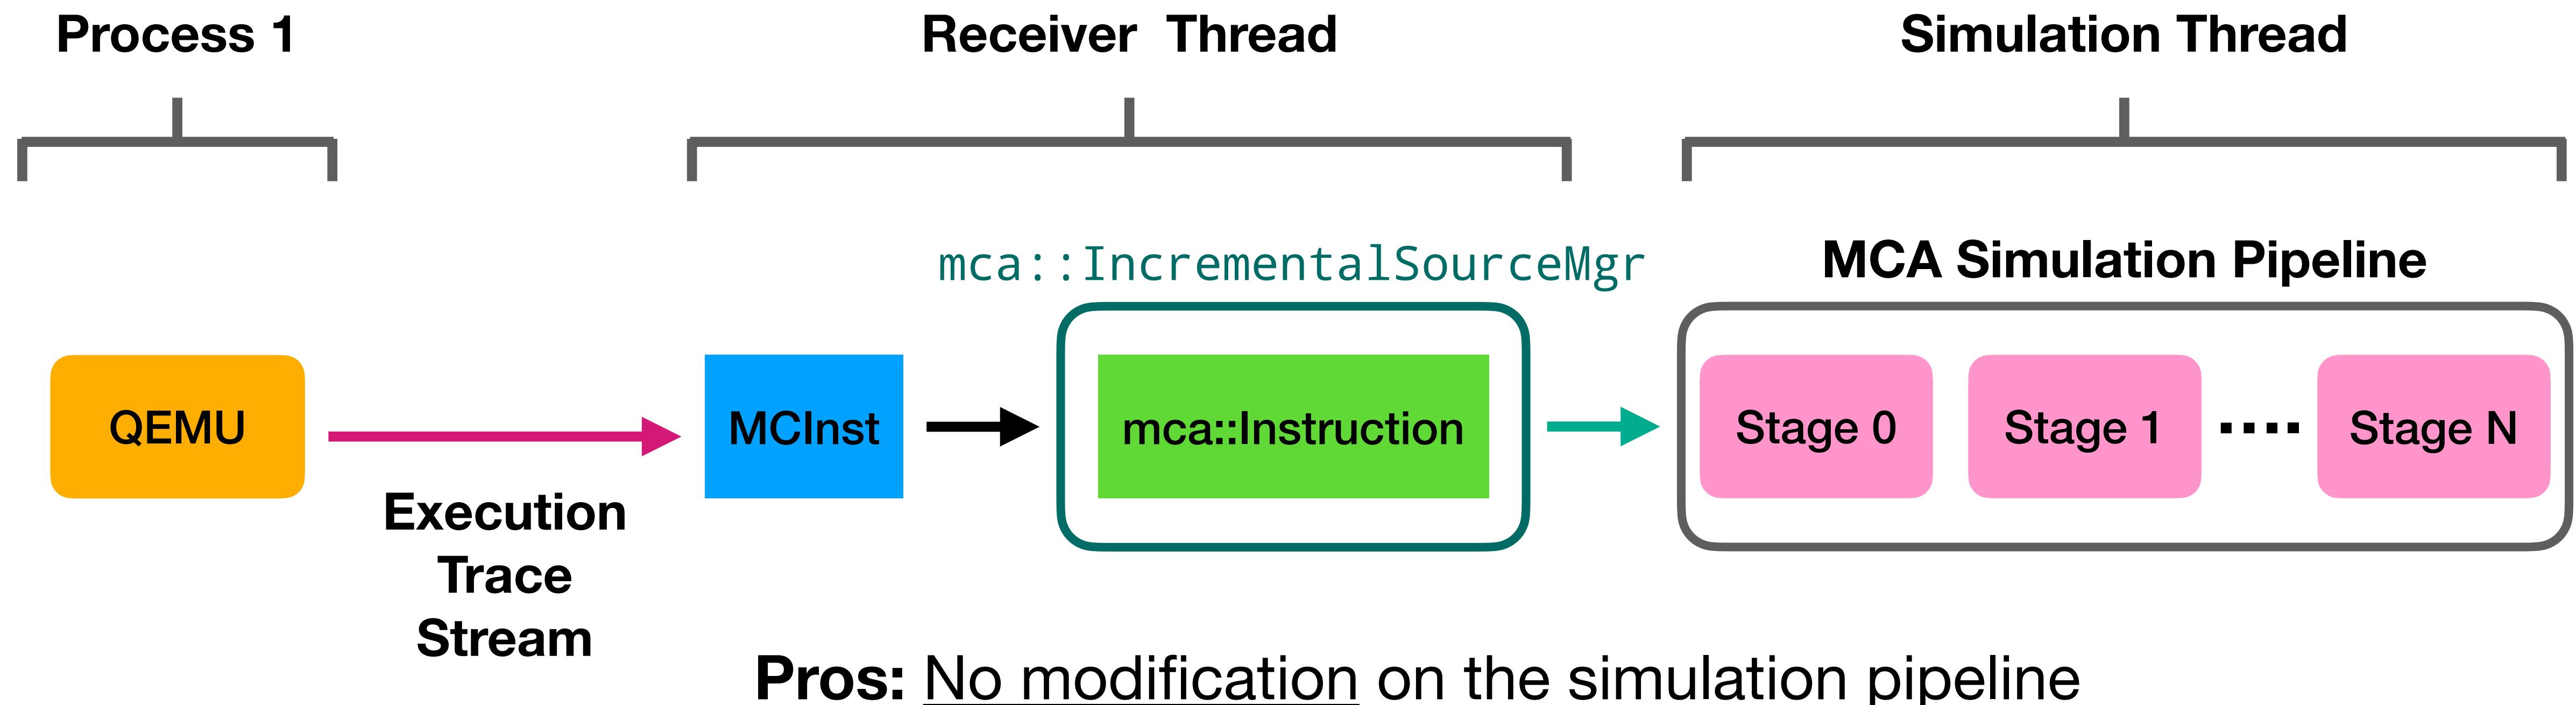
# Incremental SourceMgr

## Implement with threads: Pros & Cons



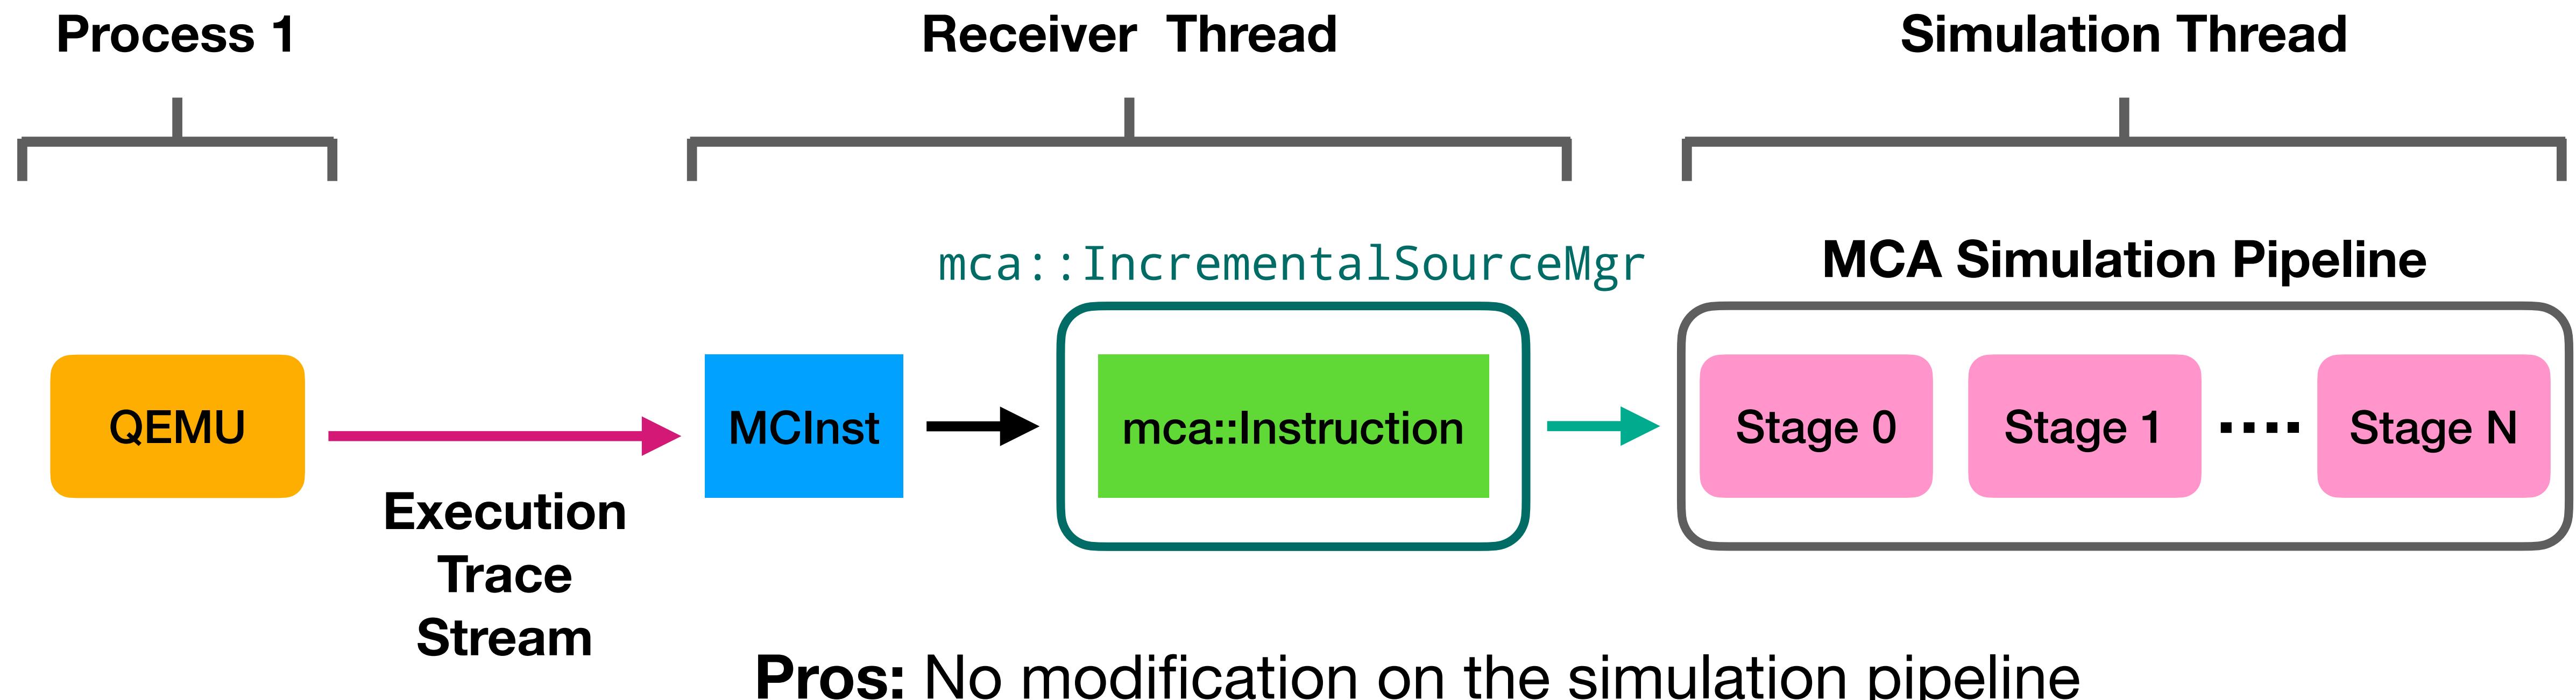
# Incremental SourceMgr

## Implement with threads: Pros & Cons



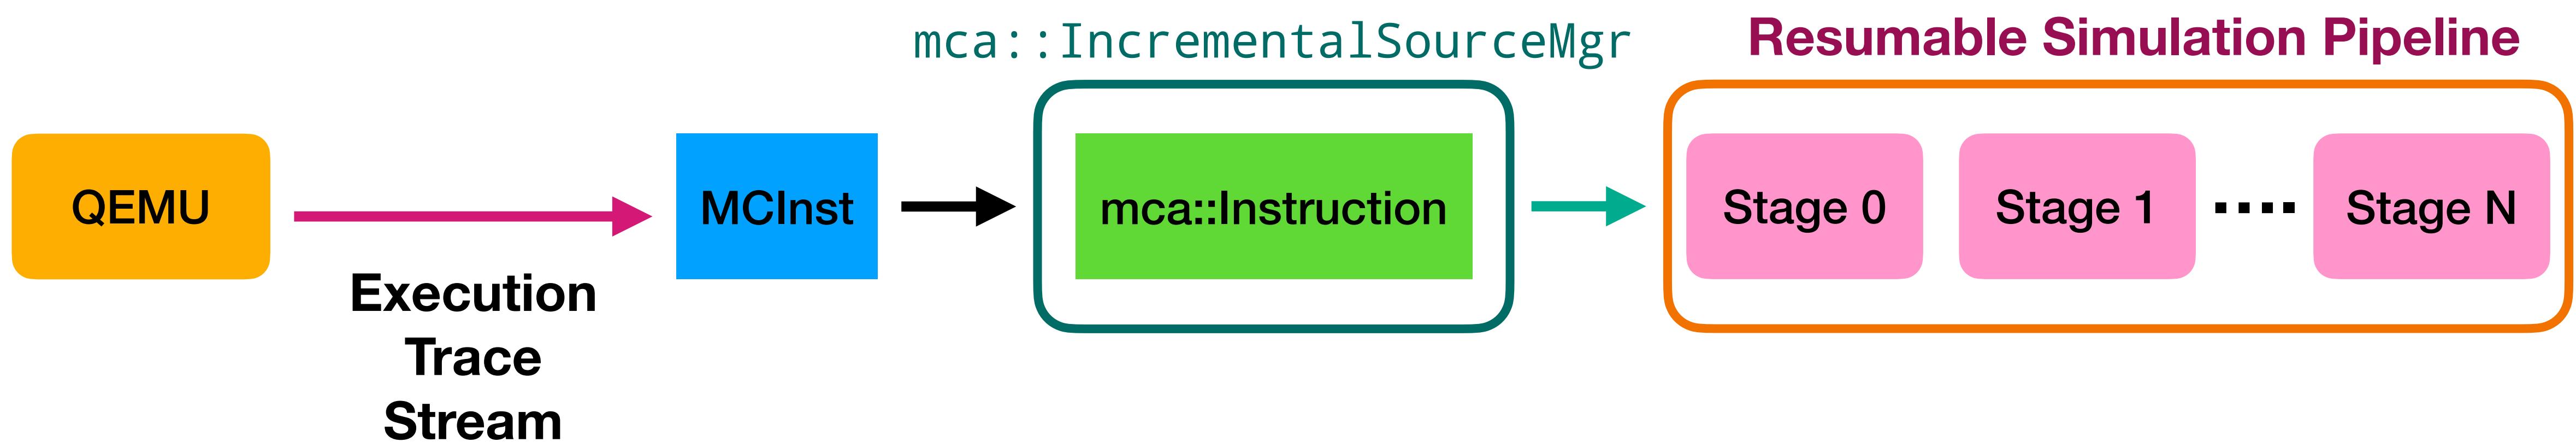
# Incremental SourceMgr

## Implement with threads: Pros & Cons



# Incremental SourceMgr

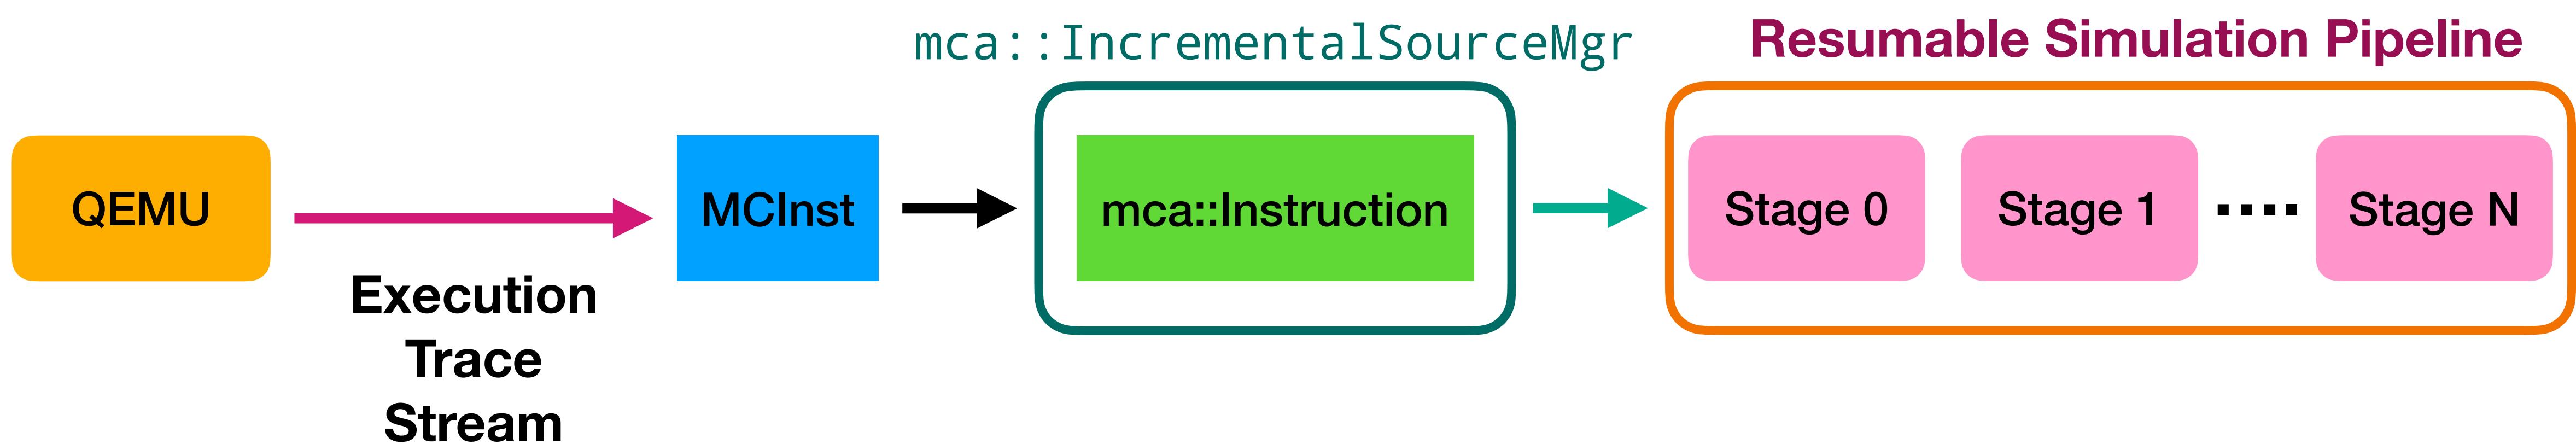
Better solution: Resumable simulation pipeline



# Incremental SourceMgr

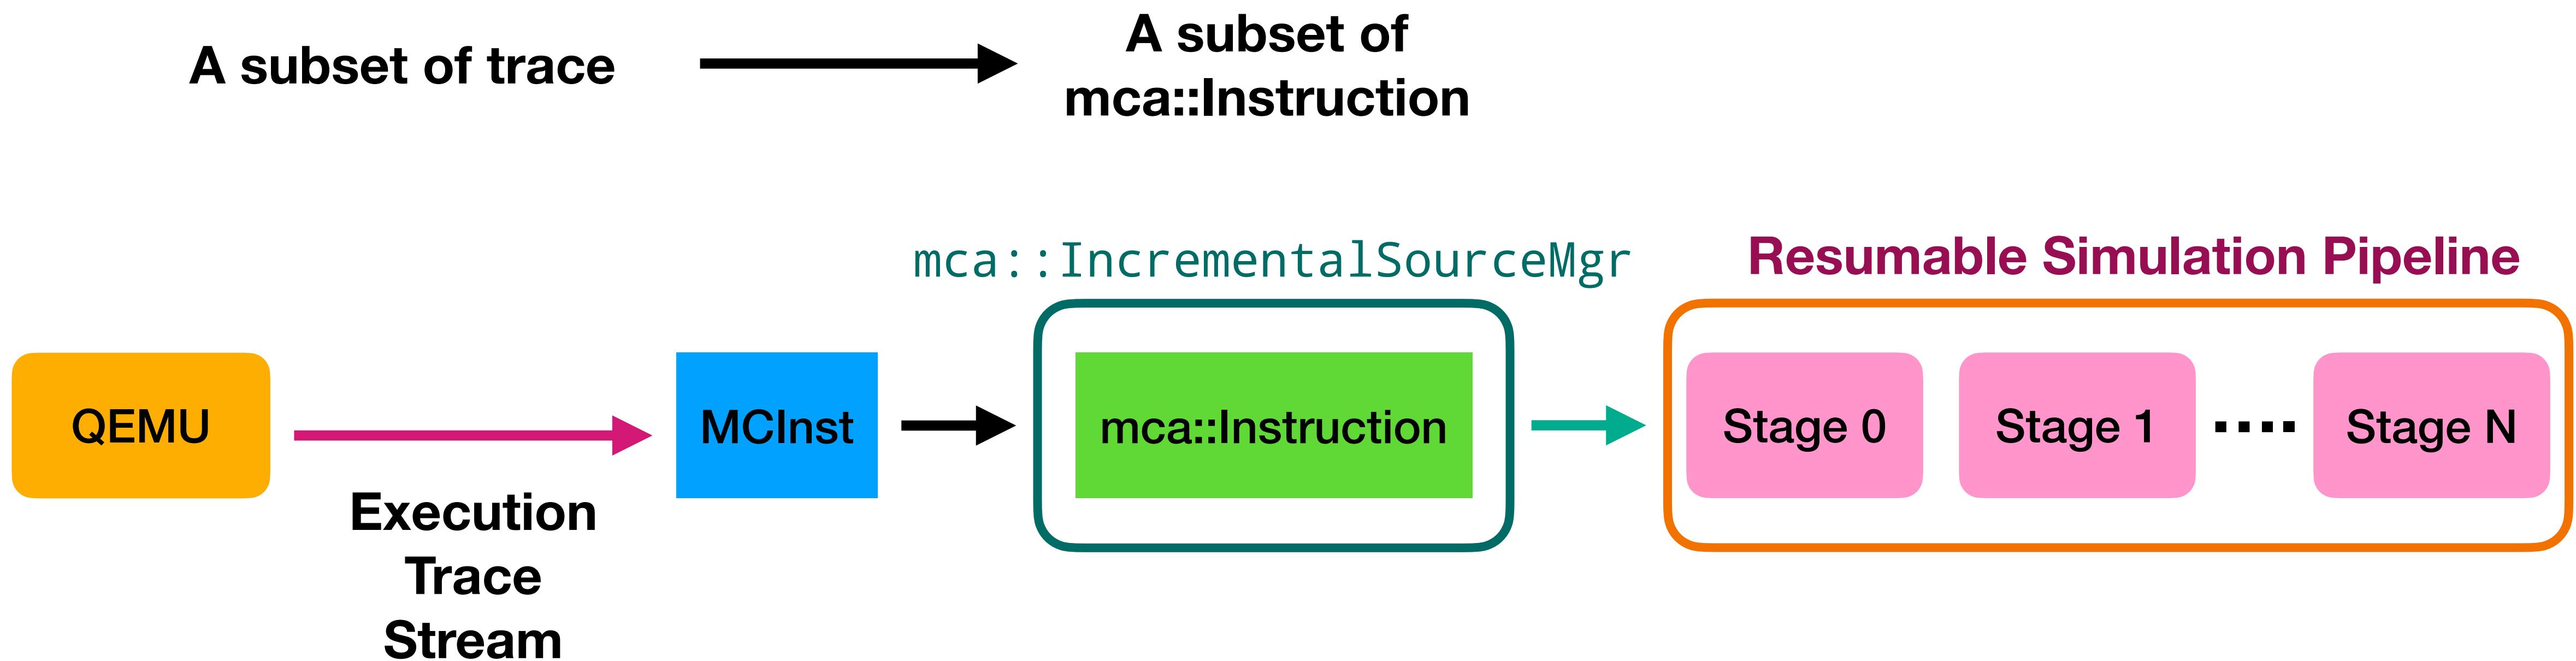
Better solution: Resumable simulation pipeline

A subset of trace



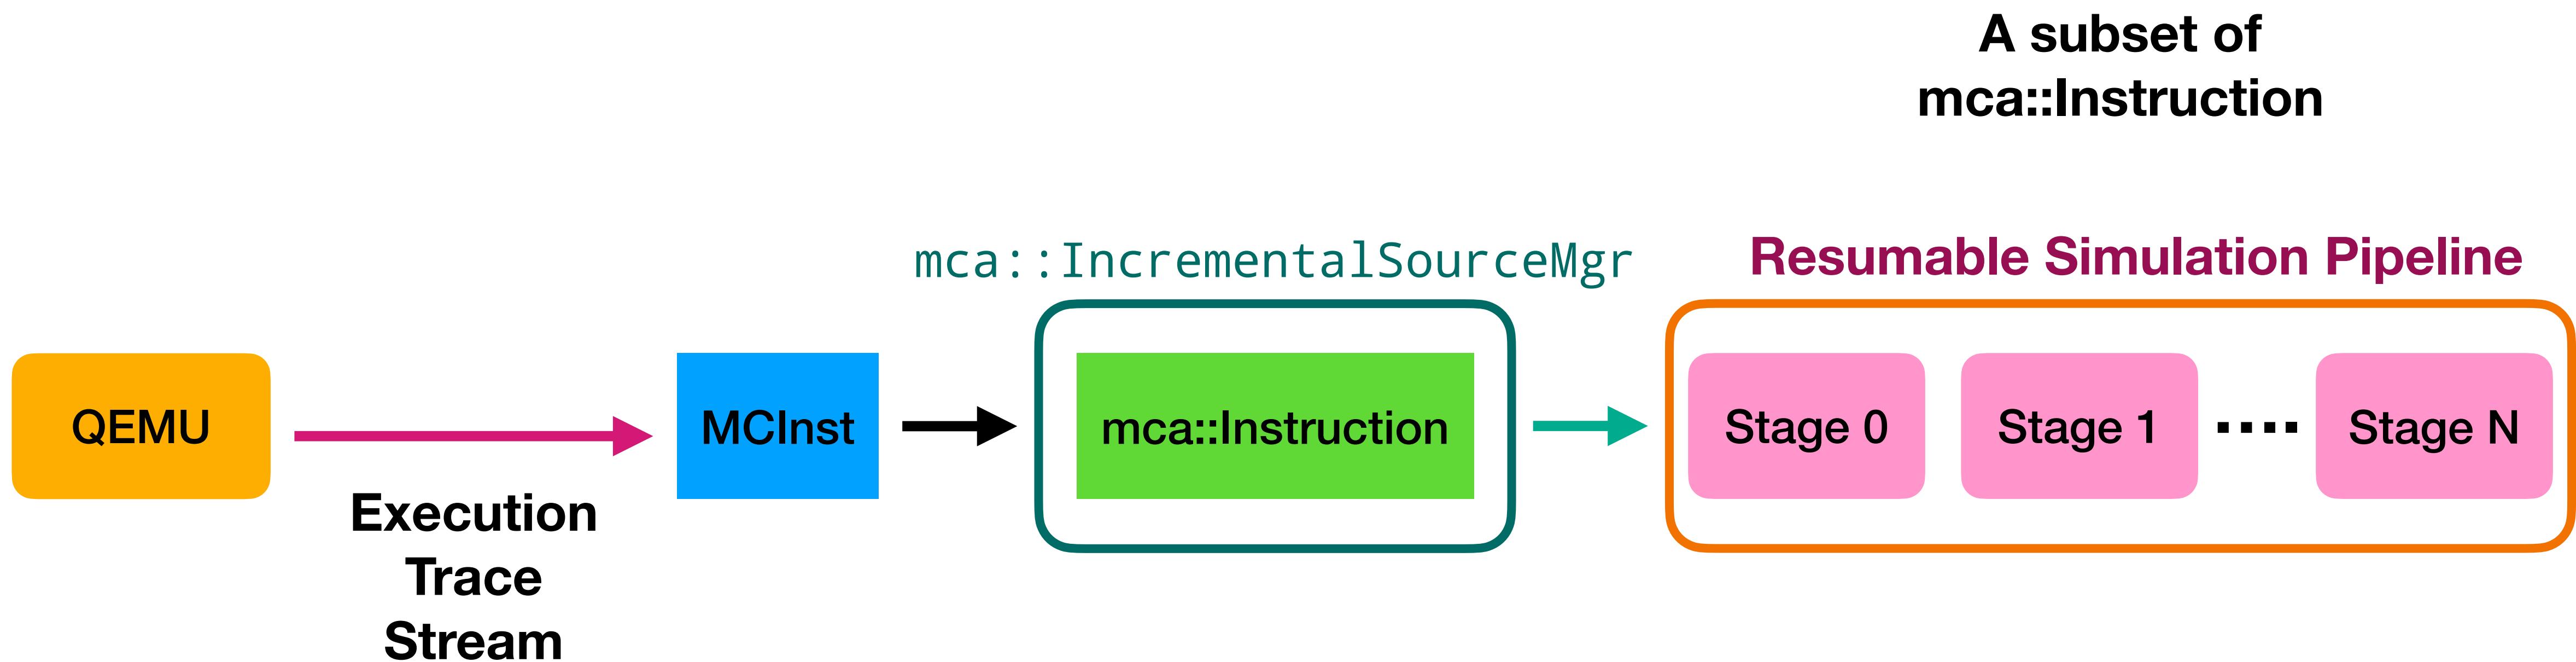
# Incremental SourceMgr

Better solution: Resumable simulation pipeline



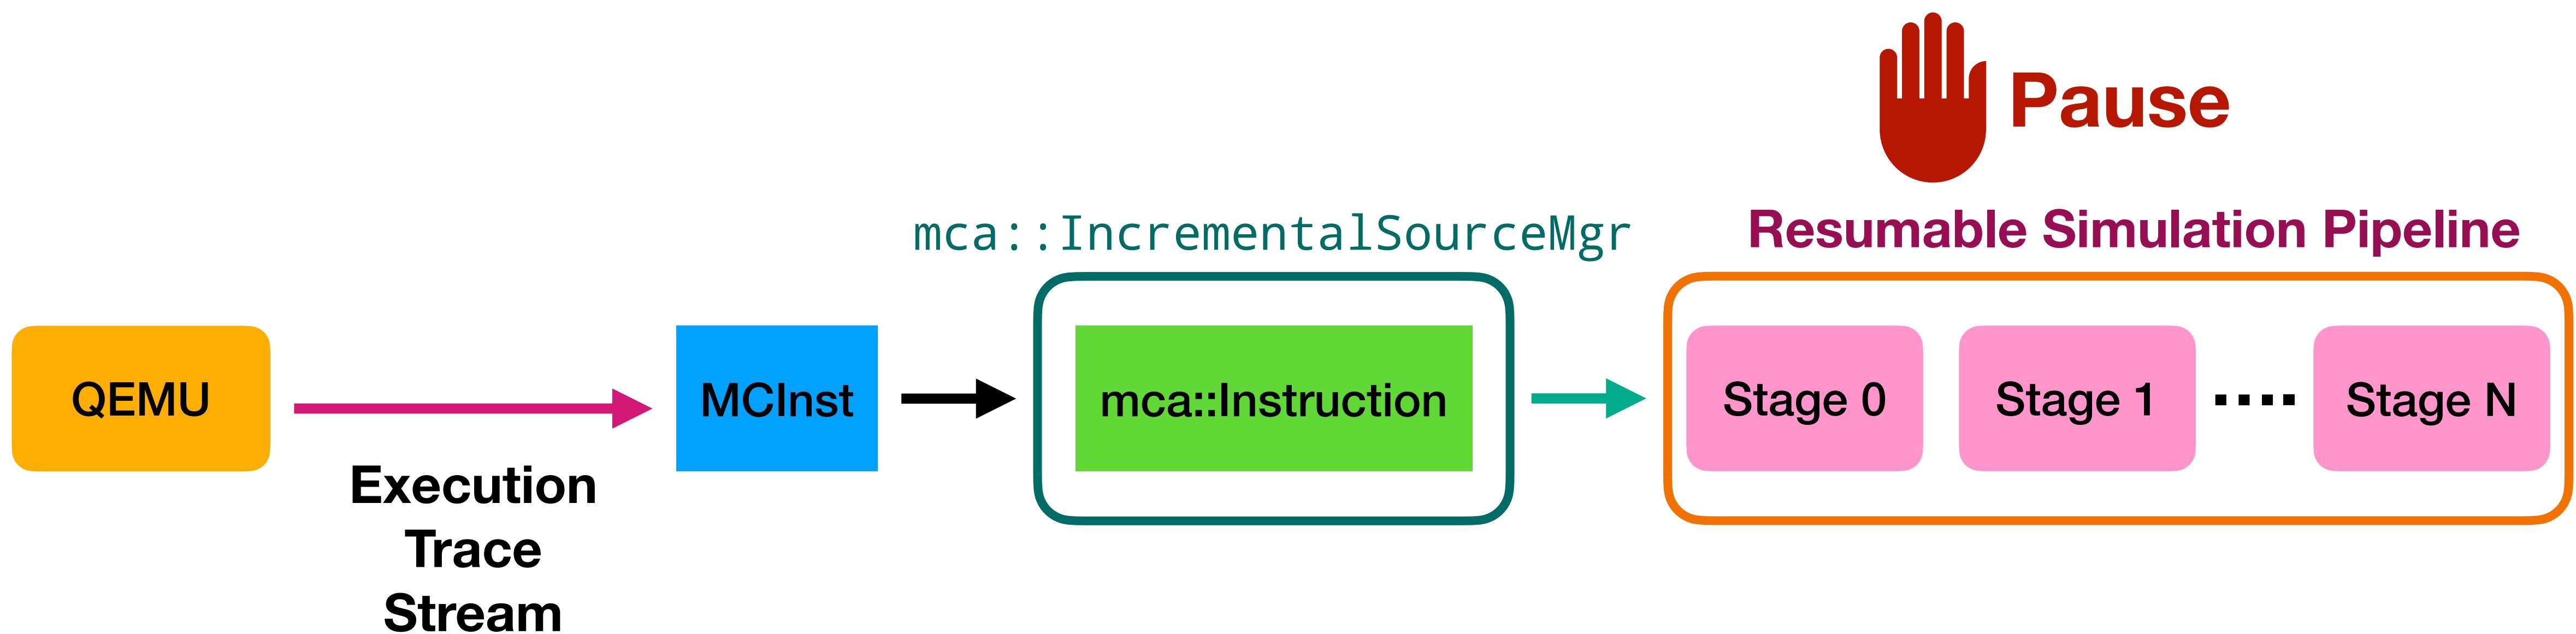
# Incremental SourceMgr

Better solution: Resumable simulation pipeline



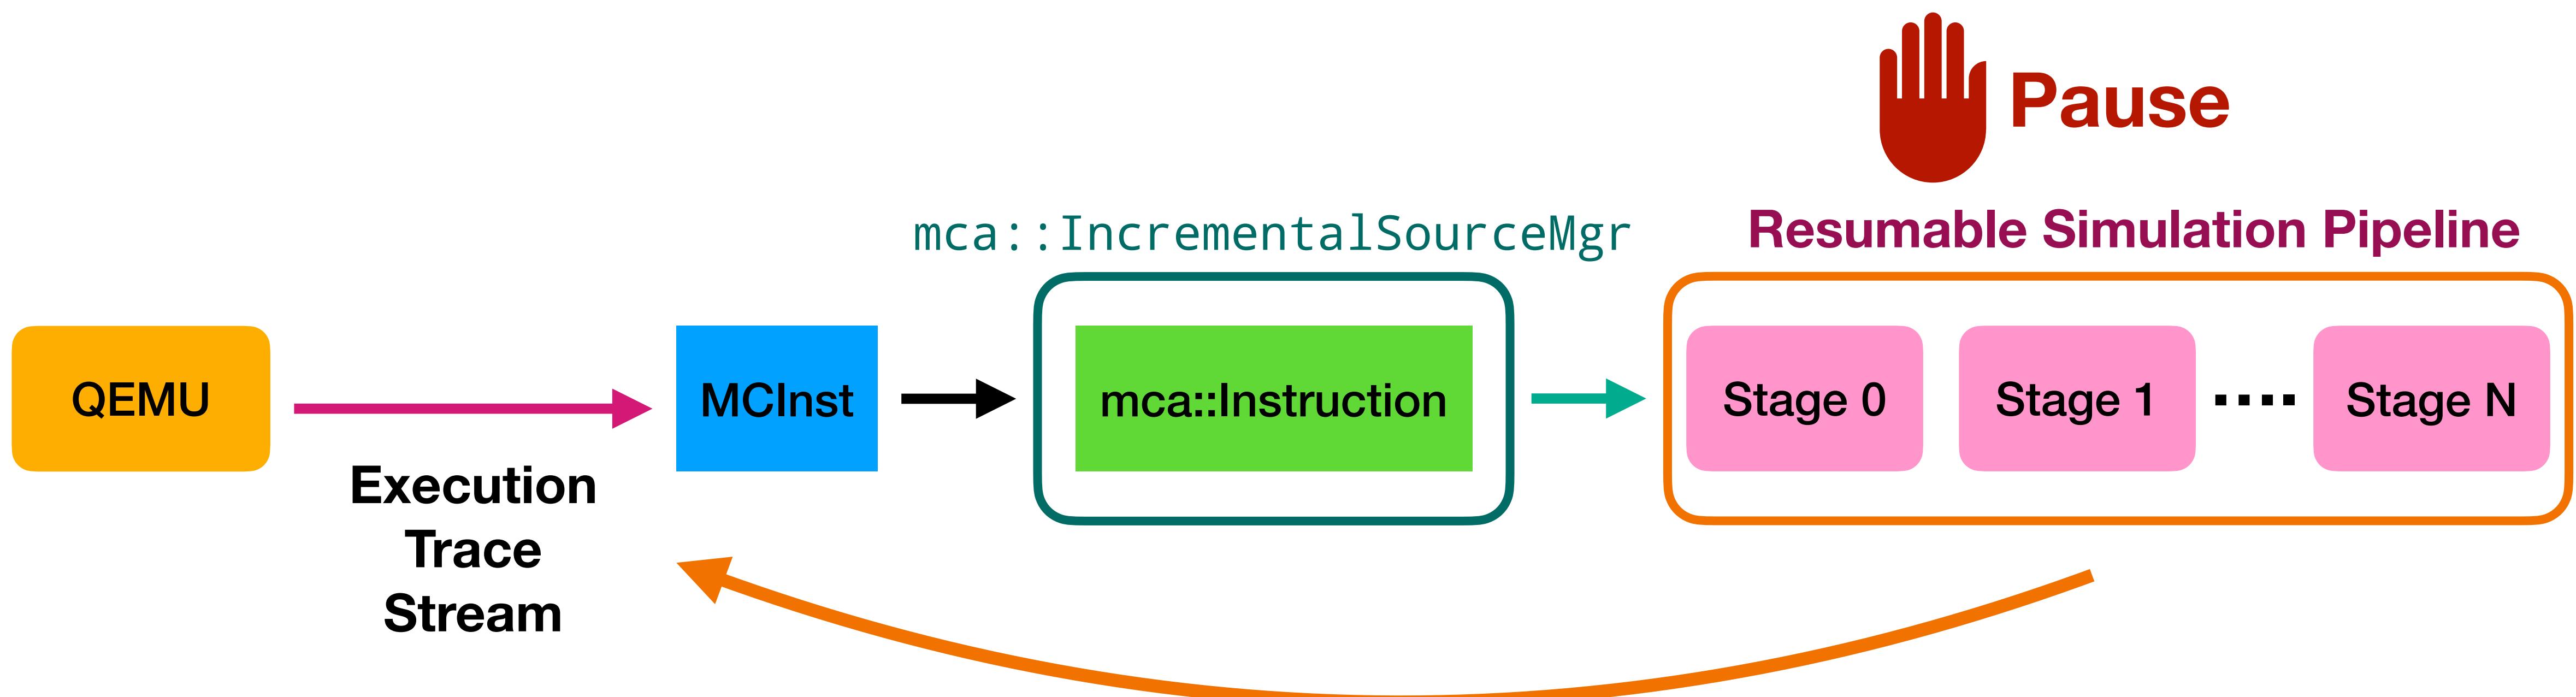
# Incremental SourceMgr

Better solution: Resumable simulation pipeline



# Incremental SourceMgr

Better solution: Resumable simulation pipeline



# Resumable simulation pipeline

# Resumable simulation pipeline

- **Save** (and restore) the analysis state from *previous* subset of instructions

# Resumable simulation pipeline

- **Save** (and restore) the analysis state from *previous* subset of instructions
- Threads are **not required** when using IncrementalSourceMgr + resumable pipeline

# Resumable simulation pipeline

- **Save** (and restore) the analysis state from *previous* subset of instructions
- Threads are **not required** when using IncrementalSourceMgr + resumable pipeline
  - Much easier to integrate into other uses

# Resumable simulation pipeline

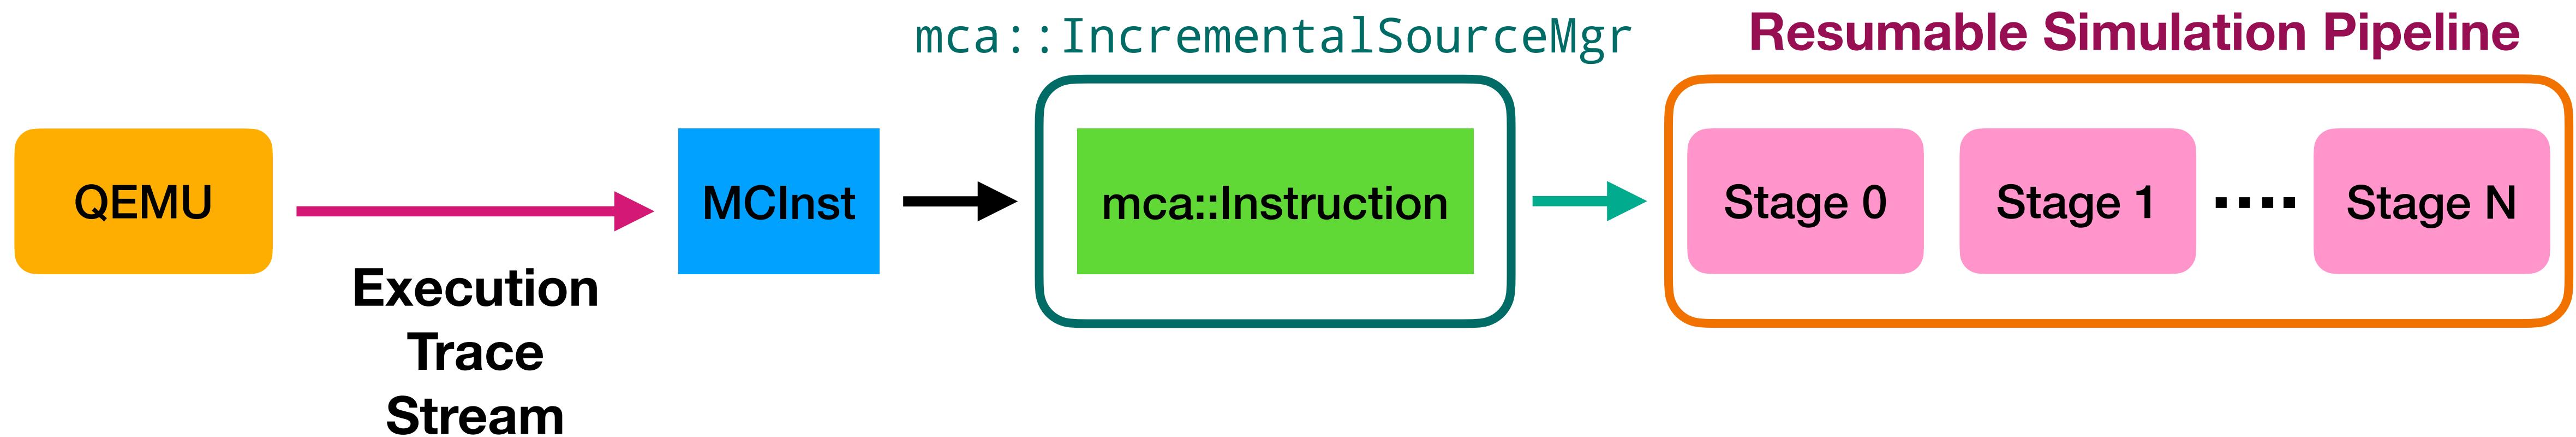
- **Save** (and restore) the analysis state from *previous* subset of instructions
- Threads are **not required** when using IncrementalSourceMgr + resumable pipeline
  - Much easier to integrate into other uses
  - You can still wrap resumable pipeline with a thread

# Resumable simulation pipeline

- **Save** (and restore) the analysis state from *previous* subset of instructions
- Threads are **not required** when using IncrementalSourceMgr + resumable pipeline
  - Much easier to integrate into other uses
  - You can still wrap resumable pipeline with a thread
- Minor downside: Modifications on the simulation pipeline

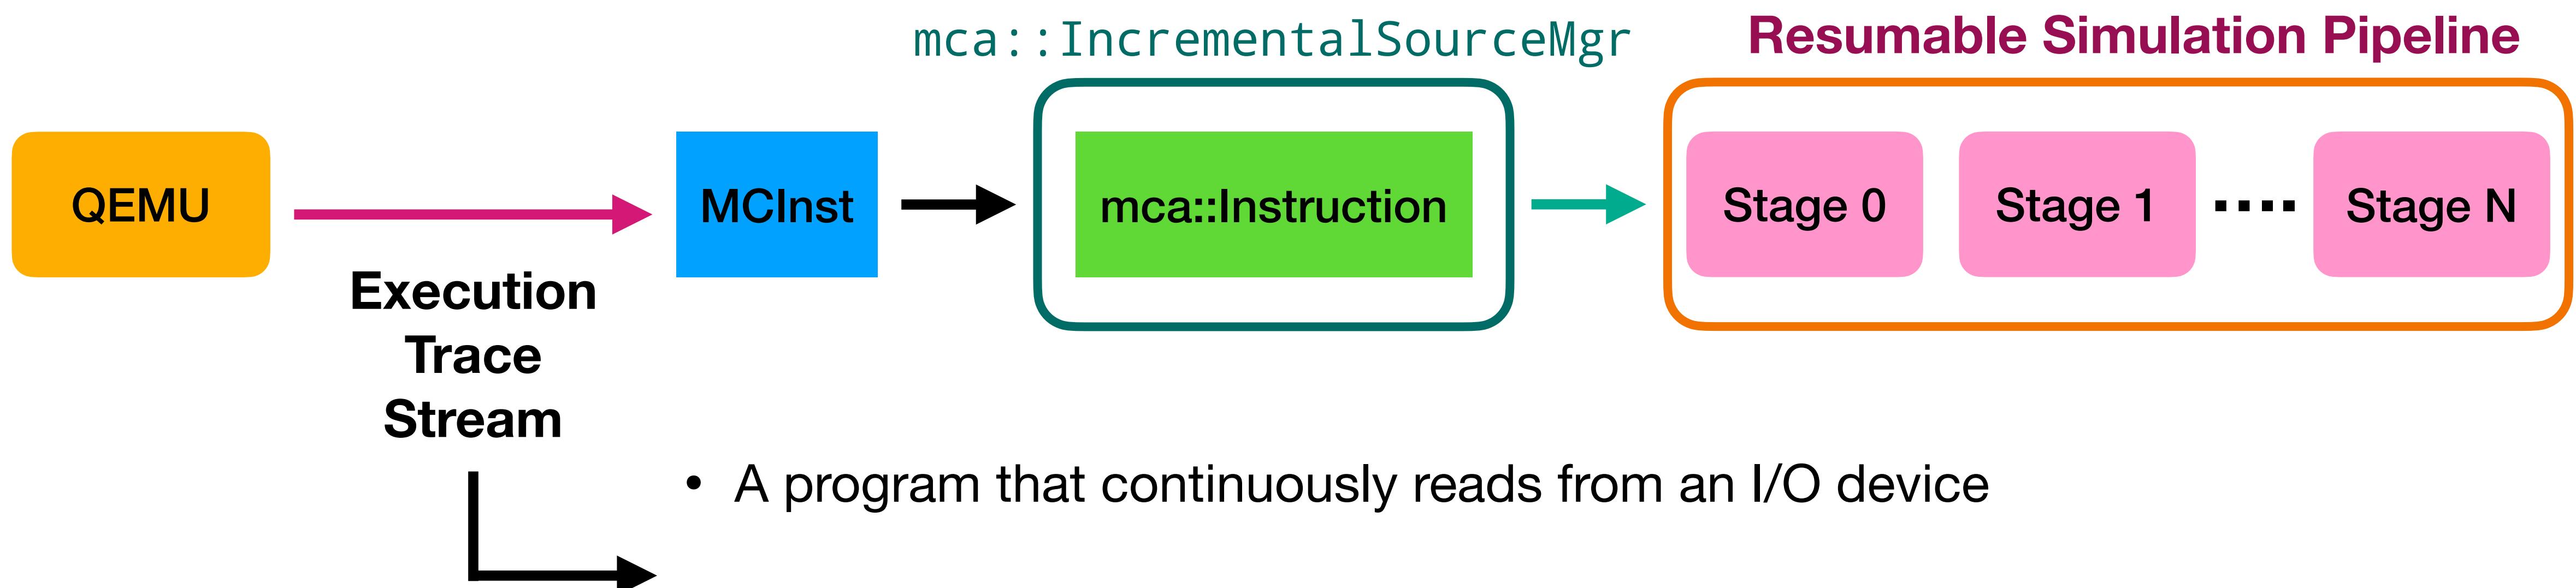
# Incremental SourceMgr + Resumable pipeline

## Put into real actions



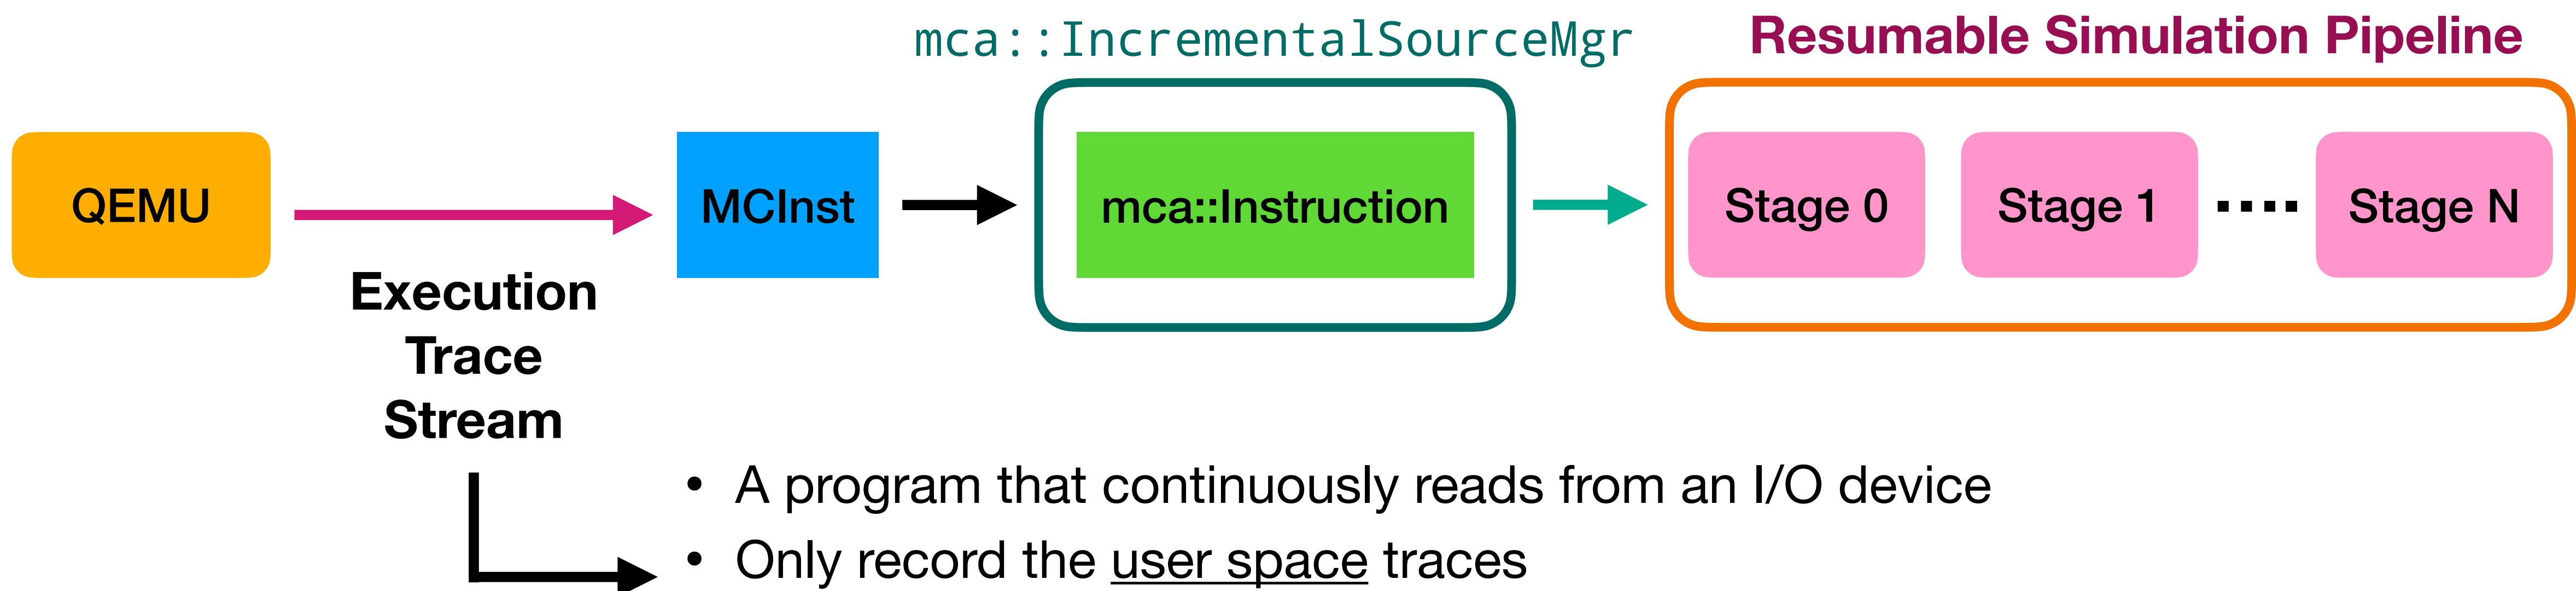
# Incremental SourceMgr + Resumable pipeline

## Put into real actions



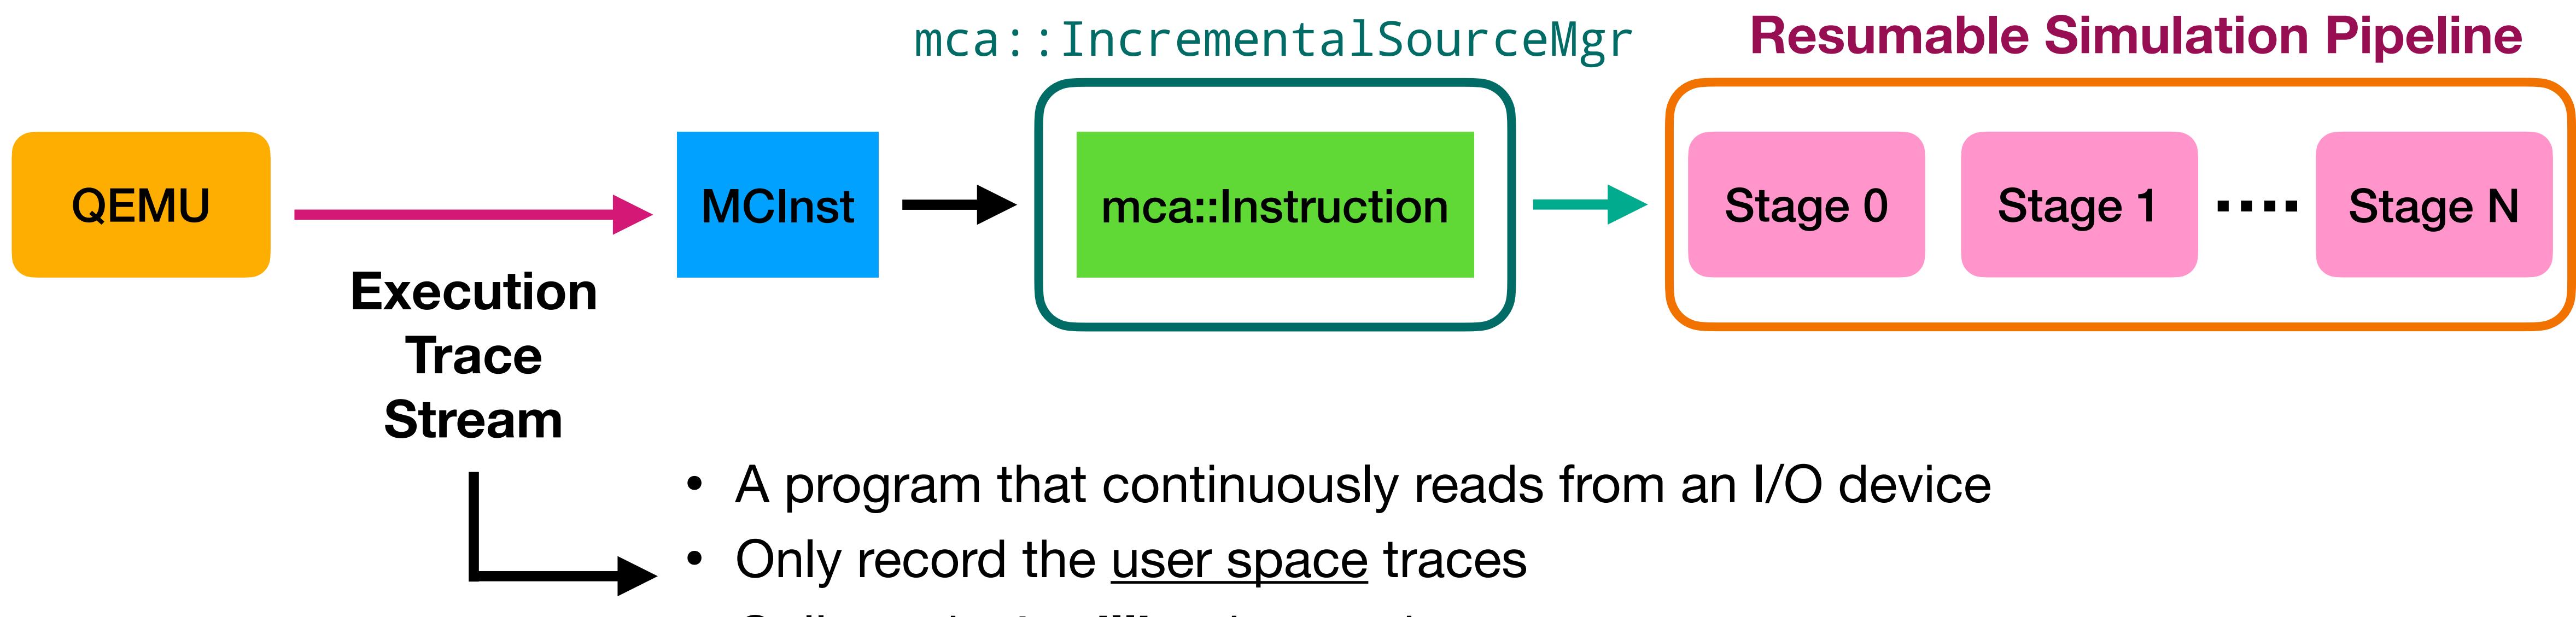
# Incremental SourceMgr + Resumable pipeline

## Put into real actions



# Incremental SourceMgr + Resumable pipeline

## Put into real actions



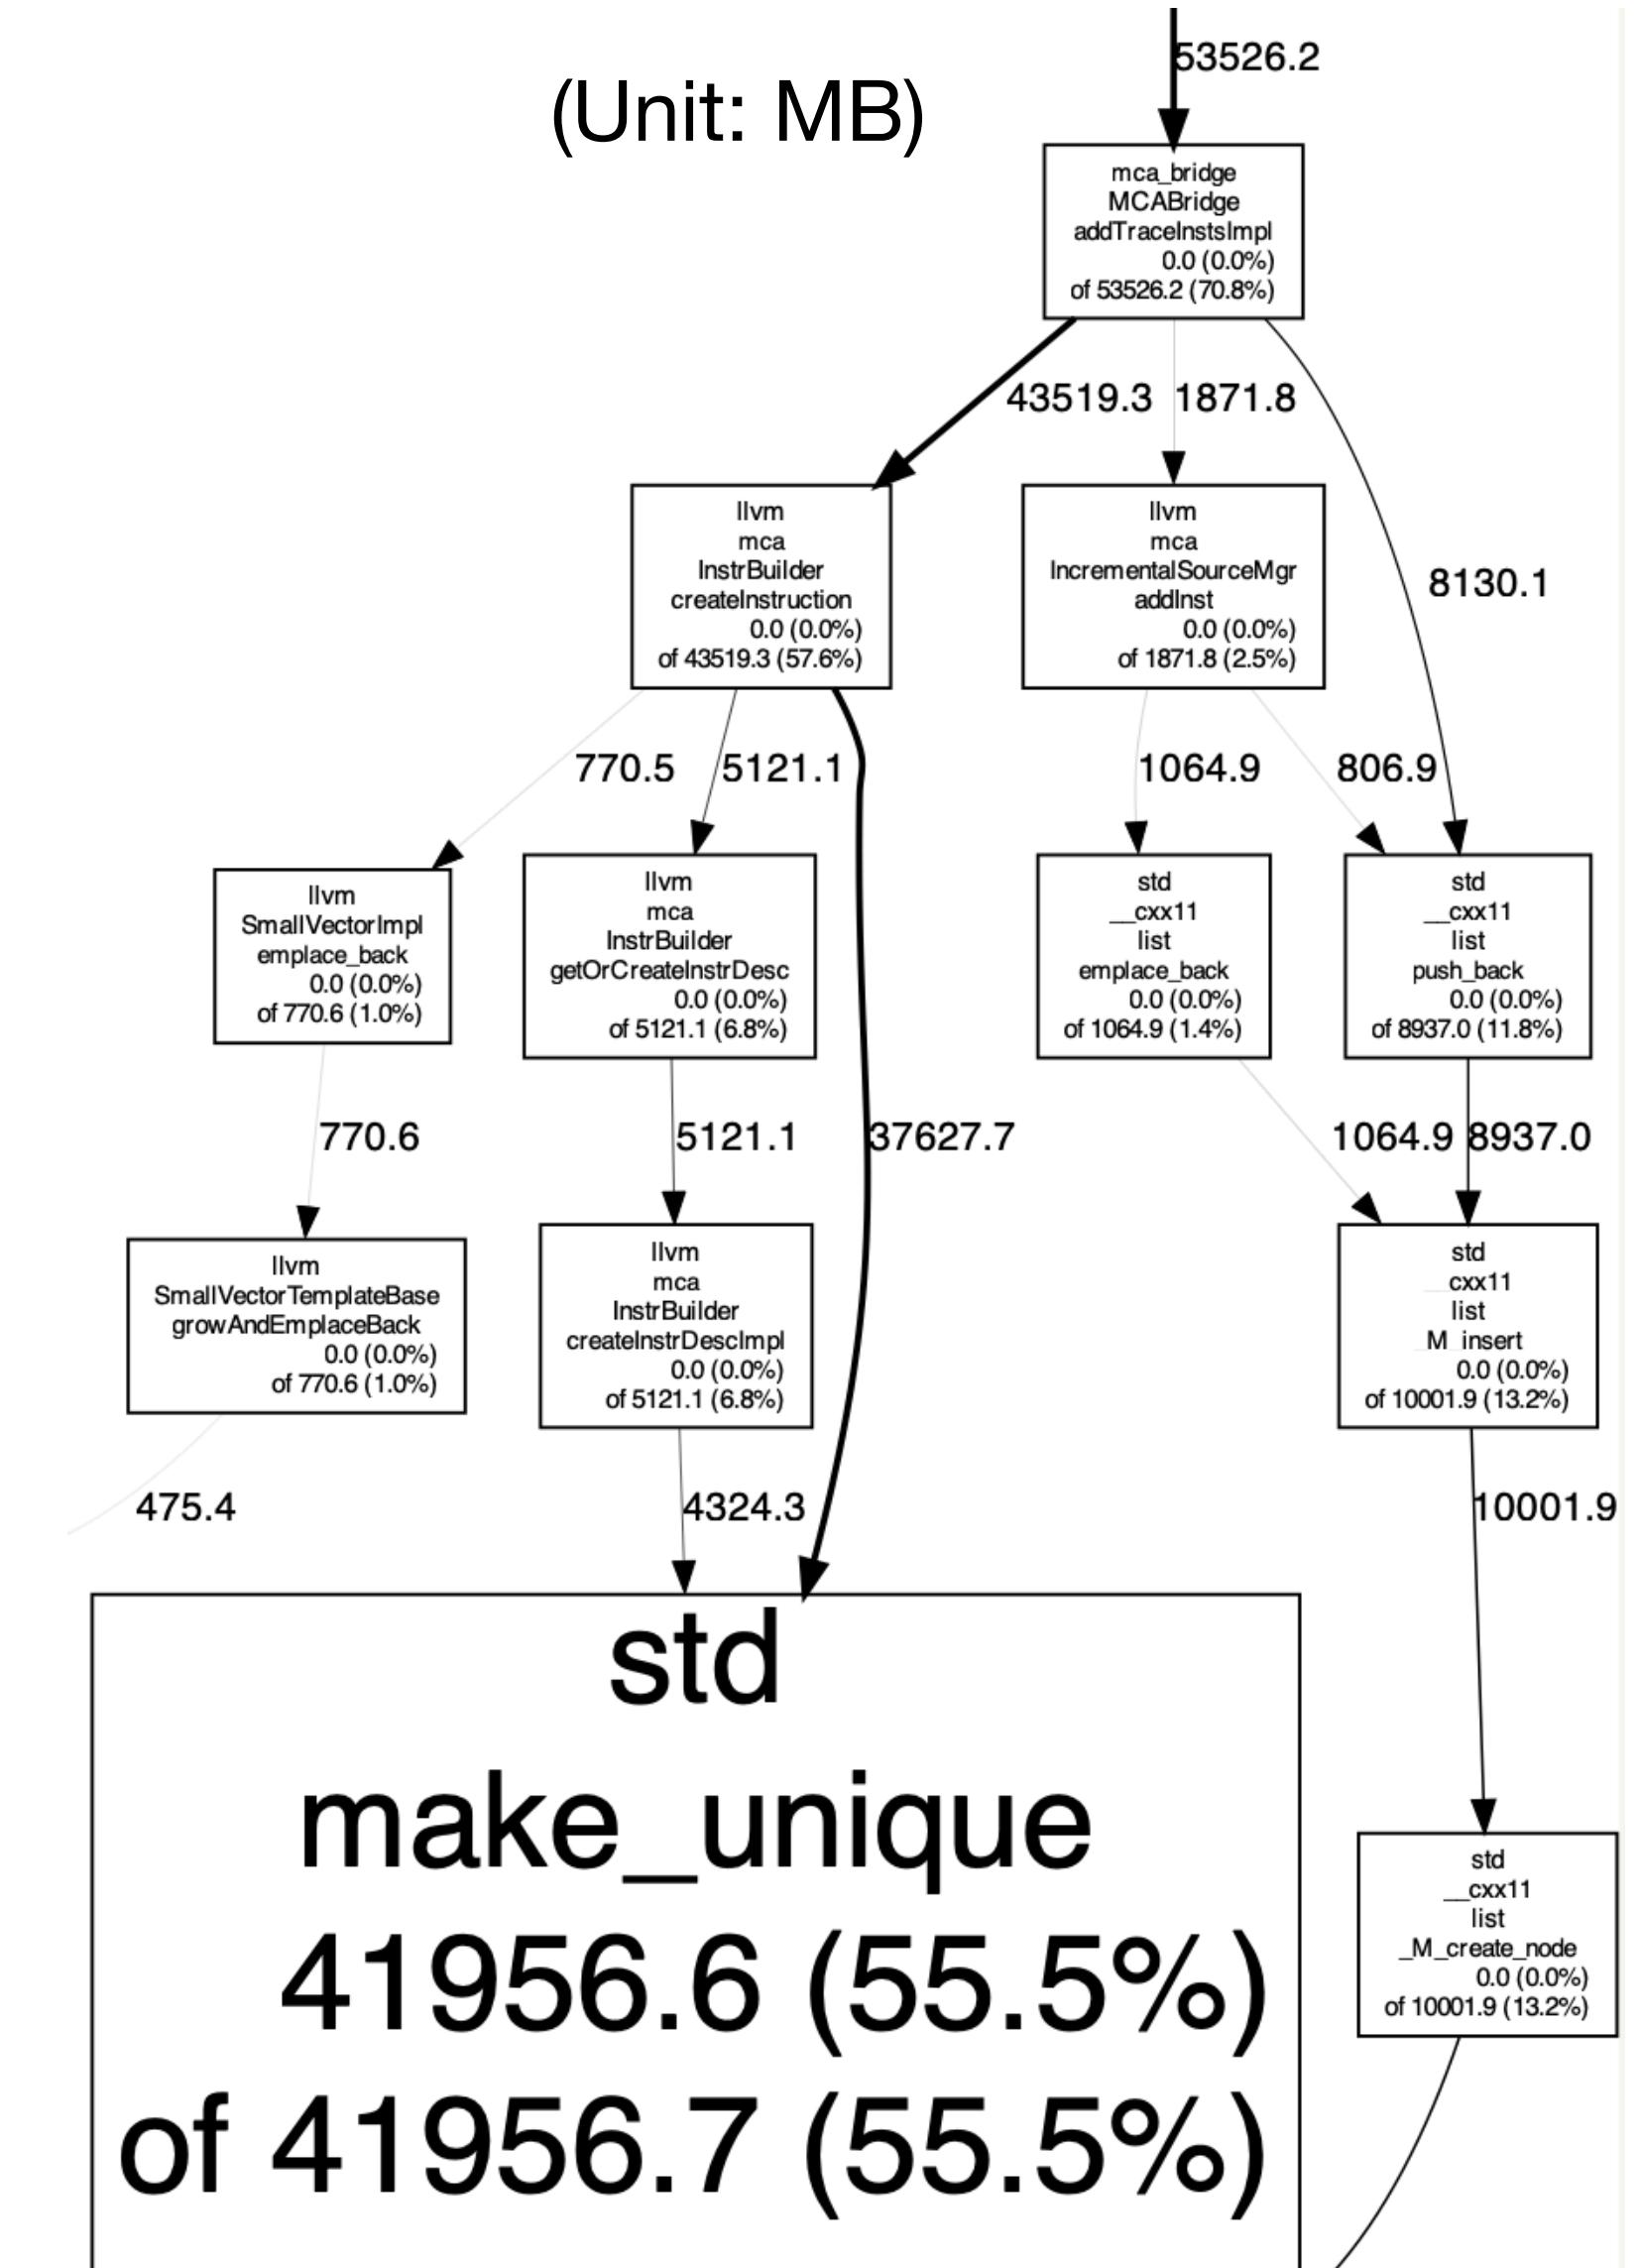
# Challenge

# Challenge

- Created a significant amount of **memory footprint**

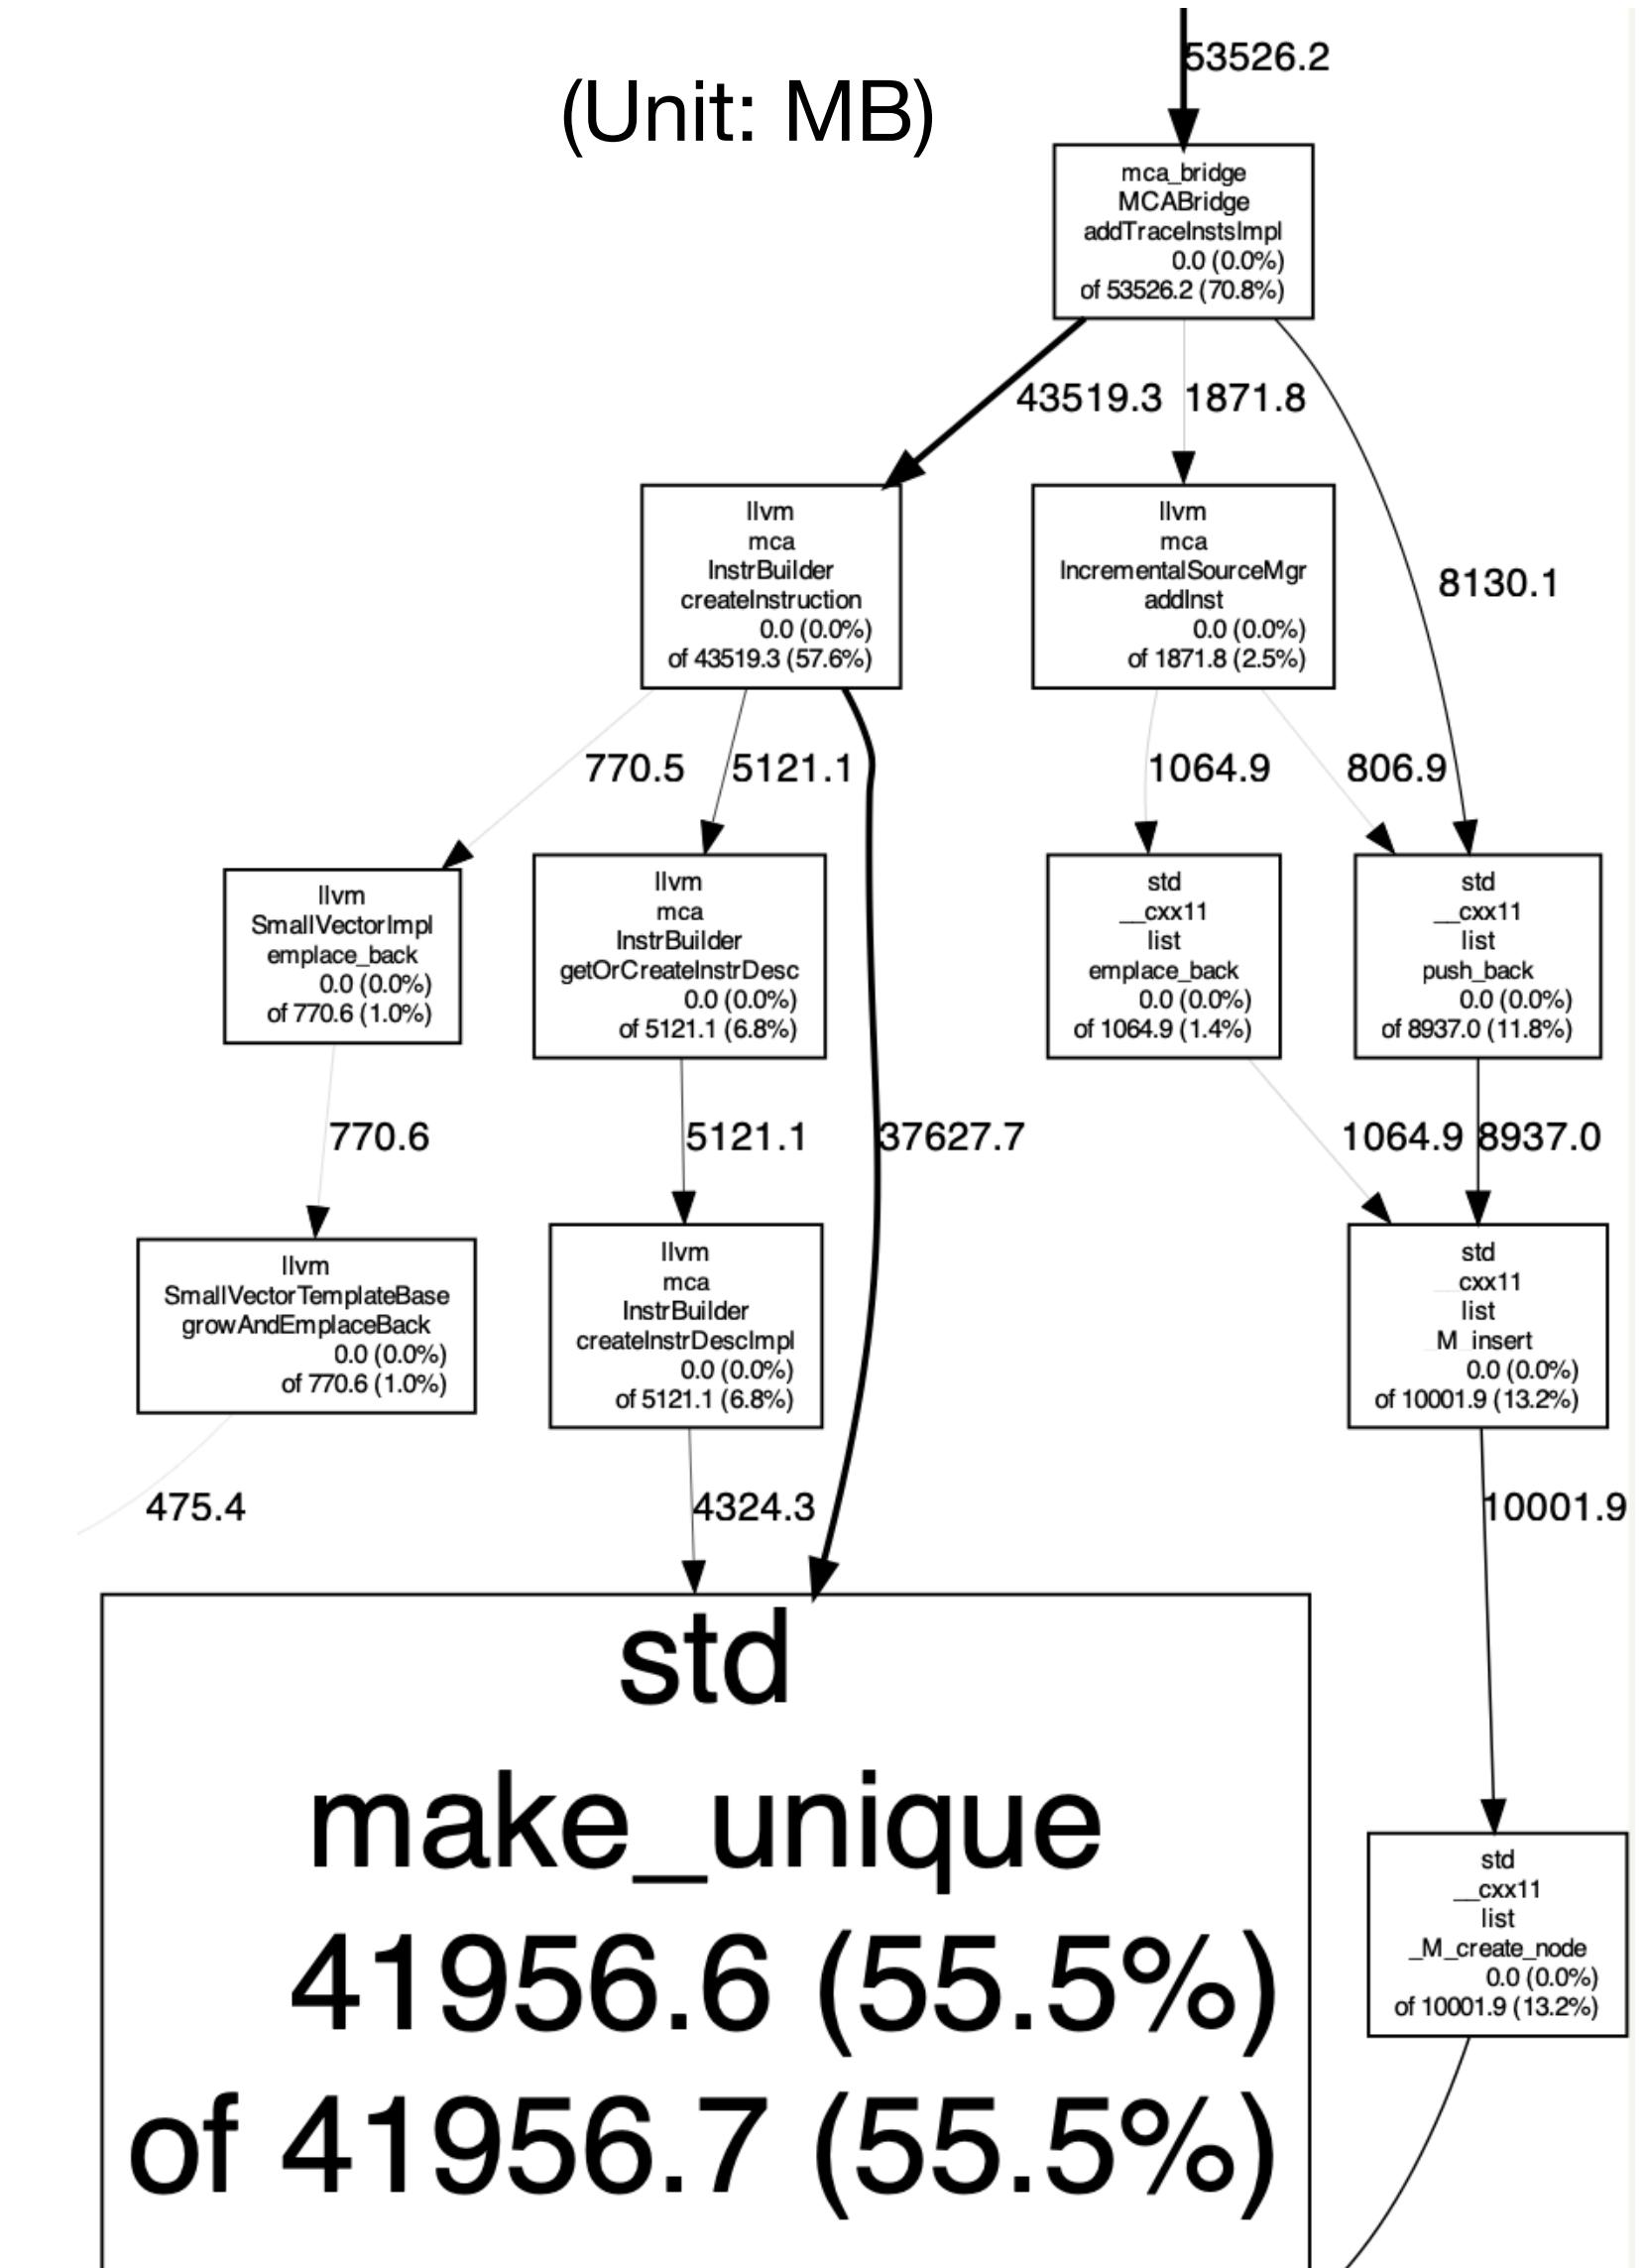
# Challenge

- Created a significant amount of **memory footprint**



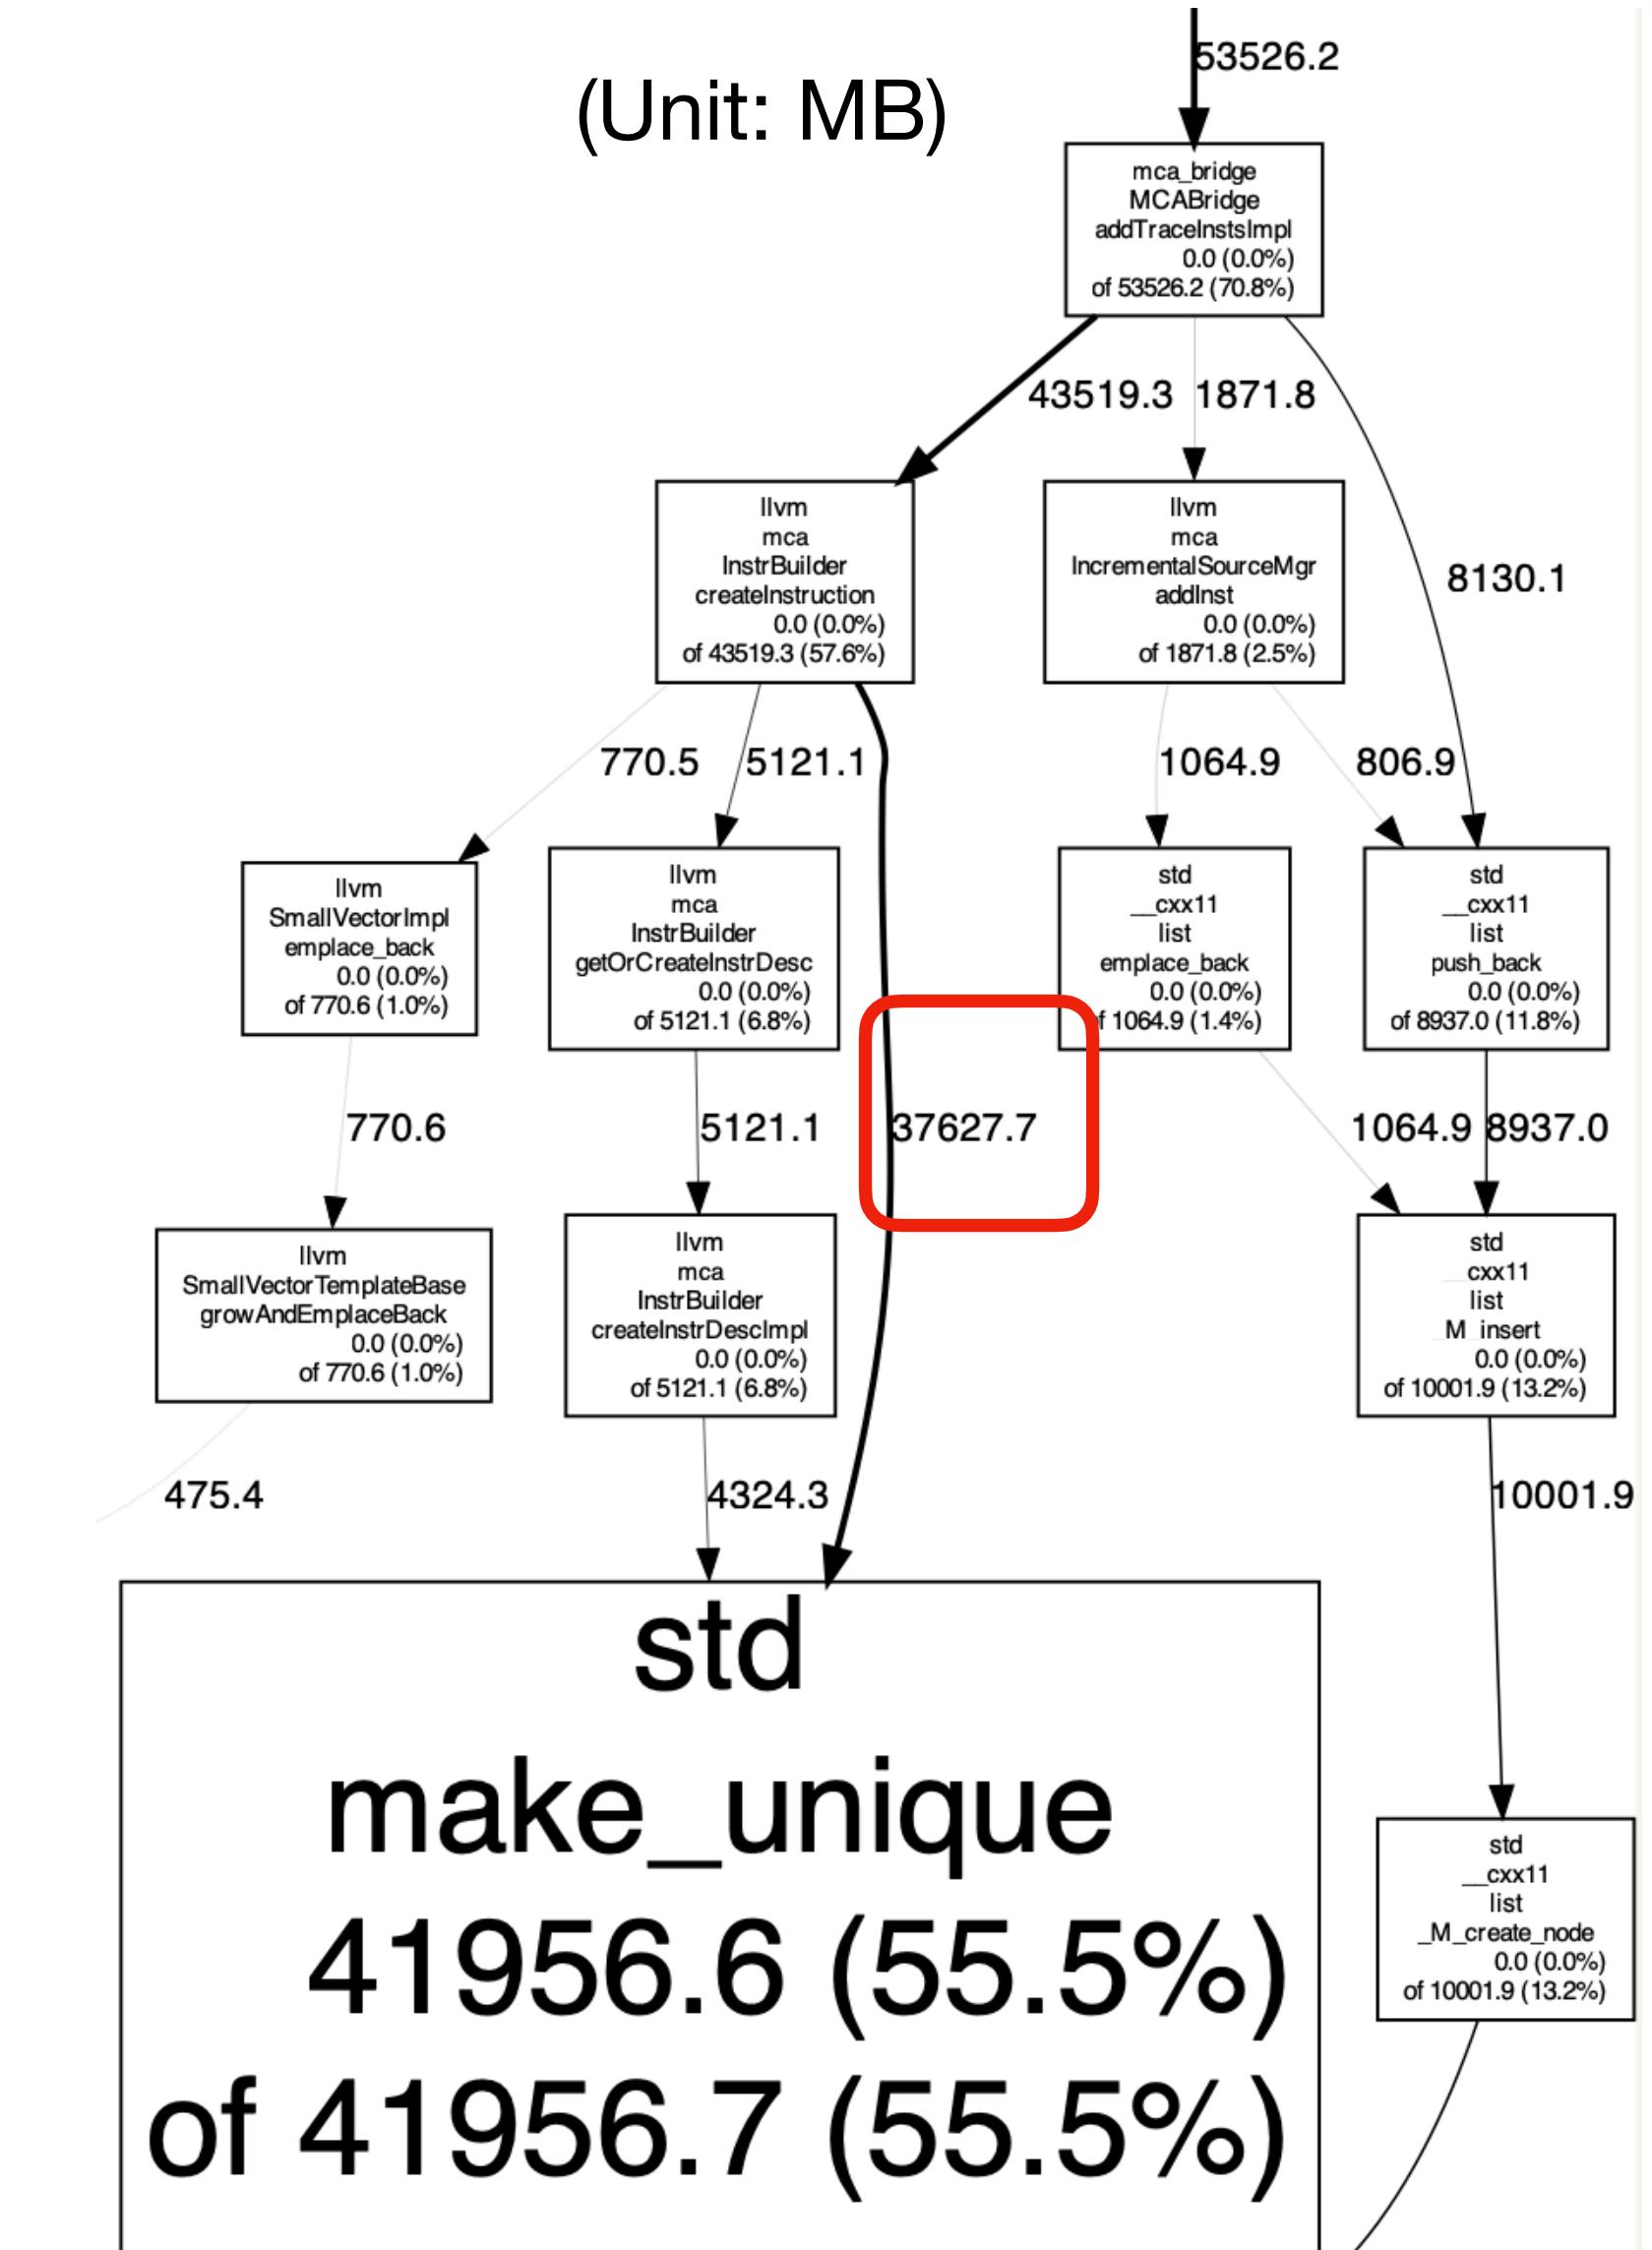
# Challenge

- Created a significant amount of **memory footprint**
- Bottleneck: **~37GB** of accumulated (virtual) memory was allocated by mca::InstrBuilder::createInstruction



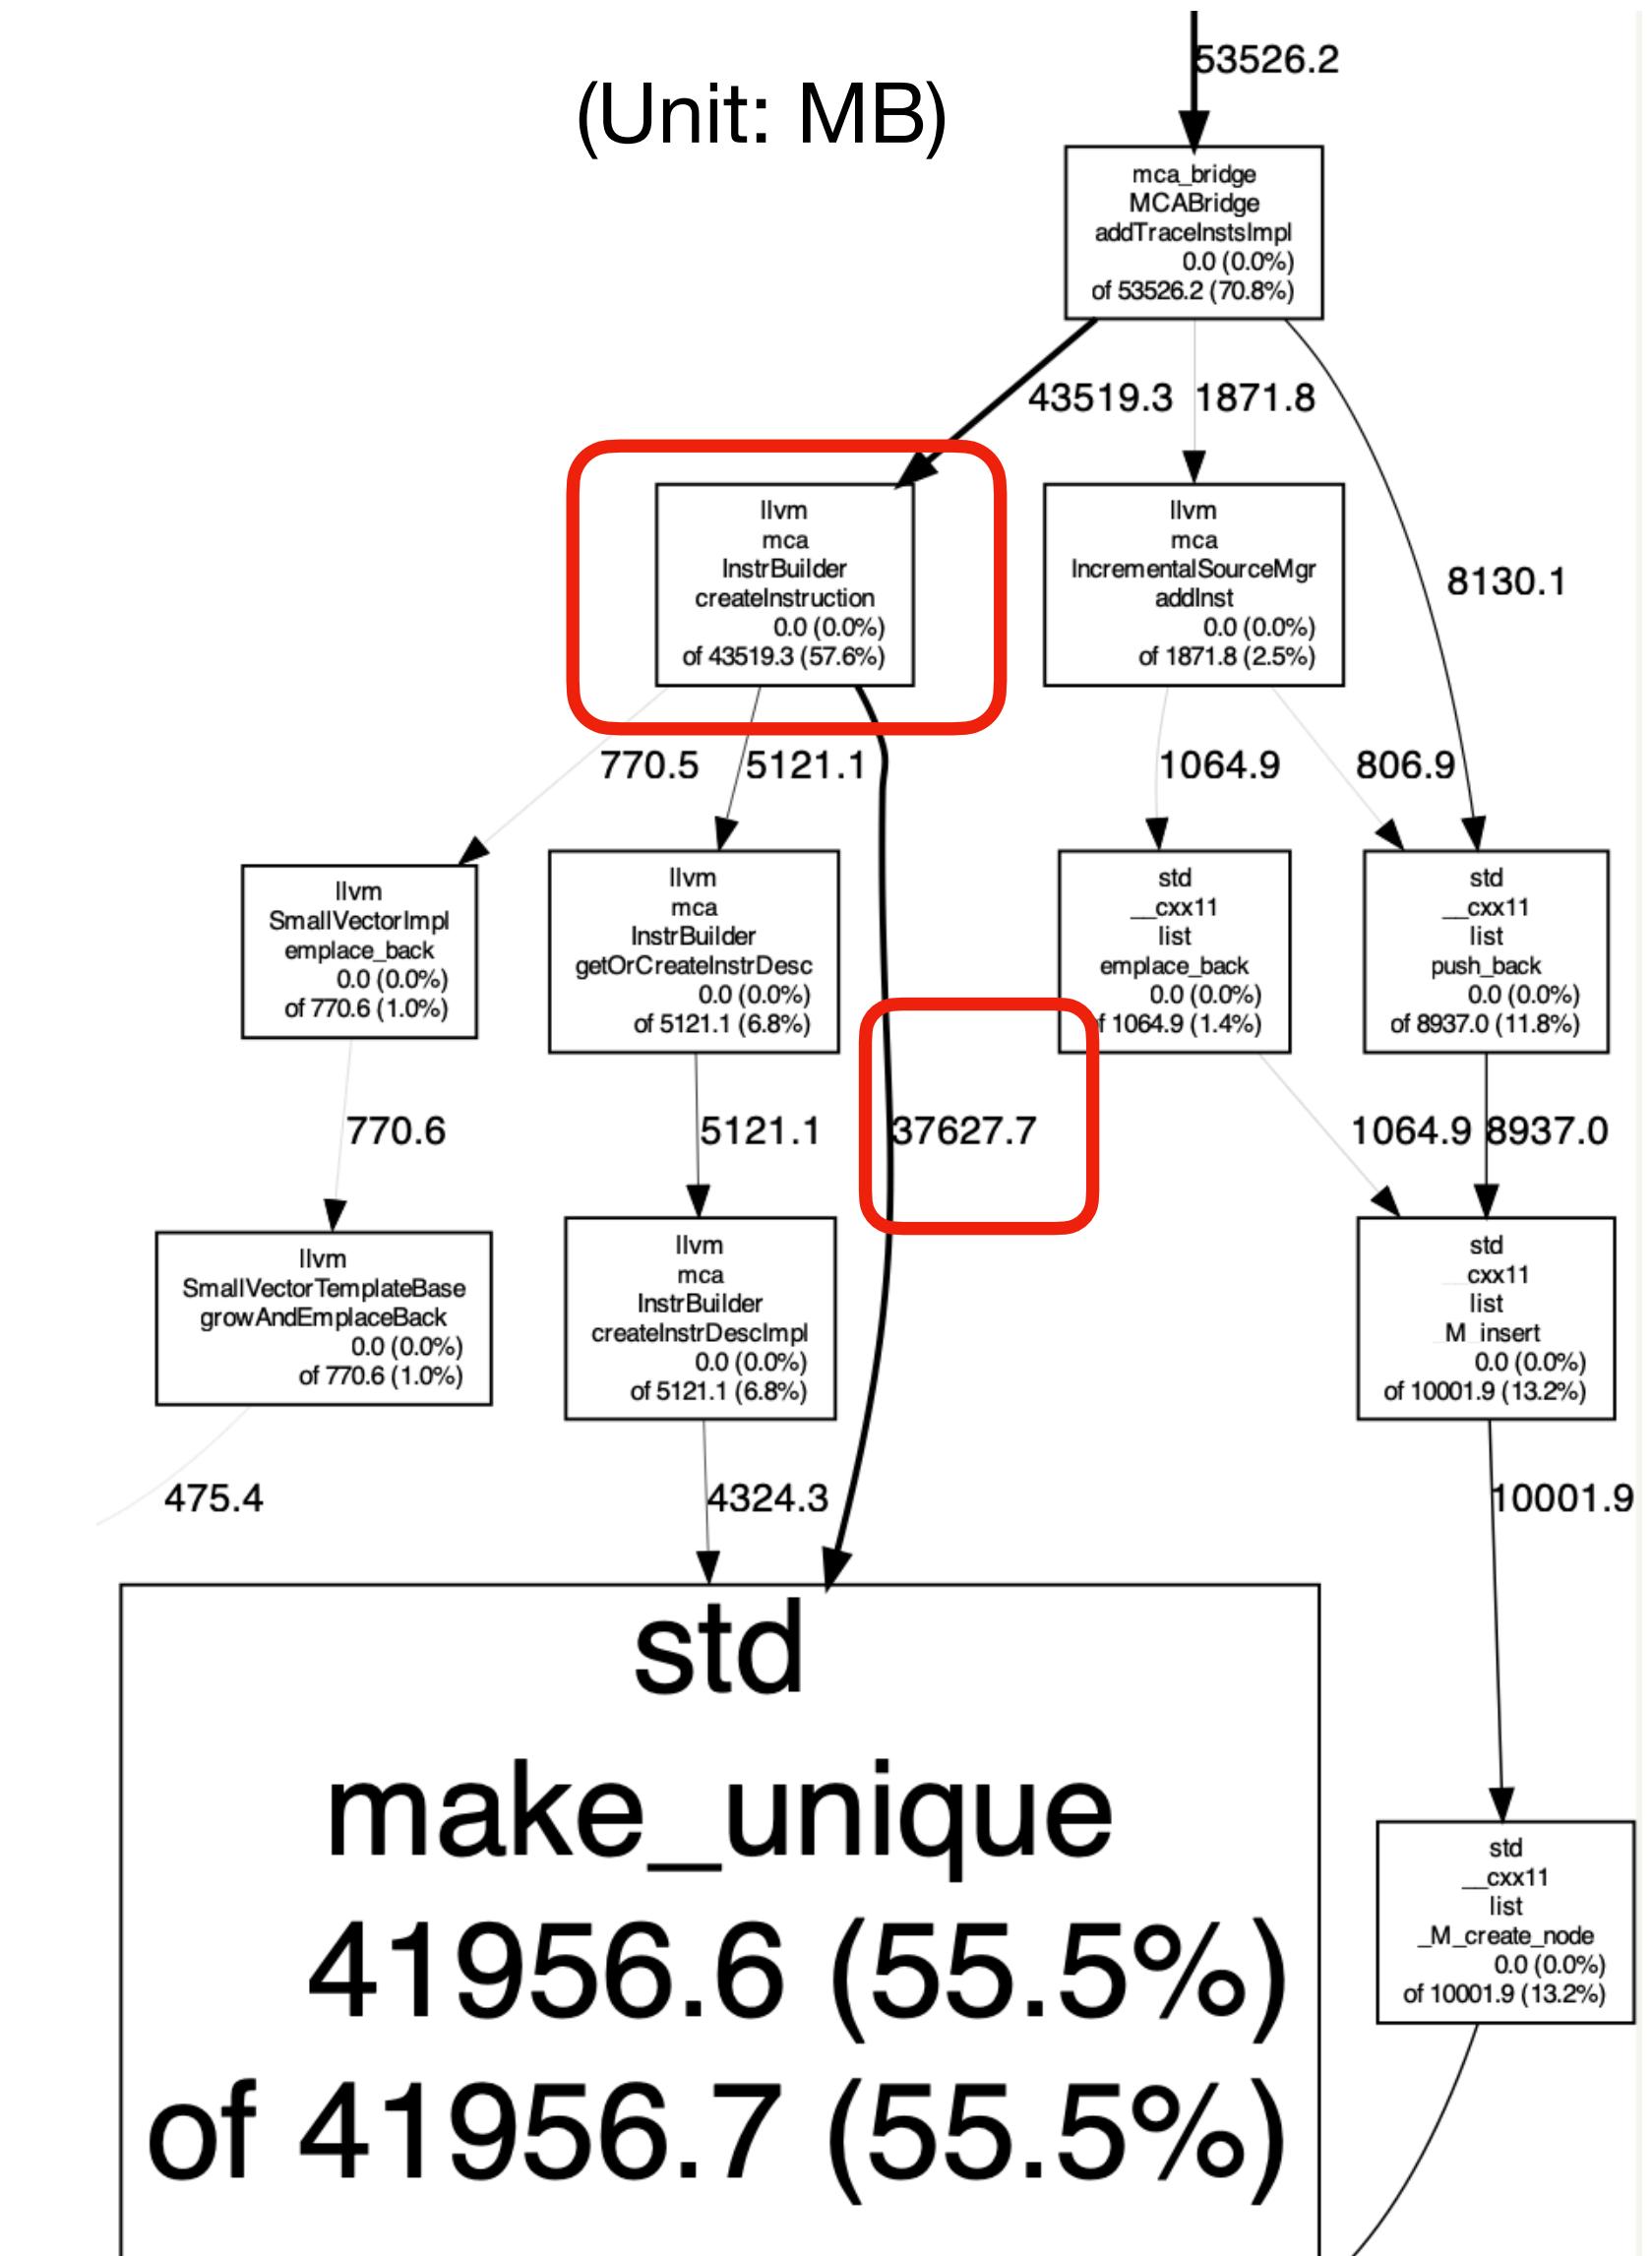
# Challenge

- Created a significant amount of **memory footprint**
- Bottleneck: **~37GB** of accumulated (virtual) memory was allocated by mca::InstrBuilder::createInstruction



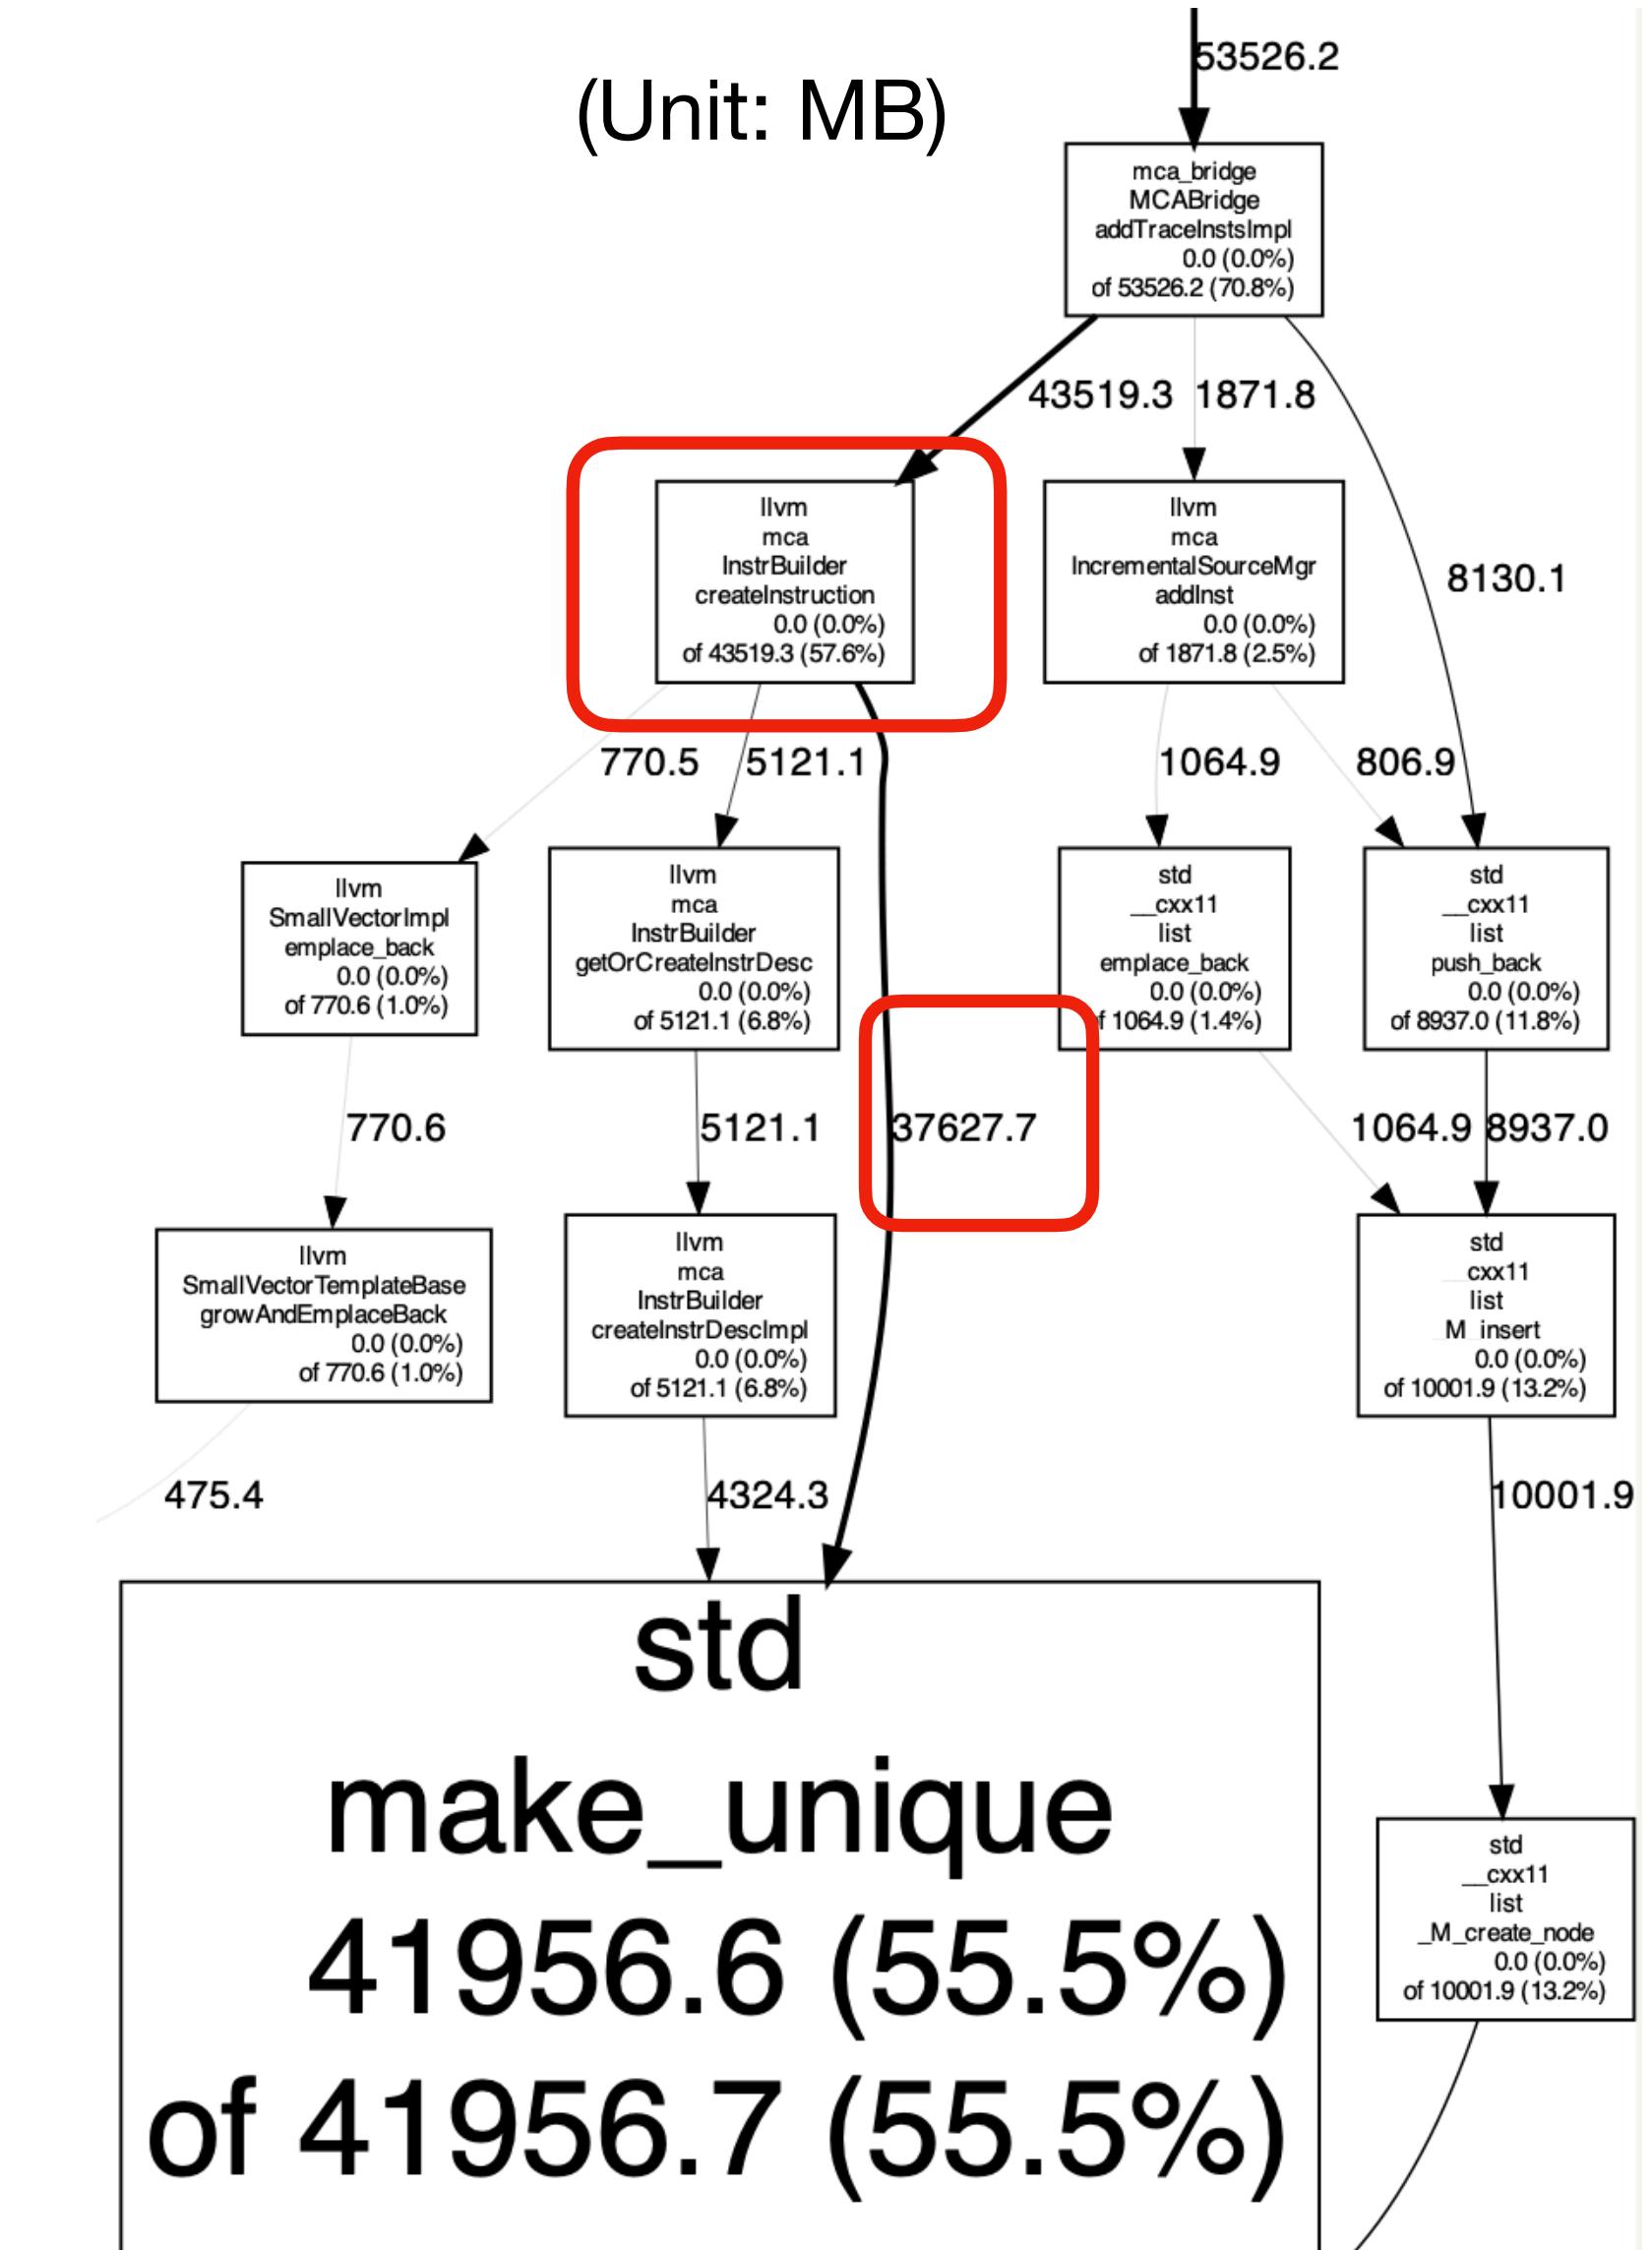
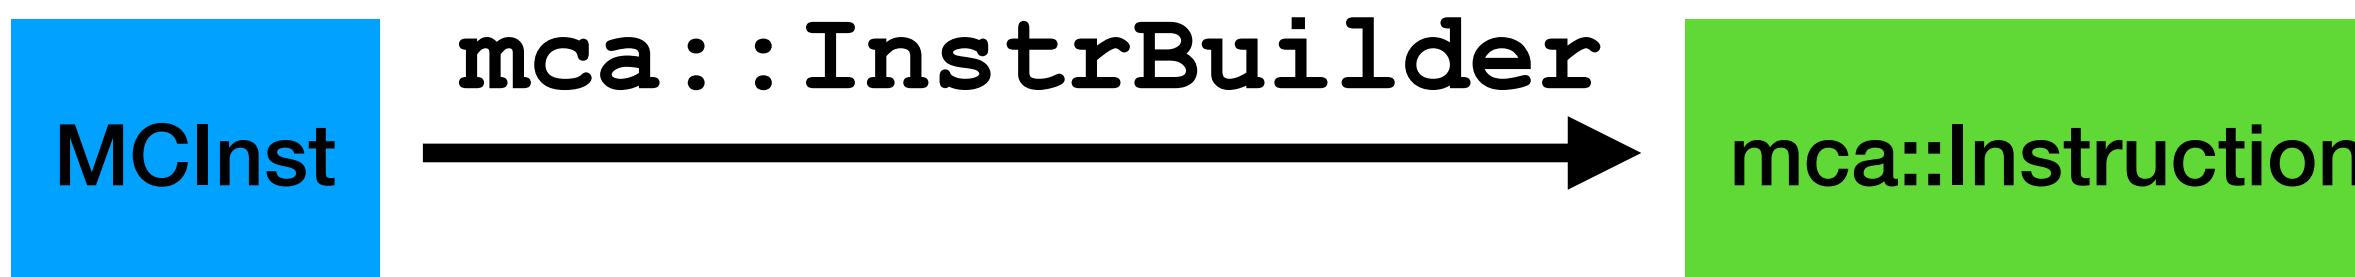
# Challenge

- Created a significant amount of **memory footprint**
- Bottleneck: **~37GB** of accumulated (virtual) memory was allocated by mca::InstrBuilder::createInstruction



# Challenge

- Created a significant amount of **memory footprint**
- Bottleneck: **~37GB** of accumulated (virtual) memory was allocated by mca::InstrBuilder::createInstruction



# **Large memory footprint**

## **Root cause**

# Large memory footprint

## Root cause

- Most of the translated `mca::Instruction` objects are **never** deallocated until the simulation is finished

# Large memory footprint

## Root cause

- Most of the translated `mca::Instruction` objects are **never** deallocated until the simulation is finished
- `mca::Instruction` is also used for tracking simulation state, so it's hard to make it immutable

# Large memory footprint

## Root cause

- Most of the translated `mca::Instruction` objects are **never** deallocated until the simulation is finished
- `mca::Instruction` is also used for tracking simulation state, so it's hard to make it immutable
- Doesn't scale really well with large input (recall: ~1 million instructions)

# Large memory footprint

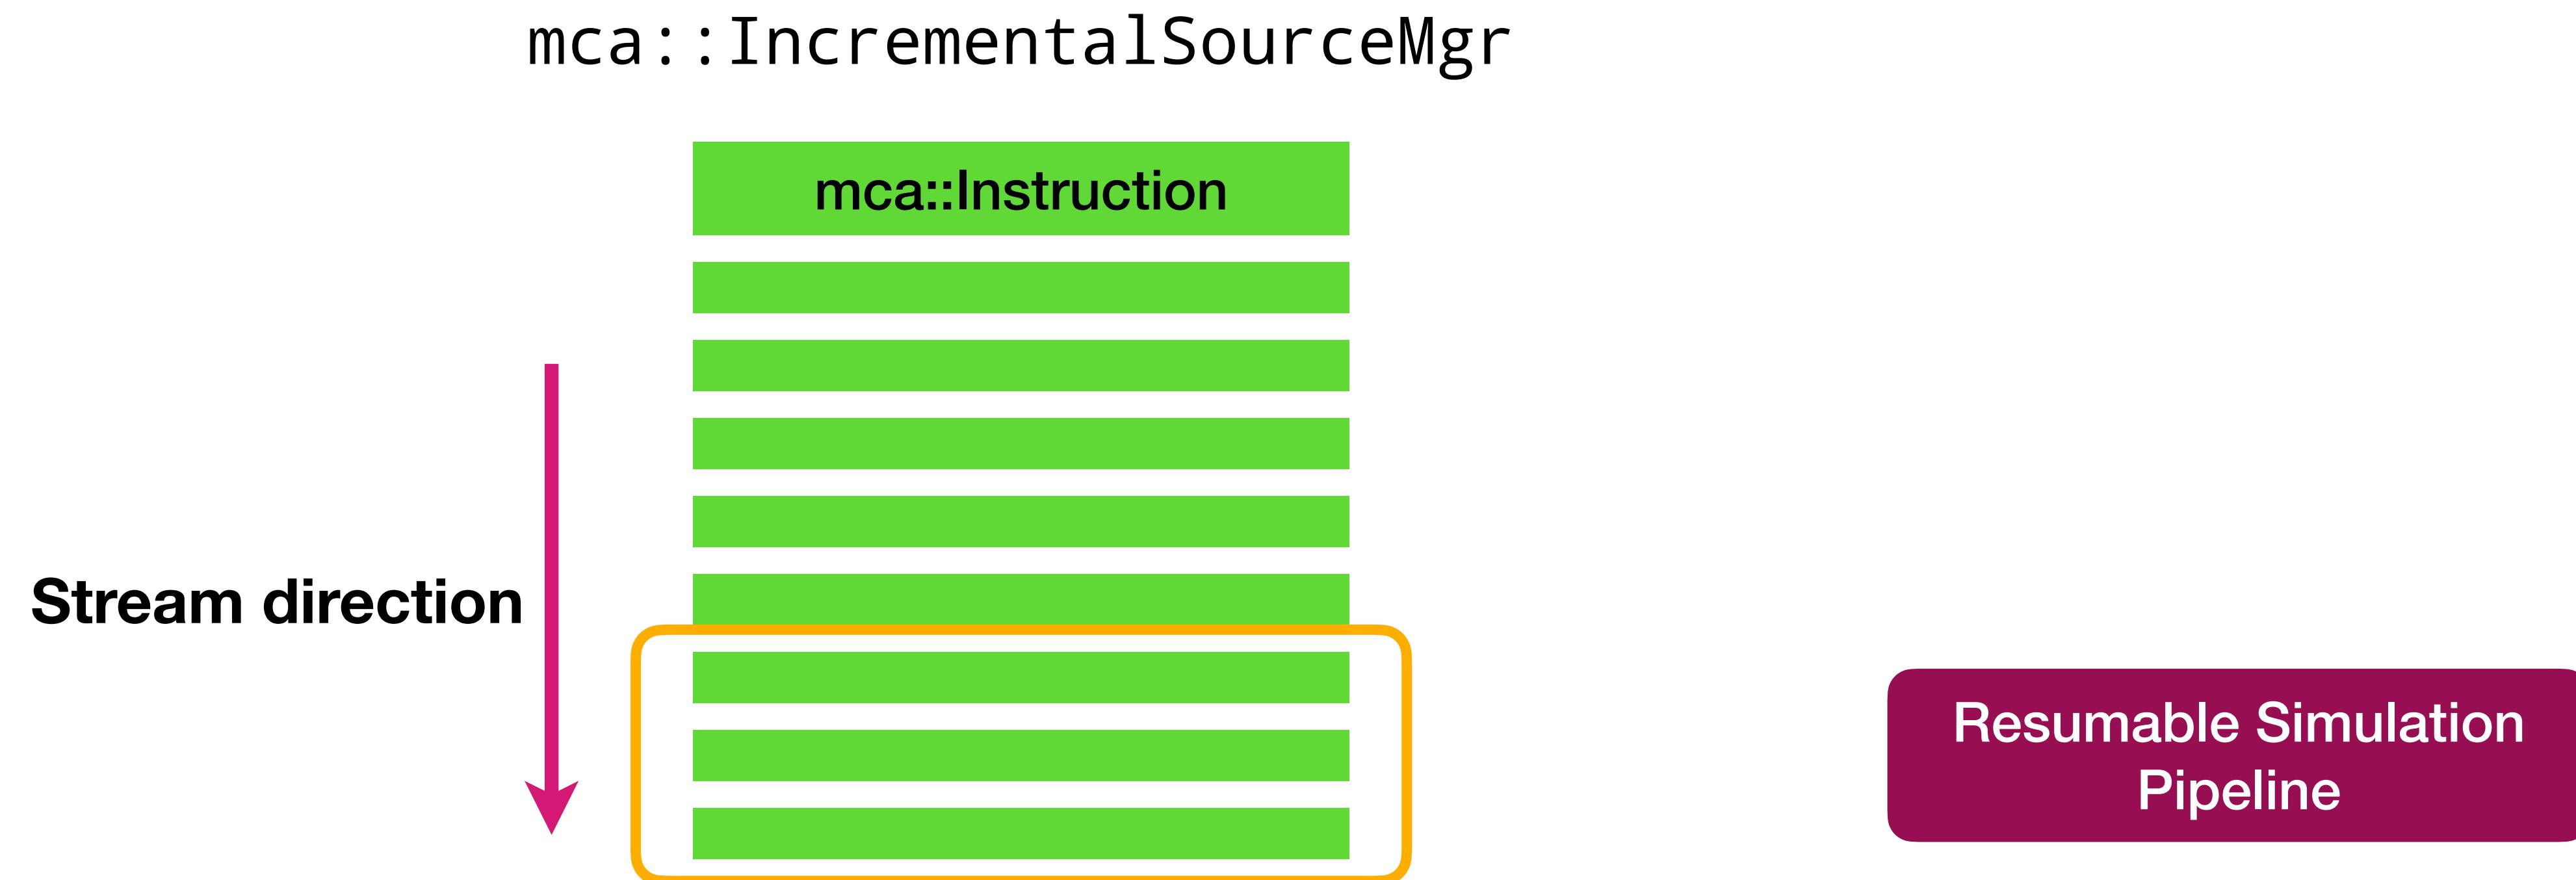
## Observation

mca::IncrementalSourceMgr



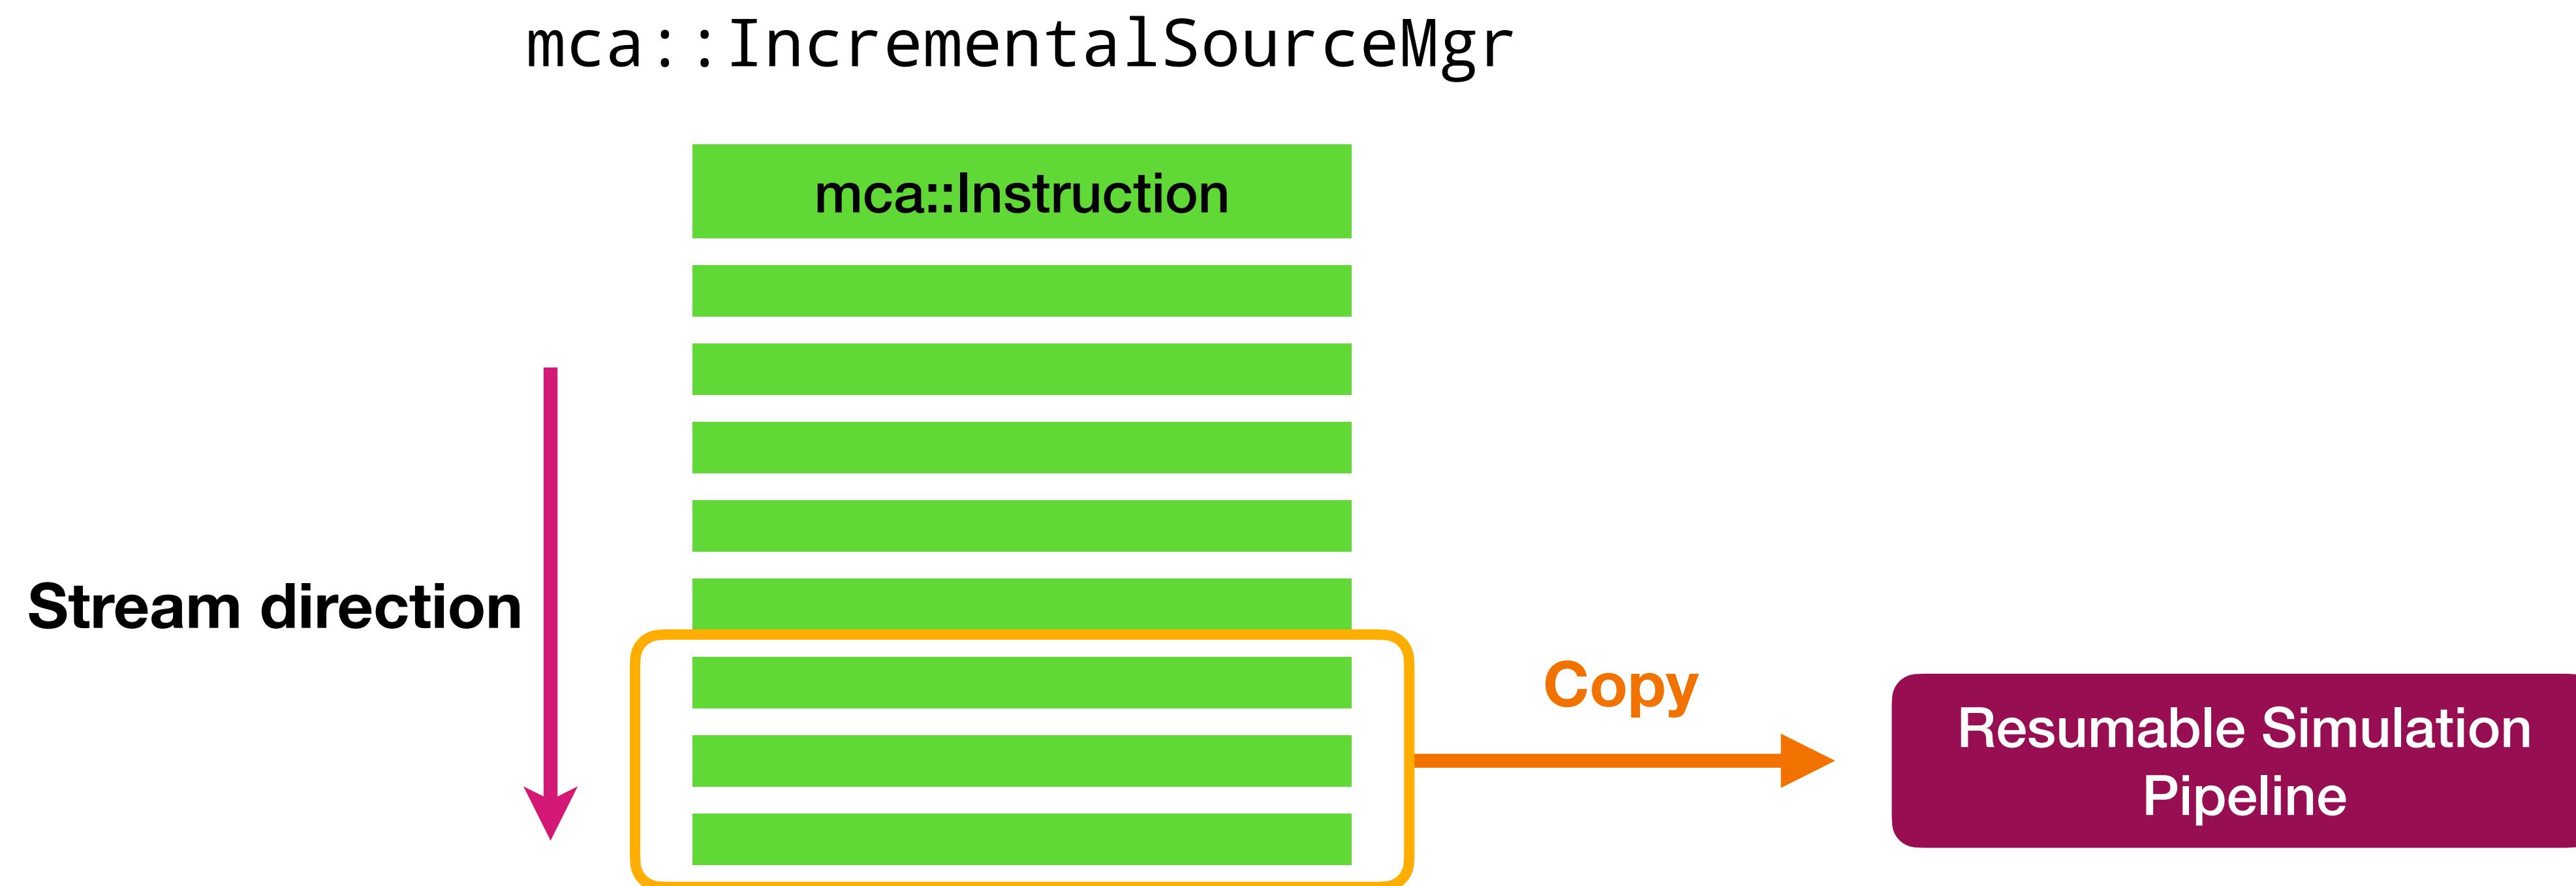
# Large memory footprint

## Observation



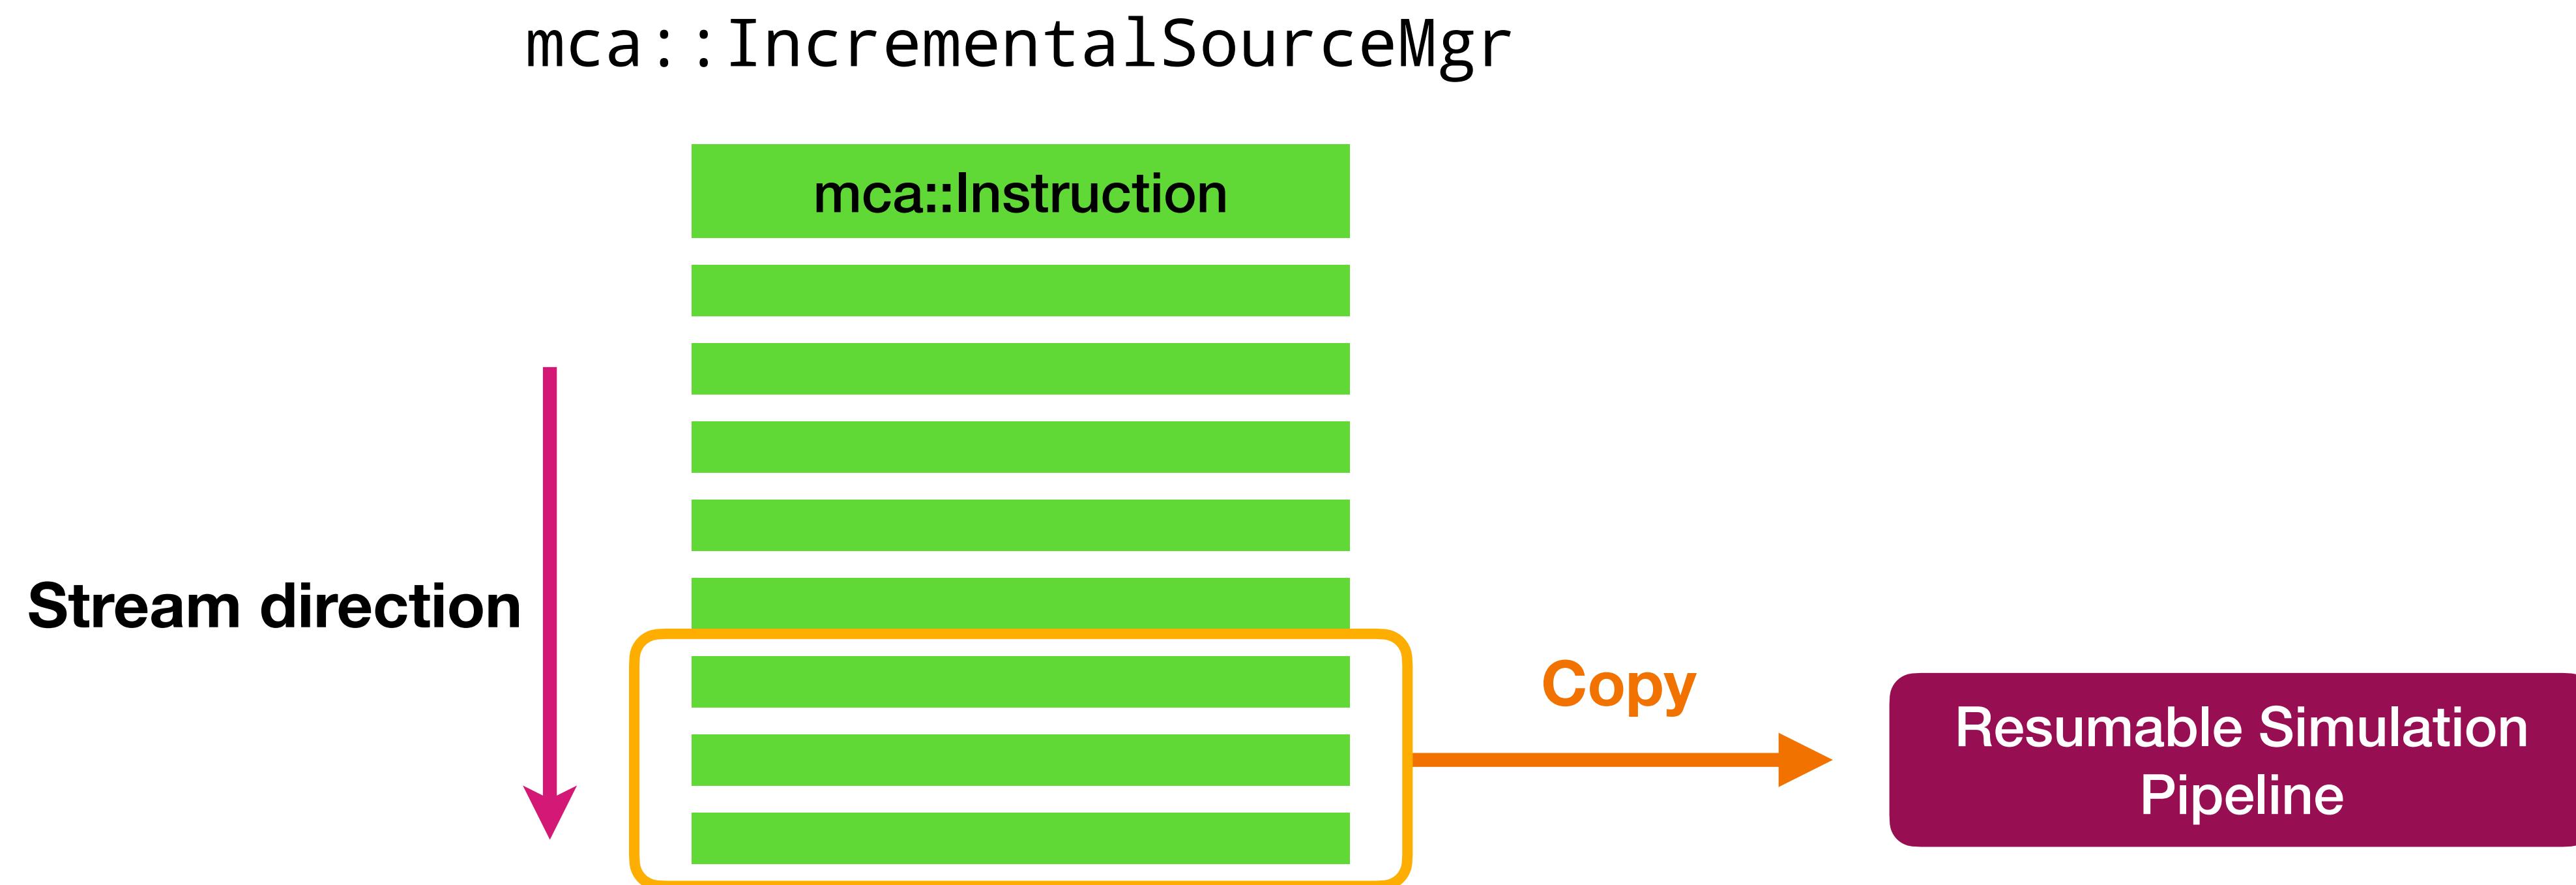
# Large memory footprint

## Observation



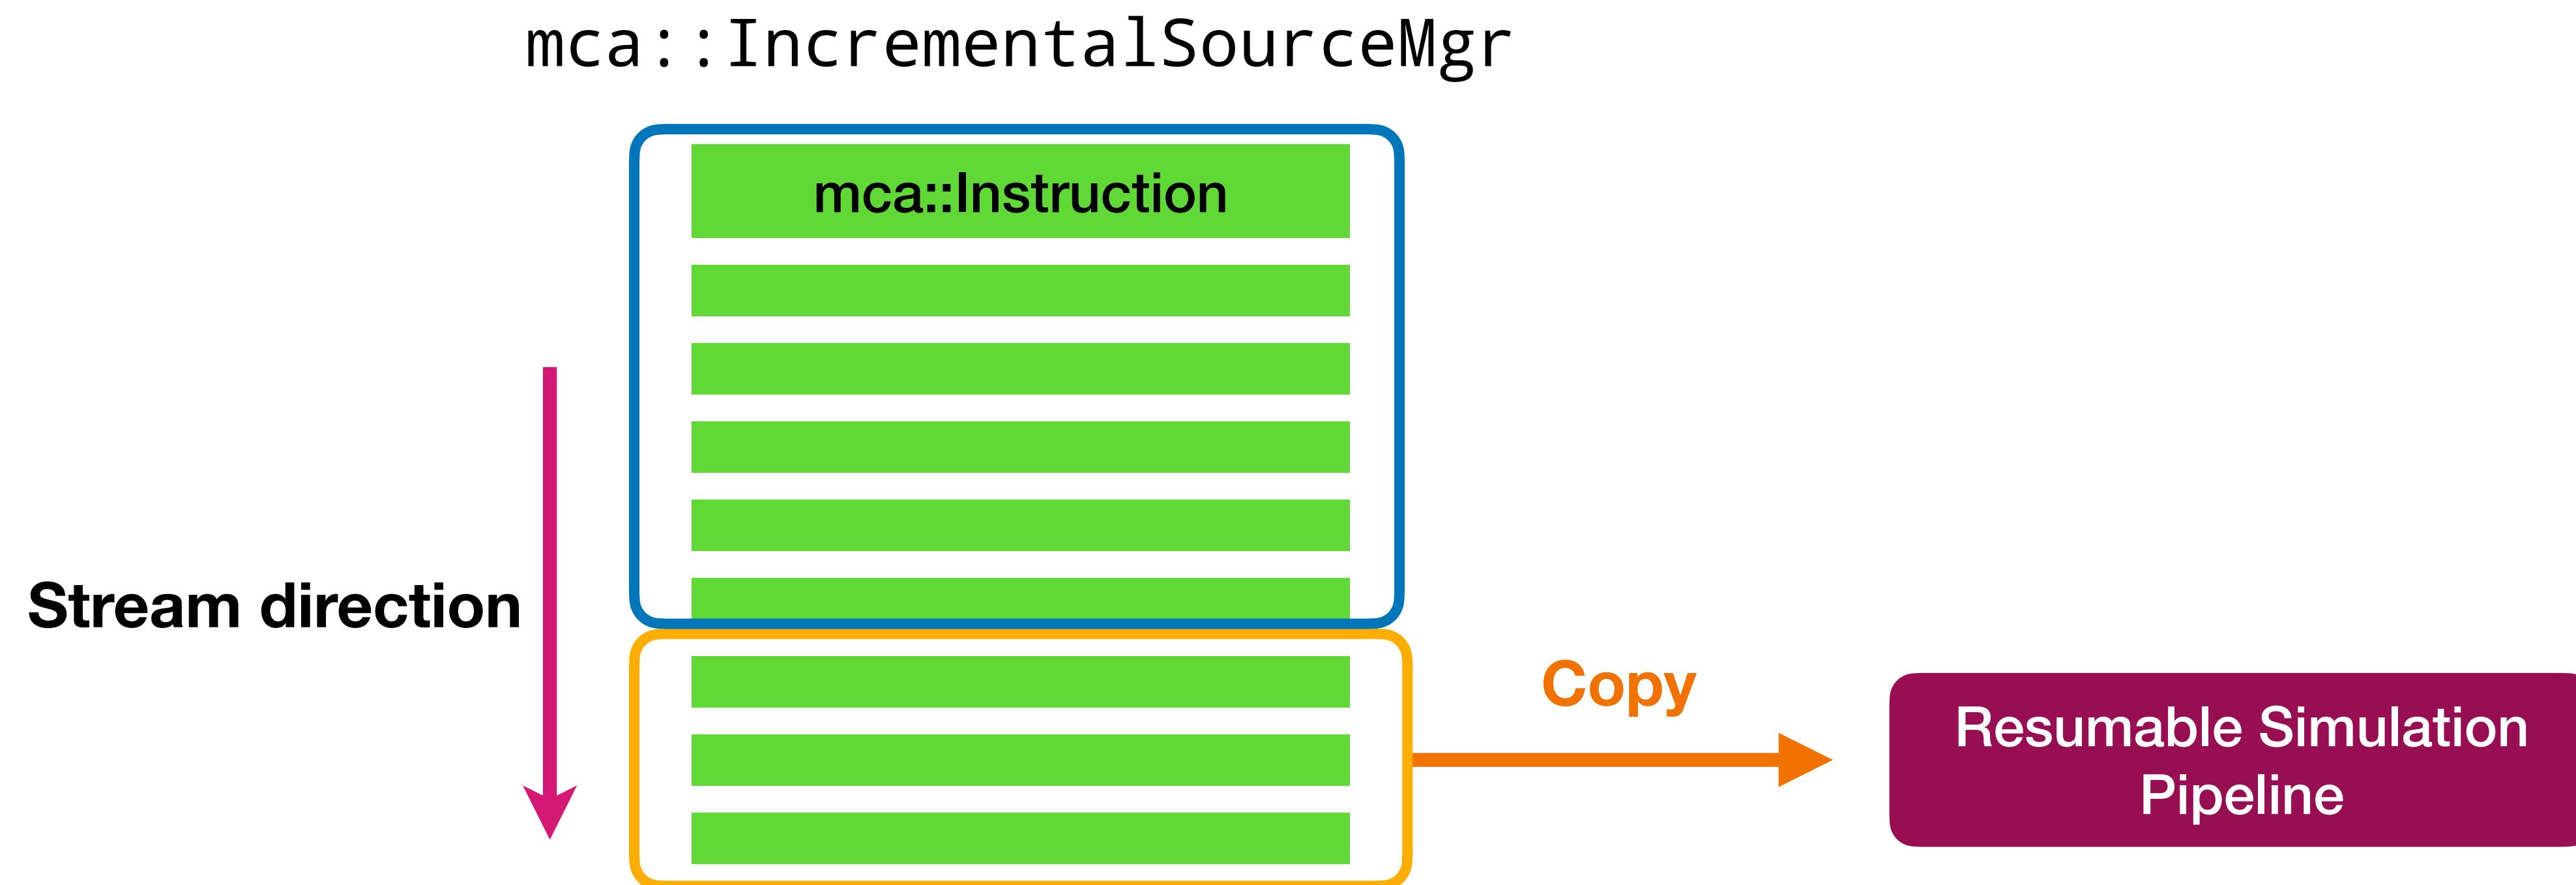
# Large memory footprint

Solution: Recycling mca::Instruction



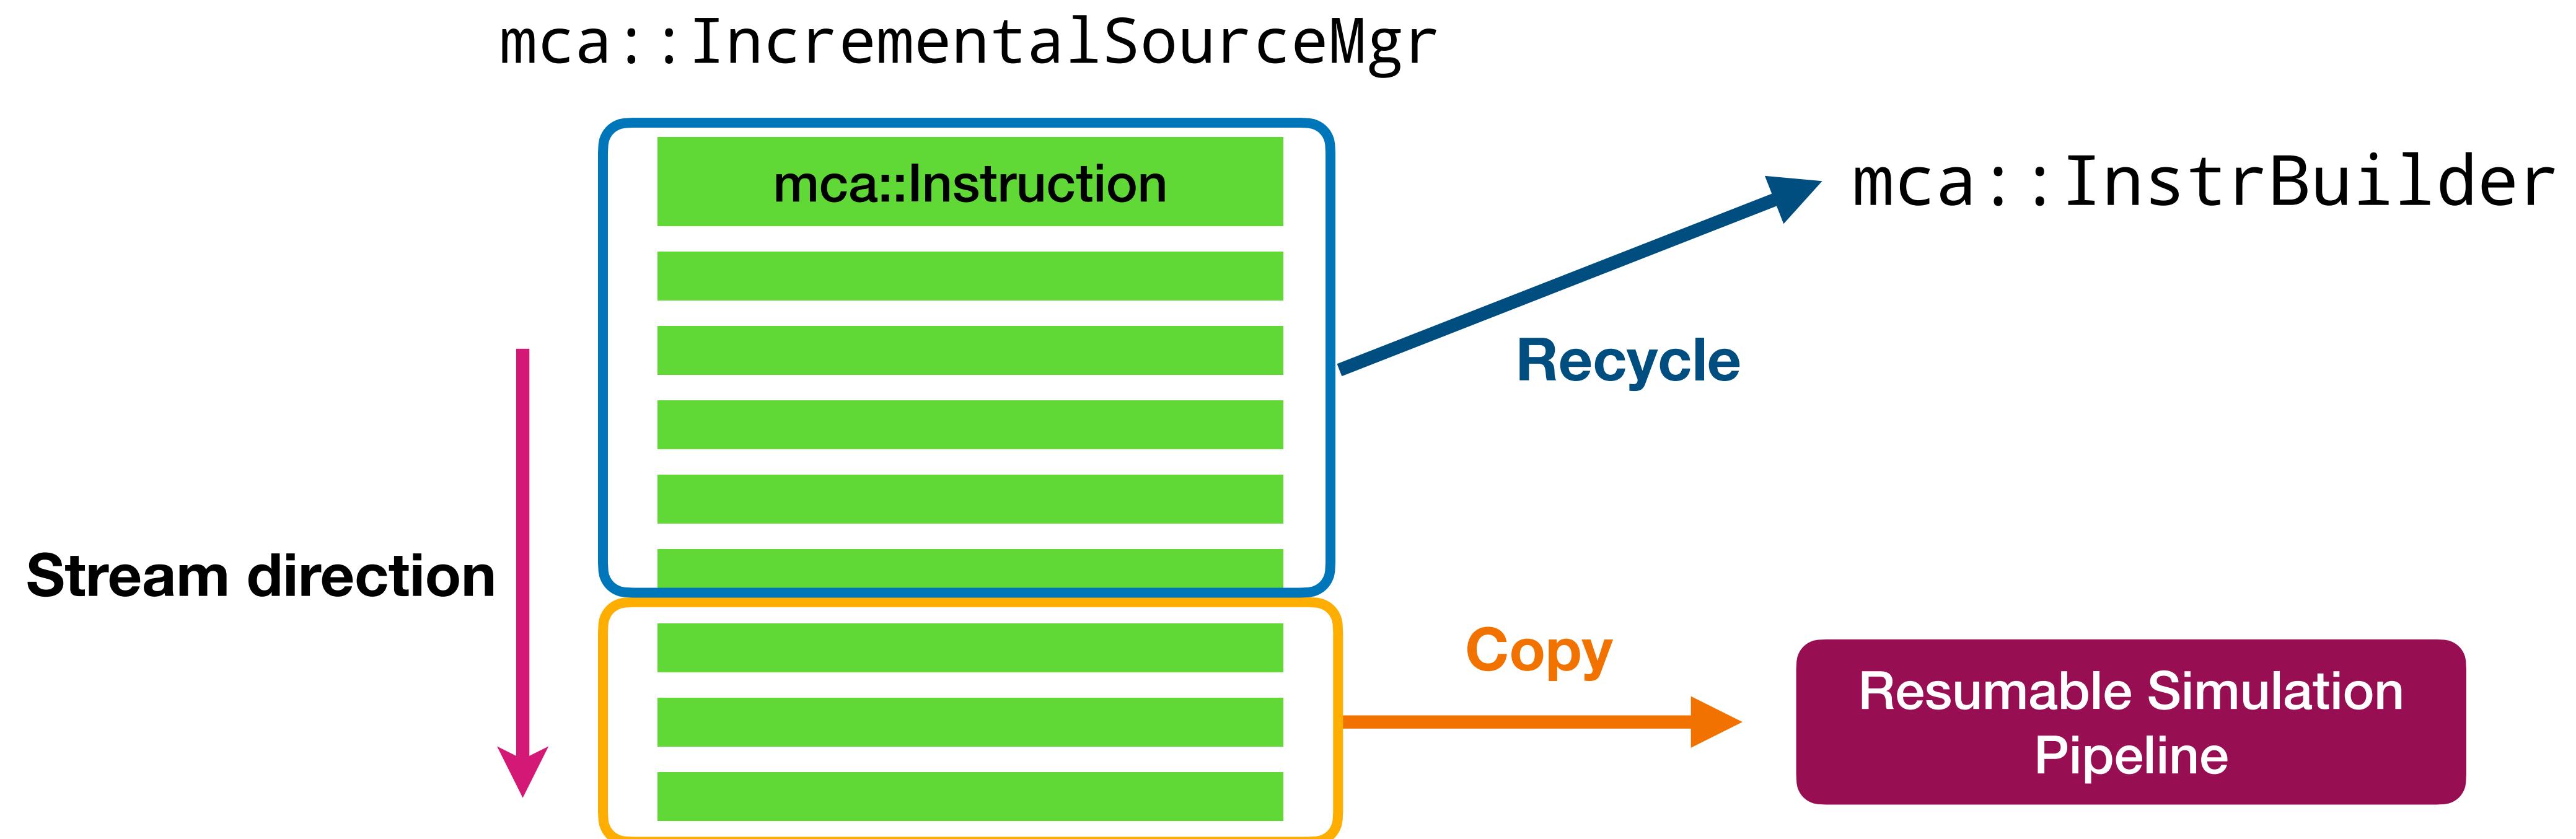
# Large memory footprint

Solution: Recycling mca::Instruction



# Large memory footprint

Solution: Recycling mca::Instruction



# 67%

**improvement on accumulated memory consumption**

**~70%**

**of the mca::Instruction objects are recycled**

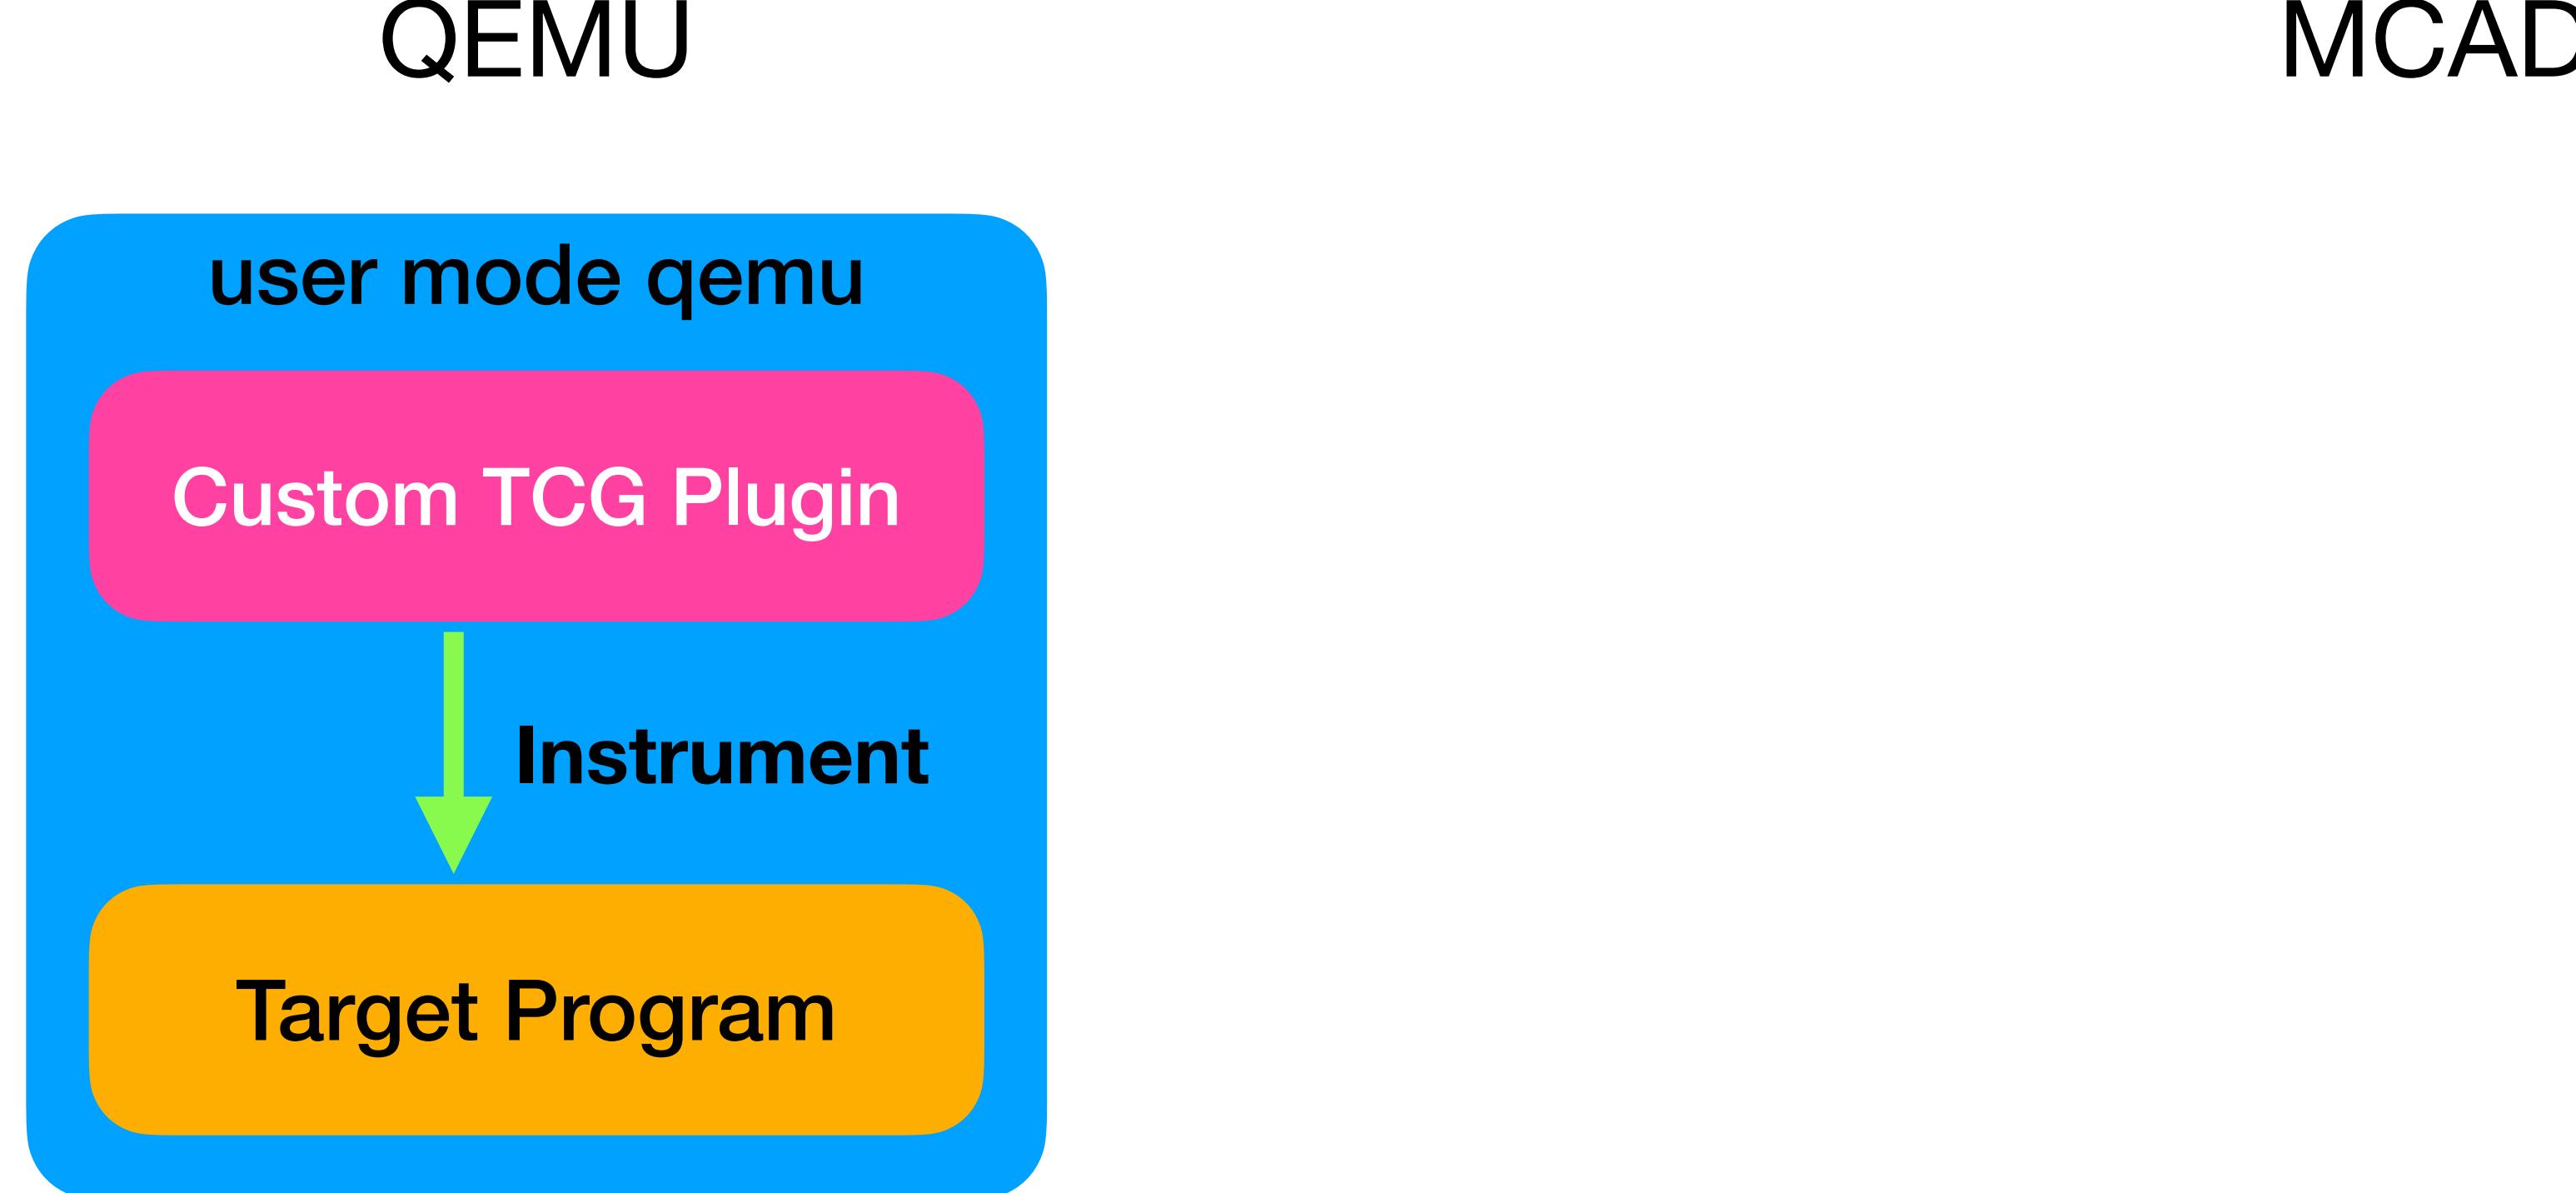
# Collecting execution traces via QEMU



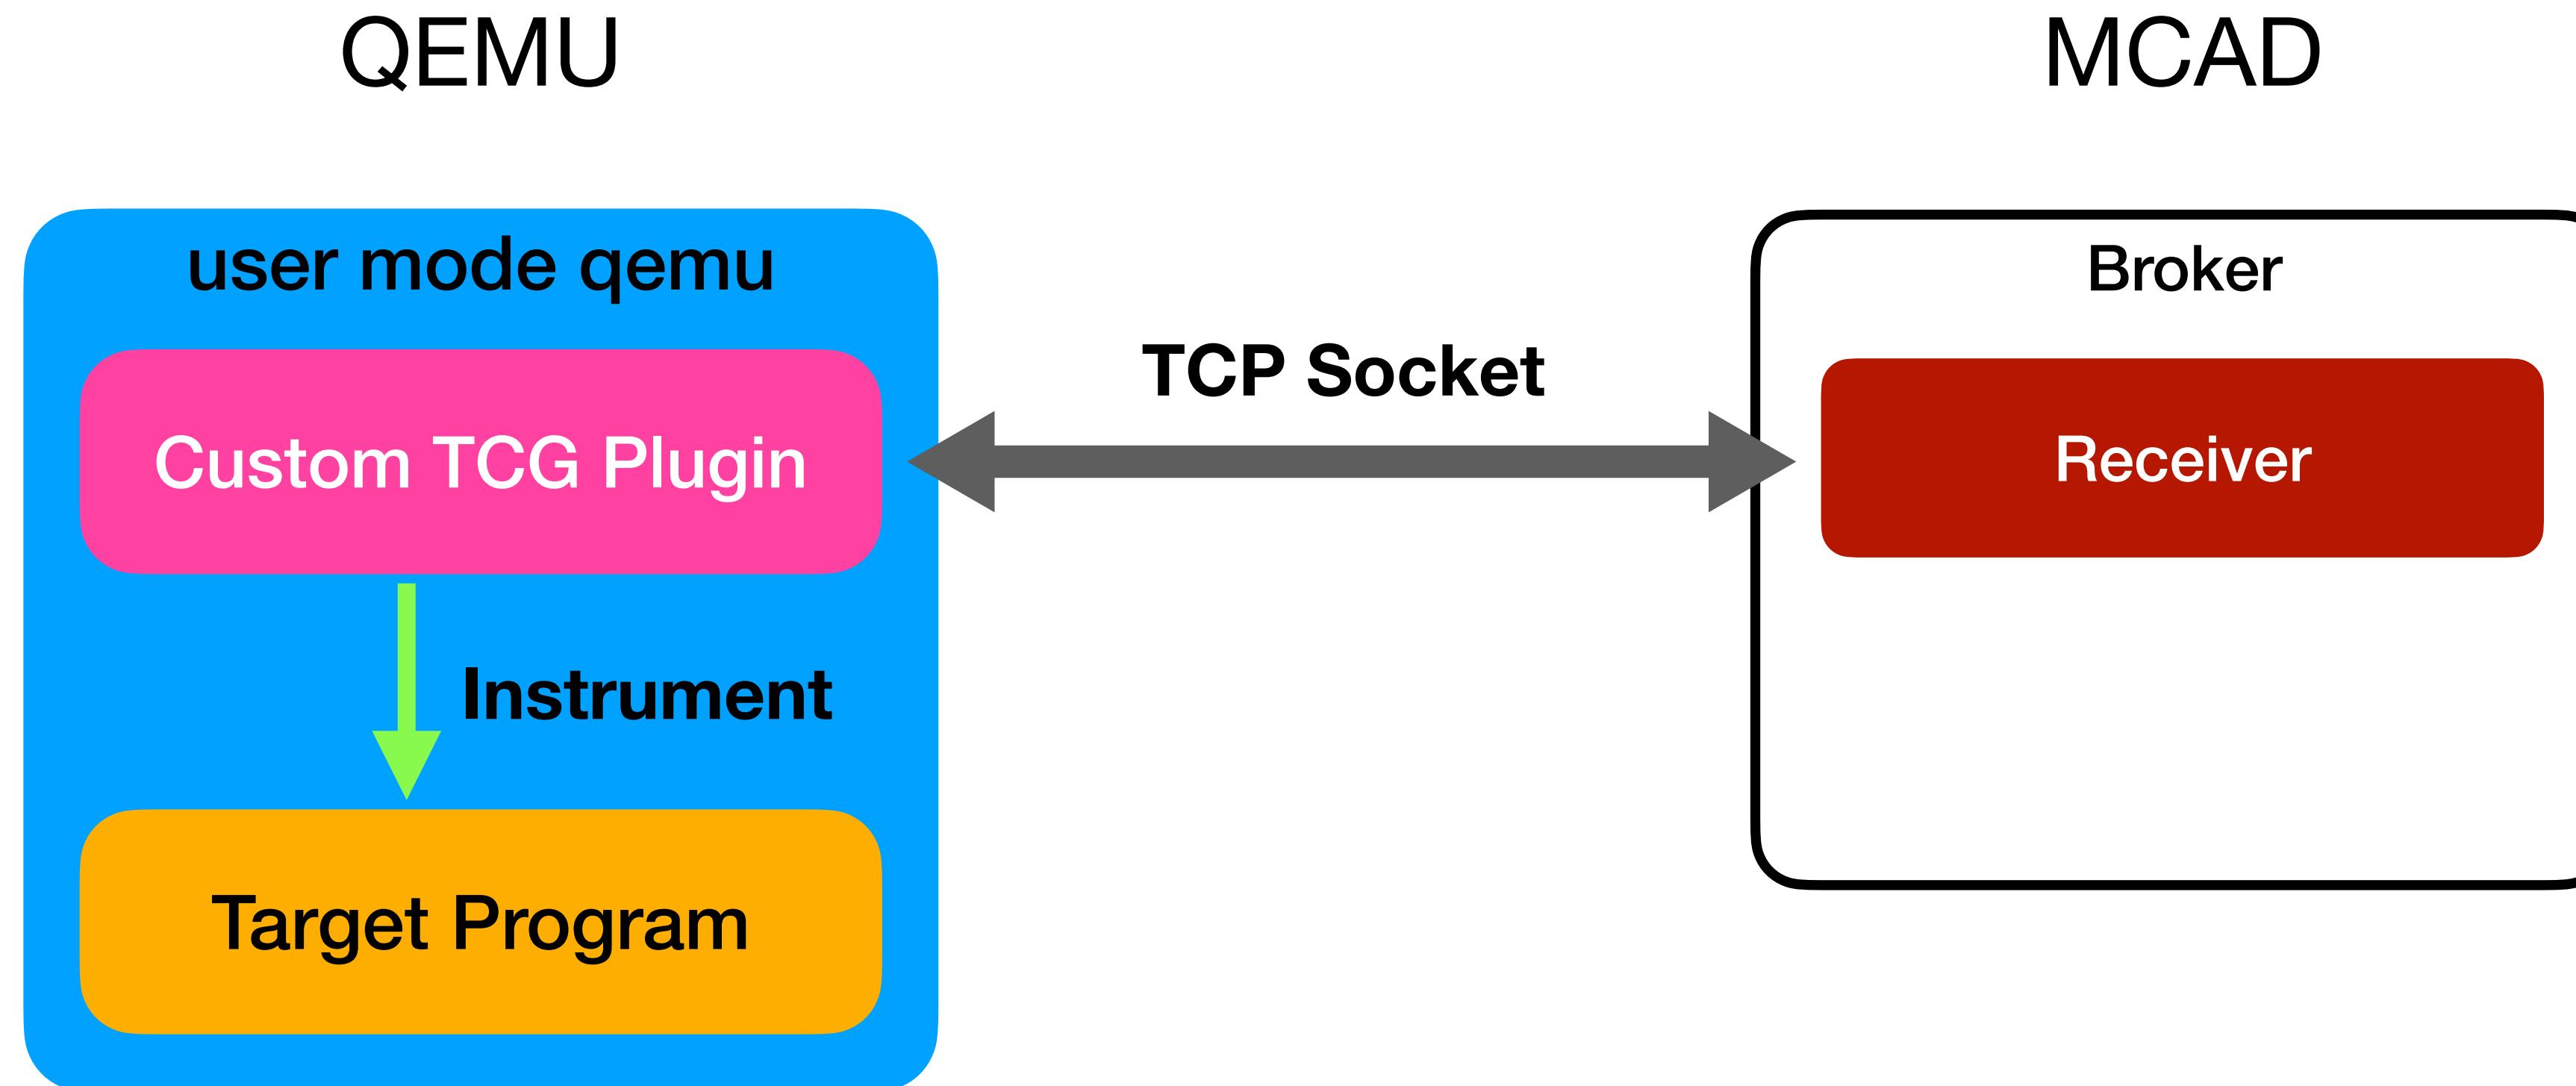
QEMU

MCAD

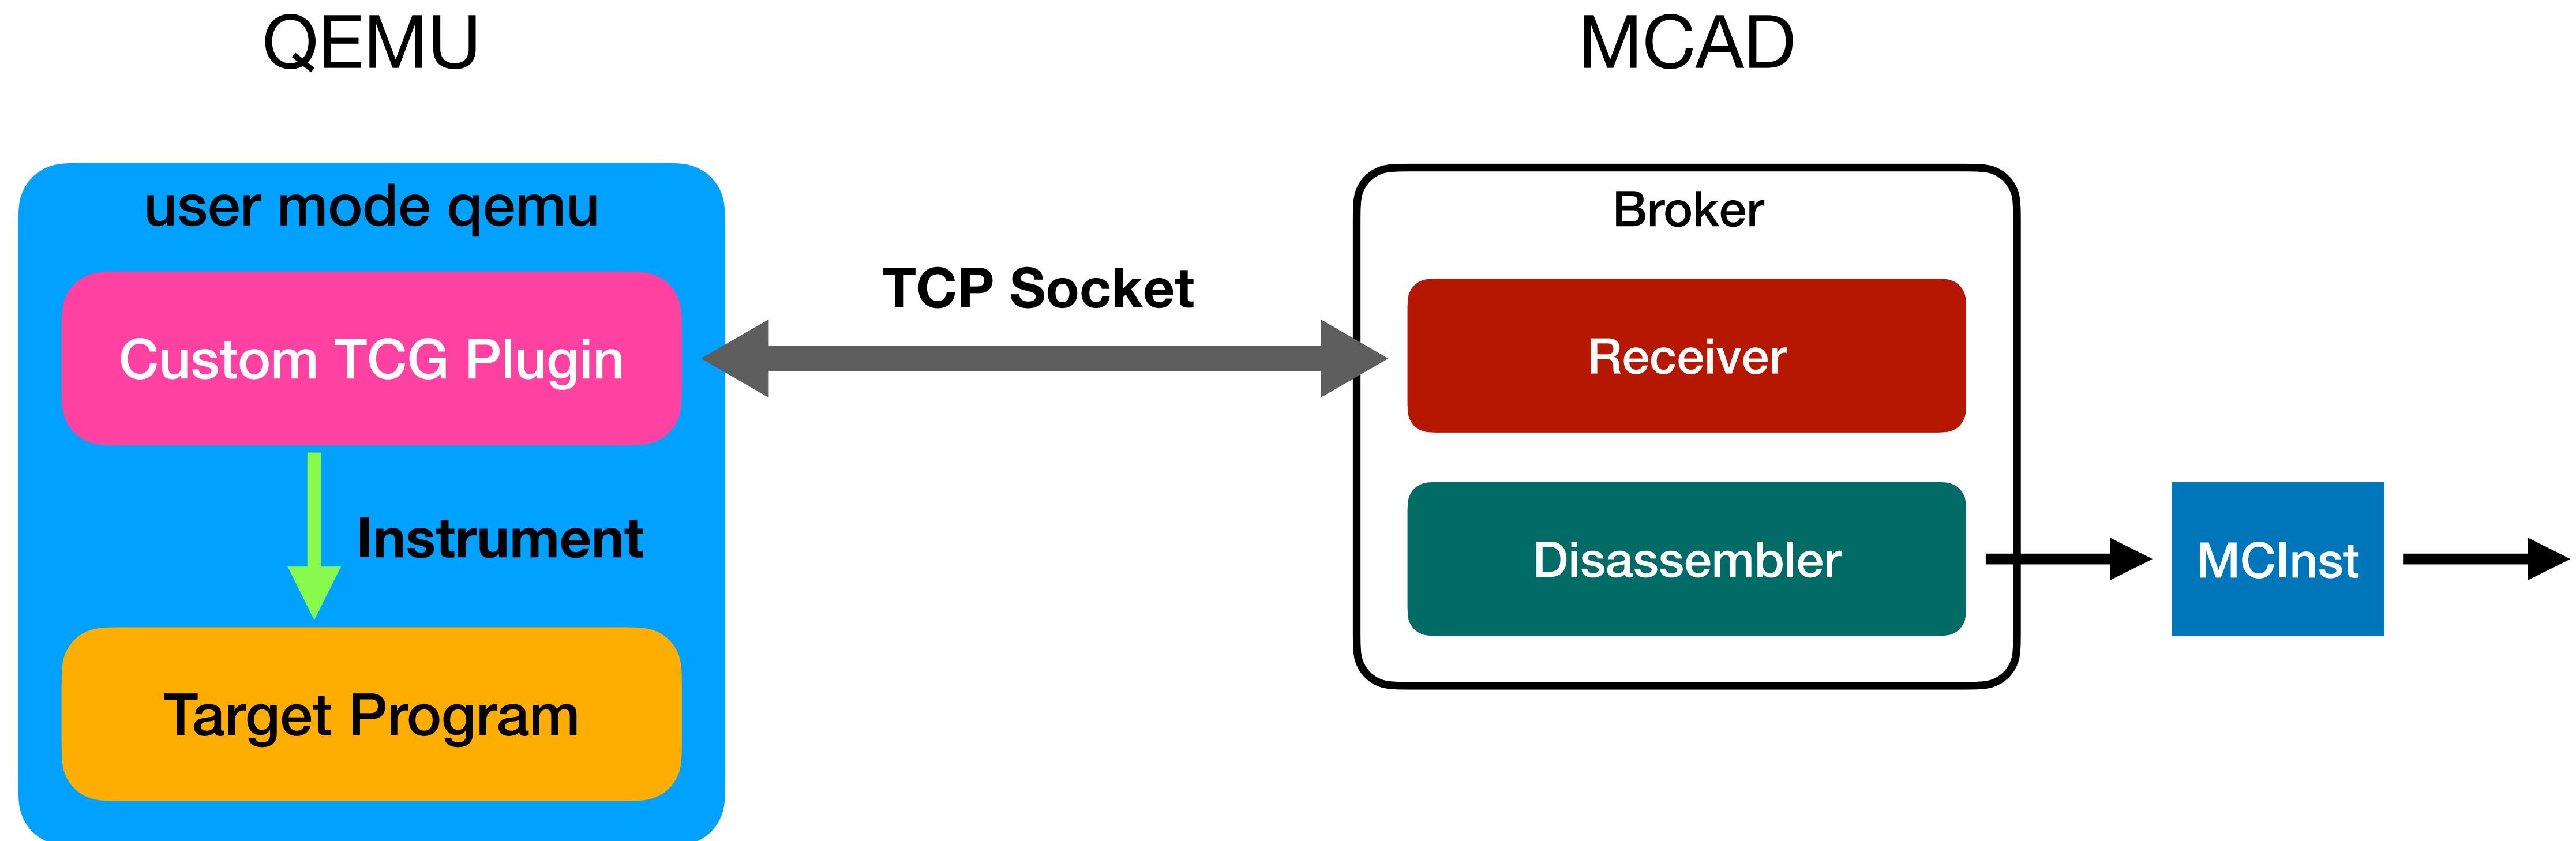
# Collecting execution traces via QEMU



# Collecting execution traces via QEMU



# Collecting execution traces via QEMU



# Custom QEMU TCG plugin

# Custom QEMU TCG plugin

- The plugin interface allows us to tap into various TCG events to collect executed instructions
  - Example: When a TCG block is translated / executed

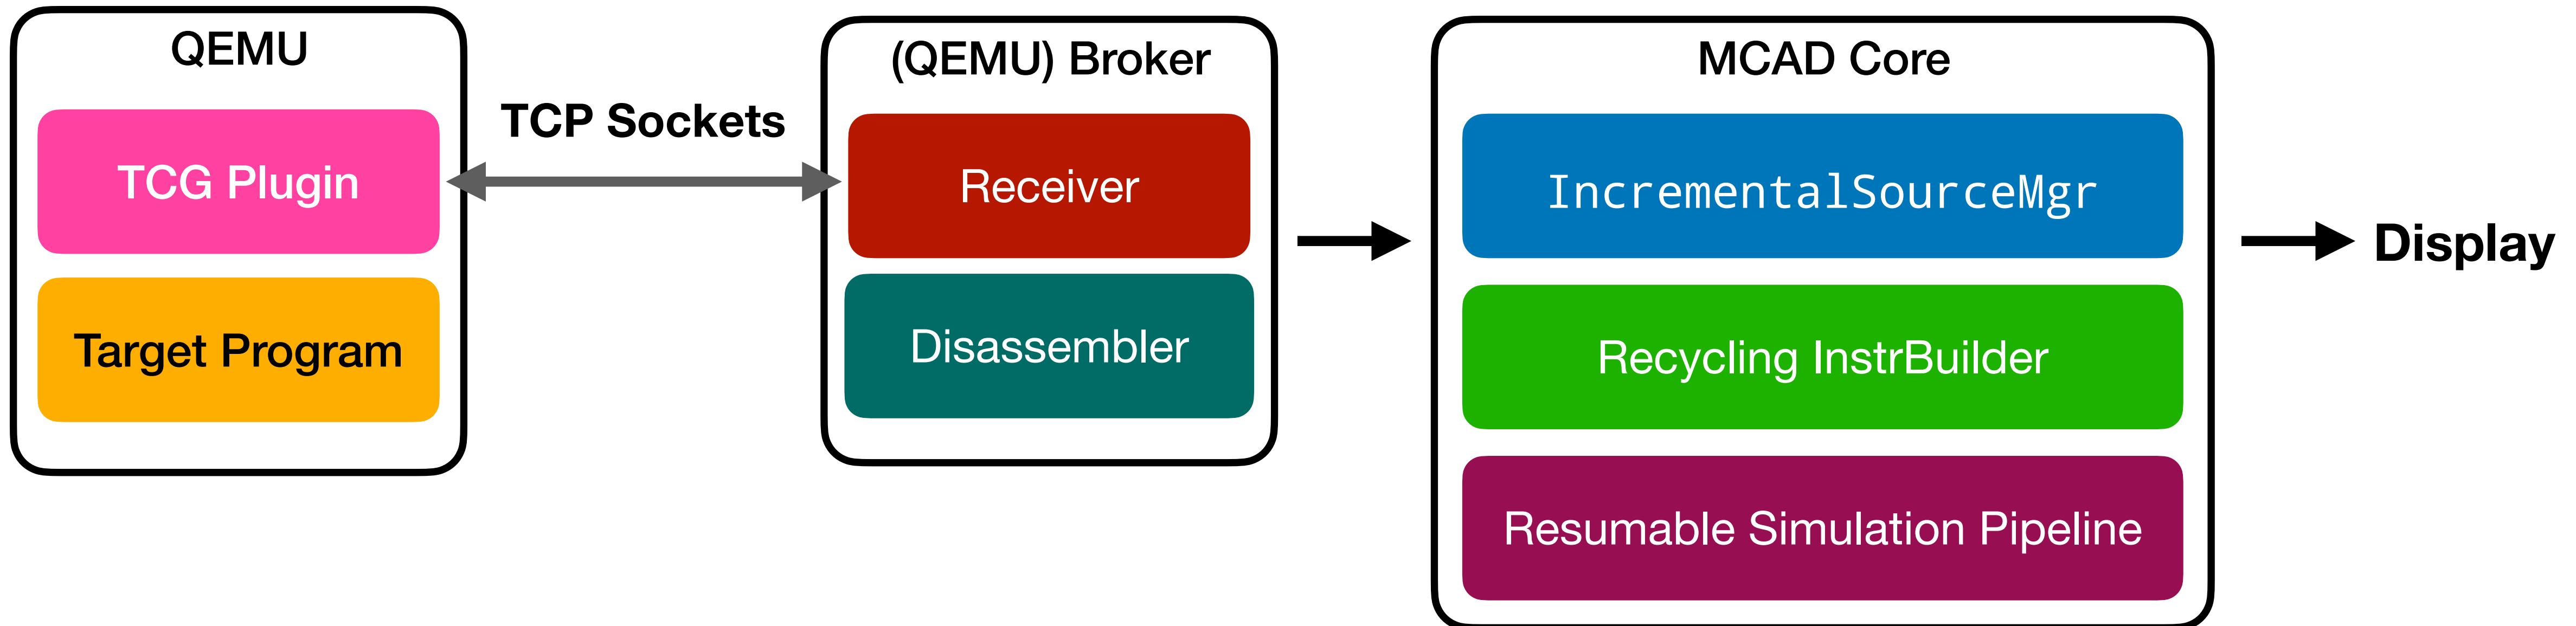
# Custom QEMU TCG plugin

- The plugin interface allows us to tap into various TCG events to collect executed instructions
  - Example: When a TCG block is translated / executed
- We also added a few plugin interfaces (not upstreamed yet)
  - Example: Retrieving CPU register values

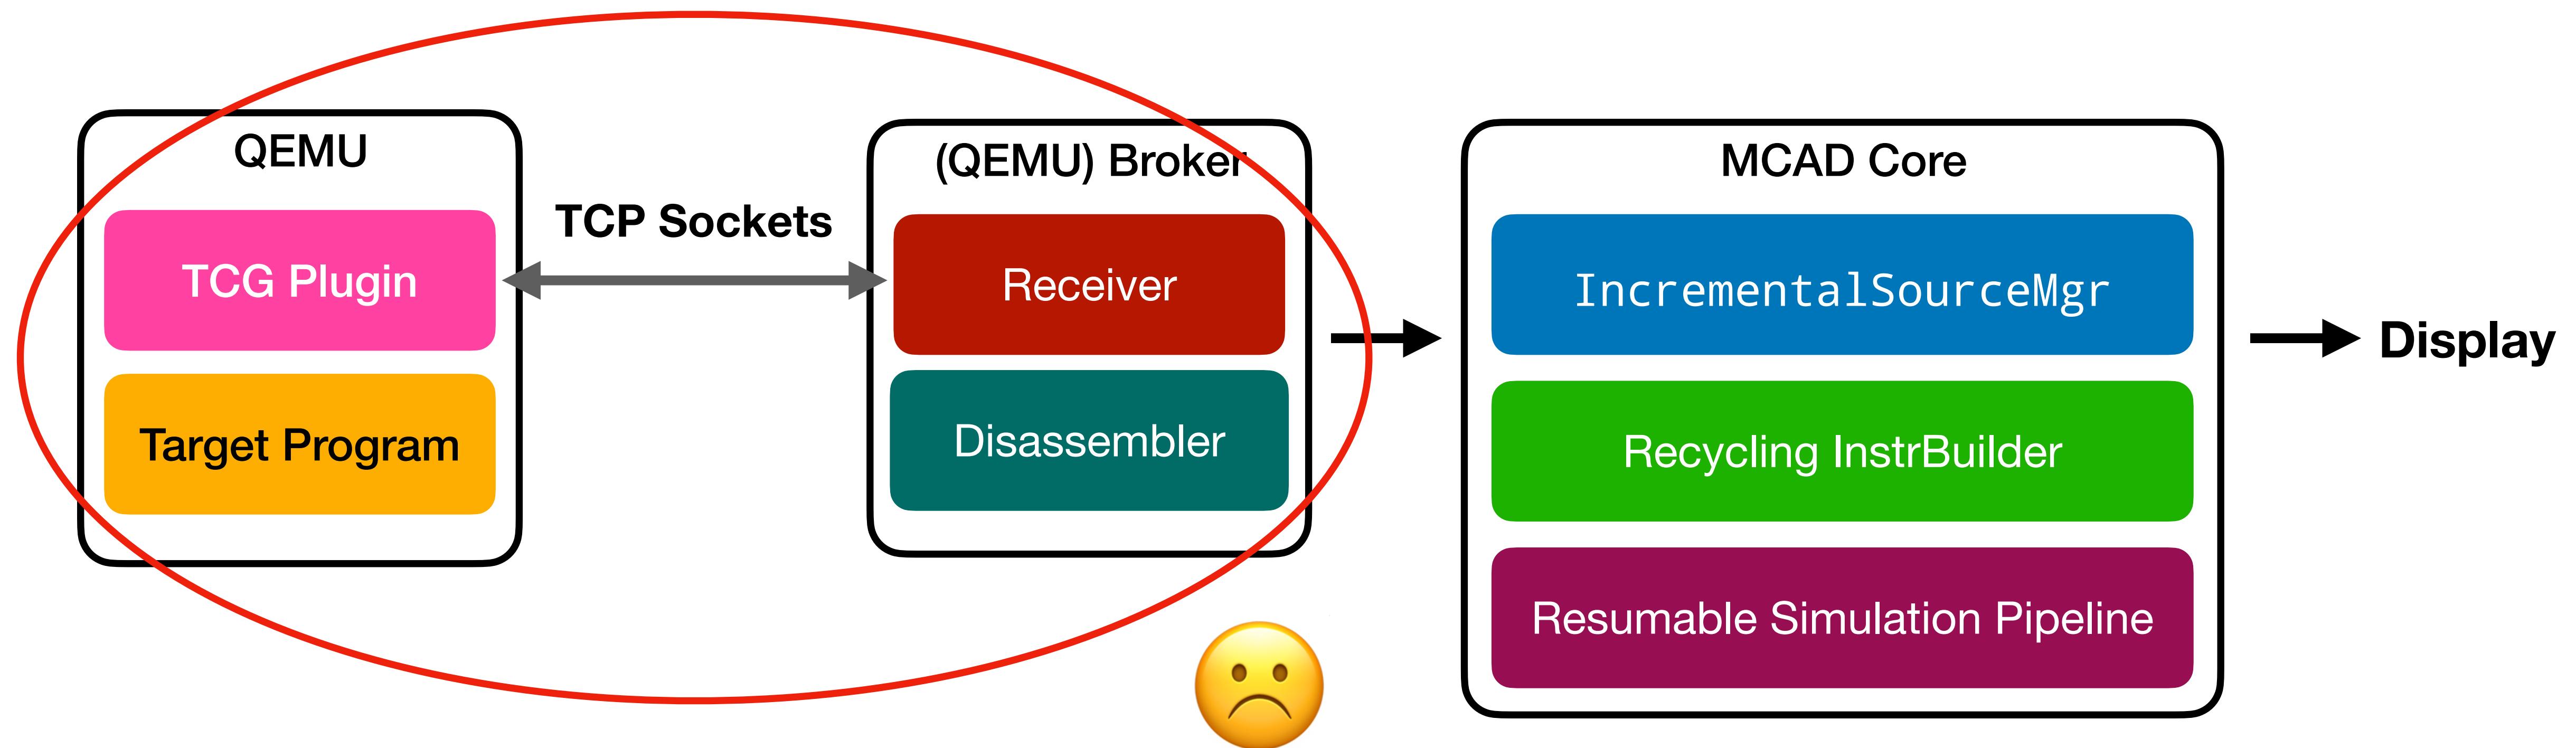
# Custom QEMU TCG plugin

- The plugin interface allows us to tap into various TCG events to collect executed instructions
  - Example: When a TCG block is translated / executed
- We also added a few plugin interfaces (not upstreamed yet)
  - Example: Retrieving CPU register values
  - Sending raw binary instructions\* through TCP sockets

# Complete structure of MCAD

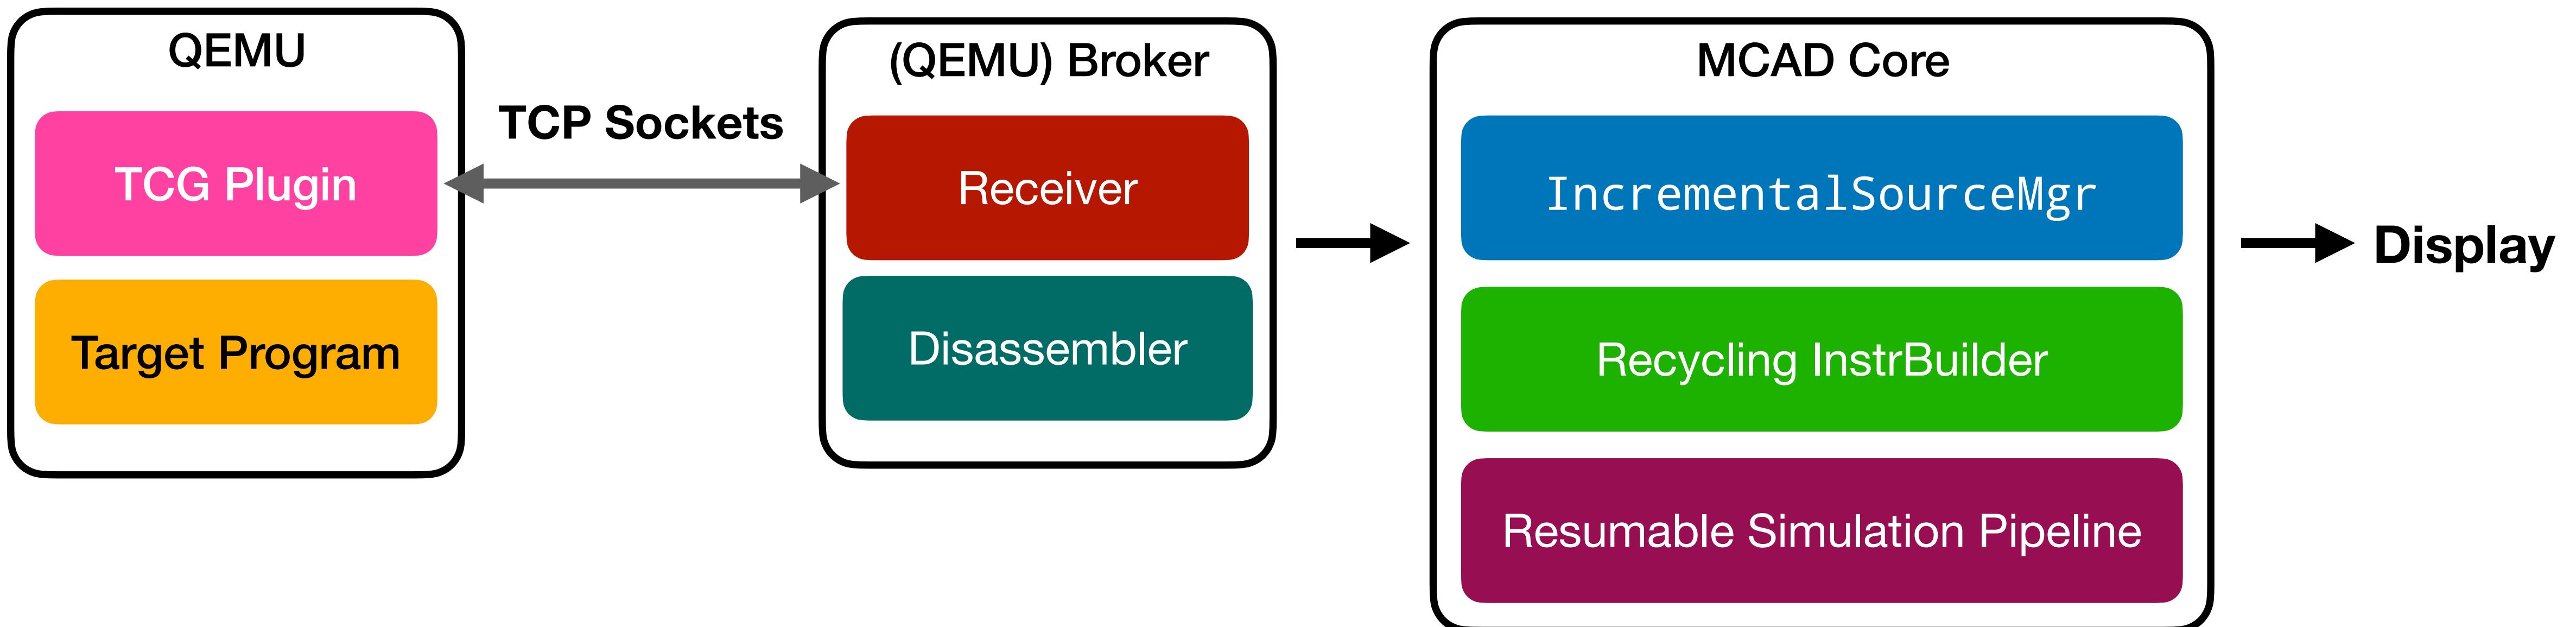


# Complete structure of MCAD



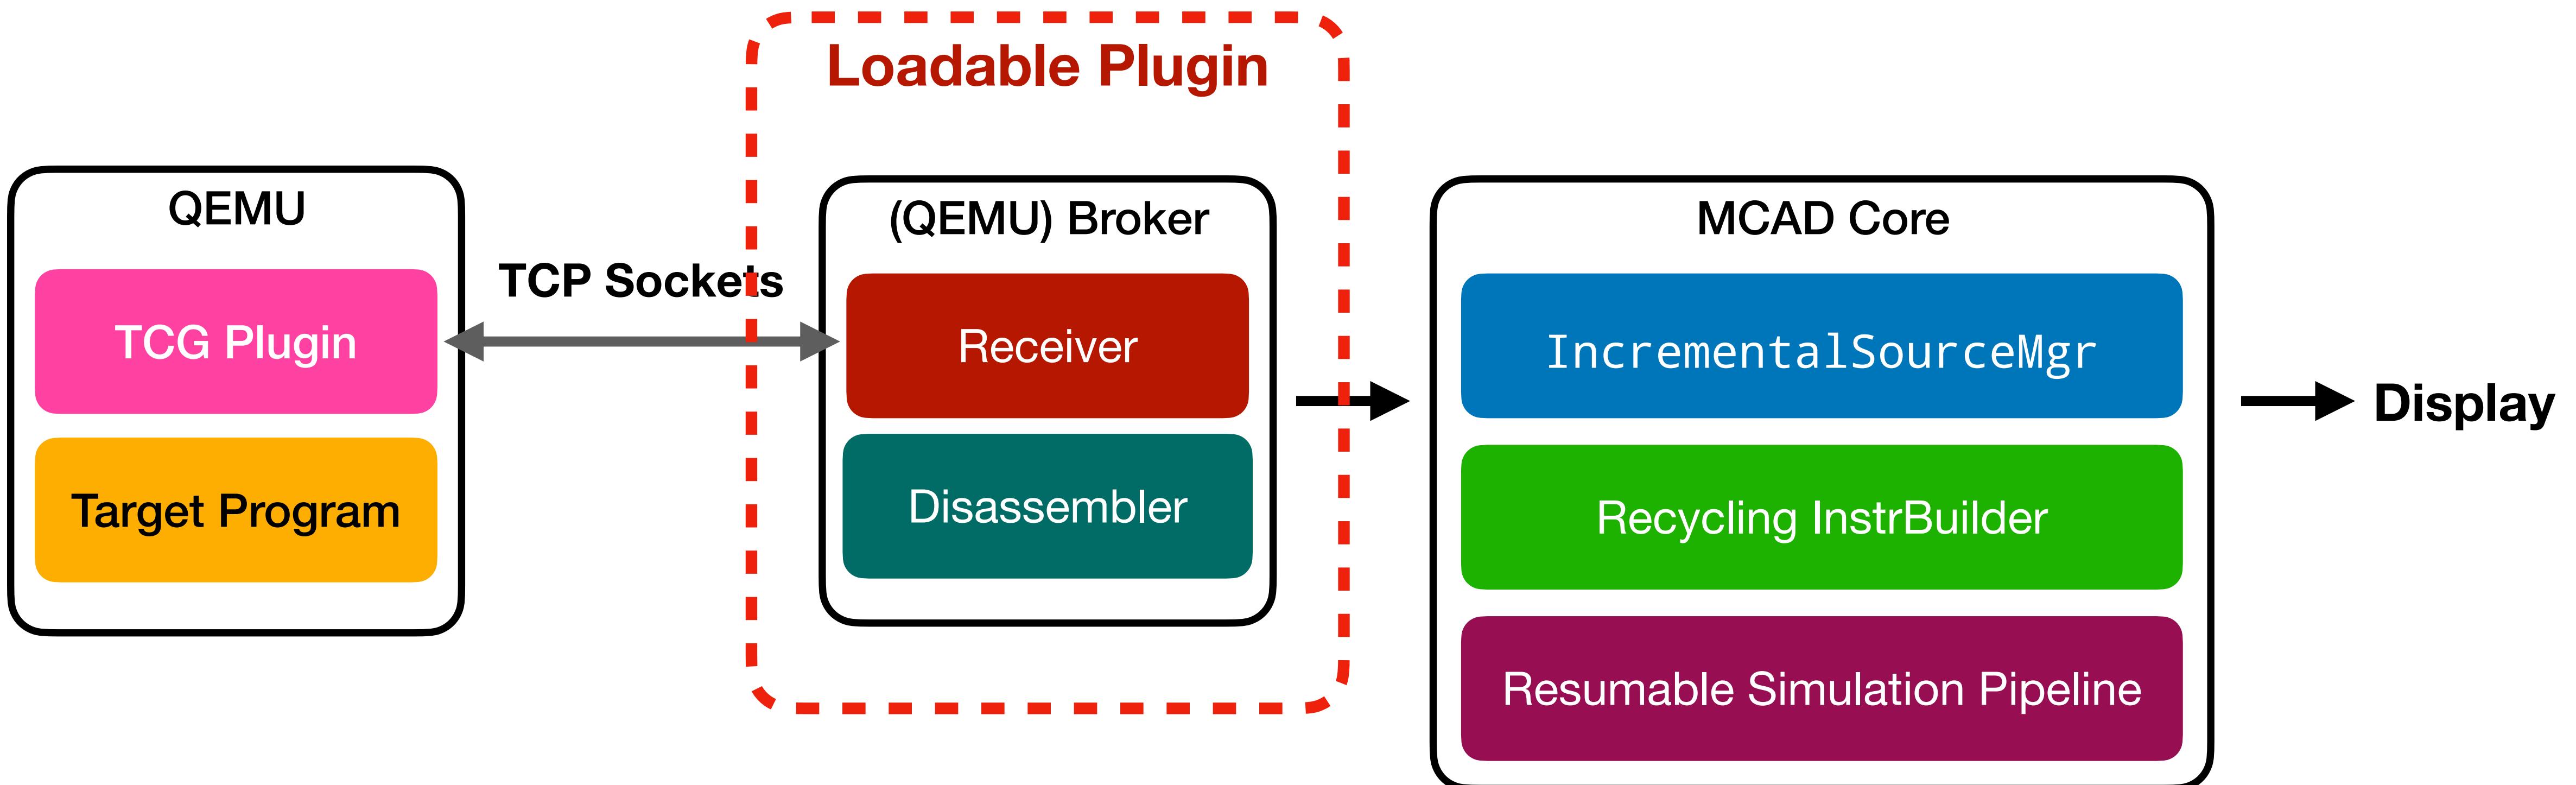
# Complete structure of MCAD

## A modular design



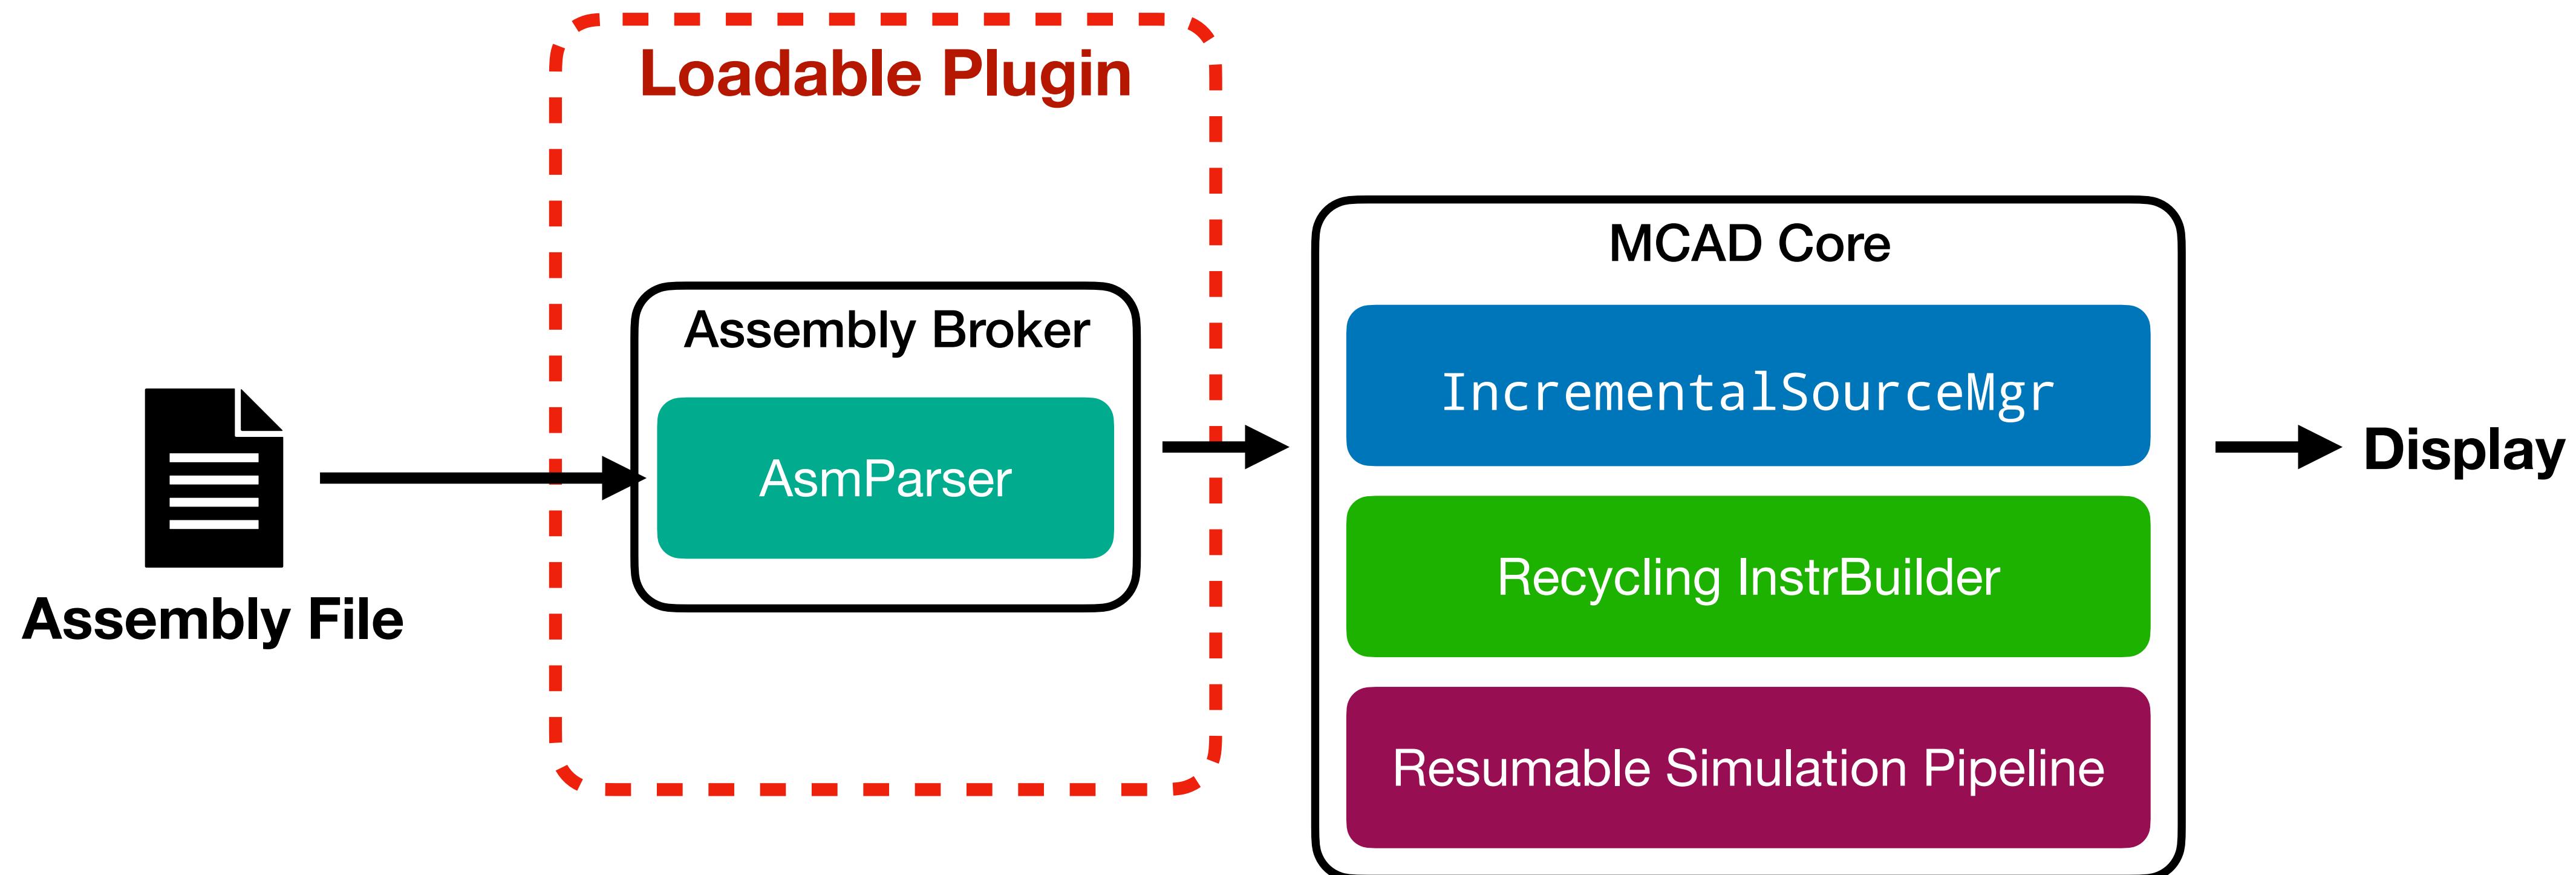
# Complete structure of MCAD

## A modular design



# Complete structure of MCAD

Example: Assembly broker plugin



# Evaluation

## Scalability compared against llvm-mca

# Evaluation

## Scalability compared against llvm-mca

- Compare MCAD against llvm-mca (i.e. baseline) on analysis speed and memory consumption

# Evaluation

## Scalability compared against llvm-mca

- Compare MCAD against llvm-mca (i.e. baseline) on analysis speed and memory consumption
- Using execution trace collected from running **x86\_64 FFmpeg 4.2** to decode a 14KB MPEG-4 video file
  - Command: `ffmpeg -i input.mp4 -f null -`
  - Size of the trace: ~27 million x86\_64 instructions

# Evaluation

## Scalability compared against llvm-mca

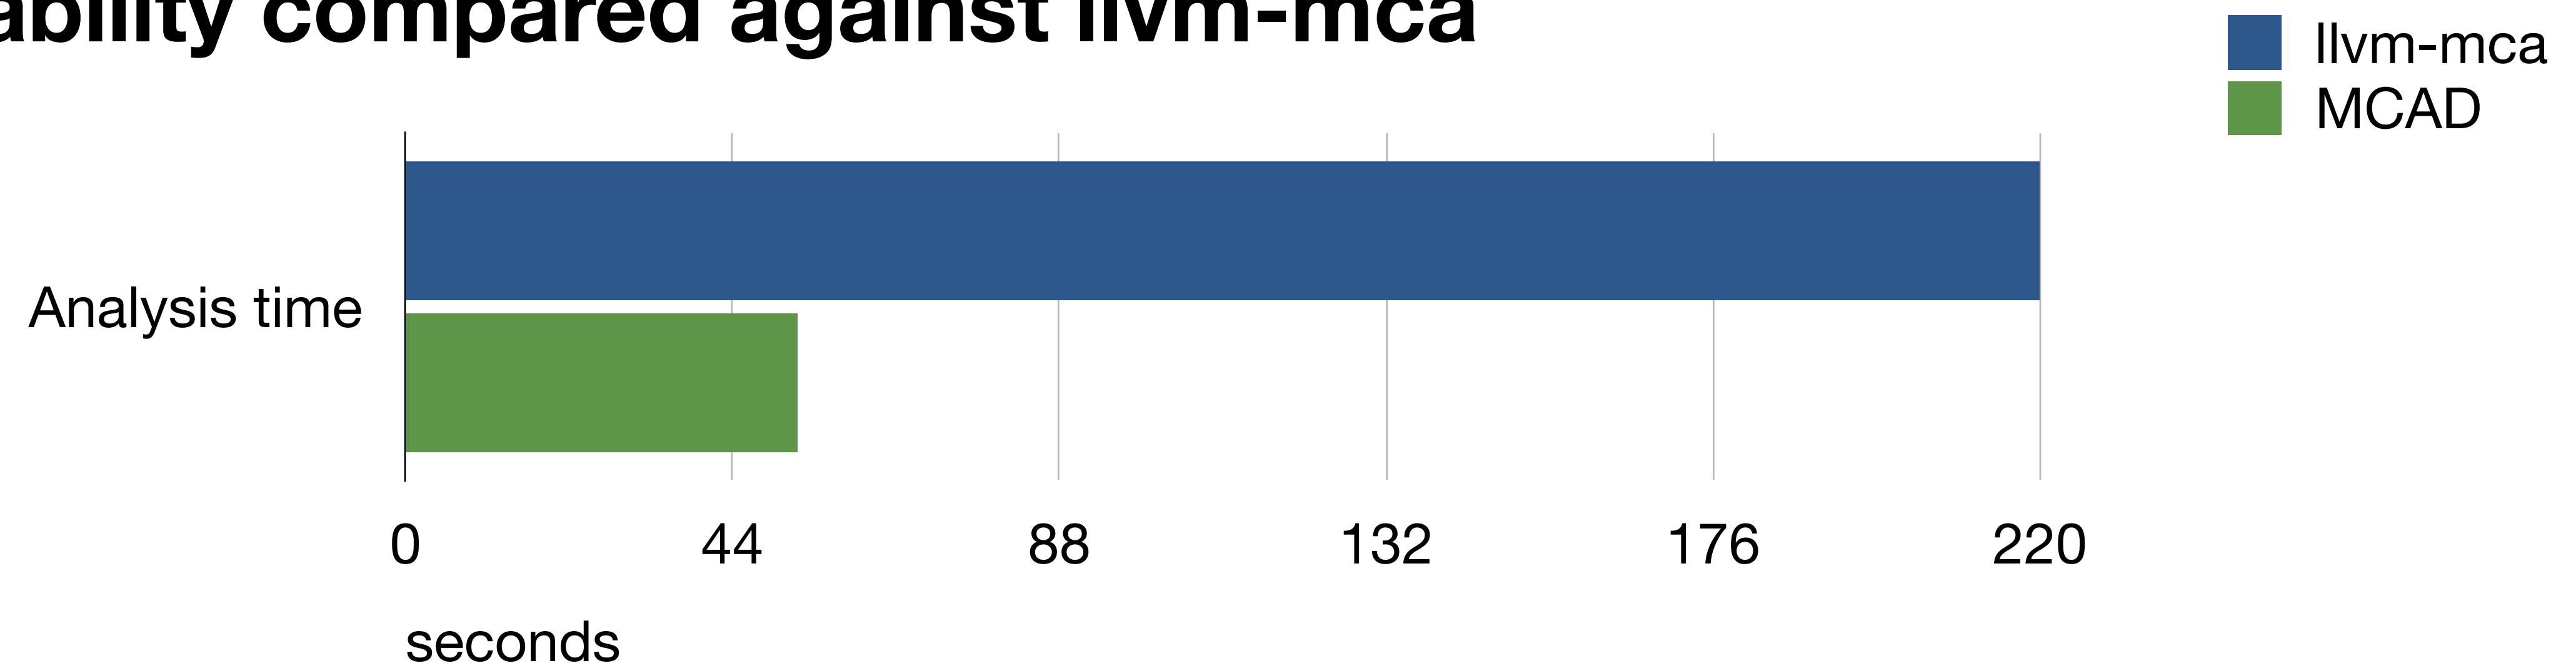
- Compare MCAD against llvm-mca (i.e. baseline) on analysis speed and memory consumption
- Using execution trace collected from running **x86\_64 FFmpeg 4.2** to decode a 14KB MPEG-4 video file
  - Command: `ffmpeg -i input.mp4 -f null -`
  - Size of the trace: ~27 million x86\_64 instructions
- For baseline, we dump the execution trace (assembly instructions) to a file before feeding into llvm-mca
  - Time measurement on baseline only accounts for llvm-mca's run time. **Excluding** the trace collection time.

# Evaluation

## Scalability compared against llvm-mca

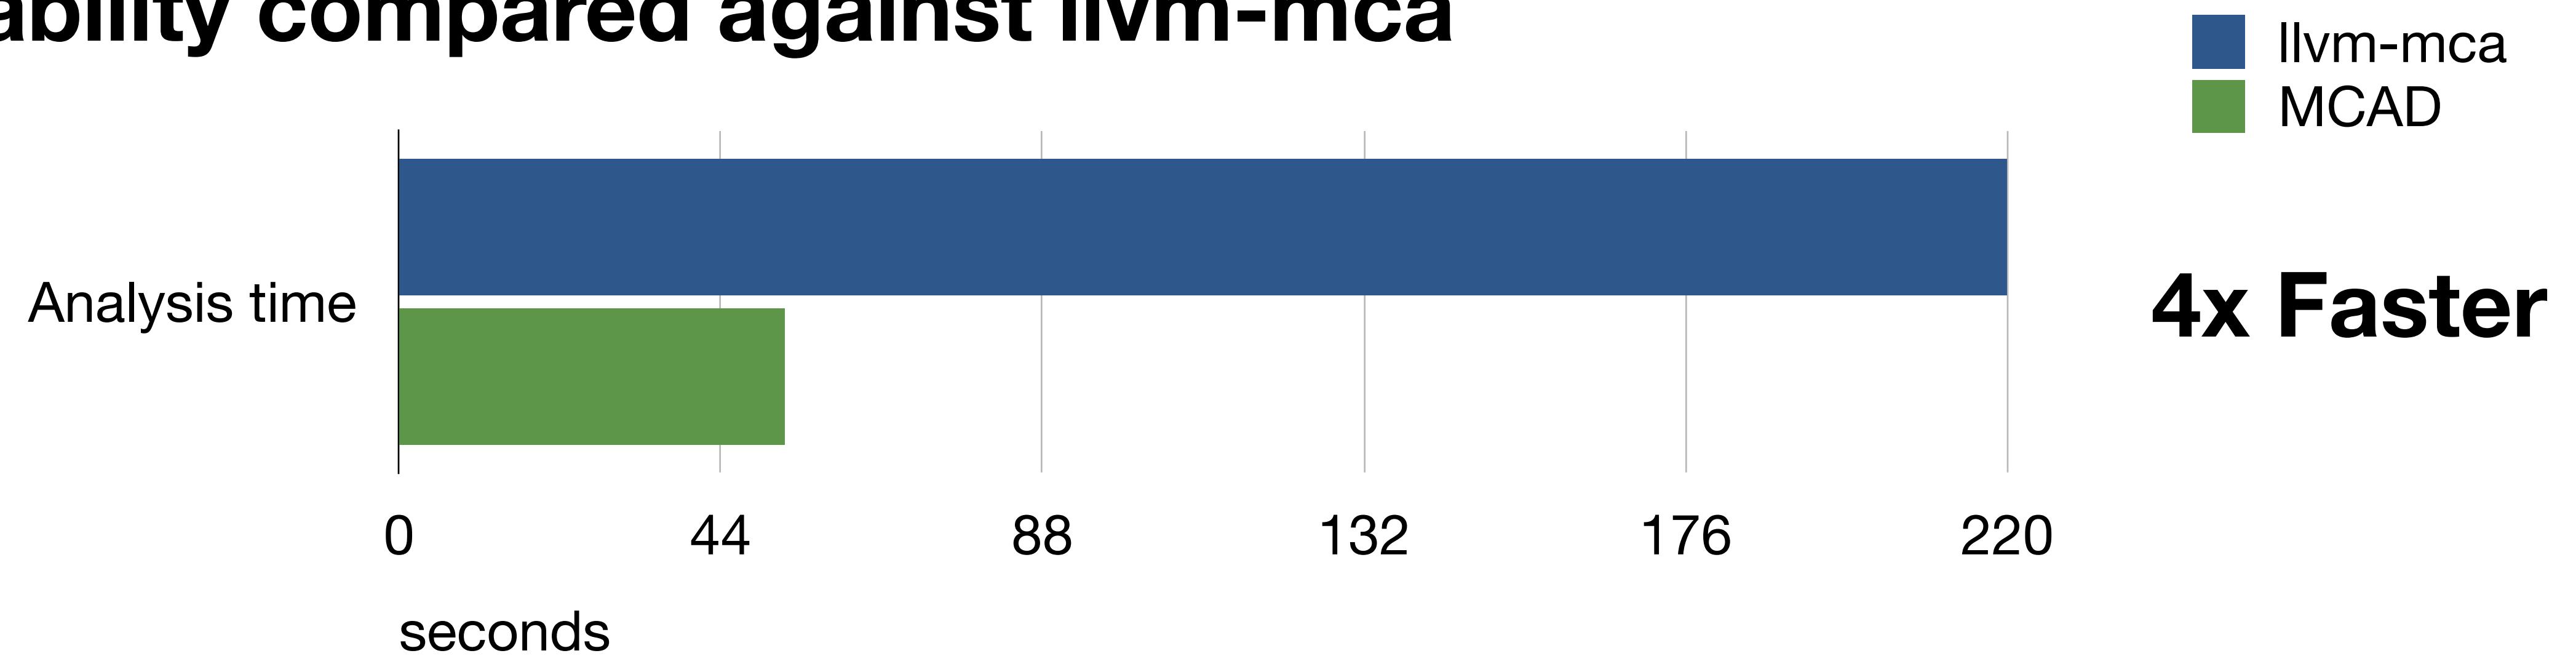
# Evaluation

## Scalability compared against llvm-mca



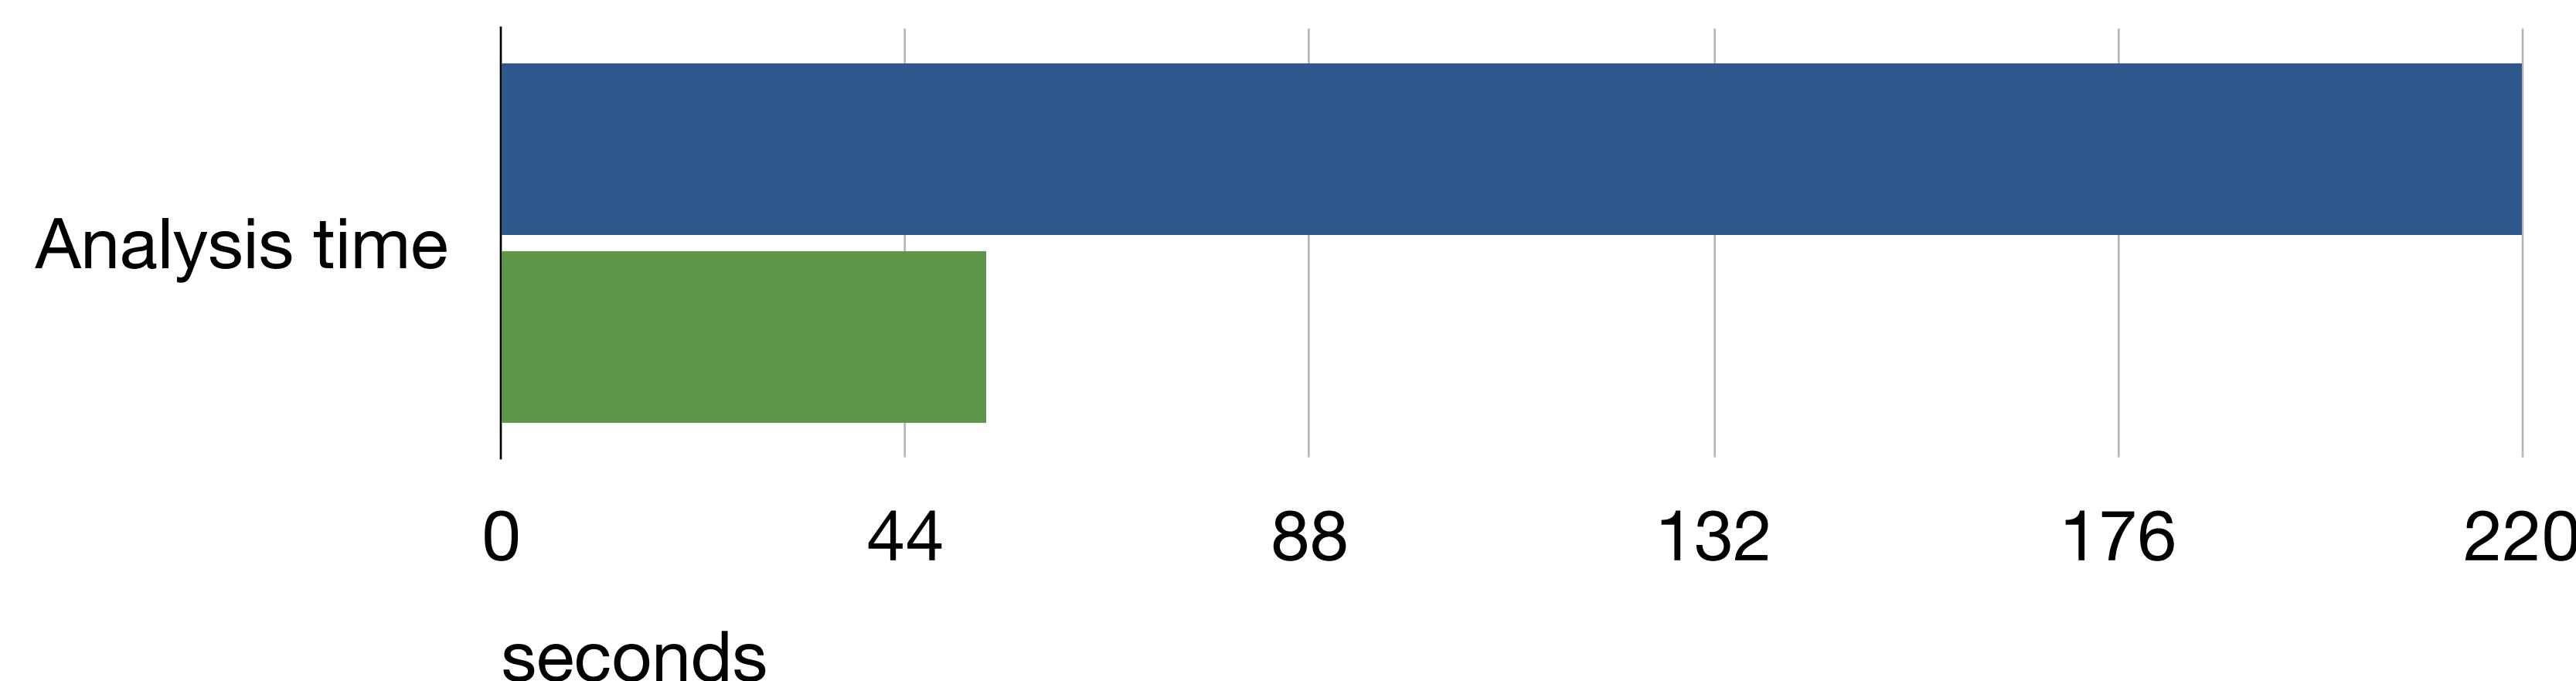
# Evaluation

## Scalability compared against llvm-mca

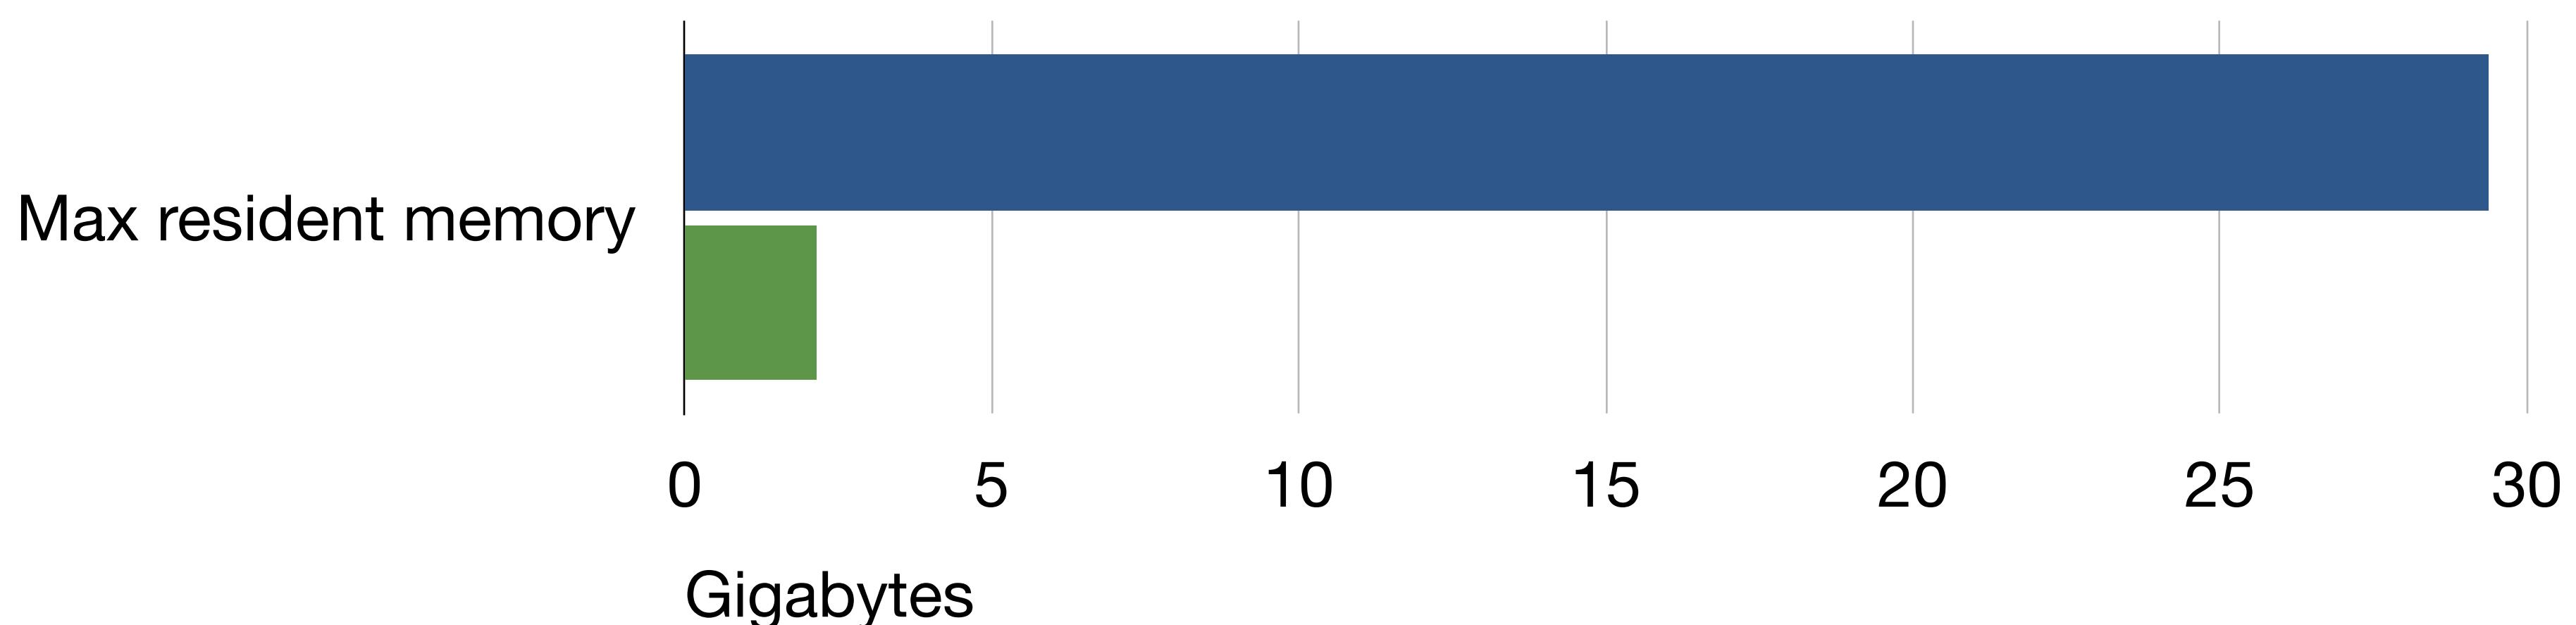


# Evaluation

## Scalability compared against llvm-mca



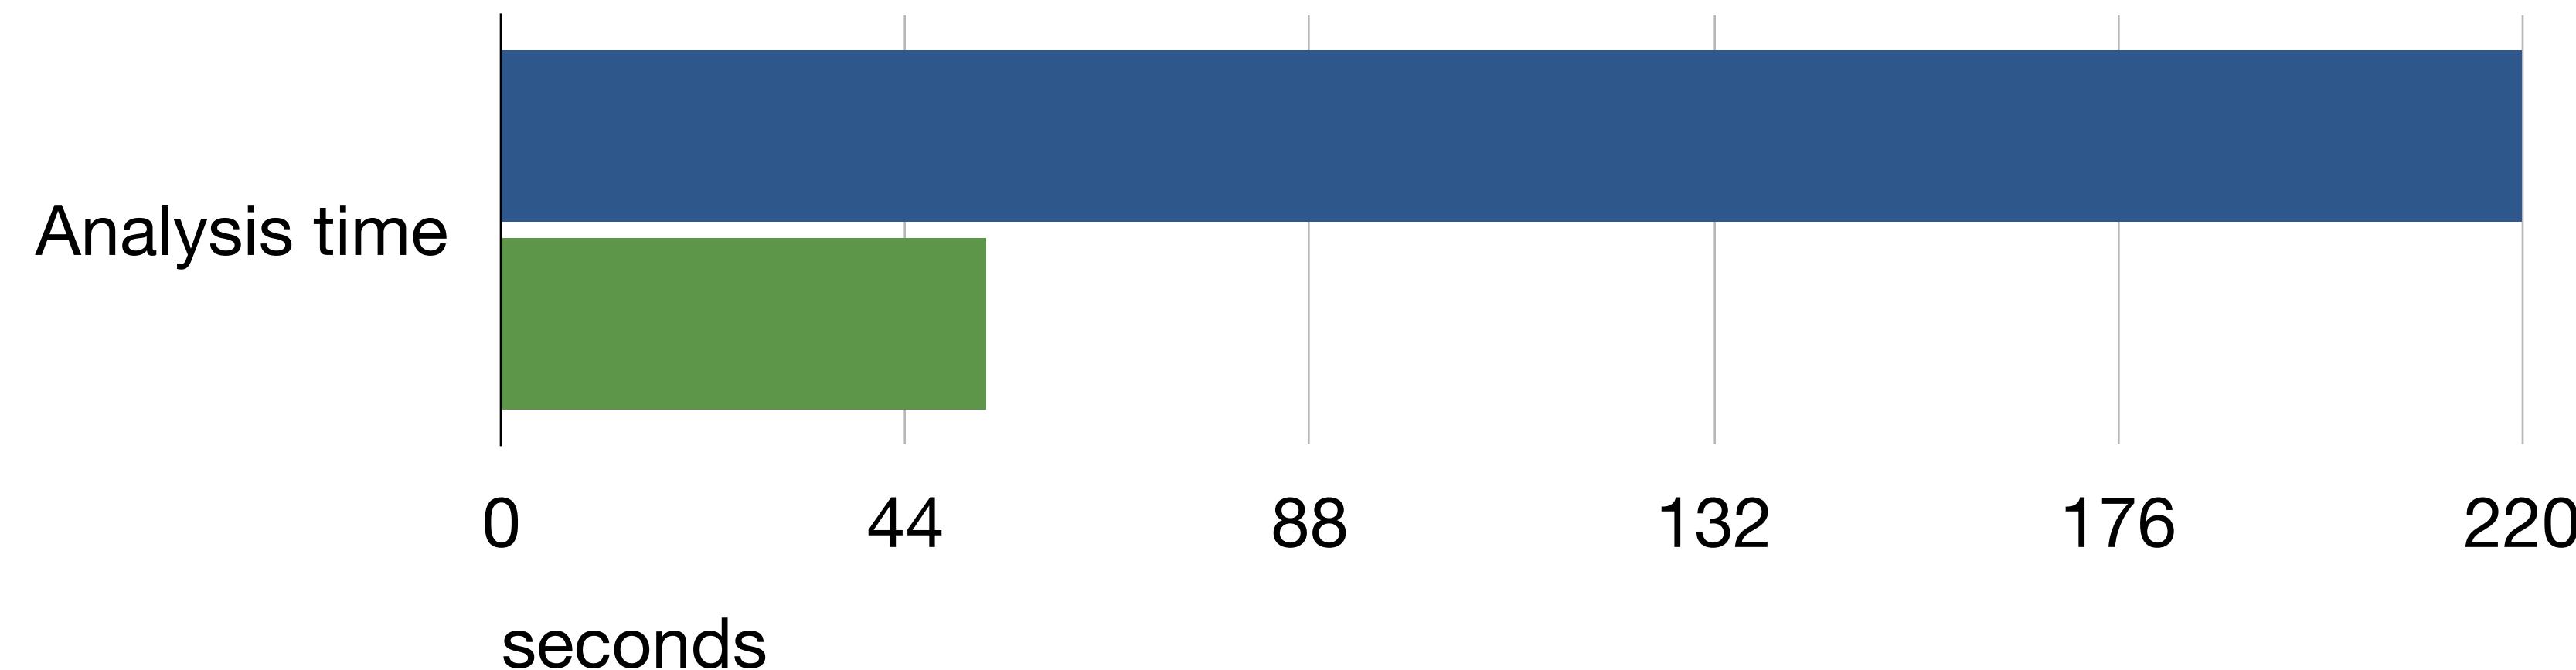
4x Faster



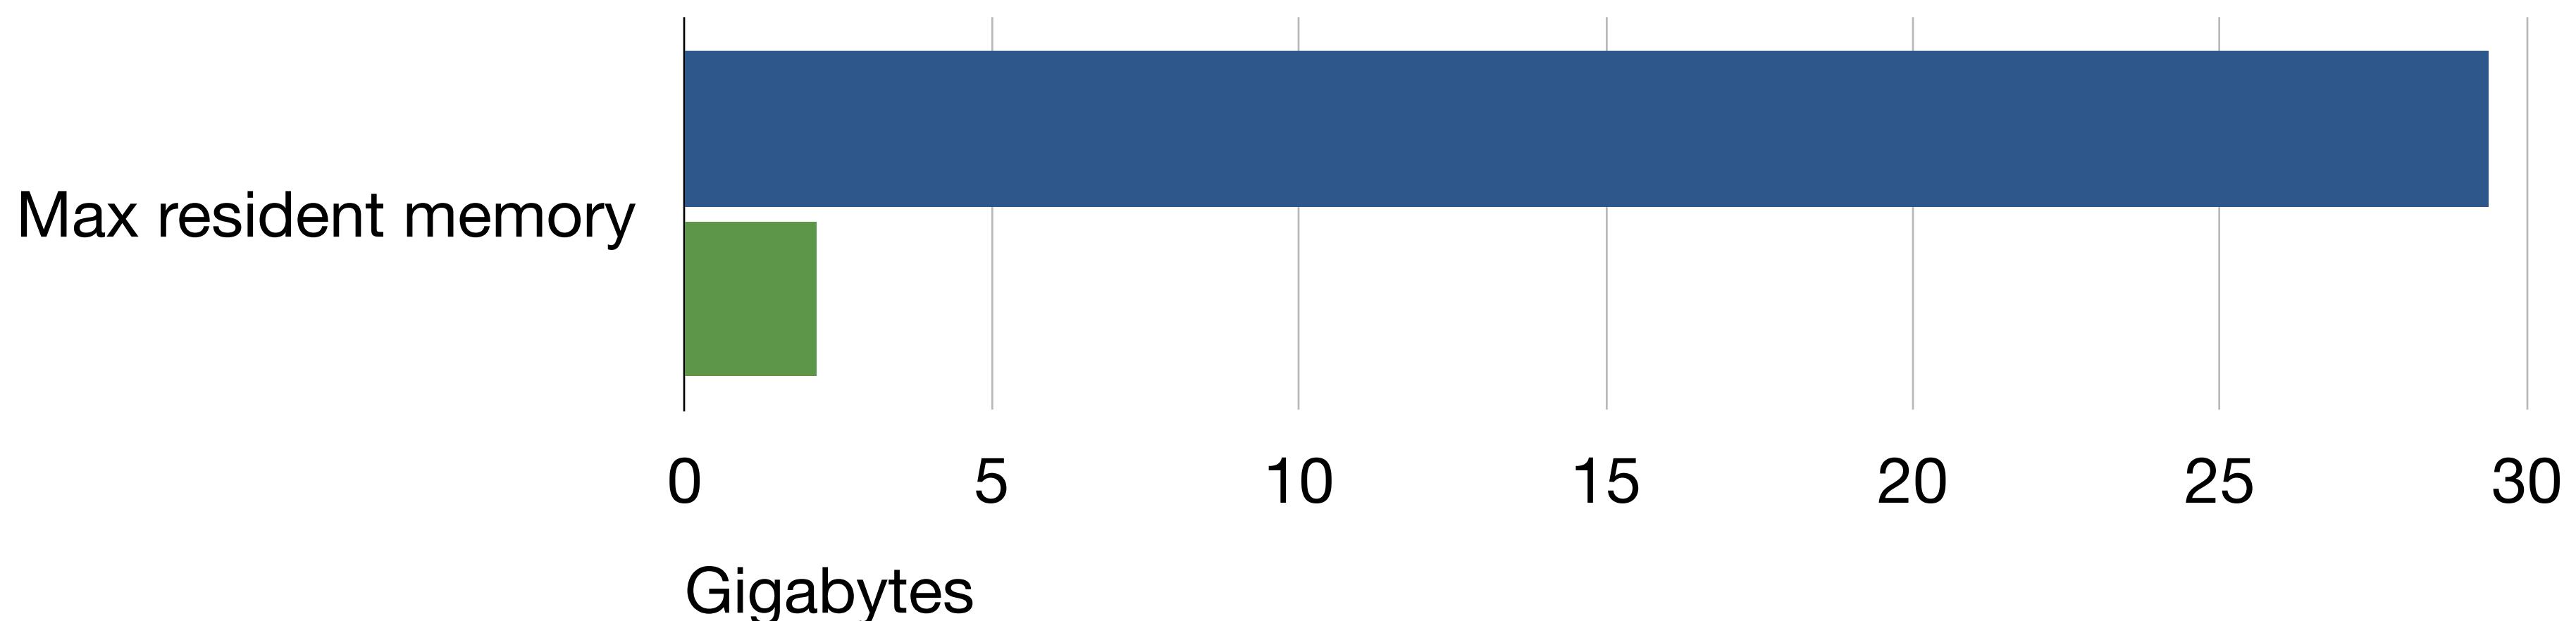
# Evaluation

## Scalability compared against llvm-mca

llvm-mca  
MCAD



**4x Faster**



**13x Less**

# Evaluation

## Scalability compared against other static throughput analysis tools

	<b>Analysis Time</b>	<b>Max Resident Memory</b>
uiCA	Timeout after 48h	113 GB
OSACA	Exit w/ error after 24h	N/A
Ithemal	Exit w/ error after 2m	N/A
<b>MCAD</b>	<b>52.69s</b>	<b>2.16 GB</b>

# Outline

Motivation

MCA Daemon (MCAD)

Future Plans

Epilogue

# // TODO

# // TODO

- More efficient ways to collect traces without QEMU

# // TODO

- More efficient ways to collect traces without QEMU
- Analyzing traces from multi-thread programs

# // TODO

- More efficient ways to collect traces without QEMU
- Analyzing traces from multi-thread programs
- Improve MCA's precision via dynamic information (e.g. memory accesses)

# // TODO

- More efficient ways to collect traces without QEMU
- Analyzing traces from multi-thread programs
- Improve MCA's precision via dynamic information (e.g. memory accesses)
- Visualizing analysis results. Or: Improve MCA's result display

# // TODO

- More efficient ways to collect traces without QEMU
- Analyzing traces from multi-thread programs
- Improve MCA's precision via dynamic information (e.g. memory accesses)
- Visualizing analysis results. Or: Improve MCA's result display
  - Example: *Loadable plugins* for custom display of the result

# // TODO

- More efficient ways to collect traces without QEMU
- Analyzing traces from multi-thread programs
- Improve MCA's precision via dynamic information (e.g. memory accesses)
- Visualizing analysis results. Or: Improve MCA's result display
  - Example: *Loadable plugins* for custom display of the result
- Going upstream: QEMU & LLVM

# Going upstream: LLVM

## The plan

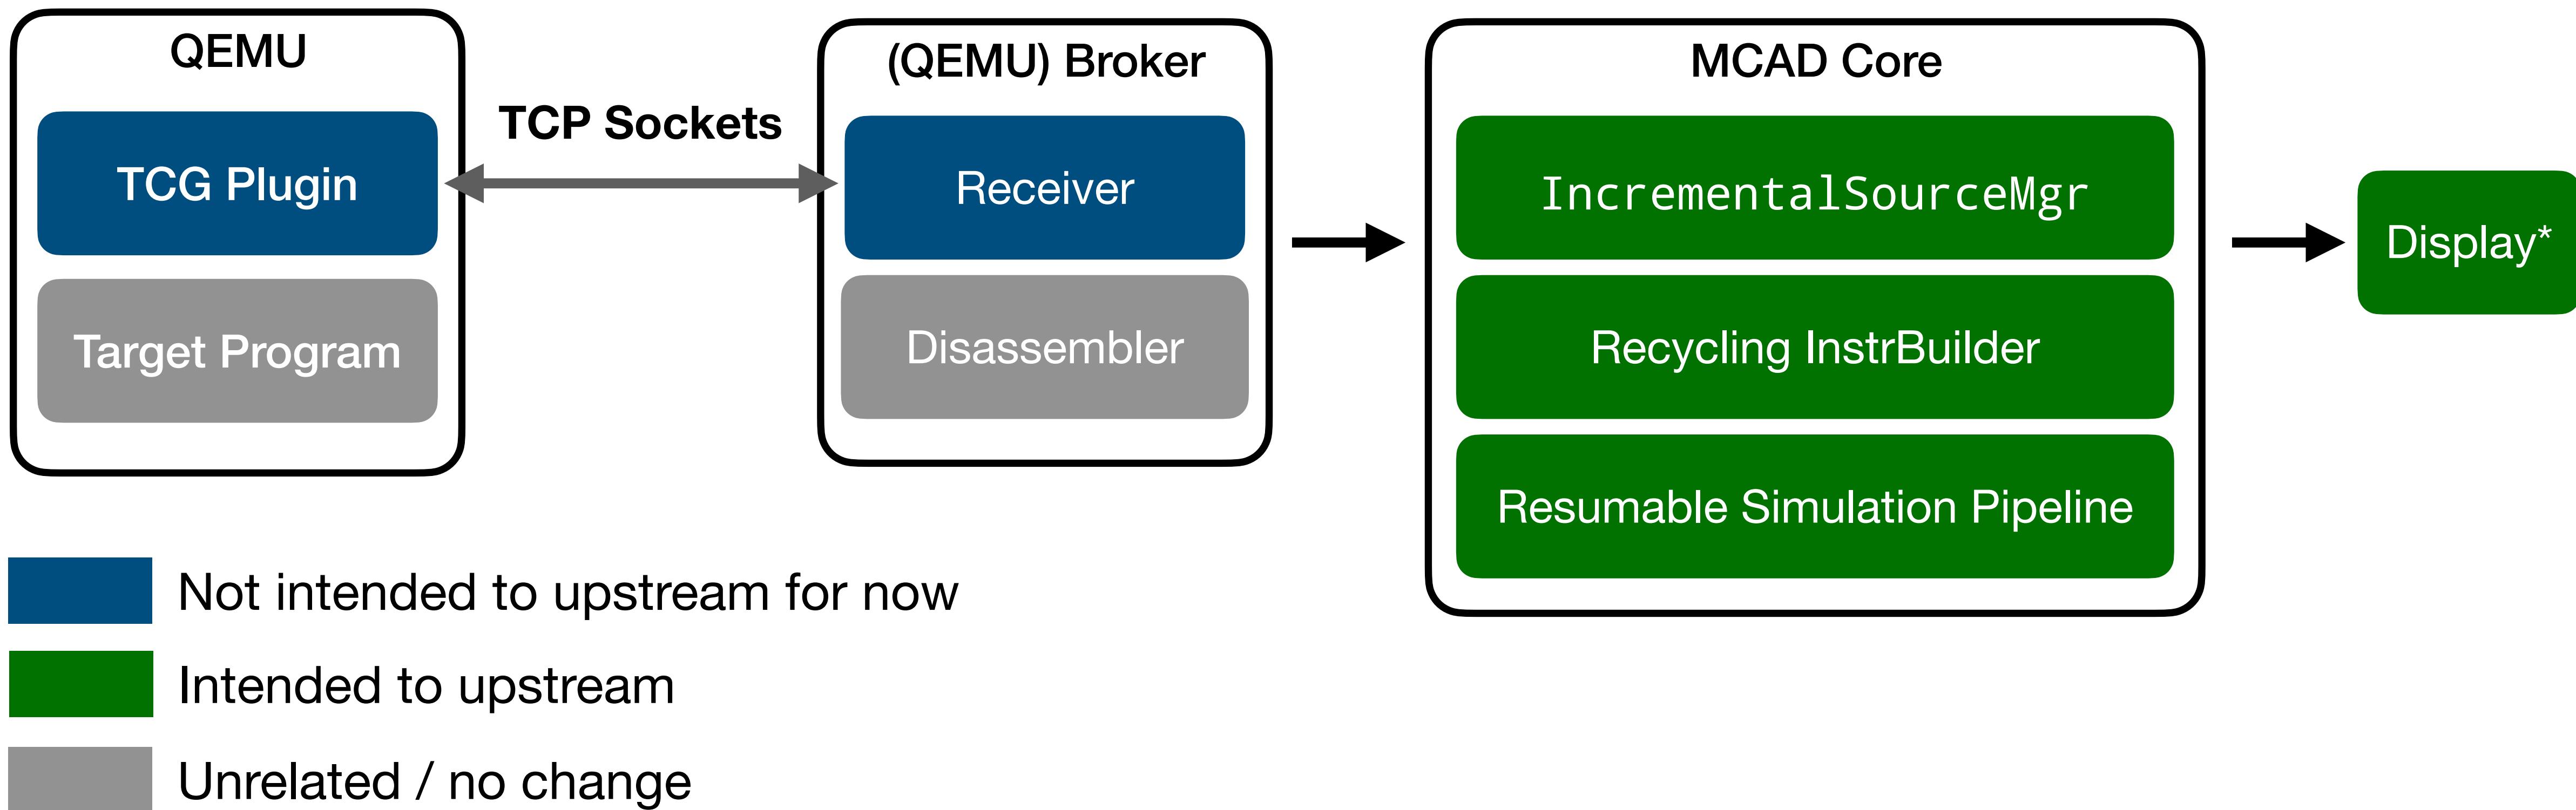
# Going upstream: LLVM

## The plan

- We would like to upstream components that are beneficial to the **core MCA libraries** first

# Going upstream: LLVM

## The plan: components to upstream



# Going upstream: LLVM

## The plan

- We would like to upstream components that are beneficial to the **core MCA libraries** first

# Going upstream: LLVM

## The plan

- We would like to upstream components that are beneficial to the **core MCA libraries** first
- QEMU broker plugin & our TCG plugin will be maintained out-of-tree

# Going upstream: LLVM

## The plan

- We would like to upstream components that are beneficial to the **core MCA libraries** first
- QEMU broker plugin & our TCG plugin will be maintained out-of-tree
- We're not sure about upstreaming rest of the tool right now

# Going upstream: LLVM

## The plan

- We would like to upstream components that are beneficial to the **core MCA libraries** first
- QEMU broker plugin & our TCG plugin will be maintained out-of-tree
- We're not sure about upstreaming rest of the tool right now
  - With the assembly broker, MCAD can be a drop-in replacement for llvm-mca...with even more features (e.g. the broker plugin infrastructure)

# Going upstream: LLVM

## The plan

- We would like to upstream components that are beneficial to the **core MCA libraries** first
- QEMU broker plugin & our TCG plugin will be maintained out-of-tree
- We're not sure about upstreaming rest of the tool right now
  - With the assembly broker, MCAD can be a drop-in replacement for llvm-mca...with even more features (e.g. the broker plugin infrastructure)
  - Some of the (advanced) interfaces in broker plugin are only used by QEMU broker. So, without the latter, it's not well tested.

# Outline

Motivation

MCA Daemon (MCAD)

Future Plans

Epilogue

# Summary

# Summary

- **MCA Daemon (MCAD)** is a high-performance *hybrid* throughput analysis tool built on top of LLVM MCA libraries

# Summary

- **MCA Daemon (MCAD)** is a high-performance *hybrid* throughput analysis tool built on top of LLVM MCA libraries
  - Online, whole-program analysis on real-world applications

# Summary

- **MCA Daemon (MCAD)** is a high-performance *hybrid* throughput analysis tool built on top of LLVM MCA libraries
  - Online, whole-program analysis on real-world applications
  - Scale up with large-scale programs with tens of **millions** of instructions

# Summary

- **MCA Daemon (MCAD)** is a high-performance *hybrid* throughput analysis tool built on top of LLVM MCA libraries
  - Online, whole-program analysis on real-world applications
  - Scale up with large-scale programs with tens of **millions** of instructions
  - We improved the performance & flexibility of core MCA libraries

# Summary

- **MCA Daemon (MCAD)** is a high-performance *hybrid* throughput analysis tool built on top of LLVM MCA libraries
  - Online, whole-program analysis on real-world applications
  - Scale up with large-scale programs with tens of **millions** of instructions
- We improved the performance & flexibility of core MCA libraries
  - We would like to merge these changes upstream to benefit the wider community

# Source code

**<https://github.com/securesystemslab/LLVM-MCA-Daemon>**



# Acknowledgements

This material is based upon work partially supported by the Defense Advanced Research Projects Agency (DARPA) under contract N66001-20-C-4027

## Special Thanks:

**Galois Inc.** (<https://galois.com>)

**Immigrant Inc.** (<https://immigrant.com>)

A wide-angle photograph of a park at sunset. In the foreground, several tall eucalyptus trees stand on a grassy hillside. A vibrant purple flowering tree is visible among them. The background features a large, dense forest of various trees under a blue sky with scattered white clouds. Sunlight filters through the trees, creating bright highlights and long shadows on the grass.

# Thank You!

## Q&A

# Appendix

# Introduction to MCA

## An example

test/tools/llvm-mca/X86/BtVer2/dot-products.s

```
vmulps    %xmm0,  %xmm1,  %xmm2  
vhaddps   %xmm2,  %xmm2,  %xmm3  
vhaddps   %xmm3,  %xmm3,  %xmm4  
  
 llvm-mca -mtriple=x86_64 -mcpu=btver2 \  
           -iterations=300 dot-products.s
```

### Timeline

Index	0123456789	0123456789	012345
[0,0]	DeeER.	.	.
[0,1]	D==eeeER	.	.
[0,2]	.D====eeeER	.	.
[1,0]	.DeeE-----R	.	.
[1,1]	. D=eeeE---R	.	.
[1,2]	. D=====eeeER	.	.
[2,0]	. DeeE-----R	.	.
[2,1]	. D=====eeeER	.	.
[2,2]	. D=====eeeER	.	.

# Introduction to MCA

## An example

test/tools/llvm-mca/X86/BtVer2/dot-products.s

```
vmulps    %xmm0,  %xmm1,  %xmm2  
vhaddps   %xmm2,  %xmm2,  %xmm3  
vhaddps   %xmm3,  %xmm3,  %xmm4  
  
 llvm-mca -mtriple=x86_64 -mcpu=btver2 \  
           -iterations=300 dot-products.s
```

### Timeline

Index	Timeline	Instruction
[0,0]	01234 . . .	vmulps
[0,1]	0123456789 . .	vhaddps
[0,2]	0123456789 . .	vhaddps
[1,0]	0123456789 . .	vmulps
[1,1]	0123456789 . .	vhaddps
[1,2]	0123456789 . .	vhaddps
[2,0]	0123456789 . .	vmulps
[2,1]	0123456789 . .	vhaddps
[2,2]	0123456789 . .	vhaddps