

Instrumentor: Easily Customizable Code Instrumentation based on LLVM

Kevin Sala (salapenades1@llnl.gov)
Johannes Doerfert (jdoerfert@llnl.gov)

9th LLVM Performance Workshop @ CGO 2025
March 1st, 2025



Instrumenting Code

- **Track runtime behavior of apps**
 - Debugging and sanitization
 - Logging of events
 - Monitor resource usage
 - Performance analysis for optimization



Instrumenting Code

- Track runtime behavior of apps
 - Debugging and sanitization
 - Logging of events
 - Monitor resource usage
 - Performance analysis for optimization

Original code:

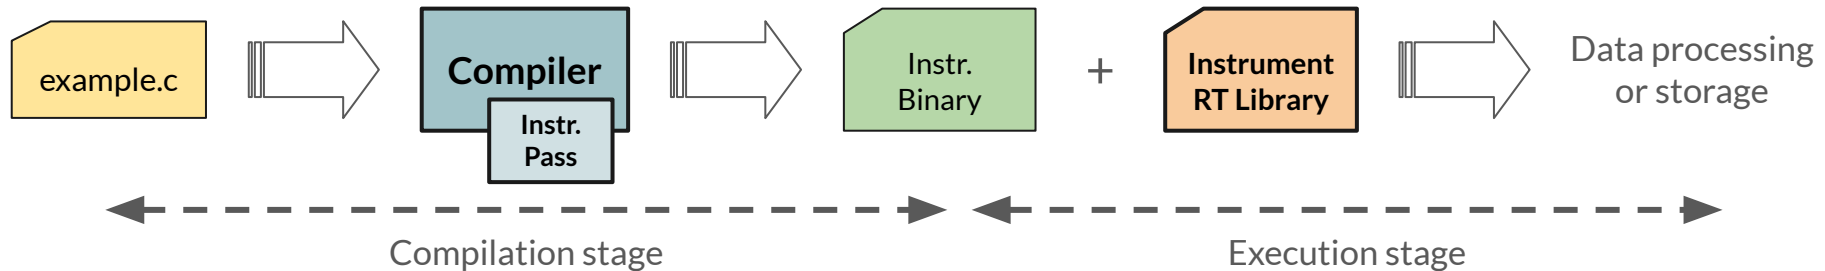
```
i32 myfunc(ptr %p) {  
    %v = load i32, ptr %p, align 8  
    store i32 10, ptr %p, align 8  
    ret i32 %v  
}
```

Instrumented code:

```
i32 myfunc(ptr %p) {  
    call void @__before_load(ptr %p, i32 4)  
    %v = load i32, ptr %p, align 8  
    store i32 10, ptr %p, align 8  
    ret i32 %v  
}
```

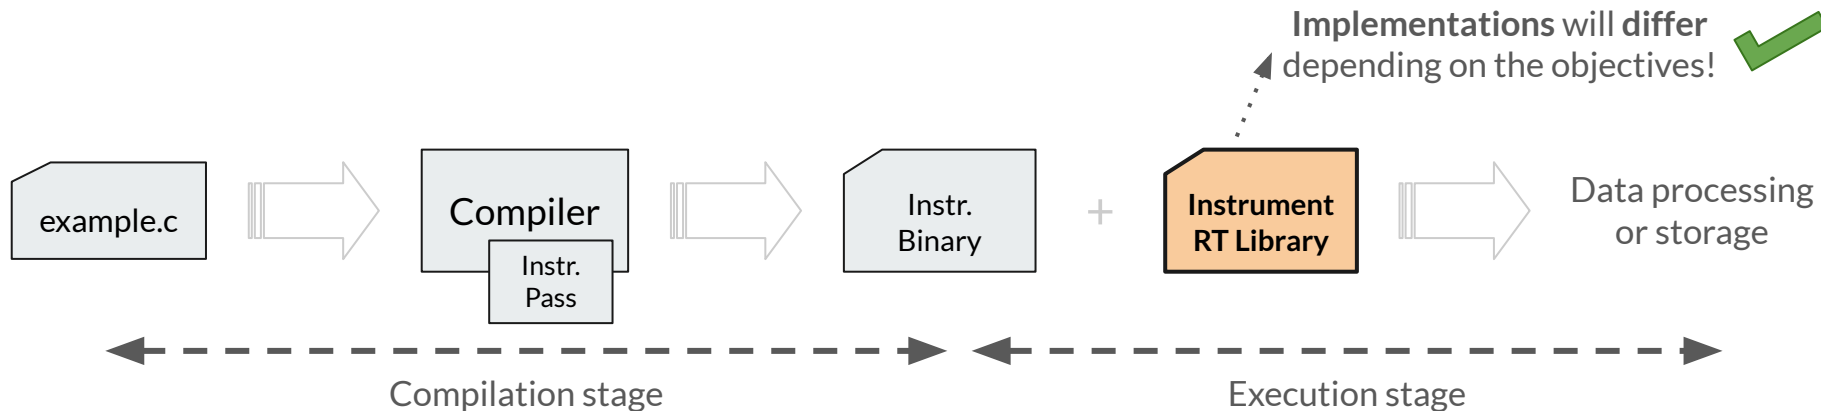
Instrumentation Support

- Two main actors
 - a. **Compiler** augments the original code with extra code
 - b. **Runtime component** receives that data during the execution



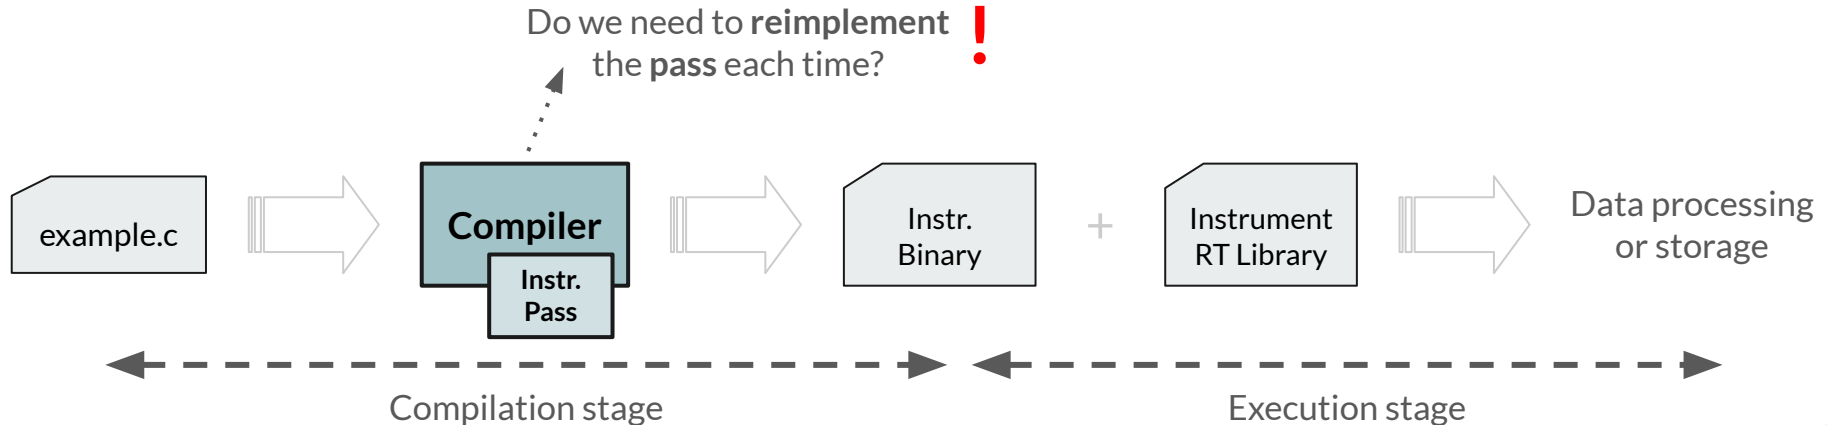
Instrumentation Support

- Two main actors
 - a. **Compiler** augments the original code with extra code
 - b. **Runtime component** receives that data during the execution



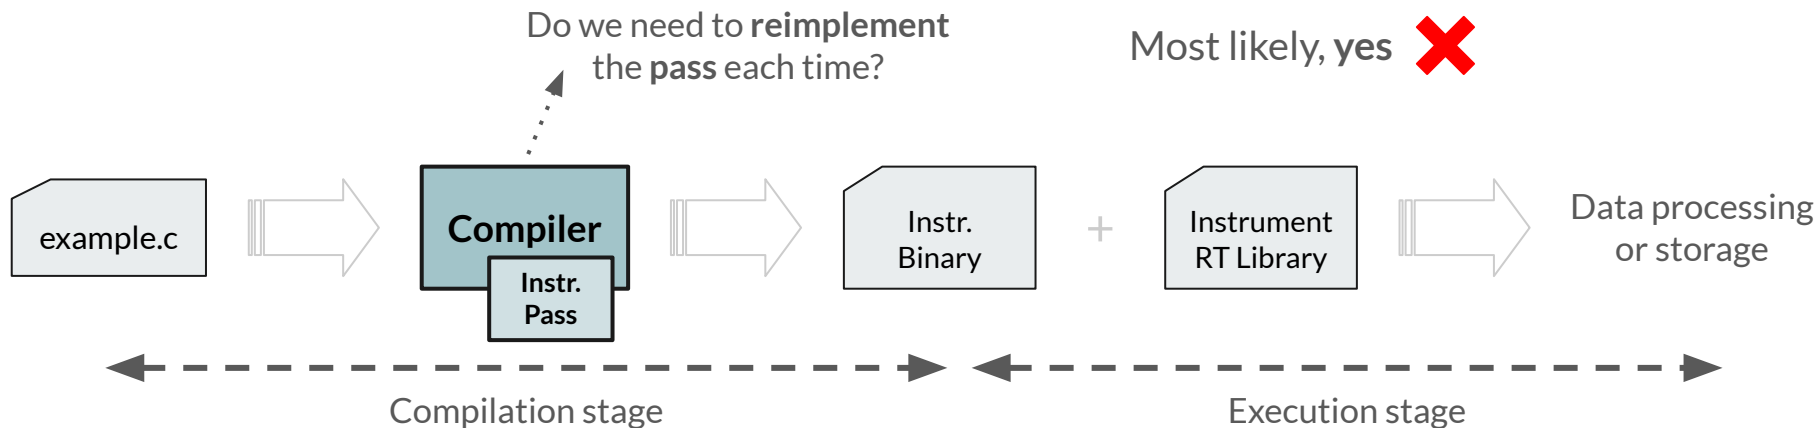
Instrumentation Support

- Two main actors
 - a. **Compiler** augments the original code with extra code
 - b. **Runtime component** receives that data during the execution




Instrumentation Support

- Two main actors
 - a. **Compiler** augments the original code with extra code
 - b. **Runtime component** receives that data during the execution



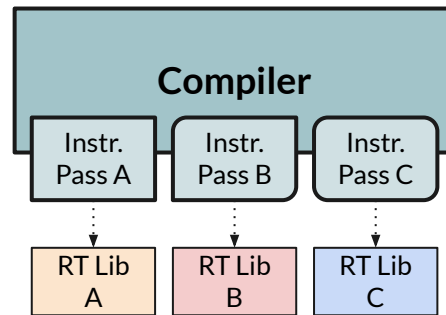
Instrumentation Support

- Compilers lack generic mechanisms for instrumenting



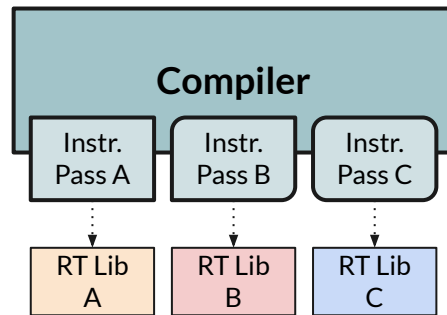
Instrumentation Support

- Compilers **lack generic mechanisms** for instrumenting
 - **Multiple passes** implement **custom** logic
 - Generally similar but quite different




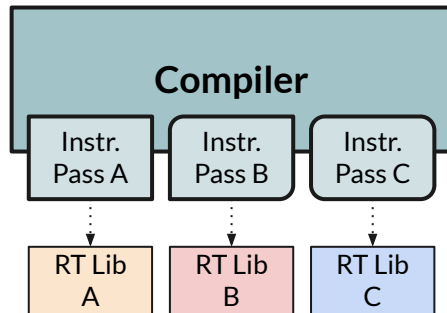
Instrumentation Support

- Compilers **lack generic mechanisms** for instrumenting
 - **Multiple passes** implement **custom** logic
 - Generally similar but quite different
- **Missing significant opportunities** like
 - Improving code **maintainability**
 - Reducing code **replication**
 - **Simplifying development** of instrumentation tools





Why not a Generic Instrumentation Pass?

- New *Instrumentor* pass in LLVM
 - Generic, customizable and extendable
 - Enabling **multiple** uses and users



Why not a Generic Instrumentation Pass?

- New *Instrumentor* pass in LLVM
 - Generic, customizable and extendable
 - Enabling **multiple** uses and users
- Exploiting the opportunities
 - Improve code **maintainability**
 - Reduce code **replication**
 - **Simplify development** of instrumentation tools



Instrumentor



Instrumentor Pass

Before *Instrumentor* pass:

```
i32 myfunc(ptr %p) {  
    %v = load i32, ptr %p, align 8  
    store i32 10, ptr %p, align 8  
    ret i32 %v  
}
```

Instrumentor Pass

```
{
  "configuration": {
    "runtime_prefix": "__instrumentor_",
    "runtime_stubs_file": "rt.c"
  },
  "instruction_pre": {
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.replace": false,
      "pointer_as": false,
      "value_size": true,
      "alignment": true,
      "is_volatile": true
    }
  }
}
```

Before *Instrumentor* pass:

```
i32 myfunc(ptr %p) {
    %v = load i32, ptr %p, align 8
    store i32 10, ptr %p, align 8
    ret i32 %v
}
```

Instrumentor Pass

```
{
  "configuration": {
    "runtime_prefix": "__instrumentor_",
    "runtime_stubs_file": "rt.c"
  },
  "instruction_pre": {
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.replace": false,
      "pointer_as": false,
      "value_size": true,
      "alignment": true,
      "is_volatile": true
    }
  }
}
```

Before *Instrumentor* pass:

```
i32 myfunc(ptr %p) {
    %v = load i32, ptr %p, align 8
    store i32 10, ptr %p, align 8
    ret i32 %v
}
```

After *Instrumentor* pass:

```
i32 myfunc(ptr %p) {
    call void @__instrumentor_pre_load(
        ptr %p, i32 4, i32 8, i32 0)
    %v = load i32, ptr %p, align 8
    store i32 10, ptr %p, align 8
    ret i32 %v
}
```


Instrumentor Pass

```
{
  "configuration": {
    "runtime_prefix": "__instrumentor_",
    "runtime_stubs_file": "rt.c"
  },
  "instruction_pre": {
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.replace": true,
      "pointer.as": false,
      "value_size": true,
      "alignment": true,
      "is_volatile": true
    }
  }
}
```


Before *Instrumentor* pass:

```
i32 myfunc(ptr %p) {
  %v = load i32, ptr %p, align 8
  store i32 10, ptr %p, align 8
  ret i32 %v
}
```

After *Instrumentor* pass:

```
i32 myfunc(ptr %p) {
  %np = call ptr @__instrumentor_pre_load(
    ptr %p, i32 4, i32 8, i32 0)
  %v = load i32, ptr %np, align 8
  store i32 10, ptr %p, align 8
  ret i32 %v
}
```

How does the Instrumentor work?



`opt -passes=instrumentor -instrumentor-read-config-file=file.json example.ll -S`

or

`clang -Xclang -finstrumentor -mllvm -instrumentor-read-config-file=file.json example.c`



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Functions
 - Global variables
 - Module
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"instruction_pre": {
  "load": {
    "enabled": true,
    "pointer": true,
    "pointer.replace": true,
    "pointer_as": true,
    "base_pointer_info": true,
    "value_size": true,
    "alignment": true,
    "value_type_id": true,
    "atomicity_ordering": true,
    "sync_scope_id": true,
    "is_volatile": true
  },
  "store": {
    "enabled": true,
    "pointer": true,
    "pointer.replace": true,
    "pointer_as": true,
    "base_pointer_info": true,
    "value": true,
    "value_size": true,
    "alignment": true,
    "value_type_id": true,
    "atomicity_ordering": true,
    "sync_scope_id": true,
    "is_volatile": true
  },
}
```



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"instruction_pre": {
  "load": {
    "enabled": true,
    "pointer": true,
    "pointer.replace": true,
    "pointer_as": true,
    "base_pointer_info": true,
    "value_size": true,
    "alignment": true,
    "value_type_id": true,
    "atomicity_ordering": true,
    "sync_scope_id": true,
    "is_volatile": true
  },
  "store": {
    "enabled": true,
    "pointer": true,
    "pointer.replace": true,
    "pointer_as": true,
    "base_pointer_info": true,
    "value": true,
    "value_size": true,
    "alignment": true,
    "value_type_id": true,
    "atomicity_ordering": true,
    "sync_scope_id": true,
    "is_volatile": true
  },
}
```



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
 - Function calls (+ inspection of args)
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"instruction_pre": {  
  "call": {  
    "enabled": true,  
    "callee": true,  
    "callee_name": true,  
    "intrinsic_id": true,  
    "allocation_info": true,  
    "num_parameters": true,  
    "parameters": true,  
    "parameters.replace": true,  
    "is_definition": true  
  },  
}
```



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
 - Function calls (+ inspection of args)
 - Allocas
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"instruction_post": {  
  "alloca": {  
    "enabled": true,  
    "address": true,  
    "address.replace": true,  
    "size": true,  
    "alignment": true  
  },  
}
```



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
 - Function calls (+ inspection of args)
 - Allocas
 - Branches, compares
 - ...
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"instruction_post": {  
  "alloca": {  
    "enabled": true,  
    "address": true,  
    "address.replace": true,  
    "size": true,  
    "alignment": true  
  },  
}
```



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
 - Function calls (+ inspection of args)
 - Allocas
 - Branches, compares
 - ...
 - Function enter/exit (+ inspect of args)
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"function_pre": {  
  "function": {  
    "enabled": true,  
    "address": true,  
    "name": true,  
    "num_arguments": true,  
    "arguments": true,  
    "arguments.replace": true  
  }  
},
```




Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
 - Function calls (+ inspection of args)
 - Allocas
 - Branches, compares
 - ...
 - Function enter/exit (+ inspect of args)
 - **Global variables**
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"global_pre": {  
  "globals": {  
    "enabled": true,  
    "address": true,  
    "address.replace": true,  
    "name": true,  
    "initial_value": true,  
    "initial_value_size": true,  
    "is_constant": true  
  }  
},
```



Instrumentor

- Instrumentation opportunities
 - Instructions
 - Loads, stores
 - Function calls (+ inspection of args)
 - Allocas
 - Branches, compares
 - ...
 - Function enter/exit (+ inspect of args)
 - Global variables
 - **Module constructor/dtor**
- Position of the instrumentation
 - Before (*pre*) and/or after (*post*)

```
"module_pre": {  
  "module": {  
    "enabled": true,  
    "module_name": true,  
    "name": true  
  }  
},  
"module_post": {  
  "module": {  
    "enabled": true,  
    "module_name": true,  
    "name": true  
  }  
},
```



Instrumentor

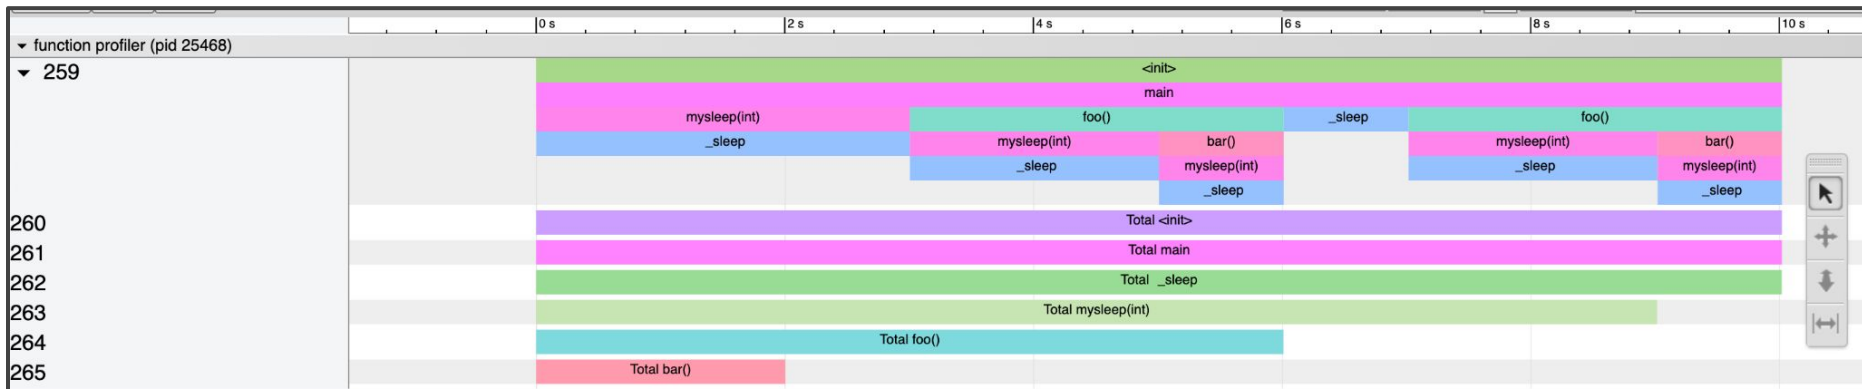
- Instrumentation opportunities
 - Instructions
 - Function enter/exit (+ inspect of args)
 - Global variables
 - Module constructor/dtor
- Other opportunities for **optimization**
 - Loop range info
 - Base pointer info

```
"special_value": {  
  "base_pointer_info": {  
    "enabled": true,  
    "base_pointer": true,  
    "base_pointer.replace": true,  
    "base_pointer_kind": true  
  },  
  "loop_value_range": {  
    "enabled": true,  
    "initial_loop_val": true,  
    "final_loop_val": true  
  }  
}
```

Use cases

Example Use: Profiler

The final result (visualized):



```
~/s/l/instrumentor >>> wc -l profiler.cpp
56 profiler.cpp
~/s/l/instrumentor >>> wc -l profiler.json
26 profiler.json
```



Example Use: Profiler

Used JSON:

```
{
  "configuration": {
    "runtime_prefix": "__instrumentor_"
  },
  "function_pre": {
    "function": {
      "enabled": true,
      "address": true,
      "name": true
    }
  },
  "instruction_pre": {
    "call": {
      "enabled": true,
      "callee": true,
      "callee_name": true
    }
  },
  "instruction_post": {
    "call": {
      "enabled": true,
      "callee": true,
      "callee_name": true
    }
  }
}
```

Example Use: Profiler

```
// LLVM Instrumentor stub runtime
#include <stdio.h>
#include "llvm/Demangle/Demangle.h"
#include "llvm/Support/Error.h"
#include "llvm/Support/TimeProfiler.h"

extern "C" {
struct __init_ty {
    __init_ty() {
        llvm::timeTraceProfilerInitialize(10, "function profiler", true);
        llvm::timeTraceProfilerBegin("<init>", "");
    }
    ~__init_ty() {
        if (has_main)
            llvm::timeTraceProfilerEnd();
        llvm::timeTraceProfilerEnd();
        if (auto Err = llvm::timeTraceProfilerWrite("prof.json", "prof.alt.json"))
            printf("Error writing out the time trace: %s\n",
                llvm::toString(std::move(Err)).c_str());
        llvm::timeTraceProfilerCleanup();
    }
    void *callee = nullptr;
    bool callee_found = false;
    bool has_main = false;
} __state;
```

Used Runtime:

```
// Continuation
void __instrumentor_pre_function(void *address, char *name) {
    if (__state.callee == address && !__state.callee_found) {
        llvm::timeTraceProfilerBegin(llvm::demangle(name), "");
        __state.callee_found = true;
    }
    if (!memcmp(name, "main", 4)) {
        __state.has_main = true;
        llvm::timeTraceProfilerBegin("main", "");
    }
}

void __instrumentor_pre_call(void *callee, char *callee_name) {
    llvm::timeTraceProfilerBegin(
        callee_name ? llvm::demangle(callee_name) : "<indirect>", "");
    if (!callee_name)
        __state.callee = callee;
}

void __instrumentor_post_call(void *callee, char *callee_name) {
    if (__state.callee_found) {
        __state.callee = nullptr;
        __state.callee_found = false;
        llvm::timeTraceProfilerEnd();
    }
    llvm::timeTraceProfilerEnd();
}
}
```

Example Use: Detect dead and redundant stores

OK

```
int A;

int main() {
    A = 0;
    A++;
    fprintf(stdout, "value of A: %d\n", A);
}
```

```
[salapenades1@tioga11]~/deadstore% ./main
value of A: 1
```

Dead Store

```
int A;

int main() {
    A = 0;
    A = 1;
}
```

```
[salapenades1@tioga11]~/deadstore% ./main
[rt] detected dead store (old: 0, new: 1)
```

Redundant Store

```
int A;

int main() {
    A = 0;
    fprintf(stdout, "value of A: %d\n", A);
    A = 0;
}
```

```
[salapenades1@tioga11]~/deadstore% ./main
[rt] detected redundant store (old: 0, new: 0)
```


Example Use: Detect dead and redundant stores

Used JSON

```
{
  "configuration": {
    "runtime_prefix": "__rt_",
    "runtime_prefix.description": "The runtime API prefix."
  },
  "instruction_pre": {
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.description": "The accessed pointer.",
      "value_size": true,
      "value_size.description": "The size of the loaded value."
    },
    "store": {
      "enabled": true,
      "pointer": true,
      "pointer.description": "The accessed pointer.",
      "value": true,
      "value.description": "The stored value.",
      "value_size": true,
      "value_size.description": "The size of the stored value."
    }
  }
}
```

Used Runtime

```
[salapenades1@tioga11]~/deadstore% wc -l rt.cpp
38 rt.cpp
```

Some extras



Extras: Use Instrumentor within LLVM

- Use Instrumentor programmatically w/o JSON file
 - Fine-grained control of what is instrumented
 - Pass custom data to RT calls

- Using class inheritance and callbacks

```
LoadIO::ConfigTy LICConfig;
LICConfig.PassPointerAS = false;
LICConfig.PassLoopValueRangeInfo = false;
LICConfig.PassValue = false;
LICConfig.ReplaceValue = false;
LICConfig.PassAlignment = false;
LICConfig.PassValueTypeId = false;
LICConfig.PassAtomicityOrdering = false;
LICConfig.PassSyncScopeId = false;
LICConfig.PassIsVolatile = false;
auto *LIC = InstrumentationConfig::allocate<LoadIO>(<*/IsPRE=*/true);
LIC->HoistKind = HOIST_MAXIMALLY;
LIC->CB = [&](Value &V) {
    return LSI.shouldInstrumentLoad(cast<LoadInst>(V), IIRB);
};
LIC->init(*this, IIRB, &LICConfig);
```

Extras: Systematic RT Functions

- Instrumentor generates systematic RT function prototypes

Runtime A

```
#include <stdint.h>
#include <stdio.h>

void __rt1_pre_load(void *pointer) {}

void __rt1_pre_store(void *pointer) {}
```

Runtime B

```
#include <stdint.h>
#include <stdio.h>

void __rt2_pre_load(void *pointer, int32_t value_size) {}

void __rt2_pre_store(void *pointer, int64_t value, int32_t value_size) {}

void __rt2_post_load(void *pointer, int64_t value, int32_t value_size) {}
```



Extras: Auto Generate RT Stub

1)

```
{
  "configuration": {
    "runtime_prefix": "--rt",
    "runtime_stubs_file": "rt.c"
  },
  "module_pre": {
    "module": {
      "enabled": true,
      "module_name": true,
      "name": true
    }
  },
  "instruction_pre": {
    "alloca": {
      "enabled": true,
      "address": true,
      "address.replace": true,
      "size": true,
      "alignment": true
    },
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.replace": true,
      "pointer_as": true,
      "value_size": true,
      "alignment": true,
      "value_type_id": true,
      "atomicity_ordering": true,
      "is_volatile": true
    }
  }
}
```



Extras: Auto Generate RT Stub

2) `opt -passes=instrumentor -instrumentor-read-config-file=file.json empty.ll -S`

1)

```
{
  "configuration": {
    "runtime_prefix": "_rt_",
    "runtime_stubs_file": "rt.c"
  },
  "module_pre": {
    "module": {
      "enabled": true,
      "module_name": true,
      "name": true
    }
  },
  "instruction_pre": {
    "alloca": {
      "enabled": true,
      "address": true,
      "address.replace": true,
      "size": true,
      "alignment": true
    },
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.replace": true,
      "pointer_as": true,
      "value_size": true,
      "alignment": true,
      "value_type_id": true,
      "atomicity_ordering": true,
      "is_volatile": true
    }
  }
}
```

Extras: Auto Generate RT Stub

2) `opt -passes=instrumentor -instrumentor-read-config-file=file.json empty.ll -S`

3) `rt.c` 

```
// LLVM Instrumentor stub runtime

#include <stdint.h>
#include <stdio.h>

void __rt_pre_module(char *module_name, char *name) {
    printf("module pre -- module_name: %s, name: %s\n", module_name, name);
}

void *__rt_pre_load(void *pointer, int32_t pointer_as, int32_t value_size,
                    int64_t alignment, int32_t value_type_id,
                    int32_t atomicity_ordering, int8_t is_volatile) {
    printf("load pre -- pointer: %p, pointer_as: %i, value_size: %i, "
           "alignment: %lli, value_type_id: %i, atomicity_ordering: %i, "
           "is_volatile: %i\n", pointer, pointer_as, value_size, alignment,
           value_type_id, atomicity_ordering, is_volatile);
    return pointer;
}

void __rt_pre_alloca(int64_t size, int64_t alignment) {
    printf("alloca pre -- size: %lli, alignment: %lli\n", size, alignment);
}
```


1)

```
{
  "configuration": {
    "runtime_prefix": "__rt_",
    "runtime_stubs_file": "rt.c"
  },
  "module_pre": {
    "module": {
      "enabled": true,
      "module_name": true,
      "name": true
    }
  },
  "instruction_pre": {
    "alloca": {
      "enabled": true,
      "address": true,
      "address.replace": true,
      "size": true,
      "alignment": true
    },
    "load": {
      "enabled": true,
      "pointer": true,
      "pointer.replace": true,
      "pointer_as": true,
      "value_size": true,
      "alignment": true,
      "value_type_id": true,
      "atomicity_ordering": true,
      "is_volatile": true
    }
  }
}
```

Complex use case: Sanitizer

Complex Use: Novel Address Sanitizer

- ASAN use extra memory and accesses
 - Requires extra memory per allocation
 - Requires 2x memory accesses
 - False negatives



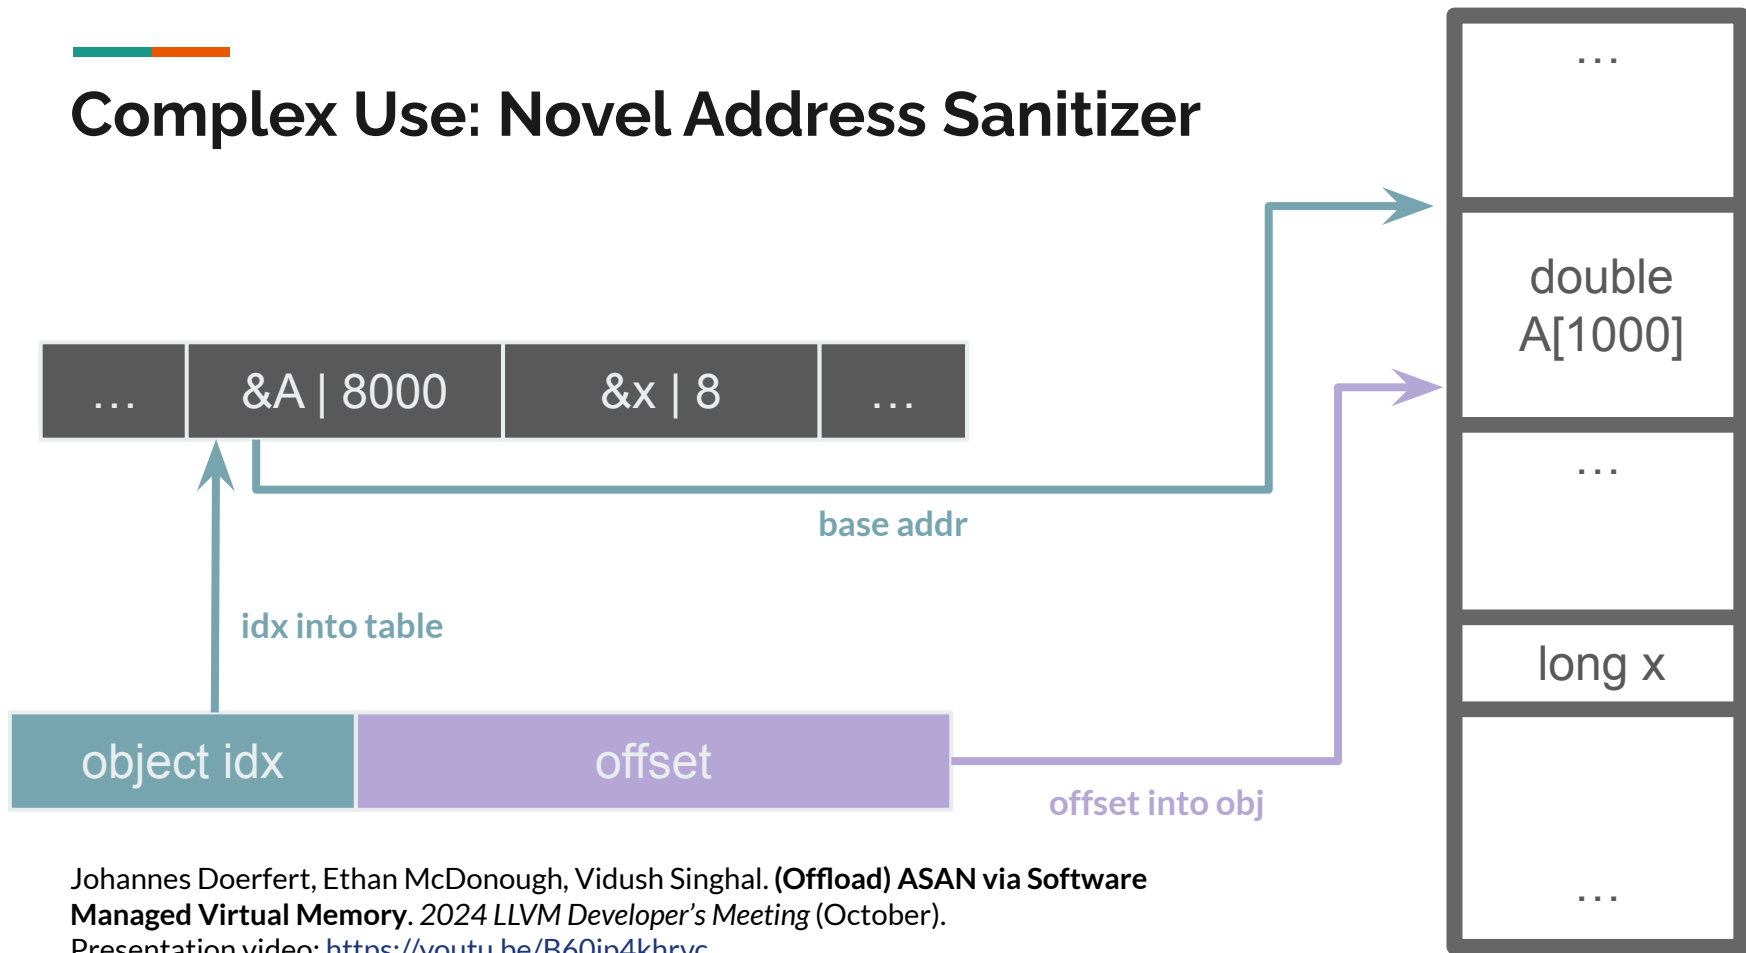


Complex Use: Novel Address Sanitizer

- **Idea:** Use virtual pointers to encode object's info
 - Replace real pointers with **virtual** pointers
 - No red zones

Complex Use: Novel Address Sanitizer

Virtual Ptr



Johannes Doerfert, Ethan McDonough, Vidush Singhal. **(Offload) ASAN via Software Managed Virtual Memory**. 2024 LLVM Developer's Meeting (October).
Presentation video: <https://youtu.be/B60jp4khrvc>

Complex Use: Novel Address Sanitizer

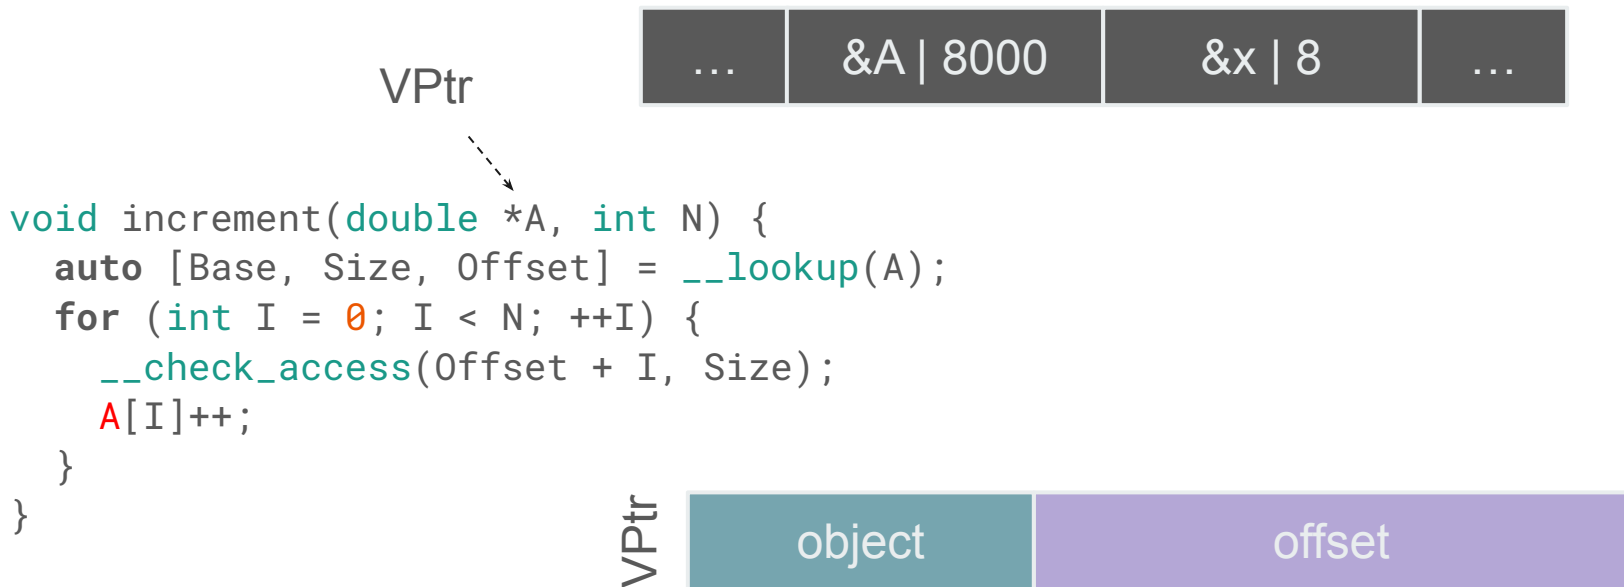


```
void increment(double *A, int N) {  
  
    for (int I = 0; I < N; ++I) {  
  
        A[I]++;  
    }  
}
```

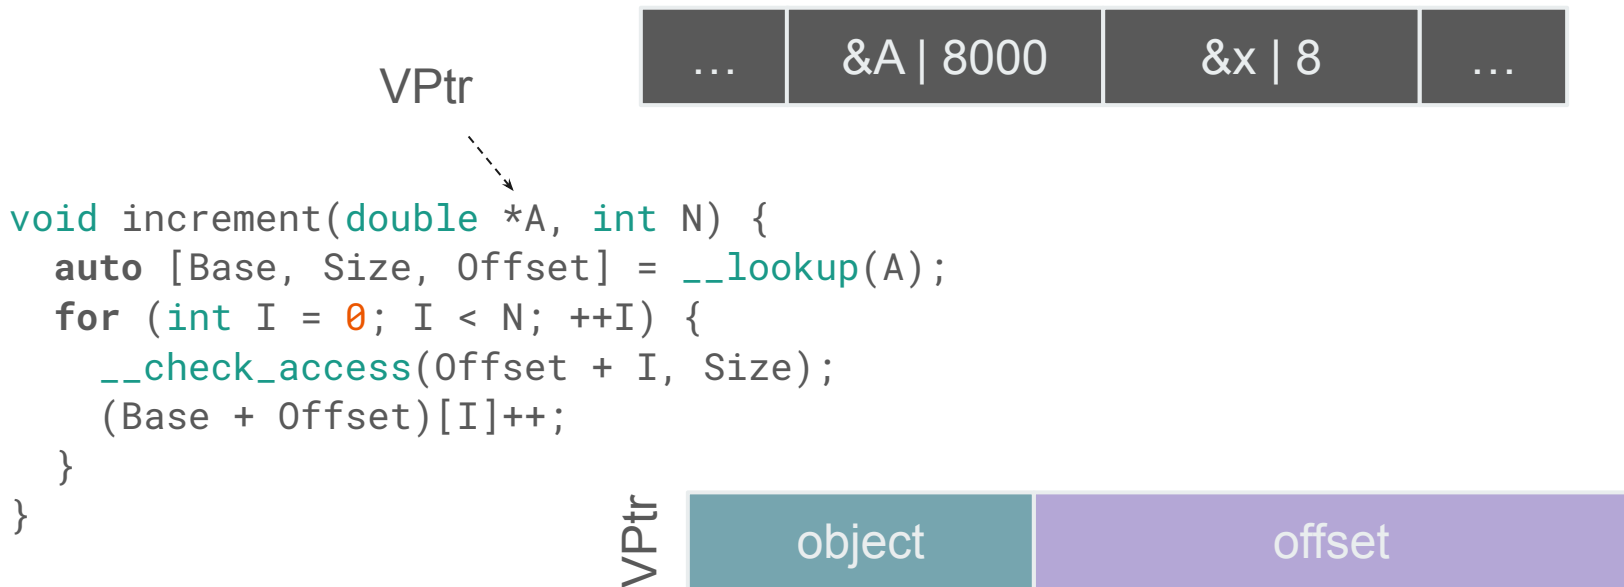
VPtr



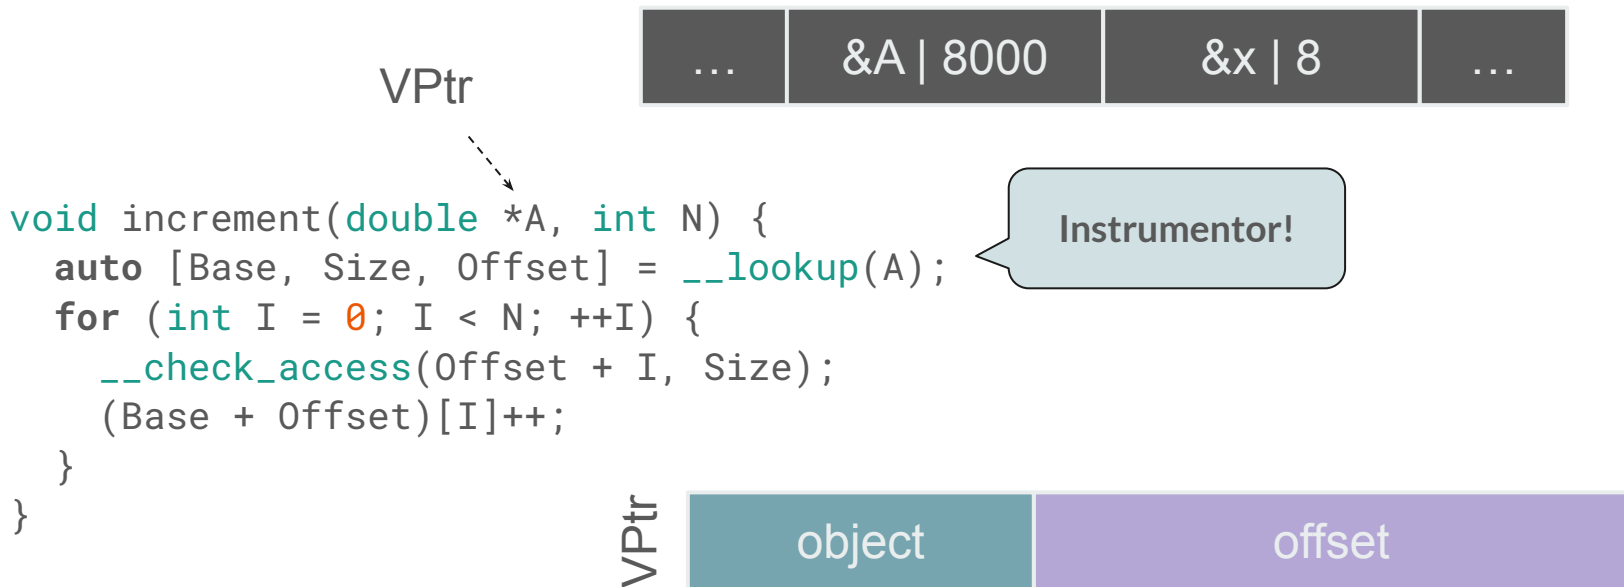
Complex Use: Novel Address Sanitizer



Complex Use: Novel Address Sanitizer



Complex Use: Novel Address Sanitizer





Conclusions

- **Instrumentor**: a customizable instrumentation based on LLVM
 - **Unified** way to **instrument** programs
 - Easy to **customize** as a user, easy to **extend** as a developer!
 - Paving the path for **future instrumentation-based tools**
- Many common use cases
 - Time profiling
 - Gather runtime information
 - etc.
- More complex use cases
 - InputGen [1]
 - Address Sanitizer (CPU and GPU code)



Thank you!

Kevin Sala (salapenades1@llnl.gov)
Johannes Doerfert (jdoerfert@llnl.gov)