






Eliminating Entire Classes of Memory Safety Vulnerabilities in C and C++

Devin Coughlin

EuroLLVM 2025



Security attacks are on the rise,
threatening both financial and
physical safety

Memory safety is a
critical challenge

Memory Safety Strategy

- Full **memory safety** requires using a **memory-safe language**
- There is **too much code** to rewrite it **all**
- Adopt memory safe languages in **new code**
 - Rewrite **high-value** codebases in a memory safe language

Metaphor: Islands of Memory Safety



- **Islands** of memory safe code...
- In a **vast ocean** of memory-unsafe C and C++.

**Must protect ocean of large existing
C/C++ codebases**

Cannot Make C/C++ Fully Memory Safe

- But can aim to **eliminate entire classes of vulnerabilities**
- If we can't eliminate, **mitigate strongly!**
- **Static and dynamic bug-finding** are useful but **not comprehensive**

Definition of Success

- **Shifting attacker behavior**
 - Drastically **increase cost** of attack development
 - **Attackers move on** to other vulnerability classes or unprotected codebases
- Ultimate sign of success
 - **Memory safety** bugs are **no longer the low-hanging fruit**
 - Attackers move on to **logic bugs**

Dimensions of Memory Safety

- Guaranteed initialization
- Bounds safety
- Lifetime safety
- Type (cast) safety
- Thread safety

Guaranteed Initialization

Guaranteed Initialization for C and C++

- Developed -ftrivial-auto-var-init extension in Clang to **protect stack**
 - Guarantee **initialization to zero**
 - Prevents **information disclosure** and many **stack grooming** attacks
- **Zero-initialize** on `free()` to protect **heap**
- Deployed on **hundreds of millions of lines of code**
- Not perfect: **Zero not intended** initial value in ~20% of cases
- **Still pretty good!**

More Information

LLVM RFC

<https://discourse.llvm.org/t/rfc-automatic-variable-initialization/50327>

Bounds Safety

Bounds Safety for C

- Developed -fbounds-safety extension for C
- Users specify buffer/bounds relations with **code annotations**
- **Run-time bounds checks** trap on out-of-bounds memory accesses
- **Compile-time rejection** of code if **bounds not determinable at run time**
- Made key design choices to lower adoption cost
 - Time to adopt: ~1 hour per 2,000 LOC
 - Adopted in **millions lines of C code**

Parameter annotations for incremental adoption

Code modification to note that count is the element count of buf

```
void fill_array_with_indices(int *bufcounted_by(count)) {  
    for (size_t i = 0; i <= count; ++i) {  
        if (buf[i] <= 0 || i >= count) trap();  
    }  
}
```


C

Compiler-generated bounds check
based on the information provided by the annotation

Use of annotations allows **preserving binary interface and signature**

Implicit wide pointers reduce annotation burden

```
void foo(int *a __counted_by(count),  
         size_t count) {  
    int *local = a;  
    ...  
}
```



Compiler
Transform

```
// Internal representation  
typedef struct {  
    int *ptr; // Pointer value  
    int *ub; // Upper bound  
    int *lb; // Lower bound  
} wide_ptr;  
  
void foo() {  
    wide_ptr local =  
        {.ptr = a, .ub = a+count, .lb = a};  
    ...  
}
```

Local variables **implicitly carry bounds**, requiring **fewer code changes**

Light-weight annotations on ABI surface

+

Implicit wide pointer types everywhere else

=

Incrementally adoptable bounds safety with reasonable annotation burden

More Information

Clang RFC

<https://discourse.llvm.org/t/rfc-enforcing-bounds-safety-in-c-fbounds-safety/70854>

EuroLLVM 2023 Keynote

<https://www.youtube.com/watch?v=RK9bfRsMdAM>

EuroLLVM 2025 Tutorial

“Adopting -fbounds-safety in Practice”

Full Implementation

<https://discourse.llvm.org/t/the-preview-of-fbounds-safety-is-now-accessible-to-the-community>

(Also in process of being upstreamed to llvm.org)

Bounds Safe Programming Model for C++

- Built -Wunsafe-buffer-usage for C++ and **hardened libc++**
- “C++ Safe Buffers” programming model
 - Compiler **rejects raw pointer arithmetic**
 - C++ standard library provides **bounds-checked buffer abstractions** (`std::span`, `std::vector`)
- Adopted on **tens of millions of lines of code**
 - **“C written in C++”**: ~2x higher adoption cost than -fbounds-safety
 - **Modern C++**: “Zoomin” at ~30,000+ lines of code per hour

Programmer **changes type** of buf to indicate it is a **contiguous span** of ints

```
void fill_array_with_indices(span<buf>, size_t count) {  
    for (size_t i = 0; i <= count; ++i) {  
        buf[i] = i; error: unsafe buffer access [-Werror,-Wunsafe-buffer-usage]  
    }  
}
```

C++

Hardened libc++ span implementation checks bounds before indexing via **operator overloading**

More Information

Clang RFC

<https://discourse.llvm.org/t/rfc-c-buffer-hardening/65734>

LLVM 2023 Talks

<https://www.youtube.com/watch?v=pQfjn7E4Qfc> (Hardened libc++)

<https://www.youtube.com/watch?v=nPRY8-FtzZg> (-Wunsafe-buffer-usage)

Full Implementation in llvm.org Open Source

Lifetime Safety

Lifetime Safety: Ref-Counting Smart Pointer Analysis

- Developed Clang Static Analyzer WebKit checkers
- Enforces a **strict programming model** (not bug finding)
- **Prevents lifetime issues** caused by misuse of **reference-counted pointers**
- Adopted on **millions of lines of code**

Programmer changes type of resource to **indicate its lifetime** must last for entire scope

C++

```
error: Local variable 'resource' is uncounted and unsafe
RefPtr<Resource> resource = someFunctionThatReturnsARef();
doSomethingThatMightReleaseResource();
resource->use();
```

WebKit RefPtr run-time implementation guarantees lifetime until destructor is called

More Information

Full implementation in llvm.org

<https://github.com/llvm/llvm-project/tree/main/clang/lib/StaticAnalyzer/Checkers/WebKit>

Lifetime Safety: Type-Isolating Allocators

- Not all code uses a reference-counting discipline
- Developed **language extensions for typed allocation**
 - C: `-ftyped-memory-operations`
 - C++: `-ftyped-cxx-new-delete`
- **Mitigates use-after-free vulnerabilities by limiting data(pointer type confusion)**
- Deployed on **hundreds of millions of lines** of userspace code
- Similar approach deployed in xnu kernel

C: Compiler Transforms Calls to Allocators

- Driven by new “typed memory operation” attribute on allocators


System Header

```
// Standard system malloc entrypoint  
void *malloc(size_t size) __attribute__((typed_memory_operation(1, malloc_typed)));
```

```
// New typed malloc entry point with extra parameter for type information  
void *malloc_typed(size_t size, tmo_type_descriptor_ref type_info);
```

Client Code

```
struct SomeType *allocation = malloc(sizeof(struct SomeType));
```



Compiles Into

```
struct SomeType *allocation = malloc_typed(sizeof(struct SomeType), SomeType info);
```

C++: Language support for typed allocation

- Proposal **P2719**: Type-aware allocation and deallocation functions
- Passes **type as a template parameter to operator**

```
template <class T>
    requires use_special_allocation_scheme<T>
void* operator new(std::type_identity<T>, std::size_t n) { ... }

template <class T>
    requires use_special_allocation_scheme<T>
void operator delete(std::type_identity<T>, void* ptr) { ... }
```

C++

More Information

Clang RFC

<https://discourse.llvm.org/t/rfc-typed-allocator-support/79720> (For C)

<https://discourse.llvm.org/t/rfc-experimental-implementation-of-p2719-type-aware-allocation-and-deallocation-functions/83876> (For C++)

LLVM 2024 Talk

<https://www.youtube.com/watch?v=GGGaiGpm5BY>

C++ WG21 Proposal

<https://github.com/ldionne/wg21/blob/main/p2719.md>

In Process of Being Upstreamed

<https://github.com/llvm/llvm-project/pull/113510> (C++)

Type (Cast) Safety

C++: Static Analysis-Enforced Programming Model

- Developed WebKit MemoryUnsafeCastChecker for C++
- **Rejects unchecked casts** unless compiler can prove safe
- Works great for C++ codebases with
 - **Run-time** representation of **types**
 - Idiom of always using **run-time-checked casts**
- **Do not have a viable approach for C codebases nor all C++ codebases**

More Information

Full implementation in [llvm.org](#)

[https://github.com/llvm/llvm-project/blob/main/clang/lib/StaticAnalyzer/Checkers/WebKit/
MemoryUnsafeCastChecker.cpp](https://github.com/llvm/llvm-project/blob/main/clang/lib/StaticAnalyzer/Checkers/WebKit/MemoryUnsafeCastChecker.cpp)

Thread Safety

No Strong Ideas for Thread Safety for C/C++

- Needs more investigation
 - -Wthread-safety useful but does not eliminate entire class
- General recommendation for thread safety is to use thread-safe language
 - **Swift Concurrency** actor model **provides data-race freedom**

Memory Safety in Heterogeneous Codebases

Metaphor: Islands of Memory Safety



- Islands of code written in memory safe languages...
- In a vast ocean of (partially) memory-unsafe C and C++.
- **Need to protect the beaches!**

Mix of Existing Memory Safety Technologies



Swift

Fully memory safe **by design**



-fbounds-safety


Bounds-safe C language extension with **annotations**



Hardened
C++

Bounds-safe via hardened libc++ and **-Wunsafe-buffer-usage**


Need to Preserve Safety Across Boundaries



Recap: Dimensions of Memory Safety

- **Guaranteed initialization:** -ftrivial-auto-var-init
- **Bounds safety:** -fbounds-safety, -Wunsafe-buffer-usage + hardened libc++
- **Lifetime safety:** typed allocators, WebKit shared pointer checks
- **Type (cast) safety:** WebKit cast static analysis, needs more investigation
- **Thread safety:** Needs more investigation

**Synthesis: “Recipe” is to Enforce
Strict Programming Models**



Ingredients for a Programming Model

- 1 Compiler checkable local rules
- 2 Developer-provided annotations
- 3 Run-time support

Ingredient 1: Compiler Checkable Local Rules

- Rules **restrict** programmer to **subset of language**
 - Programming model can make **strong guarantees**
 - More **like programming language** feature than bug-finding
- Enable programmer to **reason about rules locally**
 - Programmer should be able to easily tell whether code follows or not
- Provide **tools to check rules automatically**
 - Ensure mistakes do not fall through the cracks

Ingredient 2: Developer-Provided Annotations

- Developers provide **annotations** expressing key **expected invariants**
- Affords **modular** (assume/guarantee) local reasoning
- Analogous to **type checking**:
 - If compiler guarantees that caller passes a string as first argument, callee can assume that first parameter is a string

Ingredient 3: Run-time Support

- Sometimes **compile-time checks** are **not enough**
- Often tools **cannot reason precisely statically** about program behavior
 - Typically **undecidable!**
- Sometimes need to establish, track, and check **global** invariants
 - Global static checking is very expensive and **does not scale to large programs**
- **Design run-time abstractions** to handle difficult cases

-fbounds-safety Recipe Instantiation

- 1 Local Rule:** Must only dereference pointers with known bounds
- 2 Developer Annotations:** Bounds attributes for pointers
- 3 Run-time Support:** Codegen to trap if dereference is out of bounds

C++ Safe Buffers Recipe Instantiation

- 1 **Local Rule:** No raw pointer arithmetic
- 2 **Developer Annotations:** Adopt span utility class for buffers
- 3 **Run-time Support:** Hardened standard library checks bounds in span

WebKit Smart Pointer Safety Recipe Instantiation

- 1 **Local Rule:** No raw pointers on stack across unknown calls
- 2 **Developer Annotations:** Adopt RefPtr utility class for pointers
- 3 **Run-time Support:** RefPtr itself handles keeping raw pointer alive

Challenge for community:

To what other security problems can we apply this recipe?

Biggest Limiting Factor is Cost of Adoption

- Zero initialization and typed allocator adoption required little adopter work
 - Easy to roll out to hundreds of millions of lines of code
- Bounds safety approaches powerful but required more adoption work
 - Adoption so far in the millions of lines of code


**Must invest in tooling to lower
adoption cost**

Need Annotation Inference Tooling

- Speed up adoption with by automatically inferring annotations
- Problem:
 - Annotations can be **checked locally**, but must be **inferred globally**
 - E.g., annotations on callee depend on annotations in caller
 - Requires **whole program reasoning!**

Proposal: Summary-Based Inference Tooling

- Achieve whole-program reasoning with **thin-LTO-style summaries**:



Summary Infrastructure Broadly Applicable

- Useful for **Clang Static Analyzer**
 - Enable **cross translation unit analysis** to **scale to larger codebases**
 - Reduce false positives and false negatives
- Useful for safe C/C++ interoperation with memory-safe languages
 - Infer lifetime annotations
 - Infer escapability annotations

Summary

- Full memory safety only comes with memory safe languages
 - Must support **safe, ergonomic interoperation** with memory-safe languages
 - Must **protect ocean of existing C/C++ code** as much as possible
- Strict programming models
 - Recipe for **eliminating** many **classes of memory-safety bugs** in C/C++
 - ① Local rules + ② Developer-provided annotations + ③ Run-time support
 - **Annotation cost** is major **limiting factor** in adoption
- Invest in **annotation inference tooling**



TM and © 2025 Apple Inc. All rights reserved.