

# LLVM Support for Sub-FP8 Quantization with RISC-V Extensions for Machine Learning Models

Jhih-Kuan Lin, Fu-Jian Shen, Kathryn Chapman, Mengshiu Yu, Jenq-Kuen Lee  
Department of Computer Science, National Tsing-Hua University, Taiwan



國立清華大學  
NATIONAL TSING HUA UNIVERSITY

## Motivation

- Model quantization optimizes performance by reducing 32-/16-bit INT/FP formats to 8-, 6-, and 4-bit formats, lowering compute, memory, and transmission costs.
- The Open Compute Project (OCP) defines floating-point formats smaller than 8-bit. We developed the **sub-FP8 ISA** extension for RISC-V to enable support for these formats in scalar operations.

## Sub-FP8 Support in LLVM

- Sub-FP8 Backend & Code Generation
- Instruction Support for Sub-FP8 in Assembler
- Efficient Pattern Matching for Sub-FP8 Operations
- Sub-FP8 C/C++ Intrinsics

```
_builtin_riscv_sfadd_f8e4m3_f8e4m3(0.1f8e4m3, 0.1f8e4m3);  
%0 = call f8e4m3 @llvm.riscv.pllab.sfadd.f8e4m3.f8e4m3.f8e4m3(f8e4m3  
0xQ1D, f8e4m3 0xQ1D)  
  
f8e4m3 = llvm.riscv.pllab.sfadd TargetConstant:i64<9897>, ConstantFP:f8e4m3<APFloat(29)>,  
ConstantFP:f8e4m3<APFloat(29)>  
  
PseudoSSET_INSERT 1, 1, implicit-def $eft, implicit-def $sft  
renamable $f15_s = noexcept PseudoSFADD killed renamable $f15_s, renamable $f15_s,  
7, 1, 1, implicit killed $eft, implicit killed $sft, implicit $frm  
  
csrwi sft, 1  
csrwi eft, 1  
sfadd fa5, fa5, fa5
```

Fig. 2. Sub-FP8 Compile Flow in LLVM (from C to assembly)

## Simulation

- The instruction set architecture (ISA) was simulated, allowing for correctness verification, using **Spike**.
- Gem5** was used to execute real programs and collect research data.




Fig. 3. The results of GEMM computation ( $(256 \times 16) \times (16 \times 256)$ ) under an environment with a 32KB L1 D-cache and compilation optimization level O3. Cache miss rates across different floating-point formats are also compared.




Fig. 4. Overall architecture of our Sub-FP8 compilation and simulation flow.

## EFT & SFT

ssft/seft[3:0]	SFT/EFT
0 0 0 1	FP8(E5M2)
0 0 1 0	FP8(E4M3)
0 0 1 1	FP6(E3M2)
0 1 0 0	FP6(E2M3)
0 1 0 1	FP4(E2M1)
others	reserve




Fig. 1. Contents of the modified Floating Point Control and Status Register (FCSR)

## RISC-V Vector Extensions

- The following vector extension instructions are proposed to support the sub-FP8 format:

Class	Instruction name	Description
Arithmetic	vsfadd, vsfsub, vsfmul, vsfddiv, vsfmacc, vsfsqrt	Vector-scalar addition, subtraction, multiplication, division, and multiply-accumulate
Comparison	vsfmin, vsfmax	Compute element-wise min/max
Reduction	vsfredosum, vsfredusum, vsfremin, vsfremax, vsdot	Perform reduction operations on a vector, such as sum or min/max, to produce a scalar result
Data Movement	vfmv, vfcvt(float -> float)	Move or convert data
Other	vsfsqrjn, vsfclass,	Other operations for floating-point numbers

## Example for GEMM

- C code snippet demonstrating a **GEMM** operation using the newly added RISC-V f8e4m3 floating-point format.

```
for (int i = 0; i < MLEN; ++i)
    for (int j = 0; j < NLEN; ++j)
        for (int k = 0; k < KLEN; ++k)
        {
            golden_array[i * NLEN + j] =
__builtin_riscv_sfmmadd_f8e4m3_f8e4m3(a_array[i * KLEN + k],
b_array[j + k * NLEN], golden_array[i * NLEN + j]);
        }
```

## Hardware Reference Design Based on CVA6

- RISC-V CVA6 Core (CV64A6) (Work-in-Progress) using the transprecision-supporting FPU, CVFPU (FPnew)




Fig. 5. CVA6 processor modified to support the sub-FP8 format and accompanying instructions. The red line shows the stages of the pipeline that use the format information provided by the FCSR (sset). New features and modules with significant modifications are shown in green.