

Pass Plugins

Past // Present // Future



Stefan Gränitz EuroLLVM Berlin 15 April 2025

Agenda

1. Past: Legacy Pass Plugins
2. Present: Modern Pass Plugins
3. Future: From Passes to Extensions?
 - Motivation
 - Proposals

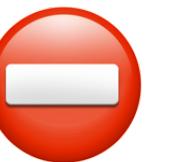
Round Tables

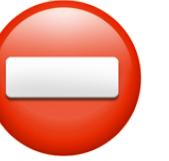


Tuesday 5 PM, right after the talk

Another one on Wednesday

Plugin term is convoluted

`ld.lld -plugin gold.so`  LTO plugin

`clang -fplugin=clad.so`  Clang frontend plugin

`clang -fpass-plugin=omvll.so`  Pass-Plugin



Past

*“Don’t dwell in the past
Don’t dream of the future
Concentrate the mind on
the present moment. Focus!*

404 @ Phabricator 😊

Past: Legacy Pass Plugins

Option `-load` dates back to the early 2000s^{1 2}

```
struct PluginLoader {  
    void operator=(const std::string &Filename);  
    static unsigned getNumPlugins();  
    static std::string& getPlugin(unsigned num);  
};  
  
static cl::opt<PluginLoader, false, cl::parser<std::string>>  
LoadOpt("load", cl::value_desc("pluginfilename"),  
        cl::desc("Load the specified plugin"));
```

Past: Legacy Pass Plugins

Plugins like Polly used static init to register new passes [1](#) [2](#)

```
class StaticInitializer {
public:
    StaticInitializer() {
        llvm::PassRegistry &Registry = *PassRegistry::getPassRegistry();
        polly::initializePollyPasses(Registry);
    }
};

static StaticInitializer InitializeEverything;
```



Present

Modern Pass Plugins

Plugin Renaissance with the New Pass Manager

Major contributions

2017 **Philip Pfaffe** adds a pass registration mechanism for Polly¹

2018 He refines it into a plugin API so that "*interaction with a plugin is always initiated from the tools perspective*"²

2020 **Serge Guelton** generalizes it, removes remaining Polly-specific code from LLVM and adds an Bye  example³

Modern Pass Plugins

On the plugin side, we implement a defined interface

```
extern "C" struct PassPluginLibraryInfo {  
    uint32_t APIVersion;  
    const char *PluginName;  
    const char *PluginVersion;  
    void (*RegisterPassBuilderCallbacks)(PassBuilder &);  
};
```

Modern Pass Plugins

```
extern "C" PassPluginLibraryInfo llvmGetPassPluginInfo() {
    return {LLVM_PLUGIN_API_VERSION, "Bye", LLVM_VERSION_STRING,
        [](PassBuilder &PB) {
            PB.registerVectorizerStartEPCallback(
                [](FunctionPassManager &PM, OptimizationLevel Level) {
                    PM.addPass(Bye());
                });
            PB.registerPipelineParsingCallback(
                [](StringRef Name, FunctionPassManager &PM, ...) {
                    if (Name == "goodbye") {
                        PM.addPass(Bye());
                        return true;
                    }
                    return false;
                });
        }});
}
```

Modern Pass Plugins

On the tools side, we load them explicitly

```
static cl::list<std::string>
PassPlugins("load-pass-plugin",
            cl::desc("Load passes from plugin library"));

for (auto &PluginFN : PassPlugins) {
    auto PassPlugin = PassPlugin::Load(PluginFN);
    if (!PassPlugin)
        continue;
    PassPlugin->registerPassBuilderCallbacks(PB);
}
```

Modern Pass Plugins: Pros

- ▶ Use `PassBuilder` the same way as in-tree tools
- ▶ Same concept in MLIR: entry-point `mlirGetPassPluginInfo()`
- ▶ Keep existing benefits:
 - C interface for plugin registration
 - Fast and easy builds against LLVM release versions

Modern Pass Plugins: Cons

Pass base class and PassBuilder definitions are C++

```
struct Bye : PassInfoMixin<Bye> {
    PreservedAnalyses run(Function &F, FunctionAnalysisManager &) {
        if (!runBye(F))
            return PreservedAnalyses::all();
        return PreservedAnalyses::none();
    }
};
```

Building Plugins correctly isn't trivial

Plugin binaries must fit target compiler's C++ ABI

Modern Pass Plugins: Tools

- ▶ `opt -load=/path/to/Bye.so -passes=goodbye`
`opt -load-pass-plugin=/path/to/Bye.so`
[docs/CommandGuide/opt.html#cmdoption-opt-load](#)
- ▶ `clang -fpass-plugin=/path/to/Bye.so`
[docs/ClangCommandLineReference.html#cmdoption-clang-fpass-plugin](#)
- ▶ `flang -fpass-plugin=/path/to/Bye.so`
[docs/FlangCommandLineReference.html#cmdoption-flang-fpass-plugin](#)
- ▶ `clang-repl -fpass-plugin=/path/to/Bye.so`
[weliveindetail.github.io/blog/post/2024/08/29/omvll-clang-repl.html](#)

Modern Pass Plugins: Tools

- ▶ `swiftc -load-pass-plugin=/path/to/bye.so`
mainline: [swiftlang/swift/pull/68985](https://github.com/swift-lang/swift/pull/68985)
- ▶ `rustc -zllvm-plugins=/path/to/bye.so`
unstable: `llvm.plugins = true` option rust-lang.zulipchat.com
- ▶ `ld.11d --load-pass-plugin=/path/to/Bye.so` (since 15.x¹)
undocumented 

Modern Pass Plugins: in the wild

Open-source projects:

- ▶ <https://github.com/llvm/llvm-project/tree/release/20.x/polly>
- ▶ <https://github.com/EnzymeAD/Enzyme>
- ▶ <https://github.com/open-obfuscator/o-mvll>

Yes, it's a niche for sure. But it might also be a chicken-egg-problem..



Future

From Passes to
Rich out-of-tree Extensions?

Motivation

Claim: There is a demand for domain-specific compiler extensions.

Evidence? Looking at sanitizers:

2017 asan,dfsan,msan,tsan,safestack,cfi,esan,scudo

2018 asan,dfsan,msan,hwasan,tsan,safestack,cfi,esan,scudo,ubsan

2019 asan,dfsan,msan,hwasan,tsan,safestack,cfi,esan,scudo,ubsan

2020 asan,dfsan,msan,hwasan,tsan,safestack,cfi,scudo,ubsan,gwp_asan

2021 asan,dfsan,msan,hwasan,tsan,safestack,cfi,scudo,ubsan,gwp_asan

2022 asan,dfsan,msan,hwasan,tsan,safestack,cfi,scudo,ubsan,gwp_asan

2023 asan,dfsan,msan,hwasan,tsan,safestack,cfi,scudo,ubsan,gwp_asan

2024 asan,dfsan,msan,hwasan,tsan,safestack,cfi,scudo,ubsan,gwp_asan

2025 asan,rtsan,dfsan,msan,hwasan,tsan,tysan,safestack,cfi,scudo,ubsan,gwp_asan,nsan

^^^^^

^^^^^

^^^^^

Sanitizers: Who are the newcomers?

We can now check for:

- ▶ TypeSanitizer: type-based aliasing violations
- ▶ NumericalStabilitySanitizer: floating point precision issues
- ▶ RealtimeSanitizer: blocking calls in code with deterministic runtime

Observations: More domain-specific + less C/C++ specific

Future: Towards Rich Out-of-tree Extensions?

Should we build everything upstream forever?

Alternative: Could we implement extensions like Sanitizers as plugins?

- [] Frontend: Attributes control where/how they apply (or not)
- [x] IR Pass: Inject instrumentation, mostly calls into a runtime library
- [] Driver: Add runtime library to the link line

Frontend with built-in Sanitizer (Realtime Sanitizer)

Clang: new [[clang::(non)blocking]] attributes translate
to built-in `llvm::Attribute::SanitizeRealtime(Blocking)`
 llvm-project/commit/f03cb005eb4b

Swift: `RTSanStandaloneSwift` package wraps C API in expression macros
 swiftpackageindex.com/package/realm-sanitizer/RTSanStandaloneSwift

Rust: `rtsan-standalone` crate wraps C API in procedural macros
 crates.io/crates/rtsan-standalone

also preparing RFC for rustc built-in support
 github.com/rust-lang/rfcs/pull/3766

Frontend without built-in Sanitizer

Clang:

- ▶ Frontend-Plugin could define attributes [[clang::(non)blocking]]
- ▶ Emit annotations instead of `llvm::Attribute::SanitizeRealtime?`
- ▶ Combine with Pass-Plugin in a single shared lib! github.com/vgvassilev/clad

Modern languages:

- ▶ Could language features emit annotations directly?



Round Table: Can we use annotations?

Or could we teach Pass-Plugins to define LLVM attributes?

How to add a runtime library to the link line?

Named metadata entries for auto-linking might help:

- ▶ `lldm.linker.options`

[docs/LangRef.html#automatic-linker-flags-named-metadata](#)

- ▶ `lldm.dependent-libraries`

[docs/LangRef.html#dependent-libs-named-metadata](#)

 **Round Table:** Can we make it consistent or find a better way?

Future: Rich out-of-tree Extensions?

Doesn't seem impossible!

- [x] Frontend: Attributes control where/how they apply (or not)
- [x] IR Pass: Inject instrumentation, mostly calls into a runtime library
- [x] Driver: Add runtime library to the link line

Realtime Sanitizer: What is the story?

- ▶ Start a **hack in a fork** and reach a PoC
- ▶ Promote in domain-specific communities and find interested contributors
adc23.sched.com/event/1PudD/radsan-a-realtime-safety-sanitizer
- ▶ Write RFC discourse.llvm.org/t/rfc-nolock-and-noalloc-attributes/76837/
- ▶ RFC considered to mature downstream or implement outside of LLVM¹
- ▶ [#92460](#) merged upstream in May 2024 and **150+ PRs** since

Realtime Sanitizer: Developer perspective

Pro upstream:

- **re-use infrastructure from other sanitizers**
- reviewers give guidance, help find issues and propose improvements
- boost reachable audience and get maximum convenience for users
- no immediate commercial interests (apparently)

Downside:

- requirements on code-quality and cross-platform support
- extra complexity from considering interference with other sanitizers

Future: Proposals

If we want to promote out-of-tree extensions, we could:

1. Provide re-usable infrastructure
2. Make it a playground to test new ideas
3. Motivate vendors to support plugins!

1. Re-usable infrastructure for out-of-tree extensions

Today:

- Plugin interface: unit-tests + Bye example with LIT tests
- Most bots don't build examples
- Most vendors don't ship examples
- Bye is quite primitive

Make it a **Reference Plugin**, that is built and deployed by default?

1. Re-usable infrastructure for out-of-tree extensions

Reference Plugin:

- (1) Complexity of real-world extension
- (2) Should work for LLVM and MLIR
- (3) Do something useful for experimentation
- (4) Consider a pure C interface?

2. Playground to test new ideas and not fork LLVM

Load a Python script and run it in as a pass?

- (1) Python is popular + real-world complexity (e.g. static libPython?)
- (2) Bindings for LLVM and MLIR
- (3) Write IR transforms without building the plugin!

Two open-source repos with proof-of-concept:

- ▶ C++ with Numba's llvmlite: github.com/weliveindetail/llvm-py-pass
- ▶ Rust with llvmcpy from rev.ng: github.com/aneeshdurg/pyllvmpass

3. Motivate vendors to support and ship plugins!

Concerns:

- Security and tampering with internals (probably Apple)
 - Would code-signing checks help?
- Compatibility, versioning and dependence (probably Rust)
 - Would a pure C API version help?

 **Round Table:** Let's keep dreaming of a bright future for a bit!



Pass Plugins

Round Tables

Tuesday 5 PM, right after the talk

Another one on Wednesday

Contact

stefan.graenitz@gmail.com

weliveindetail.github.io/blog/about/