

# Can Vectorization Slow Down Performance?

## Addressing the Challenges of Vectorizing Stride Access

October 28, 2025

Kotaro Kinoshita (@kinoshita-fj)

2025 US LLVM Developers' Meeting

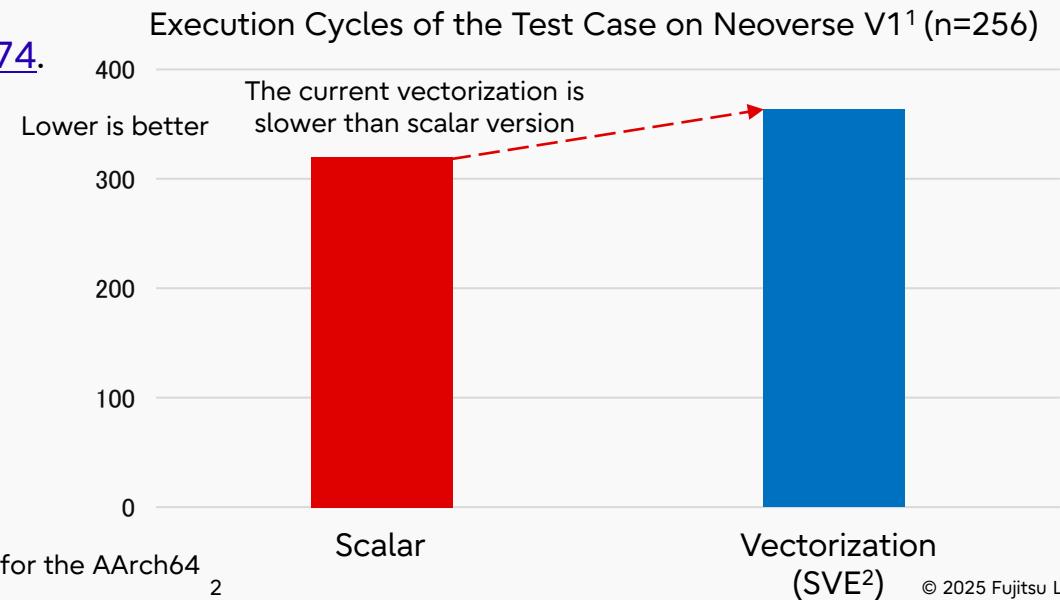


# Introduction

- Vectorization is generally expected to improve performance. However, we found performance degradation in some cases of vectorizing strided access.
- This talk focuses on one such case.
- For example, vectorizing the test case leads to performance degradation by inefficient code generation.
- We opened the issue for this [#129474](#).

Test Case

```
void func(double *a, double *b, int n) {
    for (int i = 0; i < n; i++) {
        a[i] = b[i * 10] + 1;
    }
}
```



# Current Address Calculation is Inefficient



LLVM (21.1.0) Vectorization (SVE)

```
void func(double *a, double *b, int n) {
    for (int i = 0; i < n; i++) {
        a[i] = b[i * 10] + 1;
    }
}
```

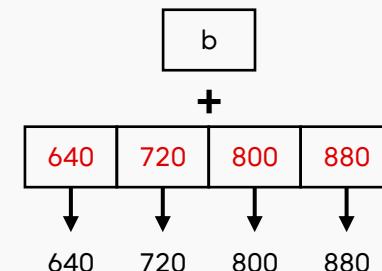
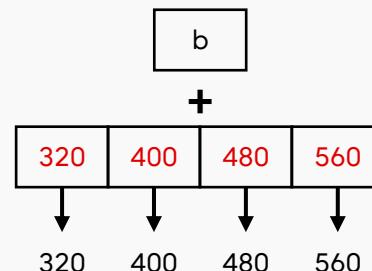
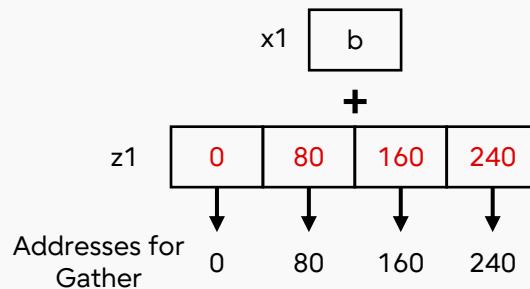


```
index z1.d, #0, #1
loop:
    add z2.d, z1.d, z0.d
    mul z1.d, z1.d, #80
    ld1d { z1.d }, p0/z, [x1, z1.d]
    ...
    mov z1.d, z2.d
    ...
```

1st Iteration

2nd Iteration

3rd Iteration



- Address calculation uses vector instructions inside the loop.

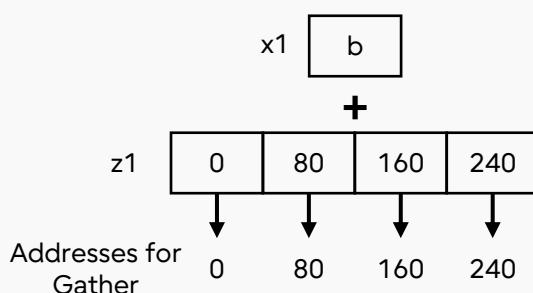
# Efficient Instructions for Strided Access

- If the pattern is identified as a strided access, efficient instructions like these can be generated.

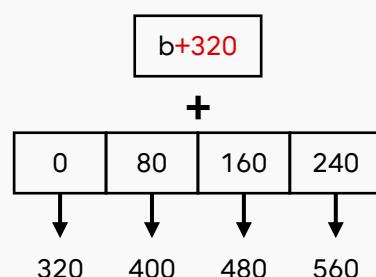
Better Vectorization

```
index z1.d, #0, #80
loop:
    ld1d { z2.d }, p0/z, [x1, z1.d]
    ...
    add x1, x1, x2
    ...
```

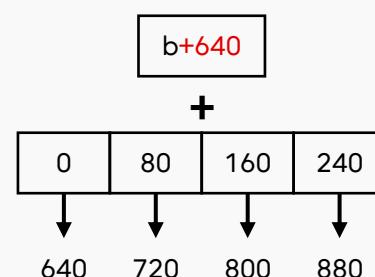
1st Iteration



2nd Iteration



3rd Iteration



- Generate offset vector outside the loop and update base with a scalar instruction.

# Improvement Status

- Convert gather loads with invariant stride into strided loads [#147297](#)
  - Detect strided access and introduce the StridedLoadRecipe in LoopVectorize.
  - Contributed by [@Mel-Chen](#) for RISC-V, which has vector strided load/store.

Input IR  
for.body:  
...  
    %idx = mul nuw nsw i64 %iv, 80  
    % gep = getelementptr inbounds nuw i8, ptr %b, i64 %idx  
    %0 = load double, ptr %gep, align 8  
...  
...



LoopVectorize (VPlan)  
vector.body:  
...  
    WIDEN ir<%0> = load ir<%gep>, stride = ir<80>, runtimeVF = vp<%1>  
...  
StridedLoadRecipe

- Improve strided access vectorization for AArch64 SVE [#164205](#)

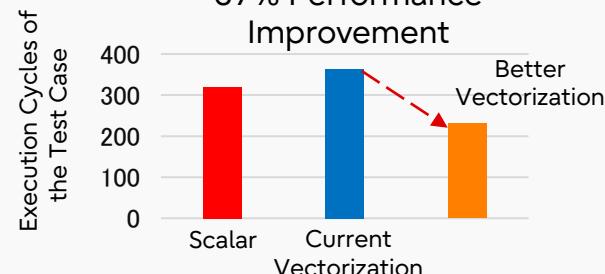
- Legalize the StridedLoadRecipe for architectures that don't have vector strided load/store instructions, such as AArch64.

Legalize StridedLoadRecipe  
vector.ph:  
...  
    EMIT vp<%1> = step-vector i64  
    EMIT vp<%2> = mul vp<%1>, ir<80>  
vector.body:  
...  
    EMIT vp<%3> = wide-ptradd ir<%gep>, vp<%3>  
    WIDEN ir<%0> = load vp<%3>  
...  
...



Better Vectorization

```
index z1.d, #0, #80
loop:
    ld1d { z2.d }, p0/z, [x1, z1.d]
    ...
    add x1, x1, x2
    ...
```



# Other Issues

- We have discovered other issues in vectorizing strided accesses.

- Variable stride widths cannot be vectorized.

- The loop is multi-versioned, and only the version for the  $m=1$  case is vectorized, while the other case ( $m \neq 1$ ) remains scalar.

## Variable Stride

```
for (int i = 0; i < n; i++) {  
    a[i] = b[i * m] + 1;  
}
```

- When IndVarSimplify widens the index of a strided access to 64-bit, the number of memory access instructions can double.
  - Related Issue [#86785](#).

# Acknowledgement



- This presentation is based on results obtained from a project, JPNP21029, subsidized by the New Energy and Industrial Technology Development Organization (NEDO).

# Thank you

