

# PyDSL

A Python Subset for a better MLIR  
programming experience: Updates on  
new features and development

James Cox  jellyfish3.14  jamesthejellyfish

Ritsuka  ritsuka314  Ritsuka314

Kevin Lee  cacophonia  KevinLeeFM

Fabien Zhao  Fabien  fabienzz

Kai-Ting Amy Wang  kai.ting.wang@huawei.com

# Background: What is PyDSL?

- Python-based language that compiles to MLIR.
- @compile decorator turns a Python function into PyDSL and automatically compiles it.
- Pipeline: parses the AST, converts it to MLIR, lowers to shared library binary file.
- Compiled PyDSL functions can be called directly from Python. Arguments are converted to `ctype`, passed to shared library file.
- Has multiple supported backends including CPU and NPU. Can also be passed through existing MLIR pipelines.

```
import numpy as np

from pydsl.frontend import compile
from pydsl.memref import MemRef
from pydsl.scf import range
from pydsl.type import SInt16, SInt32

@compile()
def f(arr: MemRef[SInt32, 8, 4], x: SInt32, y: SInt16) -> SInt32:
    z = x + y
    for i in range(8):
        for j in range(4):
            arr[i, j] += z

    return z

arr = np.zeros((8, 4), dtype=np.int32)
sm = f(arr, 7, 9)
print(arr)
print(sm)
```

# Template Support

- Templates support PyDSL types as well as Python literals.
- Replaces the `@compile()` decorator, adds template descriptors at runtime to the AST's type parameters, and then everything is handled by the standard PyDSL compilation pipeline.
- `@template()` defers compilation until the function is called, creating a binary for each set of descriptors.
- Similar in design to C++ templates.
- Function caching is still working in progress.

*adapted from tests/e2e/test\_template.py*

```
@template()  
def calc[T, N, M](mat: Tensor[T, N, M]) -> Tensor[T, N, M]:  
    n = Index(N)  
    m = Index(M)  
    mat[n - 1, m - 1] = 1  
    return mat  
  
arr = calc[F32, 10, 5](np.zeros((10, 5), dtype=np.float32))  
arr = calc[F64, 1, 40](np.zeros((1, 40), dtype=np.float64))
```

# Autotuning support

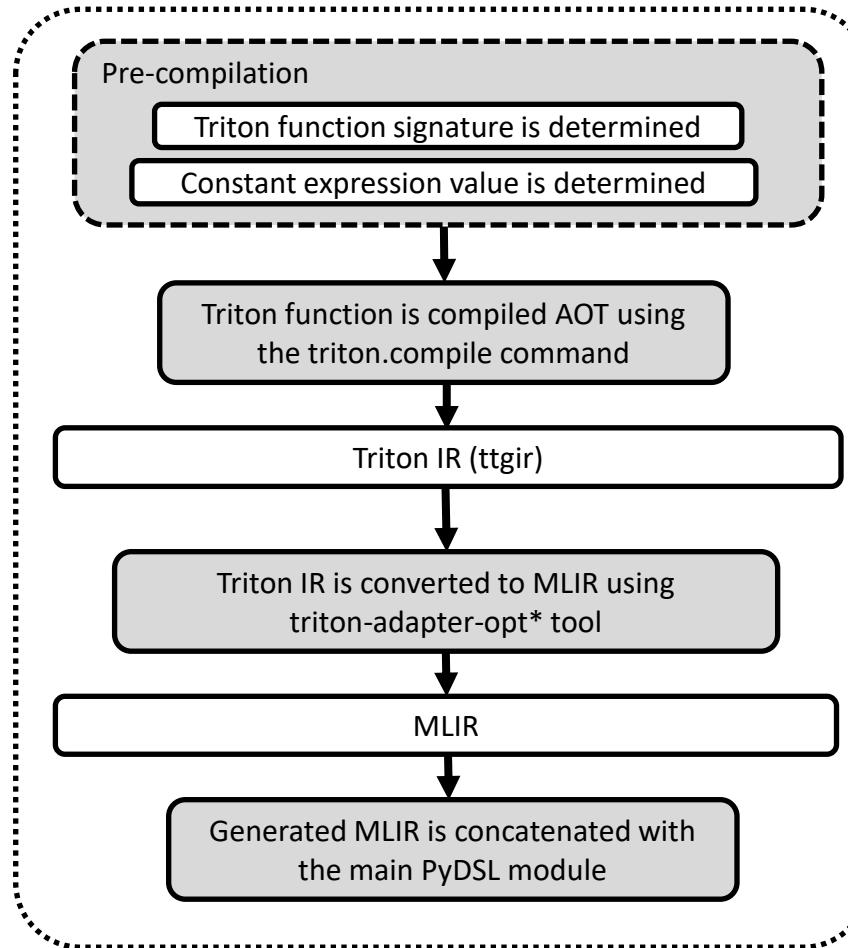
- PyDSL has autotune support for parameters such as tile sizes, types, and anything that is a variable within a kernel.
- Using the @autotune decorator, a user can specify any number of autotune parameters.
- Performs an exhaustive search, compiling each possibility. Therefore it is best for small configuration spaces.

*adapted from examples/autotune/test\_program.py*

```
autotune_values = Var("NUM", [100, 200, 600]) *  
                  TestingData([[N,A]])  
  
@autotune(autotune_values)  
def heat_fuse_tile(n: Index, A: MemRef[F32, N, N]):  
    v = Index(NUM)  
    for _ in arange(S(v)):  
        for i in arange(n):  
            for j in arange(n):  
                # -3/2 x^2 + 5/2 x + 1  
                A[i, j] = (-1.5) * A[i, j] * A[i, j] + (2.5)  
                           * A[i, j] + 1
```

# Inter-operability with Triton

- Triton JITFunction kernels can be directly called by PyDSL functions.
- Currently PyDSL handles tiling, triton kernel is executed as a single thread
- Multi-core interop support is planned for NPU target



\*triton-adapter-opt available in open-source repo <https://gitcode.com/Ascend/triton-ascend>

*from examples/triton\_interop.py:*

```
@triton.jit
def kernel(x_ptr, y_ptr, output_ptr, n_elements,
BLOCK_SIZE):
    pid = tl.program_id(axis=0)
    block_start = pid * BLOCK_SIZE
    offsets = block_start + tl.arange(0, BLOCK_SIZE)
    mask = offsets < n_elements
    x = tl.load(x_ptr + offsets, mask=mask)
    y = tl.load(y_ptr + offsets, mask=mask)
    output = x + y
    tl.store(output_ptr + offsets, output, mask=mask)
```

```
@compile()
def func_test(x: ArrayType, y: ArrayType, out: ArrayType):
    size: Index = 98432
    BLOCK_SIZE = 64
    index_size = size // BLOCK_SIZE
    for i in arange(0, index_size):
        kernel(x, y, out, size,
               BLOCK_SIZE, i)
```

# Improvements Made during the Past Year

1. Memref/Tensor indexing. Uses Python's slicing syntax to implement `memref.subview`, `tensor.extract_slice`, `tensor.insert_slice`
2. Support Memref with strided layout (not shown)
3. Powerful reduce op
  1. Reduce along multiple dimensions
  2. Specify custom combiner function using inline functions
4. CallMacros that make the language extensible
  1. Python functions that evaluate at compile time
  2. Python AST -> MLIR
  3. Most part of PyDSL is implemented via CallMacros
  4. Used to implement Ascend A3 target

---

1. 

```
t1[()] = 456      # t1: Tensor[UInt32]
t2[1] = 5          # t2: Tensor[UInt32, 2]
t2x2[1, 1] = 5    # t2x2: Tensor[UInt32, 2, 2]

t1[2:9:3, 3:6] = t1[4:7, 1:7:2]
```
3. 

```
@InlineFunction.generate()
def sum(a: UInt64, b: UInt32) -> UInt64:
    return a + b

@compile()
def f(
    arr: MemRef[UInt32, DYNAMIC, DYNAMIC, DYNAMIC],
    out: MemRef[UInt64, DYNAMIC],
):
    linalg.reduce(sum, arr, init=out, dims=[0, 2])
```
4. 

```
MemRefF32GM = MemRef.get((16, 8), F32, memory_space=NPU_AddrSpace.GM)
@compile(target_class=AscendTarget)
def f(in: MemRefF32GM, out: MemRefF32GM):
    m = alloc((16, 8), F32, memory_space=NPU_AddrSpace.UB)
    npu.load(in, m)
    npu.vexp(m, out=m)
    npu.store(m, out)

@CallMacro.generate()
def store(visitor: ToMLIRBase, src: Compiled, dst: Compiled):
    # ... Type inferencing
    rep = npu.StoreOp(result_tensor_type, lower_single(src), lower_single(dst))
```

---

# Open source development

- Everything is open-sourced on [GitHub](https://github.com/Huawei-CPLLab/PyDSL) (<https://github.com/Huawei-CPLLab/PyDSL>), with integrated CI and active development
- Everyone is free to contribute! Many issues raised to show proposed features/changes
- Language guide: PyDSL/blob/main/docs/usage.md
- Any questions regarding PyDSL can be asked on the discord: <https://tinyurl.com/PyDSLdiscord>



## CI

All checks have passed  
2 successful checks

✓ AutoTest / Code formatting check (pull\_request) Successful in 13s

✓ AutoTest / Pytest check (pull\_request) Successful in 1m

No conflicts with base branch  
Merging can be performed automatically.

Merge pull request ▾ You can also merge this with the command line. [View](#)

## Issues

Compile-time crash with linalg unary ops on integer tensors

#120 · Ritsuka314 opened 2 weeks ago

Triton interop improvements [feature](#)

#104 · Balint-R opened on Sep 3

Implement linalg.generic [feature](#)

#103 · Balint-R opened on Sep 3

Write language specification [documentation](#) [large](#)

#102 · Balint-R opened on Aug 29

Give linalg elementwise functions a proper definition and signature [refactor](#)

#101 · Balint-R opened on Aug 29

## Pull requests

fix #120: make linalg unary ops reject integer types

#121 opened 2 weeks ago by Ritsuka314

fix #43: turn on lint, code cleanup and minor refactoring

#119 opened 2 weeks ago by Ritsuka314 • Draft

fix #90: add a flag to control linalg elementwise ops cast

#118 opened 3 weeks ago by Ritsuka314

fix #55: support pos only and kw only args

#117 opened last month by Ritsuka314

add the polybench testcases to the test folder [feature](#) [tests](#)

#64 opened on Aug 11 by jamesthejellyfish