

# An investigation of missed opportunities in devirtualization

Yangguang Li, Szymon Sobieszek, Ehsan Amiri

[www.huawei.com](http://www.huawei.com)

# Whole-program devirtualization in LLVM

, typeinfo name

## typeinfo

```
@_ZTI1A = hidden constant { ptr, ptr } { ..... , ptr @_ZTS1A }, align 8
```

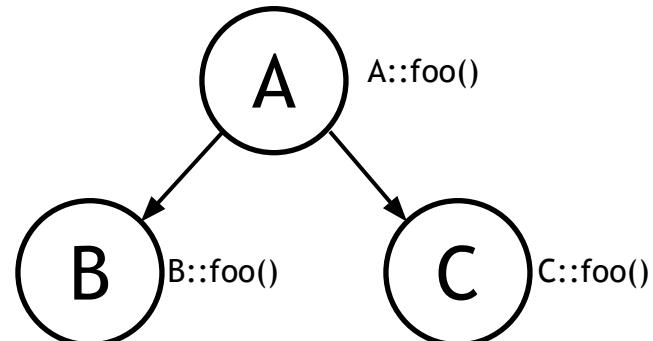
```
@_ZTS1A = hidden constant [3 x i8] c"1A\00", align 1
```

## Virtual table

## Type metadata

Pointer to virtual function A::foo.  
The offset is important

```
 @_ZTV1A = [ ....., ..... , ptr @_ZN1A3fooEv, ..... , ..... ]
 @_ZTV1B = [ ....., ..... , ptr @_ZN1B3fooEv, ..... , ..... ]
 @_ZTV1C = [ ....., ..... , ptr @_ZN1C3fooEv, ..... , ..... ]
```



```
class A {
public:
    virtual void foo() {}
};

class B : public A {
public:
    void foo() {}
};

class C : public A {
public:
    void foo() {}
};
```

WPD checks for every pair  
(type info, vtable offset) whether all  
function pointers in **compatible** vtables  
are the same.

# Whole-program devirtualization in LLVM

, type info name

type info

```
@_ZTI1A = hidden constant { ptr, ptr } { ..... , ptr @_ZTS1A }, align 8
```

```
 @_ZTS1A = hidden constant [3 x i8] c"1A\00", align 1
```

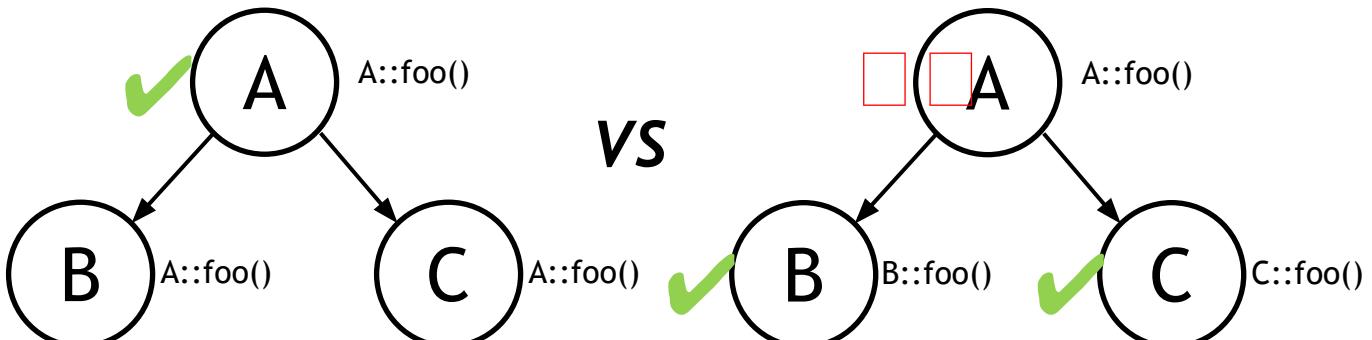
## Virtual table

## Type metadata

Typeinfo name of the class that defines the called function  
C \*c; c->foo() □ typeinfo name of A is used in the metadata

```
%vtable = load ptr, ptr %call, align 8, !tbaa !12271
%1 = tail call { ptr, i1 } @llvm.type.checked.load(ptr %vtable, i32 72, metadata !"__ZTSA"), !nosanitize !12273
%2 = extractvalue { ptr, i1 } %1, 0, !nosanitize !12273
invoke void %2(.....) .....
```

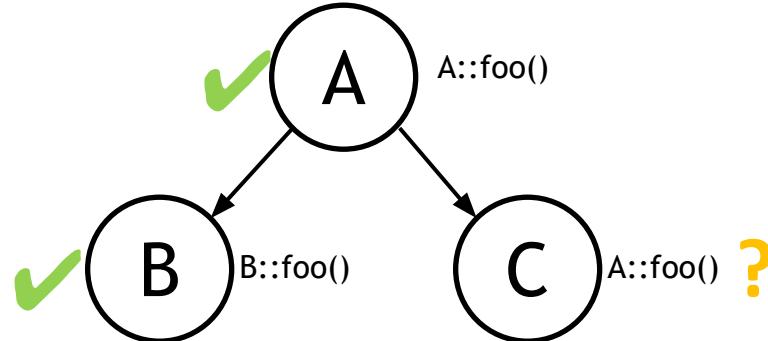
```
@_ZTV1A = [ptr null,
            ptr @_ZTI1A,
            ptr @_ZN1A3fooEv]
@_ZTV1B = [ptr null,
            ptr @_ZTI1B,
            ptr @_ZN1A3fooEv]
@_ZTV1C = [ptr null,
            ptr @_ZTI1C,
            ptr @_ZN1A3fooEv]
```



```
@_ZTV1A = [ptr null,  
           ptr @_ZTI1A,  
           ptr @_ZN1A3fooEv]  
@_ZTV1B = [ptr null,  
           ptr @_ZTI1B,  
           ptr @_ZN1B3fooEv]  
@_ZTV1C = [ptr null,  
           ptr @_ZTI1C,  
           ptr @_ZN1C3fooEv]
```

# What kind of opportunities do we miss?

```
class A {  
public:  
    virtual void foo() {}  
};  
class B : public A {  
public:  
    void foo() {}  
};  
class C : public A {};  
void bar(A *a, B *b, C *c)  
{  
    a->foo();  
    b->foo();  
    c->foo();  
}
```



But can we rely on data type of C? What about non-strict aliasing and type punning.

Devirtualization already assumes data types of objects are reliable. It is easy to construct test cases that have different behavior with and without devirtualization.

Another idea: What if we keep track of type conversions in the source code?

# What other opportunities do we miss?

- Example: We construct an object and pass it to a function as a pointer to a parent class.
- `Json_array`, `Json_object`, and `Json_string` are all leaf classes deriving from `Json_dom`.

```
void DB_restrictions::get_as_json(Json_array &restrictions_array) const {
    for (auto &revocations_itr : m_restrictions) {
        Json_array privileges;
        Json_object revocations_obj;
        Json_string db_name(revocations_itr.first.c_str());
        revocations_obj.add_clone(consts::Database, &db_name);
        ulong revokes_mask = revocations_itr.second;
        while (revokes_mask != 0){
            Json_string priv_str(get_one_priv(revokes_mask));
            privileges.append_clone(&priv_str);
        }
        revocations_obj.add_clone(consts::Privileges, &privileges);
        restrictions_array.append_clone(&revocations_obj);
    }
}
```

```
bool add_clone(const std::string &key, const Json_dom *value) {
    return value == nullptr || add_alias(key, value->clone());
}
```

# How often do these opportunities occur?

## ■ Inlining exposes these devirtualization opportunities

- › Assuming the caller keeps some information about the constructor.
- › But devirtualization exposes inlining opportunity.
- › A simple and quick experiment: Do more devirtualization after LTO inlining and repeat inlining.
- › Number outside parenthesis: All extra devirtualization after keeping track of conversions.
- › Number inside the parentheses: extra devirtualizations when we keep track of constructor calls.

■ Some of the simplest cases can be caught by CSE, if we add a CSE pass (very small fraction)

■ There are multiple ways to generalize/tune this experiment.

Benchmark	Number of devirtualized callsites		Perf gain
	Before LTO inlining	After LTO inlining	
omnetpp_r	39 (14)	194 (92)	-
xalancbmk_r	30 (10)	120 (103)	1%
parest_r	24 (24)	20 (20)	-
MySQL	272 (207)	3312 (3118)	0.4%

# Thank you

[www.huawei.com](http://www.huawei.com)

Copyright©2021 Huawei Technologies Co., Ltd. All Rights Reserved.

The information in this document may contain predictive statements including, without limitation, statements regarding the future financial and operating results, future product portfolio, new technology, etc. There are a number of factors that could cause actual results and developments to differ materially from those expressed or implied in the predictive statements. Therefore, such information is provided for reference purpose only and constitutes neither an offer nor an acceptance. Huawei may change the information at any time without notice.