

ENABLE EDGE AI PRODUCTS YOU DREAM OF

Understanding linalg.pack and linalg.unpack

28 OCTOBER 2025

MAXIMILIAN BARTEL



WHY IS A BLOCKED MATMUL IMPLEMENTATION NOT GETTING THE BEST PERFORMANCE?

Logical layout

$$\begin{array}{c} \text{A} \\ \begin{array}{|c|c|c|c|} \hline 0 & 1 & 2 & 3 \\ \hline 4 & 5 & 6 & 7 \\ \hline 8 & 9 & 10 & 11 \\ \hline 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}
 \times
 \begin{array}{c} \text{B} \\ \begin{array}{|c|c|c|c|} \hline 0 & 1 & 2 & 3 \\ \hline 4 & 5 & 6 & 7 \\ \hline 8 & 9 & 10 & 11 \\ \hline 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}
 =
 \begin{array}{c} \text{C} \\ \begin{array}{|c|c|c|c|} \hline 0 & 1 & 2 & 3 \\ \hline 4 & 5 & 6 & 7 \\ \hline 8 & 9 & 10 & 11 \\ \hline 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}$$

- Fast algorithm, but it generates cache misses:
- Reads from A are good
- Reads from B and C are not

In-memory layout

$$\begin{array}{c} \text{A} \\ \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}
 \times
 \begin{array}{c} \text{B} \\ \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}
 =
 \begin{array}{c} \text{C} \\ \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}$$

$$\begin{array}{c} \text{A} \\ \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}
 \times
 \begin{array}{c} \text{B} \\ \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}
 =
 \begin{array}{c} \text{C} \\ \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ \hline \end{array} \end{array}$$

WE CAN IMPROVE MEMORY ACCESS PATTERNS BY CHANGING THE MEMORY LAYOUT

Logical layout

$$\begin{array}{c}
 \begin{array}{c} A \\ B \end{array} \times \begin{array}{c} B \\ = \end{array} \begin{array}{c} C \end{array}
 \end{array}$$

| | | | |
|----|----|----|----|
| 0 | 1 | 4 | 5 |
| 2 | 3 | 6 | 7 |
| 8 | 9 | 12 | 13 |
| 10 | 11 | 14 | 15 |

| | | | |
|----|----|----|----|
| 0 | 4 | 1 | 5 |
| 8 | 9 | 12 | 13 |
| 2 | 6 | 3 | 7 |
| 10 | 11 | 14 | 15 |

| | | | |
|----|----|----|----|
| 0 | 1 | 4 | 5 |
| 2 | 3 | 6 | 7 |
| 8 | 9 | 12 | 13 |
| 10 | 11 | 14 | 15 |

- Changing the memory layout improves cache hit rate
- Optimal layout aligns perfectly with the reads
- This is called data tiling or packed layouts

In-memory layout

$$\begin{array}{c}
 \begin{array}{c} A \\ B \end{array} \times \begin{array}{c} B \\ = \end{array} \begin{array}{c} C \end{array}
 \end{array}$$

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 0 | 1 | 4 | 5 | 2 | 3 | 6 | 7 | 8 | 9 | 12 | 13 | 10 | 11 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|---|---|---|---|----|----|----|----|
| 0 | 4 | 1 | 5 | 8 | 9 | 12 | 13 | 2 | 6 | 3 | 7 | 10 | 11 | 14 | 15 |
|---|---|---|---|---|---|----|----|---|---|---|---|----|----|----|----|

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 0 | 1 | 4 | 5 | 2 | 3 | 6 | 7 | 8 | 9 | 12 | 13 | 10 | 11 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|

$$\begin{array}{c}
 \begin{array}{c} A \\ B \end{array} \times \begin{array}{c} B \\ = \end{array} \begin{array}{c} C \end{array}
 \end{array}$$

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 0 | 1 | 4 | 5 | 2 | 3 | 6 | 7 | 8 | 9 | 12 | 13 | 10 | 11 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|

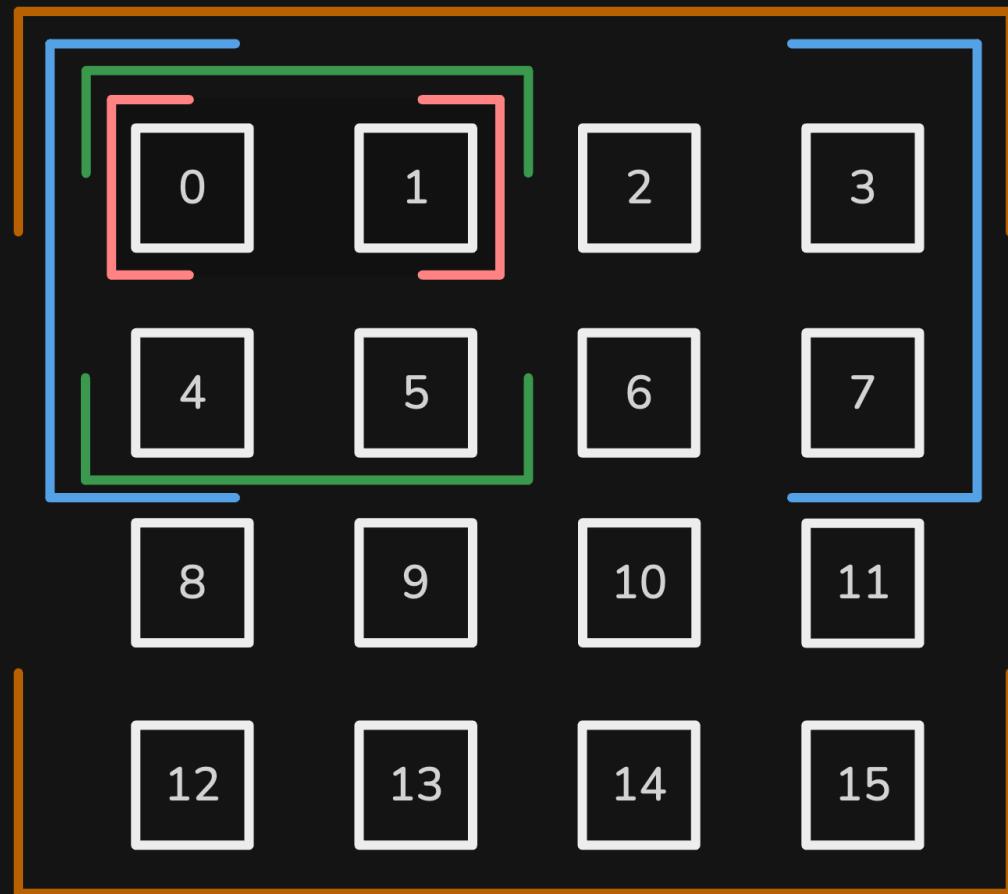
| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|----|----|---|---|---|---|----|----|----|----|
| 0 | 4 | 1 | 5 | 8 | 9 | 12 | 13 | 2 | 6 | 3 | 7 | 10 | 11 | 14 | 15 |
|---|---|---|---|---|---|----|----|---|---|---|---|----|----|----|----|

| | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 0 | 1 | 4 | 5 | 2 | 3 | 6 | 7 | 8 | 9 | 12 | 13 | 10 | 11 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|

PACKING INCREASES THE DIMENSIONALITY OF THE TENSOR

Exemplary 2D matrix

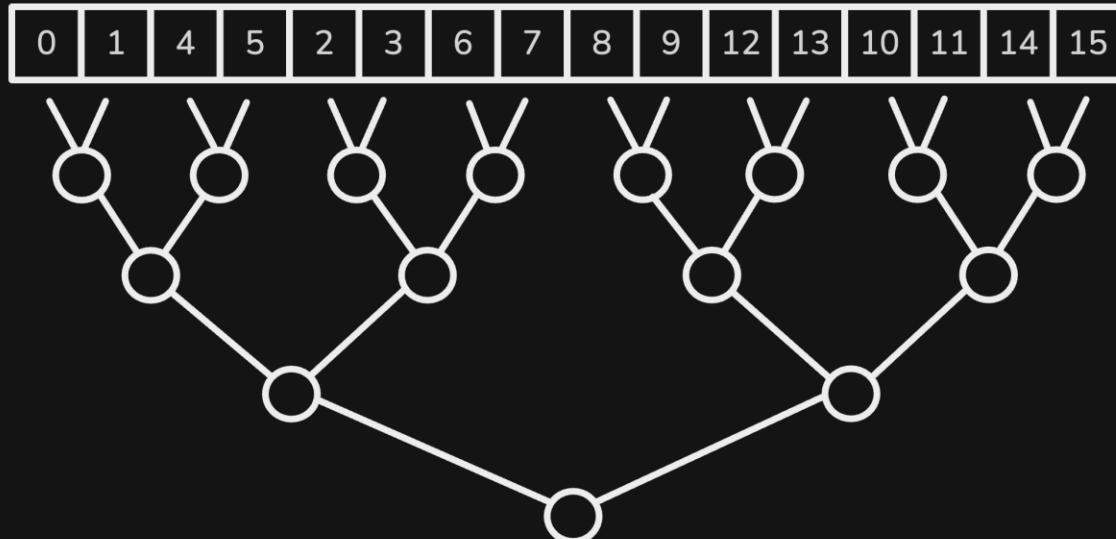
Insights



- Starting from the original logical layout, the new dimensions are marked in different colors
- We go from 4×4 to $2 \times 2 \times 2 \times 2$
- This is so far just a view

PACKING MEANS ACTIVELY FLATTEN THE LAYOUT IN MEMORY

Flattened Tensor



Insights

- Packing takes the view and flattens it
- The reduction tree visualizes the dimensions
- These data movements come at a cost!

EXAMPLE 1

A PACK OP IS ALLOWED TO PAD TO A SPECIFIC TILE SIZE

Exemplary code

```
1 func.func @simple_pad_and_pack_static_tiles(  
2     %input: tensor<3x1xf32>, ←  
3     %output: tensor<1x1x5x2xf32>,  
4     %pad: f32)  
5         -> tensor<1x1x5x2xf32> {  
6     %0 = linalg.pack %input  
7         padding_value(%pad : f32)  
8         inner_dims_pos = [0, 1]  
9         inner_tiles = [5, 2] ←  
10    into %output : tensor<5x1xf32> -> tensor<1x1x5x2xf32>  
11    return %0 : tensor<1x1x5x2xf32>  
12 }
```

Insights



EXAMPLE 2

WE ALSO NEED TO BE ABLE TO REVERT THE PADDING

Exemplary code

```
1 func.func @unpack_as_pad(  
2     %arg0: tensor<1x1x2x3xf32>, %arg1: tensor<1x2xf32>)  
3     -> tensor<1x2xf32> {  
4     %pack = linalg.unpack %arg0  
5         inner_dims_pos = [0, 1]  
6         inner_tiles = [2, 3] ←  
7         into %arg1 : tensor<1x1x2x3xf32> -> tensor<1x2xf32> ←  
8     return %pack : tensor<1x2xf32>  
9 }
```

Insights

| | | | | | |
|---|---|----|----|----|----|
| 0 | 1 | -1 | -1 | -1 | -1 |
|---|---|----|----|----|----|

| | |
|---|---|
| 0 | 1 |
|---|---|

EXAMPLE 3

DYNAMIC SHAPES ARE FULLY SUPPORTED BY THE UNPACK OP

Exemplary code

```
● ● ●  
1 func.func @unpack_fully_dynamic(  
2     %source: tensor<?x?x?x?xf32>, %dest: tensor<?x?xf32>,  
3     %tile_n : index, %tile_m : index)  
4         -> tensor<?x?xf32> {  
5     %0 = linalg.unpack %source  
6         inner_dims_pos = [0, 1]  
7         inner_tiles = [%tile_n, %tile_m]  
8     into %dest : tensor<?x?x?xf32> -> tensor<?x?xf32>  
9     return %0 : tensor<?x?xf32>  
10 }
```

Insights

- Dynamic dimensions are fully supported by this op
- The inner_tiles can take SSA values
- A custom parser treats it as part of an attribute

EXAMPLE 3

UNPACK WILL DECOMPOSE AND INSERT OPERATIONS TO GET THE SHAPES OF THE TENSORS

Exemplary code

```
● ● ●  
1 func.func @unpack_fully_dynamic(  
2     %arg0: tensor<?x?x?xf32>, %arg1: tensor<?x?xf32>,  
3     %arg2: index, %arg3: index)  
4     -> tensor<?x?xf32> {  
5     ... // Get the dimensions of the packed tensor  
6     %0 = tensor.empty(%dim, %dim_1, %dim_0, %dim_2) : tensor<?x?x?xf32>  
7     %transposed = linalg.transpose ... permutation = [0, 2, 1, 3]  
8     %collapsed = tensor.collapse_shape %transposed [[0, 1], [2, 3]]  
9     %dim_3 = tensor.dim %arg1, %c0 : tensor<?x?xf32>  
10    %dim_4 = tensor.dim %arg1, %c1 : tensor<?x?xf32>  
11    %extracted_slice = tensor.extract_slice  
12        %collapsed[0, 0] [%dim_3, %dim_4] [1, 1] : tensor<?x?xf32> to tensor<?x?xf32>  
13    %1 = linalg.copy  
14    return %1 : tensor<?x?xf32>  
15 }
```

Insights

- Decomposing dynamic shapes leads to many dim operations
- Even a padded layout is supported by dynamic shapes
- The output is parsed for its dimensions

EXAMPLE 4

PACK ALSO SUPPORTS DYNAMIC SHAPES – BUT THE LOWERING CANNOT HANDLE IT YET

Exemplary code

```
● ● ●  
1 func.func @pack_fully_dynamic(  
2     %source: tensor<?x?xf32>, %dest: tensor<?x?x?x?xf32>,  
3     %tile_n : index, %tile_m : index, %cst_0 : f32)  
4     -> tensor<?x?x?x?xf32> {  
5     %0 = linalg.pack %source  
6         padding_value(%cst_0 : f32)  
7         inner_dims_pos = [0, 1]  
8         inner_tiles = [%tile_n, %tile_m]  
9     into %dest : tensor<?x?xf32> -> tensor<?x?x?x?xf32>  
10    return %0 : tensor<?x?x?x?xf32>  
11 }
```

Insights

- This is a valid pack operation
- Decomposition should be like unpack
- However, it doesn't lower (yet)
- A lowering can be enabled by inser_slice supporting dynamic shapes

UNIT DIMENSIONS ARE THE MOST COMMON PITFALLS AND PRODUCERS OF BUGS

Exemplary code



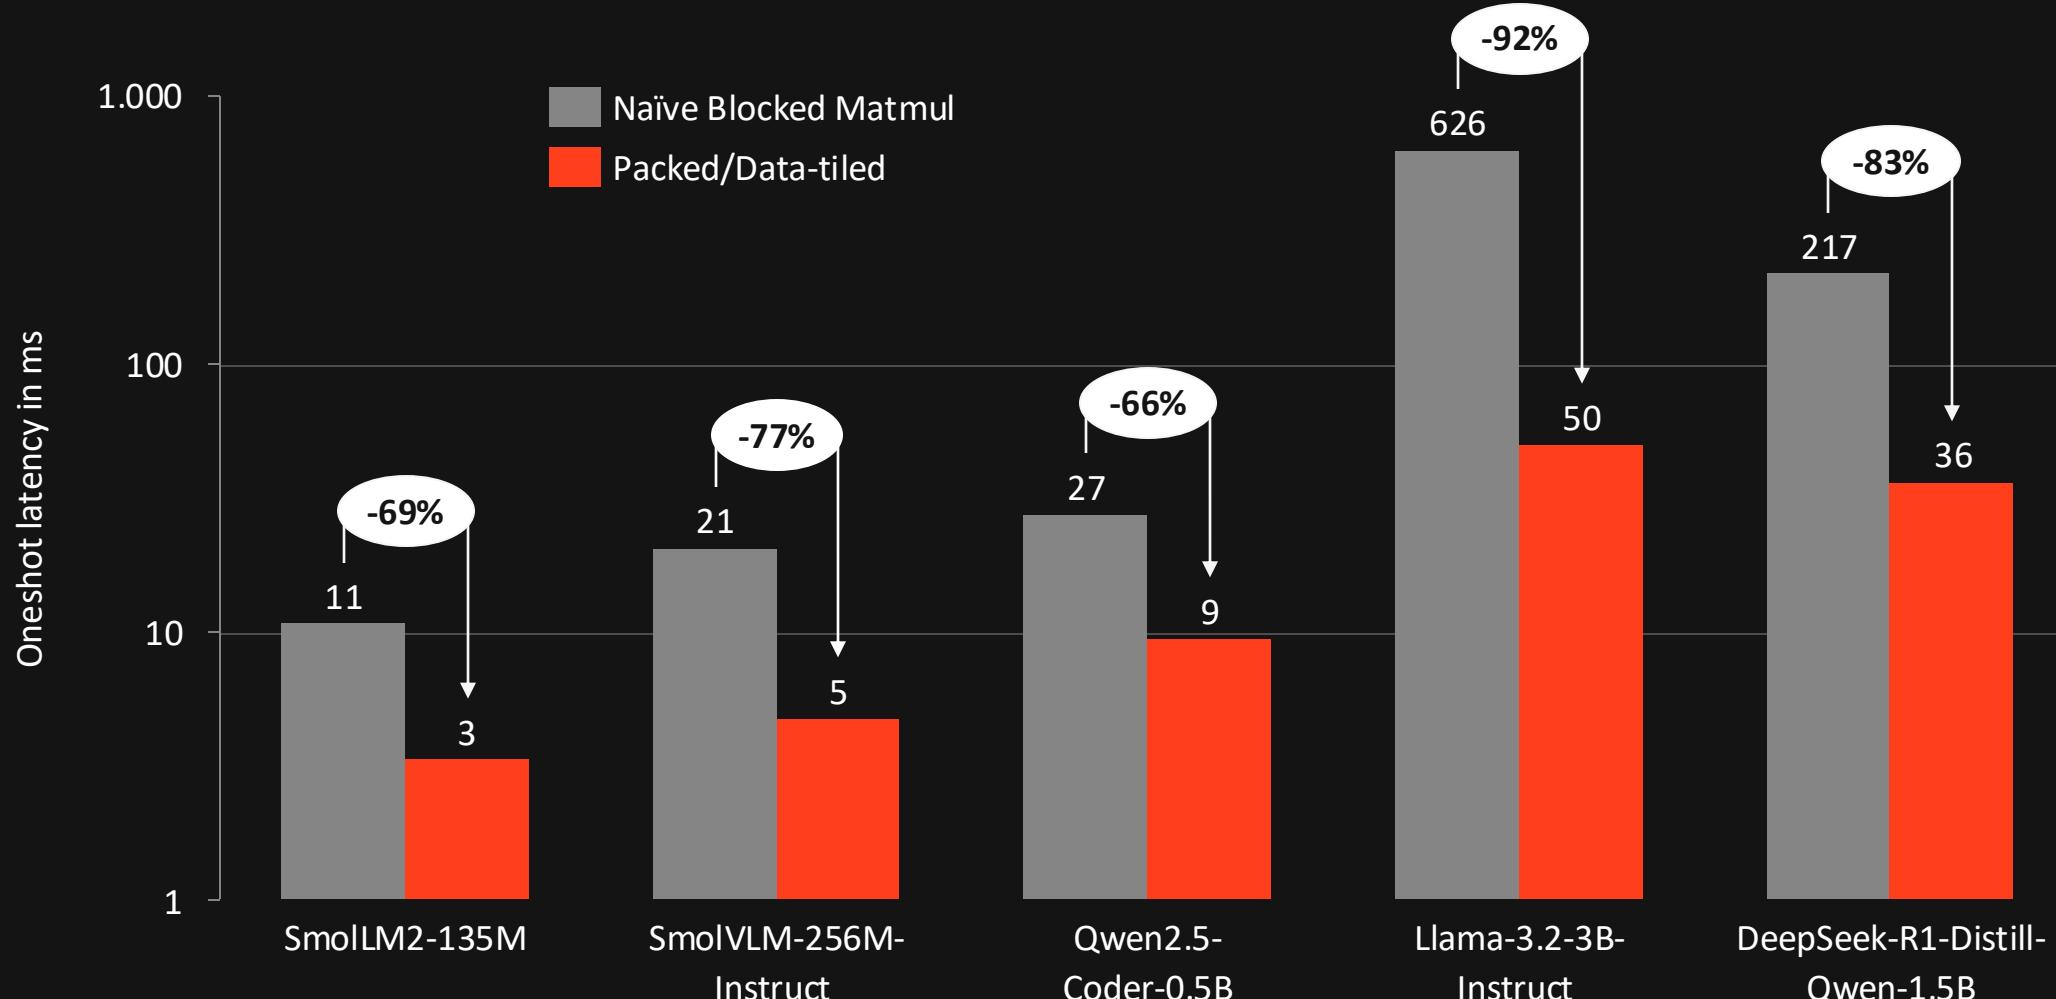
```
1 func.func @unit_dims(
2     %arg0: tensor<1x1x1x4x1xf32>, %arg1: tensor<1x1x4xf32>
3     -> tensor<1x1x1x4x1xf32> {
4     %pack = linalg.pack %arg1
5         outer_dims_perm = [1, 2, 0]
6         inner_dims_pos = [2, 0]
7         inner_tiles = [4, 1]
8         into %arg0 : tensor<1x1x4xf32> -> tensor<1x1x1x4x1xf32>
9     return %pack : tensor<1x1x1x4x1xf32>
10 }
```

Insights

- Most of the tricky behavior arises with unit dimensions
- This behavior is common after tiling
- Special handling is required to produce “better” IR in these cases
- This op eliminates data movements
- Untiled non-unit dim dimension in between unit dims are possible

FULL MODEL BENCHMARKS

WHEN RUNNING FULL MODELS WE SEE LATENCY DECREASES FROM 69% TO 92%



PACK AND UNPACK MAKE MEMORY ACCESSES VERY EFFICIENT, BUT THEY COME AT A COST



How to use pack effectively

- Packing can degrade performance for single kernels
- Only full model compilation can fuse packing operators into producers to hide the movement cost



What are unexplored paths?

- Explore data tiling for other kernels like convolutions
- Explore matrices as representations remove complex attributes (Triton)