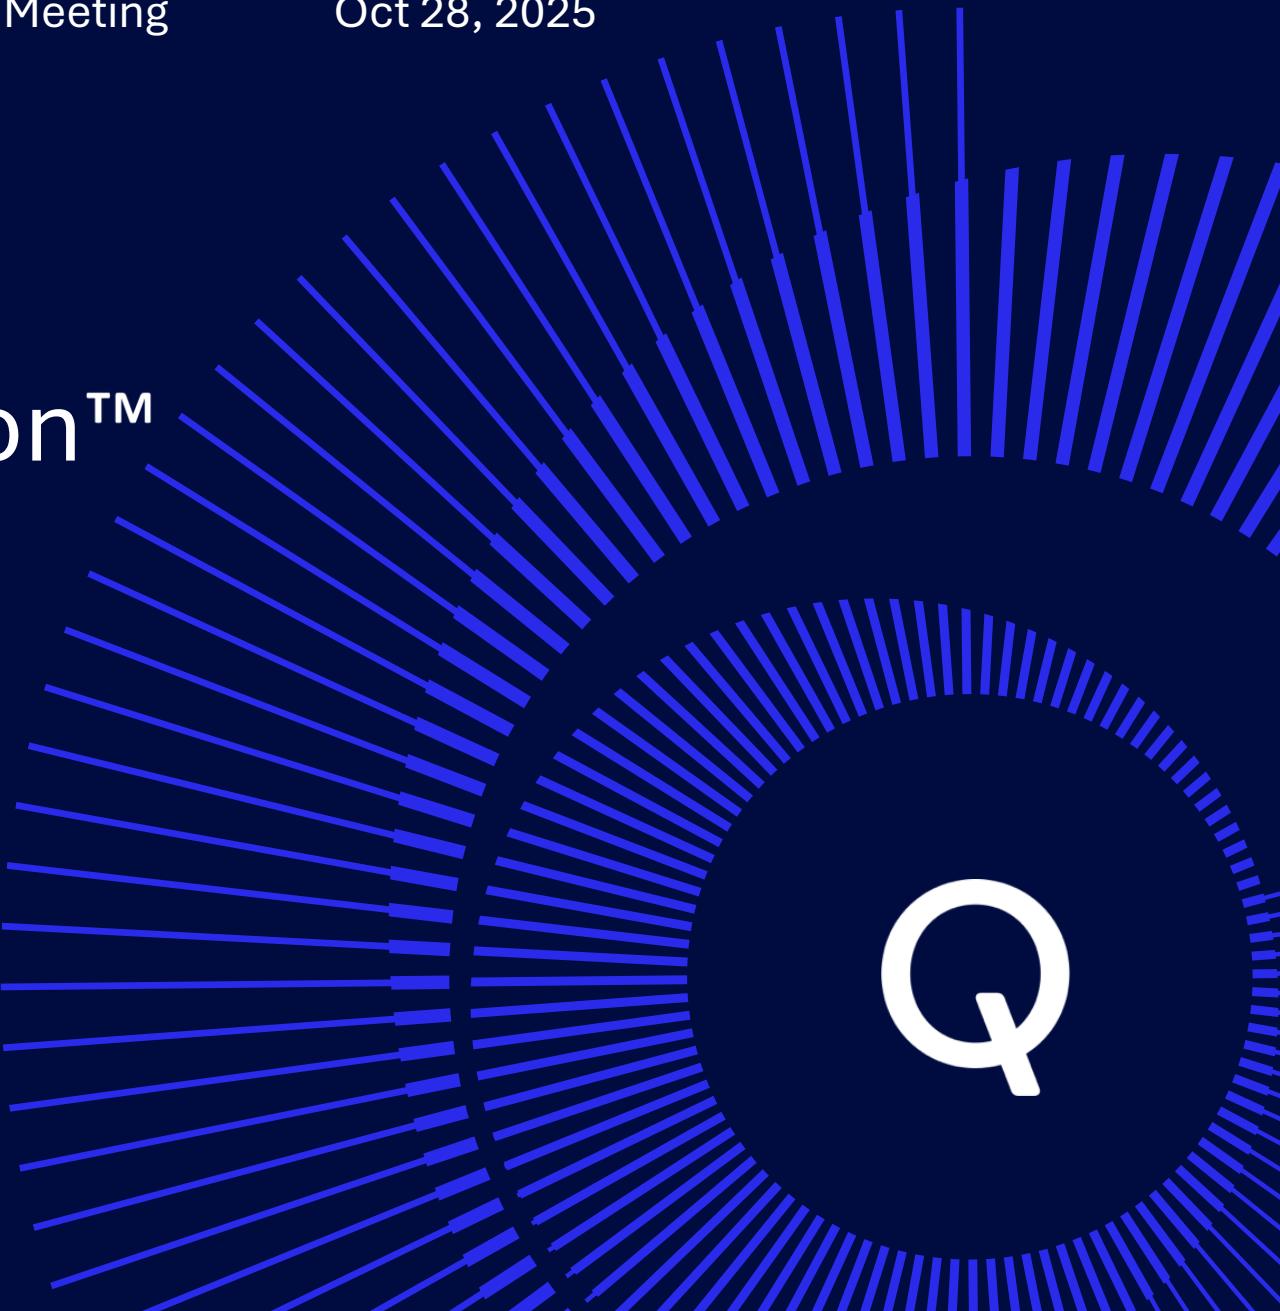


# Accelerating ML on Hexagon™ A Glimpse into Qualcomm's MLIR-based Compiler

Speakers: Franck Slama, Muthu Baskaran

Qualcomm Technologies, Inc.

Authors: Mohammed Javed Absar, Ankit Aggarwal, Iulian Brumar,  
Snigdha Suresh Dalvi, Shalini Jain, Venkat Rasagna Reddy  
Komatireddy, Mitesh Kothari, Samarth Narang, Tasmia Rahman,  
Arun Rangasamy, Abhikrant Sharma, Franck Slama, Jyotsna  
Verma, Zachary Zipper, Muthu Baskaran

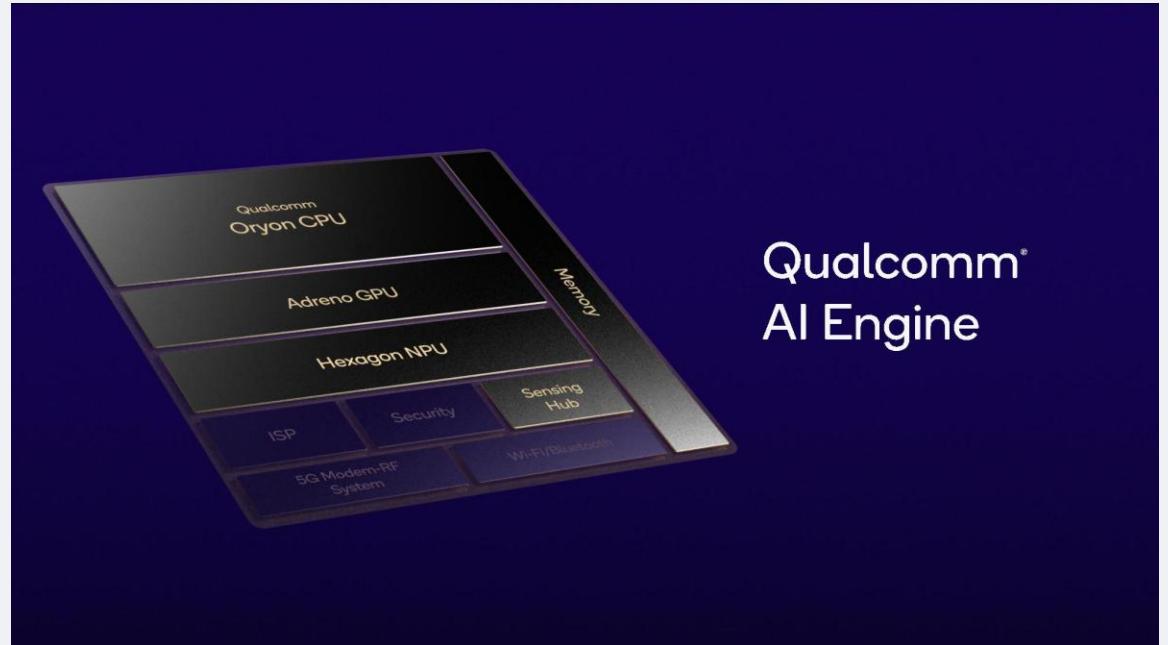


# MLIR Hexagon Compiler

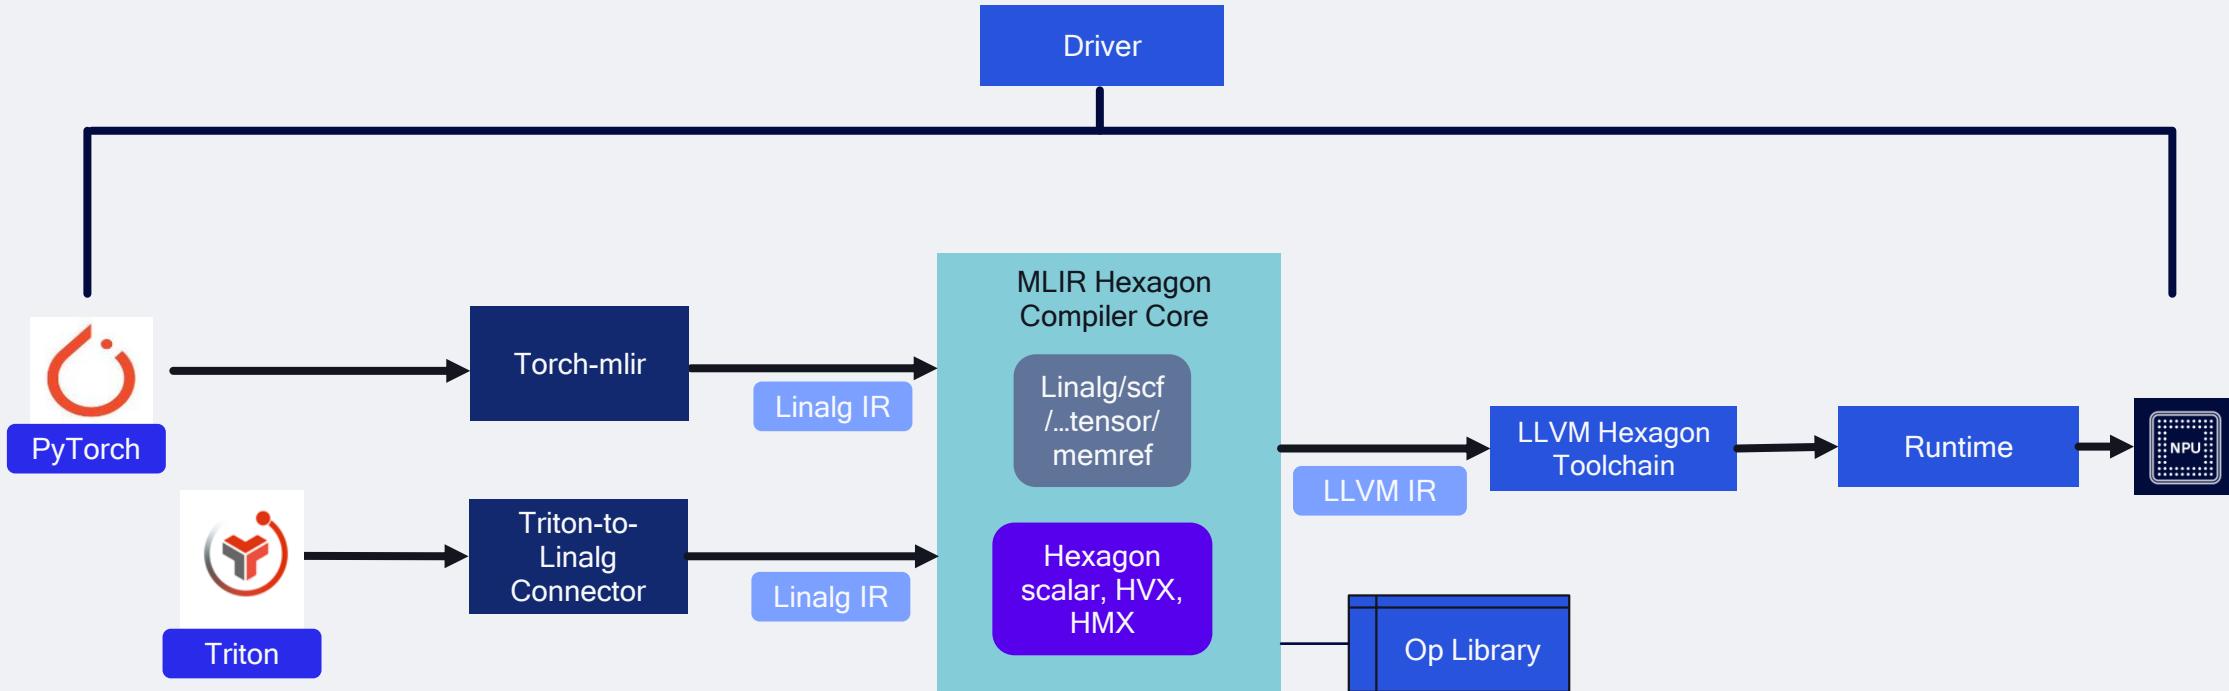
- MLIR-based AI compiler targeting Hexagon NPUs
  - Leverages advancements from both open-source and Qualcomm legacy AI compilers
- Major workstream areas
  - **MLIR Hexagon compiler core:** open-source MLIR infrastructure + Hexagon NPU pass pipeline
  - **PyTorch model compilation** and execution on Hexagon NPU targets
  - **Triton kernel compilation**
- Key results
  - **Functional MLIR Hexagon pipeline w/ developer-friendly Python driver**
  - **Successful demonstration of LLM models** w/ tiling, fusion, vectorization, and memory optimizations
  - **Successful mapping of simple to non-trivial Triton kernels** (flash attention, softmax, argmax, matmul, and many more)
  - Initial performance of Triton kernels mapped via MLIR compiler **up to 80% of hand-written kernels**

# Target Platform

- Hexagon NPUs
  - ❑ 4-way multi-threaded VLIW
  - ❑ Vector registers and dedicated memory
  - ❑ Hexagon Vector eXtensions (HVX)
    - Long vectors (128B) that can process 128 int8 elements or 64 fp16 elements or 32 fp32 elements
  - ❑ Hexagon Matrix eXtensions (HMX)

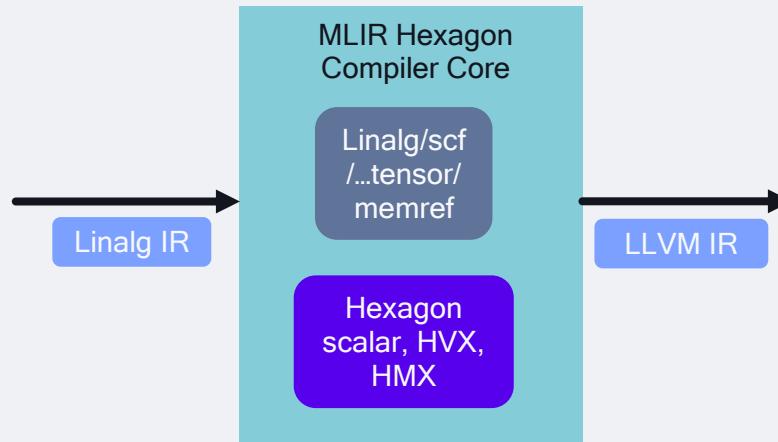


# PyTorch and Triton Workflow via Hexagon NPU Backend



- Approach: Leverage upstream Triton and MLIR developments in addition to building downstream target-specific optimizations

# Hexagon Passes for Lowering and Optimizations



Passes	Description
MLIR core built-in passes	Named to generic, add-fast-math, vector-to-scf, scf-to-cf, vector-to-llvm, math-to-llvm, canonicalize, cse, dce, ...
Fusion	Fuse linalg ops that have producer consumer relation and perform multi-user fusion with re-computation
Tiling	Tiling for scratchpad memory, vectorization, and extracting parallelism
Vectorization	Generate vectorized HVX code
Vectorized math functions	Replace llvm.math.* intrinsics with hand-optimized vectorized math library
Memory optimization passes	Scratchpad management; elimination of redundant creation of memory buffers and associated memory copies
Mapping to tensor units	Lowering to transform data layout and compute for calling tensor unit microkernels

# Illustrative Example - 1

```
%62 = linalg.generic {indexing_maps = [affine_map<(d0,
d1, d2, d3) -> (0, d1, d3)>, affine_map<(d0, d1, d2, d3)
-> (d3)>, affine_map<(d0, d1, d2, d3) -> (d3, d2)>,
affine_map<(d0, d1, d2, d3) -> (d0, d1, d2)>],
iterator_types = ["parallel", "parallel", "parallel",
"reduction"]}
ins(%X, %cst_35, %cst_36 : tensor<?x7x3072xf32>,
tensor<3072xf32>, tensor<3072x768xf32>
outs(%expanded_193 : tensor<1x7x768xf32>) {
^bb0(%in: f32, %in_374: f32, %in_375: f32, %out: f32):
%269 = arith.addf %in, %in_374 : f32
%270 = math.powf %269, %cst_11 : f32
%271 = arith.truncf %cst_4 : f64 to f32
%272 = arith.muldf %270, %271 : f32
%273 = arith.addf %269, %272 : f32
%274 = arith.truncf %cst_3 : f64 to f32
%275 = arith.muldf %273, %274 : f32
%276 = math.tanh %275 : f32
%277 = arith.addf %276, %cst_12 : f32
%278 = arith.muldf %269, %cst_10 : f32
%279 = arith.muldf %278, %277 : f32
%280 = arith.muldf %279, %in_375 : f32
%281 = arith.addf %out, %280 : f32
linalg.yield %281 : f32
} -> tensor<?x7x768xf32>
```

Fused computation  
(fusion done on 'linalg graph')

```
scf.for %i= %c0 to %cI step %istep{
  scf.for %j= %c0 to %cJ step %jstep {
    %subview = memref.subview %X[0, %i, %j]
    ... : memref<1x7x3072xf32> to memref<3072xf32,
    strided<[1], offset: ?>>
    %s_memory = hexagonmem.alloc() :
    memref<3072xf32>

    hexagonmem.copy %subview to %s_memory :
    memref<3072xf32, strided<[1], offset: ?>> to
    memref<3072xf32>

    scf.for %k= %c0 to %xK step %kstep {
      ...
      %153 = arith.addf %149, %150
      fastmath<fast> : vector<32xf32>
      %154 = math.powf %153, %cst_12
      fastmath<fast> : vector<32xf32>
      ...
    }
    hexagonmem.copy %s_memory to %subview :
    memref<3072xf32> to memref<3072xf32,
    strided<[1], offset: ?>>
    hexagonmem.dealloc %s_memory :
    memref<3072xf32>
  }
}
```

Tiling, Vectorization,  
Scratchpad memory  
alloc/dealloc and data  
copy

```
scf.for %i= %c0 to %cI step %istep{
  scf.for %j= %c0 to %cJ step %jstep {
    %subview = memref.subview %X[0, %i, %j]
    ... : memref<1x7x3072xf32> to memref<3072xf32,
    strided<[1], offset: ?>>
    %s_memory = hexagonmem.alloc() :
    memref<3072xf32>

    memref.dma_start %subview[...] %
    s_memory [...]
    memref.dma_wait ...

    scf.for %k= %c0 to %xK step %kstep {
      ...
      %153 = arith.addf %149, %150
      fastmath<fast> : vector<32xf32>
      %154 = math.powf %153, %cst_12
      fastmath<fast> : vector<32xf32>
      ...
    }
    memref.dma_start %s_memory[...]
    %subview[...]
    memref.dma_wait ...
    hexagonmem.dealloc %s_memory :
    memref<3072xf32>
  }
}
```

Tiling, Vectorization,  
Scratchpad memory  
alloc/dealloc and data  
copy w/ DMA

# Illustrative Example - 2

## Flash Attention core in Triton

```
@triton.jit
def
_attn_fwd_inner(acc, l_i, m_i, q, K_block_ptr, V_block_ptr, start_m, qk_scale,
    BLOCK_M: tl.constexpr, BLOCK_DMODEL: tl.constexpr, BLOCK_N:
    tl.constexpr,
    STAGE: tl.constexpr, offs_m: tl.constexpr, offs_n: tl.constexpr,
    N_CTX: tl.constexpr):
    lo, hi = 0, N_CTX
    K_block_ptr = tl.advance(K_block_ptr, (lo, 0))
    V_block_ptr = tl.advance(V_block_ptr, (lo, 0))
    # loop over k, v and update accumulator
    for start_n in range(lo, hi, BLOCK_N):
        start_n = tl.multiple_of(start_n, BLOCK_N)
        # -- compute qk ---
        k = tl.load(K_block_ptr)
        qk = tl.dot(q, tl.trans(k))
        m_ij = tl.maximum(m_i, tl.max(qk, 1) * qk_scale)
        qk = qk * qk_scale - m_ij[:, None]
        p = tl.math.exp2(qk)
        l_ij = tl.sum(p, 1)
        # -- update m_i and l_i
        alpha = tl.math.exp2(m_i - m_ij)
        l_i = l_i * alpha + l_ij
        # -- update output accumulator --
        acc = acc * alpha[:, None]
        # update acc
        v = tl.load(V_block_ptr)
        # p = p.to(tl.float32)
        acc = tl.dot(p, v, acc)
        # update m_i and l_i
        m_i = m_ij
        V_block_ptr = tl.advance(V_block_ptr, (BLOCK_N, 0))
        K_block_ptr = tl.advance(K_block_ptr, (BLOCK_N, 0))
    return acc, l_i, m_i
```

## Parallelism, fusion

```
%42 = linalg.generic {indexing_maps = [#map3, #map6, #map3],
    iterator_types = ["parallel", "parallel"]} ins(%37,
    %expanded_8 : tensor<256x64xf32>, tensor<256x1xf32>) outs(%3
    : tensor<256x64xf32> {
    ^bb0(%in: f32, %in_11: f32, %out: f32):
        %53 = arith.mulf %in, %29 : f32
        %54 = arith.subf %53, %in_11 : f32
        %55 = math.exp2 %54 : f32
        linalg.yield %55 : f32
    } -> tensor<256x64xf32>
```

## Vectorization

```
scf.for %arg38 = %c0 to %c256 step %c1 {
    scf.for %arg39 = %c0 to %c64 step %c32 {
        %subview = memref.subview %alloc_17[%arg38, %arg39] [1, 32] [1, 1] :
        memref<256x64xf32> to memref<32xf32, strided<[1], offset: ?>
        %subview_29 = memref.subview %expand_shape_20[%arg38, 0] [1, 1] [1, 1]
        : memref<256x1xf32> to memref<f32, strided<[], offset: ?>
        %subview_30 = memref.subview %alloc_7[%arg38, %arg39] [1, 32] [1, 1] :
        memref<256x64xf32> to memref<32xf32, strided<[1], offset: ?>
        %32 = vector.transfer_read %subview[%c0], %cst_4 {in_bounds = [true]} :
        memref<32xf32, strided<[1], offset: ?>, vector<32xf32>
        %33 = memref.load %subview_29[] : memref<f32, strided<[], offset: ?>
        %34 = vector.broadcast %33 : f32 to vector<32xf32>
        %35 = vector.broadcast %24 : f32 to vector<32xf32>
        %36 = arith.mulf %32, %35 fastmath<fast> : vector<32xf32>
        %37 = arith.subf %36, %34 fastmath<fast> : vector<32xf32>
        %38 = math.exp2 %37 fastmath<fast> : vector<32xf32>
        vector.transfer_write %38, %subview_30[%c0] {in_bounds = [true]} :
        vector<32xf32>, memref<32xf32, strided<[1], offset: ?>
    }
}
```

## Call to hand-optimized library

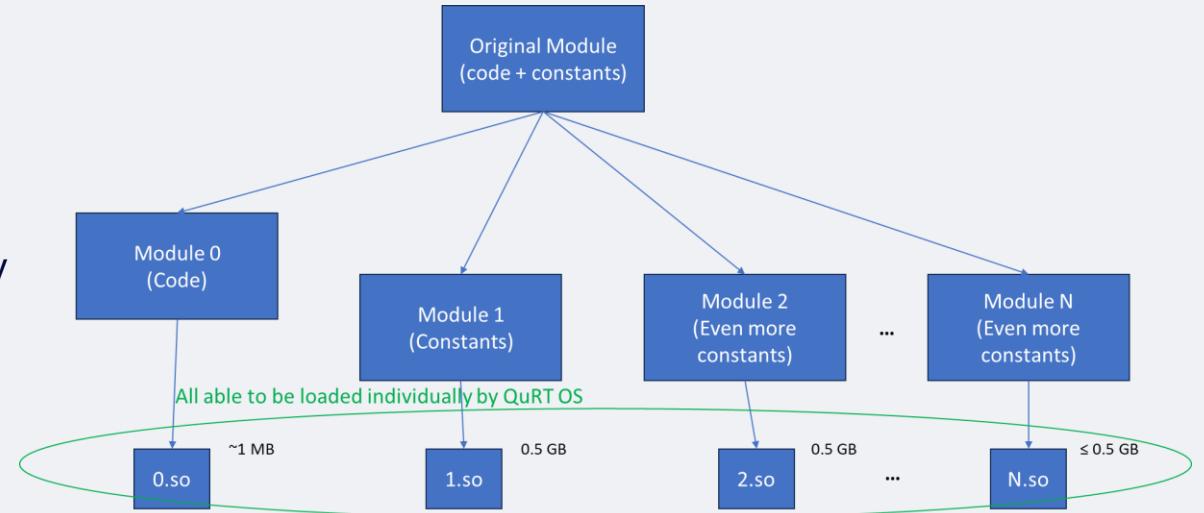
```
%525 = llvm.intr.vector.extract %523[0] : vector<32xf32> from vector<64xf32>
%526 = llvm.call @_hexagon.exp.vsf.128B(%525) : (vector<32xf32>) -> vector<32xf32>
%527 = llvm.intr.vector.insert %526, %524[0] : vector<32xf32> into vector<64xf32>
```

# Scalability: An MLIR pass to lower the constants separately

- For large models, we want to split the module into various parts, loaded independently by the OS
  - Since 99% of the memory footprint is due to constant parameters (weights, biases, factors) → need to lower separately the constants

- Module splitting implemented at the MLIR level

- New MLIR pass that separates out constants:  
1 module in, N module out
- Creation of new MLIR modules (constants-only) on the fly
- In the main module (for code):  
Replace the *arith.constant* by an extern



- Compiling (and linking) all these modules appropriately (the principal module needs to be linked against all the constants-only modules, as it uses their constants)
- Result: Ability to lower the constants separately, helps to run large models on device, on Qualcomm Hexagon NPU

# Conclusion

- MLIR-based AI compiler targeting Hexagon NPUs
  - Leverages advancements from both open-source and Qualcomm legacy AI compilers
- Demonstrated mapping of PyTorch models and Triton kernels via the compiler on to Hexagon NPUs
  - Can generate code for a variety of ML models. Currently improving the efficiency of the code being produced.
- Current capabilities include support for single NPU configurations
  - Next: Support for multi-NPU configurations
- Planning to release/open-source our compiler in December 2025

# Thank you

Nothing in these materials is an offer to sell any of the components or devices referenced herein.

© Qualcomm Technologies, Inc. and/or its affiliated companies. All Rights Reserved.

Qualcomm, Snapdragon, Hexagon, Qualcomm Spectra, Adreno, FastConnect, and Kryo are trademarks or registered trademarks of Qualcomm Incorporated.

Other products and brand names may be trademarks or registered trademarks of their respective owners.

References in this presentation to “Qualcomm” may mean Qualcomm Incorporated, Qualcomm Technologies, Inc., and/or other subsidiaries or business units within the Qualcomm corporate structure, as applicable. Qualcomm Incorporated includes our licensing business, QTL, and the vast majority of our patent portfolio. Qualcomm Technologies, Inc., a subsidiary of Qualcomm Incorporated, operates, along with its subsidiaries, substantially all of our engineering, research and development functions, and substantially all of our products and services businesses, including our QCT semiconductor business.

Snapdragon and Qualcomm branded products are products of Qualcomm Technologies, Inc. and/or its subsidiaries. Qualcomm patented technologies are licensed by Qualcomm Incorporated.

Follow us on: [in](#) [X](#) [@](#) [YouTube](#) [f](#)

For more information, visit us at [qualcomm.com](http://qualcomm.com) & [qualcomm.com/blog](http://qualcomm.com/blog)

