


# Towards Automatic Reduction of Module Bugs

Improving C-Vise

Maksim Ivanov, Google

# The Problem: Test-Case Reduction

- **Goal:** Reduce a large, failing program to a minimal, self-contained reproducer.
- **Why:** Critical for debugging compiler issues (speeds up debugging), reporting downstream issues, creating focused regression tests.



- **Challenge:** Existing tools (C-Reduce, C-Vise) are often slow. Reductions can take hours, days, or even weeks on large, real-world test cases.
- **Core Limitation:** They are fundamentally designed for single-file inputs. Multiple files are only reduced separately and the same number of files remains.

# The Specific Challenge: C++ Header Modules

Test-case reduction for C++ modules is significantly harder. A reproducer is not one file; it's a complex bundle:

- **Multiple Files:** Can involve thousands of source files, headers, and .cppmap files.
- **Multiple Commands:** Requires multiple, ordered compilation commands (to build PCMs) that must **succeed** before the final, **failing** command is run.
- **Complex Dependencies:** A bug may only manifest when a specific PCM is built and imported in a specific way.

# Our Approach: Evolving C-Vise

We chose to improve C-Vise (a Python-based C-Reduce successor) to tackle this.

- **Goal 1: Speed**

Radically improve core reduction performance, even on single-file inputs.

- **Goal 2: Versatility**

Extend the tool to natively support multi-file, multi-command reductions for C++ modules.

- **Key Insight:** Achieving Goal 1 with a new architecture directly unblocked our path to achieving Goal 2.

# Core Improvement: Hint-Based Architecture

## Old C-Vise: In-Place Modification

Heuristics (e.g., "remove function body") directly modified files. This was rigid, sequential, and hard to parallelize effectively.

```
[Heuristic] ---> [Modifies File In-Place]
```

## New C-Vise: Hint Emission

Heuristics emit "hints" (JSONs describing patches). A generic, parallel scheduler collects hints and decides how to test them.

```
[Heuristic] ---> [JSON "Hint"] --->  
[Scheduler]
```

This decouples "what to try" from "how to try it" (+binary search, etc.).

# Performance Win 1: Interleaved Execution

## Old: Sequential Passes

Runs one full heuristic to completion before starting the next. It can get "stuck" on a low-yield pass for hours.

```
Run 'lines' pass (2 hours)  
...stall...  
Run 'remove-function' pass (1 hour)  
...
```

## New: Interleaved Hints

Mixes hints from **all** heuristics in a round-robin fashion. The scheduler constantly makes progress using the best reduction available from any pass.

```
Try 'line' hint  
Try 'function' hint  
Try 'comment' hint  
Try 'line' hint (SUCCESS)  
...
```

# Performance Win 2: Folding Reductions

## Old: Wasted Parallelism

If 5 parallel jobs find 5 different successful reductions, the tool picks **one**, discards the other 4, and restarts all workers.

```
Job 1: Success (Remove line 10)
Job 2: Success (Remove line 20)
Job 3: Fail
Job 4: Success (Remove line 30)
=> Keep Job 1, Discard 2 & 4.
```

## New: Folded Reductions

We "fold" all successful, non-conflicting hints from a batch of parallel jobs into a **single** combined patch and test that.

```
Job 1: Success (Remove line 10)
Job 2: Success (Remove line 20)
Job 3: Fail
Job 4: Success (Remove line 30)
=> Fold [1, 2, 4] -> Test 1 patch.
```

This achieves a massive reduction in one step.

## Miscellaneous improvements

- Robustness:
  - hung child process termination;
  - temporary files leaks;
- Performance:
  - improved ad-hoc parsers;
  - new heuristics based on Tree-sitter parsers.

# Performance Results (Single-File)

These architectural changes resulted in a ~10x-80x speedup on our benchmark suite.

| Test Case       | C-Reduce 2.11.0 | C-Vise 2.11.0 (Old) | C-Vise (New)     | Speedup    |
|-----------------|-----------------|---------------------|------------------|------------|
| clang-363816643 | 15 hours        | 16 hours            | <b>12 min</b>    | <b>75x</b> |
| clang-383027690 | 18 min          | 2.5 hours           | <b>2 min</b>     | <b>9x</b>  |
| clang-321217557 | ?               | 30 days (est.)      | <b>8.5 hours</b> | <b>85x</b> |
| clang-329180703 | $\infty$ (hung) | 44 hours            | <b>45 min</b>    | <b>60x</b> |
| clang-410818184 | 31 hours        | 85 hours            | <b>1.5 hours</b> | <b>20x</b> |
| gcc-94937       | 24 hours        | $\infty$ (hung)     | <b>40 min</b>    | <b>35x</b> |
| gcc-92516       | 4.1 hours       | 3.2 hours           | <b>15 min</b>    | <b>13x</b> |

---

# Applying to C++ Header Modules

# Streamlined multi-file handling

## Before

- Files were processed sequentially by each pass.
- Binary search was limited to instances in a single file.
- Number of files remained constant.

## After


- Passes now operate simultaneously on all files.
- Binary search operates across file boundaries.
- Detecting file cross-references, deleting unused files/dirs.

# Compilation command reductions

Our approach - use Makefiles:


- stores compilation commands and dependencies between them;
- can also be executed by the interestingness test.

New heuristics: cmd parameter removal, target removal.



# Module map reductions


New heuristics for structured removal of contents from module map files.



# Detecting file references

Goal: new passes to delete unused files; to attempt deleting a file with all references.

How: Run the Clang preprocessor to build the graph of #include's.



# Performance results (Header Modules)

| <b>Test Case</b> | <b>input size</b> | <b>duration</b> | <b>output size</b> |
|------------------|-------------------|-----------------|--------------------|
| clang-355835505  | 37 MB, 2638 files | 2.5 hours       | 47 KB, 13 files    |

# Summary

## Done

- Re-architected C-Vise with "hints" for flexibility and parallelism, achieving **10x-80x speedup** on single-file tests.
- Efficient multi-file reduction.
- Header module aware passes.

## Future Work

- C++20 modules support.
- More heuristics.
- Improving parallelism bottlenecks.
- Reduction in the cloud.
- LLM-based heuristics and drivers.