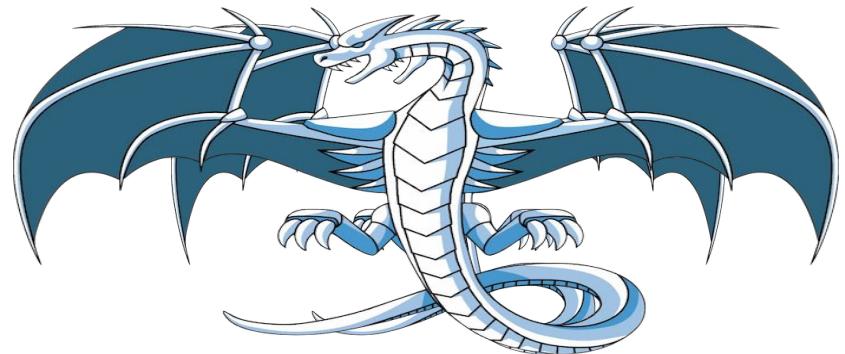


# Widen Your Char-izons

Adding wide character conversion to LLVM-libc

By: Uzair Nawaz and Sriya Pratipati



01

# Introduction



# How Do We Represent Characters/Strings?

- Strings are just arrays of bytes (`char` is 1 byte)
- Super convenient for 1-byte characters, but not straightforward when representing more complex characters (emojis, other languages, etc)

|      |      |      |      |
|------|------|------|------|
| L    | L    | V    | M    |
| 0x4C | 0x4C | 0x56 | 0x4D |

# Multibyte vs Wide Characters

- Multibyte-Character Strings
  - Characters vary in size between 1 to 4 bytes
  - Length of a string in bytes  $\neq$  # of characters in the string
  - Referenced by a char \* so possible to stop in the middle of a character
  - Typically represented by UTF-8 encoding
- Wide-Character Strings
  - Every character takes up the same number of bytes (usually 4 on most systems)
  - Easy to calculate length of string
  - Can't stop in the middle of a wide character
  - Represented by UTF-32 encoding on most systems



Multibyte representation (UTF-8) - 6 bytes

|      |      |      |          |      |      |
|------|------|------|----------|------|------|
|      |      |      | $\Sigma$ |      |      |
| 0xF0 | 0x9F | 0xA4 | 0xAF     | 0xCE | 0xA3 |

Wide character representation (UTF-32) - 8 bytes

|            |            |
|------------|------------|
|            | $\Sigma$   |
| 0x0001F921 | 0x000003A3 |

02

# Conversion Process



# UTF-8 Encoding Details

Code point → UTF-8 conversion

| First code point | Last code point | Byte 1   | Byte 2   | Byte 3   | Byte 4   |
|------------------|-----------------|----------|----------|----------|----------|
| U+0000           | U+007F          | 0yyyyyyy |          |          |          |
| U+0080           | U+07FF          | 110xxxxy | 10yyzzzz |          |          |
| U+0800           | U+FFFF          | 1110wwww | 10xxxxxy | 10yyzzzz |          |
| U+010000         | U+10FFFF        | 11110uvv | 10vvwwww | 10xxxxxy | 10yyzzzz |

# Multibyte -> Wide Character Example

Input Multibyte String:

|      |      |      |
|------|------|------|
| A    | Σ    |      |
| 0x65 | 0xCE | 0xA3 |



Partial State

|  |  |  |  |
|--|--|--|--|
|  |  |  |  |
|  |  |  |  |

Output Wide Character String:

|  |  |
|--|--|
|  |  |
|  |  |



# Multibyte -> Wide Character Example

Input Multibyte String:

|      |      |      |
|------|------|------|
| A    | Σ    |      |
| 0x65 | 0xCE | 0xA3 |



Partial State

|      |  |  |  |
|------|--|--|--|
| A    |  |  |  |
| 0x65 |  |  |  |

Output Wide Character String:

|  |  |
|--|--|
|  |  |
|  |  |



# Multibyte -> Wide Character Example

Input Multibyte String:

|      |      |      |
|------|------|------|
| A    | Σ    |      |
| 0x65 | 0xCE | 0xA3 |



Partial State

|  |  |  |  |
|--|--|--|--|
|  |  |  |  |
|  |  |  |  |

Output Wide Character String:

|            |  |
|------------|--|
| A          |  |
| 0x00000065 |  |



# Multibyte -> Wide Character Example

Input Multibyte String:

|      |      |      |
|------|------|------|
| A    | Σ    |      |
| 0x65 | 0xCE | 0xA3 |



Partial State

|      |  |  |
|------|--|--|
| Σ    |  |  |
| 0xCE |  |  |

Output Wide Character String:

|            |  |
|------------|--|
| A          |  |
| 0x00000065 |  |



# Multibyte -> Wide Character Example

Input Multibyte String:

|      |      |      |
|------|------|------|
| A    | Σ    |      |
| 0x65 | 0xCE | 0xA3 |

Partial State

|      |      |  |  |
|------|------|--|--|
| Σ    |      |  |  |
| 0xCE | 0xA3 |  |  |



Output Wide Character String:

|            |  |
|------------|--|
| A          |  |
| 0x00000065 |  |



# Multibyte -> Wide Character Example

Input Multibyte String:

|      |      |      |
|------|------|------|
| A    | Σ    |      |
| 0x65 | 0xCE | 0xA3 |

Partial State

|  |  |  |  |
|--|--|--|--|
|  |  |  |  |
|  |  |  |  |



Output Wide Character String:

|            |            |
|------------|------------|
| A          | Σ          |
| 0x00000065 | 0x000003A3 |



03

# Libc Interface



# Example of mbtowc use

```
const char* mb_str = "\ud83d\udcbb";
wchar_t wc_string[1];
mbstate_t mbs;
size_t ret = mbtowc(wc_string, &mb_str, /* max # of bytes to read */ 1,
&mbs);

ASSERT(ret == -2);
```

# Libc Interface

```
const char* mb_str = "\ud83d\udcbb";
wchar_t wc_string[1];
mbstate_t mbs;
size_t ret = mbrtowc(wc_string, &mb_str, /* max # of bytes to read */ 1,
&mbs);

ASSERT(ret == -2);

ret = mbrtowc(wc_string, &mb_str + 1, /* max # of bytes to read */ 3, &mbs);
ASSERT(ret == 3);

ASSERT(wc_string[0] == 0x0001F921);
```

# Restartable vs Non-Restartable

- Restartable
  - Takes in an input of an mbstate, can stop conversion mid-character and pick up where it left off
- Non-Restartable
  - Has its own internal state that is maintained globally on each call to the function

04

# Architecture



# `mbstate_t`

- Represents a partial conversion state
- Layout:

| Field                            | Size    |
|----------------------------------|---------|
| Partial State as UTF-32          | 32 bits |
| # bytes stored in partial state  | 8 bits  |
| # of total bytes in mb-character | 8 bits  |

# CharacterConverter Class

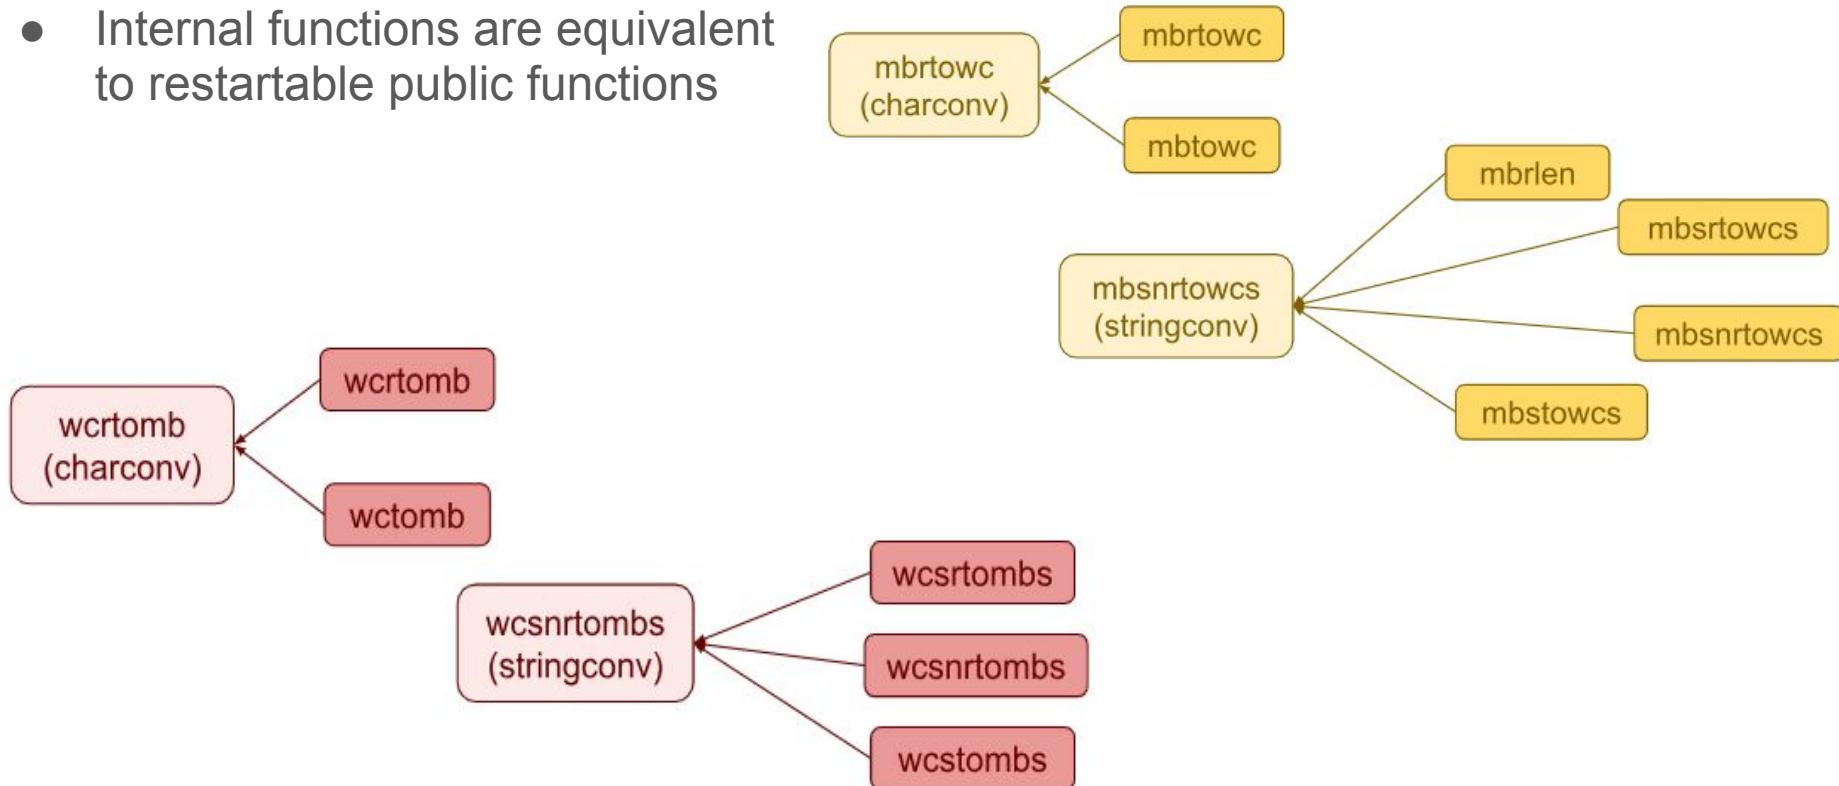
- Main internal interface to interact with mbstate
- Multibyte → Wide Character
  - `push(char8_t)`: Push a single byte from a multibyte sequence
  - `char32_t pop_utf32()`: Pop a wide character
- Wide Character → Multibyte
  - `push(char32_t)`: Push a wide character
  - `char8_t pop_utf8()`: Pop a single byte from a multibyte sequence
- Other utilities
  - `clear()`
  - `isEmpty()/isFull()`
  - `isValidState()`

# StringConverter Class

- Layer of abstraction above `CharacterConverter`
- Construct with an input string and then `pop` converted characters

# Internal Restartable Functions

- Internal functions are equivalent to restartable public functions



05

# Design Decisions



# Size of mbstate/what to store

- Final decision: 6 bytes
  - 4 bytes to hold partial conversion
  - 8 bits each for number of total bytes and bytes stored
- Alternative 1: 4 bytes to hold partial conversion
  - Have to deduce total bytes and conversion status each time
- Alternative 2: 4 bytes
  - `state[20:0]` : partial conversion (utf-32)
  - `state[22:21]`: total bytes
  - `state[28:23]`: num bits processed
  - `state[31:29]`: unused

# StringConverter Class

- The toughest design decision of the entire project
- Do we need a class to handle string conversion, or is the character converter sufficient?
- Class allows for scalability to UTF-16 conversions
- Simplifies code for internal functions

06

# Reflections



# Future Expansion

- Wide character support in `FILE`
- `wprintf`
- 16-bit wide characters using UTF-16 (for windows)
- `wctypes.h`: `iswalpha`, `iswupper/iswlower`, etc
- Widechar to floating point conversion (`wcstod`)
- Add Bazel rules for conversion functions



Thank you for listening!

