



Optimizing generic code lowering to LLVM-IR through function equivalence coalescing



A.K.A. Merging equivalent functions in the front end



TL; DR:

Templates / generics generate many functions.

**Equivalent ones can be merged / coalesced in LLVM IR
and the front end can do pretty good job at it.**



C++ has templates, Rust and Carbon have generics

Generics refers to the ability to generalize code by adding compile-time parameters.

- **Generic function:** function with at least a compile-time parameter (explicit or deduced)
- **Generic type:** type with a compile-time parameter
- **Generic interface:** interface with a compile-time parameter



C++ has templates, Rust and Carbon have generics

Generics refers to the ability to generalize code by adding compile-time parameters.

- **Generic function:** function with at least a compile-time parameter (explicit or deduced)
- **Generic type:** type with a compile-time parameter
- **Generic interface:** interface with a compile-time parameter - interfaces have functions

In Carbon, function parameters:

- Regular parameters `fn F(x: Int) -> bool;`
- **Checked generics** `fn F[T:! type](x: T) -> T;`
- Template generics `fn F[template T:! type](x: T*) -> T*;`

Rust and Carbon's checked generics

Pros:

- **Performance:** Monomorphization eliminates runtime overhead associated with generics (dynamic dispatch, type checking at runtime)
- **Safety:** Type checking at compile time


Cons:

- Increase in code size and compile time



C++: template function -> argument deduction

```
// C++  
  
template <typename T>  
  
T maximum(T a, T b) {  
  
    return (a > b) ? a : b;  
  
}  
  
void main() {  
  
    int intMax = maximum(5, 10);  
    double doubleMax = maximum(3.14, 2.718);  
}
```



Function instantiations

Carbon: checked generic function -> similar deduction

```
// C++  
  
template <typename T>  
  
T maximum(T a, T b) {  
  
    return (a > b) ? a : b;  
  
}  
  
void main() {  
  
    int intMax = maximum(5, 10);  
  
    double doubleMax = maximum(3.14, 2.718);  
  
}
```

Function instantiations

```
// Carbon  
  
fn maximum[T:! Ordered](a: T, b: T) -> T {  
  
    return if a > b then a else b;  
  
}  
  
fn main() {  
  
    var intMax: i32 = maximum(5, 10);  
  
    var doubleMax: f64 = maximum(3.14, 2.718);  
  
}
```

Function specifics

Previous example revisited with pointers

```
// C++                                     // Carbon

template <typename T>

T maximum(T a, T b) {
    return (a > b) ? a : b;
}

void main() {
    int *i_1, *i_2;
    int *intMax = maximum(i_1, i_2);
    double *d_1, *d_2;
    double *doubleMax = maximum(d_1, d_2);
}
```

Function instantiations

```
fn maximum[T:! Ordered](a: T, b: T) -> T {
    return if a > b then a else b;
}

fn main() {
    var i_1: i32*; var i_2: i32*;
    var intMax: i32* = maximum(i_1, i_2);
    var d_1: f64*; var d_2: f64*;
    var doubleMax: f64* = maximum(d_1, d_2);
```

Function specifics



Problem to solve

Function instantiations / specifics, where the arguments have different front-end types, but translate in LLVM to functions with arguments of the same LLVM type.



Problem to solve

Function instantiations / specifics, where the arguments have different front-end types, but translate in LLVM to functions with arguments of the same LLVM type.


Goal: Deduplicate or coalesce such function instantiations / specifics in LLVM-IR.




Why is this important?

Code size and compile time.

Current implementation: Carbon's front-end IR




Carbon's front-end IR: SemIR - focus for this talk



Note: A front end IR vs an AST? ... great idea!



Representation: generic function in FrontendIR



Given “specific_id”: query FrontendIR:

“Tell me more about this instruction or type for this instantiation / specific”

FrontendIR internal
data structure



Easy: compare function bodies! Done! Right?

Single generic function body

Instruction(specific_id1) == Instruction(specific_id2) ?

...

...

Type (specific_id1) == Type (specific_id2) ?

...

First consideration: the instruction can be a call.

To determine equivalence, we need to evaluate the body of *Other_F* => Build the definition of *Other_F* for the two specific_ids => Applies for transitive calls, so **need the full call graph!**

Single generic function body **F**

...

Call_To_Other_F(specific_id1) == Call_To_Other_F(specific_id2) ?

...

Type (specific_id1) == Type (specific_id2) ?

...

Second consideration: recursion.

To determine equivalence in the presence of recursion: need to compare specific_ids.

Single generic function body **F**

Call_To_F(specific_id1) == Call_To_F(specific_id2) ?

...


...

Type (specific_id1) == Type (specific_id2) ?

...

Putting the two together:

Evaluate **transitive calls** + Consider **recursion** => Need call graph **SCCs** (strongly connected components)





So far, checking LLVM type equivalence suffices

Third consideration: LLVM types may be insufficient

- Calling a function on the same “type” may not always call the same function.

```
interface I {  
    fn F();  
}  
class X {}  
impl X as I {  
    fn F() { Core.Print(1); }  
}  
class Y {}  
impl Y as I {  
    fn F() { Core.Print(2); }  
}  
  
fn G(T:! I) { T.F(); }  
  
fn Run() {  
    G(X);  
    G(Y);  
}
```

Other frontend-specific considerations

- Carbon may define generic constants in interfaces too
=> need to consider different LLVM-IR `allocas`.

```
interface I {
    let T!: type;
}

fn G(U:! I) {
    var x: U.T;
}

class C {}

impl C as I where .T = bool {}

class D {}
impl D as I where .T = i32 {}

fn H() {
    G(C);
    G(D);
}
```

All correctness considerations

For each instantiation/specific, consider:

- anything that translates to an `llvm::Type`
- any frontend specific feature that can generate different LLVM IR

Final consideration: performance

Opt for hashing aggregated information into function “fingerprints” using LLVM’s BLAKE3.

- reduce storage
- ease of comparison



High level algorithm for coalescing generic function definitions

At the top level:

1. Generate all function definitions and collect data for each one.
2. Perform coalescing logic and cleanup duplicates.



High level algorithm for coalescing generic function definitions

1. Generate all function definitions:

- Collect **function type** for each function definition into a **common fingerprint** (hash)
- Generate LLVM-IR for each instruction in the function
- If an **instruction/type is specific-dependent**, add this info into the **common fingerprint** (hash)
- When a **call** to another **specific of a generic** is found, that function will need a function definition (creates **declaration**, and marks it as **needing a definition**, will come back later)
 - Collect this info to a **specific fingerprint** (hash)
 - Collect the *non-hashed* specific_id.

High level algorithm for coalescing generic function definitions

2. Perform coalescing logic

- For each two specifics of the same generic
 - 3. **Check fingerprint equivalence**
 - If step 3. returns equivalent, a list of all **equivalent functions** in the **call graph SCC** will be given.
 - Define a **canonical specific** to use for each equivalence found, and perform replacements
 - Once replaced, the duplicated may be deleted from LLVM IR

High level algorithm for coalescing generic function definitions

3. Check fingerprint equivalence

- If **common fingerprints** are different or already known not equivalent => not equivalent
- If **specific fingerprints** are the same => equivalent
- If already assumed equivalent (in a cycle) => equivalent
- Check each of the calls to other specifics (non-hashed specific_ids):
 - If all are equivalent or assumed equivalent => equivalent
 - If any are known non-equivalent => not equivalent
 - SCC resolution: assume they're equivalent and recurse.

Reminder: Why is this important?

Code size and compile time.

Alternatives considered:

- Should it be done in the front-end or the middle end?
 - LLVM has a [MergeFunctions Pass](#)
 - When this was implemented, the guestimate was “in the front-end is likely to yield better results”.
- When to do the coalescing in the front-end?
 - Should we do a pre-processing before actually generating LLVM-IR?
 - Should we first build all the call-graph to determine which specifics are expected to have a definition
 - This will duplicate most of the lowering logic, minus the LLVM-IR creation itself
- Should this be done at all?
 - We guessed yes!



Preliminary performance results

Methodology

- Generating **stress tests** in Carbon - not representative of average compilations!
(<https://github.com/carbon-language/carbon-lang/tree/trunk/testing/base>)
- Extended current code generator with generic function definitions.
- Testing on test sizes of 512 - 1024 lines of Carbon code.

```
845 fn MyNyDnM [TTTTT!: type, UUUUU!: type, VVVV!: type] (p0sHhfwj : TTTTT, jn8Gl0Yy : UUUUU, LjXhpIoTWTTPP : VVVV) {
846     var tKPjlet0VDF: bool;
847     var xfiskftKdLaZGfihMZP8y_MZGSAif0oV6nfRZ18mfjHcd: i32;
848     var u_AkEg44: bool;
849     var LdQXJ9_9wpI: i32;
850
851     ED7ouKpZJY5(jn8Gl0Yy, LjXhpIoTWTTPP, p0sHhfwj);
852     MRRx3eIoV(LdQXJ9_9wpI, u_AkEg44, LjXhpIoTWTTPP);
853     Holg8shLaE(tKPjlet0VDF, jn8Gl0Yy, u_AkEg44);
854     xe446oHmj(xFiskftKdLaZGfihMZP8y_MZGSAif0oV6nfRZ18mfjHcd, LdQXJ9_9wpI, jn8Gl0Yy);
855     YoLY03v79B(p0sHhfwj, tKPjlet0VDF, LdQXJ9_9wpI);
856 }
857
858 fn Main_method_no_template () {
859     var lcP4TvtzNDvE: i32.*;
860     var BNanx7bT: xSqqf;
861     var oYrqnvDil_k: bool;
862     var ev68jlB4ZSkQQYb: i5msb;
863
864     Holg8shLaE(oYrqnvDil_k, BNanx7bT, ev68jlB4ZSkQQYb);
865     Rczn8GQVgx(BNanx7bT, lcP4TvtzNDvE, oYrqnvDil_k);
866     ZKvqJ37KXbqvBItbA_IpdfeS3cRY307gBzBRm(lcP4TvtzNDvE, ev68jlB4ZSkQQYb, BNanx7bT);
867     KzXhj0sPzkhQ(ev68jlB4ZSkQQYb, oYrqnvDil_k, lcP4TvtzNDvE);
868     YoLY03v79B(oYrqnvDil_k, BNanx7bT, ev68jlB4ZSkQQYb);
869 }
870
```

Code size impact with function coalescing logic

- **Balanced types** (primitive and pointers): ~44-56% reduction in:
 - LLVM-IR size before and after optimizations
 - assembly size
 - # of functions generated
- **Most pointer types:** ~98% reduction in:
 - LLVM-IR size before and after optimizations
 - assembly size
 - # of functions generated
- **MergeFunctions pass** makes a difference in **~12% of tests**
 - maximum of ~4% reduction in # of functions, average ~0.1% reduction

Compile time impact with function coalescing logic

- **Balanced types** (primitive and pointers):
 - ~0.9% increase to ~2% reduction in Carbon's **lowering** stage.
 - ~25-51% reduction in Carbon's **opt and codegen** stages (-O3), with inlining disabled.
- **Most pointer types:**
 - ~7-9% reduction in Carbon's **lowering** stage!
 - ~97% reduction in Carbon's **opt and codegen** stages (-O3), with inlining disabled.
- Note: With inlining enabled, similar compile times with the coalescing logic enabled, but compile-time explosion without it, due to aggressive inlining in the non-coalesced functions.



Status and more information

The algorithm is part of Carbon's lowering to LLVM-IR stage.

(<https://github.com/carbon-language/carbon-lang/blob/trunk/toolchain/lower>)

Algorithm and other considerations are documented in the docs:

https://github.com/carbon-language/carbon-lang/blob/trunk/toolchain/docs/coalesce_generic_lowering.md

Lots of room for further improvements!



Questions? Let's talk!

