



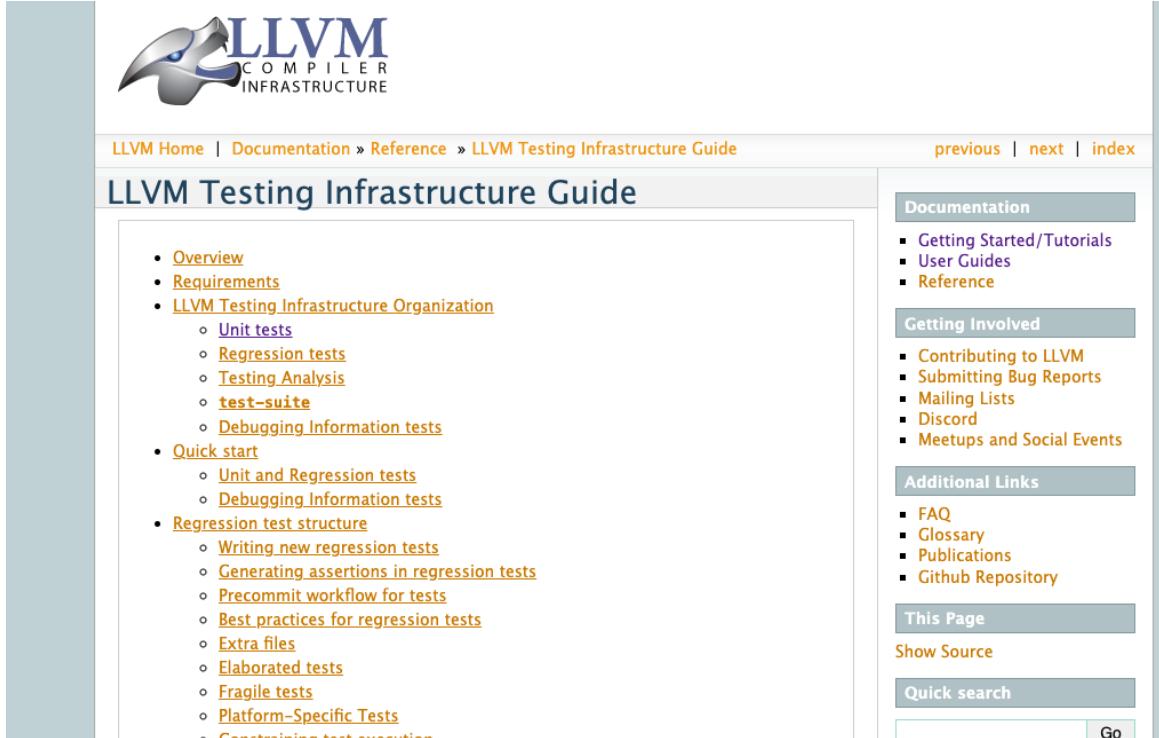
MLIR Testing Guide

What and Why?

Andrzej Warzyński
October 29th, 2025

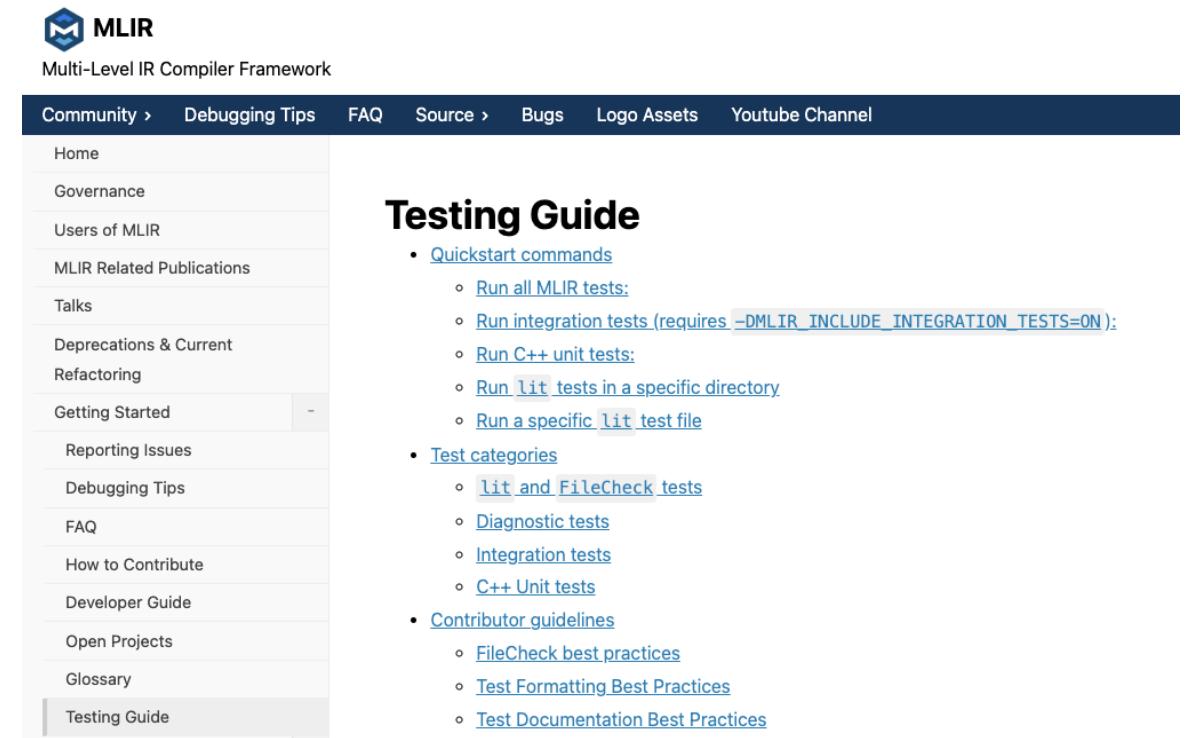
Testing Guide: LLVM + MLIR

<https://llvm.org/docs/TestingGuide.html>



The screenshot shows the LLVM Testing Infrastructure Guide. At the top, there's a navigation bar with links to LLVM Home, Documentation, Reference, and the Testing Guide. Below the title "LLVM Testing Infrastructure Guide", there's a sidebar with sections like Overview, Requirements, LLVM Testing Infrastructure Organization (with sub-links for Unit tests, Regression tests, Testing Analysis, test-suite, and Debugging Information tests), Quick start (with sub-links for Unit and Regression tests, and Debugging Information tests), Regression test structure (with sub-links for Writing new regression tests, Generating assertions in regression tests, Precommit workflow for tests, Best practices for regression tests, Extra files, Elaborated tests, Fragile tests, and Platform-Specific Tests), and Constrained test execution. The main content area contains sections for Documentation (Getting Started/Tutorials, User Guides, Reference), Getting Involved (Contributing to LLVM, Submitting Bug Reports, Mailing Lists, Discord, Meetups and Social Events), Additional Links (FAQ, Glossary, Publications, Github Repository), This Page (Show Source, Quick search), and a footer with a "Go" button.

https://mlir.llvm.org/getting_started/TestingGuide/



The screenshot shows the MLIR Testing Guide. At the top, there's a navigation bar with links to Community, Debugging Tips, FAQ, Source, Bugs, Logo Assets, and Youtube Channel. Below the title "Testing Guide", there's a sidebar with sections like Home, Governance, Users of MLIR, MLIR Related Publications, Talks, Deprecations & Current, Refactoring, Getting Started (with a dropdown menu), Reporting Issues, Debugging Tips, FAQ, How to Contribute, Developer Guide, Open Projects, Glossary, and Testing Guide. The main content area contains a list of test categories and their sub-links: Quickstart commands (Run all MLIR tests, Run integration tests (requires -DMLIR_INCLUDE_INTEGRATION_TESTS=ON), Run C++ unit tests, Run lit tests in a specific directory, Run a specific lit test file), Test categories (lit and FileCheck tests, Diagnostic tests, Integration tests, C++ Unit tests), Contributor guidelines (FileCheck best practices, Test Formatting Best Practices, Test Documentation Best Practices).

Testing Guide: LLVM + MLIR

<https://llvm.org/docs/TestingGuide.html>

The screenshot shows the LLVM Testing Infrastructure Guide. At the top, there's a navigation bar with links to LLVM Home, Documentation, Reference, and the Testing Guide. Below the title "LLVM Testing Infrastructure Guide", there's a sidebar with sections like "Documentation", "Getting Involved", "Additional Links", "This Page", and "Quick search". The main content area contains a list of topics such as Overview, Requirements, LLVM Testing Infrastructure Organization, Quick start, Regression test structure, and Platform-Specific Tests.

This presentation!

https://mlir.llvm.org/getting_started/TestingGuide/

The screenshot shows the MLIR Testing Guide. At the top, there's a navigation bar with links to Community, Debugging Tips, FAQ, Source, Bugs, Logo Assets, and Youtube Channel. The main content area has a sidebar with links to Home, Governance, Users of MLIR, MLIR Related Publications, Talks, Deprecations & Current, Refactoring, Getting Started, Reporting Issues, Debugging Tips, FAQ, How to Contribute, Developer Guide, Open Projects, Glossary, and Testing Guide. The main content area lists various test categories and their sub-links, such as Quickstart commands, Test categories, and Contributor guidelines.

The remaining slides: MLIR Testing Guide overview with LIT + FileCheck examples ☺

Why?

Tests document our code – write it as such.

- **Tests are key** - If you want to understand the code, look at the tests.
- **Test discoverability** – Well-documented tests make it easier to pair tests with patterns and understand their purpose.
- **Test consistency** – Consistent documentation and naming lowers cognitive load and helps avoid duplication.

Anything else?

- To avoid a situation where we end up with tests that we don't understand!
- To make it easier to maintain MLIR!

Only test what's required!

```
// RUN: mlir-opt %s -cse | FileCheck %s
```

Bad

```
// CHECK-LABEL: func.func @simple_constant() -> (i32, i32)
func.func @simple_constant() -> (i32, i32) {
// CHECK-NEXT: %result = arith.constant 1 : i32
// CHECK-NEXT: return %result, %result : i32, i32
// CHECK-NEXT: }

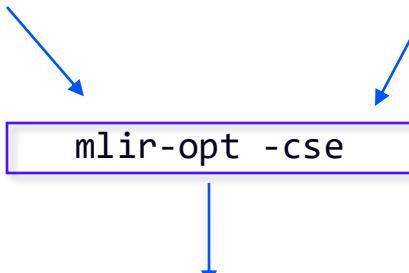
%0 = arith.constant 1 : i32
%1 = arith.constant 1 : i32
return %0, %1 : i32, i32
}
```

Good

```
// RUN: mlir-opt %s -cse | FileCheck %s
```

```
// CHECK-LABEL: func.func @simple_constant
func.func @simple_constant() -> (i32, i32) {
// CHECK-NEXT: %[[RESULT]] = arith.constant 1
// CHECK-NEXT: return %[[RESULT]], %[[RESULT]]

%0 = arith.constant 1 : i32
%1 = arith.constant 1 : i32
return %0, %1 : i32, i32
}
```



```
func.func @simple_constant() -> (i32, i32) {
    %c1_i32 = arith.constant 1 : i32
    return %c1_i32, %c1_i32 : i32, i32
}
```

Only test what's required!

Use FileCheck variables!

```
// RUN: mlir-opt %s -cse | FileCheck %s
```

```
// CHECK-LABEL: func.func @simple_constant() -> (i32, i32)
func.func @simple_constant() -> (i32, i32) {
// CHECK-NEXT: %result = arith.constant 1 : i32
// CHECK-NEXT: return %result, %result : i32, i32
// CHECK-NEXT: }
```

```
%0 = arith.constant 1 : i32
%1 = arith.constant 1 : i32
return %0, %1 : i32, i32
}
```

Bad

Good

```
// RUN: mlir-opt %s -cse | FileCheck %s
```

```
// CHECK-LABEL: func.func @simple_constant
func.func @simple_constant() -> (i32, i32) {
// CHECK-NEXT: %[[RESULT]] = arith.constant 1
// CHECK-NEXT: return %[[RESULT]], %[[RESULT]]
```

```
%0 = arith.constant 1 : i32
%1 = arith.constant 1 : i32
return %0, %1 : i32, i32
}
```

Data type?
Not relevant!

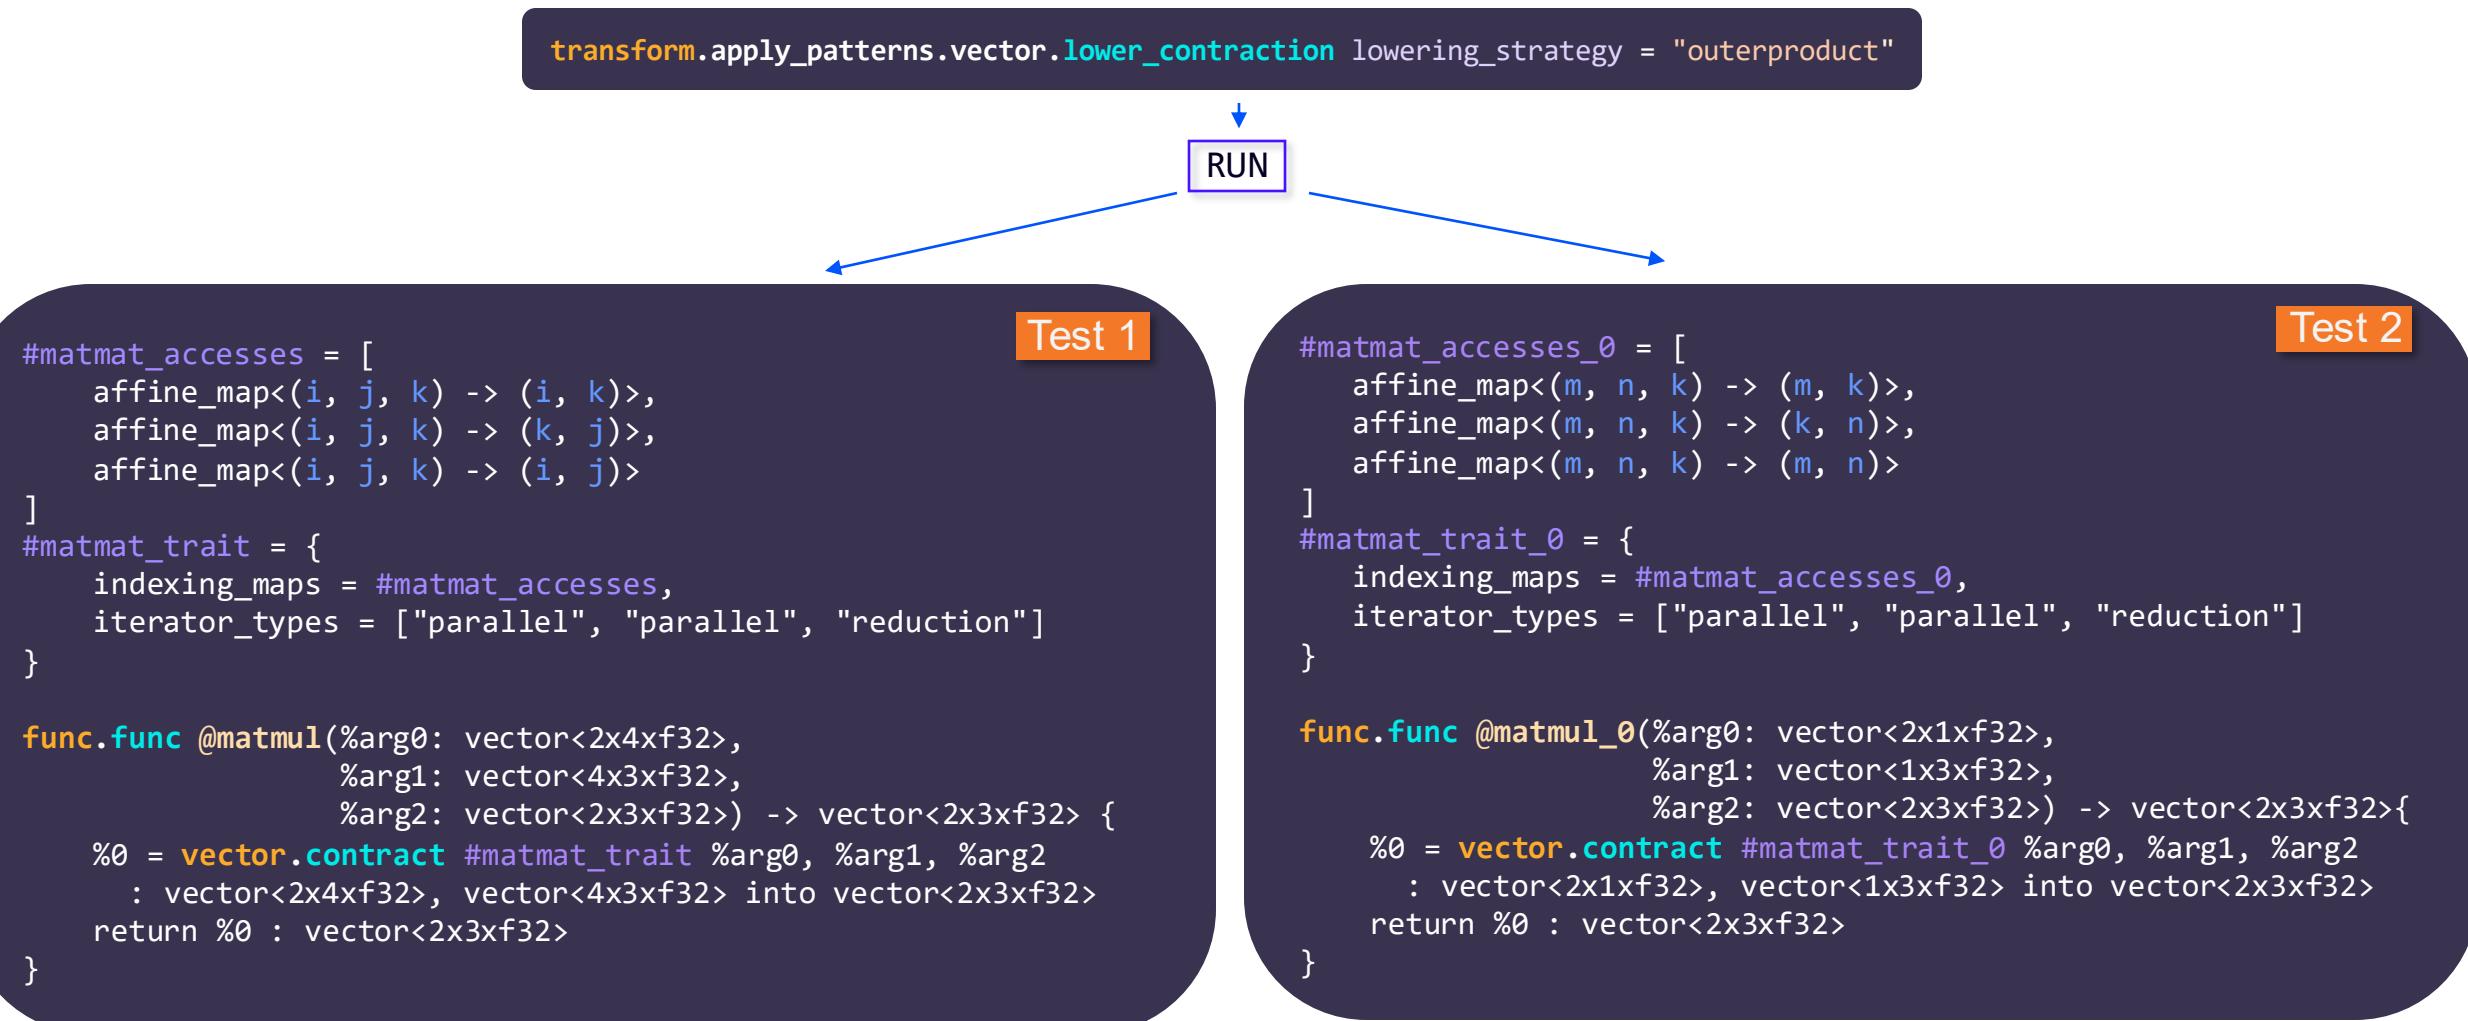
Closing '}'?
Not relevant!

mlir-opt -cse

```
func.func @simple_constant() -> (i32, i32) {
    %c1_i32 = arith.constant 1 : i32
    return %c1_i32, %c1_i32 : i32, i32
}
```

Duplicate tests in one file

Tests for `vector.contract` → `vector.outerproduct` lowering



Duplicate tests in one file

Tests for `vector.contract` → `vector.outerproduct` lowering

- If the traits are identical, what is the difference between `@matmul` and `@matmul_0`?
 - Not a problem with 1-2 tests, but there is 100s!

Identical!

(avoid duplication)

```
#matmat_accesses = [
    affine_map<(i, j, k) -> (i, k)>,
    affine_map<(i, j, k) -> (k, j)>,
    affine_map<(i, j, k) -> (i, j)>
]
#matmat_trait = {
    indexing_maps = #matmat_accesses,
    iterator_types = ["parallel", "parallel", "reduction"]
}

func.func @matmul(%arg0: vector<2x4xf32>,
                  %arg1: vector<4x3xf32>,
                  %arg2: vector<2x3xf32>) -> vector<2x3xf32> {
    %0 = vector.contract #matmat_trait %arg0, %arg1, %arg2
        : vector<2x4xf32>, vector<4x3xf32> into vector<2x3xf32>
    return %0 : vector<2x3xf32>
}
```

Test 1

==

Test 2

```
#matmat_accesses_0 = [
    affine_map<(m, n, k) -> (m, k)>,
    affine_map<(m, n, k) -> (k, n)>,
    affine_map<(m, n, k) -> (m, n)>
]
#matmat_trait_0 = {
    indexing_maps = #matmat_accesses_0,
    iterator_types = ["parallel", "parallel", "reduction"]
}

func.func @matmul_0(%arg0: vector<2x1xf32>,
                   %arg1: vector<1x3xf32>,
                   %arg2: vector<2x3xf32>) -> vector<2x3xf32>{
    %0 = vector.contract #matmat_trait_0 %arg0, %arg1, %arg2
        : vector<2x1xf32>, vector<1x3xf32> into vector<2x3xf32>
    return %0 : vector<2x3xf32>
}
```

Seemingly “similar” tests

Missed opportunities to encode helpful info in tests.

- Leverage test function names + variable names.

mlir-opt -test-vector-to-vector-lowering

```
// CHECK-LABEL: func @maskedLoad_regression_1(  
// CHECK-SAME: %[[A0:.*]]: memref<16xf32>,  
// CHECK-SAME: %[[A1:.*]]: vector<16xf32>
```

Case 1

```
// CHECK: %[[C0:.*]] = arith.constant 0 : index  
// CHECK: %[[LOAD:.*]] = vector.Load %[[A0]][%[[C]]]  
// CHECK-SAME: : memref<16xf32>, vector<16xf32>  
// CHECK: return %[[LOAD]] : vector<16xf32>
```

```
func.func @maskedload_regression_1(  
    %arg0: memref<16xf32>,  
    %arg1: vector<16xf32>) -> vector<16xf32> {  
    %c0 = arith.constant 0 : index
```

```
%vec_i1 = vector.constant_mask [16] : vector<16xi1>  
%ld = vector.maskedload %arg0[%c0], %vec_i1, %arg1  
: memref<16xf32>, vector<16xi1>, vector<16xf32>  
into vector<16xf32>
```

```
return %ld : vector<16xf32>
```

```
// CHECK-LABEL: func @maskedLoad_regression_2(  
// CHECK-SAME: %[[A0:.*]]: memref<16xi8>,  
// CHECK-SAME: %[[A1:.*]]: vector<16xi8>
```

Case 2

```
// CHECK: %[[C0:.*]] = arith.constant 0 : index  
// CHECK: %[[LOAD:.*]] = vector.Load %[[A0]][%[[C]]]  
// CHECK-SAME: : memref<16xi8>, vector<16xi8>  
// CHECK: return %[[LOAD]] : vector<16xi8>
```

```
func.func @maskedload_regression_2(  
    %arg0: memref<16xi8>,  
    %arg1: vector<16xi8>) -> vector<16xi8> {  
    %c0 = arith.constant 0 : index
```

```
%vec_i1 = vector.constant_mask [16] : vector<16xi1>  
%ld = vector.maskedload %arg0[%c0], %vec_i1, %arg1  
: memref<16xi8>, vector<16xi1>, vector<16xi8>  
into vector<16xi8>
```

```
return %ld : vector<16xi8>
```

Seemingly “similar” tests

Missed opportunities to encode helpful info in tests.

- Missed opportunities to encode helpful info in tests.
 - Why not leverage test function names + variable names?

```
// CHECK-LABEL: func @maskedLoad_regression_1(  
// CHECK-SAME: %[[A0:.*]]: memref<16xf32>,  
// CHECK-SAME: %[[A1:.*]]: vector<16xf32>
```

```
// CHECK: %[[C0:.*]] = arith.constant 0 : index  
// CHECK: %[[LOAD:.*]] = vector.Load %[[A0]][%[[C]]]  
// CHECK-SAME: : memref<16xf32>, vector<16xf32>  
// CHECK: return %[[LOAD]] : vector<16x
```

```
func.func @maskedload_regression_1(  
    %arg0: memref<16xf32>,  
    %arg1: vector<16xf32>) -> vector<16x  
%c0 = arith.constant 0 : index
```

```
%vec_i1 = vector.constant_mask [16] : vector<16xi1>  
%ld = vector.maskedload %arg0[%c0], %vec_i1, %arg1  
    : memref<16xf32>, vector<16xi1>, vector<16xf32>  
    into vector<16xf32>
```

```
return %ld : vector<16xf32>
```

Case 1

Every test is a
regression test –
encode some
unique info
instead!

Don't repeat type in
var name!

```
// CHECK-LABEL: func @maskedLoad_regression_2(  
// CHECK-SAME: %[[A0:.*]]: memref<16xi8>,  
// CHECK-SAME: %[[A1:.*]]: vector<16xi8>
```

```
// CHECK: %[[C0:.*]] = arith.constant 0 : index  
// CHECK: %[[LOAD:.*]] = vector.Load %[[A0]][%[[C]]]  
// CHECK-SAME: : memref<16xi8>, vector<16xi8>  
// CHECK: return %[[LOAD]] : vector<16xi8>
```

```
func.func @maskedload_regression_2(  
    %arg0: memref<16xi8>,  
    %arg1: vector<16xi8>) -> vector<16xi8> {  
%c0 = arith.constant 0 : index
```

```
%vec_i1 = vector.constant_mask [16] : vector<16xi1>  
%ld = vector.maskedload %arg0[%c0], %vec_i1, %arg1  
    : memref<16xi8>, vector<16xi1>, vector<16xi8>  
    into vector<16xi8>
```

```
return %ld : vector<16xi8>
```

Case 2

Are these %arg0 and
%arg1?

What's %argN?

Seemingly “similar” tests

Missed opportunities to encode helpful info in tests.

- With new, **consistent** names, the **intent** becomes clear.
 - Testing `vector.maskedload` → `vector.load` folding when the mask is **ALL_TRUE**.
 - The only difference is **f32** vs **i8**.

// CHECK-LABEL: func @maskedLoad_to_load_all_true_f32(
// CHECK-SAME: %[[BASE:.*]]: memref<16xf32>,
// CHECK-SAME: %[[PASS_THRU:.*]]: vector<16xf32>

// CHECK: %[[C0:.*]] = arith.constant 0 : index
// CHECK: %[[LOAD:.*]] = vector.Load %[[BASE]][%[[C]]]
// CHECK-SAME: : memref<16xf32>, vector<16xf32>
// CHECK: return %[[LOAD]] : vector<16xf32>

func.func @maskedload_to_load_all_true_f32 (← Updated!
 %base: memref<16xf32>,
 %pass_thru: vector<16xf32>) -> vector<16xf32> {
 %c0 = arith.constant 0 : index

 %mask = vector.constant_mask [16] : vector<16xi1>
 %ld = vector.maskedload %base[%c0], %vec_i1, %pass_thru
 : memref<16xf32>, vector<16xi1>, vector<16xf32>
 into vector<16xf32>

 return %ld : vector<16xf32>
 Updated!

Case 1

// CHECK-LABEL: func @maskedLoad_to_load_all_true_i8(
// CHECK-SAME: %[[BASE:.*]]: memref<16xi8>,
// CHECK-SAME: %[[PASS_THRU:.*]]: vector<16xi8>

// CHECK: %[[C0:.*]] = arith.constant 0 : index
// CHECK: %[[LOAD:.*]] = vector.Load %[[BASE]][%[[C]]]
// CHECK-SAME: : memref<16xi8>, vector<16xi8>
// CHECK: return %[[LOAD]] : vector<16xi8>

func.func @maskedload_to_load_all_true_i8(
 %base: memref<16xi8>,
 %pass_thru: vector<16xi8>) -> vector<16xi8> {
 %c0 = arith.constant 0 : index

 %mask = vector.constant_mask [16] : vector<16xi1>
 %ld = vector.maskedload %base[%c0], %mask, %pass_thru
 : memref<16xi8>, vector<16xi1>, vector<16xi8>
 into vector<16xi8>

 return %ld : vector<16xi8>
}

Updated!

Case 2

Seemingly “similar” tests

Missed opportunities to encode helpful info in tests.

- With new, **consistent** names, the missing cases are obvious.
 - `vector.maskedload` → `vector.load` folding when the mask is **ALL_FALSE** or **MIXED**

Case 3

```
// CHECK-LABEL: func @maskedLoad_to_load_all_false_f32(  
// CHECK-SAME: %[[BASE:.*]]: memref<16xf32>,  
// CHECK-SAME: %[[PASS_THRU:.*]]: vector<16xf32>  
  
// CHECK: return %[[PASS_THRU]] : vector<16xf32>  
  
func.func @maskedload_to_load_all_false_f32 (  
    %base: memref<16xf32>,  
    %pass_thru: vector<16xf32>) -> vector<16xf32> {  
    %c0 = arith.constant 0 : index  
  
    %mask = vector.constant_mask [0] : vector<16xi1>  
    %ld = vector.maskedload %base[%c0], %vec_i1, %pass_thru  
        : memref<16xf32>, vector<16xi1>, vector<16xf32>  
        into vector<16xf32>  
  
    return %ld : vector<16xf32>  
}
```

Case 4

```
// CHECK-LABEL: func @maskedLoad_to_load_mixed_mask_f32(  
// CHECK-SAME: %[[BASE:.*]]: memref<16xf32>,  
// CHECK-SAME: %[[PASS_THRU:.*]]: vector<16xf32>  
  
// CHECK: %[[LOAD:.*]] = vector.maskedLoad  
// CHECK: return %[[LOAD]]  
  
func.func @negative_maskedload_to_load_mixed_mask_f32 (  
    %base: memref<16xf32>,  
    %pass_thru: vector<16xf32>) -> vector<16xf32> {  
    %c0 = arith.constant 0 : index  
  
    %mask = vector.constant_mask [4] : vector<16xi1>  
    %ld = vector.maskedload %base[%c0], %vec_i1, %pass_thru  
        : memref<16xf32>, vector<16xi1>, vector<16xf32>  
        into vector<16xf32>  
  
    return %ld : vector<16xf32>  
}
```

Automation is available!

- Use `generate-test-checks.py`, but remember:
 - It is **not** authoritative about what constitutes a good test!
 - It won't fix your input IR.
 - The generated CHECK lines usually require reviewing.

```
func.func @maskedload_to_load_all_true_f32 (
    %base: memref<16xf32>,
    %pass_thru: vector<16xf32>) -> vector<16xf32> {
  %c0 = arith.constant 0 : index
  %mask = vector.constant_mask [16] : vector<16xi1>
  %ld = vector.maskedload %base[%c0], %vec_i1, %pass_thru
  : memref<16xf32>, vector<16xi1>, vector<16xf32>
  into vector<16xf32>

  return %ld : vector<16xf32>
```

Case 1

```
mlir-opt -test-vector-to-vector-lowering
| generate-test-checks.py
```

```
// CHECK-LABEL: func.func @maskedload_to_load_all_true_f32(
// CHECK-SAME:   %[ [ARG0:.*]]: memref<?xf32>,
// CHECK-SAME:   %[ [ARG1:.*]]: vector<16xf32>) -> vector<16xf32> {
// CHECK:     %[[CONSTANT_0:.*]] = arith.constant 0 : index
// CHECK:     %[[LOAD_0:.*]] = vector.load %[[ARG0]][{{\[\]}}%[[CONSTANT_0]]] : memref<?xf32>, vector<16xf32>
// CHECK:     return %[[LOAD_0]] : vector<16xf32>
// CHECK: }
```

Call for action!

Leverage MLIR's ability to capture high-level context when writing tests!

1. Follow the guideline – both when **contributing** and **reviewing** PRs.
2. Prioritize “**quality**” over “**quantity**” when writing tests!
3. Send PRs to improve existing tests!