

---

# Synthesizing Practical Transfer Functions in Dataflow Analysis

Yuyou Fan, Xuanyu Peng, Dominic Kennedy, Ben Greenman, John Regehr, Loris D'Antoni



# Dataflow Analysis Recap

# What Is Dataflow Analysis?


Information/Properties that holds among all executions

```
define i4 @func(i4 %arg0) {  
    %and3 = and i4 %arg0, 3  
    %and1 = and i4 %arg0, 1  
    %xor = xor i4 %and3, %and1  
    ret i4 %xor  
}
```




## Known Bits

Known Bits Information tries to determine if a certain bit is always one or zero among all executions.




# Known Bits

Known Bits Information tries to determine if a certain bit is always one or zero among all executions.




# Known Bits

Known Bits Information tries to determine if a certain bit is always one or zero among all executions.




# Implementation of Known Bits in LLVM



# Compute Known Bits

```
define i4 @func(i4 %arg0) {
```




Xor Truth Table on Known Bits

xor	0	1	?
0	0	1	?
1	1	0	?
?	?	?	?

# Compute Known Bits Fact by Transfer Function

```
define i4 @func(i4 %arg0) {
```



```
KnownBits &KnownBits::operator^=(const KnownBits &RHS) {
    // Result bit is 0 if both operand bits are 0 or both are 1.
    APInt Z = (Zero & RHS.Zero) | (One & RHS.One);
    // Result bit is 1 if one operand bit is 0 and the other is 1.
    One = (Zero & RHS.One) | (One & RHS.Zero);
    Zero = std::move(Z);
    return *this;
}
```

Xor Truth Table on Known Bits

xor	0	1	?
0	0	1	?
1	1	0	?
?	?	?	?

# Compute Known Bits Fact by Transfer Function

LLVM implements transfer functions for operations on Known Bits, and some operations are really complicated.

```
KnownBits computeForAddCarry(  
    KnownBits &LHS, KnownBits &RHS,  
    bool CarryZero, bool CarryOne) {  
    APIInt PossibleSumZero = LHS.getMaxValue() + RHS.getMaxValue() + !CarryZero;  
    APIInt PossibleSumOne = LHS.getMinValue() + RHS.getMinValue() + CarryOne;  
    APIInt CarryKnownZero = ~(PossibleSumZero ^ LHS.Zero ^ RHS.Zero);  
    APIInt CarryKnownOne = PossibleSumOne ^ LHS.One ^ RHS.One;  
    APIInt LHSKnownUnion = LHS.Zero | LHS.One;  
    APIInt RHSKnownUnion = RHS.Zero | RHS.One;  
    APIInt CarryKnownUnion = CarryKnownZero | CarryKnownOne;  
    APIInt Known = LHSKnownUnion & RHSKnownUnion & CarryKnownUnion;  
    KnownBits KnownOut;  
    KnownOut.Zero = ~PossibleSumZero & Known;  
    KnownOut.One = PossibleSumOne & Known;  
    return KnownOut;  
}
```



## Summary

Known Bits is a domain provided in LLVM, there are other domains provided in LLVM and MLIR too.


They implement transfer functions for operations in the specification on all domains.



# The Need for a Synthesizer

# What can our synthesizer do?

In short, it reads a specification and generates \*correct\* and \*not bad\* transfer functions automatically.





# Generate more transfer functions

Provide transfer functions for unimplemented operations.

Target Platform	Number of Intrinsics Implemented	Total Number of Intrinsics
X86	30	1713
AArch64	5	1673
RISCV	2	737

LLVM only implements a small number of intrinsics compared to the total number of intrinsics



## Generate more transfer functions

Provide transfer functions for unimplemented operations.

Dialect	Domain	Number of implemented transfer functions	Number of integer operations
Comb	Known Bits	7	20
Arith	Known Bits	N/A	N/A

Other dialects can get benefit from the synthesizer such as `wasmssa`, `emitc`, `index`...

---

## Difficulties in reusing existing transfer functions

Semantics between two operations might be different.

```
// A constant has all bits known!
if (auto constant = dyn_cast<hw::ConstantOp>(op))
    return KnownBits::makeConstant(constant.getValue());
```

---

```
%res0 = comb.shl %arg0, 5 : i4 → 0 : i4
```

```
%res1 = llvm.shl %arg0, 5 : i4 → poison
```

# Difficulties in reusing existing transfer functions

Utils

CMakeLists.txt

InferIntRangeCommon.cpp

Arith InferIntegerRange

Index InferIntegerRange


```
void arith::AndOp::inferResultRanges(ArrayRef<ConstantIntRanges> argRanges,
                                      SetIntRangeFn setResultRange) {
    setResultRange(getResult(), inferAnd(argRanges));
}

void AndOp::inferResultRanges(ArrayRef<ConstantIntRanges> argRanges,
                              SetIntRangeFn setResultRange) {
    setResultRange(getResult(),
                  inferIndexOp(inferAnd, argRanges, CmpMode::Unsigned));
}
```

---




## Prove correctness of transfer functions.

Soundness: The analysis result covers all runtime values.



# Prove correctness of transfer functions.

Soundness: The analysis result covers all runtime values.




---

## Prove correctness of transfer functions.

Precision: How many values that never occur at runtime included by the analysis result.


```
define i4 @func(i4 %arg0) {
```




# Prove correctness of transfer functions.

Precision: How many values that never occur at runtime included by the analysis result.

```
define i4 @func(i4 %arg0) {
```



The set of all runtime values






## How LLVM tests a transfer function

Let's see how LLVM test soundness of transfer functions.

```
unsigned Bits = 4;  
ForeachKnownBits(Bits, [&](const KnownBits &Known1) {  
    ForeachKnownBits(Bits, [&](const KnownBits &Known2) {  
        ForeachKnownBits(1, [&](const KnownBits &KnownCarry) {  
            .....  
            ForeachNumInKnownBits(Known1, [&](const APInt &N1) {  
                ForeachNumInKnownBits(Known2, [&](const APInt &N2) {  
                    ForeachNumInKnownBits(KnownCarry, [&](const APInt &Carry) {  
                        .....  
                }  
            }  
        }  
    }  
}
```

# How the synthesizer verifies a transfer function


With SMT semantics of the operation and transfer dialect, the synthesizer checks soundness as a SMT query:



---

# Our previous work

We defined a Transfer Dialect in MLIR and it encodes LLVM APIInt operations in SMT Dialect.



Our Previous Work:

[First-Class Verification Dialects for MLIR](#)

[2023 LLVM Dev Mtg - An SMT dialect for assigning semantics to MLIR dialects](#)



## Summary

Our goal is not to beta LLVM transfer functions.

Our goal is to provide sound and precise transfer functions in C++ for new operations or domains.




# Design of the Synthesizer

---

# Big Picture



Our synthesizer comes with a synthesis loop



---

# Input Specification

Our input specification is defined in MLIR.



---

## Find candidates by stochastic search

Because the search space is extremely large, we adopt a stochastic search strategy to explore candidate transfer functions.



## Find candidates by stochastic search

Because the search space is extremely large, we adopt a stochastic search strategy to explore candidate transfer functions.

```
.....  
APInt autogen9 = autogen3 & autogen7;  
APInt autogen10 = smax(autogen6, autogen4);  
APInt autogen11 = umin(autogen1, autogen8);  
.....
```


```
APInt autogen9 = autogen3 * autogen7;
```

Both might  
happen

```
APInt autogen10 = smax(autogen9, autogen4);
```

# Evaluate candidates


Evaluation Engine verifies the soundness of a transfer function and produce a precision score by the cost function.



---

# What candidates does the synthesizer keep?

It saves two kinds of candidates.



---

# What candidates does the synthesizer keep?

It saves two kinds of candidates.

Solution set

Sound candidates and  
with high precision

Candidates only  
sound conditionally  
and with high  
precision

# What candidates does the synthesizer keep?

It saves two kinds of candidates.

Solution set

Sound candidates and with high precision


Candidates only sound conditionally and with high precision

```
domain candidate(LHS, RHS);  
if(candidate_guard(LHS, RHS))  
    return candidate(LHS, RHS);  
return top();
```

# Maintain Solution Set

The synthesizer maintains a dynamic pool of candidate transfer functions.

Less effective ones are pruned to prevent the solution set from growing excessively.



---

# Generate result

The synthesizer generates the final solution by the meet of all transfer functions.

```
Vec<2> xor_solution(Vec<2> autogen0,Vec<2> autogen1){  
    Vec<2> autogen2 = xor_partial_solution_0(autogen0,autogen1);  
    Vec<2> autogen3 = xor_partial_solution_1(autogen0,autogen1);  
    Vec<2> autogen4 = meet(autogen2,autogen3);  
    return autogen4;  
}
```

# Generate result

The synthesizer generates the final solution by the meet of all transfer functions.


```
Vec<2> xor_solution(Vec<2> autogen0,Vec<2> autogen1){  
    Vec<2> autogen2 = xor_partial_solution_0(autogen0,autogen1);  
    Vec<2> autogen3 = xor_partial_solution_1(autogen0,autogen1);  
    Vec<2> autogen4 = meet(autogen2,autogen3);  
    return autogen4;  
}
```

```
extern "C" Vec<2> sub_solution(Vec<2> autogen0,Vec<2> autogen1){  
    Vec<2> autogen2 = sub_partial_solution_0(autogen0,autogen1);  
    Vec<2> autogen3 = sub_partial_solution_1(autogen0,autogen1);  
    Vec<2> autogen4 = sub_partial_solution_2(autogen0,autogen1);  
    Vec<2> autogen5 = sub_partial_solution_3(autogen0,autogen1);  
    Vec<2> autogen6 = sub_partial_solution_4(autogen0,autogen1);  
    Vec<2> autogen7 = sub_partial_solution_5(autogen0,autogen1);  
    Vec<2> autogen8 = sub_partial_solution_6(autogen0,autogen1);  
    Vec<2> autogen9 = sub_partial_solution_7(autogen0,autogen1);  
    Vec<2> autogen10 = sub_partial_solution_8(autogen0,autogen1);  
    Vec<2> autogen11 = sub_partial_solution_9(autogen0,autogen1);  
    Vec<2> autogen12 = sub_partial_solution_10(autogen0,autogen1);  
    Vec<2> autogen13 = sub_partial_solution_11(autogen0,autogen1);  
    Vec<2> autogen14 = sub_partial_solution_12(autogen0,autogen1);  
    Vec<2> autogen15 = sub_partial_solution_13(autogen0,autogen1);  
    Vec<2> autogen16 = sub_partial_solution_14(autogen0,autogen1);  
    Vec<2> autogen17 = sub_partial_solution_15(autogen0,autogen1);  
    Vec<2> autogen18 = meet(autogen2,autogen3);  
    Vec<2> autogen19 = meet(autogen18,autogen4);  
    Vec<2> autogen20 = meet(autogen19,autogen5);  
    Vec<2> autogen21 = meet(autogen20,autogen6);  
    Vec<2> autogen22 = meet(autogen21,autogen7);  
    Vec<2> autogen23 = meet(autogen22,autogen8);  
    Vec<2> autogen24 = meet(autogen23,autogen9);  
    Vec<2> autogen25 = meet(autogen24,autogen10);  
    Vec<2> autogen26 = meet(autogen25,autogen11);  
    Vec<2> autogen27 = meet(autogen26,autogen12);  
    Vec<2> autogen28 = meet(autogen27,autogen13);  
    Vec<2> autogen29 = meet(autogen28,autogen14);  
    Vec<2> autogen30 = meet(autogen29,autogen15);  
    Vec<2> autogen31 = meet(autogen30,autogen16);  
    Vec<2> autogen32 = meet(autogen31,autogen17);  
    return autogen32;  
}
```

# Generate result

The synthesizer generates the final solution by the meet of all transfer functions.

```
Vec<2> xor_solution(Vec<2> autogen0,Vec<2> autogen1){  
    Vec<2> autogen2 = xor_partial_solution_0(autogen0,autogen1);  
    Vec<2> autogen3 = xor_partial_solution_1(autogen0,autogen1);  
    Vec<2> autogen4 = meet(autogen2,autogen3);  
    return autogen4;  
}
```



```
extern "C" Vec<2> sub_solution(Vec<2> autogen0,Vec<2> autogen1){  
    Vec<2> autogen2 = sub_partial_solution_0(autogen0,autogen1);  
    Vec<2> autogen3 = sub_partial_solution_1(autogen0,autogen1);  
    Vec<2> autogen4 = sub_partial_solution_2(autogen0,autogen1);  
    Vec<2> autogen5 = sub_partial_solution_3(autogen0,autogen1);  
    Vec<2> autogen6 = sub_partial_solution_4(autogen0,autogen1);  
    Vec<2> autogen7 = sub_partial_solution_5(autogen0,autogen1);  
    Vec<2> autogen8 = sub_partial_solution_6(autogen0,autogen1);  
    Vec<2> autogen9 = sub_partial_solution_7(autogen0,autogen1);  
    Vec<2> autogen10 = sub_partial_solution_8(autogen0,autogen1);  
    Vec<2> autogen11 = sub_partial_solution_9(autogen0,autogen1);  
    Vec<2> autogen12 = sub_partial_solution_10(autogen0,autogen1);  
    Vec<2> autogen13 = sub_partial_solution_11(autogen0,autogen1);  
    Vec<2> autogen14 = sub_partial_solution_12(autogen0,autogen1);  
    Vec<2> autogen15 = sub_partial_solution_13(autogen0,autogen1);  
    Vec<2> autogen16 = sub_partial_solution_14(autogen0,autogen1);  
    Vec<2> autogen17 = sub_partial_solution_15(autogen0,autogen1);  
    Vec<2> autogen18 = meet(autogen2,autogen3);  
    Vec<2> autogen19 = meet(autogen18,autogen4);  
    Vec<2> autogen20 = meet(autogen19,autogen5);  
    Vec<2> autogen21 = meet(autogen20,autogen6);  
    Vec<2> autogen22 = meet(autogen21,autogen7);  
    Vec<2> autogen23 = meet(autogen22,autogen8);  
    Vec<2> autogen24 = meet(autogen23,autogen9);  
    Vec<2> autogen25 = meet(autogen24,autogen10);  
    Vec<2> autogen26 = meet(autogen25,autogen11);  
    Vec<2> autogen27 = meet(autogen26,autogen12);  
    Vec<2> autogen28 = meet(autogen27,autogen13);  
    Vec<2> autogen29 = meet(autogen28,autogen14);  
    Vec<2> autogen30 = meet(autogen29,autogen15);  
    Vec<2> autogen31 = meet(autogen30,autogen16);  
    Vec<2> autogen32 = meet(autogen31,autogen17);  
    return autogen32;  
}
```



# Experimental Results



# Synthesis Result

Here is the partial result synthesizing 39 operations on KnownBits Domain. Transfer functions are tested on random 8-bit inputs.

Operation	Ours	LLVM
And/Or/Xor	100%	100%
Modu	59%	52.7%
UShlSat	96.6%	N/A
Mul	60.6%	73.2%

The results are sensitive to random factors, leading to variability between runs.

Operation	Ours	LLVM
AvgFloorS	39.3%	100%
Shl	56.9%	96.5%
Abds	60.1%	100%
Add	58.70%	100%
Sub	60.6%	100%



## Evaluation on SPEC 2017 CPU

We also test transfer functions on SPEC 2017 CPU benchmarks. The table below compares Known Bits found by LLVM and ours.

Project	perlbench	gcc	mcf	omnetpp	xalancbmk	x264	deepsjeng	leela	xz
KLOC	362	1,304	3	134	520	96	10	21	33
LLVM	1,356,555	4,272,154	910	62,251	475,736	247,344	15,578	25,207	76,907
Ours	1,305,537	4,195,918	910	62,102	442,838	218,171	14,780	19,353	72,090
Precision Loss	3.76%	<b>1.78%</b>	0.00%	0.24%	6.29%	11.79%	5.12%	23.22%	6.26%



## Summary

Our goal is not to beat LLVM, but generate correct and precise transfer functions when it comes to a new dialect or new domain.

By giving the specification and the operation with SMT semantics, the synthesizer can generate usable and sound transfer functions used in dataflow analysis.



# Thanks for listening!



Yuyou Fan  
Fifth-year PhD student  
[yuyou.fan@utah.edu](mailto:yuyou.fan@utah.edu)  
Looking for full-time job  
around May 2026  
<- My resume.