


Specializing MLIR Data Structures: an Experiment



Compilation time matters

- Faster development cycle

Compilation time matters

- Faster development cycle
- Faster CI times

Compilation time matters

- Faster development cycle
- Faster CI times
- Happiness

Compilation time matters


- Faster development cycle
- Faster CI times
- Happiness



MLIR is known to be slow




MLIR is known to be slow



MLIR is known to be slow



MLIR is known to be slow



Compile-time is a trade-off

Compile-time is a trade-off

What are you willing to pay to get faster compile-time?

First trade-off : Extensibility vs. Performance



Machine learning



Hardware design




Fully Homomorphic
encryption

General-purpose
languages




Computer
Graphics

All abstractions using a single IR



By Alex Zinenko

Built on an extensible foundation



Built on an extensible foundation


```
%a, %b = arith.addi_extended %x, %y : i32
```

mlir::Type
mlir::Attribute
mlir::Operation
mlir::OpInterface

Dialect 1

Dialect 2

Dialect 3



Built on an extensible foundation

```
%a, %b = arith.addi_extended %x, %y : i32
```


```
func.call @foo(%a, %b, %c, %d, %e) : ...
```

mlir::Type
mlir::Attribute
mlir::Operation
mlir::OpInterface

Dialect 1

Dialect 2

Dialect 3



Built on an extensible foundation

```
%a, %b = arith.addi_extended %x, %y : i32
```

```
func.call @foo(%a, %b, %c, %d, %e) : ...
```


```
cf.br ^bb0(%x : i32)
```

mlir::Type
mlir::Attribute
mlir::Operation
mlir::OpInterface

Dialect 1

Dialect 2

Dialect 3



Built on an extensible foundation

```
%a, %b = arith.addi_extended %x, %y : i32
```

```
func.call @foo(%a, %b, %c, %d, %e) : ...
```

```
cf.br ^bb0(%x : i32)
```


```
scf.for %arg3 = %c0 to %0 step %c1 {  
    ...  
}
```

```
mlir::Type  
mlir::Attribute  
mlir::Operation  
mlir::OpInterface
```

Dialect 1

Dialect 2

Dialect 3



Built on an extensible foundation

```
%a, %b = arith.addi_extended %x, %y : i32
```

```
func.call @foo(%a, %b, %c, %d, %e) : ...
```

```
cf.br ^bb0(%x : i32)
```

```
scf.for %arg3 = %c0 to %0 step %c1 {  
    ...  
}
```


```
%x, %y = test.test(%a, %b, %c) ^bb0, ^bb1 {  
    ...  
, {  
    ...  
}
```

```
mlir::Type  
mlir::Attribute  
mlir::Operation  
mlir::OpInterface
```

Dialect 1


Dialect 2

Dialect 3



Built on an extensible foundation

```
dyn_cast<MemoryEffectOpInterface>(op);
```



First trade-off : Extensibility vs. Performance

- Operation memory size
- Interfaces and traits
- Checking operation Opcode
- Attribute dictionary

Second trade-off: Usability vs. Performance

Second trade-off: Usability vs. Performance

Operation	Complexity
Creating an operation op	$O(op)$
Inserting an operation	$O(1)$
Detaching an operation	$O(1)$
Erasing an operation op	$O(op_{rec} + op_{rec}.uses)$
Replacing an operand/block operand	$O(1)$
Replacing a value v with another value	$O(v.uses)$
Creating a block $block$	$O(block)$
Inserting a block	$O(1)$
Detaching a block	$O(1)$
Erasing a block	$O(block_{rec} + block_{rec}.uses)$
Creating a region	$O(1)$
Erasing a region	$O(region_{rec} + region_{rec}.uses)$

A pointer-heavy data structure...

```
^bb(%a):
```

```
%0 = add ..., ...
```

```
%1 = add ..., ...
```

```
%2 = add %0, ...
```

```
%3 = add %0, ...
```

```
br bb2 (%3)
```


```
^bb2(%b):
```

```
scf.for ... {
```

```
^bb3:
```

```
}
```


A pointer-heavy data structure...




...resulting in costly bookkeeping

```
setOperand:  
ldr    x8, [x0, #0x48]  
mov    w9, w1  
add    x8, x8, x9, lsl #5  
ldr    x10, [x8, #0x8]  
cbz   x10, 0x5f0  
ldr    x9, [x8]  
str    x9, [x10]  
cbz   x9, 0x5f0  
ldr    x10, [x8, #0x8]  
str    x10, [x9, #0x8]  
str    x2, [x8, #0x18]  
ldr    x9, [x2]  
stp    x9, x2, [x8]  
cbz   x9, 0x604  
str    x8, [x9, #0x8]  
str    x8, [x2]  
ret
```


Could we get the best of both worlds?



Could we get the best of both worlds?




Could we get the best of both worlds?



Operation	Complexity
Creating an operation op	$O(op)$
Inserting an operation	$O(1)$
Detaching an operation	$O(1)$
Erasing an operation op	$O(op_{rec} + op_{rec}.uses)$
Replacing an operand/block operand	$O(1)$
Replacing a value v with another value	$O(v.usages)$
Creating a block $block$	$O(block)$
Inserting a block	$O(1)$
Detaching a block	$O(1)$
Erasing a block	$O(block_{rec} + block_{rec}.uses)$
Creating a region	$O(1)$
Erasing a region	$O(region_{rec} + region_{rec}.uses)$



Could we get the best of both worlds?



Operation	Complexity
Creating an operation op	$O(op)$
Inserting an operation	$O(1)$
Detaching an operation	$O(1)$
Erasing an operation op	$O(op_{rec} + op_{rec}.uses)$
Replacing an operand/block operand	$O(1)$
Replacing a value v with another value	$O(v.usages)$
Creating a block $block$	$O(block)$
Inserting a block	$O(1)$
Detaching a block	$O(1)$
Erasing a block	$O(block_{rec} + block_{rec}.uses)$
Creating a region	$O(1)$
Erasing a region	$O(region_{rec} + region_{rec}.uses)$



What if MLIR data structures were an interface?




```
pub trait Context: Sized {  
    type Operation<'a, 'b>: Operation<'a, 'b, Self>;  
    ...  
}
```


This talk:

**What are the performance opportunities of specializing
MLIR data structures?**


Could we specialize MLIR to a dialect set?




Could we specialize MLIR to a dialect set?




Could we specialize MLIR to a dialect set?



Generating MLIR IR from TableGen?



Generating MLIR IR from TableGen?



Where are the speedup opportunities?

- Checking operation Opcode
- Accessing interfaces and traits
- Reducing memory footprint
- Fixed property size

Checking operation opcodes

```
dyn_cast<AddIOP>(op)
```

```
ldr    x8, [x0, #0x30]
ldr    x8, [x8, #0x10]
adrp   x9, 0x0
ldr    x9, [x9]
cmp    x8, x9
csel   x0, x0, xzr, eq
ret
```

```
AddIOP::opcode == op.opcode
```

```
ldr    w8, [x0, #0x20]
cmp    w8, #0x0
csel   x0, x0, xzr, eq
ret
```

Checking operation opcodes

```
for (Operation *op : /*std::vector*/ ops) {          0.31 ns/op
}
```

Checking operation opcodes

for (Operation *op : /*std::vector*/ ops) { }	0.31 ns/op
for (Operation *op : ops) isa<AddIOp>(op)	0.51 ns/op

Checking operation opcodes

```
for (Operation *op : /*std::vector*/ ops) {           0.31 ns/op
}
```

```
for (Operation *op : ops)
    isa<AddI0p>(op)                                0.51 ns/op
```

```
for (Operation *op : ops)
    op.opCode == AddI0p::opcode                      0.39 ns/op
```

Checking operation opcodes

```
for (Operation *op : /*std::vector*/ ops) {           0.31 ns/op
}
```

```
for (Operation *op : ops)
    isa<AddI0p>(op)                                2.6  ns/op
```

```
for (Operation *op : ops)
    op.opCode == AddI0p::opcode                      2.3  ns/op
```

Accessing interfaces

Accessing interfaces

μBenchmark: Interfaces vs LLVM



4x faster!



```
static int64_t getCost(llvm::Instruction *op) {
    using namespace llvm;
    switch (op->getOpcode()) {
        // Terminators
        case Instruction::Ret:    return 42;
        case Instruction::Br:     return 13;
        case Instruction::Switch: return 18;
    ...
    }
    2.7ns/op
}

for (llvm::Instruction *op : ops) {
    auto cost = getCost(op);
}

```

```
for (Operation *op : /*std::vector*/ops) {
    if (auto costIface = dyn_cast<CostModel>(op))
        auto cost = casted.getCost();
}
11.71ns/op
```

Memory footprint improvements

```
struct Operation {  
    /*64*/     Block          parent  
    /*64*/     Location        location  
    /*32*/     unsigned        orderIndex  
    /*32*/     unsigned        numResults  
    /*32*/     unsigned        numSuccs  
    /*23*/     unsigned        numRegions  
    /*1*/      bool           hasOperandStorage  
    /*8*/      unsigned        propertiesStorageSize  
    /*64*/     OperationName   name  
    /*64*/     DictionaryAttr  attrs  
}
```

Memory footprint improvements

```
struct Operation {  
    /*64*/     Block          parent  
    /*64*/     Location        location  
    /*32*/     unsigned        orderIndex  
    /*32*/     unsigned        numResults  
    /*32*/     unsigned        numSuccs  
    /*23*/     unsigned        numRegions  
    /*1*/      bool           hasOperandStorage  
    /*8*/      unsigned        propertiesStorageSize  
    /*64*/     OperationName   name  
    /*64*/     DictionaryAttr  attrs  
}
```

Memory footprint improvements

```
struct Operation {  
    /*64*/     Block          parent  
    /*64*/     Location        location  
    /*32*/     unsigned        orderIndex  
    /*32*/     unsigned        numResults  
    /*32*/     unsigned        numSuccs  
    /*23*/     unsigned        numRegions  
    /*1*/      bool           hasOperandStorage  
    /*8*/      unsigned        propertiesStorageSize  
    /*64*/     OperationName   name  
    /*64*/     DictionaryAttr  attrs  
}
```

Memory footprint improvements

```
struct Operation {  
    /*64*/     Block          parent  
    /*64*/     Location        location  
    /*32*/     unsigned        orderIndex  
    /*32*/     unsigned        numResults  
    /*32*/     unsigned        numSuccs  
    /*23*/     unsigned        numRegions  
    /*1*/      bool           hasOperandStorage  
    /*8*/      unsigned        propertiesStorageSize  
    /*64*/     OperationName   name  
    /*64*/     DictionaryAttr  attrs  
}
```

24 bytes to save per operation

Memory footprint improvements

```
for (int i = 0; i < numOps; i++) {  
    MyOp::create(builder, operands, i32);  
}
```

Memory footprint improvements

```
for (int i = 0; i < numOps; i++) {  
    MyOp::create(builder, operands, i32);  
}
```

	Before	After
0 operands	44 ns/op	
2 operands	75 ns/op	

Memory footprint improvements

```
for (int i = 0; i < numOps; i++) {  
    MyOp::create(builder, operands, i32);  
}
```

	Before	After
0 operands	44 ns/op	37 ns/op
2 operands	75 ns/op	69 ns/op

Memory footprint improvements

```
for (int i = 0; i < numOps; i++) {  
    MyOp::create(builder, operands, i32);  
}
```

	Before	After
0 operands	44 ns/op	37 ns/op
2 operands	75 ns/op	69 ns/op

10% speedup

Memory footprint improvements

my.addi %cst, %x -> my.addi %x, %cst

Memory footprint improvements

`my.addi %cst, %x` -> `my.addi %x, %cst`

	Before	After
All matches	7.1 ns/op	
No matches	3.6 ns/op	

Memory footprint improvements

my.addi %cst, %x -> my.addi %x, %cst

	Before	After
All matches	7.1 ns/op	6.1 ns/op
No matches	3.6 ns/op	3 ns/op

10/20% speedup

Fixed property size

Fixed property size

Class Operation	OperandStorage	OpProperties	Block*[]	Region[]	OpOperand[]
-----------------	----------------	--------------	----------	----------	-------------

Fixed property size

Class Operation	OperandStorage	OpProperties	Block*[]	Region[]	OpOperand[]
-----------------	----------------	--------------	----------	----------	-------------

Class Operation	OperandStorage	OpProperties	OpOperand[]
-----------------	----------------	--------------	-------------

Fixed property size

Class Operation	OperandStorage	OpProperties	Block*[]	Region[]	OpOperand[]
-----------------	----------------	--------------	----------	----------	-------------

Class Operation	OperandStorage	OpProperties	OpOperand[]
-----------------	----------------	--------------	-------------

Class Operation	OperandStorage	OpProperties	OpOperand[]
-----------------	----------------	--------------	-------------

Fixed property size

Class Operation	OperandStorage	OpProperties	Block*[]	Region[]	OpOperand[]
-----------------	----------------	--------------	----------	----------	-------------

Class Operation	OperandStorage	OpProperties	OpOperand[]
-----------------	----------------	--------------	-------------

Class Operation	OperandStorage	OpProperties	OpOperand[]
-----------------	----------------	--------------	-------------

If we fix OpProperties size, can we reuse allocations during rewrites?

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

- Create a new allocation

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

- Create a new allocation
- Initialize the new operation

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

- Create a new allocation
- Initialize the new operation
- **Replace results**

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

- Create a new allocation
- Initialize the new operation
- **Replace results**
- Erase old operands

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

- Create a new allocation
- Initialize the new operation
- **Replace results**
- Erase old operands
- Deallocate old operation

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

```
op.name = myConstantName;
op.eraseOperands(0, op.getNumOperands());
cast<MyConstant>(op).setValue(0);
```

- Create a new allocation
- Initialize the new operation
- **Replace results**
- Erase old operands
- Deallocate old operation

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

```
op.name = myConstantName;
op.eraseOperands(0, op.getNumOperands());
cast<MyConstant>(op).setValue(0);
```

- Create a new allocation
- Initialize the new operation
- **Replace results**
- Erase old operands
- Deallocate old operation
- Change name field

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

```
op.name = myConstantName;
op.eraseOperands(0, op.getNumOperands());
cast<MyConstant>(op).setValue(0);
```

- Create a new allocation
- Initialize the new operation
- **Replace results**
- Erase old operands
- Deallocate old operation
- Change name field
- Erase old operands

Fixed property size

```
auto newOp = MyConstant::create(rewriter, value);
rewriter.replaceOp(&op, newOp);
```

```
op.name = myConstantName;
op.eraseOperands(0, op.getNumOperands());
cast<MyConstant>(op).setValue(0);
```

- Create a new allocation
 - Initialize the new operation
 - **Replace results**
 - Erase old operands
 - Deallocate old operation
- Change name field
 - Erase old operands
 - Initialize the new operation

Fixed property size

`add(constant(cst1), constant(cst2)) -> constant(cst1 + cst2)`

Fixed property size

`add(constant(cst1), constant(cst2)) -> constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op

Fixed property size

`add(constant(cst1), constant(cst2)) -> constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op

Replacing ops 114 ns/op

Fixed property size

`add(constant(cst1), constant(cst2))` -> `constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op

Replacing ops 114 ns/op

Reusing storage 7.4 ns/op

Fixed property size

`add(constant(cst1), constant(cst2))` -> `constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op

Replacing ops 114 ns/op

15x speedup!!!

Reusing storage 7.4 ns/op

Takeaways from these μ -benchmarks

- We can get a few improvements from memory footprint
- Interfaces/traits heavy passes would benefit a lot
- Huge potential from reusing allocations

Could we restrict MLIR features for performance?

Operation	Complexity
Creating an operation op	$O(op)$
Inserting an operation	$O(1)$
Detaching an operation	$O(1)$
Erasing an operation op	$O(op_{rec} + op_{rec}.uses)$
Replacing an operand/block operand	$O(1)$
Replacing a value v with another value	$O(v.uses)$
Creating a block $block$	$O(block)$
Inserting a block	$O(1)$
Detaching a block	$O(1)$
Erasing a block	$O(block_{rec} + block_{rec}.uses)$
Creating a region	$O(1)$
Erasing a region	$O(region_{rec} + region_{rec}.uses)$



Where could we get performance from?

- Use-def chain updates
- Linked lists walks
- Memory optimizations

Removing the def-use chain

With def-use chain

```
setOperand:  
ldr    x8, [x0, #0x48]  
mov    w9, w1  
add    x8, x8, x9, lsl #5  
ldr    x10, [x8, #0x8]  
cbz   x10, 0x5f0  
ldr    x9, [x8]  
str    x9, [x10]  
cbz   x9, 0x5f0  
ldr    x10, [x8, #0x8]  
str    x10, [x9, #0x8]  
str    x2, [x8, #0x18]  
ldr    x9, [x2]  
stp    x9, x2, [x8]  
cbz   x9, 0x604  
str    x8, [x9, #0x8]  
str    x8, [x2]  
ret
```

Removing the def-use chain

With def-use chain

```
setOperand:  
ldr    x8, [x0, #0x48]  
mov    w9, w1  
add    x8, x8, x9, lsl #5  
ldr    x10, [x8, #0x8]  
cbz   x10, 0x5f0  
ldr    x9, [x8]  
str    x9, [x10]  
cbz   x9, 0x5f0  
ldr    x10, [x8, #0x8]  
str    x10, [x9, #0x8]  
str    x2, [x8, #0x18]  
ldr    x9, [x2]  
stp    x9, x2, [x8]  
cbz   x9, 0x604  
str    x8, [x9, #0x8]  
str    x8, [x2]  
ret
```

Without def-use chain

```
setOperand:  
ldr    x8, [x0, #0x48]  
add    x8, x8, w1, uxtw #4  
str    x2, [x8, #0x8]  
ret
```

Removing the def-use chain

MyAddIOp::create(builder, {})

45 ns/op

MyAddIOp::create(builder, {a,b,c,d,e,f,g})

77 ns/op

Removing the def-use chain

MyAddIOp::create(builder, {})

45 ns/op → **45 ns/op**

MyAddIOp::create(builder, {a,b,c,d,e,f,g})

77 ns/op → **54 ns/op**

Removing the def-use chain

MyAddIOp::create(builder, {})

45 ns/op → 45 ns/op

MyAddIOp::create(builder, {a,b,c,d,e,f,g})

77 ns/op → 54 ns/op

1.42x speedup

Removing the def-use chain

my.addi %cst, %x -> my.addi %x, %cst

	MLIR	Reduce memory	
All matches	7.1 ns/op	6.1 ns/op	
No matches	3.6 ns/op	3 ns/op	

Removing the def-use chain

my.addi %cst, %x -> my.addi %x, %cst

	MLIR	Reduce memory	Without def-use chain
All matches	7.1 ns/op	6.1 ns/op	4.8 ns/op
No matches	3.6 ns/op	3 ns/op	3.6 ns/op

1.47x speedup

Removing the def-use chain

`add(constant(cst1), constant(cst2))` -> `constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op

Replacing ops 114 ns/op

Reusing storage 7.4 ns/op

Removing the def-use chain

`add(constant(cst1), constant(cst2))` -> `constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op

Replacing ops 114 ns/op

Reusing storage 7.4 ns/op

No def-use chain 5.3 ns/op

Removing the linked-list

`add(constant(cst1), constant(cst2))` -> `constant(cst1 + cst2)`

Iterating the IR 3.6 ns/op


Replacing ops 114 ns/op

Reusing storage 7.4 ns/op


No def-use chain 5.3 ns/op

No linked-list 4.9 ns/op

"Arena" allocation for operations




"Arena" allocation for operations




- Each operation requires a new allocation
- Each pointer has to be 64 bits

"Arena" allocation for operations



"Arena" allocation for operations



- One global resizable allocation
- Pointers are offsets (possibly 32 bits)
- All operations have the same size

"Arena" allocation for operations

`add(constant(cst1), constant(cst2))` \rightarrow `constant(cst1 + cst2)`

"Arena" allocation for operations

`add(constant(cst1), constant(cst2))` \rightarrow `constant(cst1 + cst2)`

MLIR

Creating the IR 75 ns/op

Rewriting the IR 114 ns/op

"Arena" allocation for operations

`add(constant(cst1), constant(cst2))` → `constant(cst1 + cst2)`

MLIR

7 operands

Creating the IR

75 ns/op

19 ns/op

Rewriting the IR

114 ns/op

76 ns/op

"Arena" allocation for operations

`add(constant(cst1), constant(cst2))` -> `constant(cst1 + cst2)`

	MLIR	7 operands	3 operands
Creating the IR	75 ns/op	19 ns/op	15 ns/op
Rewriting the IR	114 ns/op	76 ns/op	62 ns/op

Conclusion

- Very powerful optimization opportunities when the IR size is fixed

Conclusion

- Very powerful optimization opportunities when the IR size is fixed
- Memory optimizations go a long way

Conclusion

- Very powerful optimization opportunities when the IR size is fixed
- Memory optimizations go a long way
- Maybe there is hope to have an LLVM dialect as fast as LLVM?

Conclusion

- Very powerful optimization opportunities when the IR size is fixed
- Memory optimizations go a long way
- Maybe there is hope to have an LLVM dialect as fast as LLVM?
- Can we apply some of these optimizations in MLIR already?

Conclusion

- Very powerful optimization opportunities when the IR size is fixed
- Memory optimizations go a long way
- Maybe there is hope to have an LLVM dialect as fast as LLVM?
- Can we apply some of these optimizations in MLIR already?

