

Lightweight Fault Isolation (LFI): LLVM Support for Efficient Native Code Sandboxing

Zachary Yedidia and Tal Garfinkel



Shout out

Nathan Egge, Sharjeel Khan, Daniel Moghimi, Shravan Narayan, Taehyun Noh, Abhishek Sharma, Pirama Arumuga Nainar, Derek Schuff, Colin Cross, Matthew Maurer, Dan Behrendt, Kris Adler, Wonsik Kim, Urs Hözle, Cory Baker, Matthew Sotoudeh, Eli Friedman, MaskRay, Alexis Engelke

Android



Google



Memory Safety is a Big Problem

Google Chrome: ~70% of bugs (2015–2020)

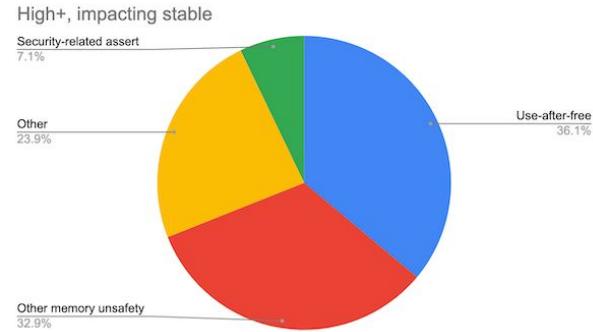


Image from the Chromium project blog
<https://www.chromium.org/Home/chromium-security/memory-safety/>

Microsoft Windows: ~70% of bugs (2006–2018)

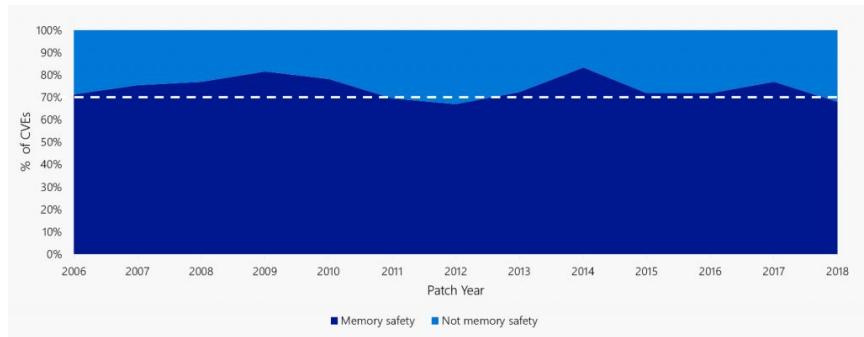


Image from the Microsoft security response center blog
<https://msrc-blog.microsoft.com/2019/07/16/a-proactive-approach-to-more-secure-code/>

Limited Options

Unsound Mitigations: ASLR, stack canaries, CFI, bounds checks.

+Works with existing code!

- Often bypassed
- Hard to reason about benefits + overheads

Rewrite code in safe language (Rust)?

+Sound security properties

- Huge engineering cost and time: rewrite, retest, support
- Lose benefits of Cooperation: expertise, shared ownership

Sound Security + Works with Existing Code?

Process-based Sandboxing?

The Good:

- + Sound Security Properties!
- + Works with existing library code!

The Bad: Expensive:

Slow context switches (IPC), Sandbox Creation, Memory Overheads

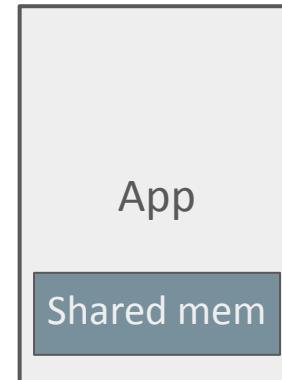
Higher Latency, Lower Throughput, Poor scalability

The Ugly: Awkward + Complex

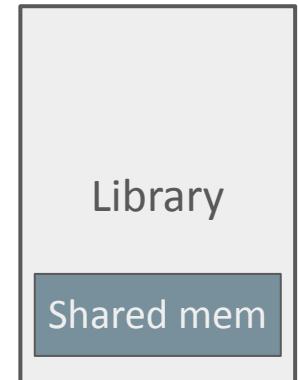
Refactoring, Developer experience 😞

application => distributed system

Process 1

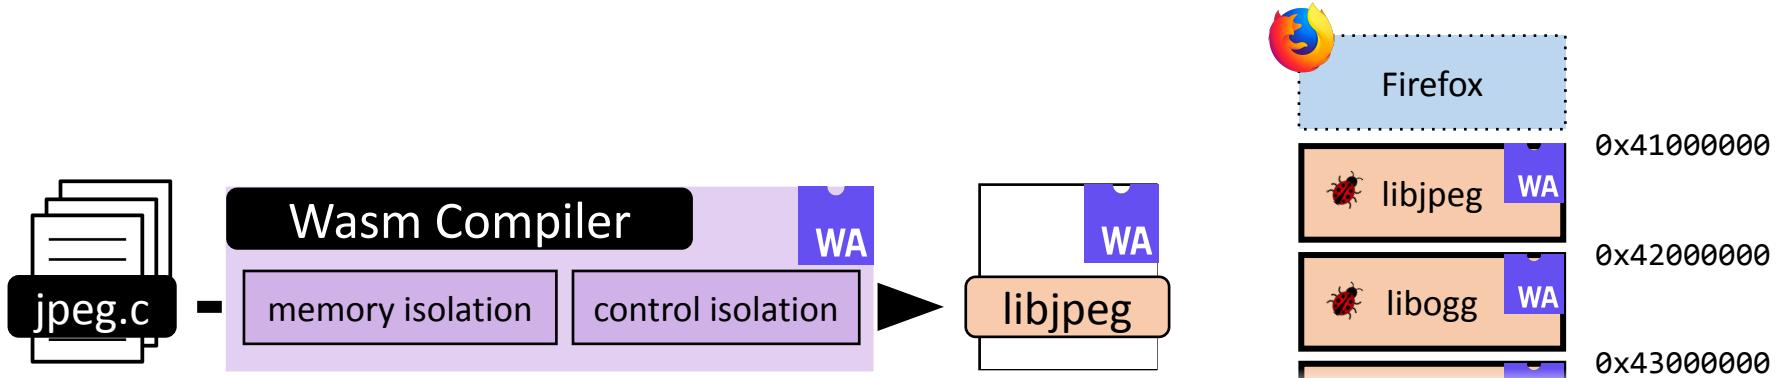


Process 2



Is there another option?

Sandboxing in Firefox Render with WebAssembly



100x faster startup

1000x faster context switches

Less memory

Securing Firefox with WebAssembly



By [Nathan Froyd](#)

Posted on February 25, 2020 in [Featured Article](#), [Firefox](#), [Rust](#), [Security](#), and [WebAssembly](#)

Protecting the security and privacy of individuals is a [central tenet](#) of Mozilla's mission, and so we constantly endeavor to make our users safer online. With a

Deployed in Firefox since 2021

Securing Firefox with WebAssembly



By [Nathan Froyd](#)

Posted on February 25, 2020 in [Featured Article](#), [Firefox](#), [Rust](#), [Security](#), and [WebAssembly](#)

Protecting the security and privacy of individuals is a [central tenet](#) of Mozilla's mission, and so we constantly endeavor to make our users safer online. With a

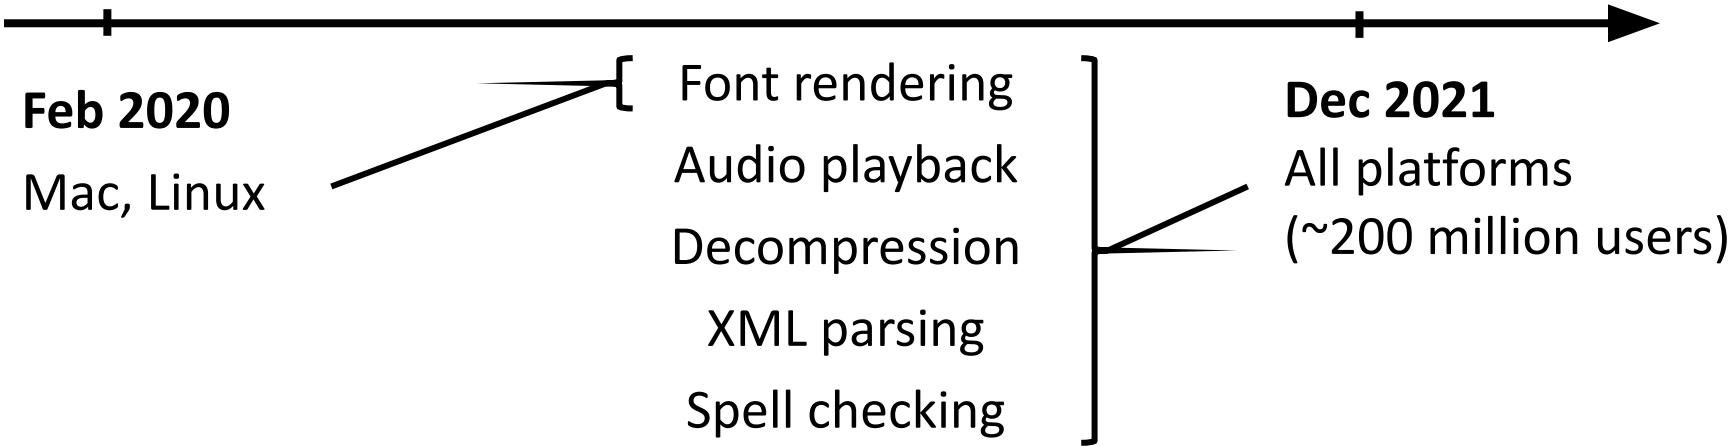
WebAssembly and Back Again: Fine-Grained Sandboxing in Firefox 95



By [Bobby Holley](#)

Posted on December 6, 2021 in [Featured Article](#), [Firefox](#), and [JavaScript](#)

In Firefox 95, we're shipping a novel sandboxing technology called [RLBox](#) —



Wasm has many limitations...

Compatibility

Very limited SIMD/intrinsics

No hand-written assembly

No dynamic code generation (JIT)

libc compat (WASI-libc – no glibc, bionic, etc.)

Limited POSIX

ABI differences (different pointer sizes)

Breaks existing tools: debuggers, perf,
sanitizers,...

Performance

Wasm2c on ARM64:

37.5% Geomean on SPEC 2017

Limits optimization ():

e.g. Precise memory traps => no
truncating SFI

Breaks key library optimizations ()

- SIMD, Hand-written Assembly, JIT

Wasm - Good for the Web, LFI - Good for library sandboxing

High-level SFI – WebAssembly

>> target sandboxed IR <<

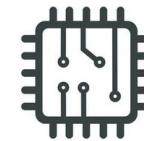
Platform Independence



Low-level SFI – LFI, (original SFI, NaCl)

>> low level instruction rewrites <<

Compatibility, Performance, Security, Simplicity

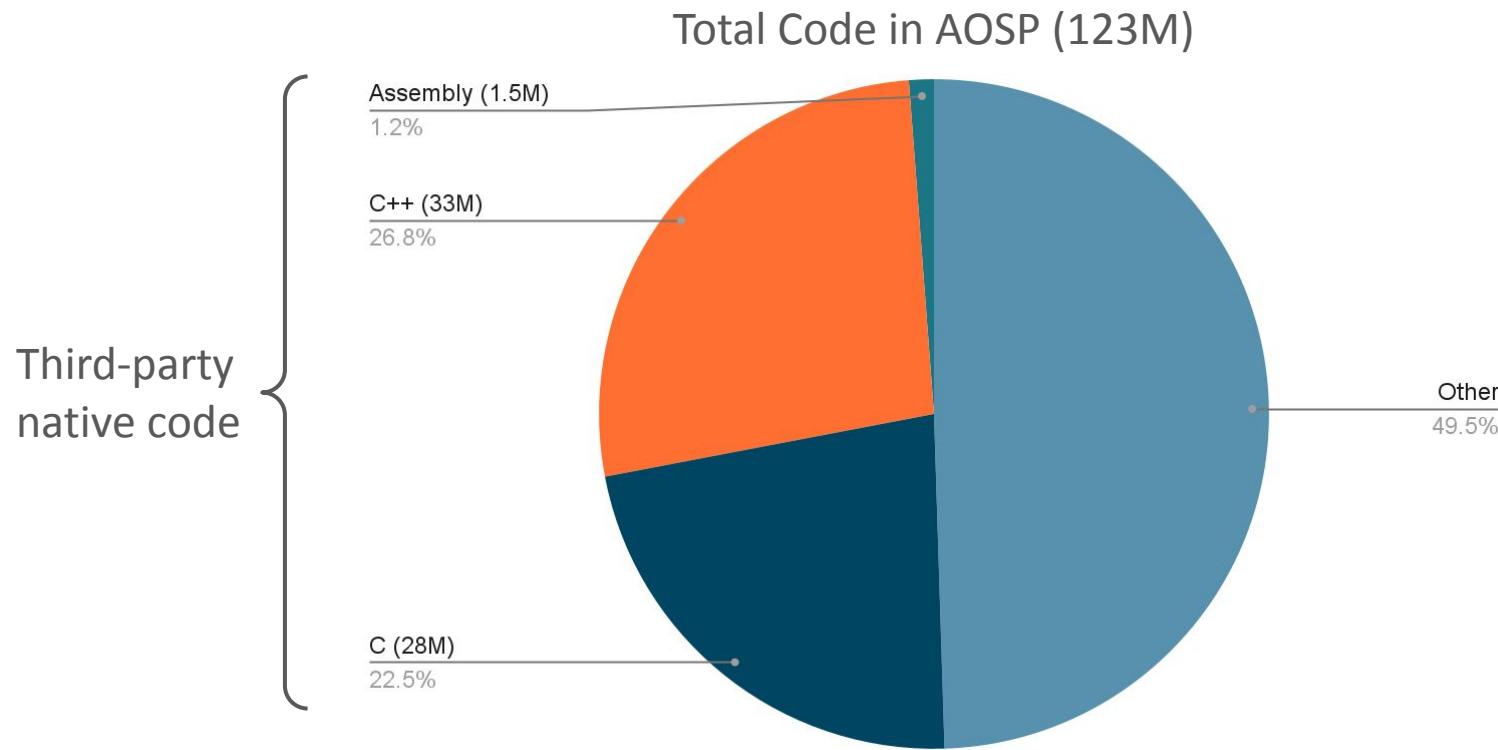


arm



Android

Half of AOSP is Unsafe Third-Party Code



Lightweight Fault Isolation (LFI)

Low-level SFI for isolating buggy or malicious C/C++/Asm.

Performance: fully exploit architecture and compiler

Compatibility: Linux API, 64-bit ABI, hand-written Asm, existing tools (gdb, perf, etc.)

Security: simple, verifiable => high assurance isolation.

Usability: easy retrofitting in existing code; minimal change to dev experience

...Another tool in the LLVM memory safety toolbox... RFC on discourse

Lightweight Fault Isolation

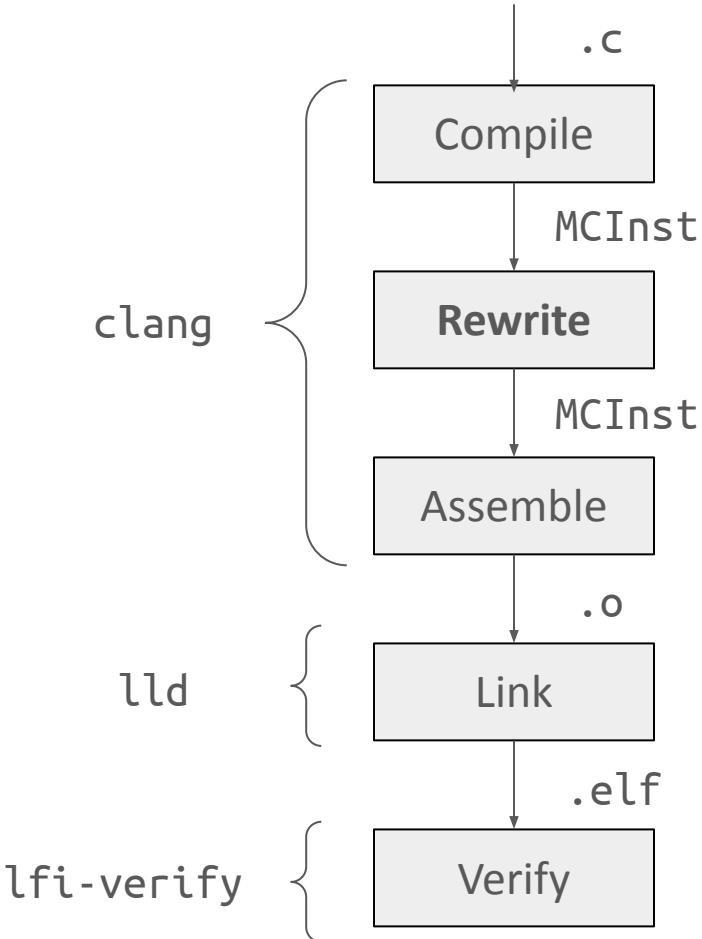
LFI Compiler Pipeline

New architecture target: aarch64_lfi.
(example: aarch64_lfi-linux-musl)

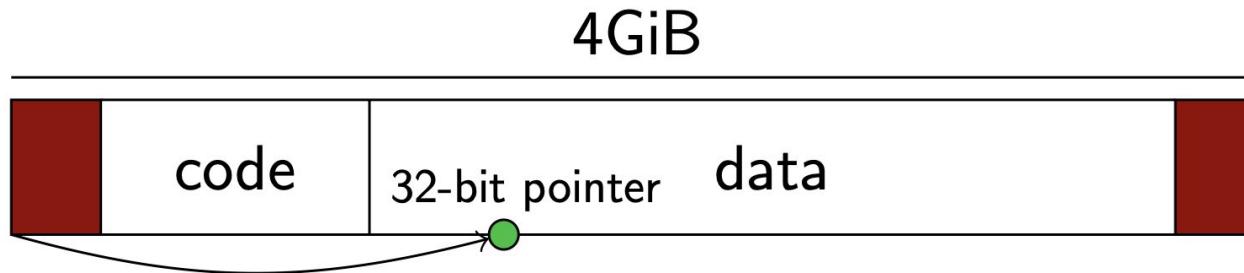
Uses custom LFI MCStreamer to apply rewrites.
Goal: automatically handle hand-written asm.
Alternative: external .s/.o rewriter.

Other compiler stages remain unchanged*.
*only need to reserve a few registers.

All sandboxed code goes through the rewriter:
Libc, dynamic linker, libc++, compiler-rt, ...



LFI Execution Environment



Runtime handles syscalls (indirect branches) and enforces strict W^X.

Dynamic codegen can be supported: Run verifier before `mprotect(PROT_EXEC)`.

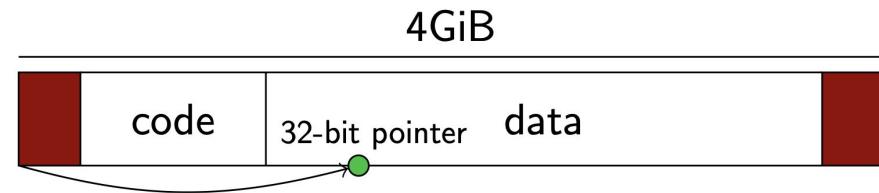
48-bit address space: supports up to 64K sandboxes (80KiB guard pages).

64-bit ABI: matches host ABI, compatible with existing code.

Rewriting Memory Accesses

Reserve x27: sandbox base.

Reserve x28: any sandbox address.



```
ldr x0, [x1]
```

```
ldr x0, [x27, w1, uxtw]
```

Mask with addressing mode

```
ldp x0, x1, [x2]
```

```
add x28, x27, w2, uxtw  
ldp x0, x1, [x28]
```

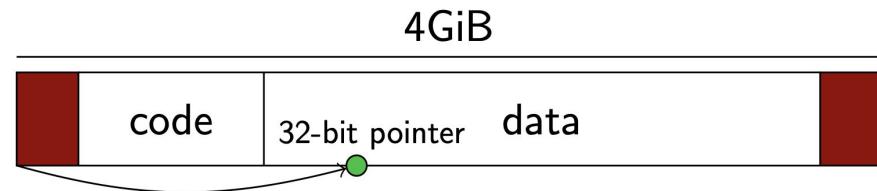
Mask with guard instruction

For experts: Arm64 fixed-width instructions → no bundling necessary. See paper for details.

Rewriting Control Flow

Mask all modifications to x28, x30.

Rewrite indirect branches to target x28.



br x0

add x28, x27, w0, uxtw
br x28

Key: Arm64 instructions are fixed-width!

ret

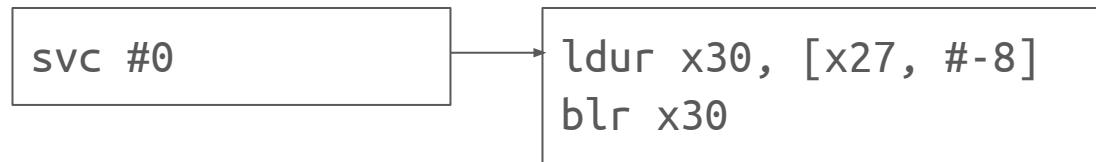
ret

Safe as long as each individual instruction is safe.

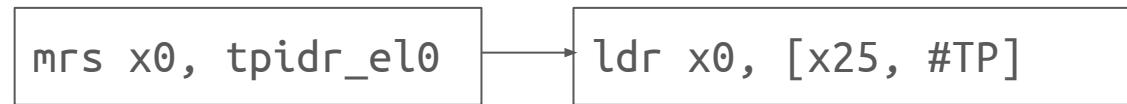
Rewriting System Call Instructions

Runtime entrypoints: placed before the first page of the sandbox (read-only).

→ x27 already points here!

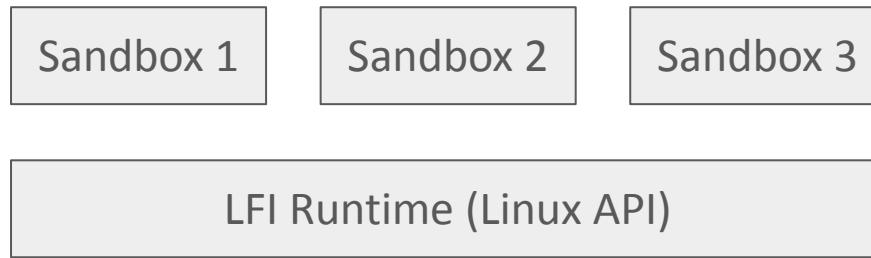


Thread-local storage: reserve a register (x25), or alternatively rewrite to runtime call.



LFI Runtime

LFI runtime: virtualizes user mode



LFI runtime: Loads sandbox, and handles system calls that come from the sandbox.

Implements a **Linux API** (similar to WASI, QEMU-user).

Many system calls are passed through with simple checks.

Putting it all together

```
$ git clone https://github.com/lua/lua
```

```
$ cd lua
```

```
$ make CC=aarch64_lfi-linux-musl-clang
```

```
...
```

```
$ lfi-run ./lua
```

```
Lua 5.5.0 Copyright (C) 1994-2025 Lua.org, PUC-Rio
```

```
> print('hello world')
```

```
hello world
```

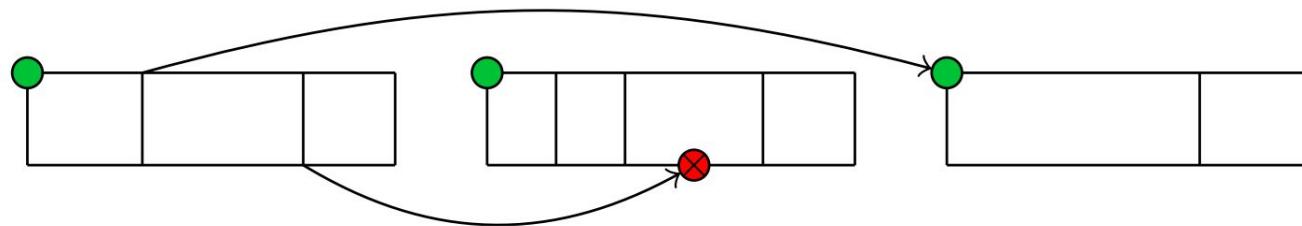
```
0000000000023900 <lua_xmove>:  
23900: eb01001f    cmp    x0, x1  
23904: 54000360    b.eq   0x23970 <lua_xmove+0x70>  
23908: 8b2042b2    add    x28, x27, w0, uxtw  
2390c: f9400a48    ldr    x8, [x28, #0x10]  
23910: 7100045f    cmp    w2, #0x1  
23914: cb22d108    sub    x8, x8, w2, sxtw #4  
23918: f9000a48    str    x8, [x28, #0x10]
```



Support for x86-64: Control Flow

Primary challenge: variable-width instructions.

- Simple solution: instruction bundles (32-byte chunks).

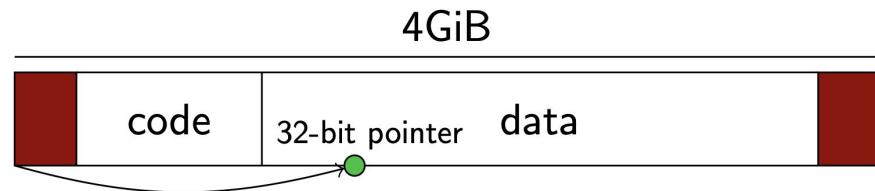


`jmpq *%rax`

`andl $0xfffffffffe0, %eax
addq %r14, %rax
jmpq *%rax`

Support for x86-64: Segmentation

Problem: no $\%rX + %eX$ addressing mode!



But there is $\%gs + %eX\dots$

Store sandbox base in $\%gs$, and use segment-relative addressing.

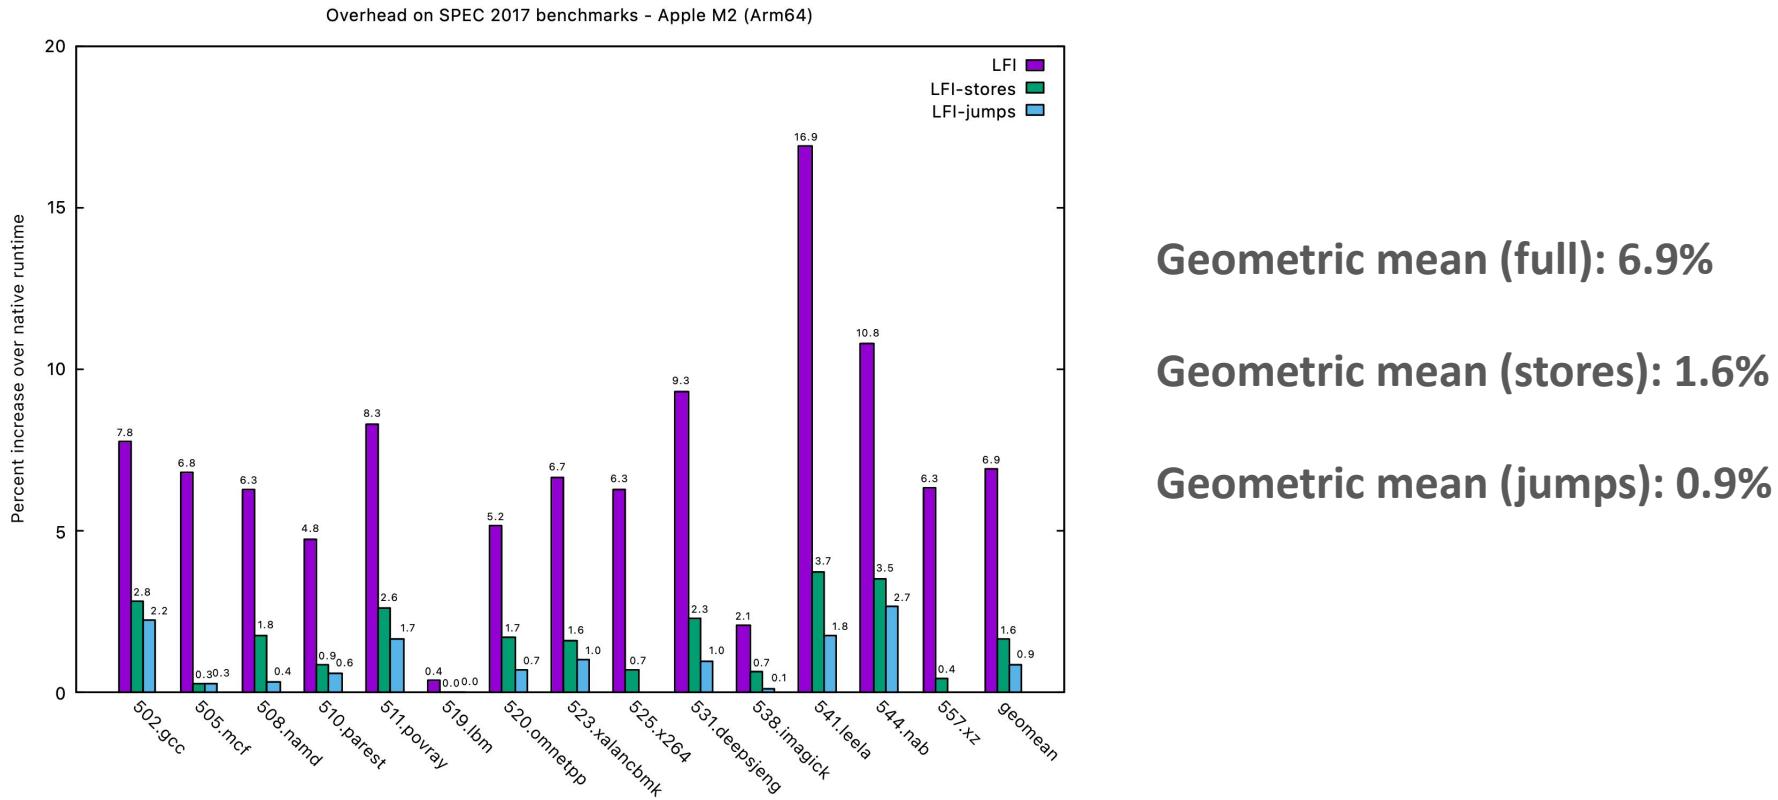
```
movq %rax, (%rdi) → movq %rax, %gs:(%edi)
```

Cuts performance overhead from ~15% to ~7%!

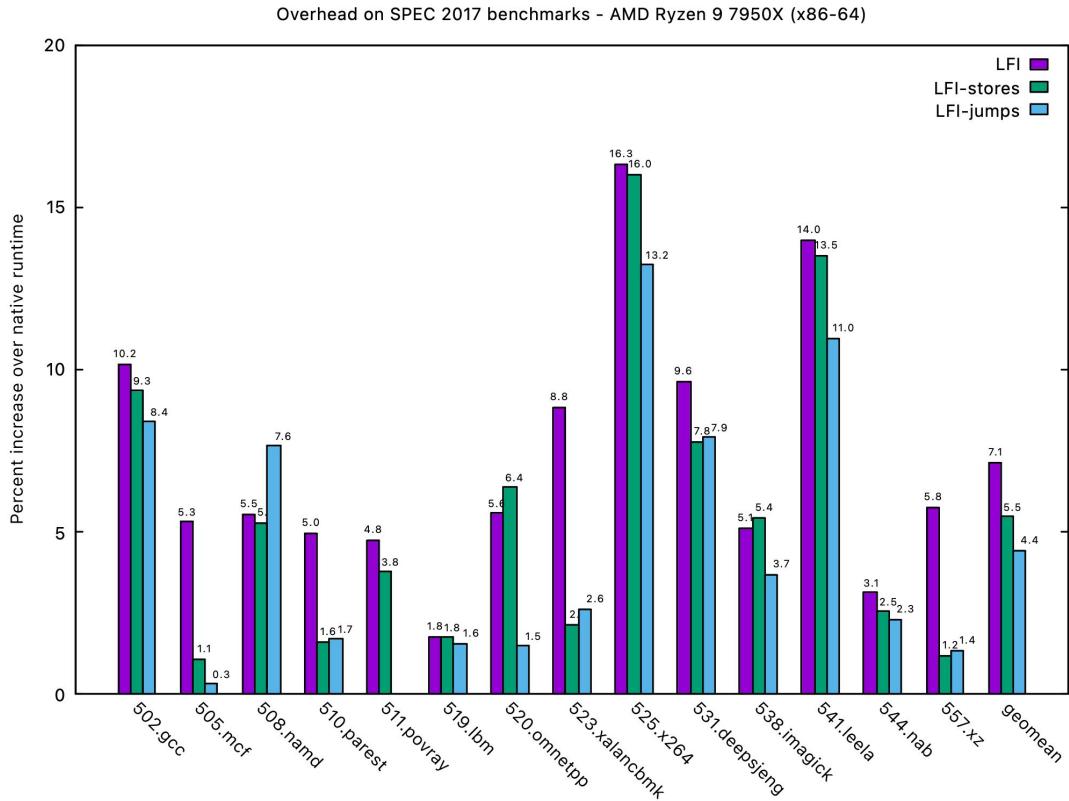
See “Segue” paper: <https://shravanrn.com/pubs/seguecg.pdf>

Performance

Performance: SPEC 2017



Performance (x86-64): SPEC 2017



Geometric mean (full): 7.1%

Geometric mean (stores): 5.5%

Geometric mean (jumps): 4.4%

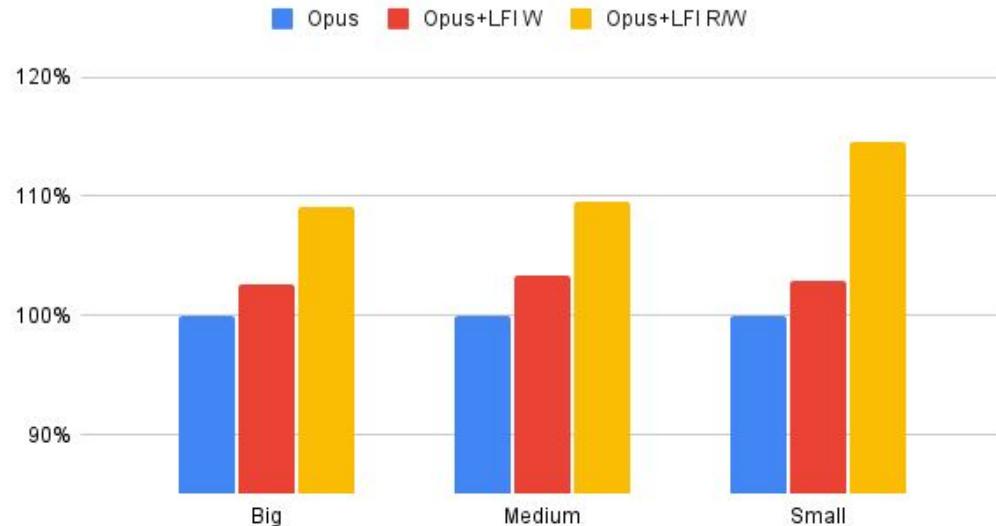


Performance: libopus

Standalone opus_demo on Pixel 9, using LFI configured as store-only and load/store.

- Single-threaded, pinned to each of the core sizes.
- Performance delta normalized to unsandboxed decoder.
- Costs higher on A520 core due to in-order execution.

Pixel 9 Opus+LFI (higher is worse)





Performance: libdav1d

dav1d is an AV1 decoder:

LOTS of hand-written assembly.

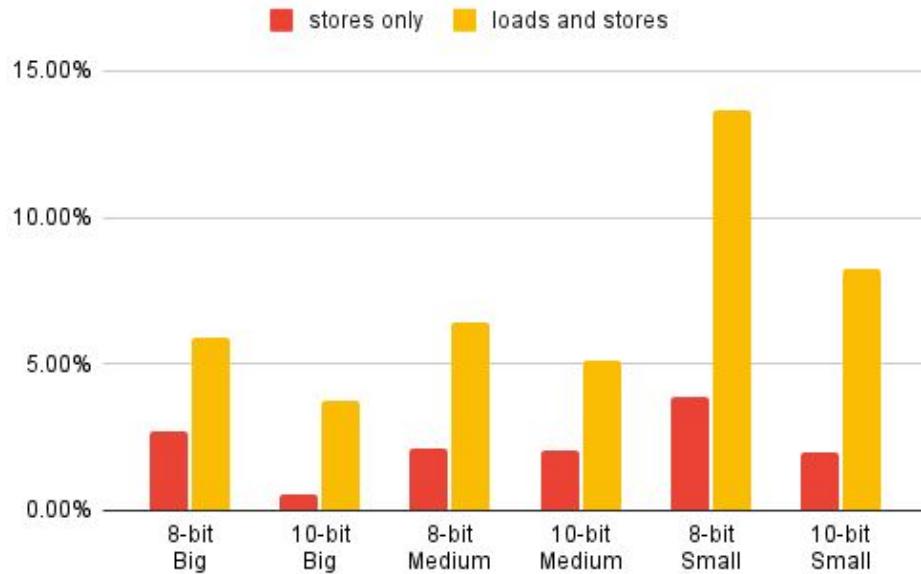
Totals grouped by language:

asm: 234001 (85.33%)

ansic: 40075 (14.61%)

sh: 161 (0.06%)

LFI dav1d - percent overhead



Performance critical, lots of asm: “stress test” for LFI!

Microbenchmark: Trampoline Overhead

Measured on Pixel 7

Platform	Cycles	Time
Function call (native)	2	0.7ns
Function call (LFI)	74	23ns
System call (native)	358	130ns
Context switch (native)	15770	6807ns

LFI: much better than inter-process communication (and less noisy).
Future: room for trampoline optimizations.

Ongoing Work



1. LFI-SpiderMonkey: Working on a secure production JIT engine.

Put the entire engine (C++ and generated code) in an LFI sandbox.

Update Gecko to use SpiderMonkey as a sandboxed library.

ium > Blink > JavaScript > Sandbox 361279118 > 40931165 > 338381304 ▾

V8 Sandbox Bypass: stack corruption due to parameter count mismatch

May 2, 2024 08:23AM

Qualcomm components

These vulnerabilities affect Qualcomm components and are described in further detail in the appropriate Qualcomm security bulletin or security alert. The severity assessment of these issues is provided directly by Qualcomm.

CVE	References	Severity	Subcomponent
CVE-2024-45569	A-377311993 QC-CR#3852339	Critical	WLAN
CVE-2024-45571	A-377313069 QC-CR#3834424	High	WLAN
CVE-2024-45582	A-377312377 QC-CR#3868093	High	Camera
CVE-2024-49832	A-377312238 QC-CR#3874301	High	Camera
CVE-2024-49833	A-377312639 QC-CR#3874372 [2] [3] [4]	High	Camera
CVE-2024-49834	A-377312055 QC-CR#3875406	High	Camera
CVE-2024-49839	A-377311997 QC-CR#385198	High	WLAN
CVE-2024-49843	A-377313194 QC-CR#3883522	High	Display

2. LFI Kernel Modules: Device driver isolation.

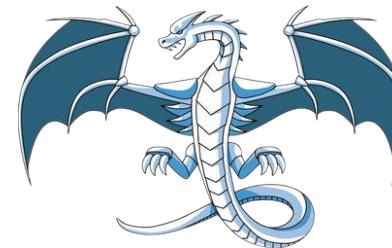
Conclusion

Lightweight Fault Isolation:

- Fast, Compatible, Simple, Secure.
- Coming soon to LLVM and Android!

How you can help!

- LFI Pilot Studies
- LFI FFI support: Java, Swift, Python,...
- LLVM implementation: optimizations, test suite, AArch64 MCInst info...



LFI

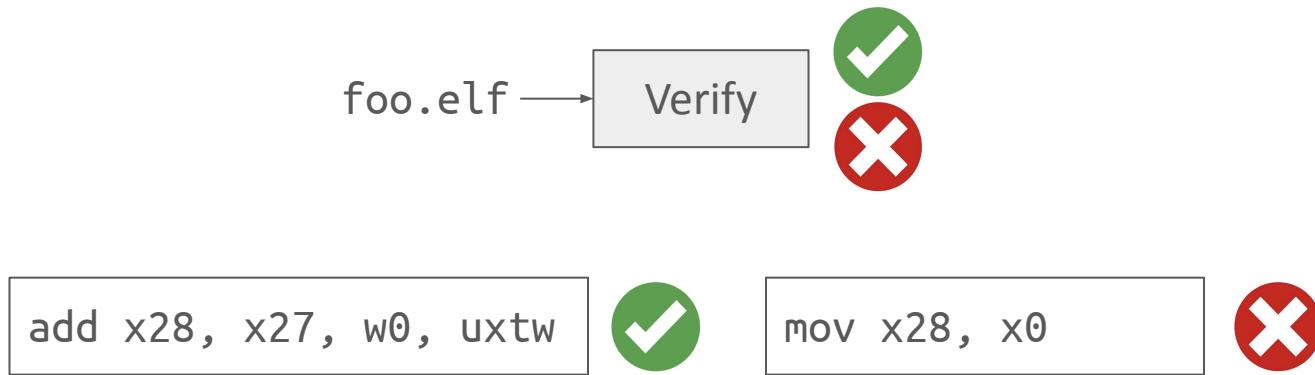
Check out RFC on Discourse!

<https://lfi-project.org>

Extra Slides

Simple Verification

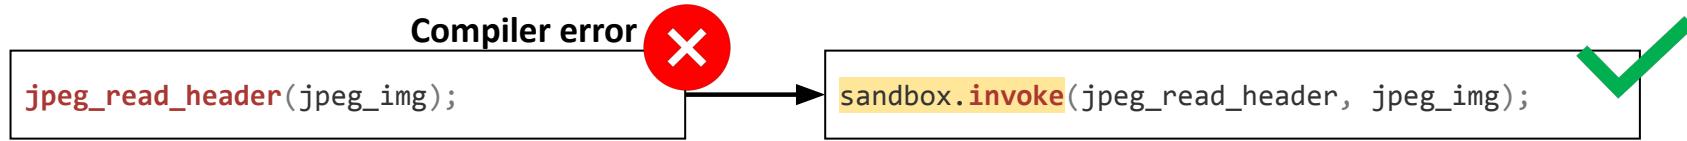
Enables: supply chain (build) safety, safe closed-source libraries, verifies rewriter.



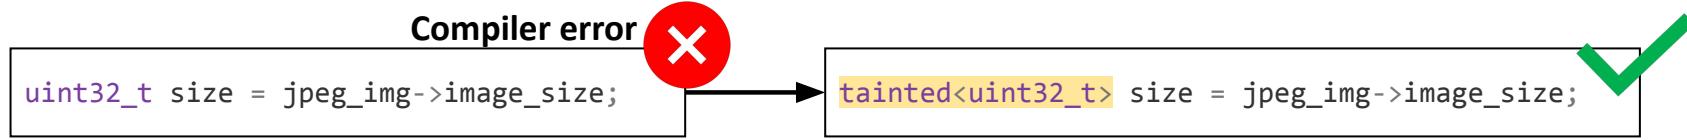
Fast (500+ MiB/s), Simple (~400 LoC), Easy to fuzz/formally verify

RLBox for Retrofitting Safely in Existing Code

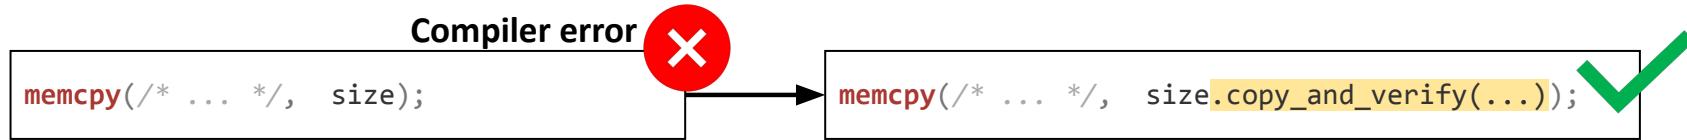
1. RLBox forces control flow to be explicit



2. RLBox forces data from the sandbox to be marked **tainted**



3. Tainted data must be checked before use



Android...

3.5 billion devices and growing

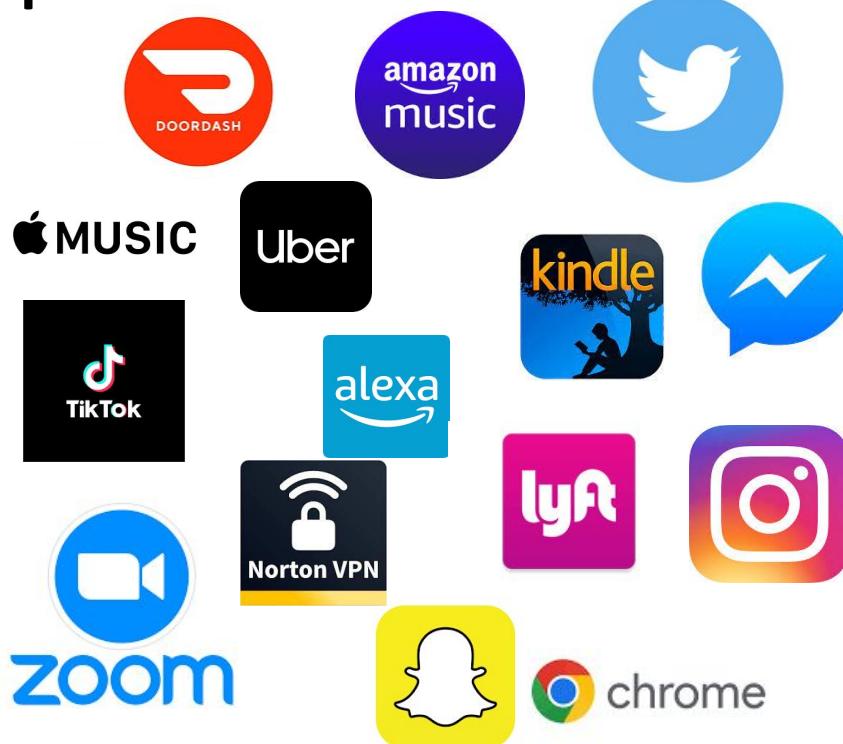
45% of global OS market

72% of global smartphone market (over 3B devices)

>45% of tablets



Lots of (Unpatched) third-party native libraries in popular Apps



Frequent vulnerabilities

App Name	Vul Lib Version	Vul Announced	TTRP (Days)	TTAF (Days)
Xbox	XML2-2.7.7	2014-11-04	12	1956
Apple Music	XML2-2.7.7	2014-11-04	12	1704
TikTok	GIFLib-5.1.1	2015-12-21	87	1429
Zoom Meetings	OpenSSL-1.0.0a	2010-08-17	91	1323
Amazon Alexa	OpenSSL-1.0.1s	2016-05-04	12	1086
Amazon Kindle	Libpng-1.6.34	2017-01-30	330	1019
StarMaker	FFmpeg-3.2	2016-12-23	4	1001
eBay	OpenCV-2.4.13	2017-08-06	41	905
Fitbit	SQLite3-3.20.1	2017-10-12	12	902
Uber	OpenCV-2.4.13	2017-08-06	41	830
Snapchat	SQLite3-3.20.1	2017-10-12	12	670
Discord	GIFLib-5.1.1	2015-12-21	87	665
Lyft	OpenCV-2.4.11	2017-08-06	41	662
Twitter	GIFLib-5.1.1	2015-12-21	87	457
Instagram	FFmpeg-2.8.0	2017-01-23	2	267

[Too Quiet in the Library: An Empirical Study of Security Updates in Android Apps' Native Code](#)