



Instruction Cost Modelling: Can we do better?

Sushant Gokhale (Nvidia), Madhur Amilkanthwar (Nvidia)

Presented by: Neil Hickey (Nvidia)

Agenda

- Introduction to Cost modelling
- What is it good at?
- What challenges does its implementation present?
- How could we extend it?

Introduction to cost modelling

Philosophy:

- Simple/ Quick approximation of LLVM-IR cost from the perspective of code gen
- High-level tunable heuristic or relative instruction costs mainly from the perspective of reciprocal throughput
- Used to guide a number of optimisations in the LLVM backend

```
for (BasicBlock &B : F) {  
    InstructionCost Cost(0);  
    for (Instruction &Inst : B) {  
        // Estimate intrinsic cost if the instruction is an intrinsic  
        if (auto *II = dyn_cast<IntrinsicInst>(&Inst)) {  
            IntrinsicCostAttributes ICA(II->getIntrinsicID(), *II);  
            Cost += TTI.getIntrinsicInstrCost(ICA, CostKind);  
        } else {  
            Cost += TTI.getInstructionCost(&Inst, CostKind);  
        }  
    }  
    OS << "Instruction Cost for BB: " << *Cost.getValue() << std::endl;  
}
```

Backend defined

- Implemented in the backend TargetTransformInfo class

```
switch (ISD) {  
default:  
    return BaseT::getArithmeticInstrCost(Opcode, Ty, CostKind, Op1Info, Op2Info);  
case ISD::SDIV:  
    if (Op2Info.isConstant() && Op2Info.isUniform() && Op2Info.isPowerOf2()) {  
        // On AArch64, scalar signed division by constants power-of-two are  
        // normally expanded to the sequence ADD + CMP + SELECT + SRA.  
        // The OperandValue properties many not be same as that of previous  
        // operation; conservatively assume OP_None.  
        InstructionCost Cost = getArithmeticInstrCost(  
            Instruction::Add, Ty, CostKind,  
            Op1Info.getNoProps(), Op2Info.getNoProps());  
        Cost += getArithmeticInstrCost(Instruction::Sub, Ty, CostKind,  
                                       Op1Info.getNoProps(), Op2Info.getNoProps());  
        Cost += getArithmeticInstrCost(  
            Instruction::Select, Ty, CostKind,  
            Op1Info.getNoProps(), Op2Info.getNoProps());  
        Cost += getArithmeticInstrCost(Instruction::AShr, Ty, CostKind,  
                                       Op1Info.getNoProps(), Op2Info.getNoProps());  
    }  
    return Cost;  
}
```

Cost modelling example


LoopVectoriser:

- Tries to decide, among different vector factors, what is the most work for minimal cost, comparing different strategies

```
// Load: Scalar load + broadcast
// Store: Scalar store + isLoopInvariantStoreValue ? 0 : extract
// FIXME: This cost is a significant under-estimate for tail folded
// memory ops.
const InstructionCost ScalarizationCost =
IsLegalToScalarize() ? getUniformMemOpCost(&I, VF)
: InstructionCost::getInvalid();

// Choose better solution for the current VF, Note that Invalid
// costs compare as maximumal large. If both are invalid, we get
// scalable invalid which signals a failure and a vectorization abort.
if (GatherScatterCost < ScalarizationCost)
    setWideningDecision(&I, VF, CM_GatherScatter, GatherScatterCost);
else
    setWideningDecision(&I, VF, CM_Scalarize, ScalarizationCost);
continue;
```

Pictorial depiction of cost calculation process



Huh, what is it good for?

Thankfully, not absolutely nothing:

- $\sim O(1)$ queries, not reliant on code gen
- Composable and overridable by different backends
- Common infrastructure shared across backends
- Used as a decision tool, doesn't have to be 100% accurate

What are the challenges

- Maintainability
- Heuristic based
- Low-level of granularity (low numbers, and integers)
- Used as a sum of IR costs.
- Happens at the IR level not at what the hardware actually executes

Maintainability

- Cost model spans 100,000+ lines across LLVM
 - LoopVectorize: 143 InstructionCost uses, 15+ cost functions
 - SLPVectorizer: 195 InstructionCost uses (26,691 lines)
 - 563 TTI method calls across all transforms
- Dual cost models in LoopVectorizer (lines 7079-7102)
 - Legacy cost model + VPlan cost model
 - Must be kept in sync with assertions
 - Comment: "TODO: Switch to more accurate costing based on VPlan"
- Vectorizer cost computations are universal/generic across a backend not specified per core
- LoopUnroll has its own summation logic, as does inliner, SLPVectorise and others
- Not a simple case of updating a single TTI hook value for a backend, but tracing usage and calculations through many lines of code

Heuristic Based

- Initially presented as a simple instruction cost metric to represent reciprocal throughput
- Has evolved through tuning to be purely a heuristic to get the desired output
- Often retroactively applied as lags behind codegen improvements
- Instructions given different costing based on optimization goals
- AArch64 reduction costs (Commit 828a867ee010, July 2025):
 - OR/XOR/AND v8i8: 15 → 5 (reduced by 10!)
 - OR/XOR/AND v16i8: 17 → 7 (reduced by 10!)
 - Reason: "Since the costs were added the codegen improved"
- AArch64 i1 reductions (Commit e79fac2968dc, June 2023):
 - v32i1: 91 → 3 (30x reduction!)
 - v64i1: 181 → 5 (36x reduction!)
 - v128i1: 362 → 9 (40x reduction!)
- AArch64 trip count threshold (Commit d945a2c9efda, Aug 2022):
 - Added `getMinTripCountTailFoldingThreshold()` to TTI interface
 - Reason: "I noticed regressions... as a solution I propose threshold"
 - Regression-driven, not hardware-based!

Add is 1

- All instructions represented as a multiple of an add instruction
- Doesn't allow finer-grained management of dispatch constraints and resource utilization
- RISCV percentage cost system (Commit 81efb825703c, July 2021):
 - Problem: RVC compressed = 0.7x, need fractional costs
 - Solution: Scale by 100 (RVI=100, RVC=70, QCEXT=150)
 - Reason: "RVC instructions... can take longer to execute... we consider that two RVC instruction are slightly more costly"
- Each target invents own scaling when integers insufficient

Sum of IR costs

- Just trivially sums up basic blocks by querying each IR instruction
- Doesn't model modern complex hardware

Can't capture complex backend fuses

Consider the output IR sequence as anticipated by the optimizer.

```
define i64 @fuse(i64 %x1, i64 %x2){  
entry:  
    %1 = mul i64 %x2, 8      ; x2 * 8  
    %2 = add i64 %x1, %1     ; x1 + (x2 * 8)  
    %addr = inttoptr i64 %2 to ptr  
    %4 = load i64, ptr %addr  
    ret i64 %4  
}
```

Fuse	
ldr	x0, [x0, x1, lsl #3]
ret	

Cost model overestimates cost of highlighted instructions that don't get materialized in final codegen !

```
Printing analysis 'Cost Model Analysis' for function 'fuse':  
Cost Model: Found an estimated cost of 1 for instruction: %0 = mul i64 %x2, 8  
Cost Model: Found an estimated cost of 1 for instruction: %1 = add i64 %x1, %0  
Cost Model: Found an estimated cost of 0 for instruction: %addr = inttoptr i64 %1 to ptr  
Cost Model: Found an estimated cost of 1 for instruction: %2 = load i64, ptr %addr, align 8
```

How can we improve things?

- Simplify the specification of instruction costs
- JSON configuration file simplifies experiments
- Type-aware overrides (i32, f32, v4i32, v2f64)
- Categories: arithmetic, memory, vector, compare_select, intrinsics
- Graceful fallback to default costs

Example JSON :


```
{  
  "instruction_costs": {  
    "arithmetic": {  
      "add": { "i32": 1, "i64": 1, "v4i32": 2 },  
      "mul": { "i32": 3, "i64": 4, "v4i32": 6 }  
    }  
  }  
}
```

Critical path counting

1. For a given IR sequence, lower to target assembly (virtual)
 - a) Use existing codegen infrastructure
 - b) Don't emit, just analyze
2. Apply target scheduling model to the basic block
 - a) Use existing MachineScheduler infrastructure
 - b) Already has: instruction latencies, resource usage, dependencies
3. Extract critical path through the BB
 - a) Not sum of all instructions (avoids ILP overcounting)
 - b) Critical path = longest dependency chain
 - c) Accounts for: instruction fusion, resource conflicts, parallelism
4. Return critical path length as cost
 - a) Maps naturally to reciprocal throughput
 - b) Already target-specific via scheduling model

Derive costs from actual assembly

- We envision a system as below



- Vision: Automatic cost derivation
- Input IR → Codegen → Measure assembly → Extract costs
- Quantum size: Basic Block
 - Large enough to capture instruction fusion, selection
 - Small enough to be reusable across programs
- Challenges to address:
 - 1. Assembly analysis complexity (fusion, macro-ops, uops)
 - 2. Mapping assembly back to IR constructs
 - 3. Handling context-dependent costs (register pressure, dependencies)
 - 4. Build time overhead
 - 5. Cache/storage for cost databases

