


Modular

BILLY ZHU
CHRIS LATTNER

OCT 28, 2025
LLVM DEV MEETING

Building Modern Language Frontends with MLIR

Lessons from Mojo's
Compile-Time
Meta-Programming



Agenda

-
- 01 Brief Mojo Primer
 - 02 Meta-Programming in Mojo
 - 03 MLIR Parametric Code Representation
 - 04 Parameter Domain Computations
 - 05 What Does This Mean in Practice?
-



Brief Mojo Primer

Modular

Why Mojo 🔥 ?

Requirement: Needed a [portable systems](#) programming language for [heterogeneous](#) compute


Goal: unlock the full power of the hardware by giving [programmers](#):

- low level control
- zero-cost abstractions
- safety guarantees

Approach: Willing to [invest](#) a lot to get the best quality, developer UX, and performance

- Mojo is a [real language](#), not a Python eDSL

```
def mandelbrot_kernel[
    width: Int # SIMD Width
](c: ComplexSIMD[float, width]) ->
SIMD[int, width]:  
  
    """A vectorized implementation of  
    the inner mandelbrot computation."""  
  
    z = ComplexSIMD[float, width](0, 0)
    iters = SIMD[DType.index, width](0)
    mask = SIMD[DType.bool, width](True)  
  
    for i in range(MAX_ITERS):
        if not any(mask):
            break  
  
        mask = z.squared_norm() <= 4
        iters = mask.select(iters + 1, iters)
        z = z.squared_add(c)  
  
    return iters
```



Modular

Blackwell Throughput vs vLLM

Note: We use vllm 0.9.0.1 as a reference in some cases cause vllm 0.10.0 crashes with memory errors due to flash infer.

153.4%

Llama3.1 - 405b arxiv-summarization
NVIDIA - B200



152.7%

Llama3.1 - 405b code_debug
NVIDIA - B200



104.3%

Llama3.1 - 405b sonnet-prefill-heavy-prefix-low
NVIDIA - B200



*Modular v25.4 vs vLLM v0.10.0

*Modular v25.4 vs vLLM v0.9.0.1

*Modular v25.4 vs vLLM v0.9.0.1

Modular

AMD MI355 Throughput vs vLLM

171%

Improved throughput

Gemma-3-27B
Sonnet Decode Heavy

Modular 25.6

AMD - MI355X

15.557 qps

vLLM 0.10.1.1

AMD - MI355X

9.921 qps

*Using the latest
docker container with
the optimized flags

Modular

How? Meta-Programming

Mojo is a “predictable” language + compiler:

- Does not embed AI or HW knowledge into the compiler itself
- Eschews “compiler magic”

Metaprogramming enables powerful features

- Specialization, codegen, autotuning, etc
- Empower Mojo programmers with “generalization” power

Mojo is cool for other reasons too:

- Powerful type system
- Fancy MLIR compiler implementation
- See many previous llvmdev talks, e.g.

“Mojo🔥: A system programming language for heterogeneous computing” at LLVM Dev 2023 [[link](#)]

```
fn fill_fib(size: Int) -> List[Int]:  
    if size <= 0:  
        return []  
    if size == 1:  
        return [0]  
    var fib: List[Int] = [0, 1]  
    for idx in range(2, size):  
        fib.append(fib[idx-2] + fib[idx-1])  
    return fib^  
  
fn main():  
    # List computed at compile-time.  
    alias a6 = fill_fib(6)  
  
    # Unrolled by iterating over List.  
    @parameter  
    for elem in a6:  
        print(elem)  
  
    # Same code computed at run-time.  
    var v6 = fill_fib(6)  
    for elem in v6:  
        print(elem)
```

C++ already does this!?

Template Meta-Programming is the existing standard in GPU programming

Terrible error messages due to “duck typing”

- Errors show up as instantiation problems

Different comp-time and runtime languages

- Hard to learn & awkward to use

Limited type system:

- Cannot use computed strings or trees ergonomically

Slow compile times, can't use debuggers on comptime code, ...

```
template <class LayoutA,
          class Offset,
          class LayoutB>
CUTE_HOST_DEVICE constexpr
auto
composition(LayoutA const& layoutA,
            Offset const& offset,
            LayoutB const& layoutB)
{
    return ComposedLayout<LayoutA, Offset, LayoutB>{layoutA, offset,
layoutB};
}

template <class A, class O, class B, class Tiler>
CUTE_HOST_DEVICE constexpr
auto
composition(ComposedLayout<A,O,B> const& a,
            Tiler const& b)
{
    return composition(a.layout_a(), a.offset(),
composition(a.layout_b(), b));
}

template <class ShapeA, class StrideA,
          class A, class O, class B>
CUTE_HOST_DEVICE constexpr
auto
composition(Layout<ShapeA,StrideA> const& a,
            ComposedLayout<A,O,B> const& b)
{
    CUTE_STATIC_ASSERT_V(b.offset() == Int<0>{}, "Require offset == 0.");

    return composition(composition(a, b.layout_a()), b.layout_b());
}
```



Meta-Programming in Mojo 🔥

Modular

Code vs Meta-Code

Normal functions have “*arguments*”:

- These are runtime values that can vary on every invocation

Mojo allows [comptime arguments]:

- Mojo instantiates these – like a template
- We refer to these as “*parameters*”

Parameters get superpowers:

- e.g. can fully unroll a loop

Expressions are written the same way in both roles

```
# Naive power function.  
fn pow(base: Int, exp: Int) -> Int:  
    res = 1  
    for i in range(exp):  
        res *= base  
    return res  
  
# Easily refactor when `exp` can be known  
# at compile time.  
fn pow_exp[exp: Int](base: Int) -> Int:  
    res = 1  
    for i in range(exp):  
        res *= base  
    return res  
  
# Can also unroll the loop at compile time.  
fn pow_exp_fast[exp: Int](base: Int) -> Int:  
    res = 1  
    @parameter  
    for i in range(exp):  
        res *= base  
    return res
```

Anything can be parameterized... by anything



Can parameterize:

- Functions, Types, Closures, etc.
- Even *arbitrary expressions*

Enables powerful user-defined libraries:

- Very important for GPU programming

Simplifies and unifies the language:

- Types are just comptime values
- Removes many special case features in other languages

```
# Complex custom datatype.
struct Layout:
    var shape: List[Int]    # has malloc
    var stride: List[Int]

    @staticmethod
    fn col_major(shape: List[Int]) -> Layout:
        return ...

struct LayoutTensor[
    dtype: DType,
    layout: Layout, # Used at comptime
    ...
]:

    # Methods can add more parameters too.
    fn load[width: Int](self, *idx: Int)
        -> SIMD[dtype, width]:
            return ...
```

How does this work?

Explicit control over expression evaluation:

- “alias” guarantees comp-time eval
- “var” is dynamic at its execution time

Meta-code is code run at comp-time

- Executed by an IR interpreter
- Supports ~arbitrary logic, including malloc

Can materialize finished values to runtime

- Shifting work to comp-time

Can also compute types in the type system!

- Mojo has powerful “dependent types”

```
# Returns a heap-allocated List.
fn fill_fib(size: Int) -> List[Int]:
    if size <= 0:
        return []
    if size == 1:
        return [0]
    var fib: List[Int] = [0, 1]
    for idx in range(2, size):
        fib.append(fib[idx-2] + fib[idx-1])
    return fib^

fn main():
    # List computed at compile-time.
    alias a6 = fill_fib(6)

    # Unrolled by iterating over List.
    @parameter
    for elem in a6:
        print(elem)

    # Another list computed at run-time.
    var v6 = fill_fib(6)
    for elem in v6:
        print(elem)
```

Polymorphic Generic IR:

Type checked *without instantiating*

What we want:

- Maintain target info for split compilation
- Good error messages, not template stack traces
- Fast compiles + don't ship source code


Type check + Generate IR **before instantiating**:

- Type checking is symbolic instead of concrete

How do we represent uninstantiated "templates" polymorphically in IR??

```
fn exp2(x: SIMD[dtype, simd_width]) -> SIMD[dtype, simd_width]:  
  
    @parameter  
    if is_nvidia_gpu():  
        @parameter  
        if dtype is DType.float16:  
            @parameter  
            if _is_sm_9x_or_newer():  
                return _call_ptx_intrinsic[  
                    scalar_instruction="ex2.approx.f16",  
                    vector2_instruction="ex2.approx.f16x2",  
                    scalar_constraints="=h,h",  
                    vector_constraints="=r,r",  
                ](x)  
            else:  
                return _call_ptx_intrinsic[  
                    instruction="ex2.approx.f16",  
                    constraints="=h,h",  
                ](x)  
        elif dtype is DType.bfloat16 and _is_sm_9x_or_newer():  
            return _call_ptx_intrinsic[  
                scalar_instruction="ex2.approx.ftz.bf16",  
                vector2_instruction="ex2.approx.ftz.bf16x2",  
                scalar_constraints="=h,h",  
                vector_constraints="=r,r",  
            ](x)  
        elif dtype is DType.float32:  
            return _call_ptx_intrinsic[  
                instruction="ex2.approx.ftz.f32",  
                constraints="=f,f",  
            ](x)
```

MLIR Parametric Code Representation



Modular

Mojo in MLIR

Uses MLIR to express manipulation of MLIR itself.

The parameterized “**base IR**” forms the basic IR building blocks. It carries application logic:

- Parameterized Types
- Parameterized Ops

The “**meta IR**” manipulates the “base IR”:

- Parametric Declarations
- Inline Instantiations

```
Mojo
fn simd_sum
  [dtype: DType, size: Int]
  (value: SIMD[dtype, size])
-> Scalar[dtype]:
  var sum: Scalar[dtype] = 0
  @parameter
  for i in range(size):
    sum += value[i]
  return sum
```

```
MLIR
lit.fn @simd_sum
<dtype: @DType, size: @Int>
(%value: @SIMD<dtype, size>
 -> @SIMD<dtype, 1> {
  ???
}
```

Parameterized Types

Parameterize

- Dig “typed holes” 

Instantiate

- Fill holes with values of the correct type 

Traditional Types

- Expect **specific kinds of Attributes** (e.g. IntegerAttr)

Parametric Ops

- Expect **attributes of a specific Type**

Traditional Type

```
def POP_SIMDType {  
    let parameters = (ins  
        IndexAttr:$size,  
        POP_DTypeConstantAttr:$dtype  
    )  
}
```

E.g.

```
%arr = pop SIMD.add(%s0, %s1)  
      : !pop SIMD<4, f32>
```

Parametric Type

```
def POP_SIMDType {  
    let parameters = (ins  
        TypedAttr:$size,  
        TypedAttr:$dtype  
    )  
}
```

E.g.

```
%arr = pop SIMD.add(%s0, %s1)  
      : !pop SIMD< : index, : dtype>
```

  `builtin.integer<2> : index`
  `kgen.param.ref<"size"> : index`
  `... : ...`

Parameterized Ops

Op Attribute/Properties can be parameterized too

A key part of programmatic codegen

Traditional Ops

- Expect **specific kinds of Attributes** (e.g. IntegerAttr)

Parametric Ops

- Expect **attributes of a specific Type**

Traditional Op

```
let arguments = (ins
  POP_SIMDType:$vector,
  IndexAttr:$idx
);

pop.simd.extract %arr {
  idx = IntegerAttr<3> : index
}
```

Parametric Op

```
let arguments = (ins
  POP_SIMDType:$vector,
  IndexTypedAttr:$idx
);

pop.simd.extract %arr {
  idx = builtin.integer<2> : index
}
  ↗
  ① builtin.integer<2> : index
  ② kgen.param.ref<"size"> : index
  ③ ...
```

Parametric Declarations

Declarations (e.g. functions, struct types) encapsulate a block of parametric IR over a set of input parameters

- The input parameters are typed and named
- The body of the declaration can refer to the input parameters

Enables reuse of the parametric body with different input parameter values

```
Parametric Function
lit.fn @simd_sum
<size: Int, dtype: DType>
(
    %value: !pop SIMD<
        #kgen.param.ref<"size">,
        #kgen.param.ref<"dtype">
    >
)
-> !pop SIMD<
    1,
    #kgen.param.ref<"dtype">
>
{
    ...
}

lit.fn @main {
    ...
    %s2 = pop SIMD.splat %scalar
        : !pop SIMD<2, f32>
    %s4 = pop SIMD.splat %scalar
        : !pop SIMD<4, f32>
    lit.call @simd_sum<2, f32>(%s2)
    lit.call @simd_sum<4, f32>(%s4)
}
```

Inline Instantiations

Instantiate a block of parametric IR inline

Parameter-If

- Select one of the blocks to instantiate based on some parameter value

Parameter-For

- Instantiate a block multiple times
- Declares an input parameter representing the value returned by an iterator
- The body will be instantiated once with each iterated value

Conditional Instantiation

```
example.before
kgen.param.if <condition> {
    example.do_something
} else {
    example.do_something_else
}
example.after
```

E.g. when `condition` is False

```
example.before
example.do_something_else
example.after
```

Repeated Instantiation

```
example.before
kgen.param.for i in iterable ... {
    example.do_something { attr = i }
}
example.after
```

E.g. when `iterable` is range(2)

```
example.before
example.do_something { attr = 0 }
example.do_something { attr = 1 }
example.after
```

Putting them together ...



Parametric IR

```
lit.fn @simd_sum
<size: @Int, dtype: @DType>
(%value: @SIMD<size, dtype>)
-> @SIMD<1, dtype> {
...
kgen.param.for iter in ... {
...
    kgen.param.call @other_fn<size>()
...
    kgen.param.if iter {
        ...
        my.param.op { dt = dtype }
        ...
    }
    kgen.deferred {
        name = "hw_specific",
        attrs = {...} }
}
%c = kgen.param.constant = <size>
...
}
```

Parameterized Ops

```
lit.fn @simd_sum
<size: @Int, dtype: @DType>
(%value: @SIMD< ,      >)
    -> @SIMD< ,      > {
    ...
    kgen.param.for iter in ... {
        ...
        kgen.param.call @other_fn<      >()
        ...
        kgen.param.if      {
            ...
            my.param.op { dt =      }
            ...
        }
        kgen.deferred {
            name = "hw.specifc",
            attrs = {      } }
    }
    %c = kgen.param.constant = <      >
    ...
}
```

Parameters

```
size: @Int  dtype: @DType
      size  dtype
      1  dtype
      iter
      size
      iter
      dtype
      ...
      size
```

MLIR enables a natural interleave of ...

Meta IR

Base IR

Parameterized Ops

```
lit.fn @simd_sum
<size: @Int, dtype: @DType>
(%value: @SIMD< ,      >)
    -> @SIMD< ,      > {
    ...
    kgen.param.for iter in ... {
        ...
        kgen.param.call @other_fn<      >()
        ...
        kgen.param.if      {
            ...
            my.param.op { dt =      }
            ...
        }
        kgen.deferred {
            name = "hw_specific",
            attrs = {      } }
    }
    %c = kgen.param.constant = <      >
    ...
}
```

Parameters

```
size: @Int  dtype: @DType
      size  dtype
      1   dtype
      iter
      size
      iter
      dtype
      ...
      size
```


Modular

Meta IR

Op Scopes

```
lit.fn @simd_sum
  <size: @Int, dtype: @DType>
    (%value: @SIMD< , >)
      -> @SIMD< , > {
        ...
        kgen.param.for iter in ... {
          ...
          kgen.param.call @other_fn<     >()
          ...
          kgen.param.if      {
            ...
            my.param.op { dt =       }
            ...
          }
          kgen.deferred {
            name = "hw_specific",
            attrs = {   } }
        }
        %c = kgen.param.constant = <     >
        ...
      }
```

Parameter Scopes




Modular

Parameter Domain Computations

8

Modular



Consequences

What are the consequences of allowing computation on parameters?

Type-checking requires determining type equality

Recall that types can be parameterized.

- Dependent Types - types depending on non-type values

Type equality requires parameter equality

When are two parameter expressions “equal”?

Different ways to write the same Type:

```
!my_dialect.array<3 + 1>
!my_dialect.array<2 * 2>
!my_dialect.array<pow(2, 2)>
...
```

Evaluation?

Evaluation is the “ultimate” judge of identity.

- Interpret the parameter expression until an irreducible value is reached

Cons:

- Slow
- Requires fully concrete expressions, which requires instantiation

Different ways to write the same Type:

```
!my_dialect.array<3 + 1>
!my_dialect.array<2 * 2>
!my_dialect.array<pow(2, 2)>
...
```

All Evaluate to:

```
!my_dialect.array<4>
```

Un-evaluatable expressions:

```
!my_dialect.array<x + y>
!my_dialect.array<y + x>
!my_dialect.array<x + 0 + y * 1>
```

Canonicalization

Canonicalize expressions into a “normal” form.

- Goal: All expressions that represent the same computation canonicalize into the same form

Requires a powerful expression [canonicalizer](#)

Complement with user-annotation to bridge any gaps

Different ways to write the same Type:

```
!my_dialect.array<x + y>
!my_dialect.array<y + x>
!my_dialect.array<x + 0 + y * 1>
...
```

All Canonicalize to:

```
!my_dialect.array<x + y>
```

Parameter Representation

Expressions in Mojo have two possible IR representations:

- As an Op:
 - Produces ssa values
- As a Typed Attribute:
 - Produces parameter values

Parameter expressions exist in attribute form

Pros:

- Parameter equality reduces to pointer equality
- Smaller memory footprint & CoW
- Concise inline representation

Source

```
(4 - 2) * 3
```

Op Repr.

```
%0 = index.constant 4 : index
%1 = index.constant 2 : index
%2 = index.constant 3 : index
%3 = index.sub %0, %1 : index
%4 = index.mul %3, %2 : index
```

Attribute Repr.

```
      MulNode
      /   \
SubNode   ConstNode(3)
 /   \
ConstNode(4) ConstNode(2)
```

Function References

Attribute expressions can still invoke user-declared functions

This completes the loop:

- Meta-Code is Code
- Code is Meta-Code

Functions can choose to participate in canonicalization, or stay symbolic

Given

```
fn pow(base: Int, exp: Int) -> Int:
```

...

Source

```
(4 - 2) * pow(2, 3)
```

Attribute Repr.

```
      MulNode
      /   \
SubNode   CallNode(@pow, 2, 3)
 /   \
ConstNode(4) ConstNode(2)
```

What Does This Mean in Practice?



Modular

Unified Libraries

Massively flattens learning curve

Write rich & expressive libraries to express meta-programming logic

Directly test library APIs at runtime:

- Deal with runtime failures instead of compile-time failures
- Debuggers work!

```
# Comptime layout algebra
alias TileType[*tile_sizes: Int] =
    LayoutTensor[
        dtype,
        Self._compute_tile_layout[
            *tile_sizes,
            ()[0],
            origin,
            address_space=address_space,
            element_layout=element_layout,
            layout_int_type=layout_int_type,
            linear_idx_type=linear_idx_type,
            masked = masked or
                _tile_is_masked[layout, *tile_sizes](),
            alignment=alignment,
        ]
]

fn tile[
    *tile_sizes: Int
](self, *tile_coords: Int) ->
    self.TileType[*tile_sizes]:

    alias num_tiles =
        stdlib.builtin.variadic_size(tile_sizes)
    alias _tiled_layout =
        Self._compute_tile_layout[*tile_sizes]()
    var offset = 0
    var runtime_shape =
        tile_type.RuntimeLayoutType.ShapeType()
    var runtime_stride =
        tile_type.RuntimeLayoutType.StrideType()
    ...
}
```

Opt-In Static Specialization

Choose how much of your data structure is static vs dynamic

- Shift info from dynamic to static for more safety guarantees
- Shift info from static to dynamic for quick prototyping and dynamic behavior

Remember: You can reuse the same libraries as you shift data between the variable-domain and the parameter-domain

More Dynamic

```
# Dynamic size & dynamic elements
struct List:
    var size: Int
    var elems: UnsafePointer[Int]
```

```
# Static size & dynamic elements
struct SizedList[size: Int]:
    var elems: UnsafePointer[Int]
```

```
# Dynamic size & static element
struct SplatList[elem: Int]:
    var size: Int
```

```
# Static size & static elements
struct StaticList[size: Int, *elems: Int]:
    # No fields.
```

More Static

Modular

What do you prefer?

Types checked at instantiation time

```
Building CUDA object
examples/46_depthwise_simt_conv2dfprop/CMakeFiles/46_depthwise
_simt_conv2dfprop.dir/depthwise_simt_conv2dfprop.cu.o
/home/clattner/cutlass/include/cutlass/conv/kernel/direct_convolutio
n.h(95): error: incomplete type is not allowed detected during:
instantiation of class
"cutlass::conv::kernel::DirectConvolutionParams<Mma_, Epilogue_,
ThreadblockSwizzle_, ConvOperator, Arguments_,
ConvOutputIteratorParameter_, ConvProblemSize_, GroupMode_,
ThreadBlockOutputShape_> [with
Mma_=cutlass::conv::threadblock::DepthwiseFpropDirectConvMultipleSt
age<ThreadblockShape,
cutlass::conv::threadblock::DepthwiseFpropActivationDirect2dConvTileAc
cessIteratorFixedStrideDilation<cutlass::MatrixShape<64, 64>,
ThreadBlockOutputShape, StrideShape, DilationShape,
cutlass::conv::TensorNHWCType<1, 10, 10, 64>, ElementInputA,
LayoutInputA,
cutlass::transform::PitchLinearStripminedThreadMap<cutlass::PitchLinear
Shape<64, 100>, 128, 4>, cutlass::AlignedArray<ElementInputA, 4, 16>>,
cutlass::transform::threadblock::RegularTileAccessIteratorDirectConv<cut
lass::MatrixShape<100, 64>, ElementInputA, cutlass::layout::RowMajor, 0,
cutlass::transform::PitchLinearStripminedThreadMap<cutlass::PitchLinear
Shape<64, 100>, 128, 4>, false, 16>, cutlass::arch::CacheOperation::Global,
cutlass::conv::threadblock::DepthwiseFpropFilterDirectConvTileAccessIter
```

Types checked without instantiating

```
fn process_prime[x: Int where is_prime(x)]():
  ...
fn main():
  alias value = get_int()

# `value` not always prime.
process_prime[value]()

@parameter
if is_prime(value):
  # This is OK.
  process_prime[value]()
```

```
$ mojo test.mojo
test.mojo:9:25: error: invalid call to 'process_prime':
unable to satisfy constraint
  process_prime[value]()
  ~~~~~^
test.mojo:1:31: note: constraint declared here
fn process_prime[x: Int where is_prime(x)]():
```

All made
possible by the
power of MLIR



Thank You

Modular
MODULAR

