

# Optimizing Flang's Optimizer

Leandro Lupori

# Flang - LLVM's Fortran Front-end

- Flang lowers Fortran to HLFIR (High-Level Fortran IR) + other MLIR dialects

- Flang


Fortran → HLFIR → FIR → MLIR

- MLIR/LLVM

MLIR → LLVM IR → Executable

# Flang - Fortran to HLFIR

```
function fact(x) result(res)
integer :: x, res
if (x <= 1) then
    res = 1
else
    res = x * fact(x - 1)
endif
end function
```



```
func.func @_QPfact(%arg0: !fir.ref<i32> {fir.bindc_name = "x"}) -> i32 {
    %1 = fir.alloca i32 {bindc_name = "res", uniq_name = "_QFFactEres"}
    %2:2 = hlfir.declare %1 {uniq_name = "_QFFactEres"} :
        (!fir.ref<i32>) -> (!fir.ref<i32>, !fir.ref<i32>)
    %3:2 = hlfir.declare %arg0 {uniq_name = "_QFFactEx"} :
        (!fir.ref<i32>) -> (!fir.ref<i32>, !fir.ref<i32>)
    %4 = fir.load %3#0 : !fir.ref<i32>
    %c1_i32 = arith.constant 1 : i32
    %5 = arith.cmpi sle, %4, %c1_i32 : i32
    fir.if %5 {
        hlfir.assign %c1_i32 to %2#0 : i32, !fir.ref<i32>
    } else {
        %7 = fir.load %3#0 : !fir.ref<i32>
        %9 = arith.subi %7, %c1_i32 : i32
        %10:3 = hlfir.associate %9 {adapt.valuebyref} :
            (i32) -> (!fir.ref<i32>, !fir.ref<i32>, i1)
        %11 = fir.call @_QPfact(%10#0) fastmath<contract> :
            (!fir.ref<i32>) -> i32
        %12 = arith.muli %7, %11 : i32
        hlfir.assign %12 to %2#0 : i32, !fir.ref<i32>
        hlfir.end_associate %10#1, %10#2 : !fir.ref<i32>, i1
    }
    %6 = fir.load %2#0 : !fir.ref<i32>
    return %6 : i32
}
```


# HLFIR/FIR Optimizations

- HLFIR preserves many details of Fortran constructs
  - For instance, variable attributes and array operations are preserved
  - Other details, not useful for optimizations or other transformations are discarded, ex:
    - The distinction between function and subroutine
- Optimizations that rely on Fortran knowledge occur in HLFIR and FIR stages
  - Fortran intrinsic functions may be replaced with inline code
  - Expressions may be simplified

# Evaluating Flang Performance

- Setup
  - SPEC CPU® 2017 (rate, 1 copy)
    - Only benchmarks that have Fortran
  - Armv9.0 - Neoverse V2
  - Flags: -O3 -flto
- Goal
  - Flang should be no more than 10% slower than GFortran on any SPEC CPU® 2017 benchmark

# LLVM 20 vs GCC 14



# Profiling 527.cam4\_r

- Perf was used to identify the functions that accounted for the most execution time
  - Ubuntu 24.04 LTS
  - perf record --cpu 1 -e cycles/period=280000000/u
  - Perf ran on cpu 0 and benchmarks on cpu 1, to minimize interference
  - 10 samples per second - CPU frequency 2.8GHz
- The performance issues were spread throughout the program

# Profiling 527.cam4\_r

## perf report -n

- Compare using samples - overhead is relative to execution time, that differs

Overhead	Samples	Command	Shared Object	Symbol
8.11%	168	cam4_r_base.llv	libm.so.6	[.] log@@GLIBC_2.29
6.90%	143	cam4_r_base.llv	libm.so.6	[.] pow@@GLIBC_2.29
6.42%	133	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMradswPradcswmx
6.23%	129	cam4_r_base.llv	libm.so.6	[.] exp@@GLIBC_2.29
4.93%	102	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMae_r_rad_propsPaer_rad_props_sw
3.62%	75	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMstratiformPstratiform_tend
3.43%	71	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMraadePradabs
3.19%	66	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMae_r_rad_propsPget_hydro_rad_props
2.80%	58	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMzm_convPientropy
2.32%	48	cam4_r_base.llv	cam4_r_base.llvm20	[.] _QMphysics_typesPphysics_update

LLVM 20

Overhead	Samples	Command	Shared Object	Symbol
8.08%	140	cam4_r_base.gcc	libm.so.6	[.] log@@GLIBC_2.29
6.46%	112	cam4_r_base.gcc	libm.so.6	[.] pow@@GLIBC_2.29
5.83%	101	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __radae_MOD_radabs.constprop.0
5.65%	98	cam4_r_base.gcc	libm.so.6	[.] exp@@GLIBC_2.29
4.10%	71	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __radsw_MOD_radcswmx.isra.0
3.87%	67	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __aer_rad_props_MOD_get_hydro_rad_props.isra.0
3.69%	64	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __zm_conv_MOD_ientropy.isra.0
3.29%	57	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __mapz_module_MOD_ppm2m.lto_priv.0
3.12%	54	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __aer_rad_props_MOD_aer_rad_props_sw.constprop.0.isra.0
2.48%	43	cam4_r_base.gcc	cam4_r_base.gcc14	[.] __cldwat_MOD_pcond.constprop.0

GCC 14

# Profiling 527.cam4\_r

## perf annotate

```
000000000326e80 <_QMaer_rad_propsPaer_rad_props_sw>:  
    stp    x29, x30, [sp, #-96]!  
    stp    x28, x27, [sp, #16]  
    mov    x29, sp  
    stp    x26, x25, [sp, #32]  
    stp    x24, x23, [sp, #48]  
    stp    x22, x21, [sp, #64]  
    stp    x20, x19, [sp, #80]  
    sub    sp, sp, #0x12, lsl #12  
    sub    sp, sp, #0x500
```

LLVM 20

- No debug info...
- Unlike GCC, LLVM sometimes loses debug info when LTO is enabled
- Disabling LTO helps in mapping assembly instructions to Fortran constructs
- Many performance issues occur even with LTO disabled

# Profiling 527.cam4\_r

perf annotate - LTO disabled

```
|     tau    (:,:,:,:) = -100._r8
398: stp    x10, x10, [x11]
      subs   x9, x9, #0x1
      stp    x10, x10, [x11, #16]
      stp    x10, x10, [x11, #32]
      stp    x10, x10, [x11, #48]
      stp    x10, x10, [x11, #64]
      stp    x10, x10, [x11, #80]
      stp    x10, x10, [x11, #96]
      stp    x10, x10, [x11, #112]
```

LLVM 20

```
|     add    x4, x22, x4
1c0: mov    x0, x3
      tau    (:,:,:,:) = -100._r8
1c4: stp    q31, q31, [x0], #32
      cmp    x1, x0
      ↑ b.ne  1c4
      add    x1, x1, #0x360
      add    x3, x3, #0x360
      cmp    x4, x1
      ↑ b.ne  1c0
```

3.70

GCC 14

# Profiling 527.cam4\_r

perf annotate - LTO disabled

- The LTO version used *stp q0, q0, [x11, #<offs>]* instead  
(X registers are 64-bit wide while Q registers are 128-bit)
- Both versions used many *stp* instructions and some *str* instructions
- LLVM didn't seem to realize that a single loop could be used to initialize the 3-dimensional array *tau*

# Digging in

- To make analysis easier a test program with a non-optimizable assignment to a 3D array (in a loop) was used

HLFIR:

```
func.func private @_QFPinit_tau() attributes {fir.host_symbol = @_QQmain, llvm.linkage = #llvm.linkage<internal>} {
    ...
    %10 = fir.shape %c4, %c27, %c19 : (index, index, index) -> !fir.shape<3>
    %11:2 = hlfir.declare %9(%10) {uniq_name = "_QFEtau"} : (!fir.ref<!fir.array<4x27x19xf64>>, !fir.shape<3>) ->
        (!fir.ref<!fir.array<4x27x19xf64>>, !fir.ref<!fir.array<4x27x19xf64>>)
    %cst = arith.constant -1.000000e+02 : f64
    hlfir.assign %cst to %11#0 : f64, !fir.ref<!fir.array<4x27x19xf64>>
    return
}
```

# Digging in

FIR:

```
func.func private @_QFPinit_tau() attributes {fir.host_symbol = @_QQmain, llvm.linkage = #llvm.linkage<internal>} {
    ...
    %9 = fir.shape %c4, %c27, %c19 : (index, index, index) -> !fir.shape<3>
    %10 = fir.declare %8(%9) {uniq_name = "_QFEtau"} : (!fir.ref<!fir.array<4x27x19xf64>>, !fir.shape<3>) ->
        !fir.ref<!fir.array<4x27x19xf64>>
    fir.do_loop %arg0 = %c1 to %c19 step %c1 unordered {
        fir.do_loop %arg1 = %c1 to %c27 step %c1 unordered {
            fir.do_loop %arg2 = %c1 to %c4 step %c1 unordered {
                %11 = fir.array_coor %10(%9) %arg2, %arg1, %arg0 :
                    (!fir.ref<!fir.array<4x27x19xf64>>, !fir.shape<3>, index, index, index) -> !fir.ref<f64>
                fir.store %cst to %11 : !fir.ref<f64>
            }
        }
    }
    return
}
```

# Digging in

LLVM:

```
.preheader2.i:                                ; preds = %.preheader2.i.preheader, %.preheader2.i
    %33 = phi i64 [ %34, %.preheader2.i ], [ 1, %.preheader2.i.preheader ]
    %.idx.i = mul nuw nsw i64 %33, 864
    %gep5.i = getelementptr i8, ptr getelementptr (i8, ptr @_QFEtau, i64 -904), i64 %.idx.i
    %gep.i = getelementptr i8, ptr %gep5.i, i64 40
    store double -1.000000e+02, ptr %gep.i, align 16, !tbaa !8
    %gep.1.i = getelementptr i8, ptr %gep5.i, i64 48
    store double -1.000000e+02, ptr %gep.1.i, align 8, !tbaa !8
    %gep.2.i = getelementptr i8, ptr %gep5.i, i64 56
    store double -1.000000e+02, ptr %gep.2.i, align 16, !tbaa !8
    ...
```

- The 2 inner loops were completely unrolled
- But this is a contiguous array
- Maybe LLVM can produce a more optimized IR if Flang lowers multidimensional array assignments using a single loop?

# Multidimensional array linearization

- [flang] Optimize assignments of multidimensional arrays [#146408](#)
- flang/lib/Optimizer/HLFIR/Transforms/OptimizedBufferization.cpp

```
llvm::LogicalResult BroadcastAssignBufferization::matchAndRewrite(
    hlfir::AssignOp assign, mlir::PatternRewriter &rewriter) const {
...
if (lhs.isSimplyContiguous() && extents.size() > 1) {
    // Flatten the array to use a single assign loop, that can be better optimized.
    ...
    hlfir::LoopNest loopNest =
        hlfir::genLoopNest(loc, builder, flatExtents, /*isUnordered=*/true,
                           flangomp::shouldUseWorkshareLowering(assign));
    builder.setInsertionPointToStart(loopNest.body);
    mlir::Value arrayElement =
        builder.create<hlfir::DesignateOp>(loc, fir::ReferenceType::get(eleTy),
                                             flatArray, loopNest.oneBasedIndices);
    builder.create<hlfir::AssignOp>(loc, rhs, arrayElement);
}
...
}
```

# Multidimensional array linearization

FIR:

```
func.func private @_QFPinit_tau() attributes {fir.host_symbol = @_QQmain, llvm.linkage = #llvm.linkage<internal>} {
    ...
    %10 = fir.shape %c4, %c27, %c19 : (index, index, index) -> !fir.shape<3>
    %11 = fir.declare %9(%10) {uniq_name = "_QFETau"} : (!fir.ref<!fir.array<4x27x19xf64>>, !fir.shape<3>) ->
        !fir.ref<!fir.array<4x27x19xf64>>
    %12 = fir.convert %11 : (!fir.ref<!fir.array<4x27x19xf64>>) -> !fir.ref<!fir.array<2052xf64>>
    fir.do_loop %arg0 = %c1 to %c2052 step %c1 unordered {
        %c2052_0 = arith.constant 2052 : index
        %13 = fir.shape %c2052_0 : (index) -> !fir.shape<1>
        %14 = fir.array_coor %12(%13) %arg0 : (!fir.ref<!fir.array<2052xf64>>, !fir.shape<1>, index) -> !fir.ref<f64>
        fir.store %cst to %14 : !fir.ref<f64>
    }
    return
}
```

# Multidimensional array linearization

LLVM:

```
vector.ph:                                ; preds = %27, %_QFPinit_tau.exit
%32 = phi i64 [ %36, %_QFPinit_tau.exit ], [ %30, %27 ]
br label %vector.body, !dbg !69

vector.body:                                ; preds = %vector.body, %vector.ph
%index = phi i64 [ 0, %vector.ph ], [ %index.next, %vector.body ]
%33 = getelementptr double, ptr @_QFEtau, i64 %index, !dbg !69
%34 = getelementptr i8, ptr %33, i64 16, !dbg !69
store <2 x double> splat (double -1.000000e+02), ptr %33, align 16, !dbg !69, !tbaa !73
store <2 x double> splat (double -1.000000e+02), ptr %34, align 16, !dbg !69, !tbaa !73
%index.next = add nuw i64 %index, 4
%35 = icmp eq i64 %index.next, 2052, !dbg !69
br i1 %35, label %_QFPinit_tau.exit, label %vector.body, !dbg !69, !llvm.loop !80

_QFPinit_tau.exit:                          ; preds = %vector.body
%36 = add nsw i64 %32, -1, !dbg !68
%37 = icmp sgt i64 %32, 1, !dbg !68
br i1 %37, label %vector.ph, label %_.crit_edge, !dbg !68
```

# Multidimensional array linearization

AArch64:

```
mov x8, #-4586634745500139520
mov x9, x0
dup v0.2d, x8
adrp x8, _QFEtau+16
add x8, x8, :lo12:_QFEtau+16
mov x10, x8
mov w11, #2052
.LBB0_7:
subs x11, x11, #4
stp q0, q0, [x10, #-16]
add x10, x10, #32
b.ne .LBB0_7
```

- Test program

```
328:    mov    x8, #0xc059000000000000          // #-4586634745500139520
        mov    w9, #0x804                      // #2052
        dup    v0.2d, x8
        add    x8, x20, #0x10
338:    subs   x9, x9, #0x4
        stp    q0, q0, [x8, #-16]
        add    x8, x8, #0x20
        b.ne   338
```

- 527.cam4\_r

# Results

## Multidimensional array linearization

- In the end, there was practically no difference in 527.cam4\_r
  - aer\_rad\_props\_sw is a large function
  - It spent less than 10% of its time in array initialization
  - There were other performance issues in aer\_rad\_props\_sw

# Results



## Multidimensional array linearization

- However, there was a substantial performance improvement in other test programs
- The number of instructions per array initialization was reduced from 38 to 8

Array size	Speedup
16KB	8%
64KB	28%
128KB	31%
256KB	32%
512KB+	0%


# Profiling 527.cam4\_r (again)

Using flame graphs to find performance differences



# Profiling 527.cam4\_r (again)

waccmx\_is



# Profiling 527.cam4\_r (again)

- **waccmx\_is = trim(name) == trim(wacmx\_opt)**

```
0.90      → bl    _FortranATrim
0.90      → bl    _FortranATrim
...
9.91      cmp   w12, w13
           cset  w12, hi // hi = pmore
...
0.90      → bl    free@plt
...
→ bl    free@plt
```

LLVM

```
waccmx_is = (trim(name) == trim(wacmx_opt))
add  x3, x3, #0x80
mov  x1, x2
mov  x2, #0x10                      // #16
→ bl  _gfortran_compare_string@plt
  cmp  w0, #0x0
end function wacmx_is
  cset w0, eq // eq = none
  ldp  x29, x30, [sp], #16
← ret
```

GCC 15

# Fortran relational operations

- Fortran character strings with fixed length are padded with blanks on the right
- Fortran 2018 language specification  
10.1.5.5.1 Interpretation of relational intrinsic operations

*“For a character relational intrinsic operation, the operands are compared one character at a time in order, beginning with the first character of each character operand. If the operands are of unequal length, the shorter operand is treated as if it were extended on the right with blanks to the length of the longer operand.”*
- Conclusion: the calls to trim() are unnecessary and can be removed

# HLFIR Expression Simplification

- The first step was to add the hlfir.char\_trim operation
  - This made it easier to identify and remove calls to trim()
  - When removal was not possible the operation was converted to a runtime call
- With it a pass to simplify the comparison of characters was written

```
trim(x) == trim(y) → x == y
```

- This optimization shouldn't cause performance regressions
  - It is unlikely to preclude others which bring more significant gains

# HLFIR Expression Simplification

- flang/lib/Optimizer/HLFIR/Transforms/ExpressionSimplification.cpp

```
llvm::LogicalResult
matchAndRewrite(hlfir::CharTrimOp trimOp,
               mlir::PatternRewriter &rewriter) const override {
    int trimUses = std::distance(trimOp->use_begin(), trimOp->use_end());
    auto cmpCharOp = getFirstUser<hlfir::CmpCharOp>(trimOp);
    auto destroyOp = getLastUser<hlfir::DestroyOp>(trimOp);
    if (!cmpCharOp || !destroyOp || trimUses != 2)
        return rewriter.notifyMatchFailure(
            trimOp, "hlfir.char_trim is not used (only) by hlfir.cmpchar");
    rewriter.eraseOp(destroyOp);
    rewriter.replaceOp(trimOp, trimOp.getChr());
    return mlir::success();
}
```

# HLFIR Expression Simplification

HLFIR:

```
%7 = hlfir.char_trim %4#0 : (!fir.boxchar<1>) -> !hlfir.expr<!fir.char<1,?>>
%8 = hlfir.char_trim %6#0 : (!fir.boxchar<1>) -> !hlfir.expr<!fir.char<1,?>>
%9 = hlfir.getCmpchar eq %7 %8 : (!hlfir.expr<!fir.char<1,?>>, !hlfir.expr<!fir.char<1,?>>) -> i1
%10 = fir.convert %9 : (i1) -> !fir.logical<4>
hlfir.assign %10 to %2#0 : !fir.logical<4>, !fir.ref<!fir.logical<4>>
hlfir.destroy %8 : !hlfir.expr<!fir.char<1,?>>
hlfir.destroy %7 : !hlfir.expr<!fir.char<1,?>>
```

Optimized HLFIR:

```
%7 = hlfir.getCmpchar eq %4#0 %6#0 : (!fir.boxchar<1>, !fir.boxchar<1>) -> i1
%8 = fir.convert %7 : (i1) -> !fir.logical<4>
hlfir.assign %8 to %2#0 : !fir.logical<4>, !fir.ref<!fir.logical<4>>
```

# Results

## HLFIR Expression Simplification

- 527.cam4\_r time went from 217 to 208.8 seconds
- Estimated 3.5% speedup  
(considering margin of error and variability)

# Function Inlining

- Challenge: different compilers inline different functions
- Considering only the time spent on *radcswmx*, it looks like GCC is twice as slow
- However, a call graph or flame graph shows this is mostly due to *raddedmx* being inlined by GCC but not by LLVM

Children	Self	Command	Shared Object	Symbol
- 9.52%	3.23%	cam4_r_base.llv	cam4_r_base.llvm_trim0	[.] _QMradswPradcswmx
- 6.30%		_QMradswPradcswmx		
+ 3.25%		_QMradswPraddedmx		
0.91%		_FortranAAssign		
0.61%		_memset_zva64		
+ 3.23%		_start		

LLVM


Children	Self	Command	Shared Object	Symbol
- 11.28%	6.73%	cam4_r_base.gcc	cam4_r_base.gcc15_o3_g	[.] __radsw_MOD_radcswmx
+ 6.73%		_start		
- 4.55%		__radsw_MOD_radcswmx		
3.14%		exp@@GLIBC_2.29		
0.51%		_memset_zva64		

GCC 15

# Community Contributions

- Other developers also contributed with Flang performance enhancements
- Considering single thread performance only, these are some recent examples:
  - [\[flang\] Simplify hlfir.index in a few limited cases](#)
  - [\[flang\]\[driver\] Accelerate complex division when -ffast-math is specified](#)
  - [\[flang\] Create TBAA subtree for COMMON block variables](#)
  - [\[flang\] Lower hlfir.cmpchar into inline implementation in simplify-hlfir-intrinsics](#)
  - [\[flang\] Inline hlfir.eoshift during HLFIR intrinsics simplification](#)
  - [\[flang\]\[runtime\] Optimize Descriptor::FixedStride\(\)](#)
  - [\[flang\]\[runtime\] Speed up initialization & destruction](#)
  - [\[flang\] Optimize redundant array repacking](#)

# LLVM vs GCC



# Future Work

- Improve the performance of 548.exchange2\_r
- Check Flang performance on SPEC CPU® v8
- Test other compiler flags: -fno-lto, -mcpu=neoverse-v2
- Check Flang OpenMP performance with SPEC speed

linaro

linaro

linaro

linaro

Thank you

leandro.lupori@linaro.org



linaro.org