



# Automating the search aspects of compiler engineering: IR for auto-tuning & beyond



Rolf Morel

[<rolf.morel@intel.com>](mailto:rolf.morel@intel.com)

LLVM Dev Meeting – Oct 2025, Santa Clara




## Premise:

Compiler engineering involves a lot tweaking of pipelines and knobs

An average plan for an average day:

1. Out of all possible optimizations, first try F. Evaluate. If perf isn't there...
2. Then tweak F's knobs. Evaluating each time.
3. Try optimization G. Evaluate.
4. Then tweak G's knobs. Evaluating each time.
5. Now try F;G. Tweak knobs. Evaluating each time.
6. Then G;F. Tweak knobs. Evaluating each time.




## Promise:

Allow for ***expressing such plans*** and ***mechanize their execution***



An average plan for an average day:

1. Out of all possible optimizations, first try optimization F. Evaluate. If perf isn't there...
2. Then tweak F's knobs. Evaluating each time.
3. Try optimization G. Evaluate.
4. Then tweak G's knobs. Evaluating each time.
5. Now try F;G. Tweak knobs. Evaluating each time.
6. Then G;F. Tweak knobs. Evaluating each time.



In this talk,

***expressing such plans*** via

- IR – an extension of *schedule* IR, i.e. pipelines as programs
  - which gives representation to the different choices we are considering

***mechanize their execution*** via

- Mapping such schedules-with-choices to constraint problems
  - so, we can hand off the combinatorial search to solvers



# Agenda

1. RECAP:
  - I. **Schedules**: an alternative to pipelines
  - II. Rewriting MLIR with its **Transform dialect**
2. IR for tuning knobs ... with joint constraints
3. Non-deterministic semantics via SMT  
... and regaining determinism through rewrites
4. IR for alternative transform sequences
5. Mechanized schedule/pipeline optimization



# RECAP: What are schedules?



**program = algorithm + schedule**



# Image processing à la Halide

## algorithm

```
blurx(x,y) = in(x-1,y)
              + in(x,y)
              + in(x+1,y)
out(x,y) = blurx(x,y-1)
           + blurx(x,y)
           + blurx(x,y+1)
```



## program

```
par for out.yo in 0..out.y.extent/4
  for out.xo in 0..out.x.extent/4
    alloc blurx[blurx.y.extent][blurx.x.extent]
    for out.yi in 0..4
      let blurx.y.min = 4*out.yo.min + out.yi.min - 1
      for blurx.y in blurx.y.min..blurx.y.max
        for blurx.xo in blurx.x.min/4..blurx.x.max/4
          vec for blurx.xi in 0..4
            blurx[blurx.y.stride*blurx.y+...] =
              in[in.y.stride*(blurx.y.min+blurx.y)
                  +4*blurx.xo+ramp(4)] + ...
          vec for out.xi in 0..4
            out[out.y.stride*(4*(out.yo-out.yo.min)+out.yi)+...] =
              blurx[blurx.y.stride*(out.yi-1-blurx.y.min)
                  + out.xi - blurx.x.min] + ...
```




# Image processing à la Halide

## algorithm (high-level code)

```
blurx(x,y) = in(x-1,y)  
            + in(x,y)  
            + in(x+1,y)  
  
out(x,y) = blurx(x,y-1)  
            + blurx(x,y)  
            + blurx(x,y+1)
```

## schedule

```
blurx: split x by 4 →  $x_o$ ,  $x_i$   
        vectorize:  $x_i$   
        store at out. $x_o$   
        compute at out. $y_i$   
  
out: split x by 4 →  $x_o$ ,  $x_i$   
     split y by 4 →  $y_o$ ,  $y_i$   
     reorder:  $y_o$ ,  $x_o$ ,  $y_i$ ,  $x_i$   
     parallelize:  $y_o$   
     vectorize:  $x_i$ 
```



## program (low-level & optimized code)

```
par for out.yo in 0..out.y.extent/4  
  for out.xo in 0..out.x.extent/4  
    alloc blurx[blurx.y.extent][blurx.x.extent]  
    for out.yi in 0..4  
      let blurx.y.min = 4*out.yo.min + out.yi.min - 1  
      for blurx.y in blurx.y.min..blurx.y.max  
        for blurx.xo in blurx.x.min/4..blurx.x.max/4  
          vec for blurx.xi in 0..4  
            blurx[blurx.y.stride*blurx.y+...] =  
              in[in.y.stride*(blurx.y.min+blurx.y)  
                  +4*blurx.xo+ramp(4)] + ...  
          vec for out.xi in 0..4  
            out[out.y.stride*(4*(out.yo-out.yo.min)+out.yi)+...] =  
              blurx[blurx.y.stride*(out.yi-1-blurx.y.min)  
                  + out.xi - blurx.x.min] + ...
```

# Many scheduling DSLs

... in support of  
optimizing BLAS / tensor programs:


LIFT & RISE  
(Steuwer et al, 2015 & 2022)

Exo  
(Ikarashi et al, 2022)

```
A = t.placeholder((1024, 1024))
B = t.placeholder((1024, 1024))
k = t.reduce_axis((0, 1024))
C = t.compute((1024, 1024), lambda y, x:
              t.sum(A[k, y] * B[k, x], axis=k))
s = t.create_schedule(C.op)
```

Fireiron  
(Hagedorn et al, 2020)

Tiramisu  
(Baghdadi et al, 2019)



TVM (Chen et al, 2018)

All share the same *core idea*:  
declarative descriptions of how to transform code

# **RECAP:**

# **What is MLIR's Transform Dialect?**




# MLIR's Transform meta-Dialect

*transform IR* describes transformations on *payload IR*

**payload IR:**

```
func.func @gemv(%α, %A, %x, %β, %y) {  
    %βy = linalg.generic attrs { iter_types = ["par"] } ... {  
        %βy_elem = arith.mulf %β_elem, %y_elem  
        linalg.yield %βy_elem : f32  
    } -> tensor<?xf32>  
    ...  
    %αAx_plus_βy = linalg.generic  
        { iter_types = ["par", "reduction"] } ... {  
            ...  
        } -> tensor<x?xf32>  
    func.return %αAx_plus_βy : tensor<x?xf32>  
}
```



**Transformed payload IR:**


```
func.func @gemv(%α, %A, %x, %β, %y) {  
    %dim = tensor.dim %y, %c0  
    %βy = scf.for %i = %c0 to %dim step %c4 ... {  
        %ex_slice = tensor.extract_slice ...  
        ... = linalg.generic attrs {iter_types = ["par"]} {  
            %βy_slice_elem = arith.mulf %β_elem, %y_slice_elem  
            linalg.yield %βy_slice_elem : f32  
        } -> tensor<?xf32>  
        %in_slice = tensor.insert_slice ...  
        scf.yield %in_slice  
    }  
    ...  
    %αAx_plus_βy = linalg.generic  
        { iter_types = ["par", "reduction"] } ... {  
            ...  
        } -> tensor<x?xf32>  
    func.return %αAx_plus_βy : tensor<x?xf32>  
}
```

**transform IR:**


```
transform.sequence(%func: !transform.any_op) {  
    %elemwise = transform.match attrs { iter_types = ["par"] } %func  
    transform.tile_using_for %elemwise [4]  
}
```




# Progressive lowering/optimization through the dialects




# Progressive lowering/optimization through the dialects




# Progressive lowering/optimization through the dialects




# Progressive lowering/optimization through the dialects



# Progressive lowering/optimization through the dialects, schedule by schedule



# Progressive lowering/optimization through the dialects, via an overall schedule



```
transform.sequence {  
    // @lower_to_linalg  
    ...  
    %1 = transform.apply_registered_pass  
        "tosa-to-linalg-pipeline" to %0  
    ...  
    // @tile_outer_dims  
    ...  
    %2 = transform.tile_using_for %1 [4]  
    ...  
    // @vectorize  
    ...  
    %3 = transform.vectorize %2 [4, 8]  
    ...  
}
```

.mlir  
kernel as tosa-program

Transform  
interpreter

.mlir  
vector/affine/scf  
-program

# What about tuning schedules?



# Tuning schedules: selecting tile sizes

```
transform.sequence(%mod) {  
    ...  
    %target = transform.match "linalg.matmul" %mod  
    %tiled, %cache_loops:2 = transform.tile_using_for tile_sizes [16,32] %target  
    ...  
    %tiled2, %reg_loops:2 = transform.tile_using_for tile_sizes [4,6] %tiled  
    ...  
}
```

.mlir



# Tuning schedules: selecting tile sizes

```
transform.sequence(%mod) {  
    ...  
    %target = transform.match "linalg.matmul" %mod  
    %tiled, %cache_loops:2 = transform.tile_using_for tile_sizes [32,32] %target  
    ...  
    %tiled2, %reg_loops:2 = transform.tile_using_for tile_sizes [4,6] %tiled  
    ...  
}
```

.mlir



# Tuning schedules: selecting tile sizes

```
transform.sequence(%mod) {  
    ...  
    %target = transform.match "linalg.matmul" %mod  
    %tiled, %cache_loops:2 = transform.tile_using_for tile_sizes [32,64] %target  
    ...  
    %tiled2, %reg_loops:2 = transform.tile_using_for tile_sizes [5,5] %tiled  
    ...  
}
```

.mlir



# Schedules with choices: “*reifying*” knobs

```
transform.sequence(%mod) {  
    ...  
    %target = transform.match "linalg.matmul" %mod  
    %m_cache = transform.tune.knob<"m_cache"> options = [16, 32, 64]  
    %n_cache = transform.tune.knob<"n_cache"> options = [32, 48, 64, 80]  
    %tiled, %cache_loops:2 = transform.tile_using_for tile_sizes [%m_cache,%n_cache] %target  
    ...  
    %m_n_reg = transform.tune.knob<"m_n_reg"> options = [[4,6],[5,5],[3,8],[2,12]]  
    %m_reg, %n_reg = transform.param.split_handle %m_n_reg  
    %tiled2, %reg_loops:2 = transform.tile_using_for tile_sizes [%m_reg,%n_reg] %tiled  
    ...  
}
```


From which we can automatically extract the parameter space:

$$\begin{aligned} \text{m\_cache} &\in \{ 16, 32, 64 \} \\ \text{n\_cache} &\in \{ 32, 48, 64, 80 \} \\ \text{m\_n\_reg} &\in \{ [4,6], [5,5], [3,8], [2,12] \} \end{aligned}$$



# Rewriting knobs: regaining determinism

```
%m_cache = transform.tune.knob<"m_cache"> options = [16, 32, 64] .mlir
%n_cache = transform.tune.knob<"n_cache"> options = [32, 48, 64, 80]
%tiled, %cache_loops:2 = transform.tile_using_for tile_sizes [%m_cache,%n_cache] %target
...
%m_n_reg = transform.tune.knob<"m_n_reg"> options = [[4,6],[5,5],[3,8],[2,12]]
%m_reg, %n_reg = transform.param.split_handle %m_n_reg
%tiled2, %reg_loops:2 = transform.tile_using_for tile_sizes [%m_reg,%n_reg] %tiled
```




```
%m_cache = transform.param.constant 32 .mlir
%n_cache = transform.param.constant 64
%tiled, %cache_loops:2 = transform.tile_using_for tile_sizes [%m_cache,%n_cache] %target
...
%m_n_reg = transform.param.constant [3,8]
%m_reg, %n_reg = transform.param.split_handle %m_n_reg
%tiled2, %reg_loops:2 = transform.tile_using_for tile_sizes [%m_reg,%n_reg] %tiled
```


# Non-deterministic semantics via SMT




# Non-deterministic semantics via SMT-dialect



# Non-deterministic semantics via SMT-dialect



# Non-deterministic semantics via SMT-dialect



# Tuning schedules: joint choices via SMT

Can hand-code the known valid combinations:

```
...  
%m_n_reg = transform.tune.knob<"m_n_reg"> options = [[4,6],[5,5],[3,8],[2,12]]  
%m_reg, %n_reg = transform.param.split_handle %m_n_reg  
...
```

Or can express actual constraint, e.g.  $m_{\text{reg}} * n_{\text{reg}} \leq 25$ , over independent variables:

```
...  
%m_reg = transform.tune.knob<"m_reg"> options = range<2,12> : !transform.param<i64>  
%n_reg = transform.tune.knob<"n_reg"> options = range<2,12> : !transform.param<i64>  
transform.smt.constrain_params(%m_reg, %n_reg) {  
  ^bb(%m_var: !smt.int, %n_var: !smt.int):  
    %m_times_n = smt.int.mul %m_var, %n_var  
    %c25 = smt.int.constant 25  
    %prod_le_25 = smt.int.cmp le %m_times_n, %c25  
    smt.assert %prod_le_25  
}  
...
```



# Tuning schedules: joint choices as SMT

lower / project

```
%m_reg = transform.tune.knob<“m_reg”> .mlir
options = range<2,12> : !transform.param<i64>
%n_reg = transform.tune.knob<“n_reg”>
options = range<2,12> : !transform.param<i64>
transform.smt.constrain_params(%m_reg, %n_reg) {
^bb(%m_var: !smt.int, %n_var: !smt.int):
  %m_times_n = smt.int.mul %m_var, %n_var
  %c25 = smt.int.constant 25
  %prod_le_25 = smt.int.cmp le %m_times_n %c25
  smt.assert %prod_le_25
}
```


```
%m_reg = smt.declare_const : !smt.int .mlir
%m_ge_2 = smt.int.cmp le %c2, %m_reg
%m_le_12 = smt.int.cmp le %m_reg, %c12
%m_reg_options = smt.and %m_ge_2, %m_le_12
smt.assert %m_reg_options
%n_reg = smt.declare_const : !smt.int
...
%n_reg_options = smt.and %n_ge_2, %n_le_12
smt.assert %n_reg_options
%m_times_n = smt.int.mul %m_reg, %n_reg
%c25 = smt.int.constant 25
%mn_le_25 = smt.int.cmp le %m_times_n %c25
smt.assert %prod_le_25
```

```
(declare-const m_reg Int) .smtlib
(declare-const n_reg Int)
(assert (and (≤ 2 m_reg) (≤ m_reg 12)))
(assert (and (≤ 2 n_reg) (≤ n_reg 12)))
(assert (≤ (* m_reg n_reg) 25))
```


convert



# Tuning schedules: joint choices as SMT



# Tuning schedules: joint choices as SMT



# Beyond tuning knobs: alternative transform sequences



# Choice among transforms: parallel or seq

```
...A...
%tiled, %forall_loop =
  transform.tile_using_forall
    tile_sizes [%m,%n] %target
...B...
```

.mlir

or

```
...A...
%tiled, %for_loops:2 =
  transform.tile_using_for
    tile_sizes [%m,%n] %target
...B...
```

.mlir

Parameter/knob tuning does *not* allow for expressing choice between different transforms!



# Choice among transforms: parallel or seq

```
...A...  
%tiled, %forall_loop =  
  transform.tile_using_forall  
    tile_sizes [%m,%n] %target  
...B...
```

.mlir

```
...A...  
%tiled, %for_loops:2 =  
  transform.tile_using_for  
    tile_sizes [%m,%n] %target  
...B...
```

.mlir

or

*can be represented by*

```
...A...  
%tiled_target, %loop = transform.tune.alternatives<“par_or_seq”> {  
  %tiled, %forall_loop = transform.tile_using_forall tile_sizes [%m,%n] %target  
  transform.yield %tiled, %forall_loop  
}, {  
  %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [%m,%n] %target  
  transform.yield %tiled, %for_loops#0  
}  
...B...
```

.mlir

*which rewrites to*

```
...A...  
%tiled_target, %loop =  
  transform.tile_using_forall  
    tile_sizes [%m,%n] %target  
...B...
```

.mlir

Or

```
...A...  
%tiled_target, %loop, %_ =  
  transform.tile_using_for  
    tile_sizes [%m,%n] %target  
...B...
```

.mlir



# Sub-schedule choice: mapping to SMT

lower / project

```
.mlir
%res = transform.tune.alternatives<"subsched"> {
  ...subsched0...
  transform.yield %subsched0_result
}, {
  ...subsched1...
  transform.yield %subsched1_result
}, {
  ...
}, {
  ...subschedN...
  transform.yield %subschedN_result
}
```

```
%subsched = smt.declare_const : !smt.int .mlir
%subsched_lb = smt.int.cmp le %c0, %subsched
%subsched_ub = smt.int.cmp le %subsched, %cN
%subsched_options = smt.and %subsched_lb,
                     %subsched_ub
smt.assert %subsched_options

%subsched_region0 = smt.eq %subsched, %c0
// SMT-ops for subsched0
// ... conditional on %subsched == 0

%subsched_region1 = smt.eq %subsched, %c1
// SMT-ops for subsched1
// ... conditional on %subsched == 1
...
%subsched_regionN = smt.eq %subsched, %cN
// SMT-ops for subschedN
// ... conditional on %subsched == N
```



# Sub-schedule choice: path dependence

```
%tiled_target, %loop = transform.tune.alternatives<“par_or_seq”> {  
    %x = transform.tune.knob<“x”> options = [4,8,16]  
    %tiled, %forall_loop = transform.tile_using_forall tile_sizes [%x,4] %target  
    transform.yield %tiled, %forall_loop  
, {  
    %y = transform.tune.knob<“y”> options = [2,4]  
    %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [8,%y] %target  
    transform.yield %tiled, %for_loops#0  
}
```

lower  
/project

```
%x = smt.declare_const : !smt.int  
...  
%x_options = smt.or %x_eq_4, %x_eq_8, %x_eq_16
```

```
smt.assert %x_options  
%y = smt.declare_const : !smt.int  
...  
%y_options = smt.or %y_eq_2, %y_eq_4
```

```
smt.assert %y_options
```

# Sub-schedule choice: path dependence

```
%tiled_target, %loop = transform.tune.alternatives<“par_or_seq”> {  
    %x = transform.tune.knob<“x”> options = [4,8,16]  
    %tiled, %forall_loop = transform.tile_using_forall tile_sizes [%x,4] %target  
    transform.yield %tiled, %forall_loop  
, {  
    %y = transform.tune.knob<“y”> options = [2,4]  
    %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [8,%y] %target  
    transform.yield %tiled, %for_loops#0  
}
```

lower  
/project

```
%par_or_seq = smt.declare_const : !smt.int  
%lb = smt.int.cmp ge %par_or_seq, %c0  
%ub = smt.int.cmp le %par_or_seq, %c1  
%par_or_seq_options = smt.and %lb, %ub  
smt.assert %par_or_seq_options  
  
%par_or_seq_region0 = smt.eq %par_or_seq, %c0  
  
%par_or_seq_region1 = smt.eq %par_or_seq, %c1
```

```
%x = smt.declare_const : !smt.int  
...  
%x_options = smt.or %x_eq_4, %x_eq_8, %x_eq_16  
%conditional_x_options = smt.implies %par_or_seq_region0,  
                        %x_options  
smt.assert %x_options %conditional_x_options  
%y = smt.declare_const : !smt.int  
...  
%y_options = smt.or %y_eq_2, %y_eq_4  
%conditional_y_options = smt.implies %par_or_seq_region1,  
                        %y_options  
smt.assert %y_options %conditional_y_options
```

# Tied choices across sub-schedules

```
%tiled_target, %forall_loop =  
  transform.tile_using_forall  
    tile_sizes [4,8] %target  
...  
// do something with %tiled_target  
...  
transform.loop.forall_to_parallel %forall_loop
```

.mlir

*or*

```
%tiled_target, %loops:2 =  
  transform.tile_using_for  
    tile_sizes [4,8] %target  
...  
// do same something with %tiled_target  
...  
transform.loop.coalesce %loops#0
```

.mlir



# Tied choices across sub-schedules

```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop, %marker = transform.tune.alternatives<"par_or_seq"> {  
    %tiled, %forall_loop = transform.tile_using_forall tile_sizes [4,8] %target  
    transform.yield %tiled, %forall_loop, %marked_par  
, {  
    %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [4,8] %target  
    transform.yield %tiled, %for_loops#0, %marked_seq  
}  
...  
// do something with %tiled_target  
...  
transform.tune.alternatives<"par_or_seq_reprise"> {  
    transform.param.cmpi eq %marker, %marked_par // Crash if not equal!!!  
    transform.loop.forall_to_parallel %loop  
    transform.yield  
, {  
    transform.param.cmpi eq %marker, %marked_seq // Crash if not equal!!!  
    transform.loop.coalesce %loop  
    transform.yield  
}
```

.mlir



# Tied choices across sub-schedules

```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop, %marker = transform.tune.alternatives<"par_or_seq"> {  
    %tiled, %forall_loop = transform.tile_using_forall tile_sizes [4,8] %target  
    transform.yield %tiled, %forall_loop, %marked_par  
}, {  
    %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [4,8] %target  
    transform.yield %tiled, %for_loops#0, %marked_seq  
}  
... // do something with %tiled_target  
transform.tune.alternatives<"par_or_seq_reprise"> {  
    transform.param.cmpi eq %marker, %marked_par  
    transform.loop.foreach_to_parallel %loop  
    transform.yield  
}, {  
    transform.param.cmpi eq %marker, %marked_seq  
    transform.loop.coalesce %loop  
    transform.yield  
}
```

lower  
/project

```
%marked_par = smt.int.constant 1  
%marked_seq = smt.int.constant 0  
%par_or_seq = smt.declare_const : !smt.int  
...  
smt.assert %par_or_seq_options  
%par_or_seq_region0 = smt.eq %par_or_seq, %c0  
%par_or_seq_region1 = smt.eq %par_or_seq, %c1  
%marker = smt.declare_const : !smt.int  
%marker_eq_marked_par = smt.eq %marker, %marked_par  
%par_or_seq_region0_marking = smt.implies %par_or_seq_region0,  
                                %marker_eq_marked_par  
smt.assert %region0_marking  
%marker_eq_marked_seq = smt.eq %marker, %marked_seq  
%par_or_seq_region1_marking = smt.implies %par_or_seq_region1,  
                                %marker_eq_marked_seq  
smt.assert %par_or_seq_region1_marking  
%par_or_seq_reprise = smt.declare_const : !smt.int  
...  
smt.assert %par_or_seq_reprise_options  
%par_or_seq_reprise_region0 = smt.eq %par_or_seq_reprise, %c0  
%par_or_seq_reprise_region1 = smt.eq %par_or_seq_reprise, %c1  
%reprise_region0_cond = smt.implies %par_or_seq_reprise_region0,  
                        %marker_eq_marked_par  
smt.assert %reprise_region0_cond  
%reprise_region1_cond = smt.implies %par_or_seq_reprise_region1,  
                        %marker_eq_marked_seq  
smt.assert %reprise_region1_cond
```

```
(declare-const par_or_seq Int) .smtlib  
(declare-const marker Int)  
(declare-const par_or_seq_reprise Int)  
...  
(assert (implies (= par_or_seq 0) (= marker 1)))  
(assert (implies (= par_or_seq 1) (= marker 0)))  
(assert (implies (= par_or_seq_reprise 0) (= marker 1)))  
(assert (implies (= par_or_seq_reprise 1) (= marker 0)))
```

convert



# Tied choices across sub-schedules

```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop, %marker = transform.tune.alternatives<"par_or_seq"> {  
    %tiled, %forall_loop = transform.tile_using_forall tile_sizes [4,8] %target  
    transform.yield %tiled, %forall_loop, %marked_par  
, {  
    %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [4,8] %target  
    transform.yield %tiled, %for_loops#0, %marked_seq  
}  
... // do something with %tiled_target  
transform.tune.alternatives<"par_or_seq_reprise"> {  
    transform.param.cmpi eq %marker, %marked_par  
    transform.loop.foreach_to_parallel %loop  
    transform.yield  
, {  
    transform.param.cmpi eq %marker, %marked_seq  
    transform.loop.coalesce %loop  
    transform.yield  
}
```

```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop = transform.tile_using_forall tile_sizes [4,8] %target  
... // do something with %tiled_target  
transform.param.cmpi eq %marked_par, %marked_par  
transform.loop.foreach_to_parallel %loop
```

```
%marker = smt.declare_const : !smt.int  
%marker_eq_marked_par = smt.eq %marker, %marked_par  
%par_or_seq_region0_marking = smt.implies %par_or_seq_region0,  
                                %marker_eq_marked_par  
smt.assert %region0_marking  
%marker_eq_marked_seq_par = smt.eq %marker, %marked_seq_par  
%par_or_seq_region1_marking = smt.implies %par_or_seq_region1,  
                                %marker_eq_marked_seq_par  
smt.assert %par_or_seq_region1_marking
```

```
%par_or_seq_reprise = smt.declare_const : !smt.int  
par_or_seq = 0  
marker = 1  
par_or_seq_reprise = 0
```

```
smt.assert par_or_seq_reprise_options  
%par_or_seq_reprise_region0 = smt.eq %par_or_seq_reprise, %c0  
%par_or_seq_reprise_region1 = smt.eq %par_or_seq_reprise, %c1  
%par_or_seq_reprise_region0_marking = smt.implies %par_or_seq_reprise_region0,  
                                %marker_eq_marked_par  
smt.assert %region0_marking  
%par_or_seq_reprise_region1_marking = smt.implies %par_or_seq_reprise_region1,  
                                %marker_eq_marked_seq_par
```

```
(check-sat) .smtlib  
(get-assignment)
```

lower  
/project  
&convert



```
(declare-const par_or_seq Int)  
(declare-const marker Int)  
(declare-const par_or_seq_reprise Int)  
...  
(assert (implies (= par_or_seq 0) (= marker 1)))  
(assert (implies (= par_or_seq 1) (= marker 0)))  
(assert (implies (= par_or_seq_reprise 0) (= marker 1)))  
(assert (implies (= par_or_seq_reprise 1) (= marker 0)))
```

convert

rewrite


SMT solver

# Tied choices across sub-schedules


```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop, %marker = transform.tune.alternatives<"par_or_seq"> {  
    %tiled, %forall_loop = transform.tile_using_forall tile_sizes [4,8] %target  
    transform.yield %tiled, %forall_loop, %marked_par  
, {  
    %tiled, %for_loops:2 = transform.tile_using_for tile_sizes [4,8] %target  
    transform.yield %tiled, %for_loops#0, %marked_seq  
}  
... // do something with %tiled_target  
transform.tune.alternatives<"par_or_seq_reprise"> {  
    transform.param.cmpi eq %marker, %marked_par  
    transform.loop.foreach_to_parallel %loop  
    transform.yield  
, {  
    transform.param.cmpi eq %marker, %marked_seq  
    transform.loop.coalesce %loop  
    transform.yield  
}
```

```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop = transform.tile_using_forall tile_sizes [4,8] %target  
... // do something with %tiled_target  
transform.param.cmpi eq %marked_par, %marked_par  
transform.loop.foreach_to_parallel %loop
```

```
%marked_par = transform.param.constant 1  
%marked_seq = transform.param.constant 0  
%tiled_target, %loop, %_ = transform.tile_using_for tile_sizes [4,8] %target  
... // do something with %tiled_target  
transform.param.cmpi eq %marked_seq, %marked_seq  
transform.loop.coalesce %loop
```




# An average plan for an average day ... in IR




```
...  
%tensorized = transform.apply_registered_pass "lower-to-linalg" to %target  
%bufferized = transform.apply_registered_pass "bufferize" to %tensorized  
%x = transform.tune.knob<"w"> options = [4,8,16]  
%y = transform.tune.knob<"z"> options = ...  
%optimized = transform.tune.alternatives<"F_or_no_F"> {  
    %pre, %G_performed = transform.tune.alternatives<"G_before_F_or_not"> {  
        %Ged = transform.apply_registered_pass "G" with options { "w" = %w, "z" = %z } to %bufferized  
        transform.yield %Ged, %c1  
    }, {  
        transform.yield %bufferized, %c0  
    }  
    %a = transform.tune.knob<"x"> options = ["off","on"]  
    %b = transform.tune.knob<"y"> options = ...  
    %Fed = transform.apply_registered_pass "F" with options { "x" = %x, "y" = %y } to %bufferized  
    %res = transform.tune.alternatives<"G_after_F_or_not"> {  
        transform.param.cmpi eq %G_performed, %c0 // crash in case G got performed before F  
        %Ged = transform.apply_registered_pass "G" with options { "w" = %w, "z" = %z } to %Fed  
        transform.yield %Ged  
    }, {  
        transform.yield %Fed  
    }  
    transform.yield %res  
}, { // Only G, no F  
    %Ged = transform.apply_registered_pass "G" with options { "w" = %w, "z" = %z } to %bufferized  
    transform.yield %Ged  
}  
%vectorized = transform.apply_registered_pass "vectorize" to %optimized  
%llvmfed = transform.apply_registered_pass "convert-to-llvm" to %vectorized  
...
```



# Mechanized optimization of schedules



# Mechanized optimization of schedules



# In summary

Choices for parameters on transforms can be “reified” into IR

- This gives knob ops and non-deterministic reading of schedules

Choice among transform sequences can likewise be reified

- This goes beyond (parameter) auto-tuning, without the intractability of auto-scheduling

Joint constraints allow for precise control over search space

- So that engineers can specify which schedules they are willing to consider

Search for find constraint-satisfying assignments can be handed off, e.g., to a SMT solver

- Whereupon the assignments can be used to rewrite to executable schedules

# Questions?

