

Iago: AI Driven Superoptimization for LLVM

Iago



Manasij Mukherjee
NVIDIA

John Regehr
University of Utah

Talk Organization

Motivation

Why solve this problem?

Design

What worked and what doesn't?

Results

How well does it work?

What is Superoptimization?

Can a piece of code be improved?

Specifically, by **searching** for a refinement.

$$\neg A \wedge \neg B \Rightarrow \begin{matrix} ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \\ ? & ? & ? & ? & ? \end{matrix}$$

```
define i32 @foo(i32 %x) {  
entry:  
    %a = and i32 %x, 1431655765  
    %b = lshr i32 %x, 1  
    %c = and i32 %b, 1431655765  
    %d = add i32 %a, %c  
    %e = and i32 %d, 858993459  
    %f = lshr i32 %d, 2  
    %g = and i32 %f, 858993459  
    %h = add i32 %e, %g  
    %i = and i32 %h, 252645135  
    %j = lshr i32 %h, 4  
    %k = and i32 %j, 252645135  
    %l = add i32 %i, %k  
    %m = and i32 %l, 16711935  
    %n = lshr i32 %l, 8  
    %o = and i32 %n, 16711935  
    %p = add i32 %m, %o  
    %q = and i32 %p, 65535  
    %r = lshr i32 %p, 16  
    %s = and i32 %r, 65535  
    %t = add i32 %q, %s  
    ret i32 %t  
}
```

Searching for a refinement

Enumerate a bunch
of candidates

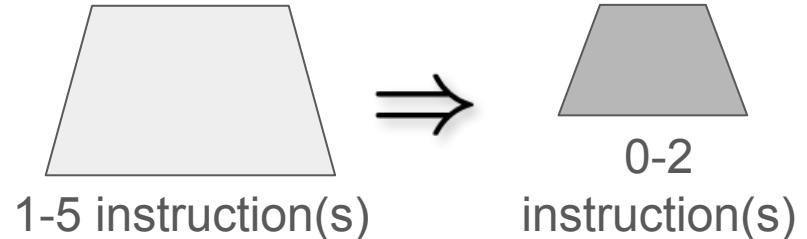


Attempt to verify each

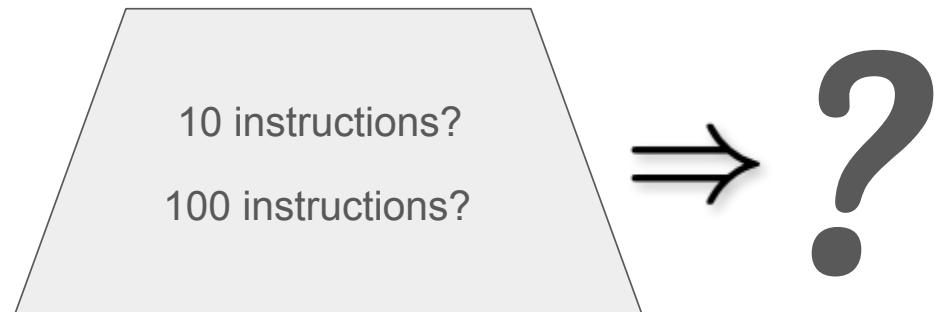
*Superoptimizing sqlite3 (53.94k lines of code)
takes ~24 hours with 128 cores with Souper*

What's on the table?

What we usually put into
InstCombine today:

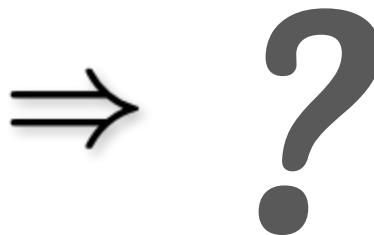


One off pattern in
your important code
that is not generally
useful



```
define i32 @foo(i32 %x) {  
entry:  
    %a = and i32 %x, 1431655765  
    %b = lshr i32 %x, 1  
    %c = and i32 %b, 1431655765  
    %d = add i32 %a, %c  
    %e = and i32 %d, 858993459  
    %f = lshr i32 %d, 2  
    %g = and i32 %f, 858993459  
    %h = add i32 %e, %g  
    %i = and i32 %h, 252645135  
    %j = lshr i32 %h, 4  
    %k = and i32 %j, 252645135  
    %l = add i32 %i, %k  
    %m = and i32 %l, 16711935  
    %n = lshr i32 %l, 8  
    %o = and i32 %n, 16711935  
    %p = add i32 %m, %o  
    %q = and i32 %p, 65535  
    %r = lshr i32 %p, 16  
    %s = and i32 %r, 65535  
    %t = add i32 %q, %s  
    ret i32 %t  
}
```

Do we really need
to optimize long
sequences?



Superoptimization examples

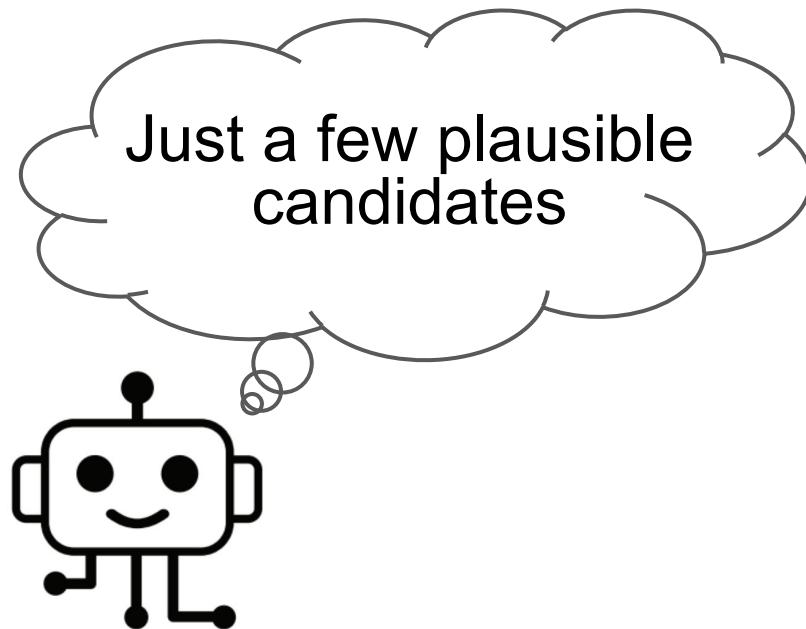
```
define i32 @foo(i32 %x) {  
entry:  
    %a = and i32 %x, 1431655765  
    %b = lshr i32 %x, 1  
    %c = and i32 %b, 1431655765  
    %d = add i32 %a, %c  
    %e = and i32 %d, 858993459  
    %f = lshr i32 %d, 2  
    %g = and i32 %f, 858993459  
    %h = add i32 %e, %g  
    %i = and i32 %h, 252645135  
    %j = lshr i32 %h, 4  
    %k = and i32 %j, 252645135  
    %l = add i32 %i, %k  
    %m = and i32 %l, 16711935  
    %n = lshr i32 %l, 8  
    %o = and i32 %n, 16711935  
    %p = add i32 %m, %o  
    %q = and i32 %p, 65535  
    %r = lshr i32 %p, 16  
    %s = and i32 %r, 65535  
    %t = add i32 %q, %s  
    ret i32 %t  
}
```

ctpop %x

```
define i32 @src(i1 %c) {  
    %8 = zext i1 %c to i32  
    %9 = select i1 %c, i32 2, i32 0  
    %10 = or i32 %8, %9  
    %11 = select i1 %c, i32 4, i32 0  
    %12 = or i32 %10, %11  
    %13 = select i1 %c, i32 8, i32 0  
    %14 = or i32 %12, %13  
    %15 = select i1 %c, i32 256, i32 0  
    %16 = or i32 %14, %15  
    %17 = select i1 %c, i32 512, i32 0  
    %18 = or i32 %16, %17  
    %19 = select i1 %c, i32 1024, i32 0  
    %20 = or i32 %18, %19  
    %21 = select i1 %c, i32 2048, i32 0  
    %22 = or i32 %20, %21  
    ret i32 %22  
}
```

select i1 %c, i32 3855, i32 0

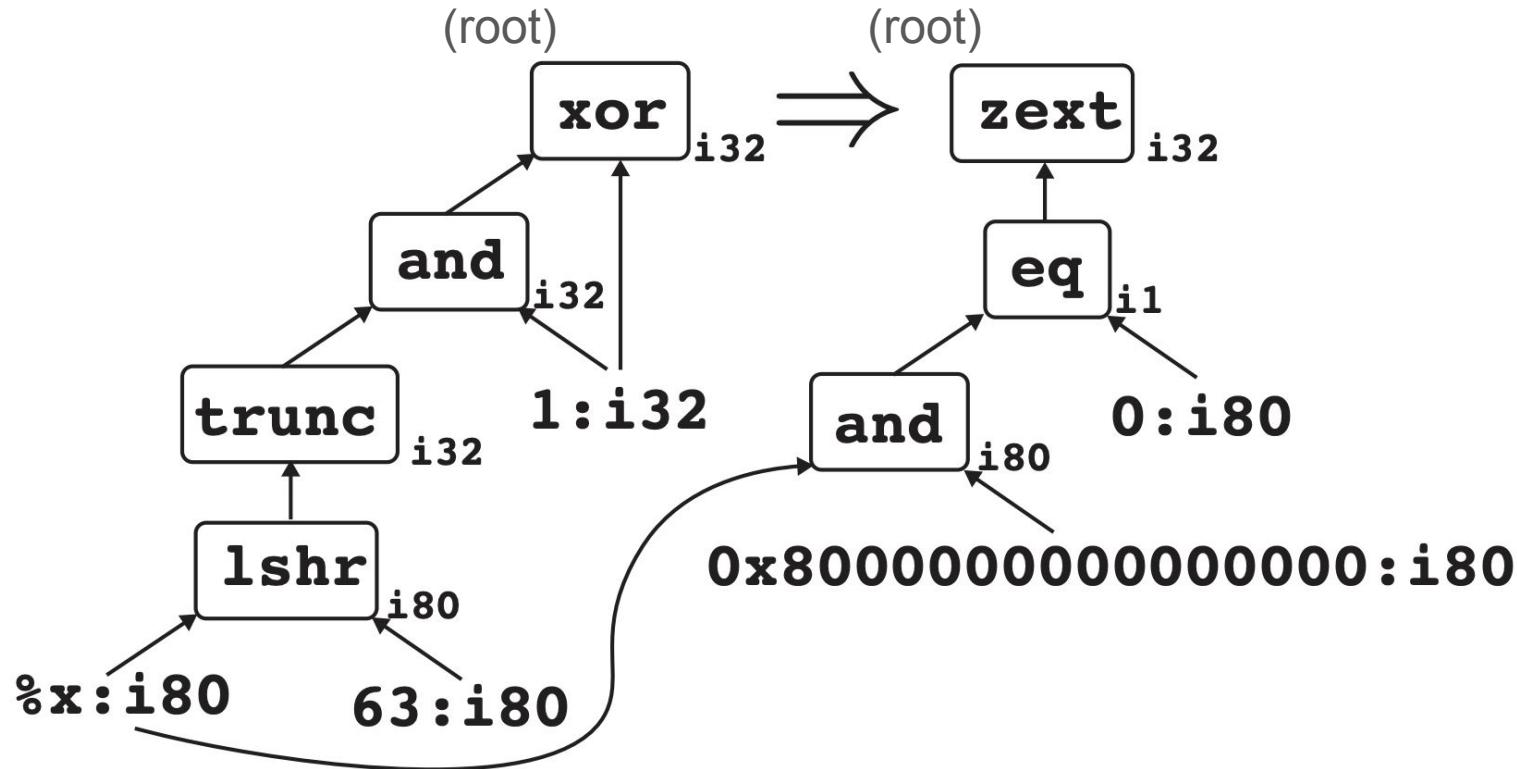
Why replace the enumeration with AI?



Reach far deeper
into the search space

Attempt to verify each

Example - guess what this is!



Superoptimization → Compilers

Can a piece of code be improved?

Can we improve the compiler such that -

Hydra

- it improves this piece of code?
- ... it improves other, similar, pieces of code?
- ... and does not regress anything else?

Talk Organization

Motivation

Why solve this problem?

Design

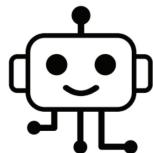
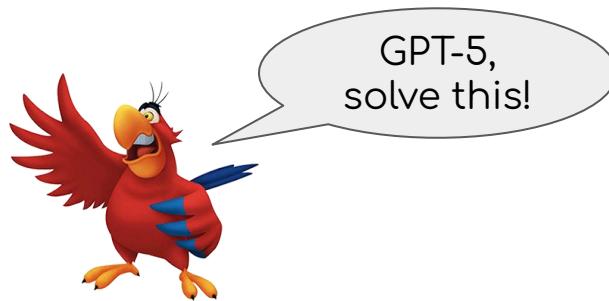
What worked and what doesn't?

Results

How well does it work?

Baseline: just asking the LLM!

$$\neg A \wedge \neg B \Rightarrow ?$$



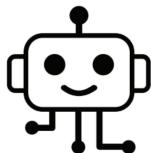
The expression you've given seems to be an incomplete logical expression. The " \Rightarrow " symbol typically indicates implication, suggesting that ...

Two significant issues

$$\neg A \wedge \neg B \Rightarrow ?$$



GPT-5,
solve this!



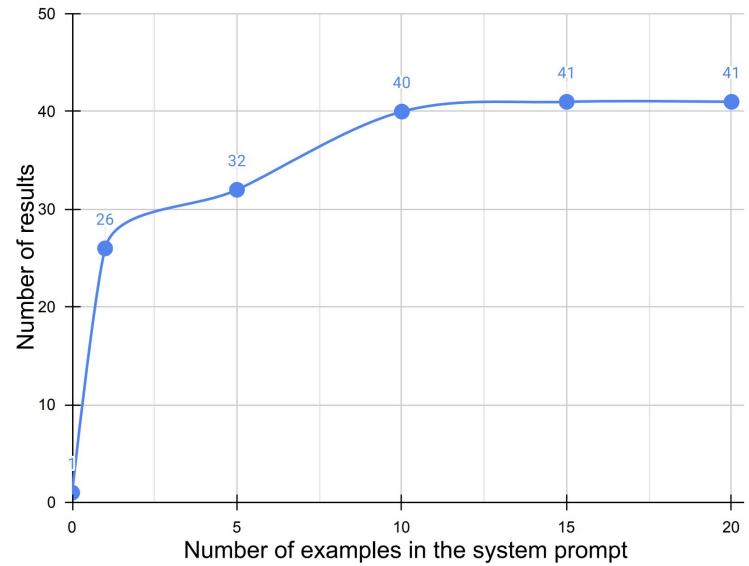
The expression you've given seems to be an incomplete logical expression. The " \Rightarrow " symbol typically indicates implication, suggesting that ...

Wrong syntax!

Wrong semantics!

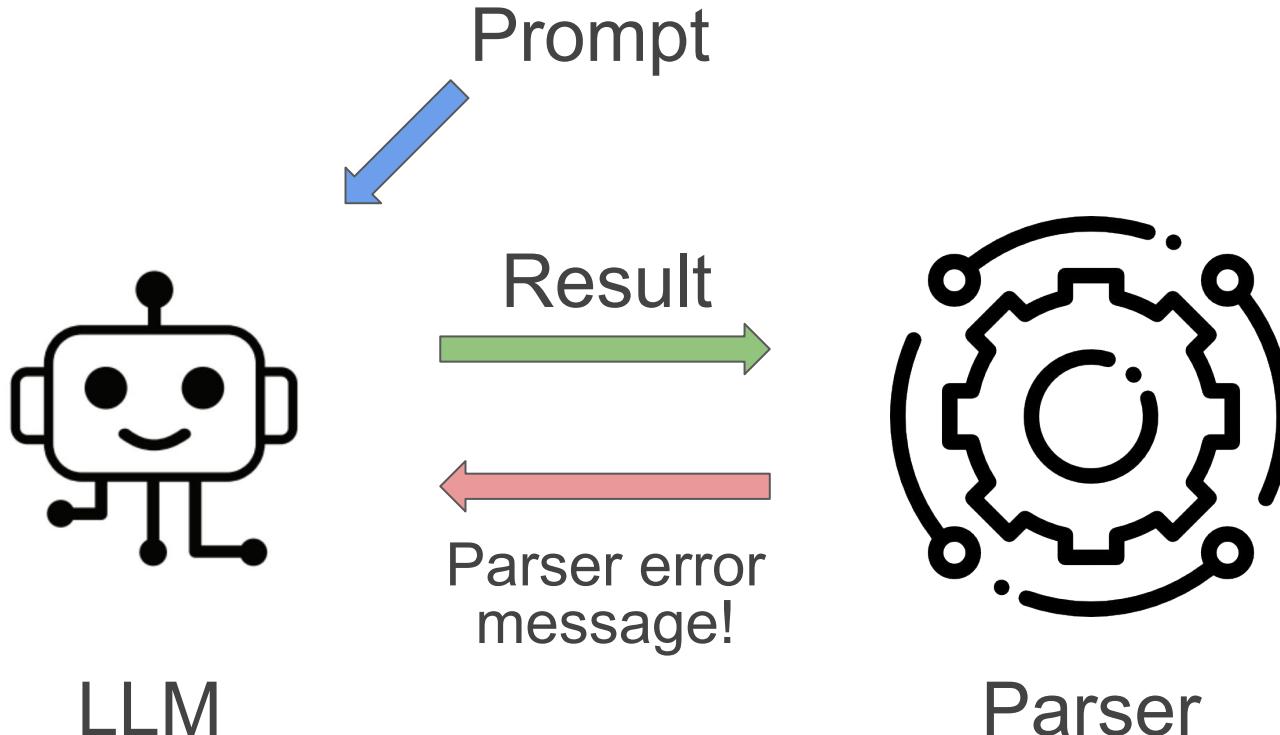
Iago's System Prompt

- Explain the idea of peephole optimizations
- Describe available operations
- Negative statements about what to avoid
- Examples



No correctness guarantee from an LLM

Getting the right syntax



Fixing seemingly consistent issues

LLM results kept redefining values

```
%x = add i32 %a, %b  
%x = add i32 %x, 1
```



```
%x = add i32 %a, %b  
%x' = add i32 %x, 1
```

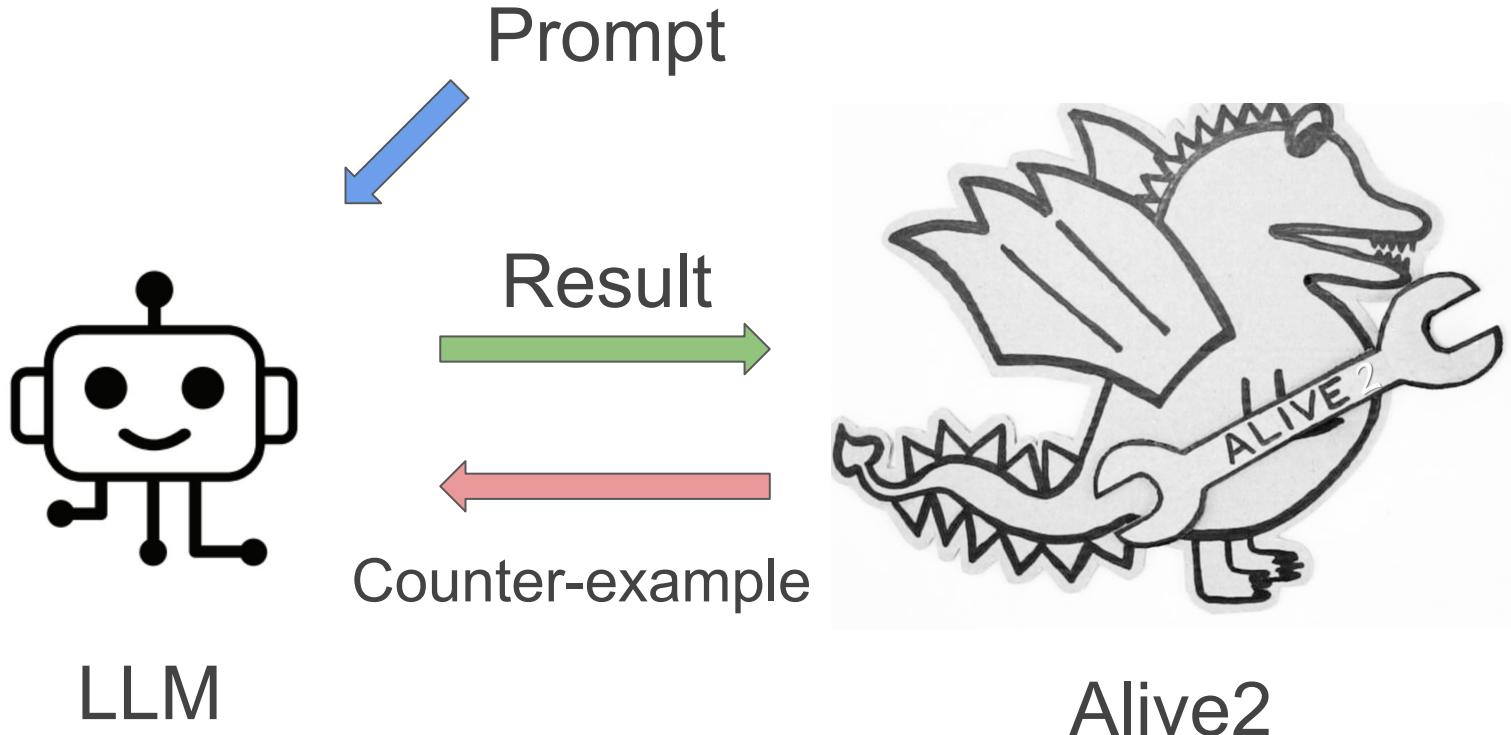
Wrong types

```
%x = or i32 %a, %b  
%c = select i32 %x, ...
```



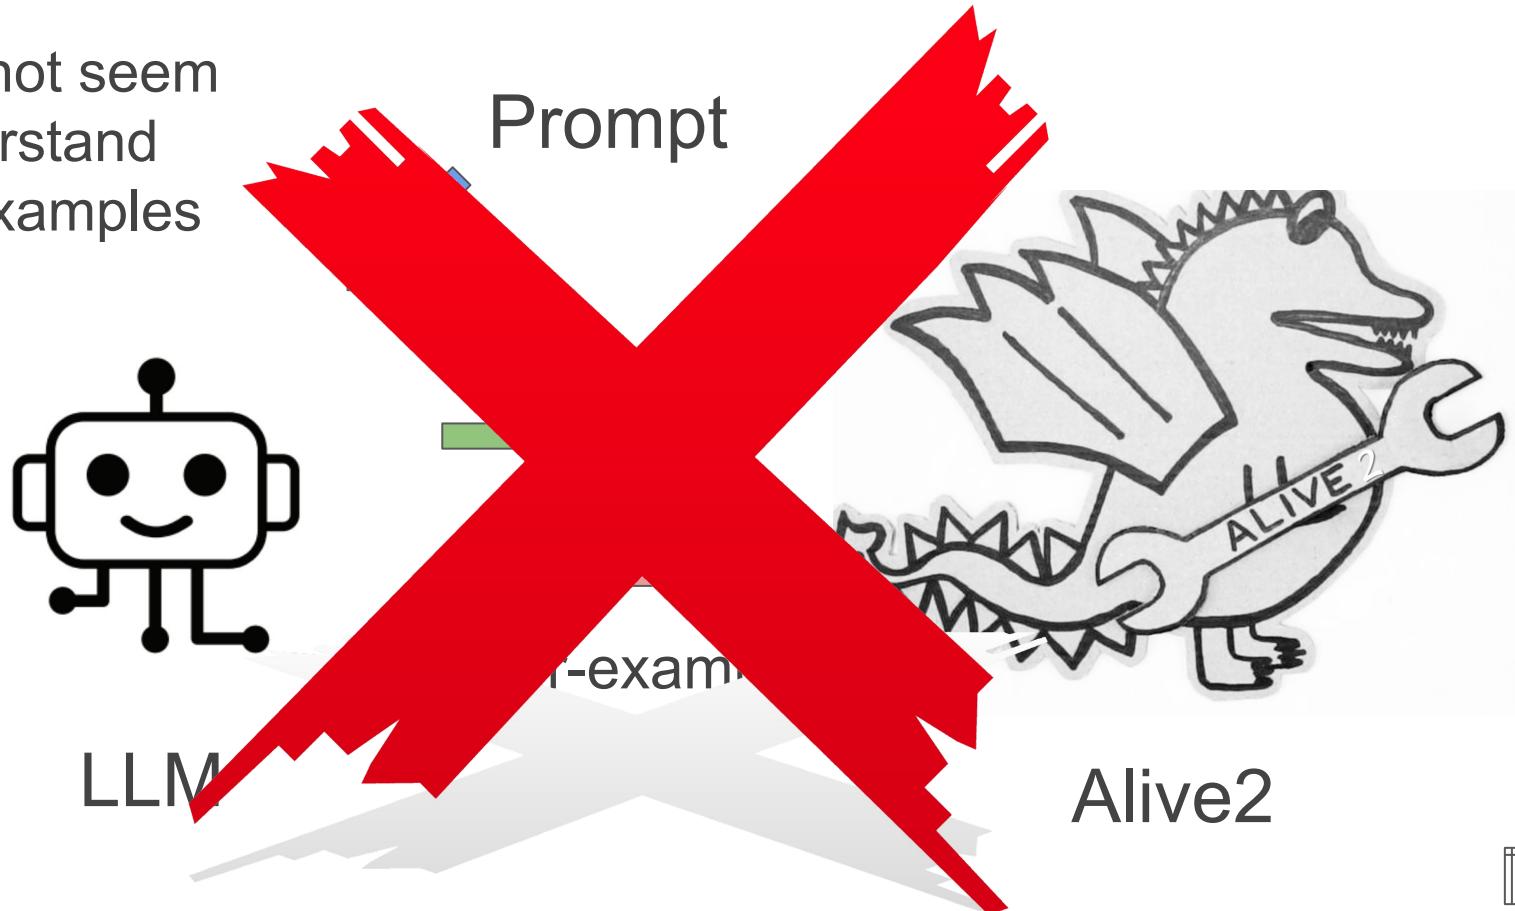
Insert appropriate icmp
or trunc

What about correctness?

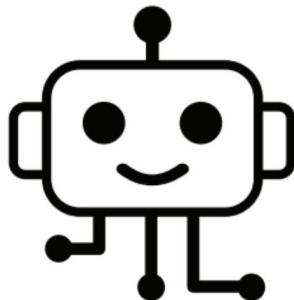


Does not work, at all.

LLMs do not seem
to understand
counterexamples



Works, sort of...



LLM

BLAH

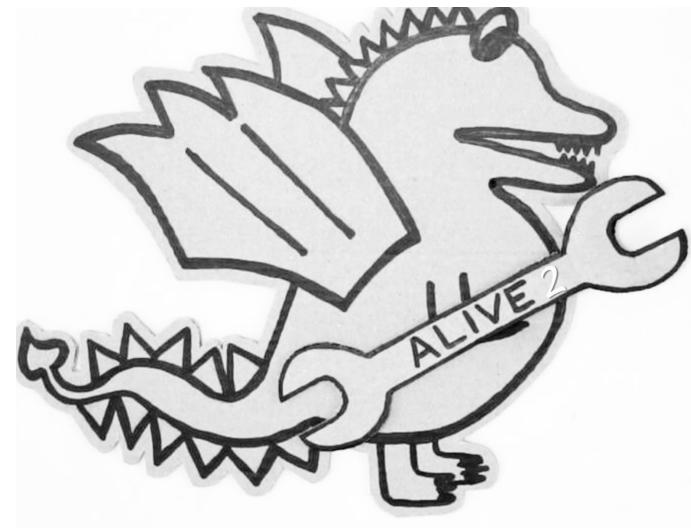


Just try again!

BLAAAAAH



Just try again!



Alive2

Can we do better?

Formal Methods to the rescue again!



$$\neg A \wedge \neg B \Rightarrow ?$$

No correctness guarantee from an LLM

Verify Results

with

Alive2 + Z3

Fix Results

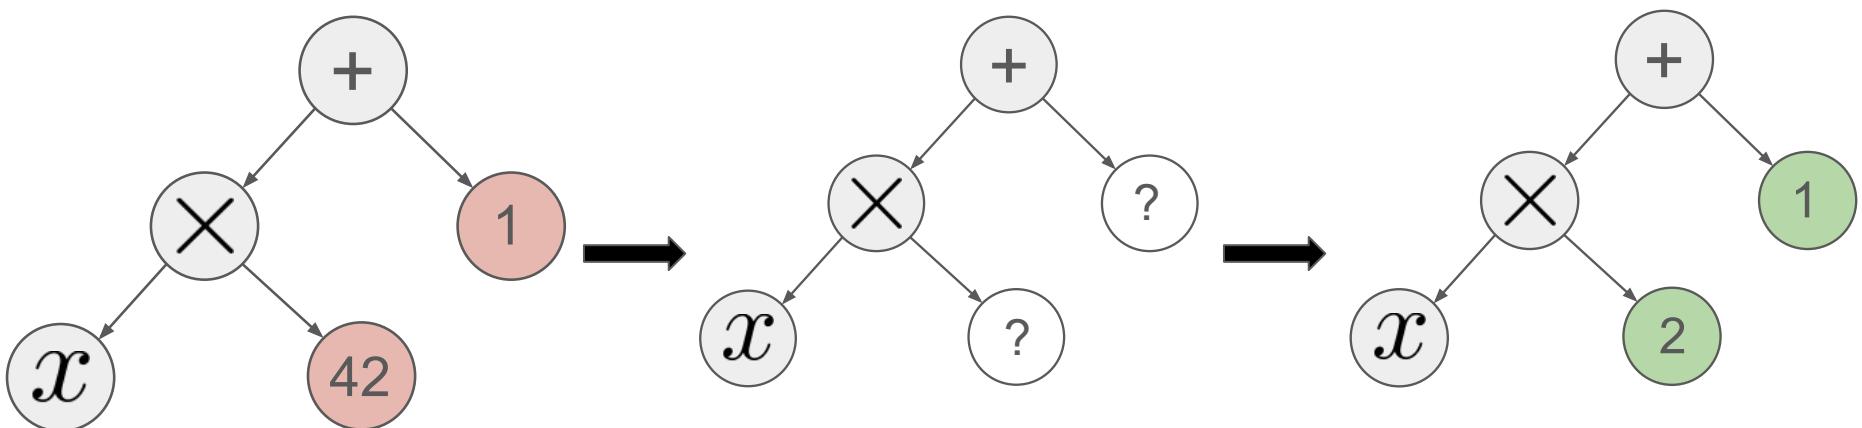
with

CEGIS



Fixing Incorrect Results

Extract a sketch



Fill in the blanks with an
SMT solver

How do we fill in the blanks with an SMT solver?

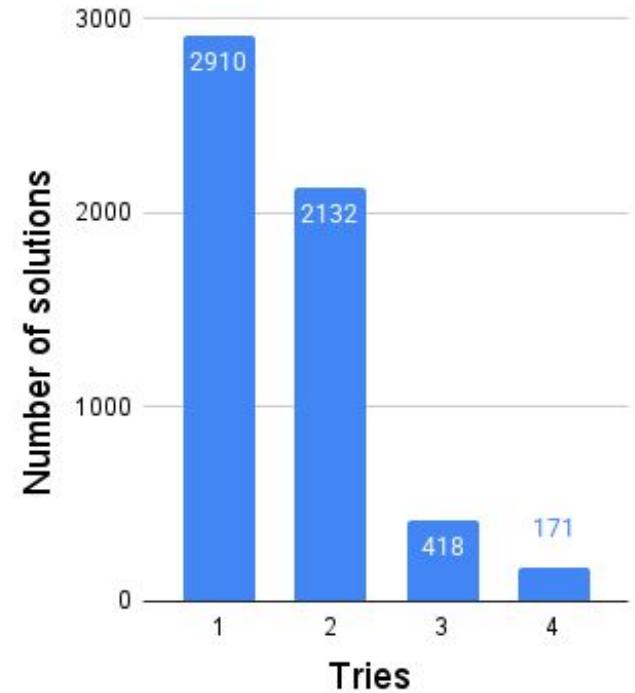
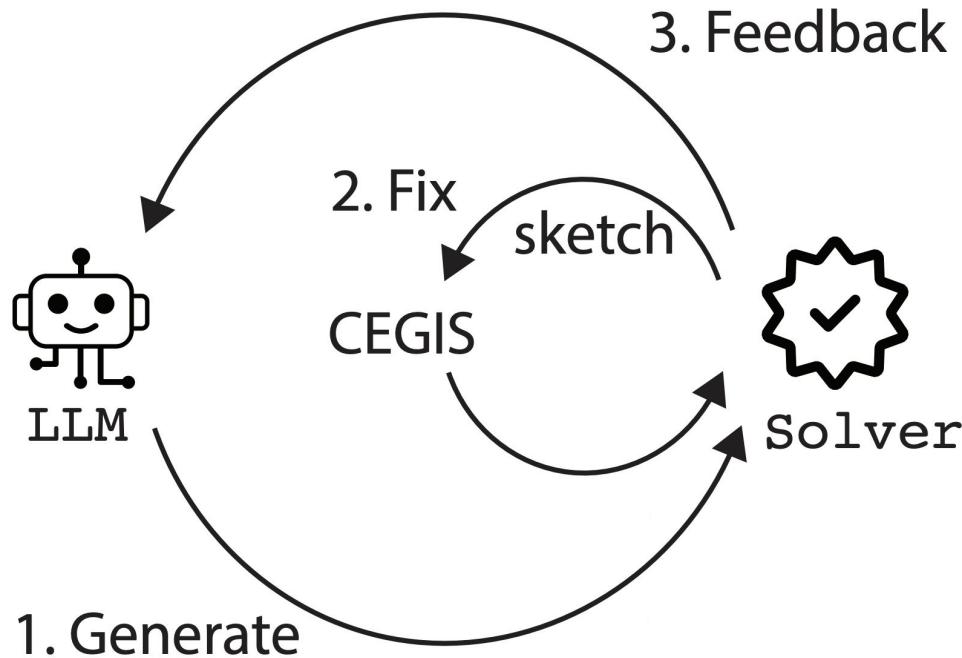
This is the logical formula in question

$$\exists C. \forall x. f(x, C) = g(x, C)$$

Counterexample Guided Inductive Synthesis

[Solving Exists/Forall Problems With Yices](#) [Dutertre]

Recap : lago's Synthesis Loops

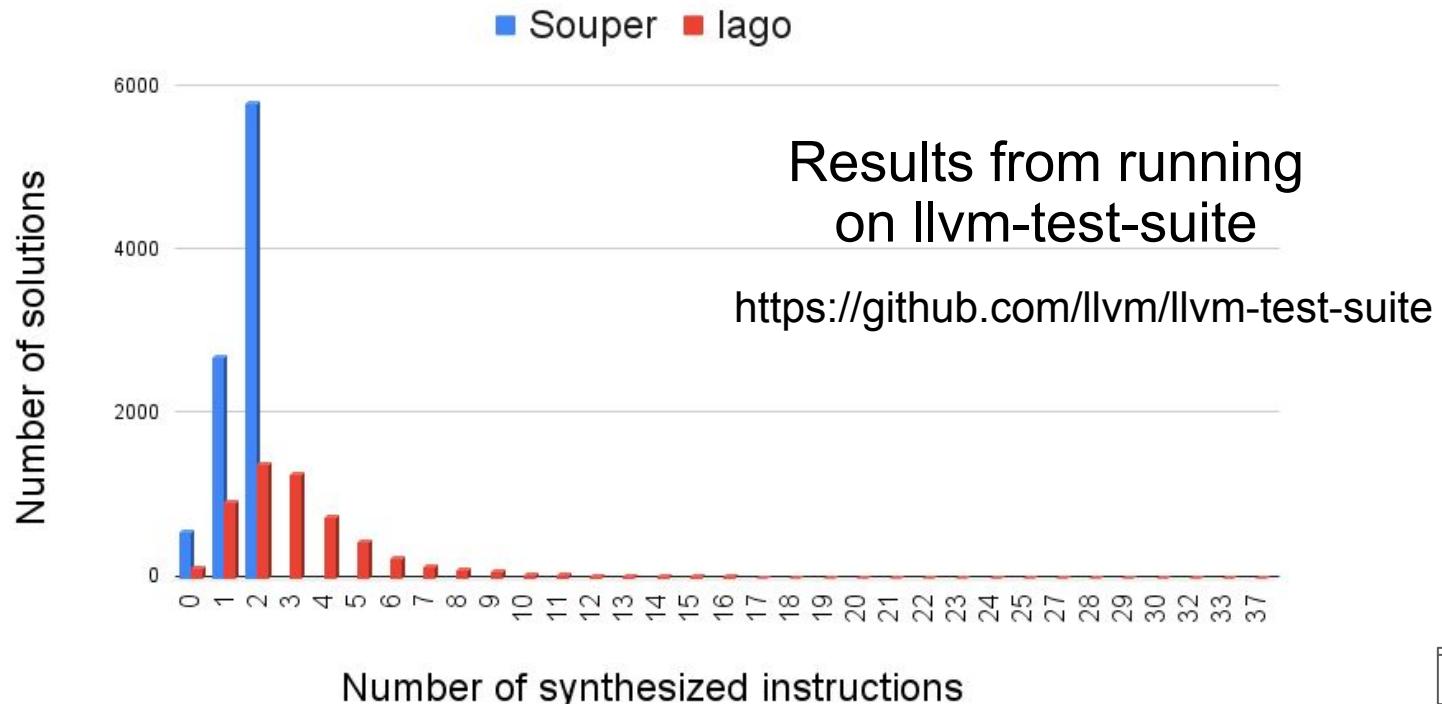


All results formally verified

What about completeness?

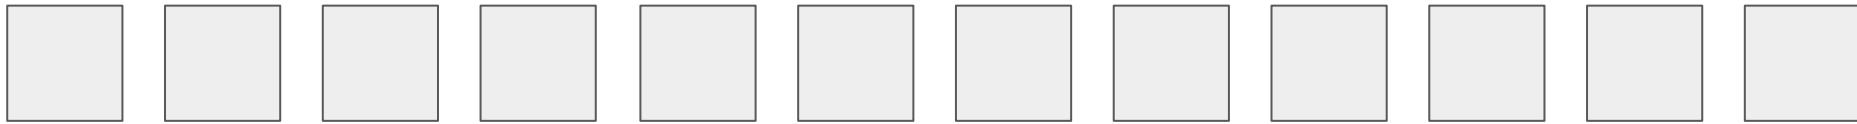
Can we guarantee finding a refinement if one exists ?

No.



Towards Completeness

First level of search tree, incomplete candidates.



Pruning, remove infeasible candidates



Fits in an LLM context window now!

Iago is sound but not complete.

Monkey's Paw

We have wished for :

Correct syntax

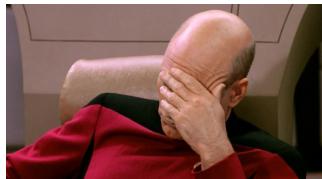
Must be a valid refinement

$x == y \models$

$(x \& \sim y) == 0$



$x == y$



```
%0:i64 = var
%1:i64 = var
%2:i64 = xor -1:i64, %1
%3:i64 = and %0, %2
%4:i1 = eq 0:i64, %3
infer %4
%5:i1 = eq %0, %1
pc %5 1:i1
result %5
```

Talk Organization

Motivation

Why solve this problem?

Design

What worked and what doesn't?

Results

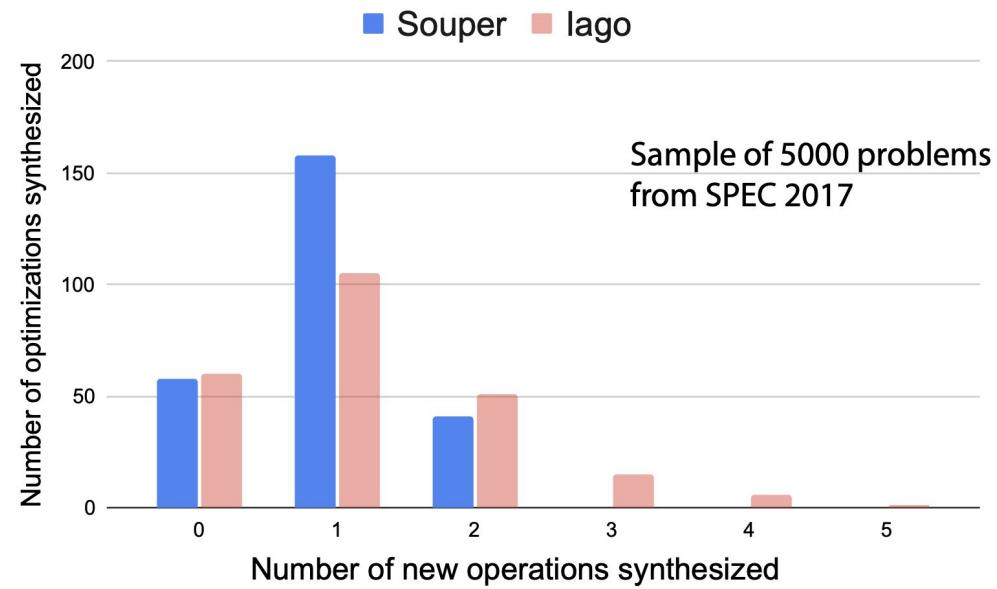
How well does it work?

Iago results summary

More complex optimizations, complements Souper.

35.7% results are new!

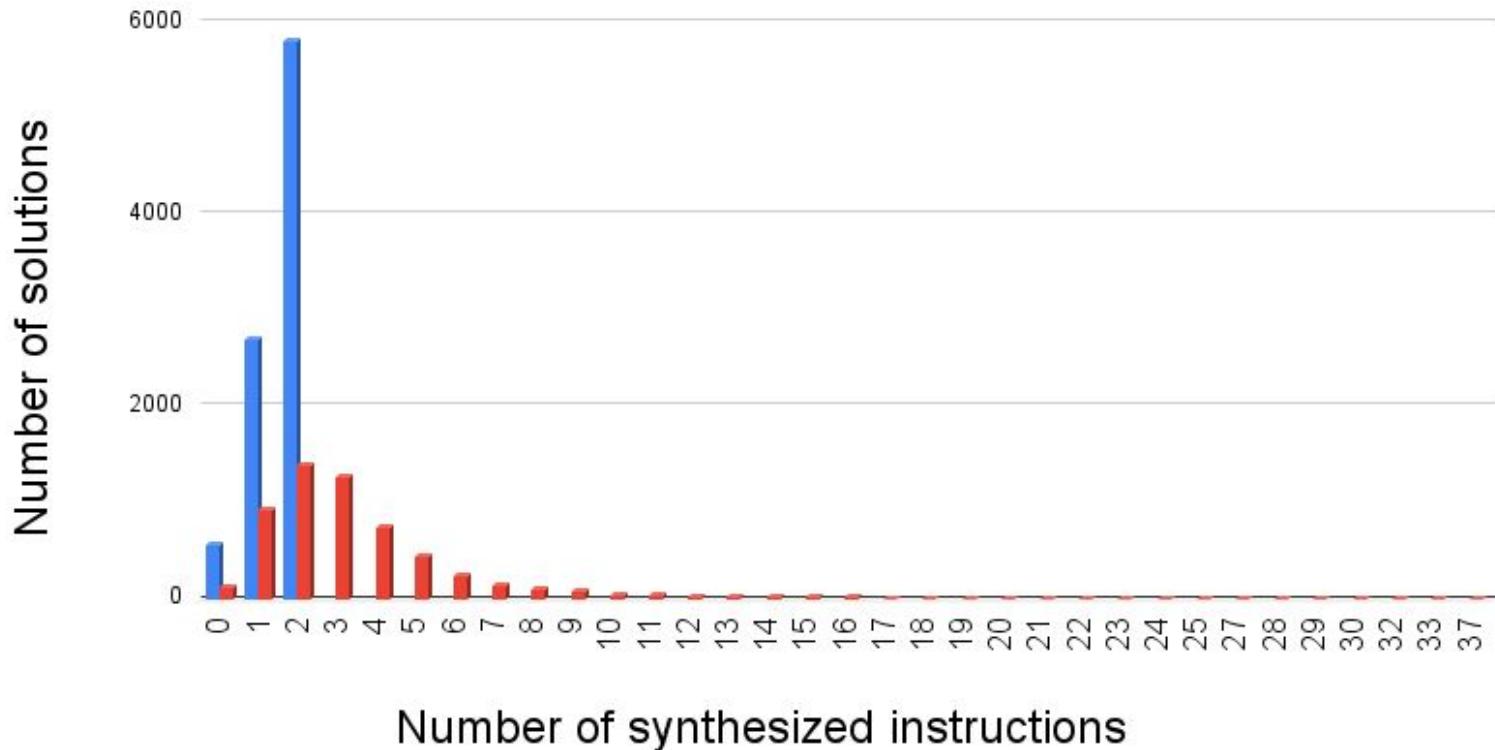
7.4% fewer total results



Results on llvm-test-suite

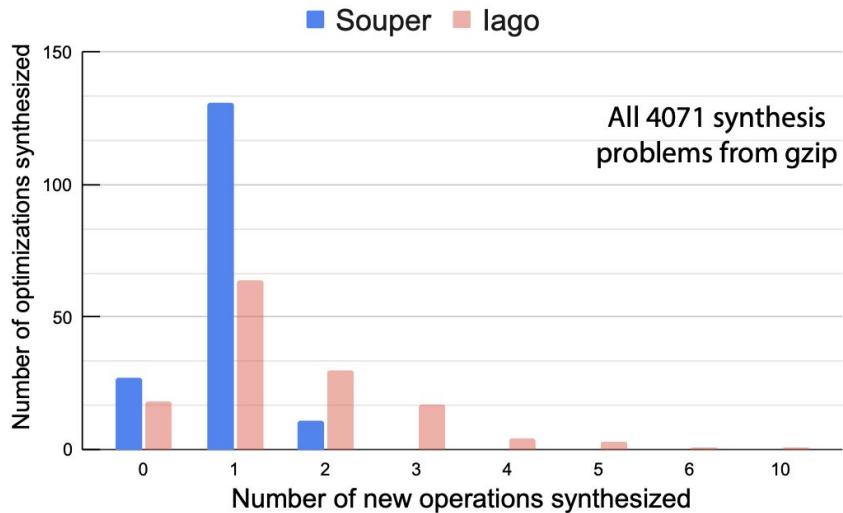
<https://github.com/llvm/llvm-test-suite>

■ Souper ■ lago

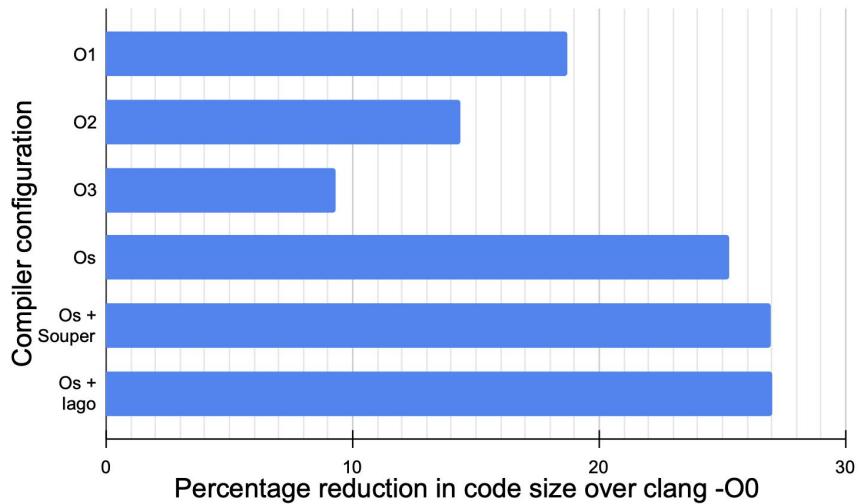


Iago results summary

Finds optimizations in gzip
a 30 year old codebase.



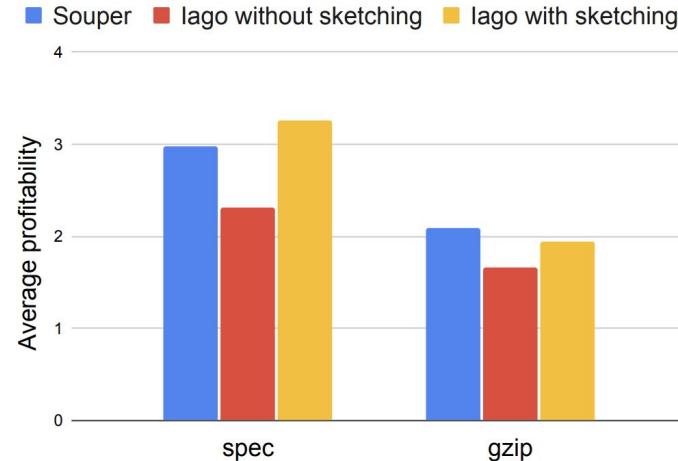
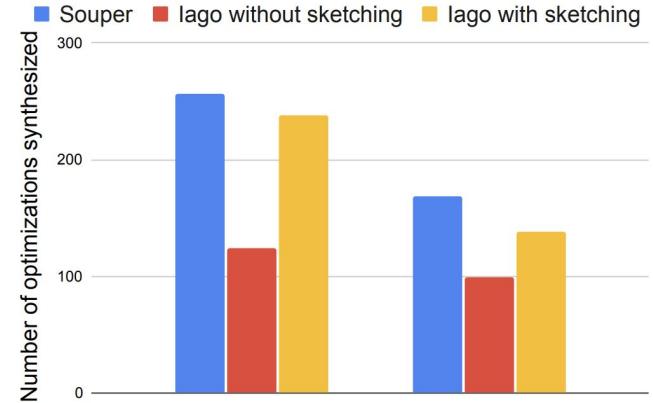
Matches Souper's code size reduction for gzip



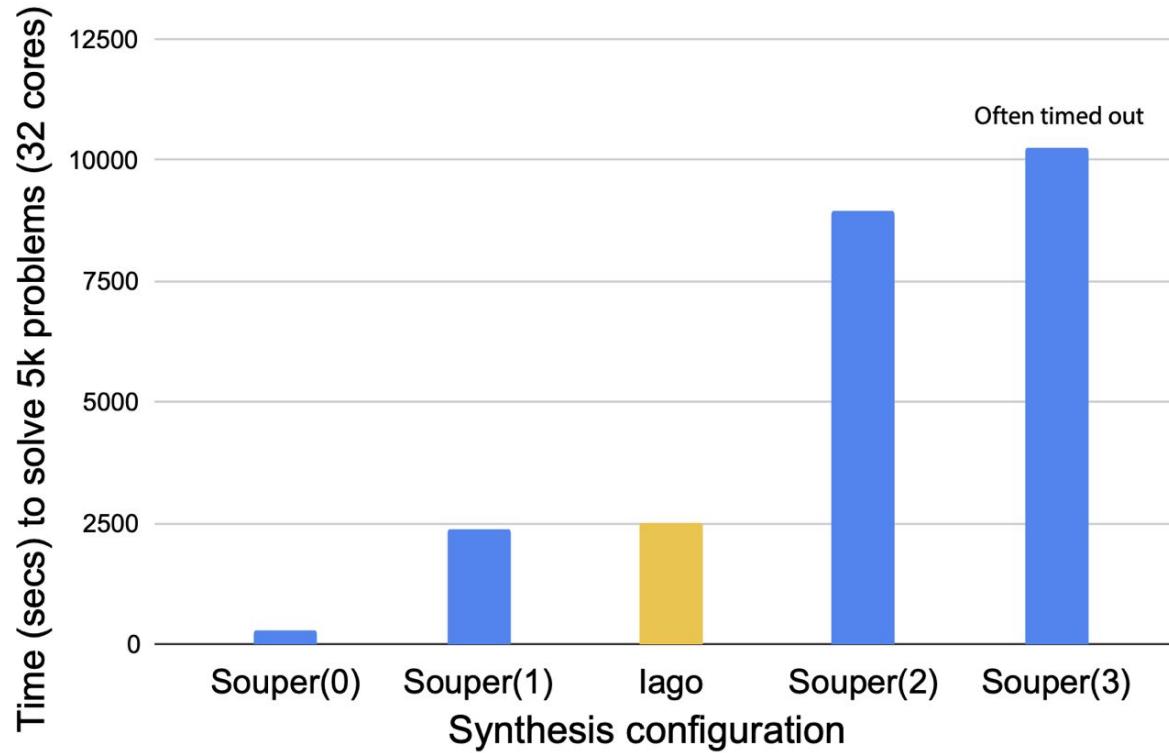
Does sketching help?

More results

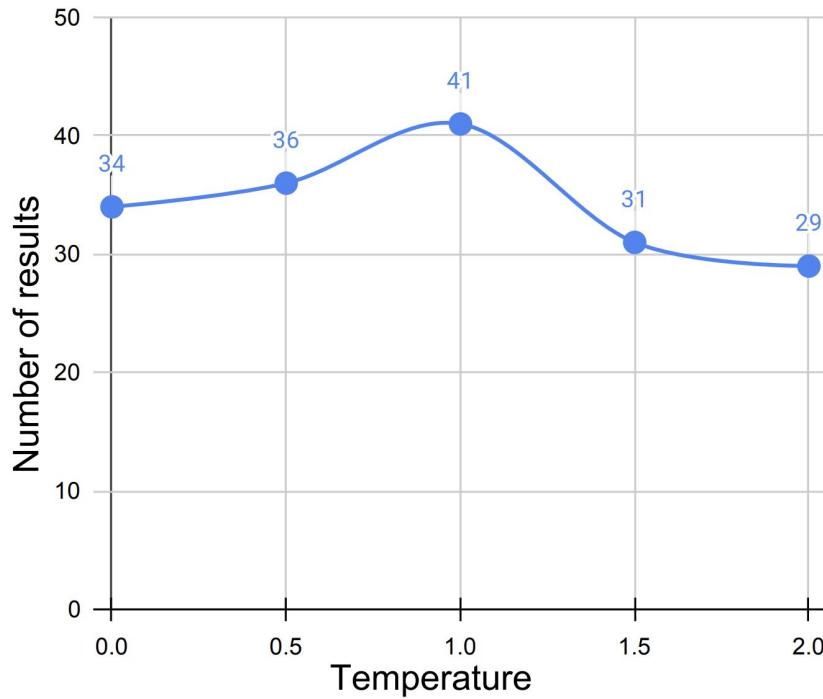
More profitable results



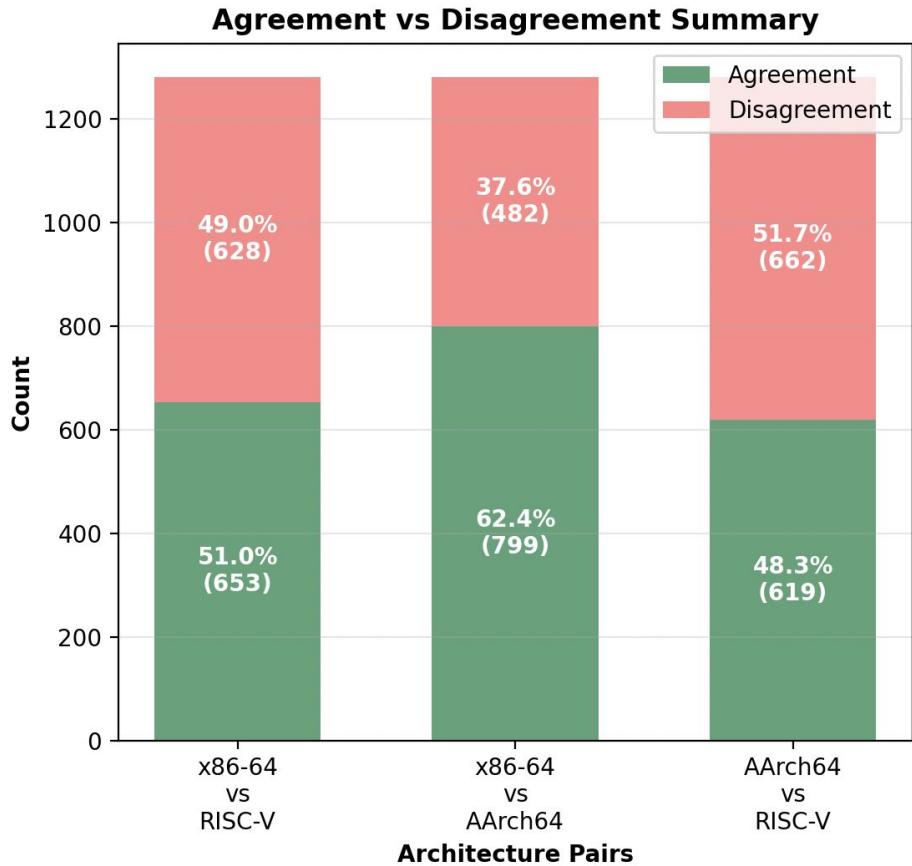
Is it faster than Souper?



GPT-4 Temperature Parameter



Do different backends agree on Iago's optimizations?



sub nuw i64 31, %x

xor i64 %x, 31

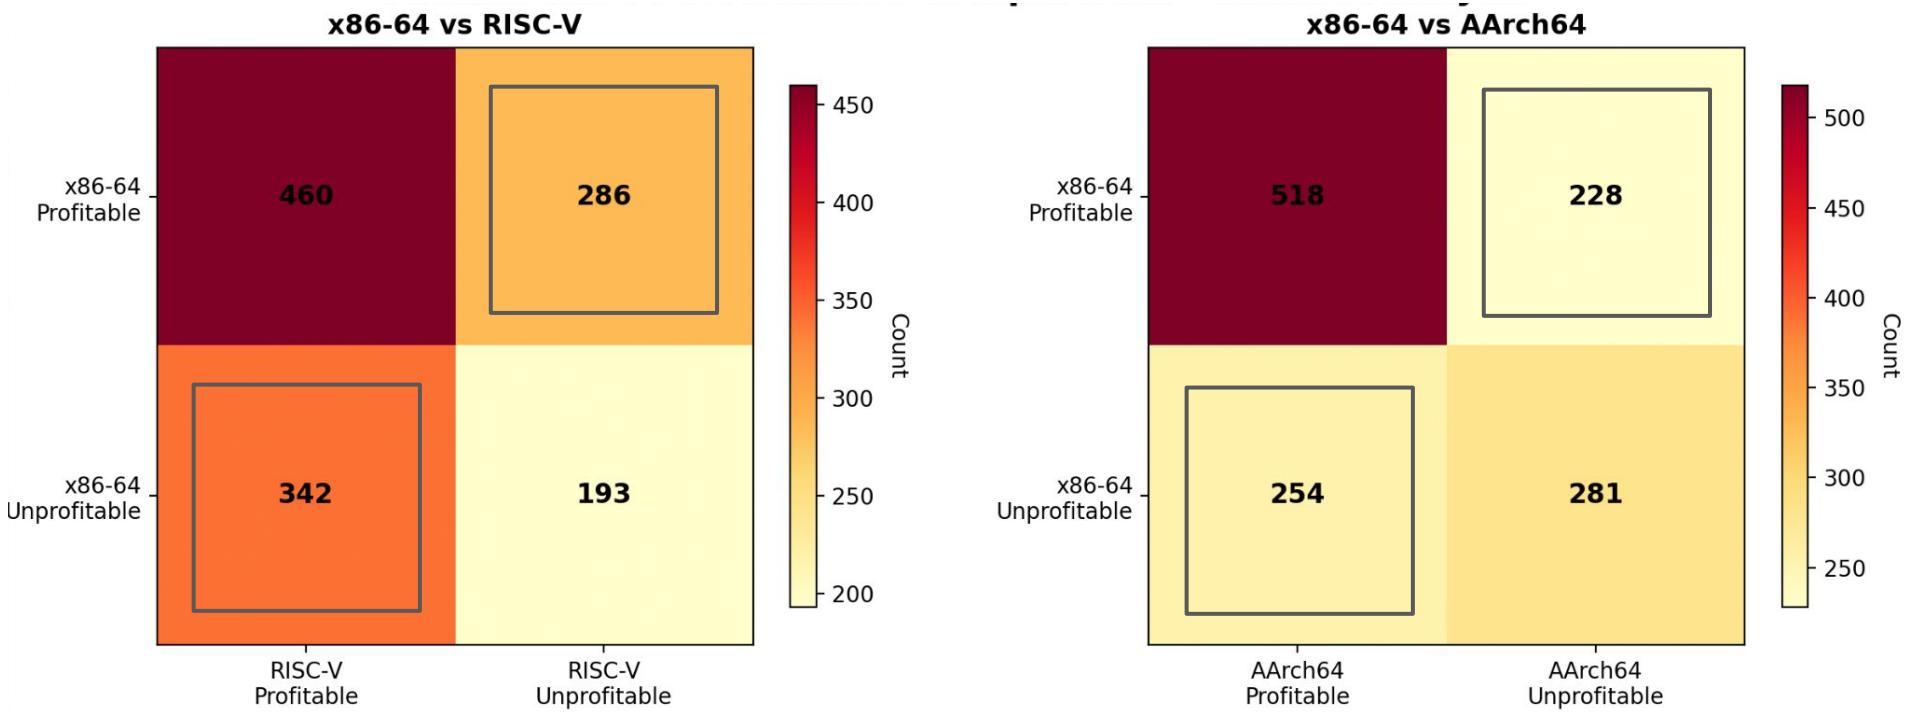
x86-64

riscv

mov rax, rdi
xor rax, 31

xori a0, a0, 31

Breakdown - x86-64 vs RISC-V and x86-64 vs AArch64



Profitable in this context - produces fewer asm instructions

Takeaways implementing these optimizations

```
define i32 @src(i1 %c) {  
    %8 = zext i1 %c to i32  
    %9 = select i1 %c, i32 2, i32 0  
    %10 = or i32 %8, %9  
    %11 = select i1 %c, i32 4, i32 0  
    %12 = or i32 %10, %11  
    %13 = select i1 %c, i32 8, i32 0  
    %14 = or i32 %12, %13  
    %15 = select i1 %c, i32 256, i32 0  
    %16 = or i32 %14, %15  
    %17 = select i1 %c, i32 512, i32 0  
    %18 = or i32 %16, %17  
    %19 = select i1 %c, i32 1024, i32 0  
    %20 = or i32 %18, %19  
    %21 = select i1 %c, i32 2048, i32 0  
    %22 = or i32 %20, %21  
    ret i32 %22  
}
```

Often multiple transforms in disguise

```
define i32 @tgt(i1 %c) #0 {  
    %8 = select i1 %c, i32 3855, i32 0  
    ret i32 %8  
}
```

Takeaways

Takeaways implementing these optimizations

Conflicts with canonicalization

Termination issues!

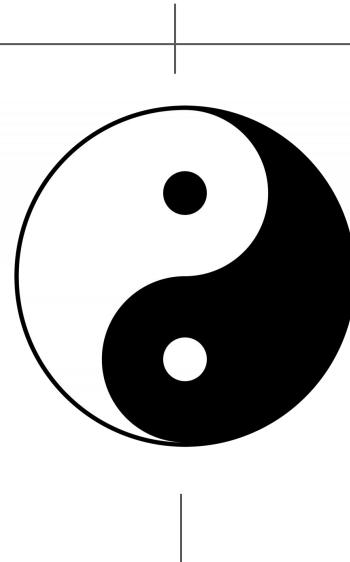
```
zext(xor 1:i1, b:i1)
=> 1 - zext(b)
```

Picking the right subset is tricky

Symbiosis between AI and Formal Methods

Great at generating text

No correctness guarantees



Exhaustive synthesis

Always correct

Some final - REDO

	src:	
!(x0:i8 <u 9)	andi	a0, a0, 255
=	slli	a0, a0, 3
zext(((zext(x0) <<nw 3)	addi	a0, a0, -65
+nsw 0xFFFFFBF))	slli	a0, a0, 32
=>	srlt	a0, a0, 32
(zext(x0) << 3)	andi	a0, a0, 255
+ 0xFFFFFFFFFFFFFFBF	slli	a0, a0, 3
	addi	a0, a0, -65
	tgt:	

	src:	
(trunc((x0:i80 >>l 63))	mov	rax, rdi
& 1) ^ 1	shr	rax, 63
=>	xor	eax, 1
zext((
(x0 & 0x8000000000000000) == 0))	tgt:	
	xor	eax, eax
	bt	rdi, 63
	setae	al

	src:	
	mov	eax, edi
	mov	ecx, edi
	xor	edx, edx
	div	esi
	sub	rcx, rdx
	mov	rax, rcx
zext(x:i32) -		
zext((x %u y:i32))		
=>		
zext((x - (x %u y)))	tgt:	
	mov	eax, edi
	xor	edx, edx
	div	esi
	sub	edi, edx
	mov	rax, rdi



Takeaways

AI for compiler work - Trust, but verify.

Formal methods to verify and fix results.



Scan the QR code to connect with us.
NVIDIA is the engine of the world's AI infrastructure.
We are the world leader in accelerated computing.

Check out our open career opportunities here:
<https://www.nvidia.com/en-us/about-nvidia/careers/>

