

# Using and Improving Optimization Remarks

*Remarkable compiler insights for performance tuning*

Tobias Stadler

`mail@stadler-tobias.de`

in collaboration with Florian Hahn, Jon Roelofs, Arnold Schwaighofer, Gerolf Hoflehner © Apple

LLVM Developers' Meeting '25, Santa Clara, California

2025-10-28

# Motivation

```
float plsVectorizeThis(std::span<float> vals) {
    float res = 0;
    for(auto x : vals) {
        // e.g. reduce: a + b; calc: x * x;
        res = reduce(res, calc(x));
    }
    return res;
}
```

```
foo.cpp:9:12: remark: loop not vectorized: cannot prove it is safe to reorder floating-point
operations; allow reordering by specifying '#pragma clang loop vectorize(enable)' before
the loop or by providing the compiler option '-ffast-math' [-Rpass-analysis=loop-
vectorize]
```

```
9 |     return a + b;
|           ^
```

```
foo.cpp:13:11: remark: loop not vectorized: call instruction cannot be vectorized
[-Rpass-analysis=loop-vectorize]
```

```
13 |     res = reduce(res, calc(x));
```

*Actionability? Discoverability?*

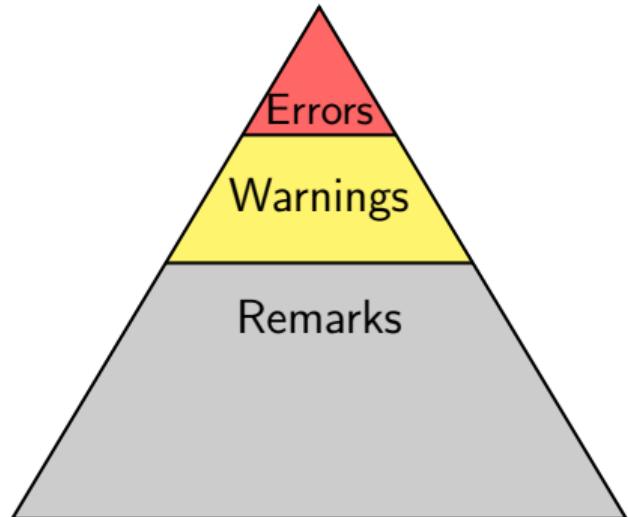
# opt-viewer Example

```
python llvm/tools/opt-viewer/opt-viewer.py <yaml files>
```

```
10         float plsVectorizeThis(std::span<float> vals) {  
11             regalloc 2 virtual registers copies 2.000000e+00 total copies cost generated in function  
12             prologepilog 24 stack bytes in function '_Z16plsVectorizeThisSt4spanIfLm18446744073709551615EE'  
13             asm-printer 22 instructions in function  
14                 float res = 0;  
15                 for (auto x : vals) {  
16                     inline 'std::span<float, 18446744073709551615ul>::begin() const' inlined into 'plsVectori...  
17                     inline 'std::span<float, 18446744073709551615ul>::end() const' inlined into 'plsVectorize...  
18                     inline '_ZN9_gnu_cxxeqlPfSt4spanIfLm18446744073709551615EEEbRKNS_17_normal_i...'...  
19                     inline '_gnu_cxx::normal_iterator<float*, std::span<float, 18446744073709551615ul>...  
20                     inline '_gnu_cxx::normal_iterator<float*, std::span<float, 18446744073709551615ul>...  
21                         TTI advising against unrolling the loop because it contains a call  
22                         gvn load of type float not eliminated because it is clobbered by call  
23                         loop-vectorize loop not vectorized: value that could not be identified as reduction is used outside the loop  
24                         loop-vectorize loop not vectorized: instruction cannot be vectorized  
25                         loop-vectorize loop not vectorized  
26                         slp-vectorizer Cannot SLP vectorize list: vectorization was impossible with available vectorization f...  
27                         asm-printer + BasicBlock: for.body  
28             res = reduce(res, calc(x));  
29             inline reduce(float, float) will not be inlined into plsVectorizeThis(std::span<float, 1844674407...  
30             inline 'calc(float)' inlined into 'plsVectorizeThis(std::span<float, 184467440737095...  
31             loop-vectorize loop not vectorized: call instruction cannot be vectorized  
32         }  
33         return res;  
34     }
```

*Digestability? Scalability?*

# Remarks in the Diagnostic Hierarchy



- ▶ Intention: 3rd diagnostic class
- ▶ Optimizer, Backend
  
- ▶ Passed: Applied optimizations
- ▶ Missed: Untaken optimization opportunities (e.g. unsafe)
- ▶ Analysis: Heuristics, statistics, ...

# User Groups

## Compiler Users

- ▶ Understand missed optimizations
- ▶ Guide the compiler
- ▶ Communicate program invariants
- ▶ Override heuristics

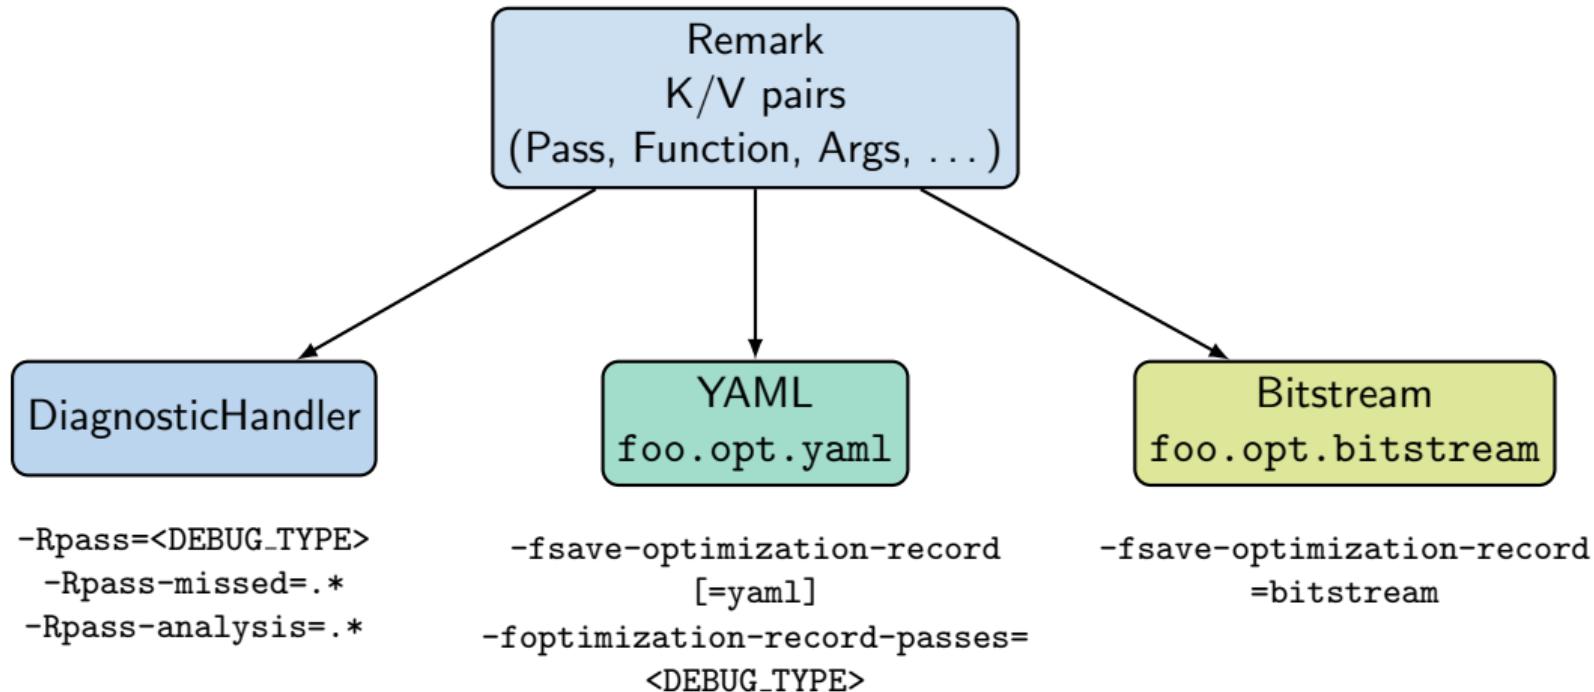
*Remarks as teaching tool:  
user intent ↔ compiler capabilities*

## Compiler Developers

- ▶ Debug optimization passes
- ▶ Tune heuristics
- ▶ Understand regressions
- ▶ Estimate blast radius

*Remarks as telemetry:  
local development & CI tracking*

# Remark Formats



# Remark Landscape

## Generation → Processing → Displaying

- |   |  |   |
|---|--|---|
| <ul style="list-style-type: none"><li>▶ Clang</li><li>▶ MLIR</li><li>▶ Swift</li><li>▶ Inlining</li><li>▶ Vectorization</li><li>▶ Unrolling</li><li>▶ ...</li></ul> | <ul style="list-style-type: none"><li>▶ libRemarks</li><li>▶ llvm-remarkutil</li><li>▶ dsymutil (MachO)</li><li>▶ opt-stats.py</li><li>▶ opt-diff.py</li></ul> | <ul style="list-style-type: none"><li>▶ Diagnostics</li><li>▶ opt-viewer.py</li><li>▶ llvm-opt-report</li><li>▶ Compiler Explorer</li></ul> |
|---|--|---|

*Fragmentation?*

# Bitstream Remarks

## Motivation: Scalability

- ▶ YAML does not scale (entire code bases, collection in CI, ...)

## Issues

- ▶ Supported on MachO only
- ▶ String table in object file
- ▶ Incompatible with opt

## Overhauling the Bitstream Format

- ▶ Bitstream files now fully standalone (Drop-in YAML replacement!)
- ▶ Format auto-detection
- ▶ More efficient encoding: 40% file size reduction
- ▶ Embed blobs: e.g. LLVM BitCode

Regarding removal of SerializerMode, see docs

```
Expected<LLVMRemarkFileHandle> setupLLVMOptimizationRemarks(...);
```

## llvm-remarkutil filter: The multi-tool for bulk processing

```
llvm-remarkutil filter --remark-type=missed --pass-name=inline -o missed.yaml  
foo.opt.bitstream bar.opt.bitstream
```

- ▶ Filter by function, pass, remark name, type, arguments
- ▶ Regex support (`--rfunction`, `--rpass-name`, ...)
- ▶ Automatic format conversion (YAML ↔ Bitstream)
- ▶ Merge multiple remark files
- ▶ `--exclude`
- ▶ `--dedupe`, `--sort`

## llvm-remarkutil summary: Summarizing Remarks

```
llvm-remarkutil summary -o summary.yaml main.opt.bitstream
```

- ▶ Summarize multiple remarks into new remarks using different strategies
- ▶ Current strategies:
  - ▶ --inline-callees: Per-callee inlining statistics
  - ▶ Contributions welcome!
- ▶ Input remark retention: --keep=none|used|all

```
template <class, class _Cp, bool _IsConst, class _Tp, class _Proj, __enable_if_t<__is_identity<_Proj>::value, int> = 0>
_LIBCPP_HIDE_FROM_ABI _LIBCPP_CONSTEXPR_SINCE_CXX20 __iter_diff_t<__bit_iterator<_Cp, _IsConst> >
__count(__bit_iterator<_Cp, _IsConst> __first, __bit_iterator<_Cp, _IsConst> __last, const _Tp& __value, _Proj&) {
    - Incoming Calls (Inlined: 1, TooCostly: 2)
        Least profitable (cost=255, threshold=250)
        Most profitable (cost=30, threshold=250)
```

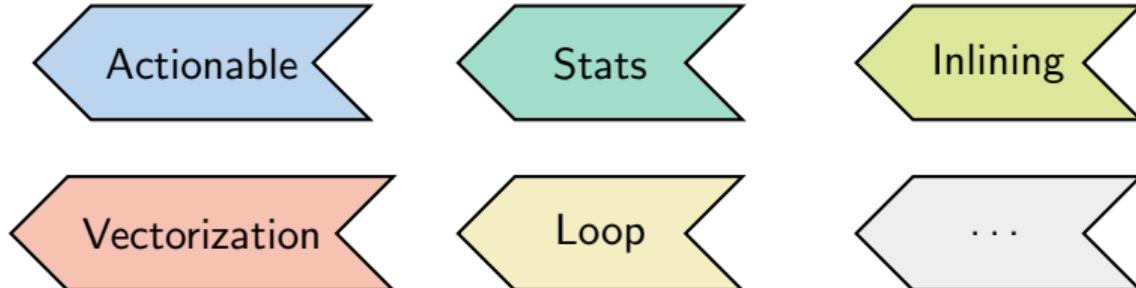
## {Function,Pass-}Local ADT Statistics

- ▶ Fine-grained STATISTIC implementation (using `thread_local`, pass instrumentation)

```
clang++ -fsave-optimization-record -mllvm --stats -O3 main.cpp
```

```
--- !Analysis
Pass: sroa
Name: PassStatistics
Function: _ZNSt6bitsetILm64EEC2Ey
Args:
- MaxPartitionsPerAlloca: '1'
- MaxUsesPerAllocaPartition: '2'
- NumAllocaPartitionUses: '4'
- NumAllocaPartitions: '2'
- NumAllocasAnalyzed: '2'
- NumDeleted: '4'
- NumPromoted: '2'
...
```

## Remark Tags: Different Needs, Different Remarks



- ▶ Multiple tags per Remark
- ▶ Extensible
- ▶ Orthogonal to remark type
- ▶ Obviate the need for sub-types: `AnalysisFPCommute`, `AnalysisAliasing`

## What's next?

More upstreaming tbd:

[github.com/tobias-stadler/llvm-project/tree/remarks-future](https://github.com/tobias-stadler/llvm-project/tree/remarks-future)

- ▶ Expose tags to OptimizationRemarkEmitter, start tagging remarks
- ▶ Warning-like diagnostic categories based on tags (-Rvectorize, ...)
- ▶ **Tag-aware diff tool** (e.g. subtract Stats)
- ▶ Python bindings for libRemarks

## Summary

- ▶ Remarks close the feedback loop between compiler and developer
- ▶ Challenge: How to extract the specific information we need?
- ▶ In need of better infrastructure, tooling, and more actionable remarks
- ▶ Bitstream format now as drop-in replacement for YAML
- ▶ New tools: `llvm-remarkutil` filter, summary
- ▶ Per-function/pass STATISTICS
- ▶ Future: Remark tags for improved organization, filtering, diffing,...
- ▶ Contributions welcome!

Thanks!