



# Scheduling Model in LLVM

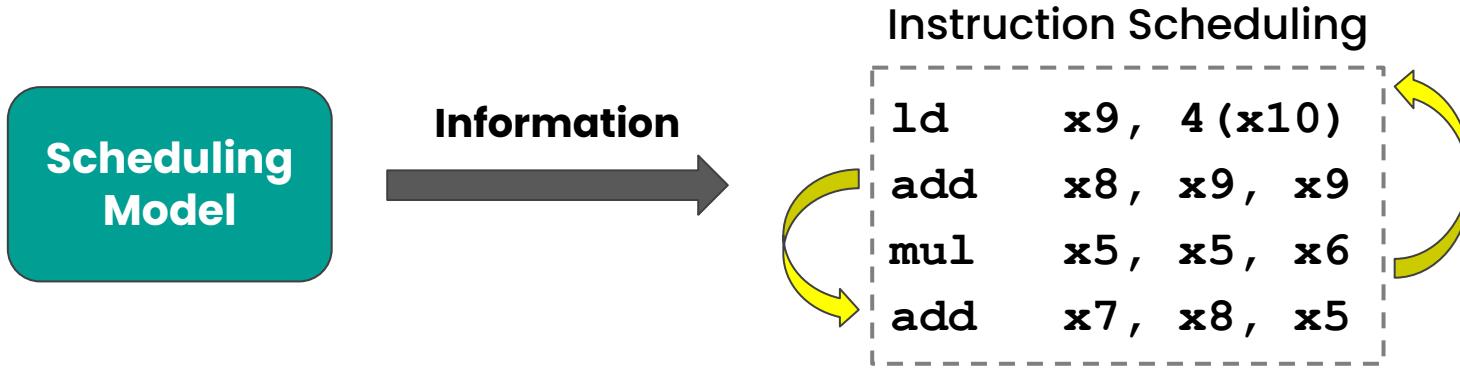
Past, Present, and Future

By Min Hsu @ LLVM Dev Meeting 2025

# By the end of this tutorial...

- What is LLVM's Scheduling Model
- How do *other* parts of LLVM use Scheduling Model
- What are Scheduling Model's connections with hardware
- How can we improve this framework

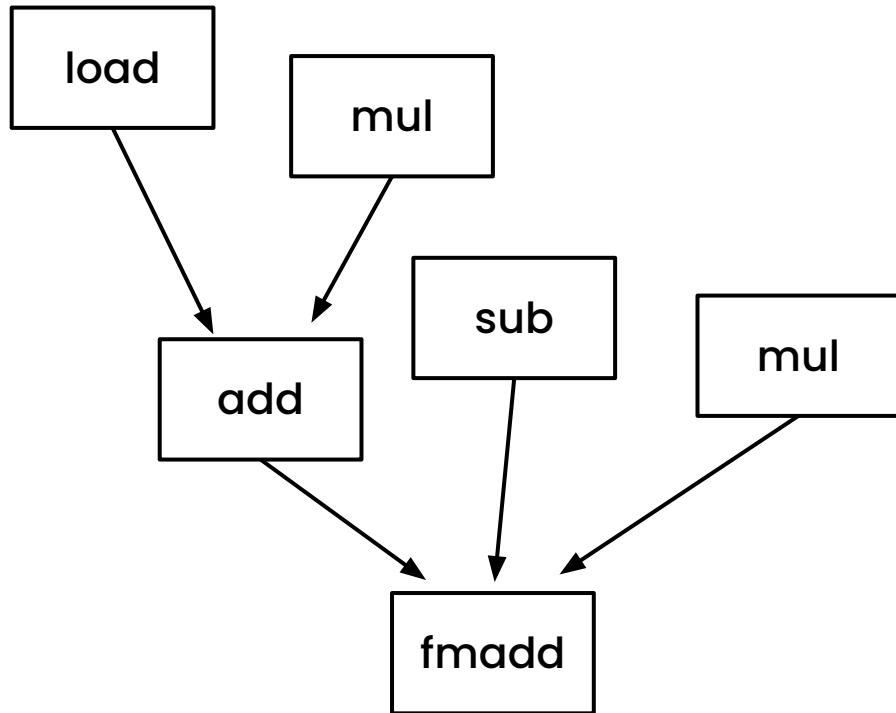
# Scheduling Model: The Inception



- Increase instruction level parallelism
- Reduce the number of register spillings

# Instruction Scheduling

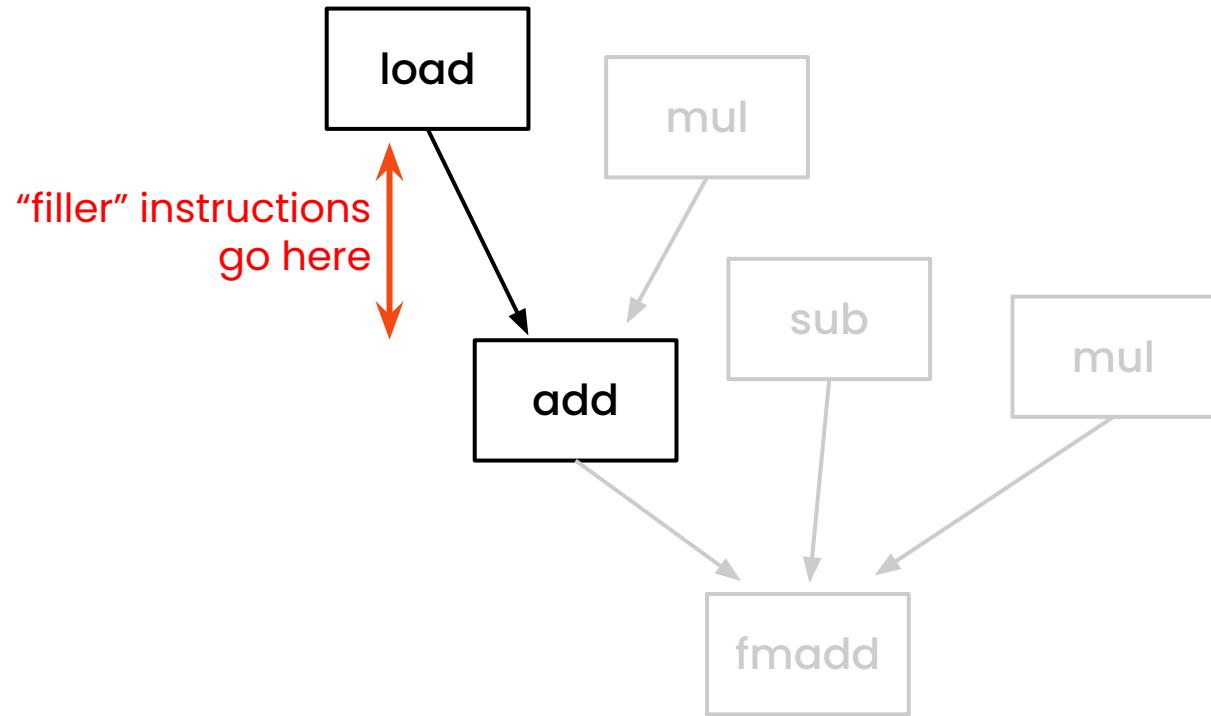
Increase Instruction Level Parallelism



Avoid *stalling* the processor pipeline

# Instruction Scheduling

Increase Instruction Level Parallelism

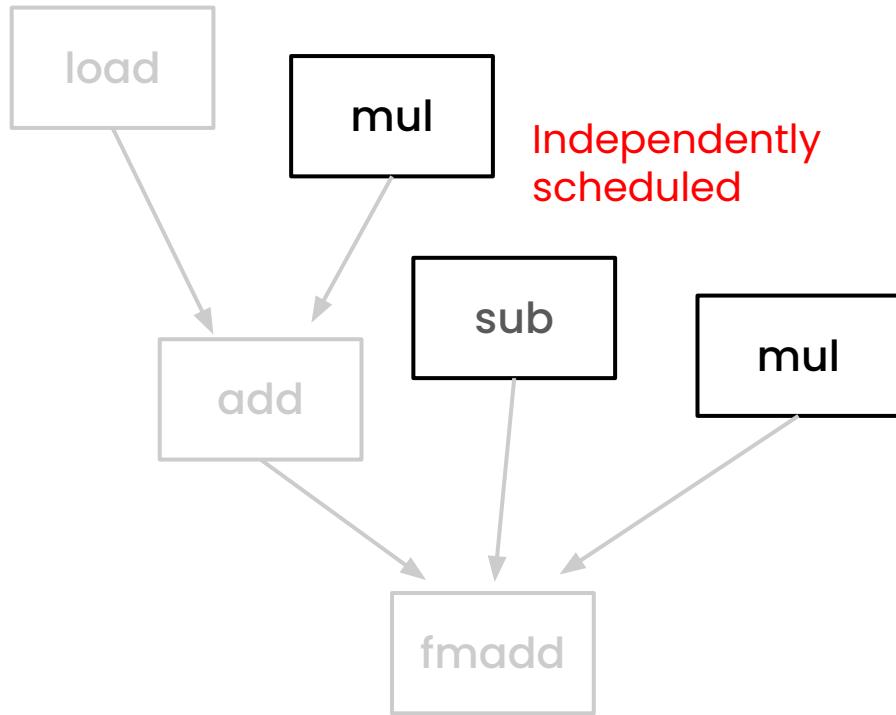


Avoid *stalling* the processor pipeline

- Make the pipeline **busy**

# Instruction Scheduling

Increase Instruction Level Parallelism

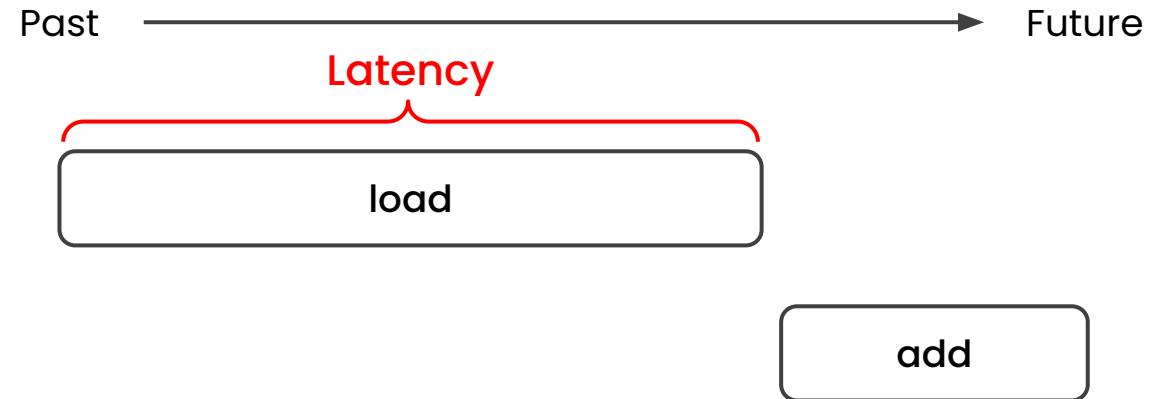
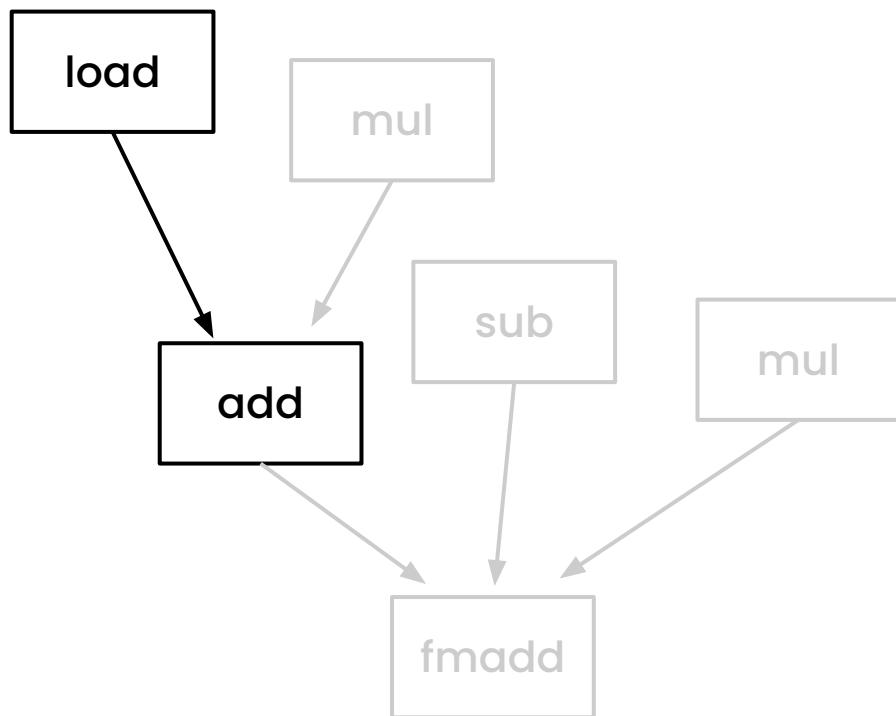


Avoid *stalling* the processor pipeline

- Make the pipeline **busy**
- Avoid multiple instructions competing for the same processor **resource**

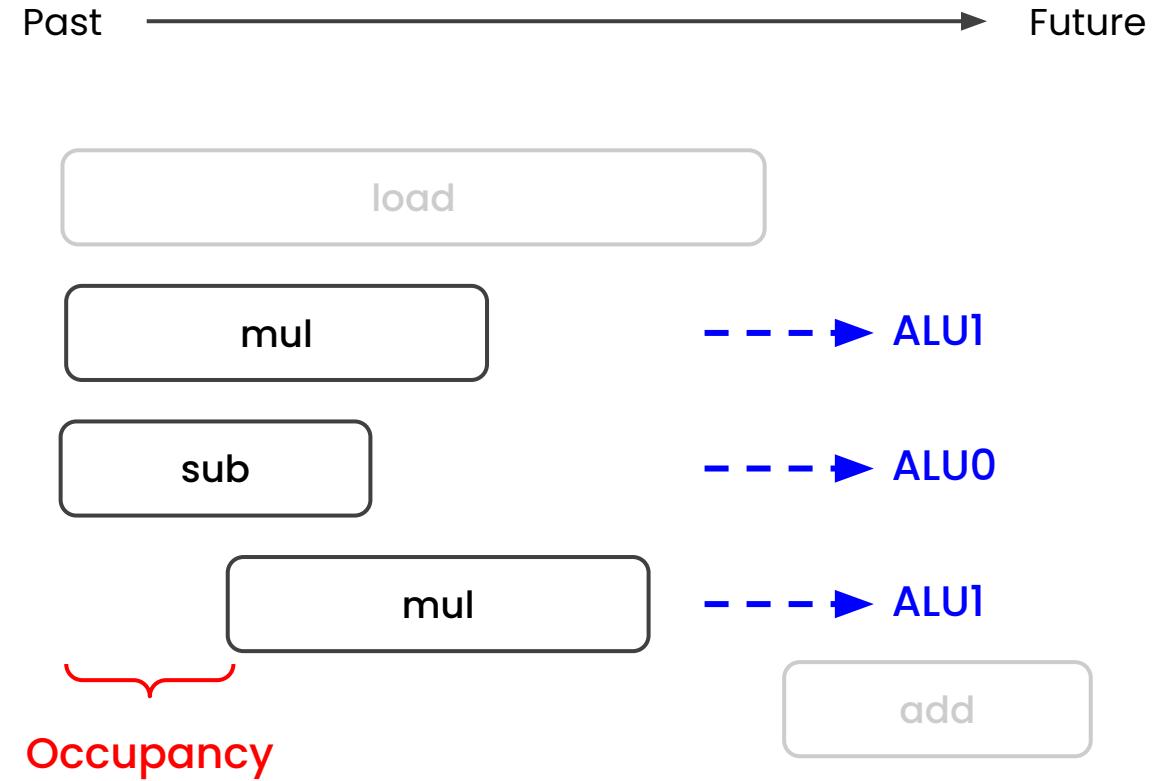
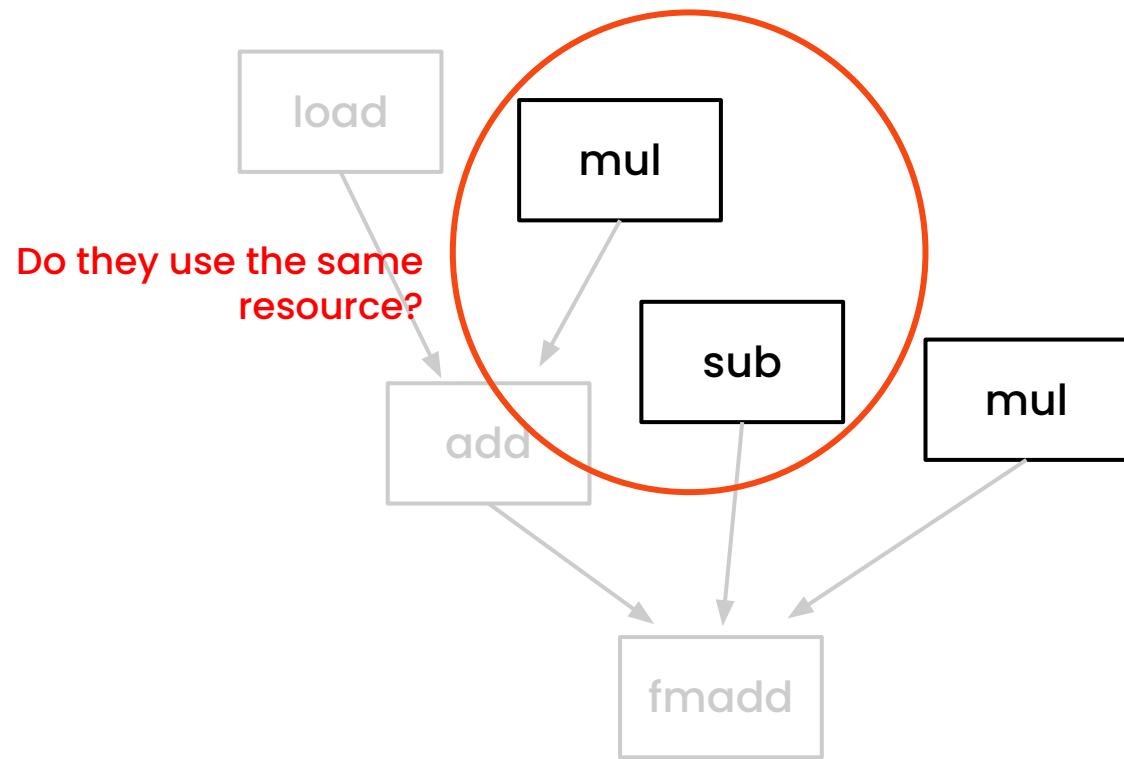
# Instruction Scheduling

Increase Instruction Level Parallelism



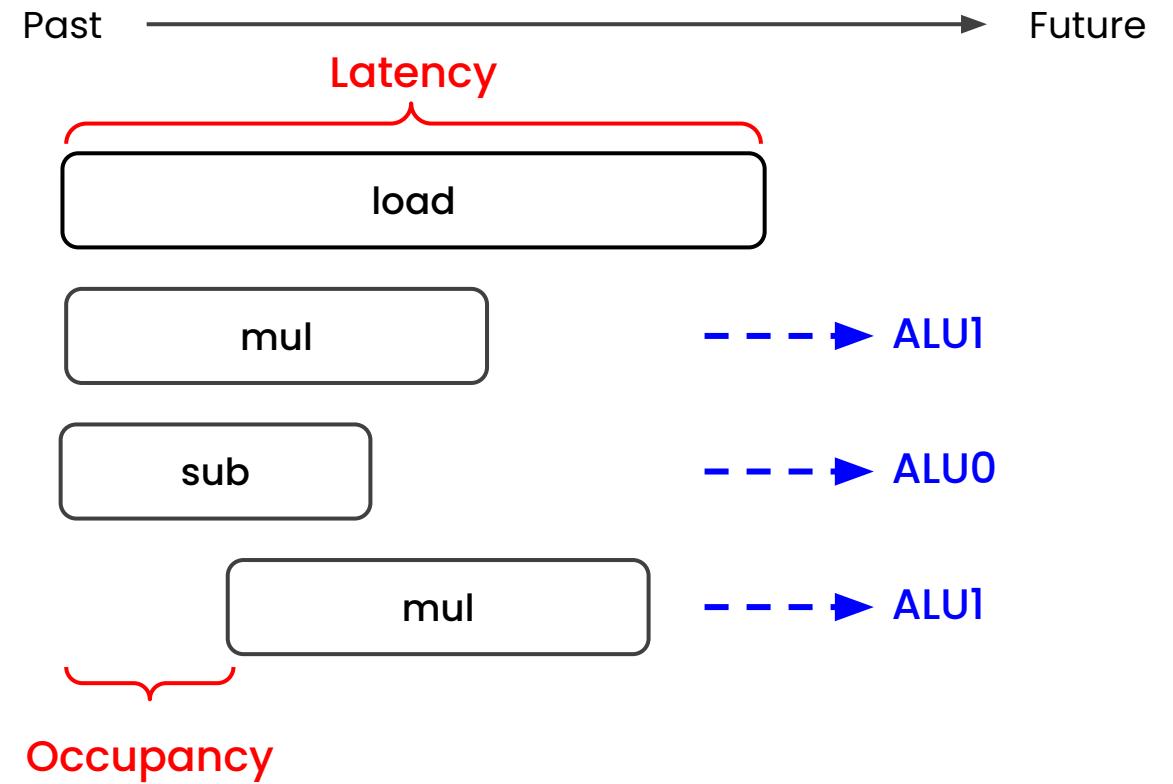
# Instruction Scheduling

Increase Instruction Level Parallelism



# Latency, Resource, and Occupancy

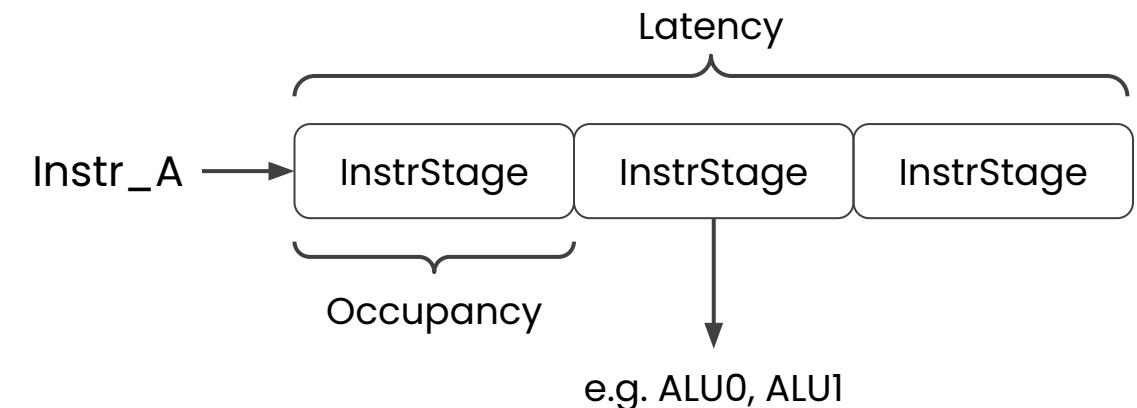
Instruction	Latency (cycles)	Resource	Occupancy (cycles)
load	10	LSU	2
mul	4	ALU1	2
sub	3	ALU0	1



# Itinerary Scheduling Model

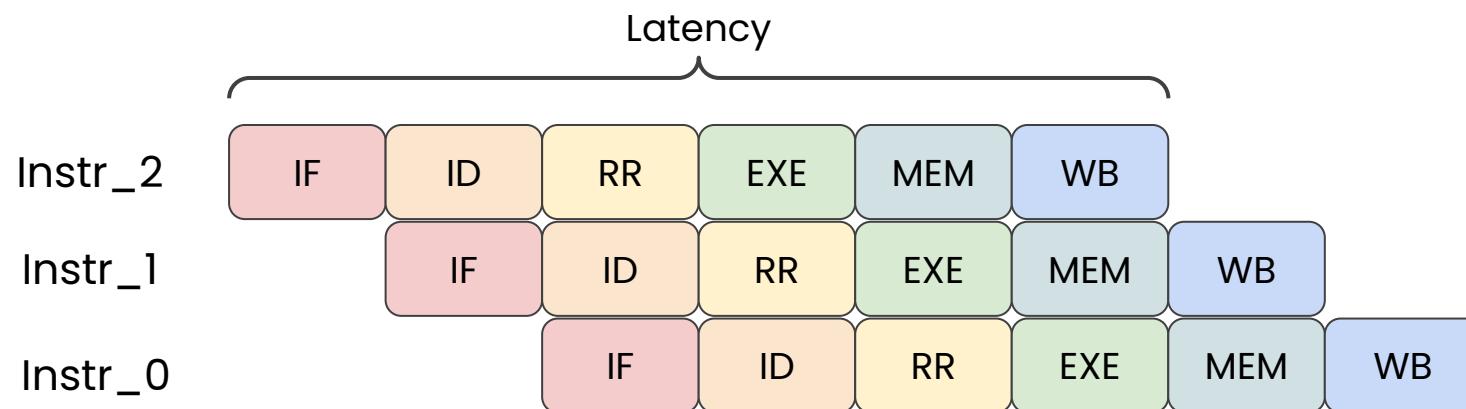
Legacy scheduling model framework

- Split instruction execution into *stages*
- Duration (i.e. occupancy) and resources used in each stage
- When will the result be available (i.e. latency)



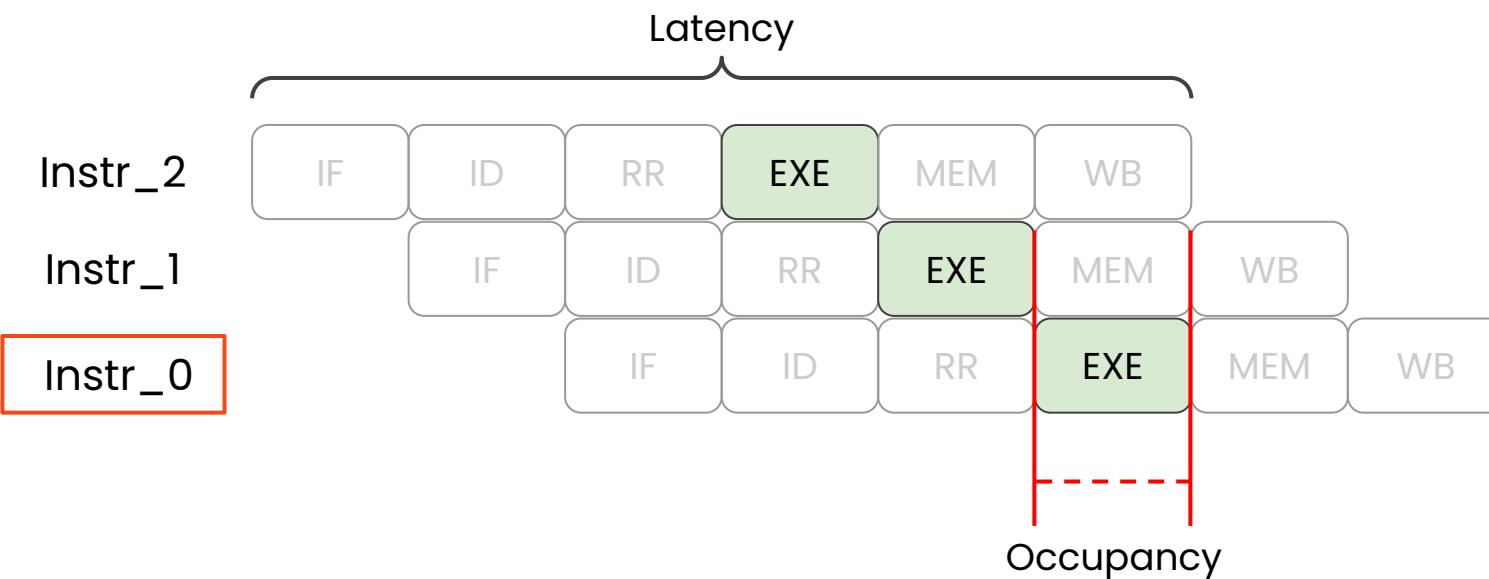
# Instruction Stages: From a Hardware Perspective

- IF: Instruction Fetch
- ID: Instruction Decode
- RR: Register Read
- EXE: Execution
- MEM: Memory
- WB: Write Back

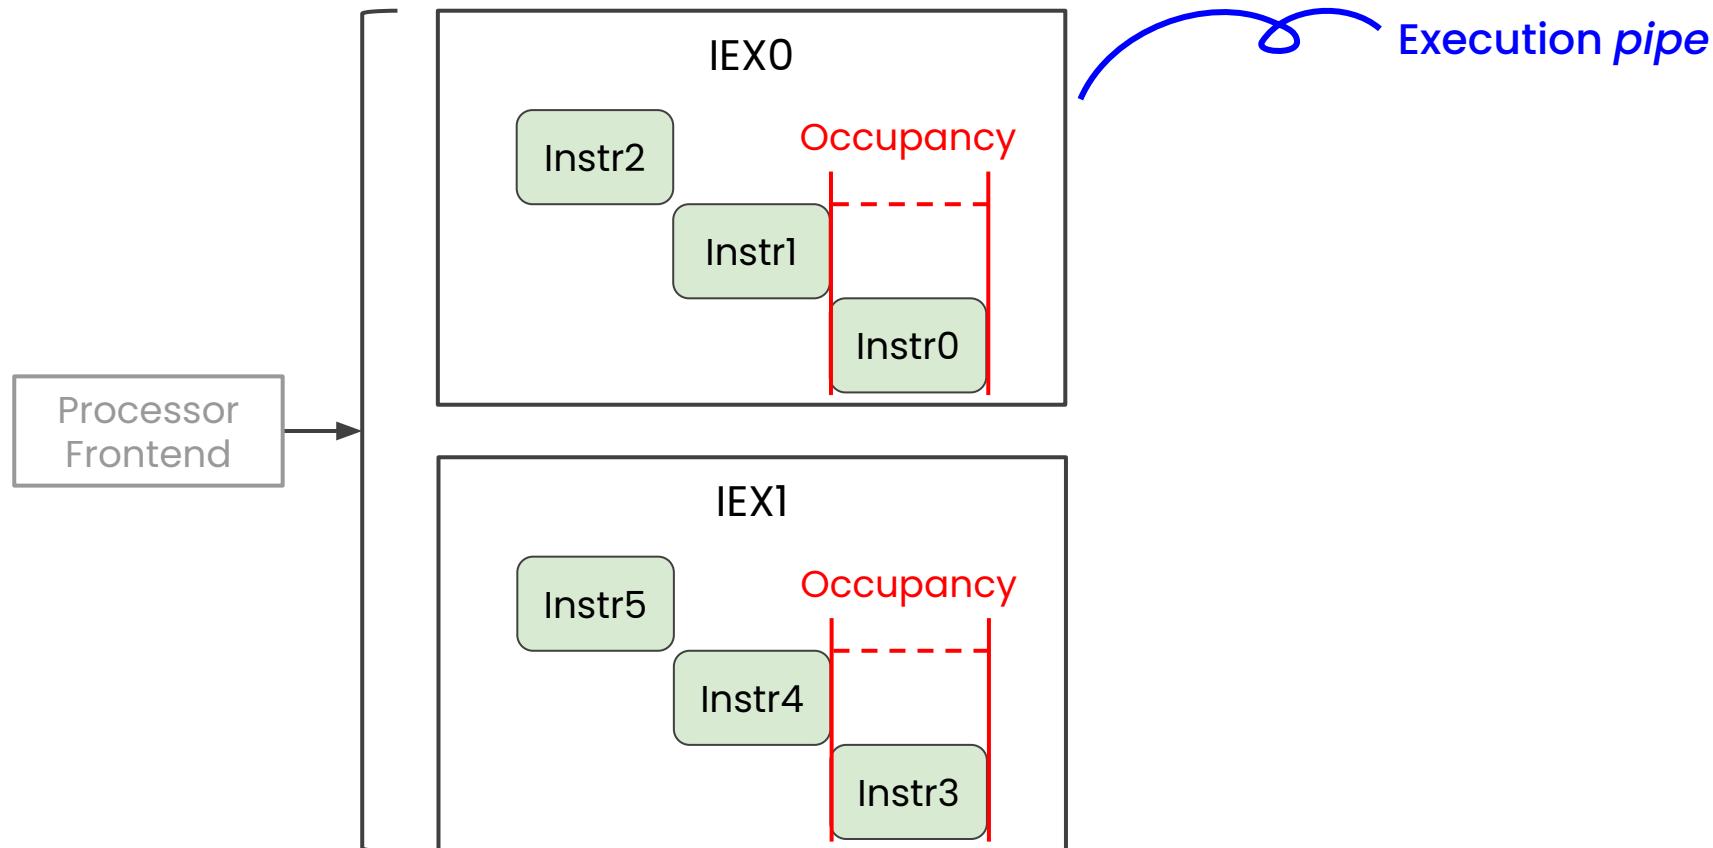


# Instruction Stages: From a Hardware Perspective

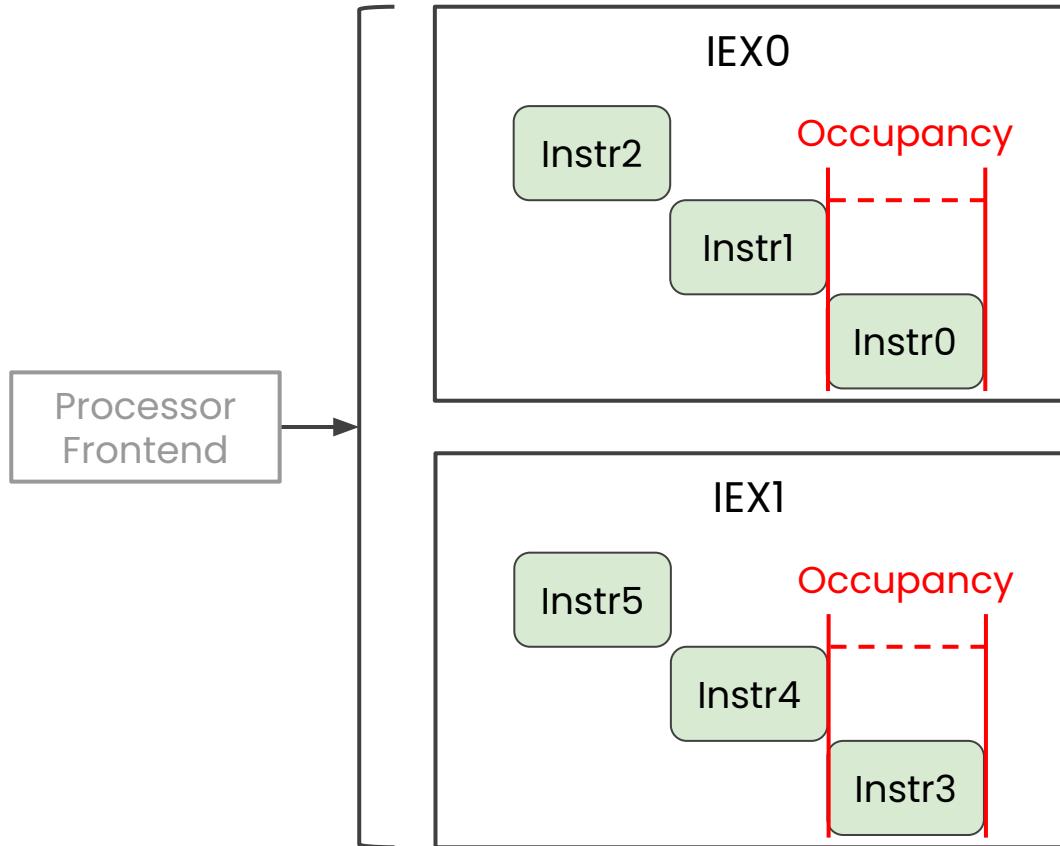
- IF: Instruction Fetch
- ID: Instruction Decode
- RR: Register Read
- EXE: Execution
- MEM: Memory
- WB: Write Back



# Instruction Stages with Superscalar Processors



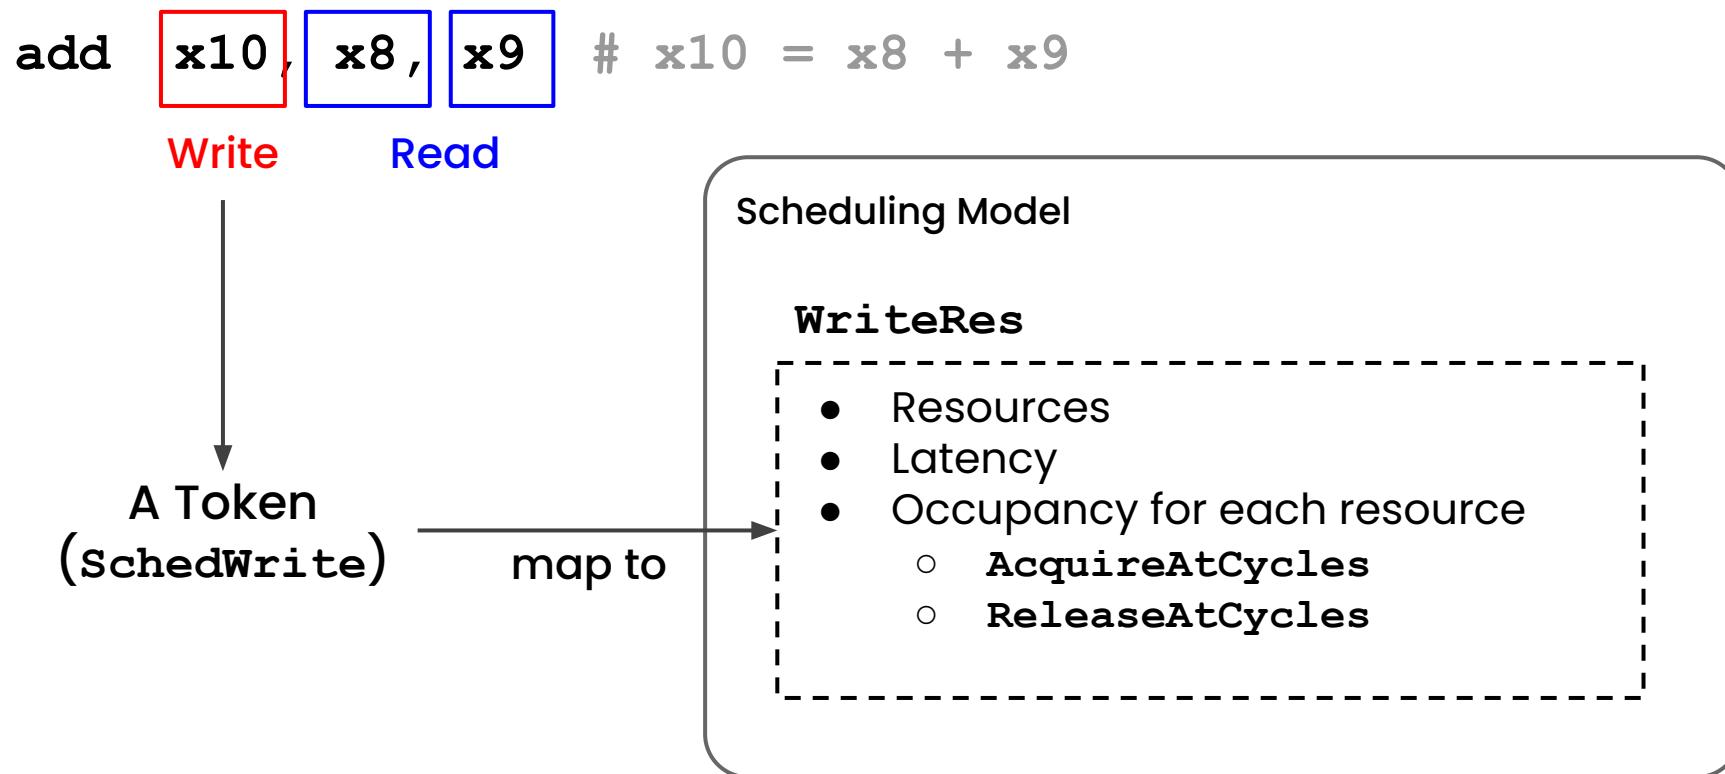
# Instruction Stages with Superscalar Processors



	Resource	Occupancy
Instr0	IEX0	2
Instr1	IEX0	2
Instr2	IEX0, IEX1	1
Instr3	IEX0, IEX1	1
Instr4	IEX1	3
Instr5	IEX0, IEX1	1

# Instruction Scheduling Model

Contemporary scheduling model framework



# Instruction Scheduling Model

## An Example

```
// In RISCVInstrInfo.td

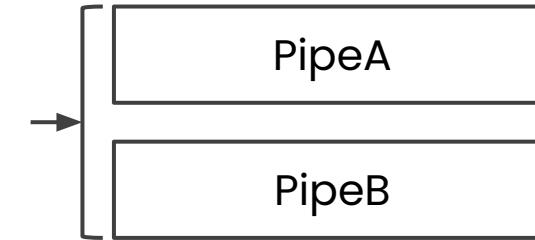
def ADD : ALU_rr<0b0000000, 0b000, "add", Commutable=1>,
          Sched<[WriteIALU, ReadIALU, ReadIALU]>;
```

```
// In RISCVSchedSiFive7.td*           Resource
                                         ↓
def : WriteRes<WriteIALU, [PipeAB]> {
    let Latency = 3;
    let AcquireAtCycles = [0];
    let ReleaseAtCycles = [1]; } Occupancy
}
```

# Instruction Scheduling Model

Processor resources

```
// In RISCVSchedSiFive7.td*
def PipeA : ProcResource<1>;
def PipeB : ProcResource<1>;
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
```



```
def : WriteRes<WriteIALU, [PipeAB]> {
    let Latency = 3;
    let AcquireAtCycles = [0];
    let ReleaseAtCycles = [1];
}
```

# Instruction Scheduling Model

Processor resources

```
// In RISCVSchedSiFive7.td*
```

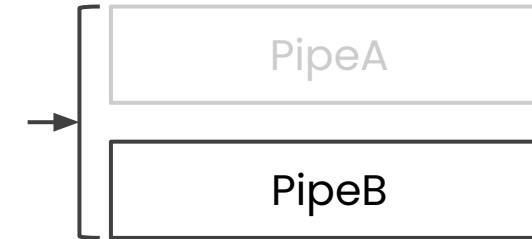
```
def PipeA : ProcResource<1>;
```

```
def PipeB : ProcResource<1>;
```

```
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
```

```
def : WriteRes<WriteIALU, [PipeAB]>;
```

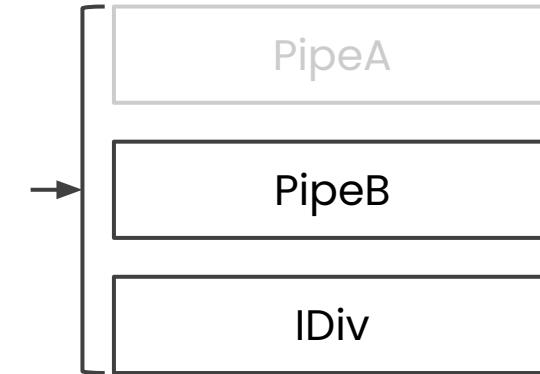
```
def : WriteRes<WriteIMul, [PipeB]>; ← Can only run on PipeB
```



# Instruction Scheduling Model

Processor resources

```
// In RISCVSchedSiFive7.td*
def PipeA : ProcResource<1>;
def PipeB : ProcResource<1>;
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
```

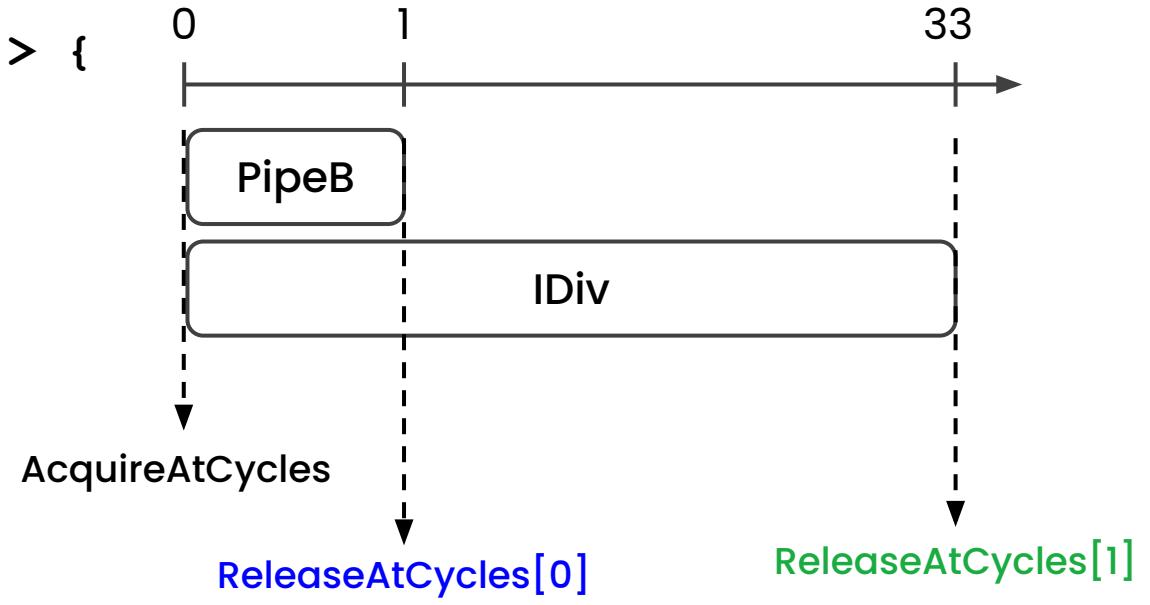


```
def : WriteRes<WriteIALU, [PipeAB]>;
def : WriteRes<WriteIMul, [PipeB]>;
def : WriteRes<WriteIDiv, [PipeB, IDiv]> ← Need PipeB and IDiv
```

# Instruction Scheduling Model

Occupancy: Acquire/ReleaseAtCycles

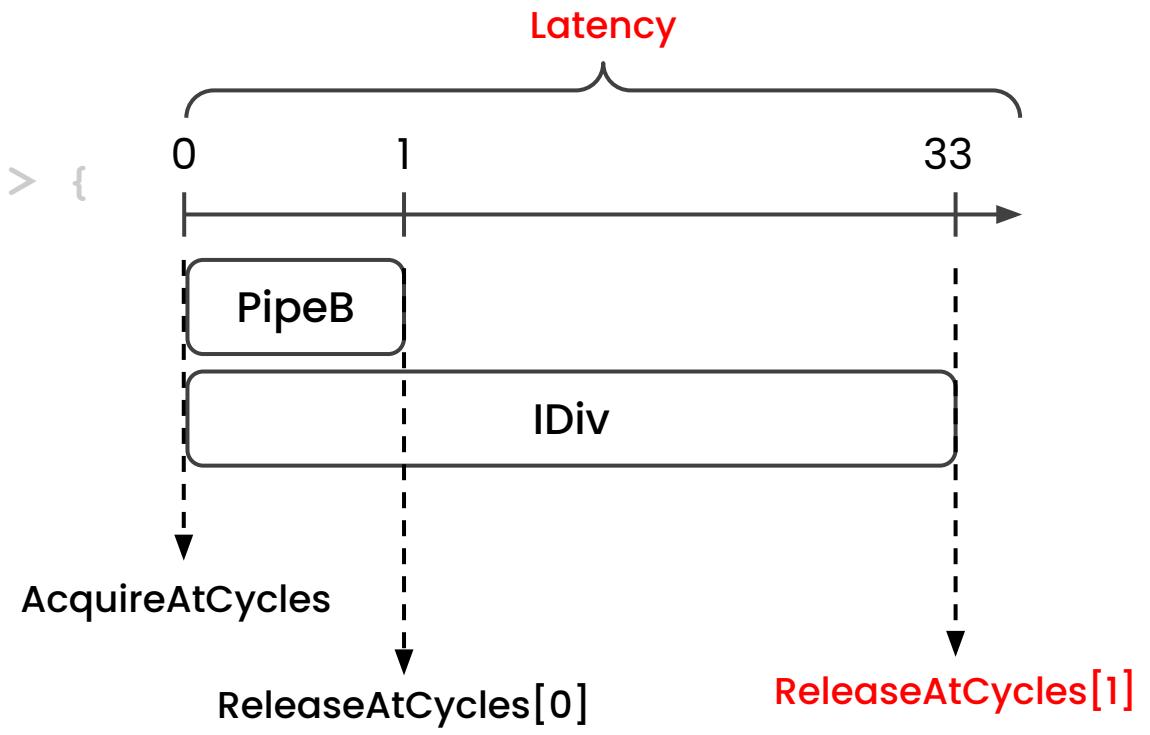
```
def : WriteRes<WriteIDiv, [PipeB, IDiv]> {  
    let Latency = 34;  
    let AcquireAtCycles = [0, 0];  
    let ReleaseAtCycles = [1, 33];  
}
```



# Instruction Scheduling Model

*Unpipelined instructions*

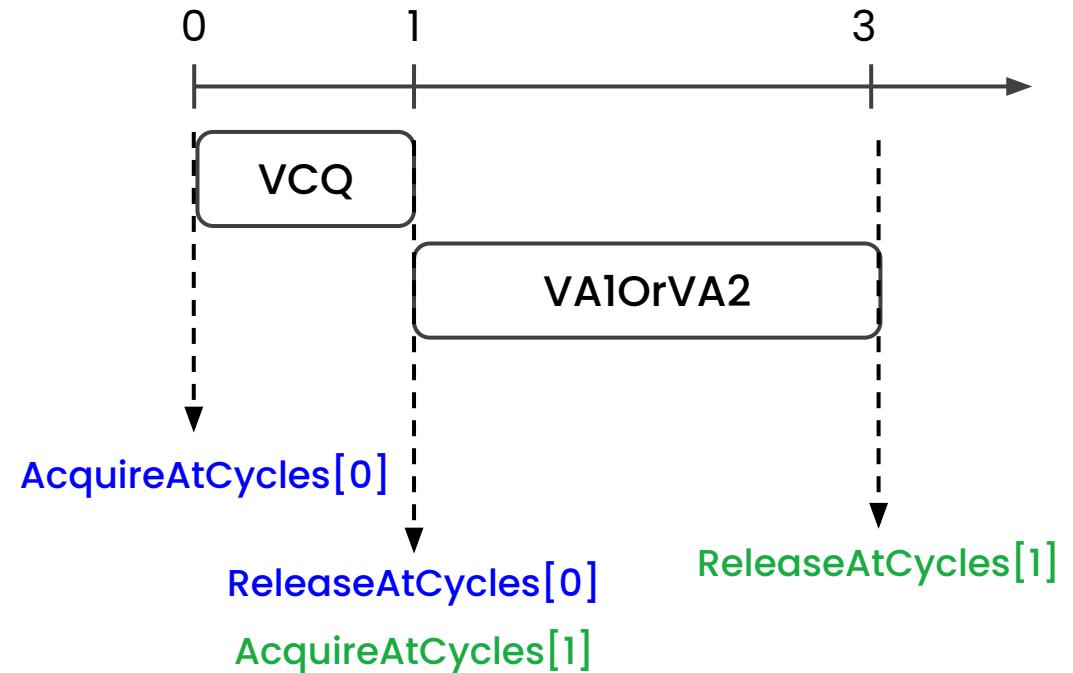
```
def : WriteRes<WriteIDiv, [PipeB, IDiv]> {  
    let Latency = 34;  
    let AcquireAtCycles = [0, 0];  
    let ReleaseAtCycles = [1, 33];  
}
```



# Instruction Scheduling Model

Occupancy: Acquire/ReleaseAtCycles

```
def : WriteRes<WriteVIALUV_M1,  
           [VCQ, VA1OrVA2]> {  
  
    let AcquireAtCycles = [0, 1];  
  
    let ReleaseAtCycles = [1, 3];  
  
}
```

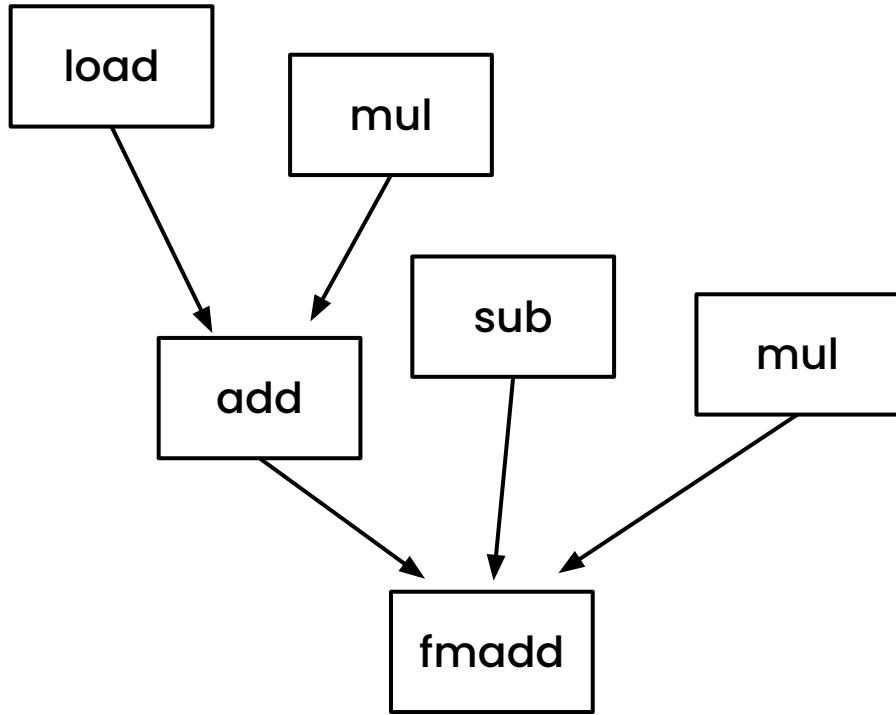


To Learn More:

<https://www.youtube.com/watch?v=XWBVLcdzmFg>

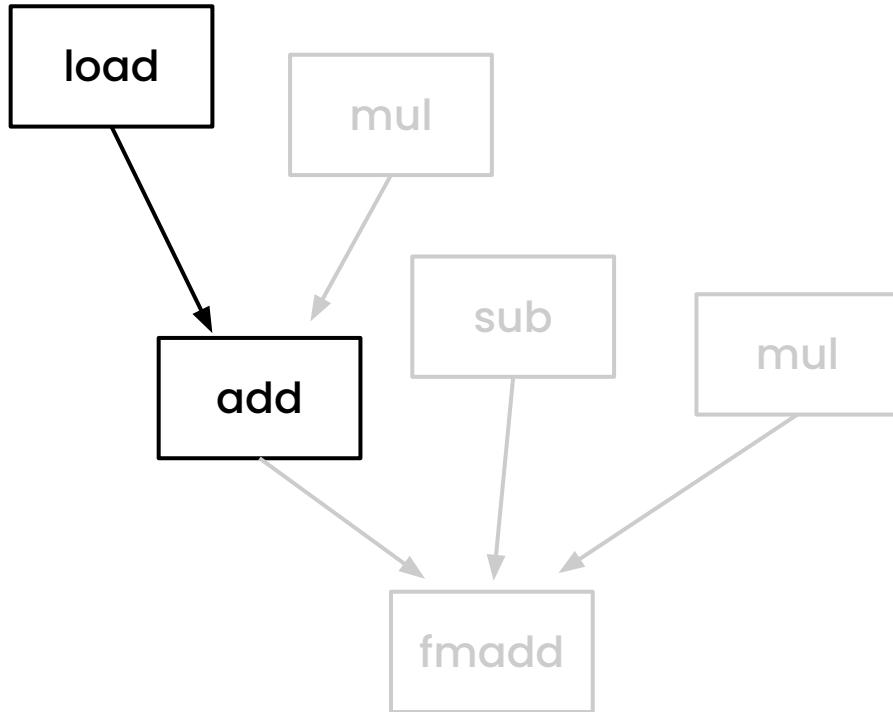
Instruction	Latency	Resource	Occupancy
load	10	LSU	2
mul	4	ALU0	2
sub	3	ALU1	1

```
def : WriteRes<WriteIMul, [ALU0]> {
    let Latency = 4;
    let AcquireAtCycles = [0];
    let ReleaseAtCycles = [2];
}
```



Instruction	Latency	Resource	Occupancy
load	10	LSU	2
mul	4	ALU0	2
sub	3	ALU1	1

```
def : WriteRes<WriteIMul, [ALU0]> {
    let Latency = 4;
    let AcquireAtCycles = [0];
    let ReleaseAtCycles = [2];
}
```



ld **x8**, 0 (**x8**)

add **x12**, **x9**, **x8**

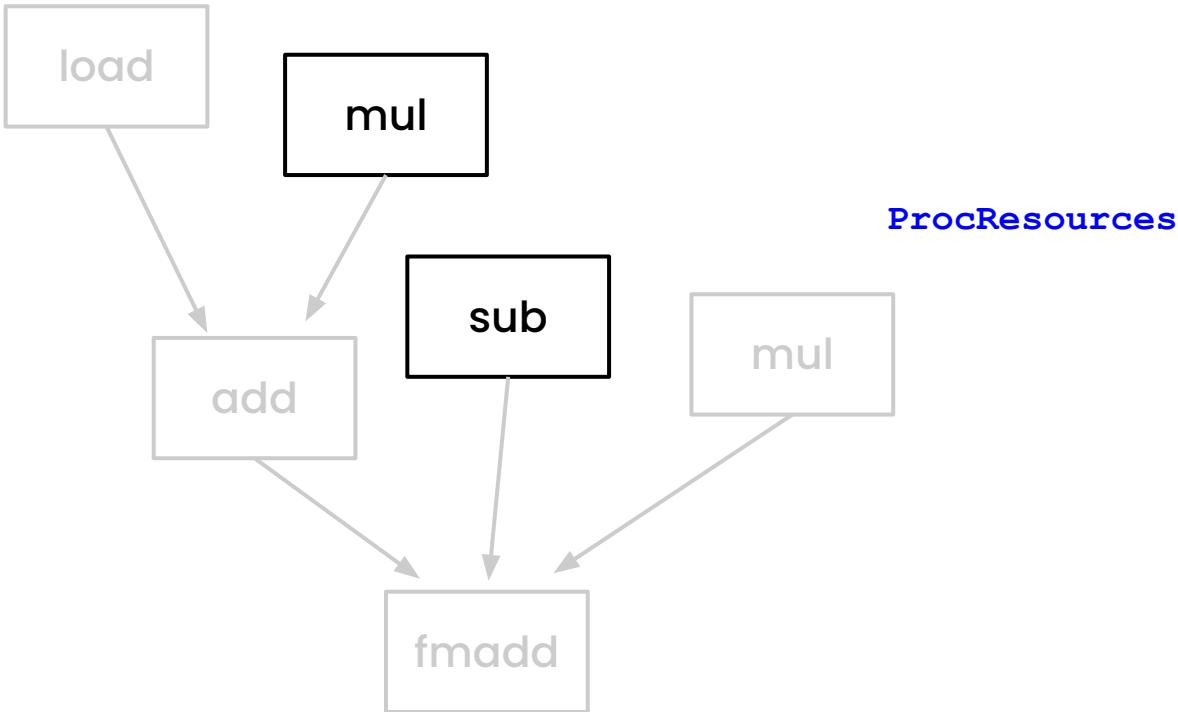
--

Latency

Avoid Data Hazard

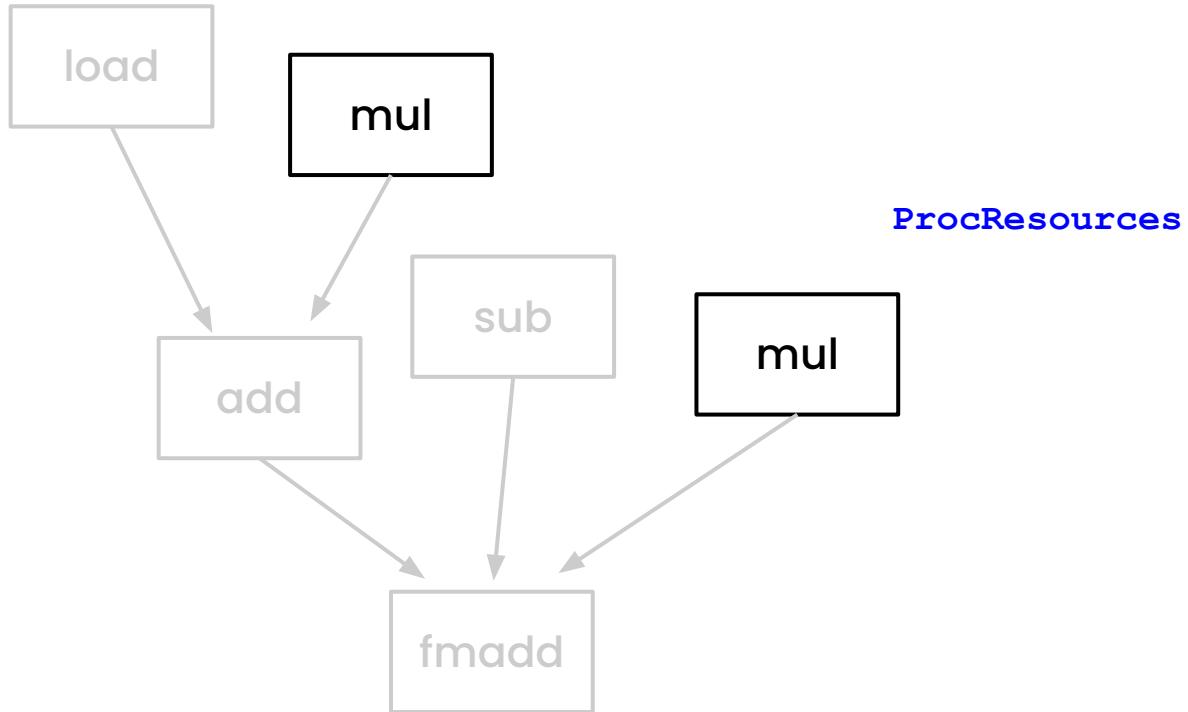
Instruction	Latency	Resource	Occupancy
load	10	LSU	2
mul	4	ALU0	2
sub	3	ALU1	1

```
def : WriteRes<WriteIMul, [ALU0]> {
    let Latency = 4;
    let AcquireAtCycles = [0];
    let ReleaseAtCycles = [2];
}
```



```
ld x8, 0(x8)
mul x9, x9, x7
sub x10, x10, x7
add x12, x9, x8
```

Instruction	Latency	Resource	Occupancy
load	10	LSU	2
mul	4	ALU0	2
sub	3	ALU1	1



```

def : WriteRes<WriteIMul, [ALU0]> {
    let Latency = 4;
    let AcquireAtCycles = [0];
    let ReleaseAtCycles = [2];
}
  
```

ld x8, 0(x8)

mul x9, x9, x7

sub x10, x10, x7

mul x11, x11, x7

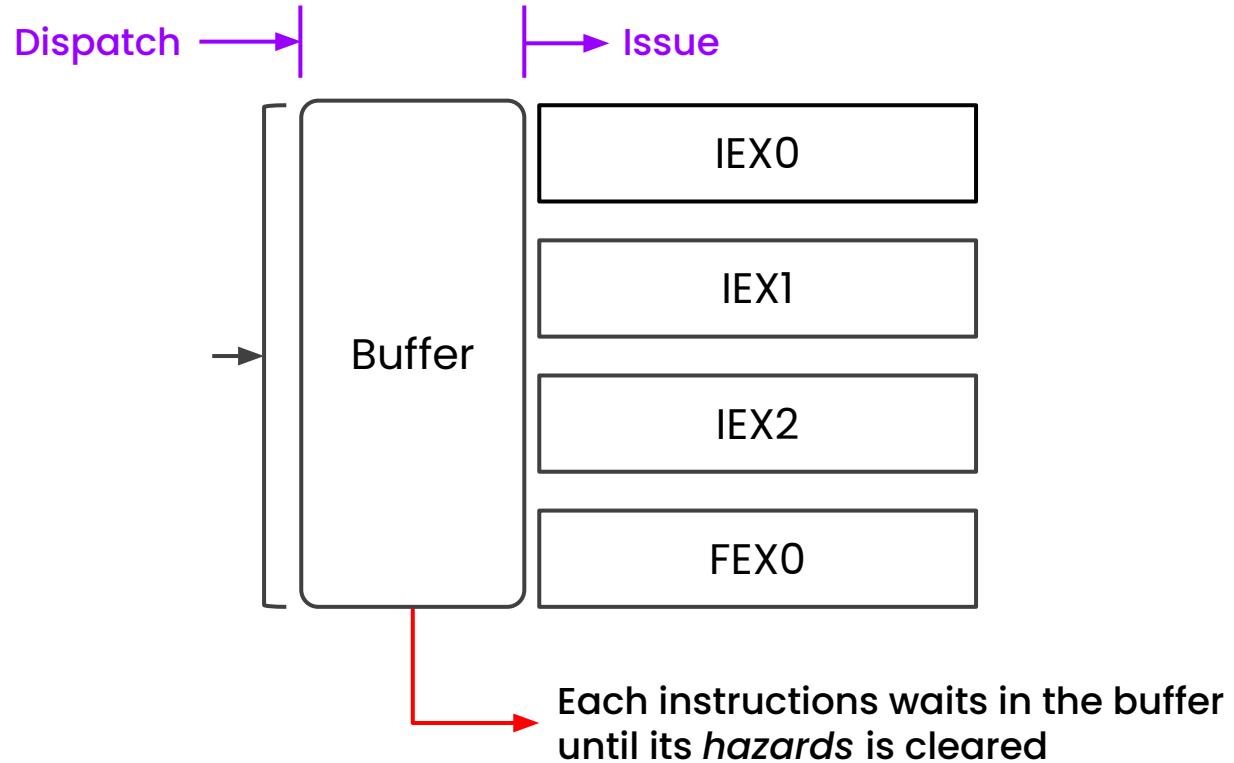
add x12, x9, x8

AcquireAtCycles  
ReleaseAtCycles

Avoid Structural Hazard

# Out-of-Order Execution

Hardware solution for avoiding hazards

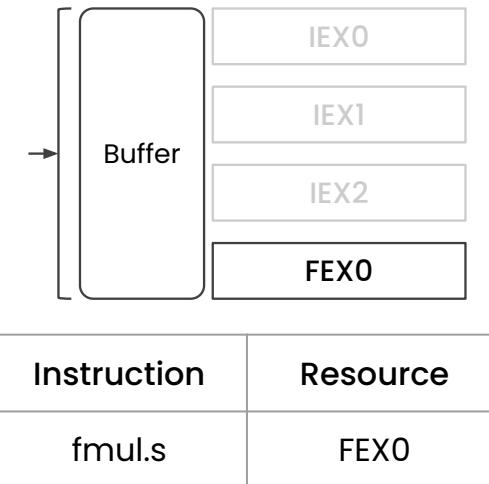


# Out-of-Order Execution

Case study with llvm-mca: *data hazard*

```
$ llvm-mca -mtriple=riscv64 -mcpu=sifive-p470 -timeline
```

Index	0123456789	Cycles
[0, 0]	DeeeeER .	fmul.s ft0, ft0, ft1
[0, 1]	D====eeeeER	fmul.s ft2, ft2, ft0
[0, 2]	D=eeeeE---R	fmul.s ft3, ft3, ft4
Simulation timeline		



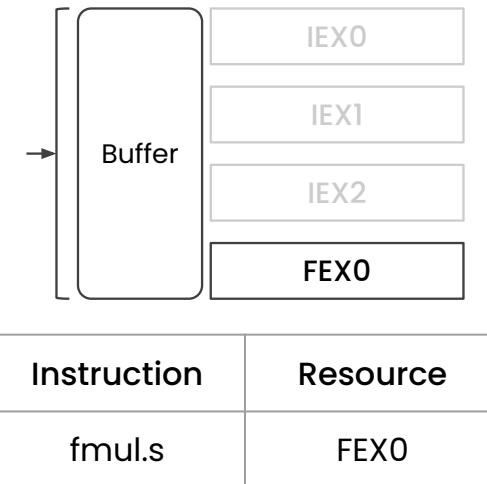
- **D**: dispatch
- **=**: stall
- **e**: executing
- **E**: end execution
- **R**: retire

# Out-of-Order Execution

Case study with llvm-mca: *data hazard*

```
$ llvm-mca -mtriple=riscv64 -mcpu=sifive-p470 -timeline
```

Index	0123456789										
[0, 0]	DeeeeER .	fmul.s	ft0,	ft0,	ft1						
[0, 1]	D=====eeeeER	fmul.s	ft2,	ft2,	ft0						
[0, 2]	D=eeeeE---R	fmul.s	ft3,	ft3,	ft4						



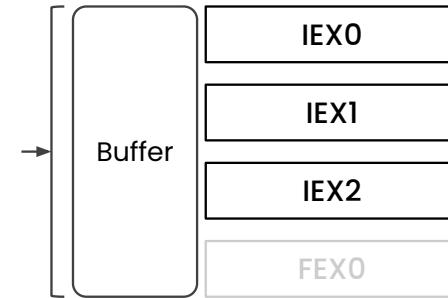
- **D:** dispatch
- **=:** stall
- **e:** executing
- **E:** end execution
- **R:** retire

# Out-of-Order Execution

Case study with llvm-mca: *structural hazard*

```
$ llvm-mca -mtriple=riscv64 -mcpu=sifive-p470 -timeline
```

Index	0123456	
[0, 0]	DeeER..	mul s0, s0, s1
[0, 1]	D=eeER.	mul a0, a0, a1
[0, 2]	D==eeER	mul a2, a2, a3
[0, 3]	.DeE--R	add s9, s9, s10
[0, 4]	.DeE--R	add s11, s11, t3



Instruction	Resource
add	IEX0, IEX1, IEX2
mul	IEX2

- **D:** dispatch
- **=:** stall
- **e:** executing
- **E:** end execution
- **R:** retire

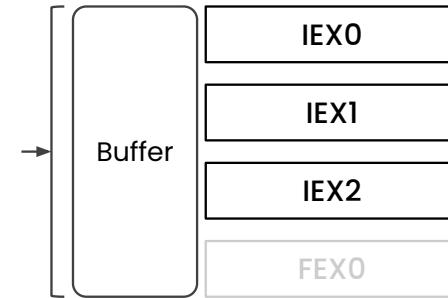
# Out-of-Order Execution

Case study with llvm-mca: *structural* hazard

```
$ llvm-mca -mtriple=riscv64 -mcpu=sifive-p470 -timeline
```

Index	0123456	
[0, 0]	DeeER..	mul s0, s0, s1
[0, 1]	D=eeER.	mul a0, a0, a1
[0, 2]	D==eEER	mul a2, a2, a3
[0, 3]	.DeE--R	add s9, s9, s10
[0, 4]	.DeE--R	add s11, s11, t3

- **D**: dispatch
- **=**: stall
- **e**: executing
- **E**: end execution
- **R**: retire



Instruction	Resource
add	IEX0, IEX1, IEX2
mul	IEX2

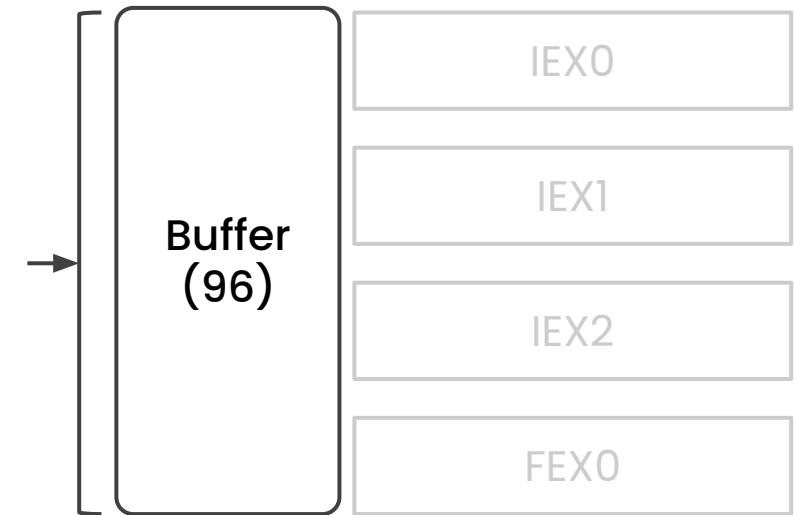
# Scheduling Model: Buffer Sizes

Unified Reservation Station

```
// In RISCVSchedSiFiveP400.td*

def SiFiveP400Model : SchedMachineModel {
    let MicroOpBufferSize = 96;
}

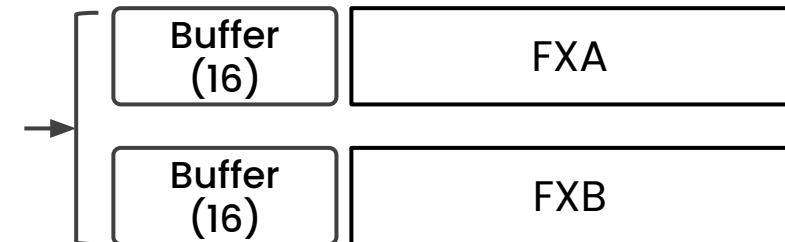
let BufferSize = -1 in {
    def IEXQ0  : ProcResource<1>;
    def IEXQ1  : ProcResource<1>;
    def IEXQ2  : ProcResource<1>;
    def FEXQ0  : ProcResource<1>;
}
```



# Scheduling Model: Buffer Sizes

Decoupled Reservation Station

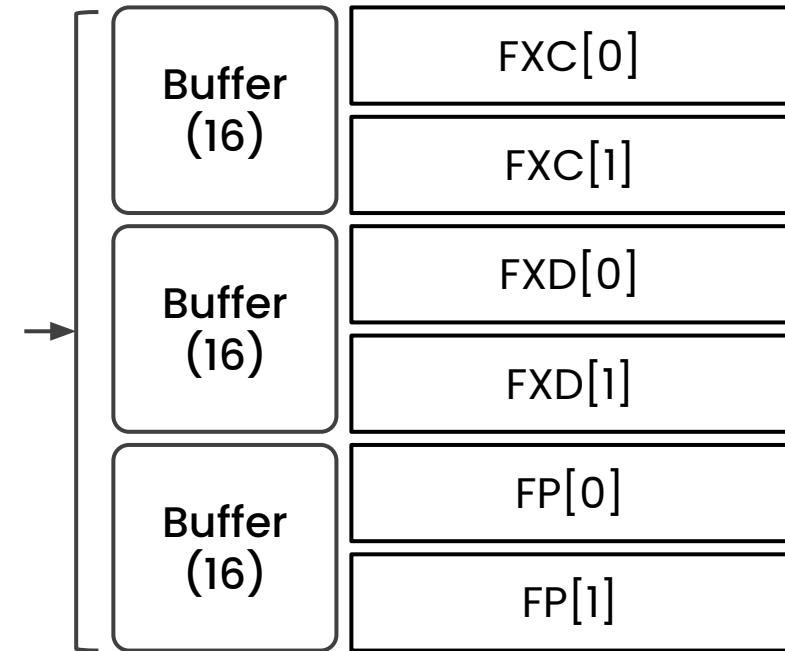
```
// In RISCVSchedTTAscalonD8.td*
let BufferSize = 16 in {
    def AscalonFXA : ProcResource<1>;
    def AscalonFXB : ProcResource<1>;
    def AscalonFXC : ProcResource<2>;
    def AscalonFXD : ProcResource<2>;
    def AscalonFP  : ProcResource<2>;
}
```



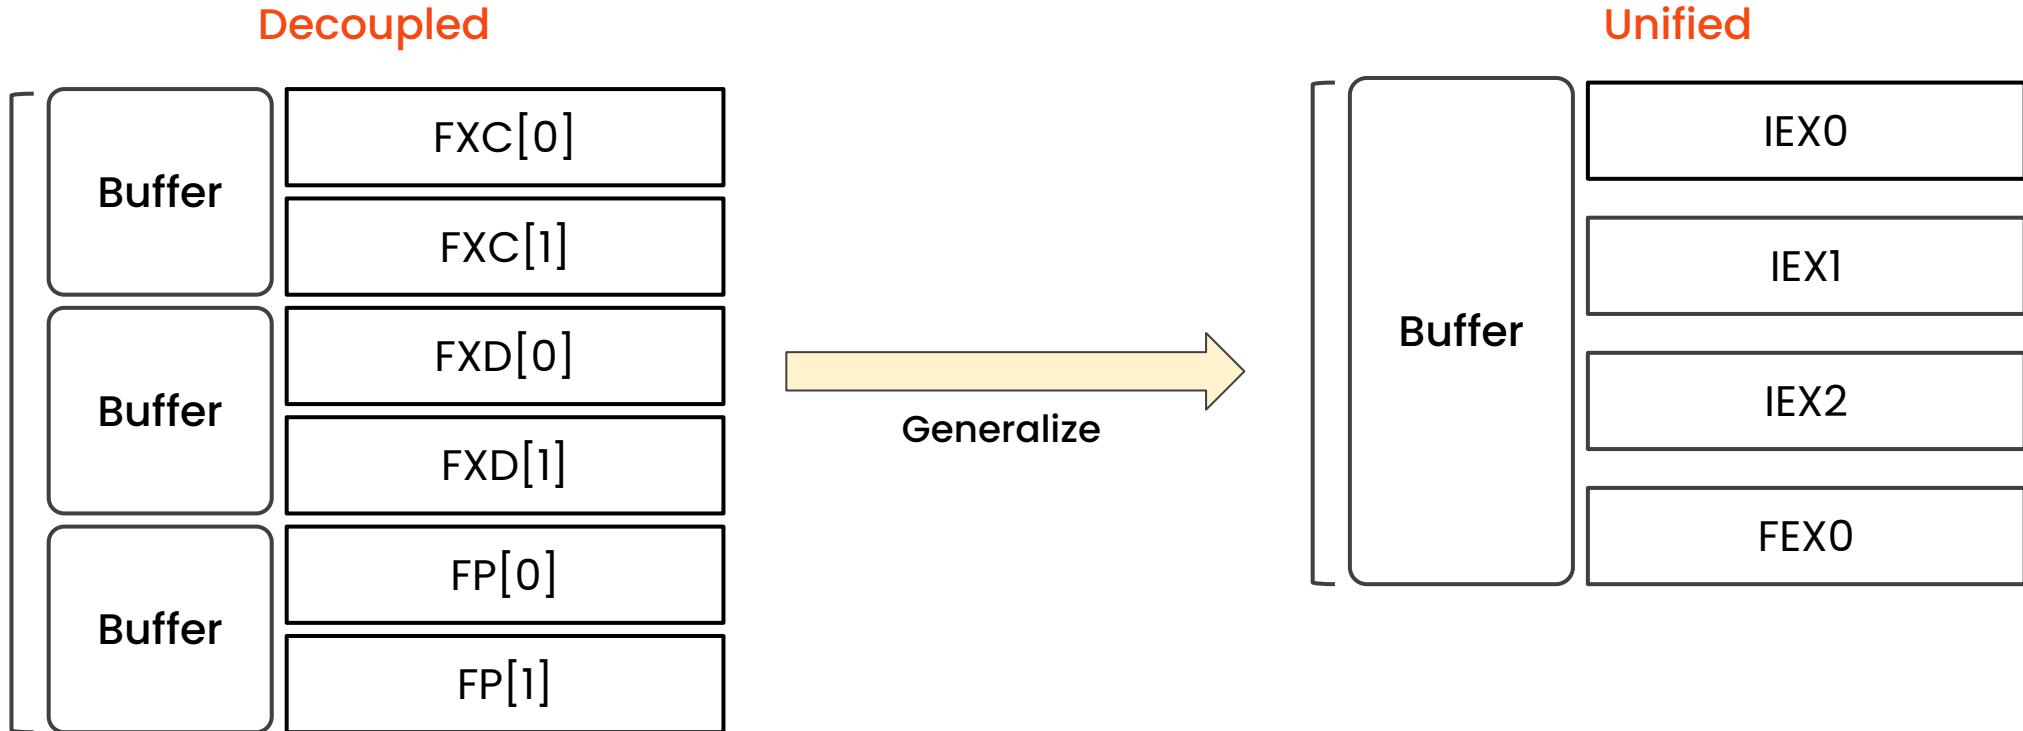
# Scheduling Model: Buffer Sizes

Decoupled Reservation Station

```
// In RISCVSchedTTAscalonD8.td*
let BufferSize = 16 in {
    def AscalonFXA : ProcResource<1>;
    def AscalonFXB : ProcResource<1>;
    def AscalonFXC : ProcResource<2>;
    def AscalonFXD : ProcResource<2>;
    def AscalonFP  : ProcResource<2>;
}
```



# Scheduling Model: Buffer Sizes



# Scheduling Model: Buffer Sizes

*In-order cores*

```
// In RISCVSchedSiFive7.td*
let BufferSize = 0 in {
    def PipeA : ProcResource<1>;
    def PipeB : ProcResource<1>;
}
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
```

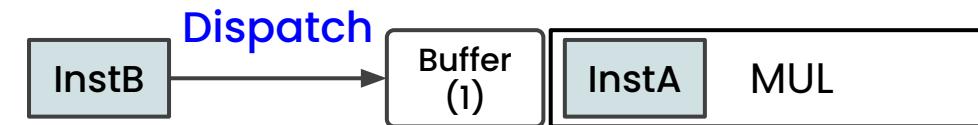
No issue/dispatch



# Scheduling Model: Buffer Sizes

*In-order cores*

```
// In RISCVSchedSiFive7.td*
let BufferSize = 0 in {
    def PipeA : ProcResource<1>;
    def PipeB : ProcResource<1>;
}
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
```

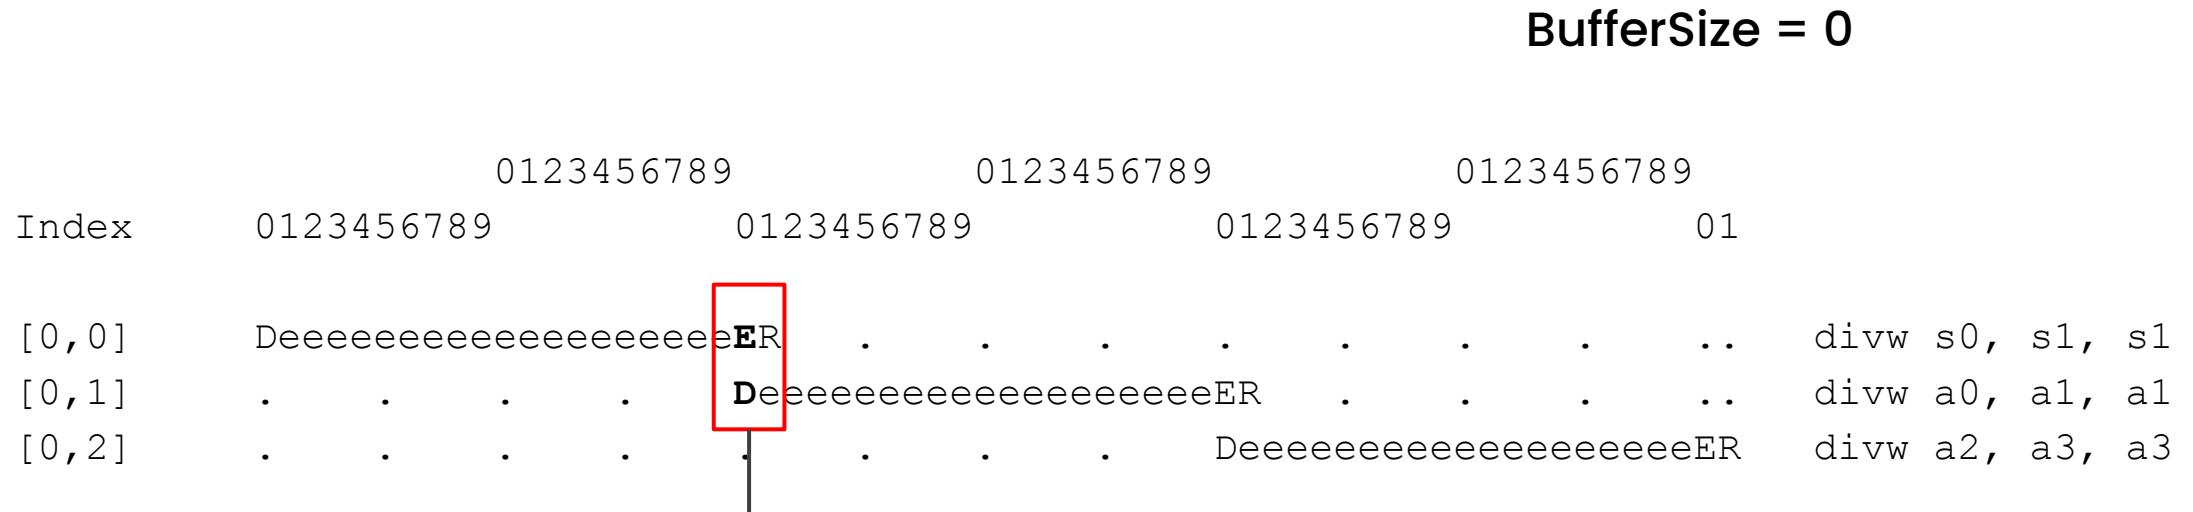


```
// In RISCVSchedSyntacoreSCR7.td*
def ALU_MUL_IS : ProcResource<1> { let BufferSize = 8; }
def ALU_DIV_IS : ProcResource<1> { let BufferSize = 8; }

def MUL : ProcResource<1> { let BufferSize = 1; }
def DIV : ProcResource<1> { let BufferSize = 1; }
```

# Scheduling Model: Buffer Sizes

In-order cores



- **D:** dispatch
- **=:** stall
- **e:** executing
- **E:** end execution
- **R:** retire

# Scheduling Model: Buffer Sizes

In-order cores

```
$ llvm-mca -mtriple=riscv64 -mcpu=syntaxcore-scr7 -timeline
```

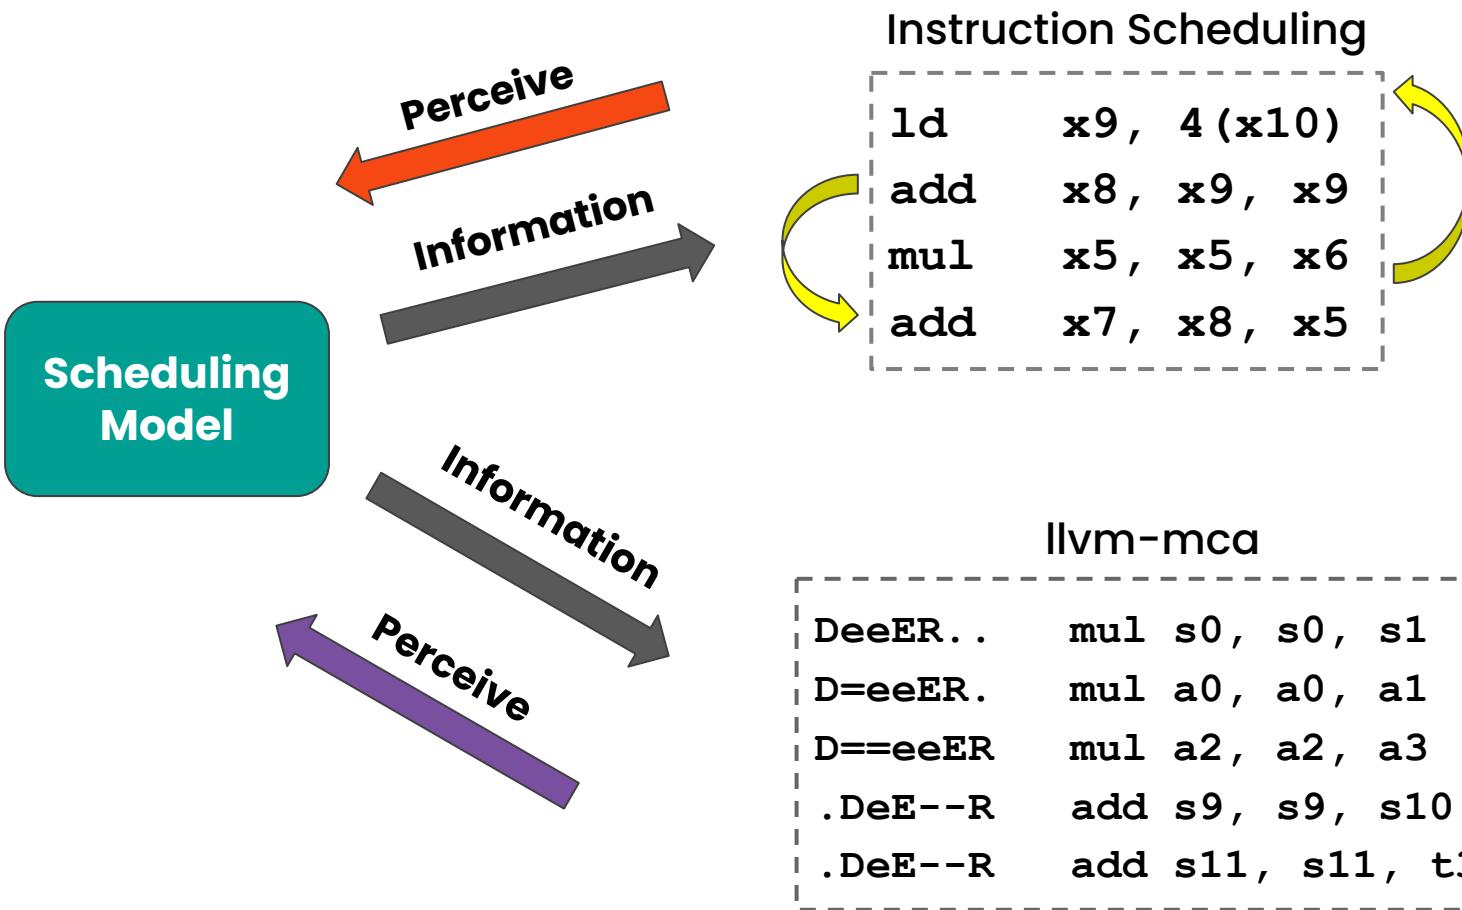
**BufferSize = 1**

	0123456789	0123456789	0123456789	
Index	0123456789	0123456789	0123456789	
[0, 0]	DeeeeeeeeeeeeeeeER	.	.	divw s0, s1, s1
[0, 1]	.D=====eeeeeeeeeeeeeeeER	.	.	divw a0, a1, a1
[0, 2]	.	D=====eeeeeeeeeeeeeeeER	.	divw a2, a3, a3



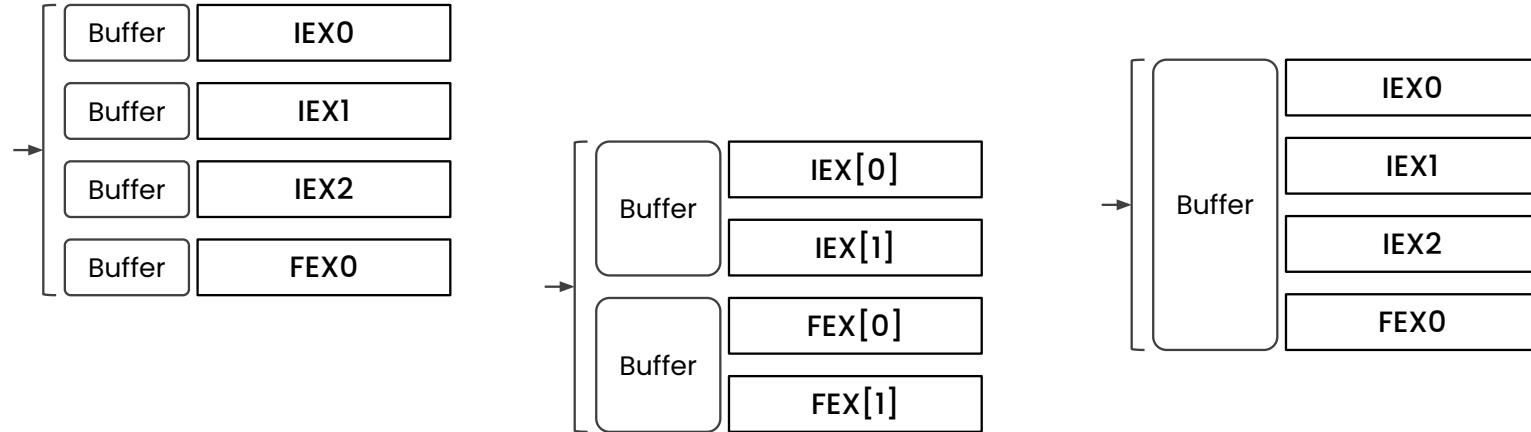
Dispatch right after the first instruction was *issued*, wait inside the buffer

- **D:** dispatch
- **=:** stall
- **e:** executing
- **E:** end execution
- **R:** retire



# Q: How do you leverage different OoO Buffer Sizes?

llvm-mca:



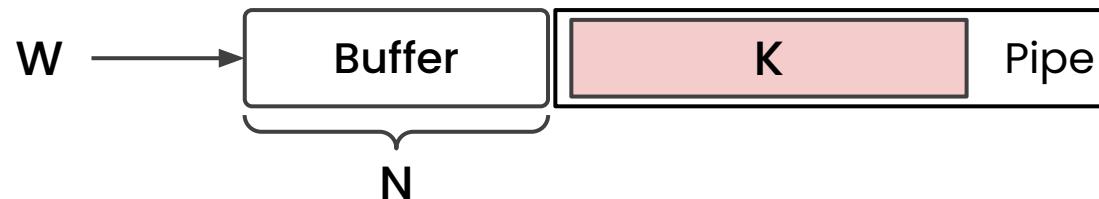
Machine Scheduler: “That’s the neat part, *WE DON’T*”

# Machine Scheduler: Out-of-Order Cores

- Does the *minimal* amount of efforts to predict hazards
  - Rationale: Nearly **impossible** to predict during compile-time

# Machine Scheduler: Out-of-Order Cores

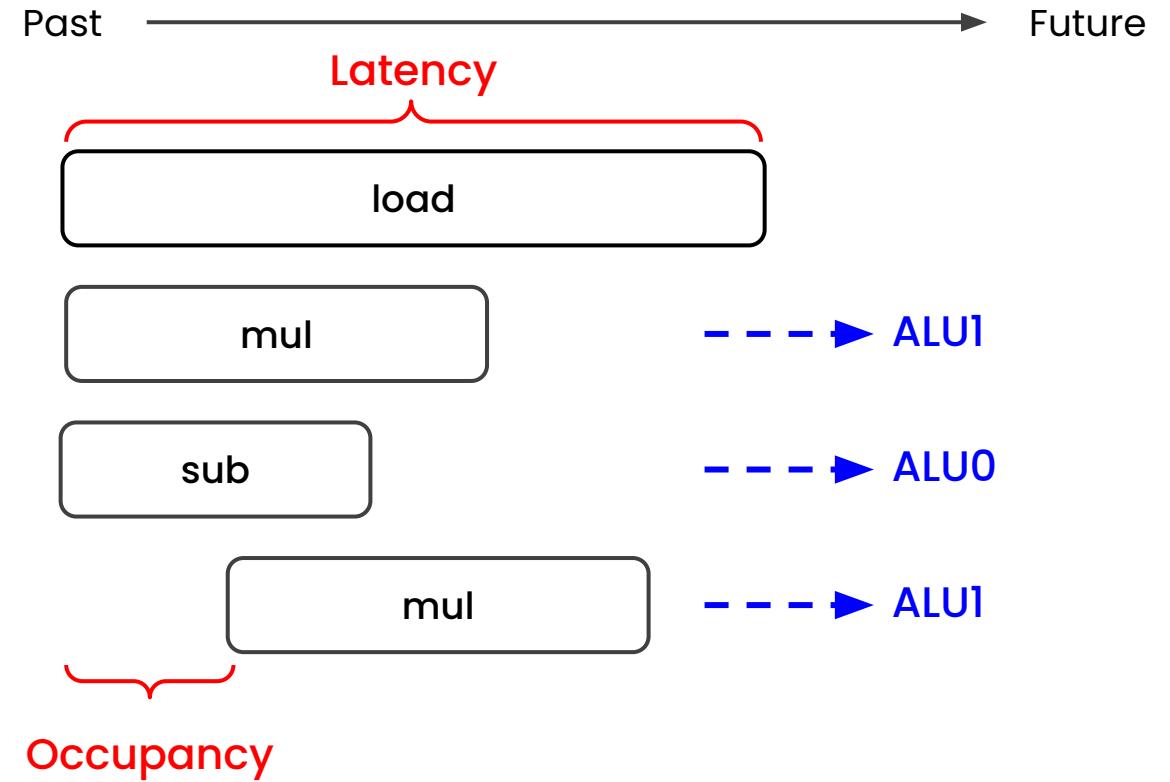
- Does the *minimal* amount of efforts to predict hazards
  - Rationale: Nearly impossible to predict during compile-time
- The condition where a buffer might be full:  $N < W*K - 1$ 
  - $N$ : buffer size
  - $W$ : number of instructions that go into this buffer per cycle
  - $K$ : *largest occupancy* that can execute in this pipe



# Machine Scheduler: In-Order Cores

MicroOpBufferSize & BufferSize = 0

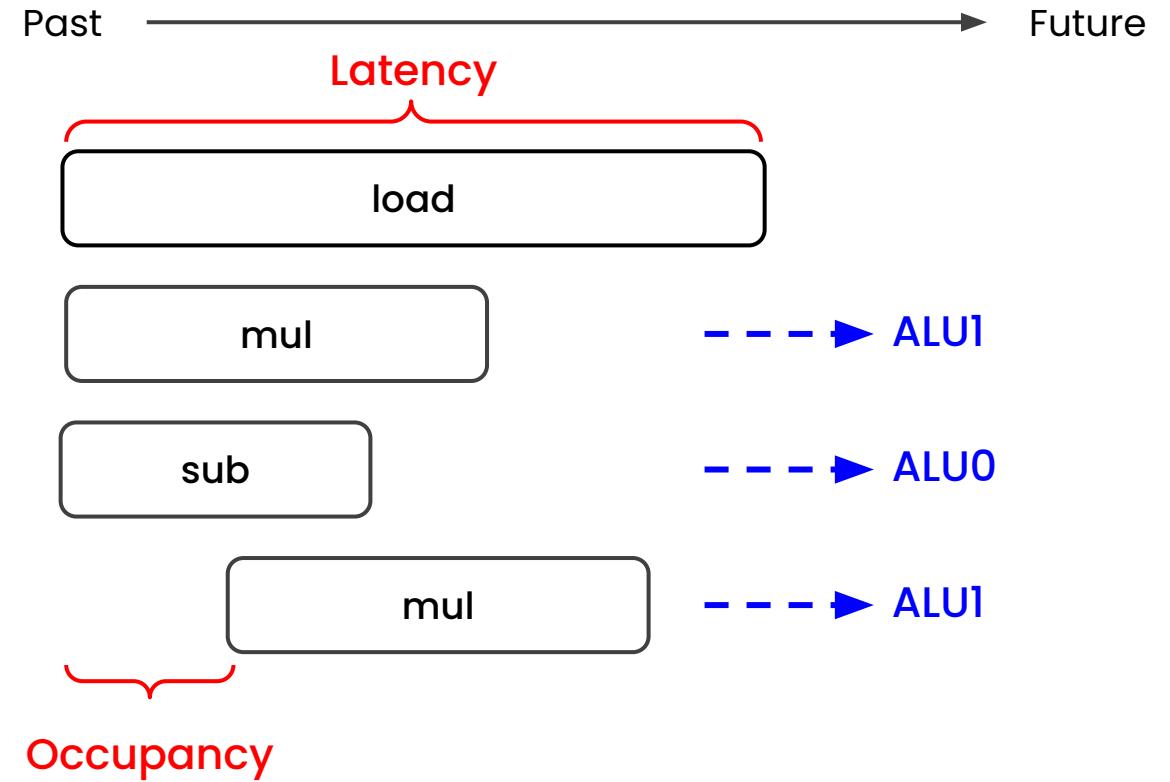
- We can predict stallings more accurately during *compile-time*



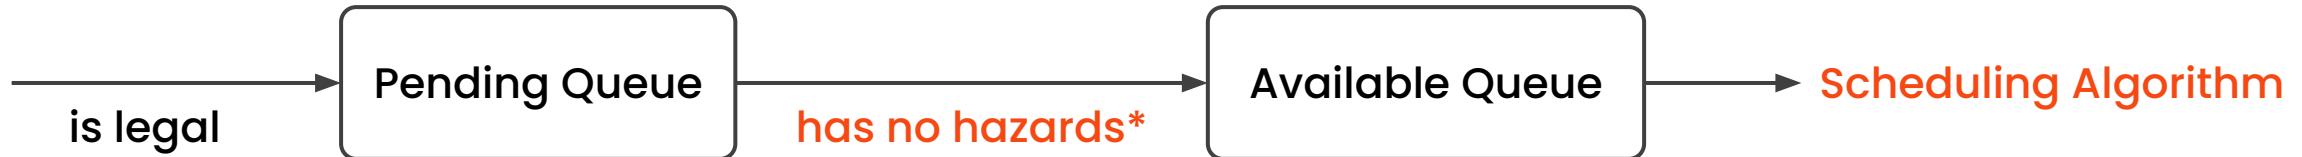
# Machine Scheduler: In-Order Cores

MicroOpBufferSize & BufferSize = 0

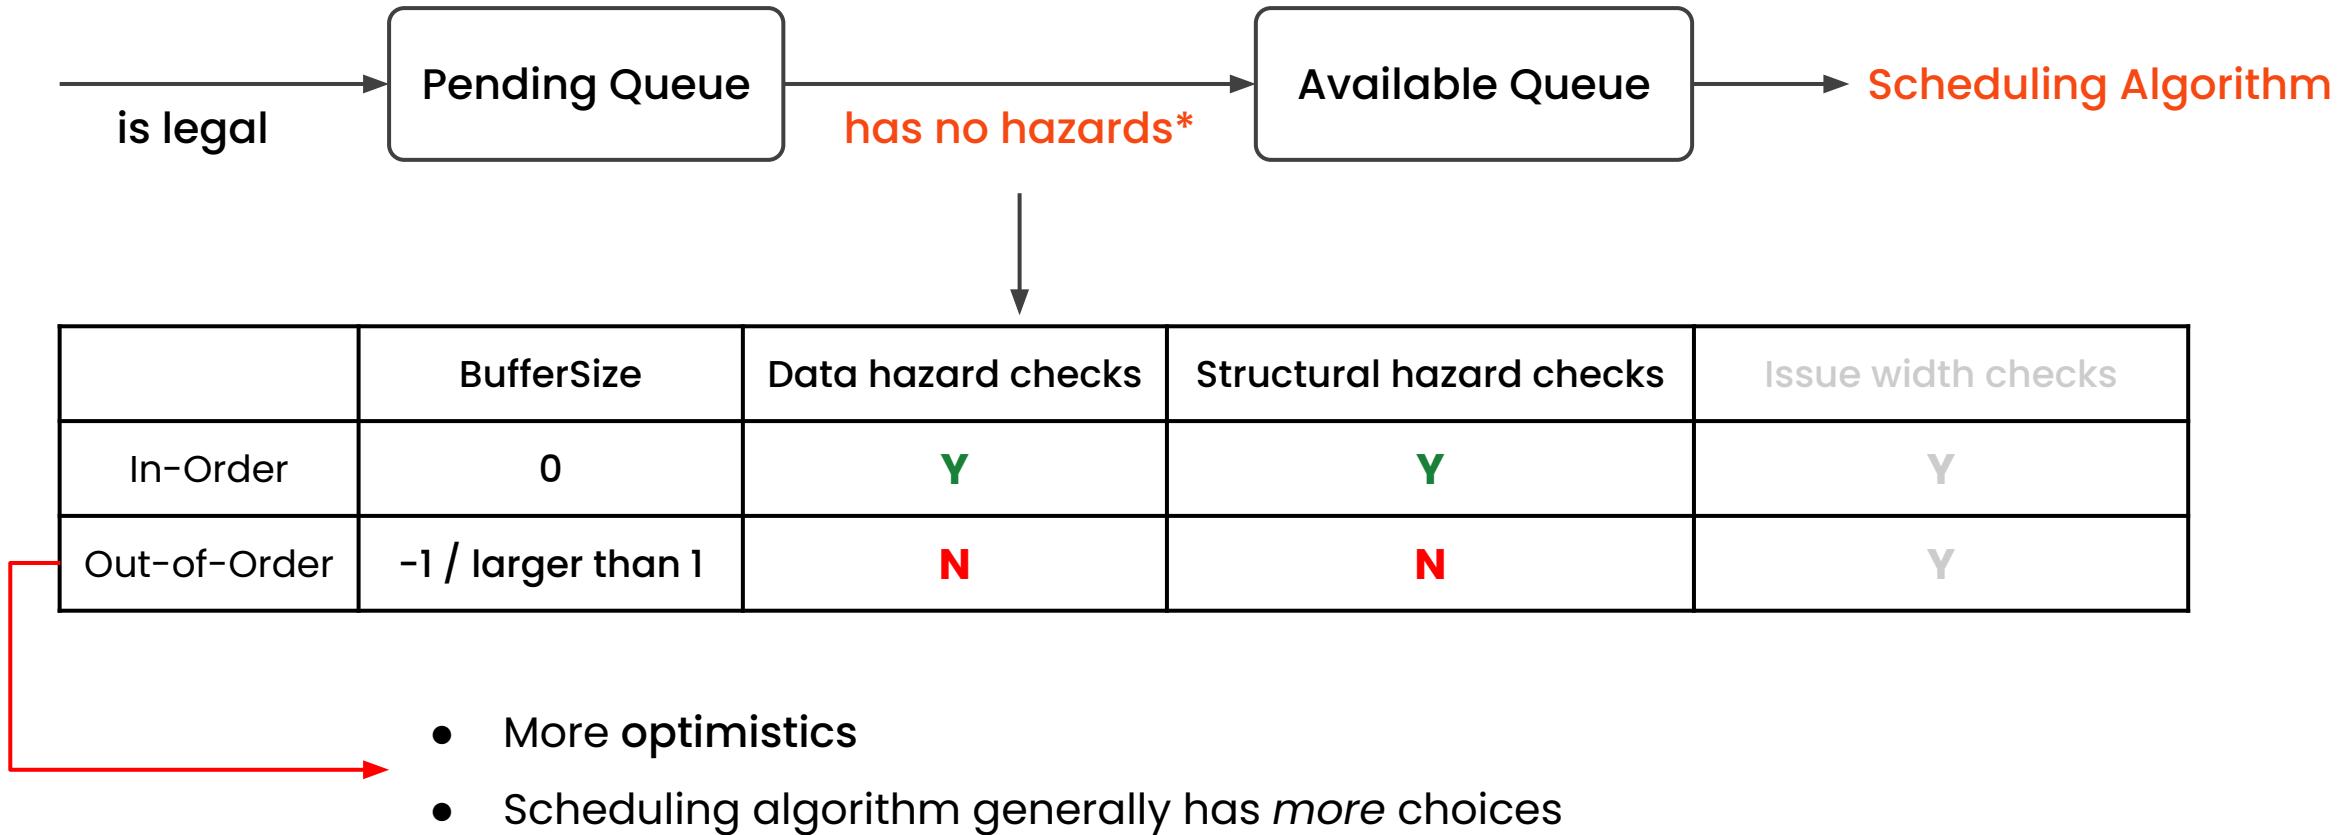
- We can predict stallings more accurately during *compile-time*
- Machine Scheduler doesn't even consider an instruction if it might induce any kind of hazards



# Machine Scheduler: Hazards

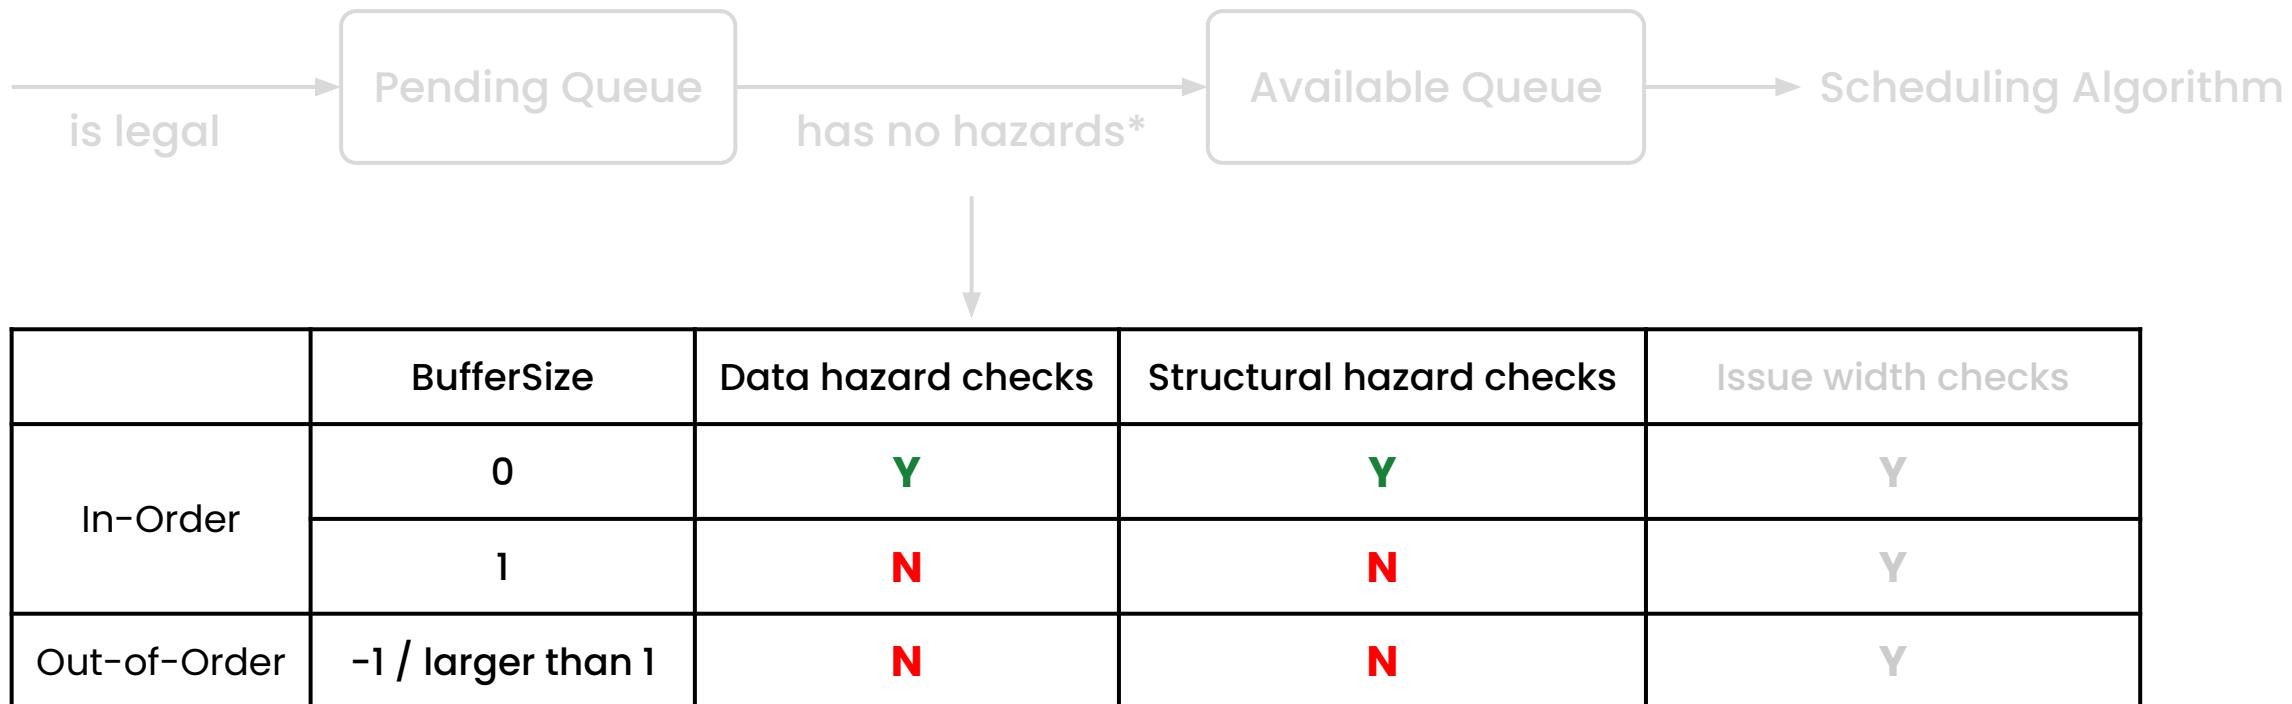


# Machine Scheduler: Hazards



# Machine Scheduler: BufferSize = 1

*In-Order core...scheduled out-of-order-ly?*



# Machine Scheduler: The *Duality* of BufferSize = 1

Too optimistic on a capability  
(i.e. out-of-order-ness) that  
doesn't even exist



FWIW...

Machine Scheduler does  
consider *data hazards* when it's  
scheduling for BufferSize = 1. But  
not structural hazards

# Machine Scheduler: The *Duality* of BufferSize = 1

Too optimistic on a capability  
(i.e. out-of-order-ness) that  
doesn't even exist



An escape hatch when **in-order** scheduling (i.e. BufferSize = 0) is  
too *strict* or *pessimistic*

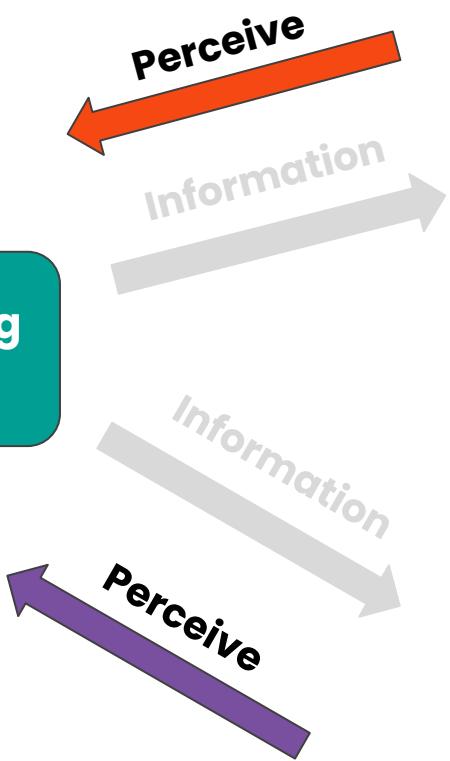
FWIW...

Machine Scheduler does  
consider *data hazards* when it's  
scheduling for BufferSize = 1. But  
not structural hazards

**Example:**

Sometimes we want to schedule  
aggressively for **register pressure**  
even though there is a hazard

## Scheduling Model



## Instruction Scheduling

ld	x9, 4 (x10)
add	x8, x9, x9
mul	x5, x5, x6
add	x7, x8, x5

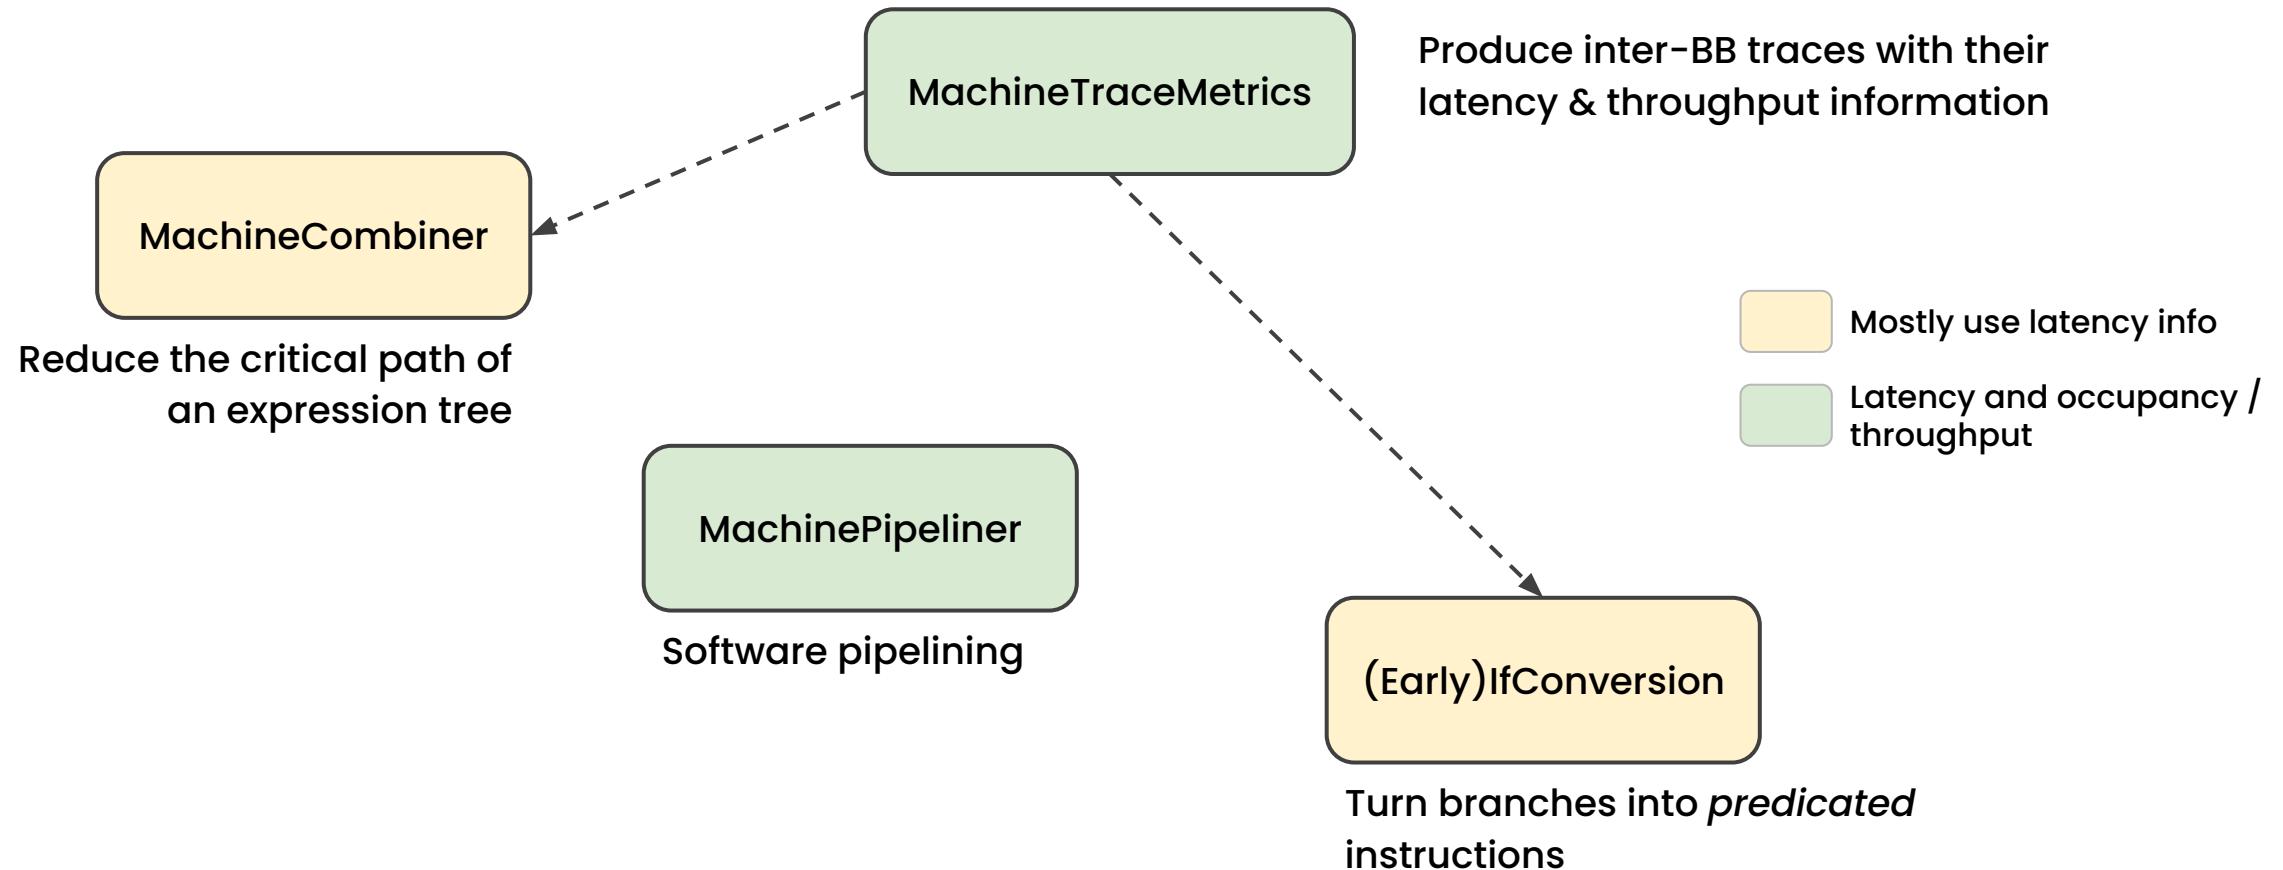
## llvm-mca

DeeER..	mul s0, s0, s1
D=eeER.	mul a0, a0, a1
D==eeER	mul a2, a2, a3
.DeE--R	add s9, s9, s10
.DeE--R	add s11, s11, t3

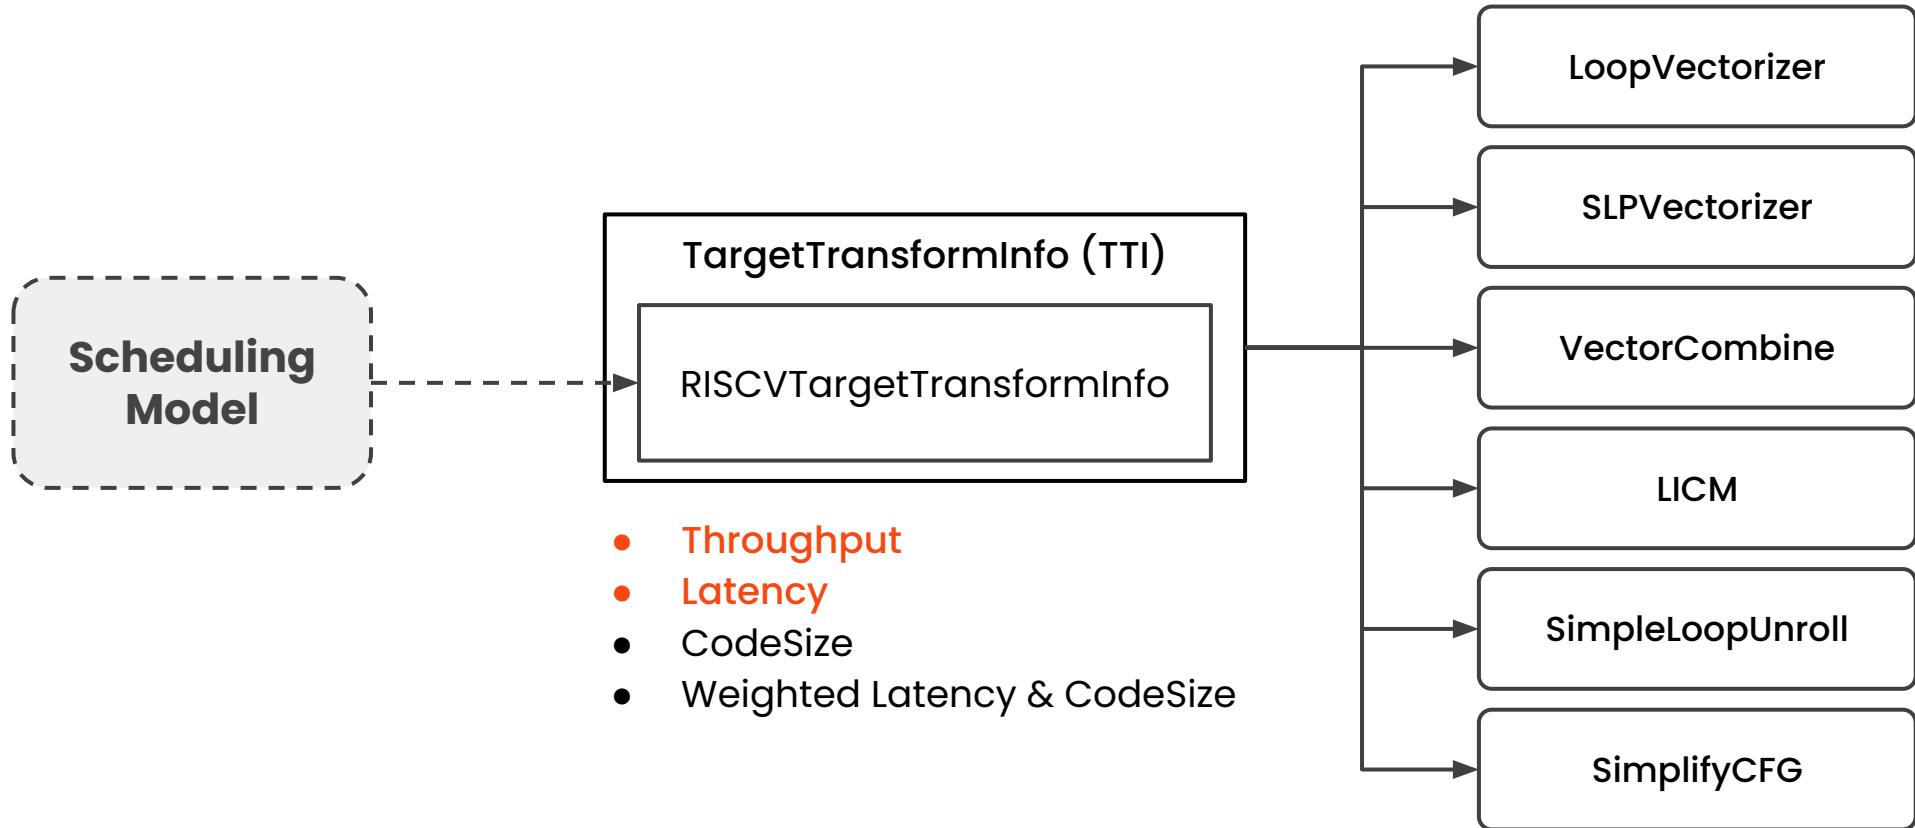


Can we use Scheduling Model in  
More Places?

# Status Quo



# Potential User: TTI Cost Model



# Potential User: TTI Cost Model

MachineInstr  
MCInst

Scheduling  
Model

## Pros

- Scheduling model can be the *centralized* place to provide such kinds of information

## Cons

- Mapping LLVM IR instructions to low-level ones can be (really) challenging
- Necessity: how many *non-trivial* instructions will actually benefit from this?

LLVM IR

TargetTransformInfo (TTI)

RISCVTargetTransformInfo

- Throughput
- Latency
- CodeSize
- Weighted Latency & CodeSize

# Other Potential Uses

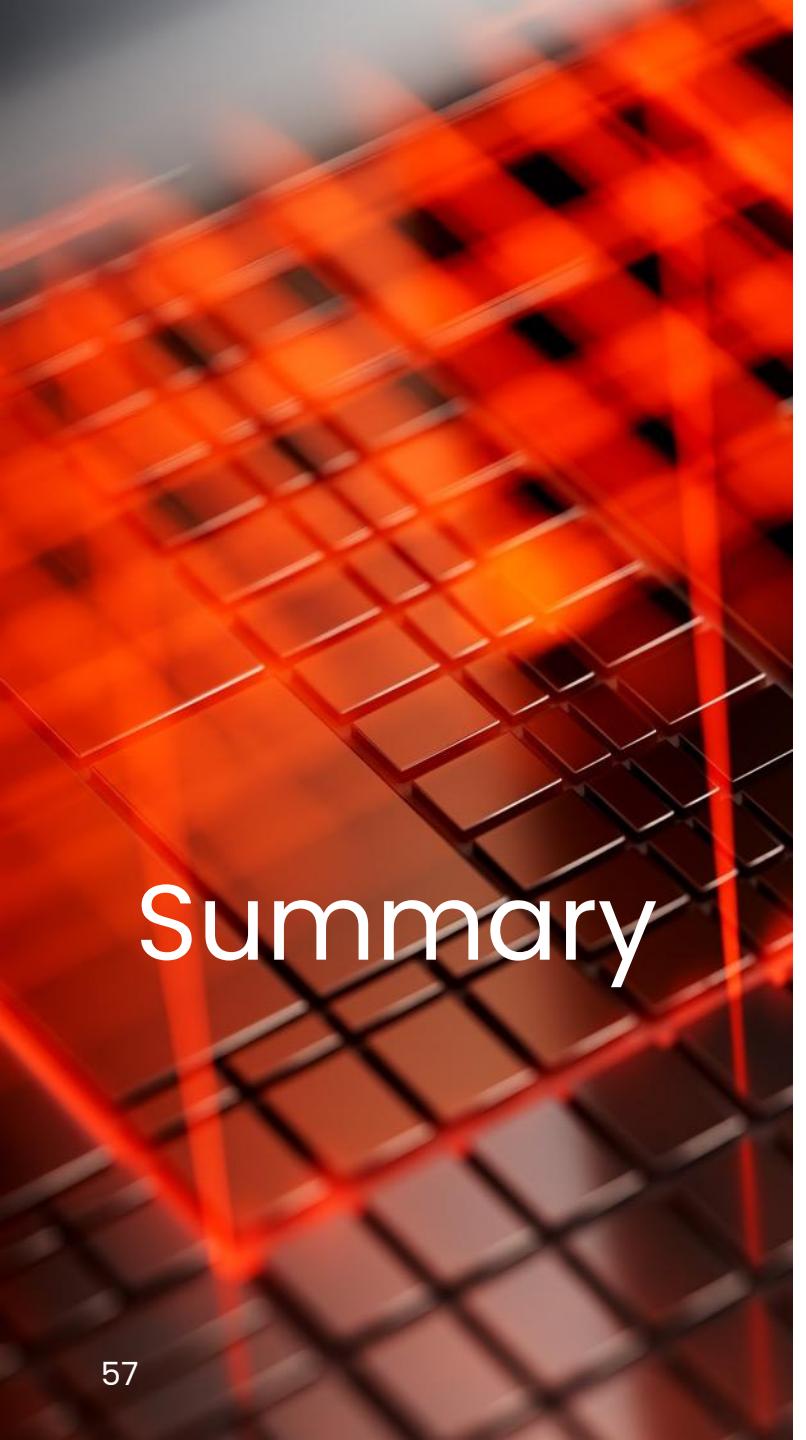
- Software pipelining “lite”: scheduling-model-guided *loop unrolling*
- *Validate* instruction scheduling with LLVM MCA
- Sort **ISel patterns** by latency

commit 3622f15491d4ae3b27f3399031933bd888e20cf0  
Author: Chris Lattner <sabre@nondot.org>  
Date: Wed Sep 28 17:57:56 2005 +0000

Prefer cheaper patterns to more expensive ones.  
Print the costs to the generated  
file

llvm-svn: 23492

```
/// Compute the number of instructions for this pattern.
/// This is a temporary hack. We should really include the instruction
/// latencies in this calculation.
static unsigned getResultPatternCost(const TreePatternNode &P,
                                     const CodeGenDAGPatterns &CGP)
```



# Summary

What does LLVM's Scheduling Model provide

Scheduling Model is all about how it's used and perceived

Potential areas we can apply Scheduling Model to



# THANK YOU

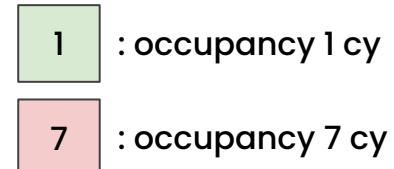
©2025 SiFive, Inc. All rights reserved.  
SiFive, HiFive and the SiFive logo are trademarks or registered trademarks of SiFive, Inc. Certain products or brand names that are not SiFive's could be trademarks or registered trademarks of their respective owners.



# Appendix

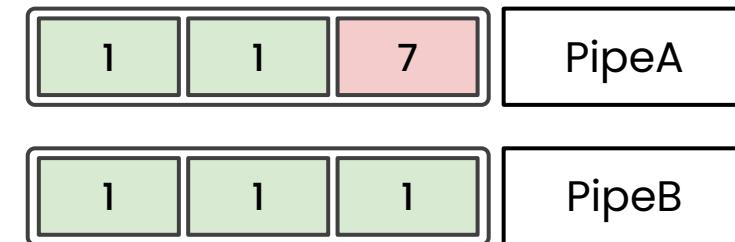
# Scheduling Model: Buffer Sizes

Buffer per pipe v.s. Pipes sharing the same buffer



```
let BufferSize = 3 in {
    def PipeA : ProcResource<1>;
    def PipeB : ProcResource<1>;
}
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
def : WriteRes<WriteIALU, [PipeAB]>;
```

Cycle = 0



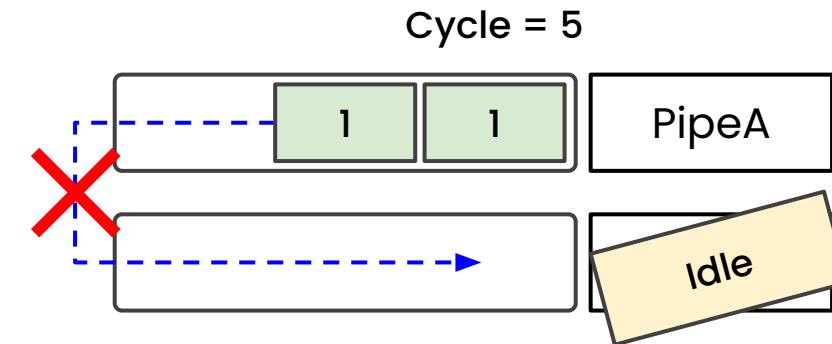
```
let BufferSize = 6 in {
    def PipeAB : ProcResource<2>;
}
def : WriteRes<WriteIALU, [PipeAB]>;
```

# Scheduling Model: Buffer Sizes

Buffer per pipe v.s. Pipes sharing the same buffer

1	: occupancy 1 cy
7	: occupancy 7 cy

```
let BufferSize = 3 in {
    def PipeA : ProcResource<1>;
    def PipeB : ProcResource<1>;
}
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
def : WriteRes<WriteIALU, [PipeAB]>;
```

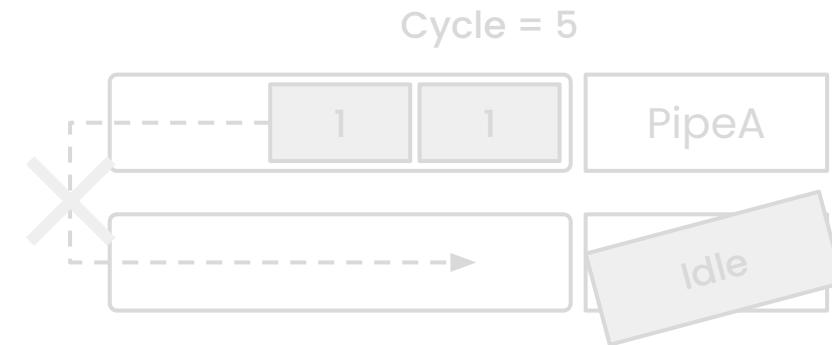
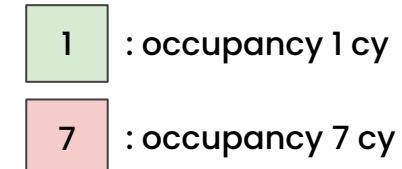


```
let BufferSize = 6 in {
    def PipeAB : ProcResource<2>;
}
def : WriteRes<WriteIALU, [PipeAB]>;
```

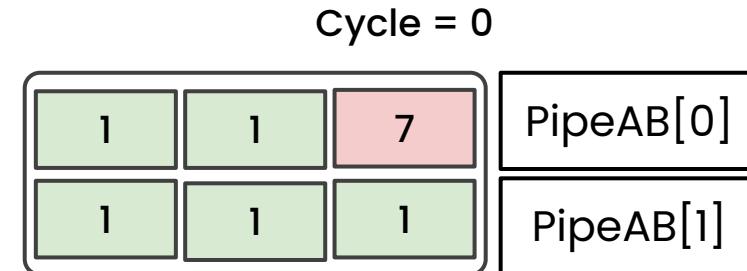
# Scheduling Model: Buffer Sizes

Buffer per pipe v.s. Pipes sharing the same buffer

```
let BufferSize = 3 in {
    def PipeA : ProcResource<1>;
    def PipeB : ProcResource<1>;
}
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
def : WriteRes<WriteIALU, [PipeAB]>;
```



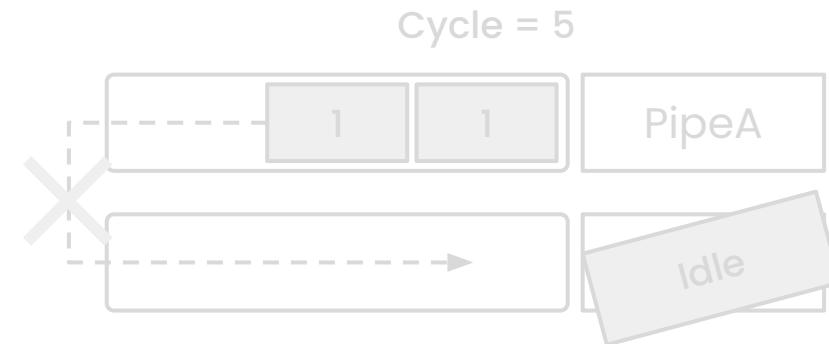
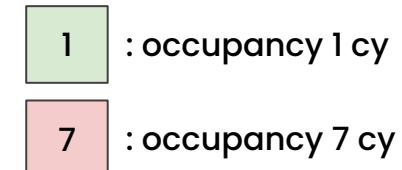
```
let BufferSize = 6 in {
    def PipeAB : ProcResource<2>;
}
def : WriteRes<WriteIALU, [PipeAB]>;
```



# Scheduling Model: Buffer Sizes

Buffer per pipe v.s. Pipes sharing the same buffer

```
let BufferSize = 3 in {
    def PipeA : ProcResource<1>;
    def PipeB : ProcResource<1>;
}
def PipeAB : ProcResGroup<[PipeA, PipeB]>;
def : WriteRes<WriteIALU, [PipeAB]>;
```



```
let BufferSize = 6 in {
    def PipeAB : ProcResource<2>;
}
def : WriteRes<WriteIALU, [PipeAB]>;
```

