



Equipping LLVM/OpenMP with Advanced OpenMP Offloading GPU Features

10th LLVM Performance Workshop
January 31, 2026

Kevin Sala | Postdoctoral Researcher
salapenades1@lbl.gov

Chaitanya Sankisa, Krzysztof Parzyszek, Michael Klemm | AMD Collaborators

Prepared by LLNL under Contract
DE-AC52-07NA27344.


Accelerating with OpenMP

- The *target* model allows **accelerating** regions on devices
 - Most constructs are available within target regions
 - Broad **vendor** support
 - Broad **language** support: C, C++, Fortran, Python (prototype)

OpenMP

```
#pragma omp target teams distribute parallel for \
    num_teams(N/256) thread_limit(256) \
    map(tofrom: Matrix[0:N*N*N]) collapse(3)
for (int z = 0; z < N; z++)
    for (int y = 0; y < N; y++)
        for (int x = 0; x < N; x++)
            Matrix[x + y*N + z*N*N] *= Val;
```

OpenMP
app




Problems

- The *target* model has **limitations** that hinder broad adoption
 - **Missing** essential GPU-specific features
 - Noticeable **overhead** in target regions (depending on features)
 - Struggle with **hierarchical parallelism**

OpenMP

```
#pragma omp target teams distribute parallel for \
    num_teams(N/256) thread_limit(256) \
    map(tofrom: Matrix[0:N*N*N]) collapse(3)
for (int z = 0; z < N; z++)
    for (int y = 0; y < N; y++)
        for (int x = 0; x < N; x++)
            Matrix[x + y*N + z*N*N] *= Val;
```

OpenMP
app





OpenMP Committee is working on that...


- Recent and upcoming features
 - **Static *team-local* memory** → Available OpenMP 6.0
 - **Dynamic *team-local* memory** → Approved OpenMP 6.1
 - **Multi-dimensional grid programming** → Under discussion
 - Reducing target region overheads → Under discussion
 - **Stream-like data dependencies** → Under discussion
 - Etc.

Overview of some recent features

- This talk will cover...
 - **Static** *team-local* memory
 - **Dynamic** *team-local* memory
 - **Multi-dimensional grid** programming
- **Overview** of such features
- **Implementation** in LLVM/OpenMP
- Early **evaluation** in two benchmarks

Groupprivate Directive

- Exposes **static shared / local** memory
 - Each *contention group* its **own copy**
 - File-scope, namespace-scope or static block-scope variables




Available
OpenMP 6.0

```
void func(int *sum, int tid) {  
    static int tmp1[1000];  
    #pragma omp groupprivate(tmp1)  
  
    #pragma omp for  
    for (int i = 0; i < 1000; i++)  
        tmp1[i] = tid + i;  
  
    #pragma omp for reduction(+: sum)  
    for (int i = 0; i < 1000; i++)  
        sum += tmp1[i];  
}  
  
int main() {  
    int sums[10];  
  
    #pragma omp target teams num_teams(10) thread_limit(256)  
    #pragma omp parallel  
    func(&sums[omp_get_team_num()], omp_get_thread_num());  
}
```

Groupprivate Directive

- Similar to CUDA and HIP

```
--global__ void func(int *sum, size_t N) {  
    __shared__ int tmp1[1000];  
    ...  
}  
  
func<<<10,256>>>(sum, N);
```




```
void func(int *sum, int tid) {  
    static int tmp1[1000];  
    #pragma omp groupprivate(tmp1)  
  
#pragma omp for  
for (int i = 0; i < 1000; i++)  
    tmp1[i] = tid + i;  
  
#pragma omp for reduction(+: sum)  
for (int i = 0; i < 1000; i++)  
    sum += tmp1[i];  
  
}  
  
int main() {  
    int sums[10];  
  
    #pragma omp target teams num_teams(10) thread_limit(256)  
    #pragma omp parallel  
    func(&sums[omp_get_team_num()], omp_get_thread_num());  
}
```

Available
OpenMP 6.0

Dynamic Groupprivate Clause

- Exposes **dynamic shared / local memory**
 - Requests a **groupprivate buffer** with runtime-determined size
 - Each *contention group* its **own buffer copy**



Approved
OpenMP 6.1

```
void func(int *sum, int tid, int N) {
    int *tmp = omp_get_dyn_groupprivate_ptr();

#pragma omp for
for (int i = 0; i < N; i++)
    tmp[i] = tid + i;

#pragma omp for reduction(+: sum)
for (int i = 0; i < N; i++)
    sum += tmp[i];
}

int main(int argc, char **argv) {
    int N = atoi(argv[1]);
    int sums[10];


#pragma omp target teams num_teams(10) thread_limit(256) \
    dyn_groupprivate(N*sizeof(int))
#pragma omp parallel
func(&sums[omp_get_team_num()], omp_get_thread_num(), N);
}
```

Dynamic Groupprivate Clause

- Somewhat similar to CUDA and HIP

```
__global__ void func(int *sum, size_t N) {
    extern __shared__ int tmp2[];
    ...
}

func<<<10, 256, N*sizeof(double)>>>(sum, N);
```



```
void func(int *sum, int tid, int N) {
    int *tmp2 = omp_get_dyn_groupprivate_ptr();

#pragma omp for
for (int i = 0; i < N; i++)
    tmp2[i] = tid + i;

#pragma omp for reduction(+: sum)
for (int i = 0; i < N; i++)
    sum += tmp2[i];
```


```
int main(int argc, char **argv) {
    int N = atoi(argv[1]);
    int sums[10];

#pragma omp target teams num_teams(10) thread_limit(256) \
    dyn_groupprivate(N*sizeof(int))
#pragma omp parallel
func(&sums[omp_get_team_num()], omp_get_thread_num(), N);
}
```

Approved
OpenMP 6.1

Dynamic Groupprivate Clause

- Groupprivate memspace is **limited**
- The **fallback** modifier determines what to do when the limit is reached:
 - **Abort** the execution
 - Return **null** pointer
 - Use memory from a **default memspace** (e.g., global)



Approved
OpenMP 6.1

```
#pragma omp target dyn_groupprivate(fallback(abort): N)
{
}
```

```
#pragma omp target dyn_groupprivate(fallback(null): N)
{
}
```

```
#pragma omp target dyn_groupprivate(fallback(default_mem): N)
{
}
```

```
#pragma omp target dyn_groupprivate(N)
{
    // default_mem is the default when not specified
}
```

Multi-dimensional grid programming

- Most accelerated applications use **grid programming**
- Single Program Multiple Data (**SPMD**)

CUDA / HIP

```
__global__ kernel(double *matrix, size_t N) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    int z = blockIdx.z * blockDim.z + threadIdx.z;

    if (x < N && y < N && z < N)
        matrix[x + y * N + z * N * N] = ...;
}

dim3 nblocks(N/256,N/256,N/256);
dim3 nthreads(256,256,256);
kernel<<<nblocks,nthreads>>>(matrix, N);
```

OpenMP 6.0

```
#pragma omp target teams distribute parallel for \
    num_teams(N/256) thread_limit(256) collapse(3)
for (int z = 0; z < N; ++z)
    for (int y = 0; y < N; ++y)
        for (int x = 0; x < N; ++x)
            matrix[x + y * N + z * N * N] = ...;
```

Multi-dimensional grid programming

- Most accelerated applications use **grid programming**
- Single Program Multiple Data (**SPMD**)

CUDA / HIP

```
__global__ kernel(double *matrix, size_t N) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    int z = blockIdx.z * blockDim.z + threadIdx.z;

    if (x < N && y < N && z < N)
        matrix[x + y * N + z * N * N] = ...;
}

dim3 nblocks(N/256,N/256,N/256);
dim3 nthreads(256,256,256);
kernel<<<nblocks,nthreads>>>(matrix, N);
```

Future OpenMP

```
#pragma omp target teams num_teams(dims(3): N/256,N/256,N/256) \
    thread_limit(dims(3): 256,256,256)
#pragma omp parallel
{
    int x = omp_get_team_num_dim(1) * omp_get_num_teams_dim(1)
        + omp_get_thread_num_dim(1);
    int y = omp_get_team_num_dim(2) * omp_get_num_teams_dim(2)
        + omp_get_thread_num_dim(2);
    int z = omp_get_team_num_dim(3) * omp_get_num_teams_dim(3)
        + omp_get_thread_num_dim(3);

    matrix[x + y * N + z * N * N] = ...;
}
```

Under
Discussion

Multi-dimensional grid programming

- Most accelerated applications use **grid programming**
- Single Program Multiple Data (**SPMD**)

Future OpenMP

```
#pragma omp target teams num_teams(dims(3): N/256,N/256,N/256) \
    thread_limit(dims(3): 256,256,256)
#pragma omp parallel
{
    int x = omp_get_team_num_dim(1) * omp_get_num_teams_dim(1)
        + omp_get_thread_num_dim(1);
    int y = omp_get_team_num_dim(2) * omp_get_num_teams_dim(2)
        + omp_get_thread_num_dim(2);
    int z = omp_get_team_num_dim(3) * omp_get_num_teams_dim(3)
        + omp_get_thread_num_dim(3);

    matrix[x + y * N + z * N * N] = ...;
}
```

Compatibility with **unidimensional** world

```
#pragma omp target teams \
    num_teams(dims(3): N/256,N/256,N/256) \
    thread_limit(dims(3): 256,256,256)
#pragma omp parallel
{
    int team = omp_get_team_num();
    int thread = omp_get_thread_num();
    int nthreads = omp_get_num_threads();

    #pragma omp single
    { ... }
}
```

Transparent flattening of identifiers/sizes

Implementation in LLVM/Clang

- **Multi-dimensional grid**
 - Support **dims** modifier in *num_teams* and *thread_limit* clauses
 - Support in LLVM/OpenMP DeviceRTL
 - **Linearization** implemented in DeviceRTL
 - Some cases not supported yet
- Full implementation of **Dynamic Groupprivate**
- **Groupprivate** not exposed as **#pragma** yet, using instead:



```
static __attribute__((address_space(3), aligned(16), loader_uninitialized)) double tmp[1024];
```

- [PR #69018](#) not evaluated yet (author @Ritanya-B-Bharadwaj)

Evaluation

- Two benchmarks
 - **Matmul** (manual implementation)
 - **Stencil3d** (extracted from HeCBench)
- Common characteristics
 - **Multi-dimensional grid** programming
 - Heavy use of **shared/local** memory
- Two platforms in LLNL
 - NVIDIA H100
 - AMD MI300A

Evaluation




Lessons learned from the implementation

- Make sure the compiler can determine the correct address space in the user code
 - **addrspace(3)** in nvidia and amdgpu
- This is problematic:


```
void *omp_get_dyn_groupprivate_ptr(...) {  
    return (!RT.was_fallback) ? RT.shmem_ptr  
                            : RT.fallback_glb_ptr;  
}
```

addrspace(3) vs. addrspace(1) at runtime => generic addrspace(0)



Lessons learned from the implementation

- Make sure to forward shared memory pointer alignment to user code
 - e.g., **16-byte** alignment in nvidia





Fortran

- AMD folks are working on the LLVM Flang implementation
 - Krishna Chaitanya Sankisa
 - Krzysztof Parzyszek
 - Michael Klemm
- More results coming soon!

Conclusions


- OpenMP *target* provides
 - **Standard** for acceleration with **multi-vendor** support
 - **Language portability:** C, C++, Fortran, Python (prototype)
 - **Abstraction** over complex hardware and software
- Recent and upcoming OpenMP features
 - **Get** the same **performance** with native offloading APIs
 - **Performance portability**
 - **Reduce** porting **effort**
- Positive results for **LLVM/OpenMP** in C/C++
- Future work: full evaluation in larger applications

Thank you!

Kevin Sala

salapenades1@llnl.gov

Watch the OpenMP tech talk on YouTube!



<https://www.youtube.com/watch?v=xOsoajfdIww>

Thank you for participating in LLVM Performance!

- Invite speakers to send copy of their slides
- Feedback to improve the workshop?