




rocMLIR: High-Performance ML Compilation for AMD GPUs with MLIR

Pablo Antonio Martinez
pablo.martinezsanchez@amd.com

EuroLLVM 2026
April 14, 2026

Agenda

1. Introduction and motivation
2. Design
3. Fusions
4. Optimizations
 - XCDs remapping
 - DirectToLDS
 - SplitK
5. Evaluation
6. Conclusions and future work


1. Introduction

The compilation flow

- MIGraphX: AMD's graph compiler (it may emit MLIR)
- TOSA: Intermediate IR
- **rock**: The rocMLIR dialect

See more about MIGraphX here:

<https://github.com/ROCm/AMDMIGraphX>




1. Introduction

Hardware targets

Gaming GPUs (RDNA)



Server GPUs (CDNA)




We support all GPUs that ROCm support


1. Introduction

What kernels do we support


Matrix multiplication



Convolutions



Attention



Also:


- gemm+gemm
- conv+gemm
- Backward convolution

Also different variants (not exhaustive list):

- KV-Cache
- Grouped convolutions
- Scaled GEMMs
- ...

1. Introduction

Quick background on GPUs



2. Design

IR Design

```
$ rocmlir-gen --arch gfx950 --num_cu 256 --operation gemm -t f32 -g 1 -m 1024 -n 1024 -k 1024
```

2. Design

IR Design

```
$ rocmlir-gen --arch gfx950 --num_cu 256 --operation gemm -t f32 -g 1 -m 1024 -n 1024 -k 1024
```

```
module attributes {
  func.func @rock_gemm(
    %arg0: memref<1048576xf32>,
    %arg1: memref<1048576xf32>,
    %arg2: memref<1048576xf32>)
  attributes
  {
    enable_splitk_for_tuning,
    kernel,
    mhal.arch = "amdgcN-amd-amdhsa:gfx950",
    num_chiplets = 8 : i64,
    num_cu = 256 : i64
  }
  %0 = rock.transform %arg0 by #transform1 : memref<1048576xf32> to memref<1x1024x1024xf32>
  %1 = rock.transform %arg1 by #transform2 : memref<1048576xf32> to memref<1x1024x1024xf32>
  %2 = rock.transform %arg2 by #transform2 : memref<1048576xf32> to memref<1x1024x1024xf32>
  rock.gemm %2 = %0 * %1 features = #gemm_features storeMethod = set :
    memref<1x1024x1024xf32> = memref<1x1024x1024xf32> * memref<1x1024x1024xf32>
  return
}
```


2. Design

IR Design

```
$ rocmlir-gen --arch gfx950 --num_cu 256 --operation gemm -t f32 -g 1 -m 1024 -n 1024 -k 1024
```

```
module attributes {
  func.func @rock_gemm(
    %arg0: memref<1048576xf32>,
    %arg1: memref<1048576xf32>,
    %arg2: memref<1048576xf32>)
  attributes
  {
    enable_splitk_for_tuning,
    kernel,
    mhal.arch = "amdgcN-amd-amdhsa:gfx950",
    num_chiplets = 8 : i64,
    num_cu = 256 : i64
  }
  %0 = rock.transform %arg0 by #transform1 : memref<1048576xf32> to memref<1x1024x1024xf32>
  %1 = rock.transform %arg1 by #transform2 : memref<1048576xf32> to memref<1x1024x1024xf32>
  %2 = rock.transform %arg2 by #transform2 : memref<1048576xf32> to memref<1x1024x1024xf32>
  rock.gemm %2 = %0 * %1 features = #gemm_features storeMethod = set :
    memref<1x1024x1024xf32> = memref<1x1024x1024xf32> * memref<1x1024x1024xf32>
  return
}
```

2. Design

IR Design

```
#map = affine_map<(d0, d1, d2) -> (d1 * 1024 + d2)>
#transform1 = #rock.transform_map<#map by
[
  <Unmerge{1024, 1024} ["m", "k"] at [1, 2] -> ["raw"] at [0]>,
  <AddDim{1} ["g"] at [0] -> [] at []>
] bounds = [1, 1024, 1024] -> [1048576]>
#transform2 = #rock.transform_map<#map by
[
  <Unmerge{1024, 1024} ["k", "n"] at [1, 2] -> ["raw"] at [0]>,
  <AddDim{1} ["g"] at [0] -> [] at []>
] bounds = [1, 1024, 1024] -> [1048576]>
#transform3 = #rock.transform_map<#map by
[
  <Unmerge{1024, 1024} ["m", "n"] at [1, 2] -> ["raw"] at [0]>,
  <AddDim{1} ["g"] at [0] -> [] at []>
] bounds = [1, 1024, 1024] -> [1048576]>
```

```
module attributes {
  func.func @rock_gemm(...) {
    %0 = rock.transform %arg0 by #transform1 : memref<1048576xf32> to memref<1x1024x1024xf32>
    %1 = rock.transform %arg1 by #transform2 : memref<1048576xf32> to memref<1x1024x1024xf32>
    %2 = rock.transform %arg2 by #transform2 : memref<1048576xf32> to memref<1x1024x1024xf32>
    rock.gemm %2 = %0 * %1 features = #gemm_features storeMethod = set ...
    return
  }
}
```

See Krzysztof presentation for more details:

<https://youtu.be/xrMJis3Rak0>

Title: *Coordinate Transformations in AMD's rocMLIR*

2. Design

Lowering pipeline

Our GPU pipeline looks like this:

```
funcPm.addPass(rock::createRockAffixTuningParametersPass(
    rock::RockAffixTuningParametersPassOptions{options.tuningFallback}));
funcPm.addPass(rock::createRockConvToGemmPass());
funcPm.addPass(rock::createRockGemmLinalgSplitkNormalizationPass());
funcPm.addPass(rock::createRockGemmToGridwisePass());
funcPm.addPass(rock::createRockRegularizePass());
funcPm.addPass(rock::createRockShuffleGemmForReductions());
funcPm.addPass(rock::createRockGridwiseGemmToBlockwisePass());
funcPm.addPass(rock::createRockBlockwiseLoadTileToThreadwisePass());

funcPm.addPass(rock::createRockLinalgAlignPass());
funcPm.addPass(rock::createRockBlockwiseGemmToThreadwisePass());
funcPm.addPass(rock::createRockPipelinePass());
funcPm.addPass(createCanonicalizerPass());
funcPm.addPass(createConvertLinalgToAffineLoopsPass());
funcPm.addPass(rock::createRockVectorizeFusionsPass());
funcPm.addPass(rock::createRockAddAsyncWaitPass());

funcPm.addPass(rock::createRockAnnotateLivenessPass());
funcPm.addPass(rock::createRockReuseLDSPass());
funcPm.addPass(rock::createRockOutputSwizzlePass());
funcPm.addPass(rock::createRockAnnotateLivenessPass());
funcPm.addPass(rock::createRockReuseLDSPass());
```

2. Design

Lowering pipeline

IR after `RockAffixTuningParametersPass`

```
func.func @rock_gemm(...) {  
  ...  
  rock.gemm %2 = %0 * %1 features = #gemm_features storeMethod = set  
  {  
    derivedBlockSize = 256 : i32,  
    params = #rock.accel_gemm_params<  
      kpackPerBlock = 8,  
      mPerBlock = 32,  
      nPerBlock = 64,  
      kpack = 8,  
      mPerWave = 32,  
      nPerWave = 16,  
      mnPerXdl = 16,  
      splitKFactor = 1,  
      scheduleVersion = 1,  
      outputSwizzle = 2,  
      wavesPerEU = 0,  
      gridGroupSize = 0,  
      forceUnroll = true  
    >  
  } : memref<1x1024x1024xf32> = memref<1x1024x1024xf32> * memref<1x1024x1024xf32>  
  return  
}
```

2. Design

Lowering pipeline

IR after `RockGridwiseGemmToBlockwisePass` (part 1)

```
%4 = rock.workgroup_id : index
%5 = rock.workitem_id : index

...

%c4 = arith.constant 4 : index
%c128 = arith.constant 128 : index
%6 = arith.remui %4, %c4 : index
%7 = arith.divui %4, %c4 : index
%8 = arith.muli %6, %c128 : index
%9 = arith.addi %7, %8 : index
%c511 = arith.constant 511 : index
%10 = arith.cmpi sgt, %4, %c511 : index
%11 = arith.select %10, %4, %9 : index

...

%22 = rock.alloc() : memref<8192xi8, #gpu.address_space<workgroup>>
%23 = rock.alloc() : memref<16384xi8, #gpu.address_space<workgroup>>

...

%24 = rock.alloc() : memref<16xf32, #gpu.address_space<private>>
%25 = rock.alloc() : memref<16xf32, #gpu.address_space<private>>
```

2. Design

Lowering pipeline

IR after `RockGridwiseGemmToBlockwisePass` (part 2)

```
scf.for %arg3 = %c0_3 to %c16_2 step %c1 {
  rock.blockwise_load_tile %1[%arg3, %12, %19, %21, %5] LDS -> %23 -> %25
  rock.blockwise_load_tile %3[%arg3, %12, %19, %21, %5] LDS -> %22 -> %24
  rock.lds_barrier
  rock.stage {
    rock.blockwise_gemm_accel %26 += %24 from %view * %25 from %view_1
    features = {#gemm_features} : ...
    rock.yield
  } {name = "MMA"}
  rock.lds_barrier
} {pipeline = #rock.pipeline<2>}
%27 = rock.alloc() : memref<8xf32, #regs>
rock.transforming_for {forceUnroll, useIndexDiffs}
  (%arg3) = [](%c0), (%arg4) = [#transform_map4](%c0_4)
  (%arg5, %arg6) = validity bounds [2] strides [1] {
    %28 = memref.load %26[%arg3] : memref<2xvector<4xf32>, #regs>>
    rock.in_bounds_store %28 -> %27[%arg4] : vector<4xf32> -> memref<8xf32, #lds>>,
index
    rock.yield
  }
}
rock.threadwise_write_all {forceUnroll, useIndexDiffs} %27 ... :
  memref<8xf32, #regs> -> memref<1x1024x1024xf32>
```

2. Design

Lowering pipeline

IR after `RockBlockwiseLoadTileToThreadwisePass` (GlobalRead)

```
scf.for %arg3 = %c0 3 to %c16 2 step %c1 {
  rock.stage {
    %33 = rock.transform %1 by #transform_map5 :
      memref<1x1024x1024xf32> to memref<16x1x32x16x256x8xf32>
    %34 = rock.threadwise_read_into [](%33) [%arg3, %12, %19, %21, %5] -> %29 :
      memref<16x1x32x16x256x8xf32> -> memref<8xf32, #regs>, vector<8xi1>

    %36 = rock.transform %2 by #transform_map7 : memref<1x1024x1024xf32> to
      memref<16x1x32x16x256x16xf32>
    %37 = rock.threadwise_read_into [](%36) [%arg3, %12, %19, %21, %5] -> %27 :
      memref<16x1x32x16x256x16xf32> -> memref<16xf32, #regs>, vector<16xi1>
    rock.yield
  } {name = "GlobalRead"}

  rock.stage { ... } {name = "LDSWrite"}

  rock.lhs_barrier

  rock.stage { ... } {name = "MMA"}

  rock.lhs_barrier
}
{pipeline = #rock.pipeline<2>}
```

2. Design

Lowering pipeline

IR after `RockBlockwiseLoadTileToThreadwisePass` (LDSWrite)

```
scf.for %arg3 = %c0_3 to %c16_2 step %c1 {
  rock.stage { ... } {name = "GlobalRead"}
```

```
  rock.stage {
    %32 = rock.workitem_id : index
    %34 = rock.transform %29 by #transform_map9 : memref<8xf32, #regs> to
      memref<4x2xf32, #regs>
    %36 = rock.transform %30 by #transform_map11 : memref<8xf32, #regs> to
      memref<4x2xf32, #regs>

    %42 = rock.transform
    rock.threadwise_copy %34 -> %36 : memref<4x2xf32, #regs> -> memref<4x2xf32, #regs>
    rock.threadwise_write_all {forceUnroll, useIndexDiffs} %30 -> [ ](%42) [%32]
      by set : memref<8xf32, #regs> -> memref<256x8xvector<4xf32>, #lds>

    // Same chain of transforms for B tile
    rock.yield
  } {name = "LDSWrite"}
```

```
  rock.lds_barrier
```

```
  rock.stage { ... } {name = "MMA"}
```

```
  rock.lds_barrier
```

```
} {pipeline = #rock.pipeline<2>}
```


2. Design

Lowering pipeline

IR after `RockBlockwiseLoadTileToThreadwisePass` (MMA)

```
scf.for %arg3 = %c0_3 to %c16_2 step %c1 {
  rock.stage { ... } {name = "GlobalRead"}
  rock.stage { ... } {name = "LDWrite"}
```

```
  rock.lds_barrier
```

```
  rock.stage {
    rock.blockwise_gemm_accel %26 += %24 from %view * %25 from %view_1
    features = #gemm features
: memref<2xvector<4xf32>, #regs> +=
    memref<16xf32, #regs> from
    memref<256xvector<8xf32>, #lds> *
    memref<16xf32, #regs> from memref<512xvector<8xf32>, #lds>
    rock.yield
  } {name = "MMA"}
```

```
  rock.lds_barrier
```

```
} {pipeline = #rock.pipeline<2>}
```

2. Design

Lowering pipeline

IR after we reach amdgpu (upstream dialect)

```
amdgpu.gather_to_lds %memspacecast[%227], %view_2[%228] : f128,
    memref<1048576xf32, #gpu.address_space<global>>,
    memref<8192xf32, #gpu.address_space<workgroup>>
```

NOTE: gather_to_lds is generated only if DirectToLDS optimization is enabled

```
affine.for %arg5 = 0 to 16 {
    %174 = memref.load %18[%arg5] : memref<16xf32, #regs>
    %175 = memref.load %19[%arg5] : memref<16xf32, #regs>
    %176 = memref.load %20[%arg4] : memref<2xvector<4xf32>, #regs>
    %177 = amdgpu.mfma 16x16x4 %174 * %175 + %176 blgp = none : f32, f32, vector<4xf32>
    memref.store %177, %20[%arg4] : memref<2xvector<4xf32>, #regs>
}
```

```
amdgpu.lds_barrier
```

2. Design

Lowering pipeline

IR after we reach rocdl (upstream dialect)

```
rocdl.mfma.f32.16x16x4f32 %494, %497, %500, %501, %502, %503 :  
    (f32, f32, vector<4xf32>, i32, i32, i32) -> vector<4xf32>
```


```
rocdl.s.waitcnt -7937  
rocdl.s.barrier
```

The rest of the IR is in llvm dialect, so we can lower all to LLVM...

... and then, to assembly!


3. Fusions

3.1 Fusions





3. Fusions

3.1 Fusions




3. Fusions

3.1 Fusions



3. Fusions

Input vs Output fusions



3. Fusions

Input fusion

```
func.func @input_fusion(
  %arg0: memref<1x1024x1024xf32>,
  %arg1: memref<1x1024x1024xf32>,
  %arg2: memref<1x1024x1024xf32>)
{
  %cst = arith.constant 4.000000e+00 : f32
  %alloc = memref.alloc() : memref<1x1024x1024xf32>

  linalg.generic {
    indexing_maps = [#map, #map],
    iterator_types = ["parallel", "parallel", "parallel"]}
    ins(%arg0 : memref<1x1024x1024xf32>) outs(%alloc : memref<1x1024x1024xf32>) {
^bb0(%in: f32, %out: f32):
      %0 = arith.addf %in, %cst : f32
      linalg.yield %0 : f32
    }

  rock.gemm %arg2 = %alloc * %arg1 {arch = "gfx950", params = #general_gemm_params} :
    memref<1x1024x1024xf32> =
    memref<1x1024x1024xf32> * memref<1x1024x1024xf32>

  return
}
```


3. Fusions

3.1 Input fusion

```
scf.for %arg3 = %c0_11 to %c64_10 step %c1 {
  rock.stage {
```

```
    ...
    %53 = rock.threadwise_read_into [ ](%52) [%arg3, %19, %26, %28, %6] -> %29 :
        memref<64x1x8x8x128x16xf32> -> memref<16xf32, #regs>, vector<16xi1>
    %54 = rock.alloc() : memref<16xf32, #regs>
    linalg.generic {indexing_maps = [#map, #map], iterator_types = ["parallel"]}
        ins(%29 : memref<16xf32, #regs>) outs(%54 : memref<16xf32, #regs>) {
    ^bb0(%in: f32, %out: f32):
        %57 = arith.addf %in, %cst_1 : f32
        linalg.yield %57 : f32
    }
    ...
  }
```

```
} {name = "GlobalRead"}
rock.stage {
```

```
    ...
    rock.threadwise_write_all %54 -> %43[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
    rock.threadwise_write_all %55 -> %47[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
  } {name = "LDSWrite"}
```

```
rock.stage {
  rock.blockwise_gemm %4 += %43 * %47 {params = #rock.general_gemm_params} :
    memref<8x16xf32, #regs> += memref<16x128x1xf32, #lds> * memref<16x128x1xf32, #lds>
  } {name = "MMA"}
} {pipeline = #rock.pipeline<2>}
```

Fusion happens after ToBlockwise pass,
in a pass called RockLinalgAlignPass

This is actually **loop fusion** because the
add is fused within the kernel loop

3. Fusions

3.1 Input fusion

```
scf.for %arg3 = %c0_11 to %c64_10 step %c1 {
  rock.stage {
```

```
    ...
    %53 = rock.threadwise_read_into [ ](%52) [%arg3, %19, %26, %28, %6] -> %29 :
        memref<64x1x8x8x128x16xf32> -> memref<16xf32, #regs>, vector<16xi1>
```

Load A tile from
global memory
to registers

```
    %54 = rock.alloc() : memref<16xf32, #regs>
    linalg.generic {indexing_maps = [#map, #map], iterator_types = ["parallel"]}
      ins(%29 : memref<16xf32, #regs>) outs(%54 : memref<16xf32, #regs>) {
```

```
      ^bb0(%in: f32, %out: f32):
        %57 = arith.addf %in, %cst_1 : f32
        linalg.yield %57 : f32
      }
```

```
    } {name = "GlobalRead"}
  rock.stage {
```

```
    ...
    rock.threadwise_write_all %54 -> %43[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
    rock.threadwise_write_all %55 -> %47[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
```

```
  } {name = "LDSWrite"}
```

```
  rock.stage {
    rock.blockwise_gemm %4 += %43 * %47 {params = #rock.general_gemm_params} :
      memref<8x16xf32, #regs> += memref<16x128x1xf32, #lds> * memref<16x128x1xf32, #lds>
```

```
  } {name = "MMA"}
```

```
} {pipeline = #rock.pipeline<2>}
```

3. Fusions

3.1 Input fusion

```
scf.for %arg3 = %c0_11 to %c64_10 step %c1 {
  rock.stage {
```

```
    ...
    %53 = rock.threadwise_read_into [ ](%52) [%arg3, %19, %26, %28, %6] -> %29 :
      memref<64x1x8x8x128x16xf32> -> memref<16xf32, #regs>, vector<16xi1>
```

```
    %54 = rock.alloc() : memref<16xf32, #regs>
    linalg.generic {indexing_maps = [#map, #map], iterator_types = ["parallel"]}
      ins(%29 : memref<16xf32, #regs>) outs(%54 : memref<16xf32, #regs>) {
    ^bb0(%in: f32, %out: f32):
      %57 = arith.addf %in, %cst_1 : f32
      linalg.yield %57 : f32
    }
  }
```

Add happens
in registers
(fused within
the kernel loop)

```
    ...
  } {name = "GlobalRead"}
  rock.stage {
```

```
    ...
    rock.threadwise_write_all %54 -> %43[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
    rock.threadwise_write_all %55 -> %47[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
  } {name = "LDSWrite"}
```

```
  rock.stage {
    rock.blockwise_gemm %4 += %43 * %47 {params = #rock.general_gemm_params} :
      memref<8x16xf32, #regs> += memref<16x128x1xf32, #lds> * memref<16x128x1xf32, #lds>
  } {name = "MMA"}
} {pipeline = #rock.pipeline<2>}
```

3. Fusions

3.1 Input fusion

```
scf.for %arg3 = %c0_11 to %c64_10 step %c1 {
  rock.stage {
    ...
    %53 = rock.threadwise_read_into [ ](%52) [%arg3, %19, %26, %28, %6] -> %29 :
      memref<64x1x8x8x128x16xf32> -> memref<16xf32, #regs>, vector<16xi1>
    %54 = rock.alloc() : memref<16xf32, #regs>
    linalg.generic {indexing_maps = [#map, #map], iterator_types = ["parallel"]}
      ins(%29 : memref<16xf32, #regs>) outs(%54 : memref<16xf32, #regs>) {
    ^bb0(%in: f32, %out: f32):
      %57 = arith.addf %in, %cst_1 : f32
      linalg.yield %57 : f32
    }
  }
}
```

A' and B tiles (%54 and %55) are moved
from registers into LDS memory

```
...
rock.threadwise_write_all %54 -> %43[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
rock.threadwise_write_all %55 -> %47[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
} {name = "LDSWrite"}
rock.stage {
  rock.blockwise_gemm %4 += %43 * %47 {params = #rock.general_gemm_params} :
    memref<8x16xf32, #regs> += memref<16x128x1xf32, #lds> * memref<16x128x1xf32, #lds>
} {name = "MMA"}
} {pipeline = #rock.pipeline<2>}
```

3. Fusions

3.1 Input fusion

```

scf.for %arg3 = %c0_11 to %c64_10 step %c1 {
  rock.stage {
    ...
    %53 = rock.threadwise_read_into [ ](%52) [%arg3, %19, %26, %28, %6] -> %29 :
      memref<64x1x8x8x128x16xf32> -> memref<16xf32, #regs>, vector<16xi1>
    %54 = rock.alloc() : memref<16xf32, #regs>
    linalg.generic {indexing_maps = [#map, #map], iterator_types = ["parallel"]}
      ins(%29 : memref<16xf32, #regs>) outs(%54 : memref<16xf32, #regs>) {
    ^bb0(%in: f32, %out: f32):
      %57 = arith.addf %in, %cst_1 : f32
      linalg.yield %57 : f32
    }
    ...
  } {name = "GlobalRead"}
  rock.stage {
    ...
    rock.threadwise_write_all %54 -> %43[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
    rock.threadwise_write_all %55 -> %47[%6]: memref<16xf32, #regs> -> memref<128x16xf32, #lds>
  } {name = "LDSWrite"}
  rock.stage {
    rock.blockwise_gemm %4 += %43 * %47 {params = #rock.general_gemm_params} :
      memref<8x16xf32, #regs> += memref<16x128x1xf32, #lds> * memref<16x128x1xf32, #lds>
  } {name = "MMA"}
} {pipeline = #rock.pipeline<2>}

```

A' tile (%43) has the addf computed

3. Fusions

3.2 Output fusion

```
func.func @output_fusion(
  %arg0: memref<1x2x320xf32>,
  %arg1: memref<1x2x1280xf32>,
  %arg2: memref<1x1280x320xf32>,
  %arg3: memref<1x2x320xf32>)
{
  %alloc = memref.alloc() {alignment = 64 : i64} : memref<1x2x320xf32>
  rock.gemm %alloc = %arg1 * %arg2 features = #gemm_features :
    memref<1x2x320xf32> = memref<1x2x1280xf32> * memref<1x1280x320xf32>
  %0 = rock.transform %alloc by #transform1 : memref<1x2x320xf32> to memref<2x320xf32>
  %1 = rock.transform %arg0 by #transform1 : memref<1x2x320xf32> to memref<2x320xf32>
  %alloc_0 = memref.alloc() {alignment = 64 : i64} : memref<2x320xf32>
  linalg.generic {indexing_maps = [#map1, #map1, #map1], iterator_types = ["parallel",
"parallel"]} ins(%0, %1 : memref<2x320xf32>, memref<2x320xf32>) outs(%alloc_0 :
memref<2x320xf32>) {
    ^bb0(%in: f32, %in_1: f32, %out: f32):
      %3 = arith.addf %in, %in_1 : f32
      linalg.yield %3 : f32
  }
  %2 = rock.transform %alloc_0 by #transform2 : memref<2x320xf32> to memref<1x2x320xf32>
  memref.copy %2, %arg3 : memref<1x2x320xf32> to memref<1x2x320xf32>
  return
}
```

3. Fusions

3.2 Output fusion

```
scf.for %arg4 = %c0_3 to %c20 step %c1_2 {
  rock.stage { ... } {name = "GlobalRead"}
  rock.stage { ... } {name = "LDWrite"}
  rock.lds_barrier
  rock.stage { ... } {name = "MMA"}
  rock.lds_barrier
} {pipeline = #rock.pipeline<2>}
```

This is the matmul kernel. It produces the output tile in %25, which resides in registers

...

```
rock.threadwise_read_into [...](%31) [%6, %13, %15, %5] -> %29 :
  memref<1x16x32xf32> -> memref<4xf32, #regs>
```

...

```
linalg.generic {
  indexing_maps = [#map30, #map30, #map30],
  iterator_types = ["parallel"]
  ins[%25, %29 : memref<4xf32, #regs>, memref<4xf32, #regs>)
  outs(%28 : memref<4xf32, #regs>) {
    ^bb0(%in: f32, %in_6: f32, %out: f32):
      %37 = arith.addf %in, %in_6 : f32
      linalg.yield %37 : f32
  }
}
```

...

3. Fusions

3.2 Output fusion

```
scf.for %arg4 = %c0_3 to %c20 step %c1_2 {
  rock.stage { ... } {name = "GlobalRead"}
  rock.stage { ... } {name = "LDWrite"}
  rock.lds_barrier
  rock.stage { ... } {name = "MMA"}
  rock.lds_barrier
} {pipeline = #rock.pipeline<2>}
...
```

```
rock.threadwise_read_into [...](%31) [%6, %13, %15, %5] -> %29 :
  memref<1x16x32xf32> -> memref<4xf32, #regs>
```

The bias tile is loaded from global into registers

```
...
linalg.generic {
  indexing_maps = [#map30, #map30, #map30],
  iterator_types = ["parallel"]
  ins(%25, %29 : memref<4xf32, #regs>, memref<4xf32, #regs>)
  outs(%28 : memref<4xf32, #regs>) {
    ^bb0(%in: f32, %in_6: f32, %out: f32):
      %37 = arith.addf %in, %in_6 : f32
      linalg.yield %37 : f32
  }
}
...
```


3. Fusions

3.2 Output fusion

```
scf.for %arg4 = %c0_3 to %c20 step %c1_2 {
  rock.stage { ... } {name = "GlobalRead"}
  rock.stage { ... } {name = "LDWrite"}
  rock.lds_barrier
  rock.stage { ... } {name = "MMA"}
  rock.lds_barrier
} {pipeline = #rock.pipeline<2>}
...
```

```
rock.threadwise_read_into [...](%31) [%6, %13, %15, %5] -> %29 :
  memref<1x16x32xf32> -> memref<4xf32, #regs>
```

...

```
linalg.generic {
  indexing_maps = [#map30, #map30, #map30],
  iterator_types = ["parallel"]}
ins(%25, %29 : memref<4xf32, #regs>, memref<4xf32, #regs>)
outs(%28 : memref<4xf32, #regs>) {
```

```
^bb0(%in: f32, %in_6: f32, %out: f32):
  %37 = arith.addf %in, %in_6 : f32
  linalg.yield %37 : f32
}
```


... After the bias we'll write %28 (the output tile) to global memory

The bias happens outside of the kernel loop, yet it's computed in registers, so no trip to global memory is performed here

4. Optimizations

XCDs remapping


- MI350X is divided into 8 XCDs
- Each XCD has:
 - Its own HBM slice
 - Its own L2 cache (4MB)
 - 32 CUs



4. Optimizations


XCDs remapping

- Let's imagine we launch a kernel with 16 blocks.
- Each of them will be scheduled to a different XCD in a round-robin fashion




4. Optimizations

XCDs remapping




4. Optimizations

XCDs remapping



4. Optimizations

XCDs remapping




4. Optimizations

XCDs remapping

- Consecutive workgroups usually work on consecutive data.

However, they are scheduled far away (for instance, WG0 and WG1 are in different XCDs)

- This is very inefficient due to:
 - Inter XCD communication
 - L2 trashing



4. Optimizations

XCDs remapping

- We can't change how the workgroups are mapped to the XCDs
- But can remap our workgroup id, so that each workgroup works with another part of the data, such that the data locality is maximized

Let's assume:

- GPU with 4 XCD and 32 CUs
- A matrix that is subdivided in 6x6 sub-tiles
- No remapping would look like this:

0	2	0	2	0	2
1	3	1	3	1	3
2	0	2	0	2	0
3	1	3	1	3	1
0	2	0	2	0	2
1	3	1	3	1	3

(Each number means that such sub-tile is assigned to the XCD with that ID)

4. Optimizations

XCDs remapping

- Our remapping consist on creating groups
- The group size is determined by a heuristic:

$$groupSize = \sqrt{\frac{numCU}{numXCD}} * \frac{bitWidthOut}{bitWidthIn} = \sqrt{\frac{32}{4}} * \frac{32}{32} = 2$$

0	0	3	3	2	2
0	0	3	3	2	2
1	1	0	0	3	3
1	1	0	0	3	3
2	2	1	1	0	1
2	2	1	1	2	3

4. Optimizations

XCDs remapping

- Our remapping consist on creating groups
- The group size is determined by an heuristic:

- NOTE: This mapping is used for all kernels. We could specialize remapping for specific kernels or use cases (e.g., KV cache)

$$groupSize = \sqrt{\frac{numCU}{numXCD}} * \frac{bitWidthOut}{bitWidthIn} = \sqrt{\frac{32}{4}} * \frac{32}{32} = 2$$

0	0	3	3	2	2
0	0	3	3	2	2
1	1	0	0	3	3
1	1	0	0	3	3
2	2	1	1	0	1
2	2	1	1	2	3

4. Optimizations

DirectToLDS

How does the picture look like? (without DirectToLDS optimization)

```
scf.for {
  rock.stage {
    rock.threadwise_read_into
    rock.threadwise_read_into
  } {name = "GlobalRead"}

  rock.stage {
    rock.threadwise_write_all
    rock.threadwise_write_all
  } {name = "LDSWrite"}

  rock.lds_barrier

  rock.stage {
    rock.blockwise_gemm_accel
  } {name = "MMA"}

  rock.lds_barrier
} {pipeline = #rock.pipeline<2>}

rock.threadwise_write_all
```

Reads A and B tile from global memory into registers

4. Optimizations

DirectToLDS

How does the picture look like? (without DirectToLDS optimization)

```
scf.for {
  rock.stage {
    rock.threadwise_read_into
    rock.threadwise_read_into
  } {name = "GlobalRead"}
```

```
rock.stage {
  rock.threadwise_write_all
  rock.threadwise_write_all
} {name = "LDSWrite"}
```

Copies A and B tile from tiles into LDS

```
rock.lds_barrier
```

```
rock.stage {
  rock.blockwise_gemm_accel
} {name = "MMA"}
```

```
rock.lds_barrier
} {pipeline = #rock.pipeline<2>}
```

```
rock.threadwise_write_all
```

4. Optimizations

DirectToLDS

How does the picture look like? (without DirectToLDS optimization)

```
scf.for {
  rock.stage {
    rock.threadwise_read_into
    rock.threadwise_read_into
  } {name = "GlobalRead"}
```

```
  rock.stage {
    rock.threadwise_write_all
    rock.threadwise_write_all
  } {name = "LDSWrite"}
```

```
  rock.lds_barrier
```

```
  rock.stage {
    rock.blockwise_gemm_accel
  } {name = "MMA"}
```

Computes matmul by reading from LDS and writing into registers

```
  rock.lds_barrier
} {pipeline = #rock.pipeline<2>}
```

```
rock.threadwise_write_all
```

4. Optimizations

DirectToLDS

How does the picture look like? (without DirectToLDS optimization)

```
scf.for {  
  rock.stage {  
    rock.threadwise_read_into  
    rock.threadwise_read_into  
  } {name = "GlobalRead"}  
  
  rock.stage {  
    rock.threadwise_write_all  
    rock.threadwise_write_all  
  } {name = "LDSWrite"}  
  
  rock.lds_barrier  
  
  rock.stage {  
    rock.blockwise_gemm_accel  
  } {name = "MMA"}  
  
  rock.lds_barrier  
} {pipeline = #rock.pipeline<2>}
```

`rock.threadwise_write_all`

Reads from registers and writes the output back into global memory

4. Optimizations

DirectToLDS

Modern AMD GPUs have instructions that allows you to load from global memory to LDS directly, so we can have something like:

```
scf.for {
  rock.stage {
    rock.threadwise_read_into
    rock.threadwise_read_into
    rock.yield
  } {name = "GlobalRead"}

  rock.lds_barrier

  rock.stage {
    rock.blockwise_gemm_accel
    rock.yield
  } {name = "MMA"}

  rock.lds_barrier

} {pipeline = #rock.pipeline<2>}

rock.threadwise_write_all
```

This now loads from global memory into LDS directly (DirectToLDS)

This is the same, load from LDS and store to registers

This is the same, loads from registers and stores to global memory

4. Optimizations

DirectToLDS

Modern AMD GPUs have instructions that allows you to load from global memory to LDS directly, so we can have something like:

`rock.threadwise_read_into`



`amdgpu.gather_to_lds`



`rocdl.load.to_lds`




`global_load_lds_dword`
`global_load_lds_dwordx3`
`global_load_lds_dwordx4`

The rock op lowers to amdgpu, rocdl and later to the DirectToLDS assembly load instruction

4. Optimizations

GEMM specific: SplitK




Let's imagine a GPU with 8 CUs and a matmul kernel with: $M=N=4$ and $K=8$


Let's assume tile sizes for M and N are 4. Therefore, we can only use 1 CU (out of 8) because M and N dimensions are small

4. Optimizations

GEMM specific: SplitK



Let's imagine a GPU with 8 CUs and a matmul kernel with: $M=N=4$ and $K=8$



With SplitK=8 we split the K dimension 8 times, thus allowing $8\times$ more parallelism, using the 8 CUs effectively

5. Evaluation

Software:

- rocMLIR
- hipBLASLt (Matmul)
- CKTile (Attention)

We used ROCm 7.2

We show the average over 3 independent runs

We evaluate 2 kernels: Matmul and Attention


Exhaustive tuning mode was used in rocMLIR

Hardware:

- MI300X (gfx942), MI350X (gfx950)
- RDNA3 (gfx1101)

5. Evaluation (Matmul)


RDNA 3



- 60 configs tested:
- hipBLASLt is faster 16% of the configs
 - rocMLIR is faster 76% of the configs

5. Evaluation (Matmul)

MI300X




60 configs tested:

- hipBLASLt is faster 13% of the configs
- rocMLIR is faster 86% of the configs

5. Evaluation (Matmul)

MI350X




60 configs tested:

- hipBLASLt is faster 20% of the configs
- rocMLIR is faster 75% of the configs

5. Evaluation (Attention)

MI300X



397 configs tested:

- CKTile is faster 56% of the configs
- rocMLIR is faster 30% of the configs

6. Conclusions and future work

Conclusions

- rocMLIR combines upstream MLIR dialects with downstream GPU-focused dialects:
- Our work upstream:
 - We have fixed several bugs in different parts of MLIR
 - We have added support for different amdgpu and rocdl ops
- Upstreaming candidates:
 - Transforms: The main benefit is exploited in the rock dialect, so upstreaming does not seem trivial
 - Fusions: Too much dependent on our internal IR design, upstreaming not appropriate
- rocMLIR is open source, check it out: <https://github.com/ROCm/rocMLIR>

6. Conclusions and future work

Conclusions

- rocMLIR combines upstream MLIR dialects with downstream GPU-focused dialects
- Our work upstream:
 - We have fixed several bugs in different parts of MLIR
 - We have added support for different amdgpu and rocdl ops
- Upstreaming candidates:
 - Transforms: The main benefit is exploited in the rock dialect, so upstreaming does not seem trivial
 - Fusions: Too much dependent on our internal IR design, upstreaming not appropriate
- rocMLIR is open source, check it out: <https://github.com/ROCm/rocMLIR>


We are hiring!

If you are interested, make sure to visit AMD's stand during the conference for more information

6. Conclusions and future work

Future work

Integrating linalg as middle layer
between graph IRs and rock



Using Triton as the backend of rocMLIR



